

# 畳み込みニューラルネット トワークの基礎と応用

# Agenda

19:00 ~ 19:05 opening

19:05 ~ 19:20 ニューラルネットワークの基礎

19:20 ~ 19:30 **walkthrough** Kerasを用いたコーディングサンプル

19:30 ~ 19:45 置み込みニューラルネットワークの基礎

19:45 ~ 20:00 **walkthrough** CNNのフィルタの理解

20:00 ~ 20:10 休憩

20:10 ~ 20:25 CNNアーキテクチャの紹介

20:25 ~ 20:40 Skymind Intelligent Layer (SKIL) の解説

20:40 ~ 20:55 **walkthrough** SKILを用いたモデルのデプロイ

20:55 ~ 21:00 closing



# Activity Recognition



```
if(speed<4){  
    status=WALKING;  
}
```

# Activity Recognition



```
if(speed<4){  
    status=WALKING;  
}
```

```
if(speed<4){  
    status=WALKING;  
} else {  
    status=RUNNING;  
}
```

# Activity Recognition



```
if(speed<4){  
    status=WALKING;  
}
```



```
if(speed<4){  
    status=WALKING;  
} else {  
    status=RUNNING;  
}
```



```
if(speed<4){  
    status=WALKING;  
} else if(speed<12){  
    status=RUNNING;  
} else {  
    status=BIKING;  
}
```

# Activity Recognition



```
if(speed<4){  
    status=WALKING;  
}
```



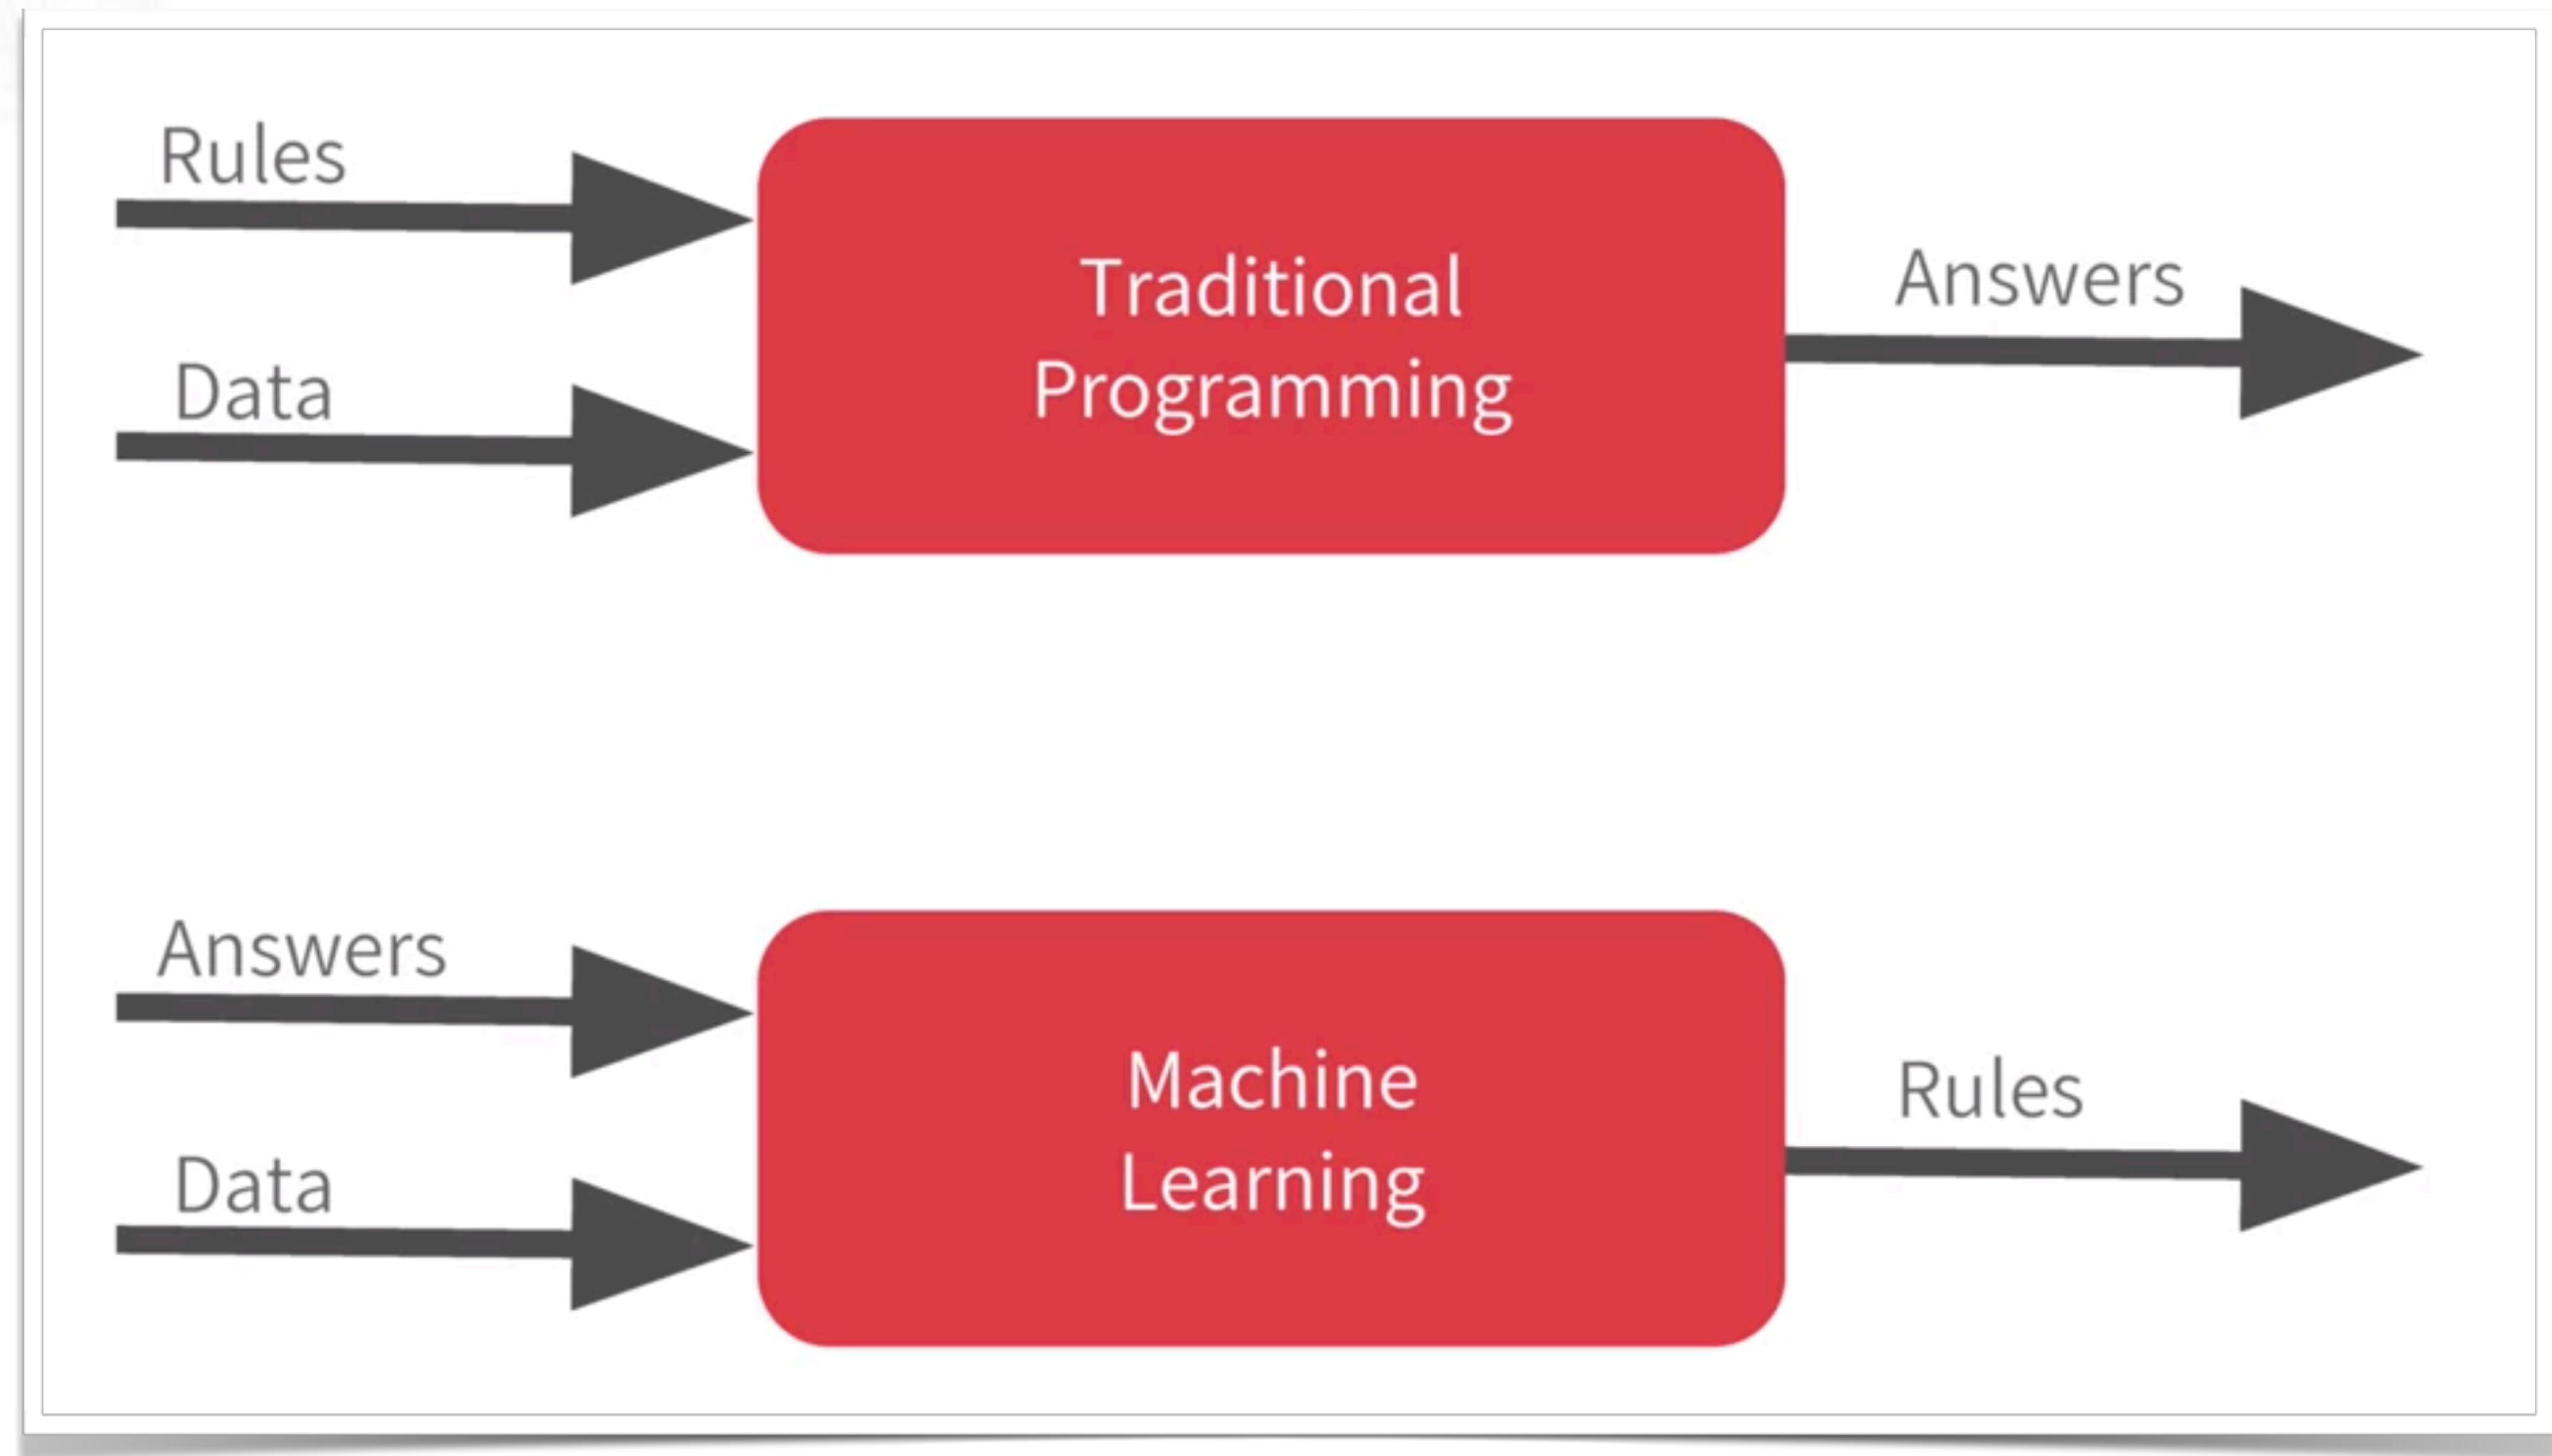
```
if(speed<4){  
    status=WALKING;  
} else {  
    status=RUNNING;  
}
```



```
if(speed<4){  
    status=WALKING;  
} else if(speed<12){  
    status=RUNNING;  
} else {  
    status=BIKING;  
}
```



// Oh crap



# Activity Recognition



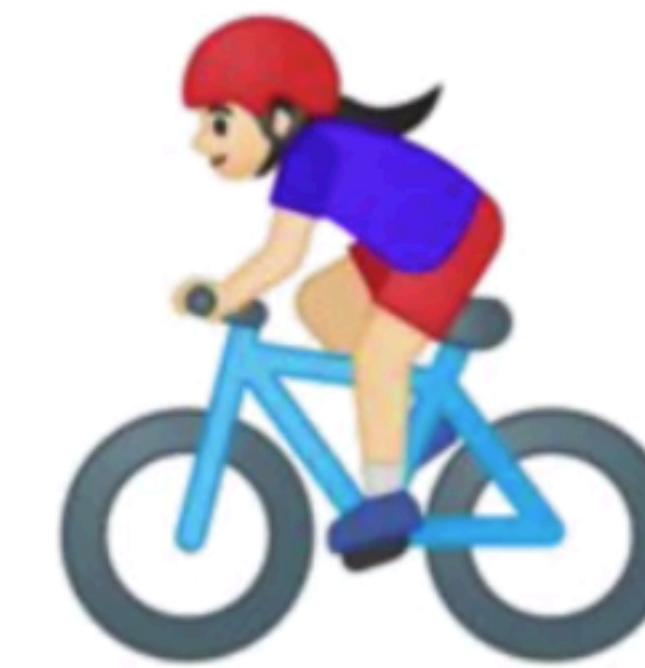
0101001010100101010  
1001010101001011101  
0100101010010101001  
0101001010100101010

Label = WALKING



1010100101001010101  
0101010010010010001  
0010011111010101111  
1010100100111101011

Label = RUNNING



1001010011111010101  
1101010111010101110  
1010101111010101011  
1111110001111010101

Label = BIKING



111111111010011101  
0011111010111110101  
0101110101010101110  
1010101010100111110

Label = GOLFING  
(Sort of)

# 機械学習の"Hello World"

$$\begin{aligned}x &= -2, -1, 0, 1, 2, 3, 4 \\y &= -3, -1, 1, 3, 5, 7, 9\end{aligned}$$

$$y = f(x)$$

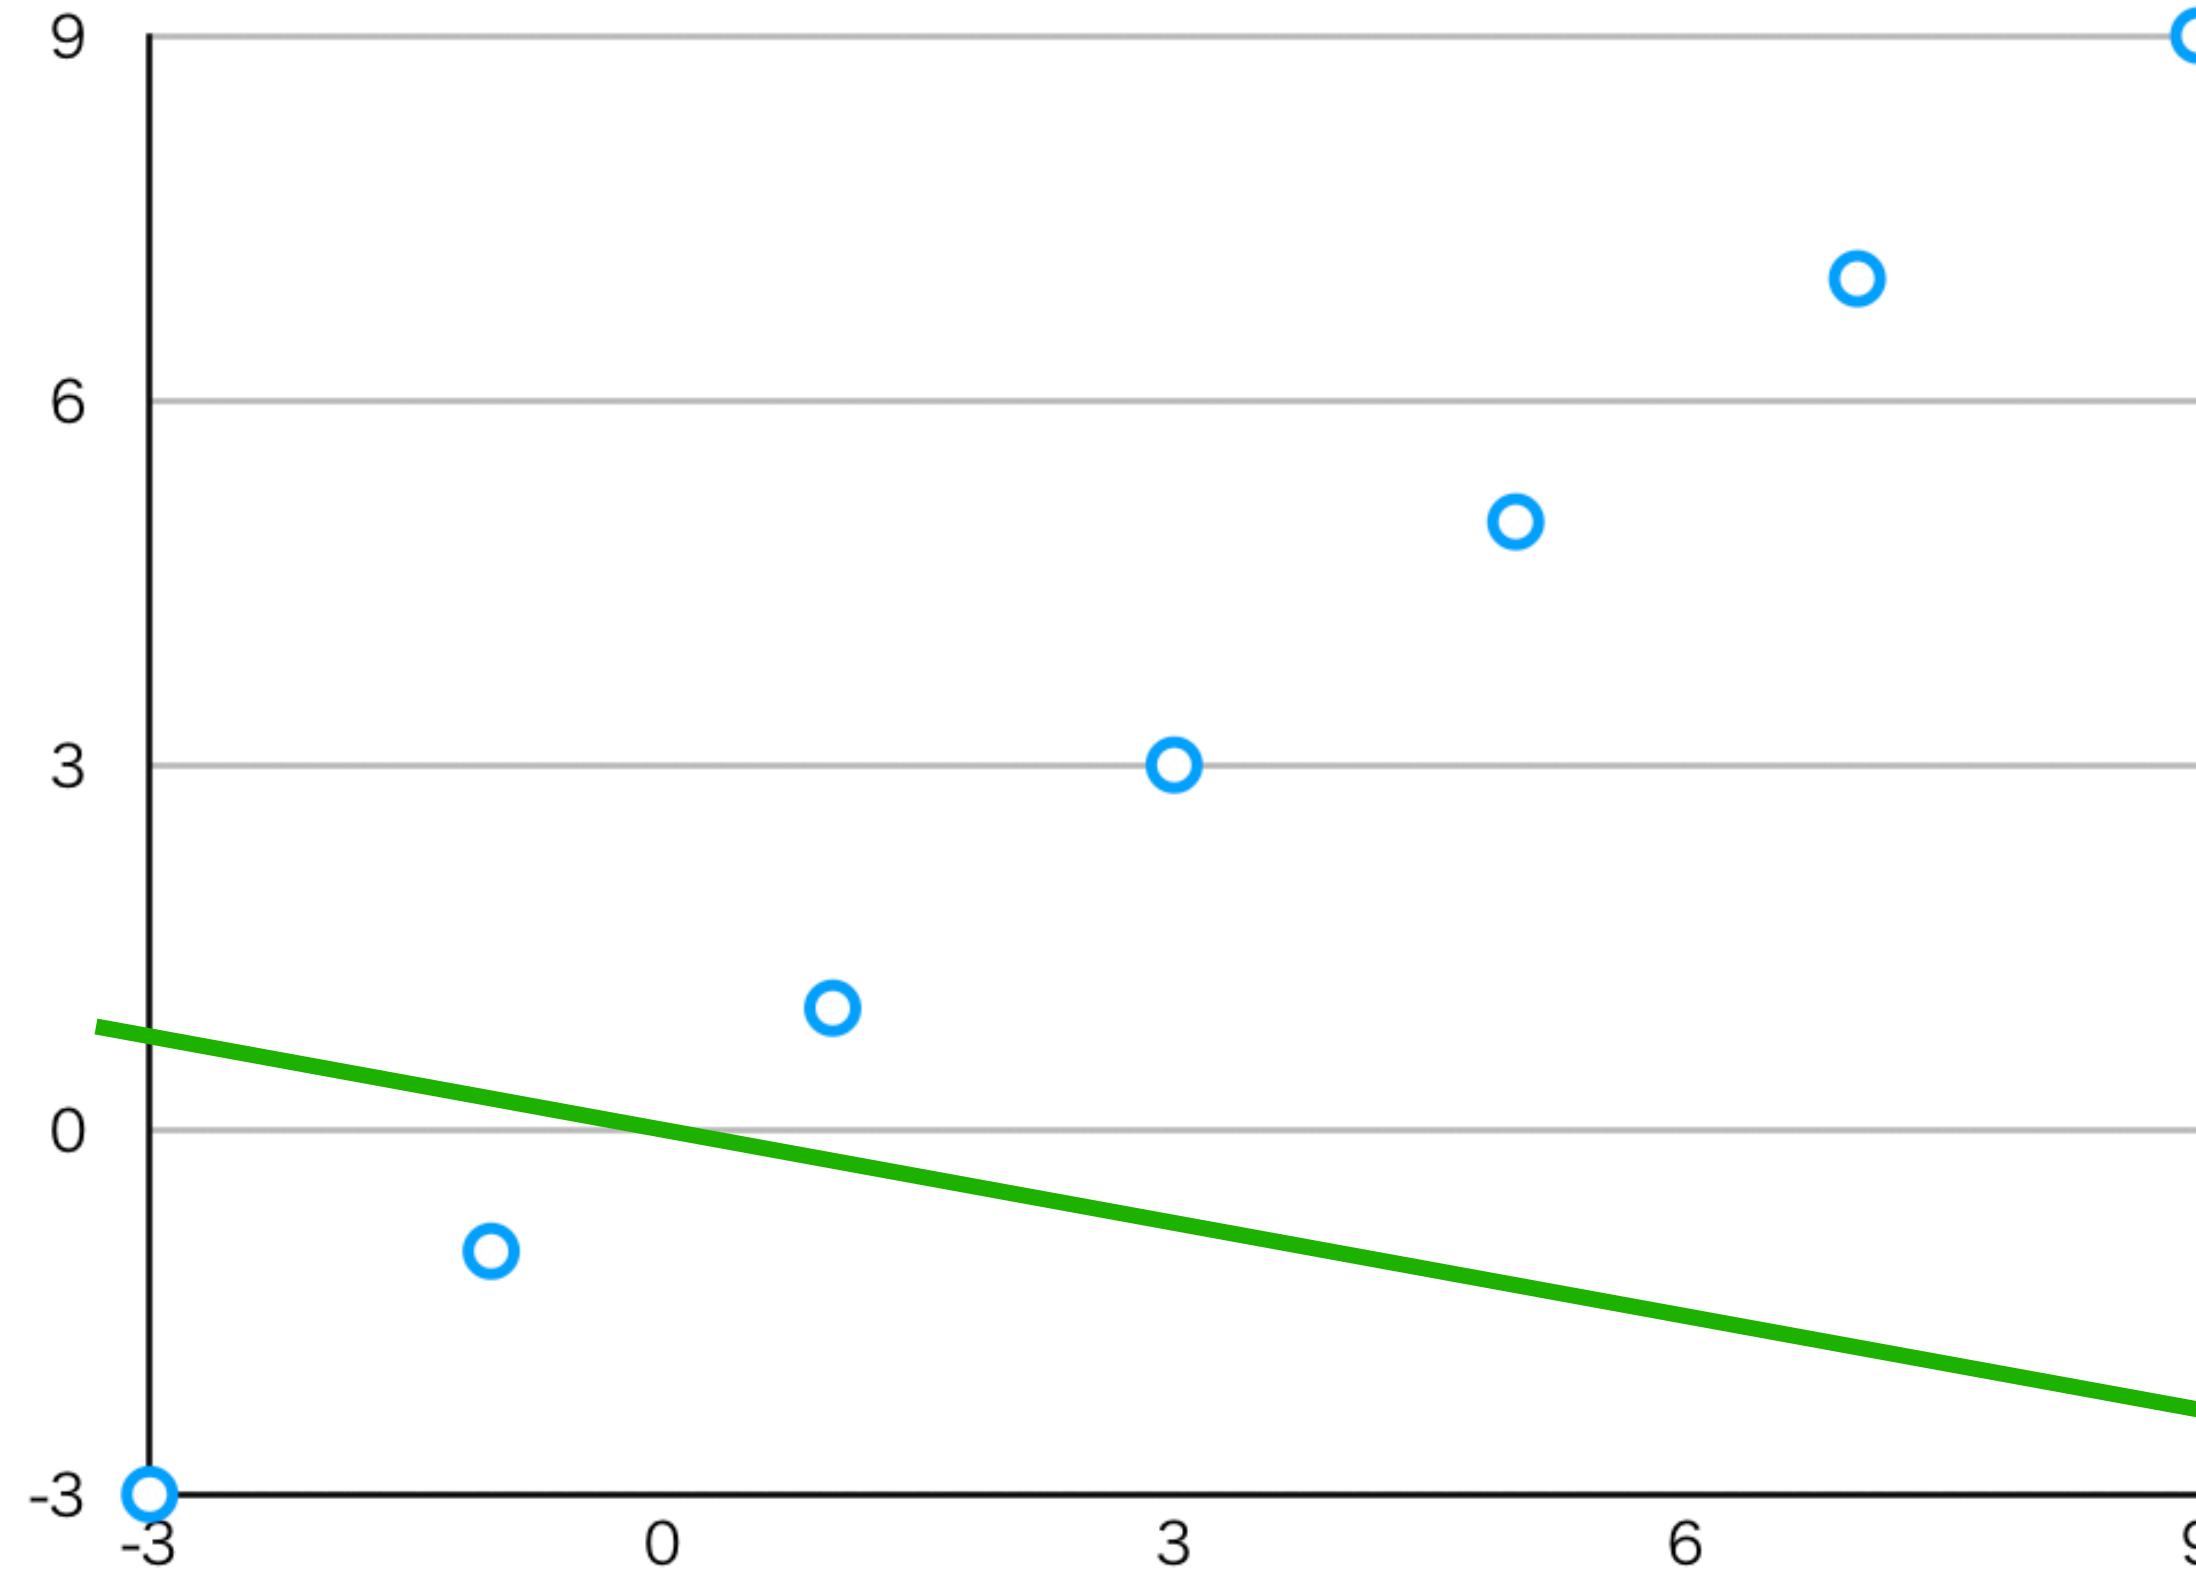
# 機械学習の"Hello World"

$$\begin{aligned}x &= -2, -1, 0, 1, 2, 3, 4 \\y &= -3, -1, 1, 3, 5, 7, 9\end{aligned}$$

$$y = f(x) = 2x + 1$$

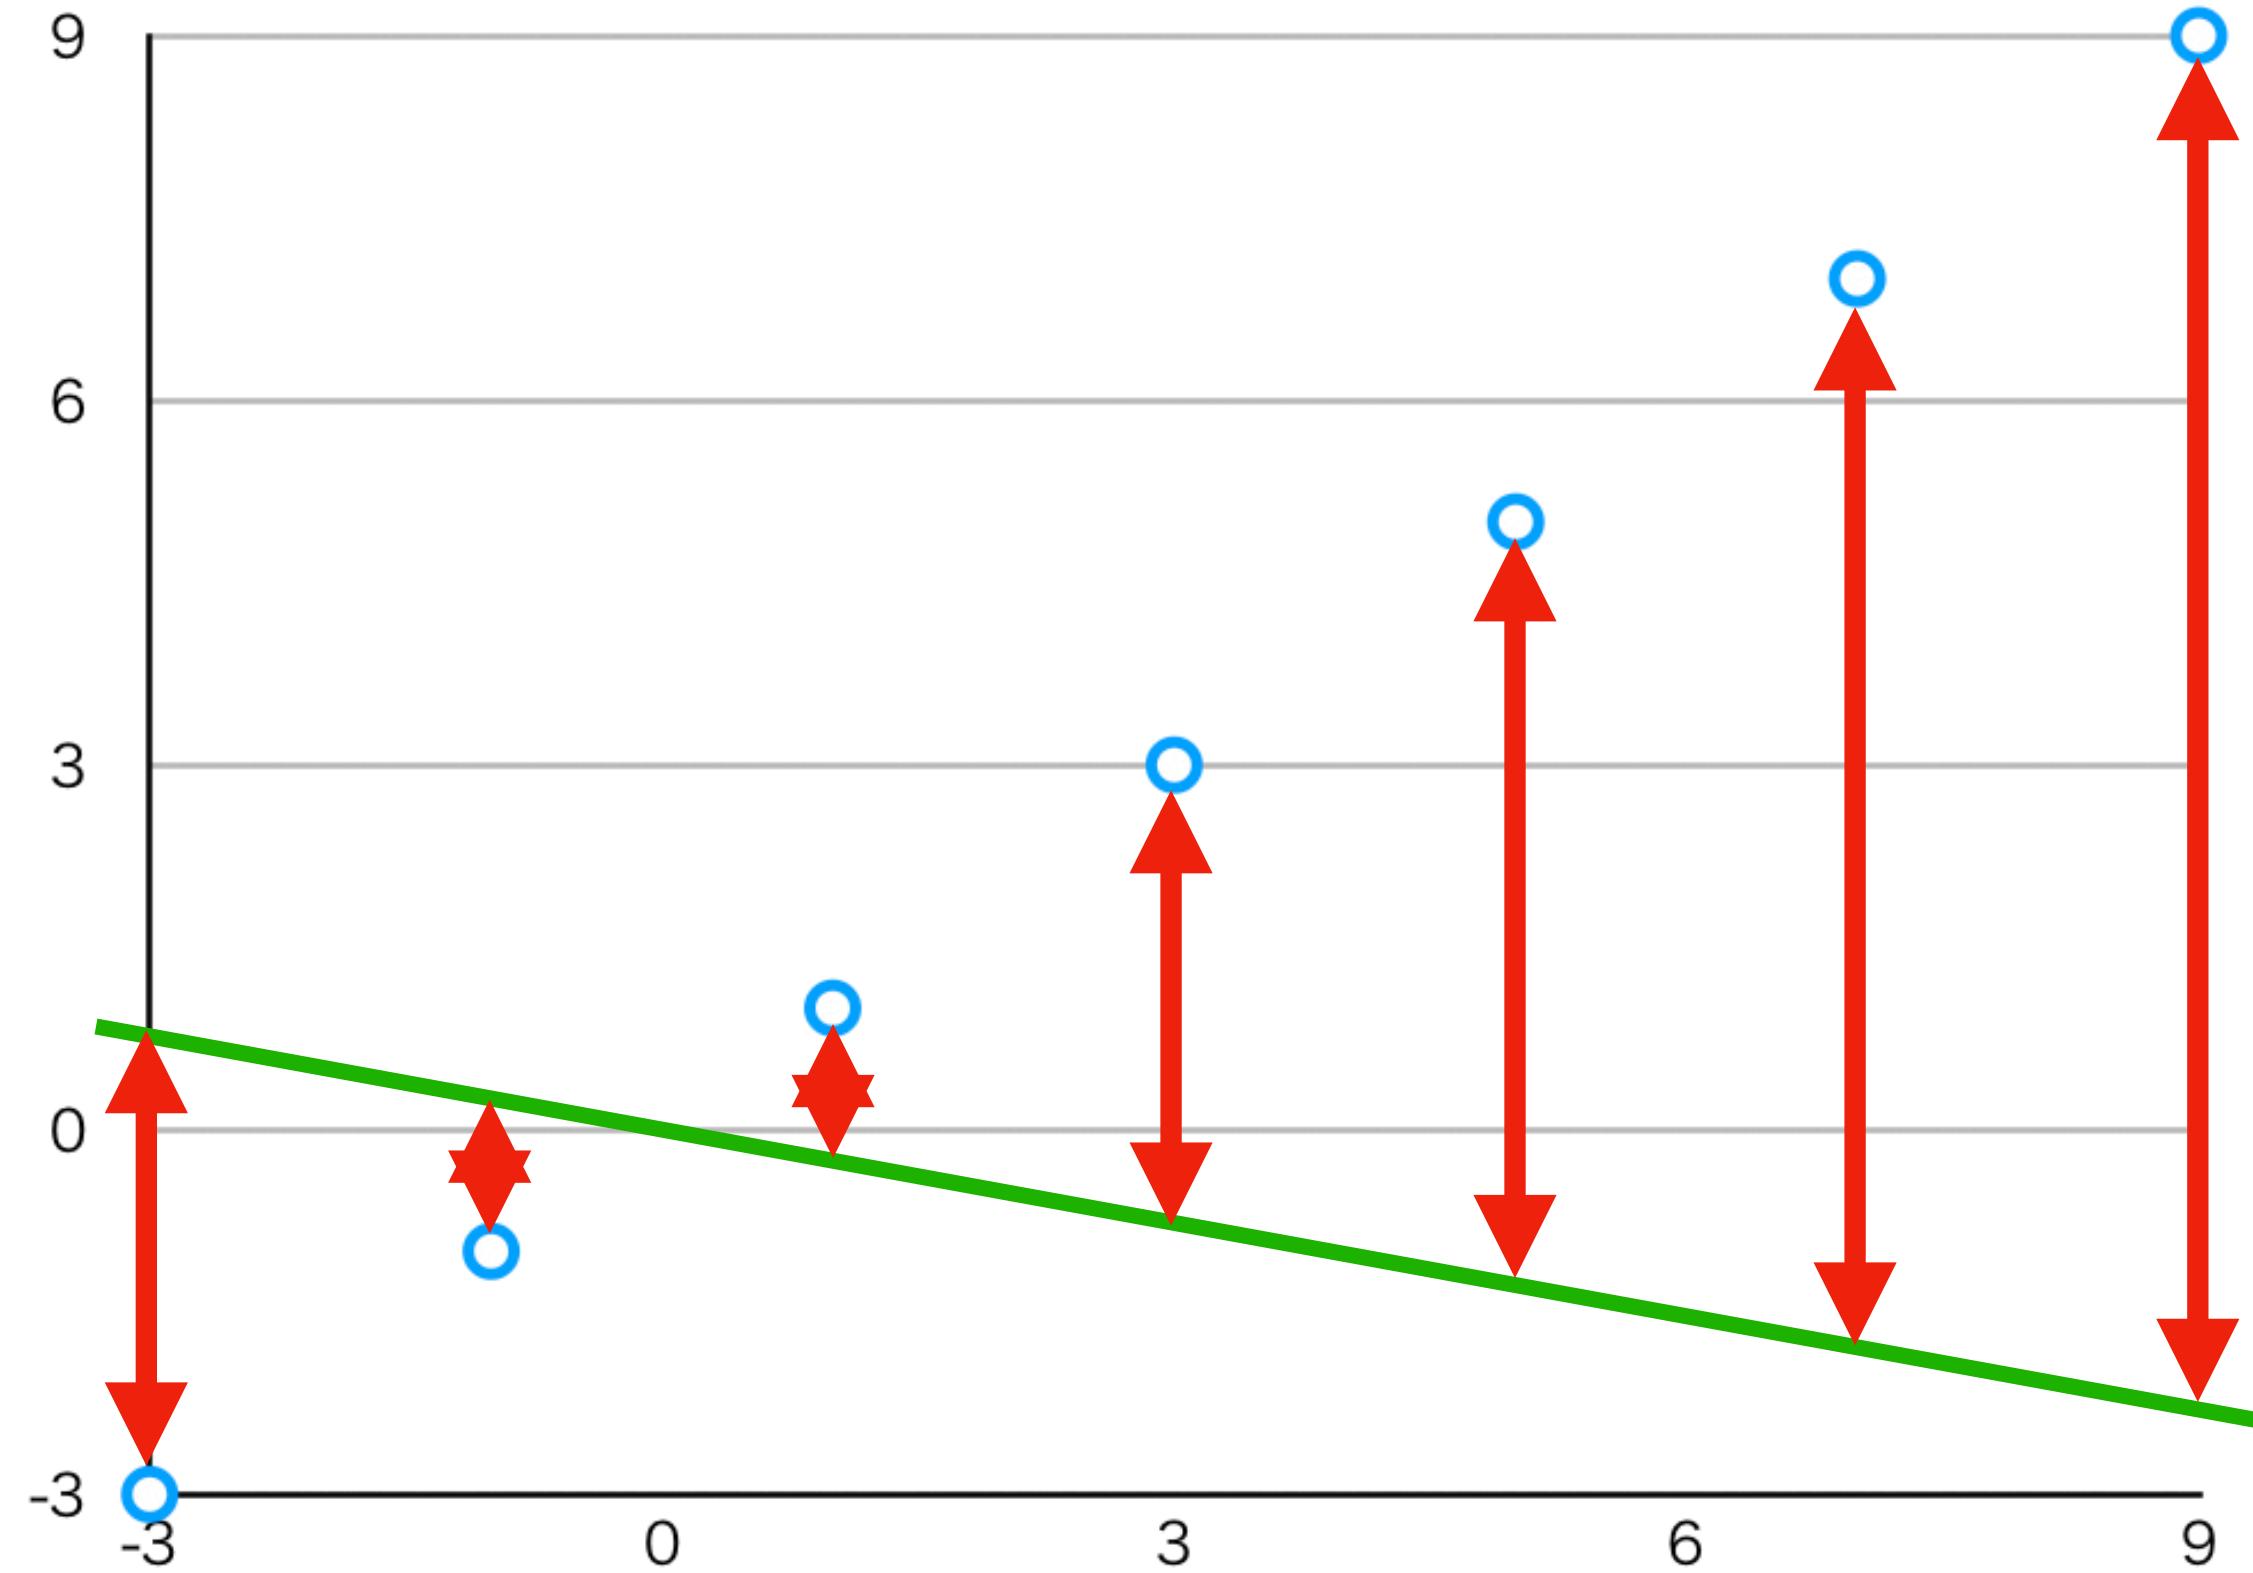
# 機械学習のアプローチ

- 適当にモデルを初期化 ( $y_{\_}=ax+b$ )

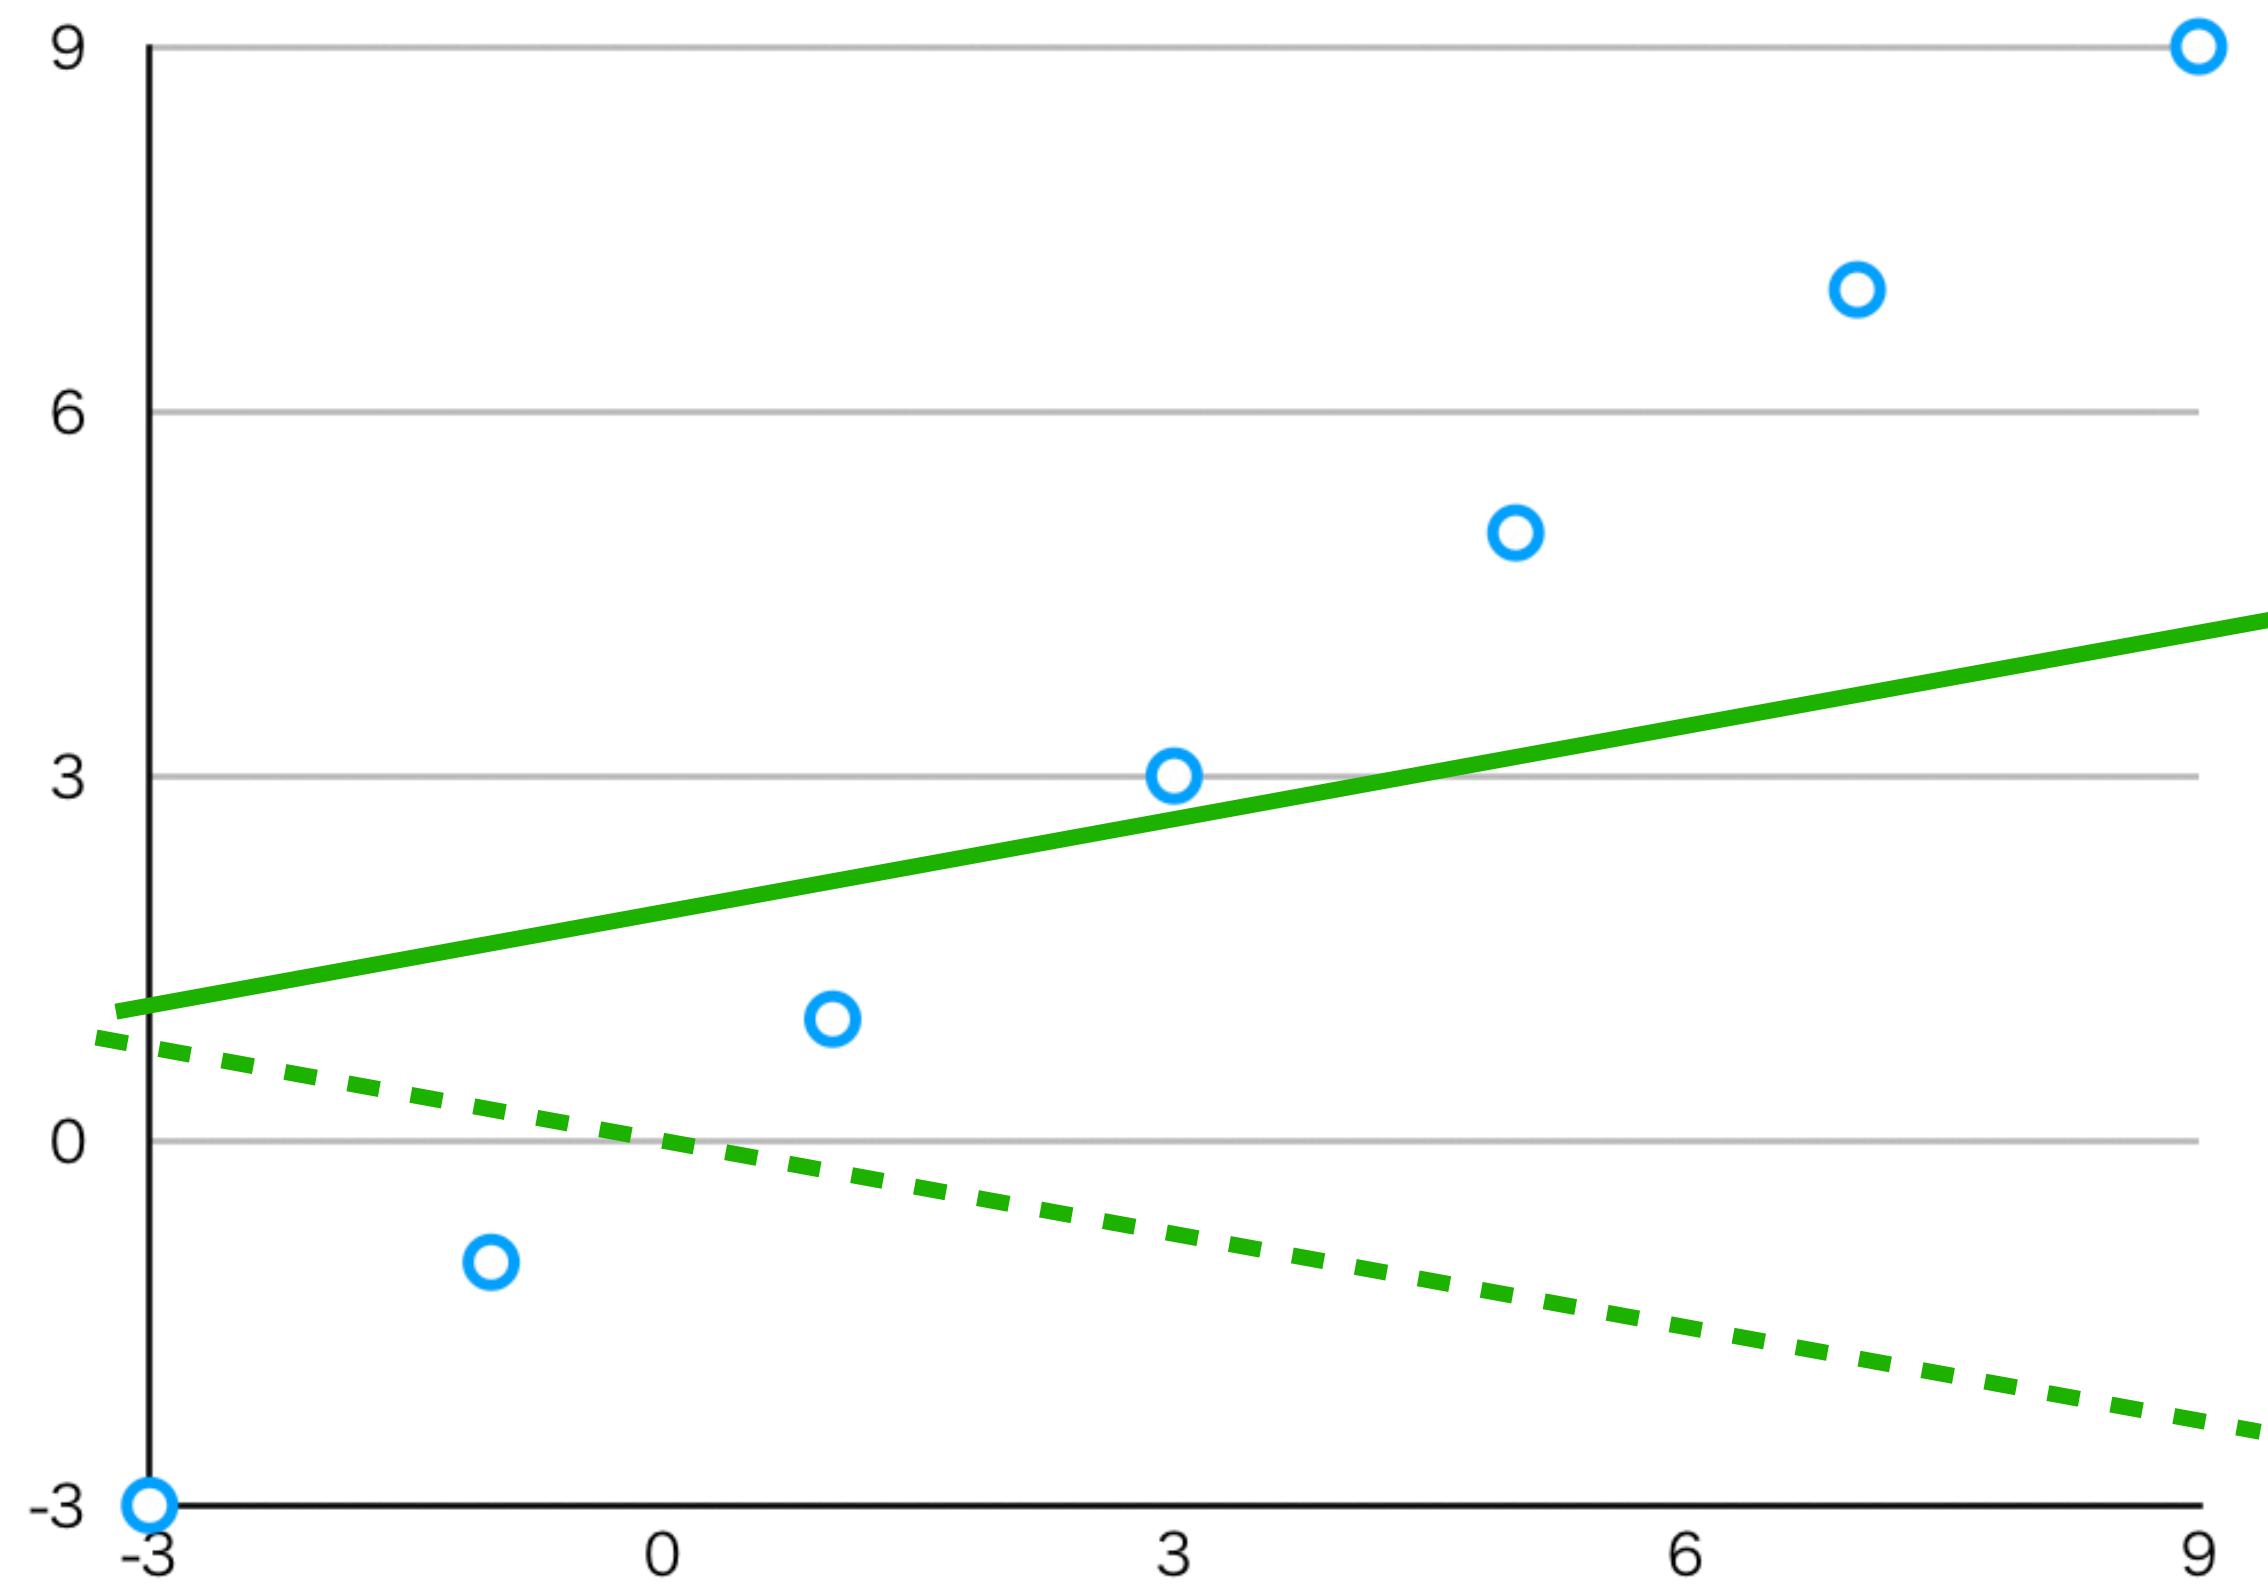


# 機械学習のアプローチ

- 適当にモデルを初期化 ( $y_{\_} = ax + b$ )
- 誤差（損失）を計算 ( $L = 1/N \sum (y - y_{\_})^2$ )

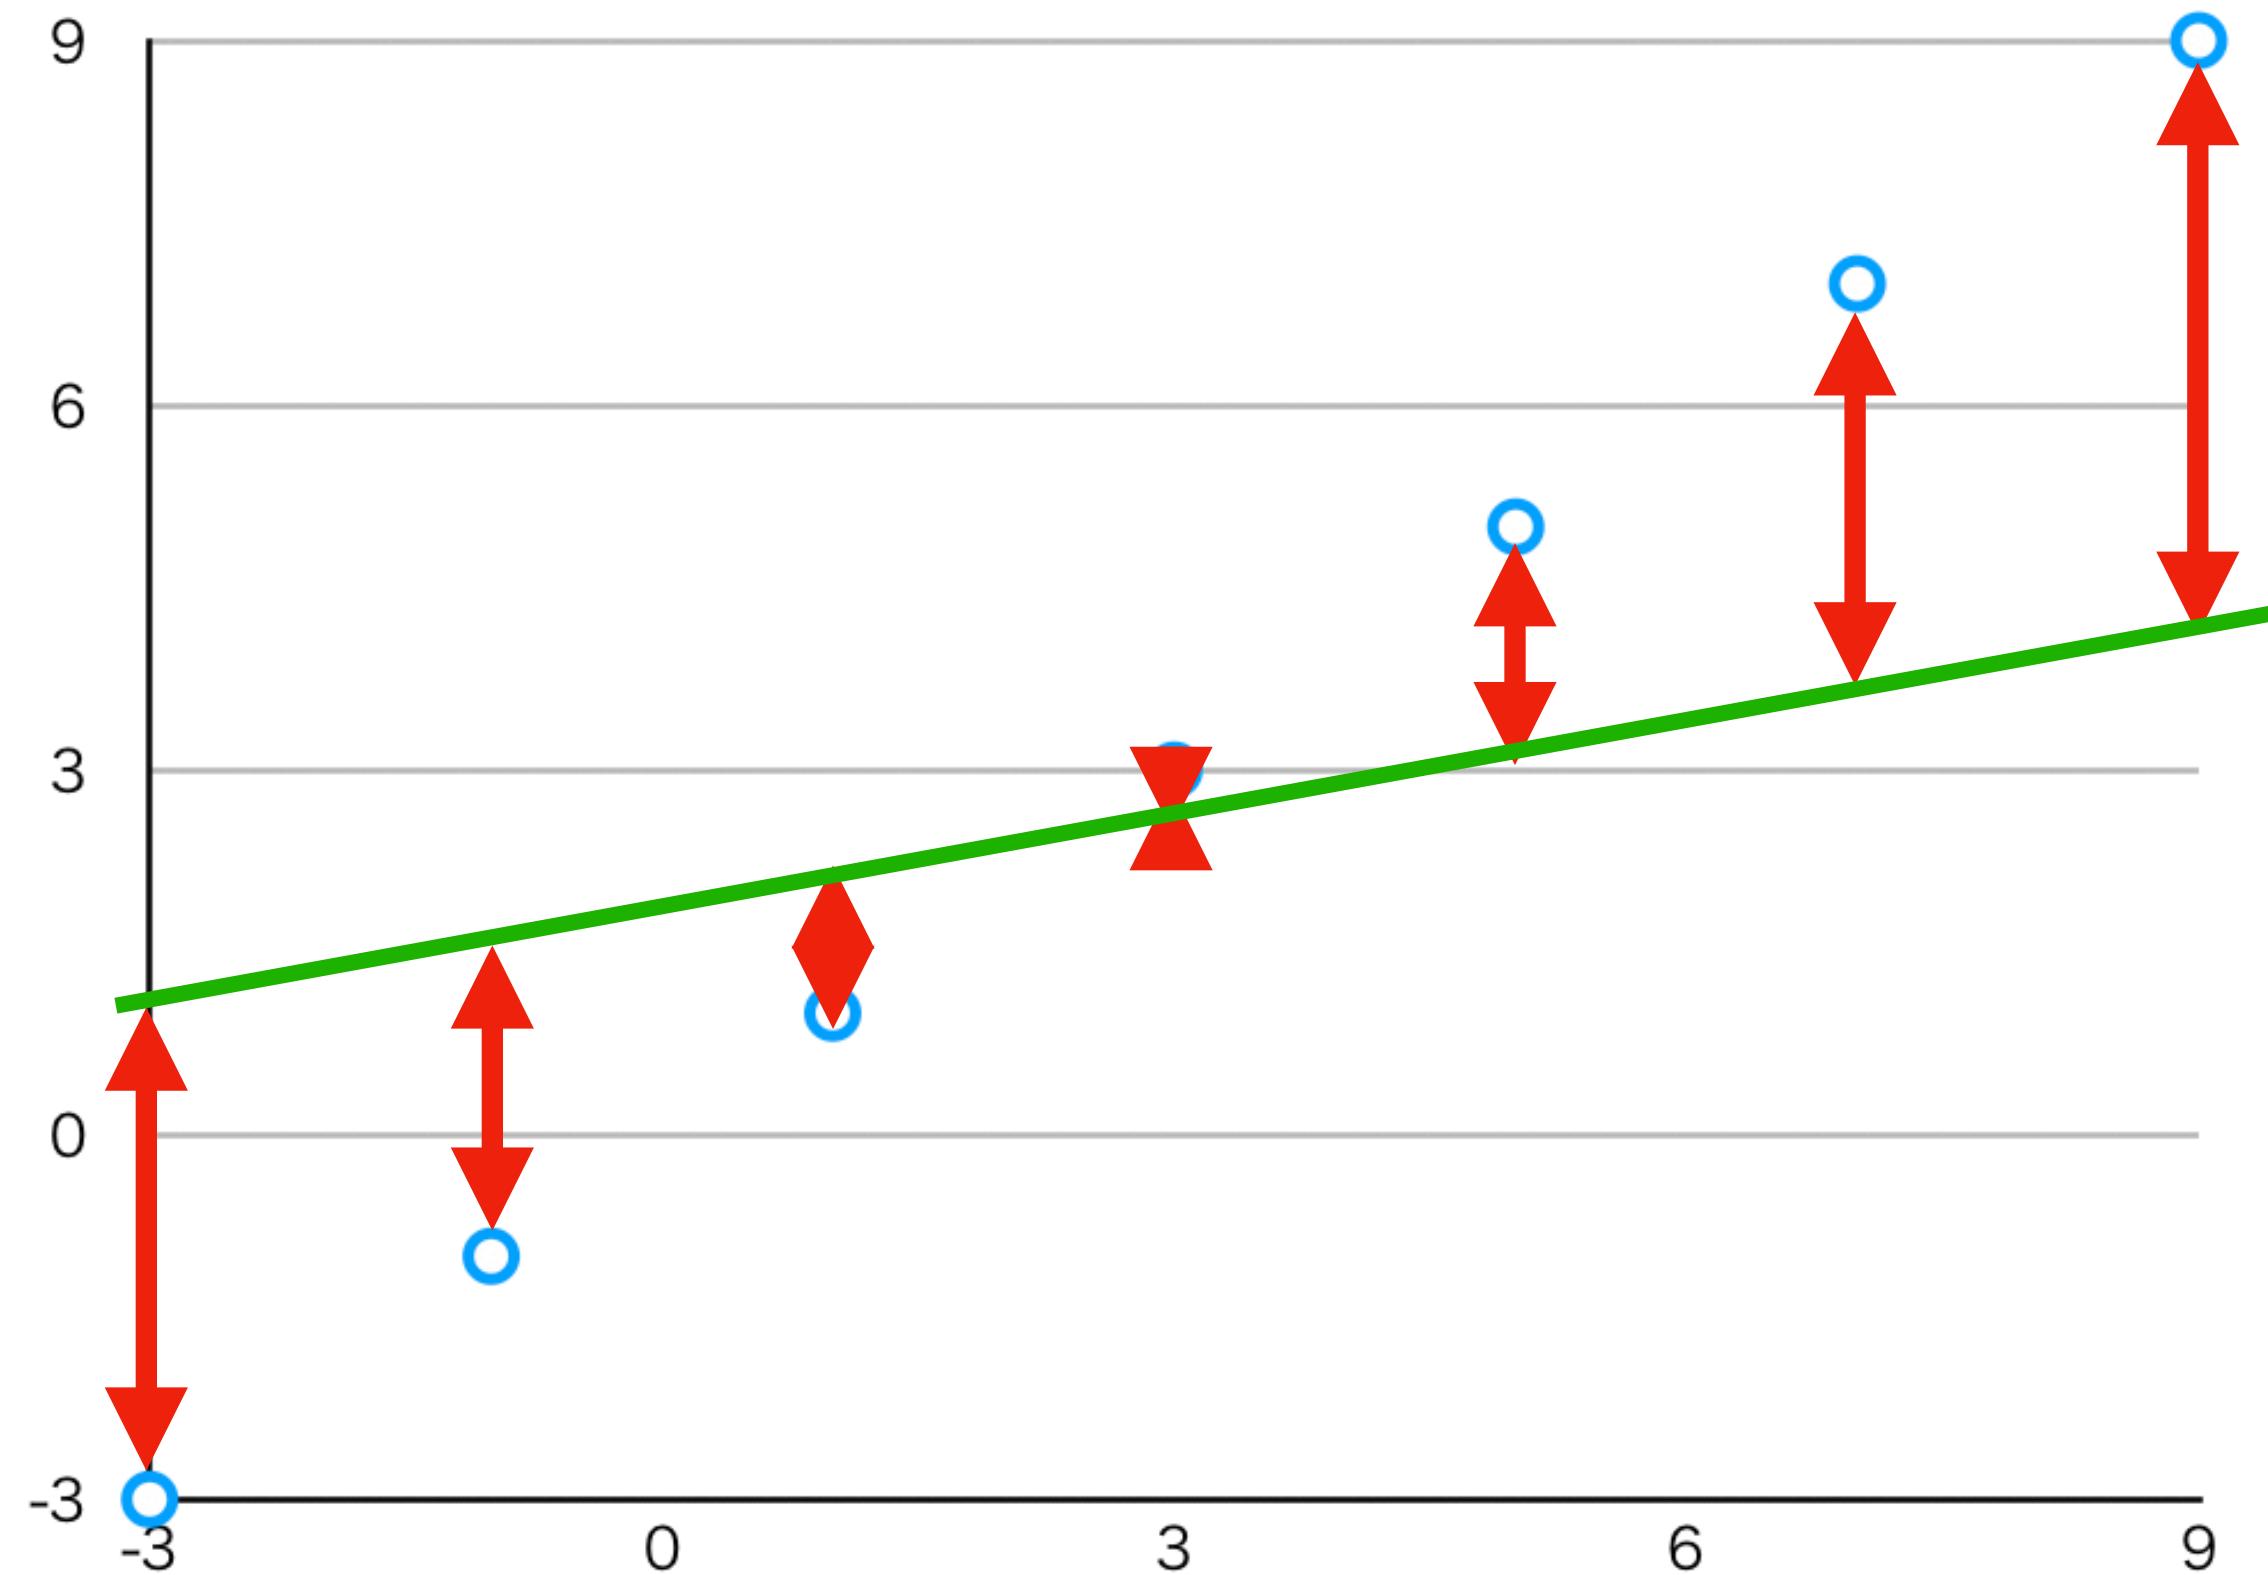


# 機械学習のアプローチ



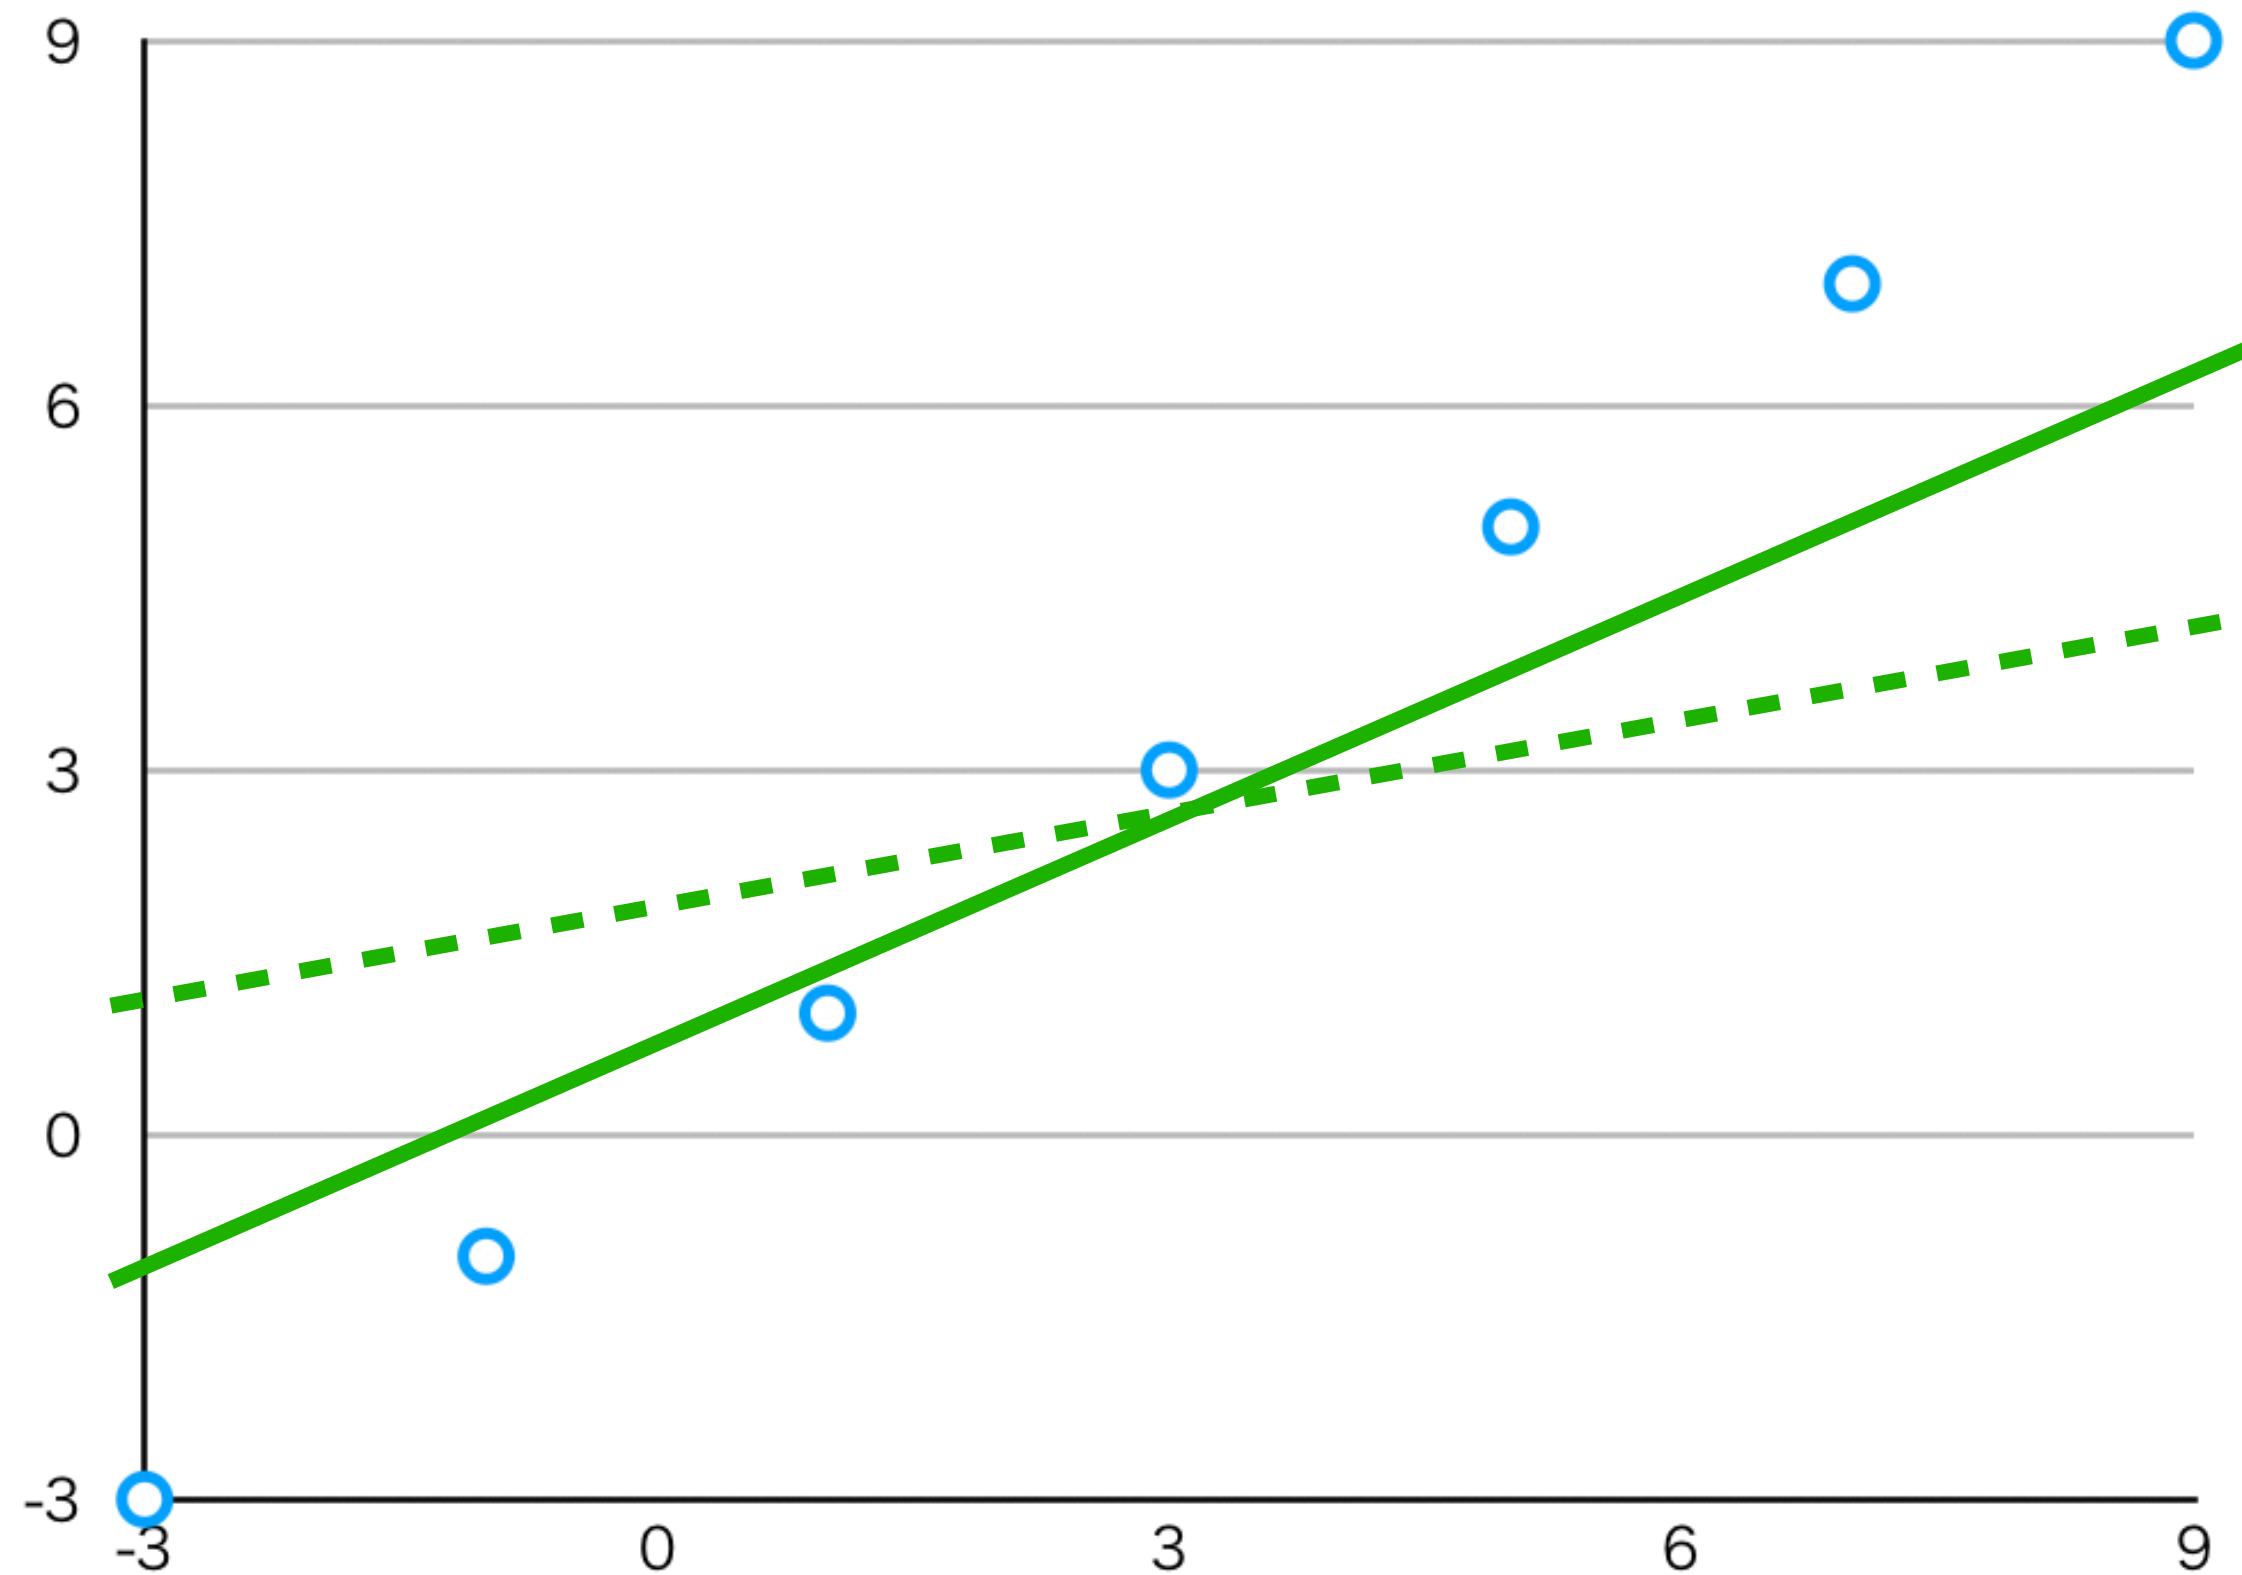
- 適当にモデルを初期化 ( $y_- = ax + b$ )
- 誤差（損失）を計算 ( $L = \frac{1}{N} \sum (y - y_-)^2$ )
- 誤差が小さくなるようにパラメータ ( $a, b$ ) を少し更新 ( $a \leftarrow a - \eta \frac{\partial L}{\partial a}$ )

# 機械学習のアプローチ



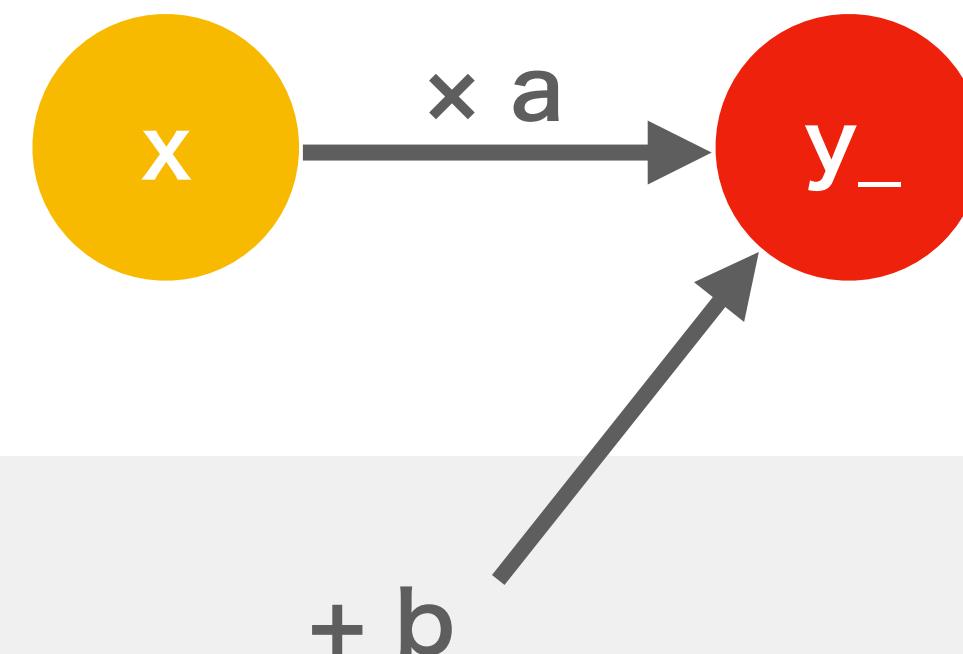
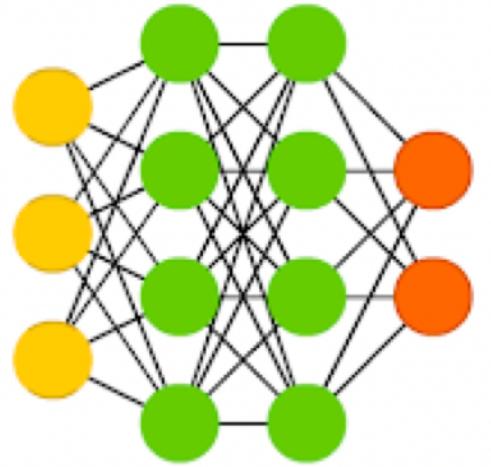
- 適当にモデルを初期化 ( $y_ = ax + b$ )
- 誤差（損失）を計算 ( $L = 1/N \sum (y - y_)^2$ )
- 誤差が小さくなるようにパラメータ ( $a, b$ ) を少し更新 ( $a \leftarrow a - \eta \partial L / \partial a$ )
- 誤差（損失）を計算 ( $L = 1/N \sum (y - y_)^2$ )

# 機械学習のアプローチ



- 適当にモデルを初期化 ( $y_{\_} = ax + b$ )
- 誤差（損失）を計算 ( $L = 1/N \sum (y - y_{\_})^2$ )
- 誤差が小さくなるようにパラメータ ( $a, b$ ) を少し更新 ( $a \leftarrow a - \eta \frac{\partial L}{\partial a}$ )
- 誤差（損失）を計算 ( $L = 1/N \sum (y - y_{\_})^2$ )
- 誤差が小さくなるようにパラメータ ( $a, b$ ) を少し更新 ( $a \leftarrow a - \eta \frac{\partial L}{\partial a}$ )
- ...

```
1 import tensorflow as tf
2
3
4 model = tf.keras.Sequential([tf.keras.layers.Dense(units=1, input_shape=[1])])
5 model.compile(optimizer='sgd', loss='mean_squared_error')
6
7 xs = np.array([-1.0, 0.0, 1.0, 2.0, 3.0, 4.0], dtype=float)
8 ys = np.array([-3.0, -1.0, 1.0, 3.0, 5.0, 7.0], dtype=float)
9
10 model.fit(xs, ys, epochs=500)
11
12 >>>print(model.predict([10.0]))
13 [[18.977331]]
```



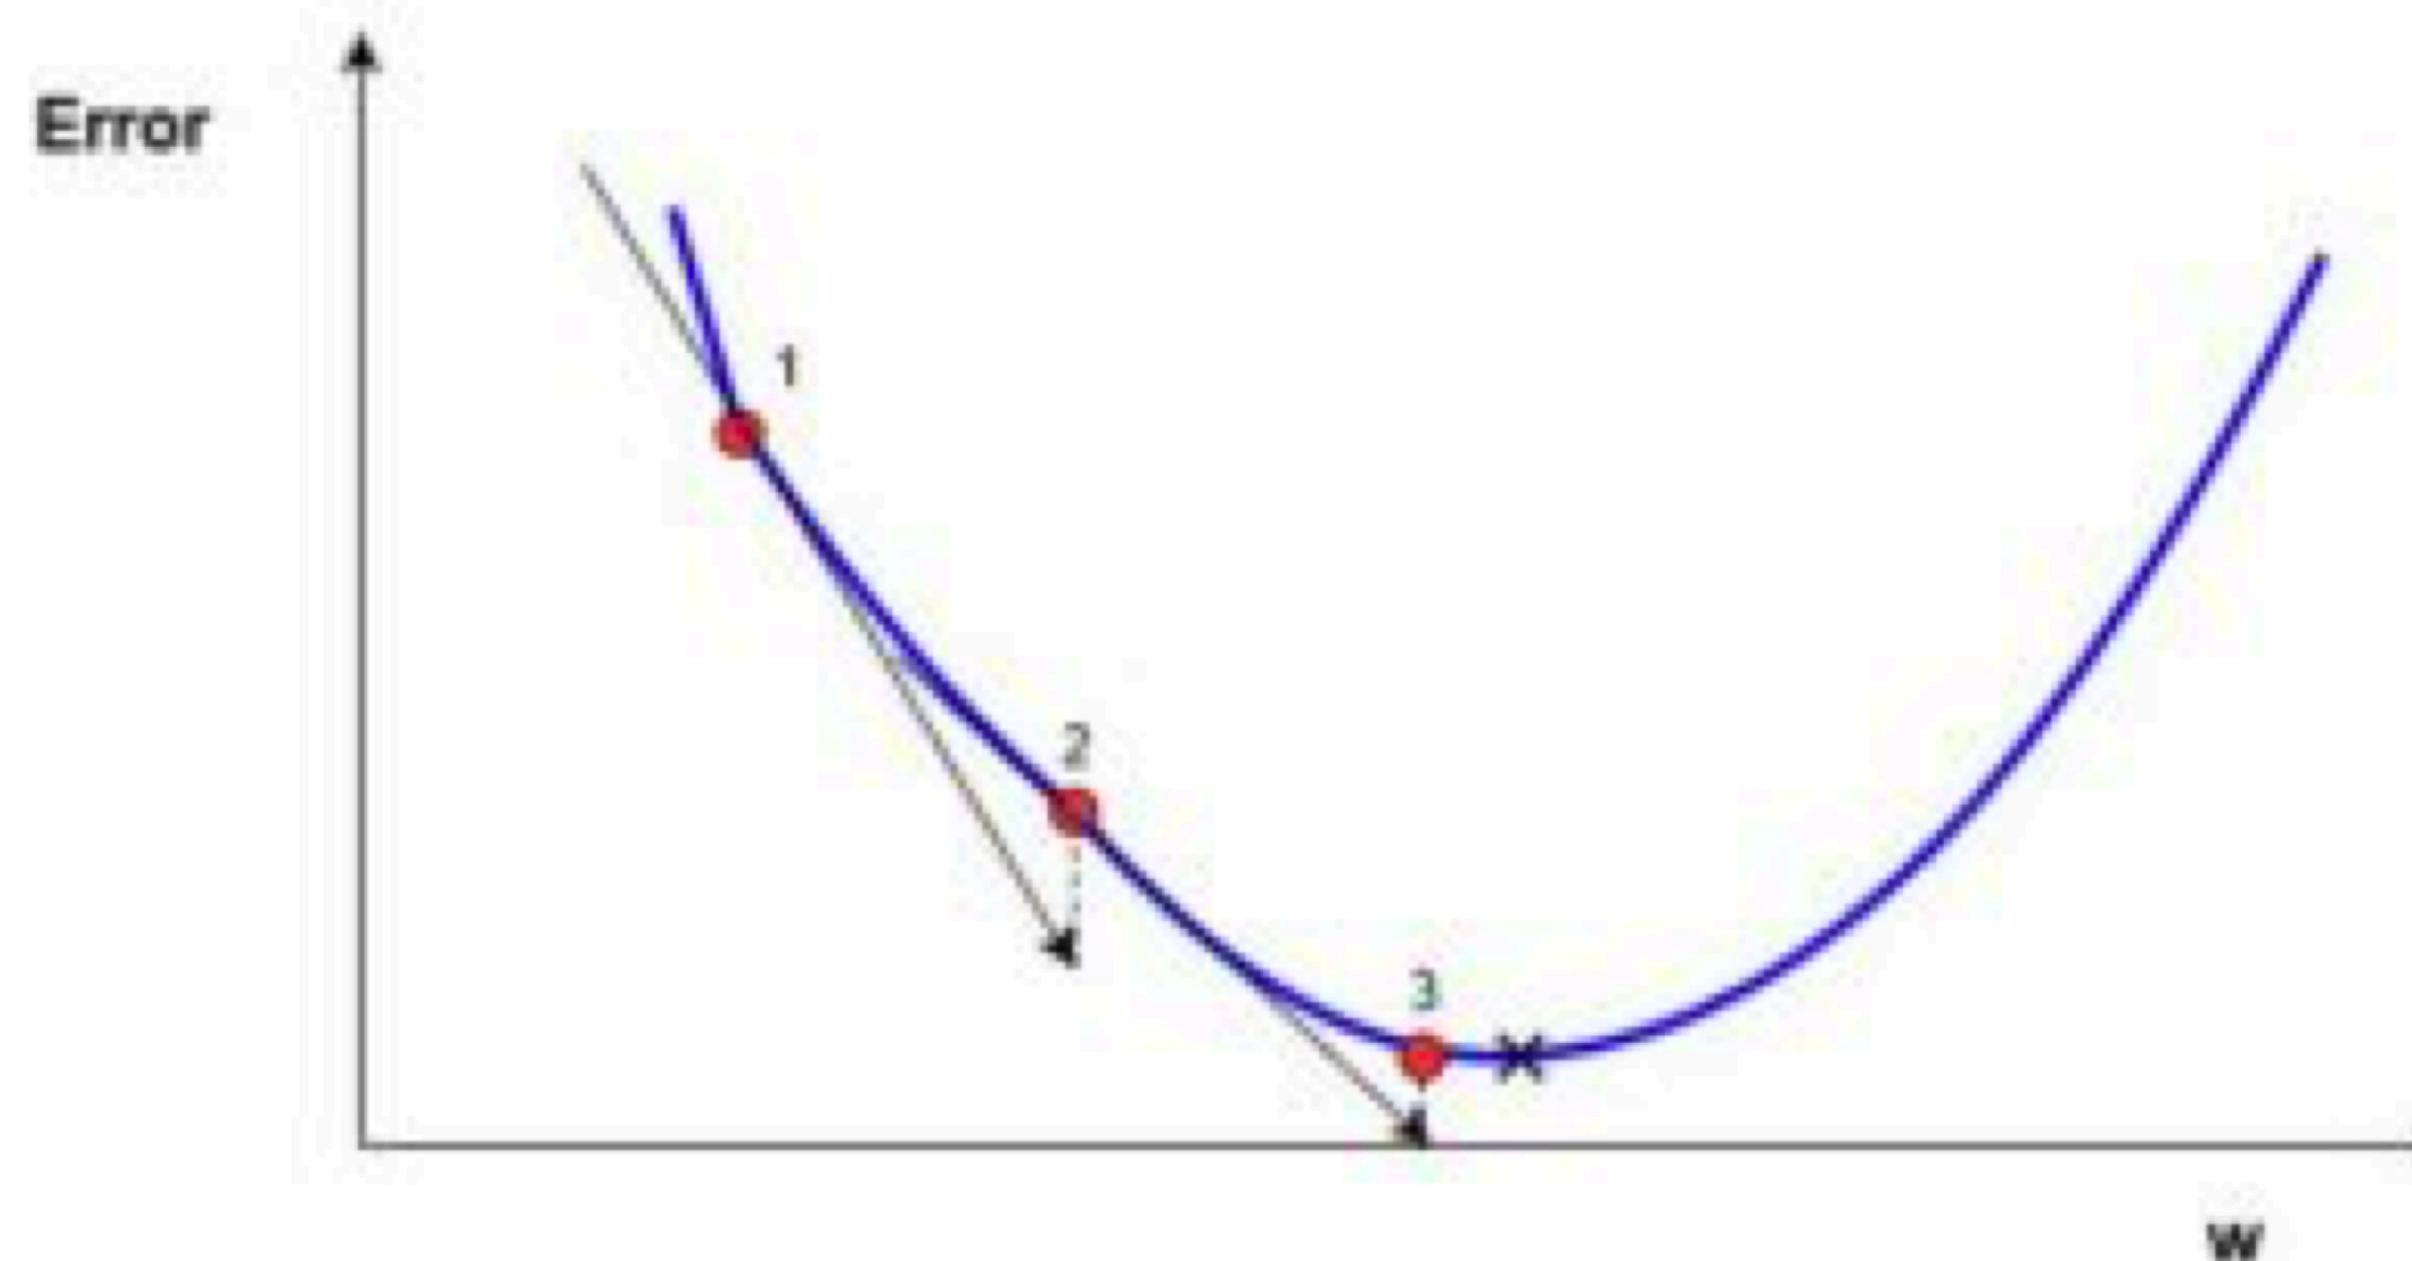
```
1 import tensorflow as tf
2
3
4 model = tf.keras.Sequential([tf.keras.layers.Dense(units=1, input_shape=[1])])
5 model.compile(optimizer='sgd', loss='mean_squared_error')
6
7 xs = np.array([-1.0, 0.0, 1.0, 2.0, 3.0, 4.0], dtype=float)
8 ys = np.array([-3.0, -1.0, 1.0, 3.0, 5.0, 7.0], dtype=float)
9
10 model.fit(xs, ys, epochs=500)
11
12 >>>print(model.predict([10.0]))
13 [[18.977331]]
```

```
1 import tensorflow as tf
2
3
4 model = tf.keras.Sequential([tf.keras.layers.Dense(units=1, input_shape=[1])])
5 model.compile(optimizer='sgd', loss='mean_squared_error')
6
7 xs = np.array([-1.0, 0.0, 1.0, 2.0, 3.0, 4.0], dtype=float)
8 ys = np.array([-3.0, -1.0, 1.0, 3.0, 5.0, 7.0], dtype=float)
9
10 model.fit(xs, ys, epochs=500)
11
12 >>>print(model.predict([10.0]))
13 [[18.977331]]
```

確率的勾配降下法 SGD, stochastic gradient descent.

デファクトスタンダード

(改良版 : momentum, Nesterovの加速法, RMSProp, Adam, ...)



$$w \leftarrow w - \eta \frac{\partial L}{\partial w}$$

# 勾配降下法



$$w \leftarrow w - \eta \frac{\partial L}{\partial w}$$

「目隠しで足元の勾配情報のみを使って山の頂上を目指すようなもの」

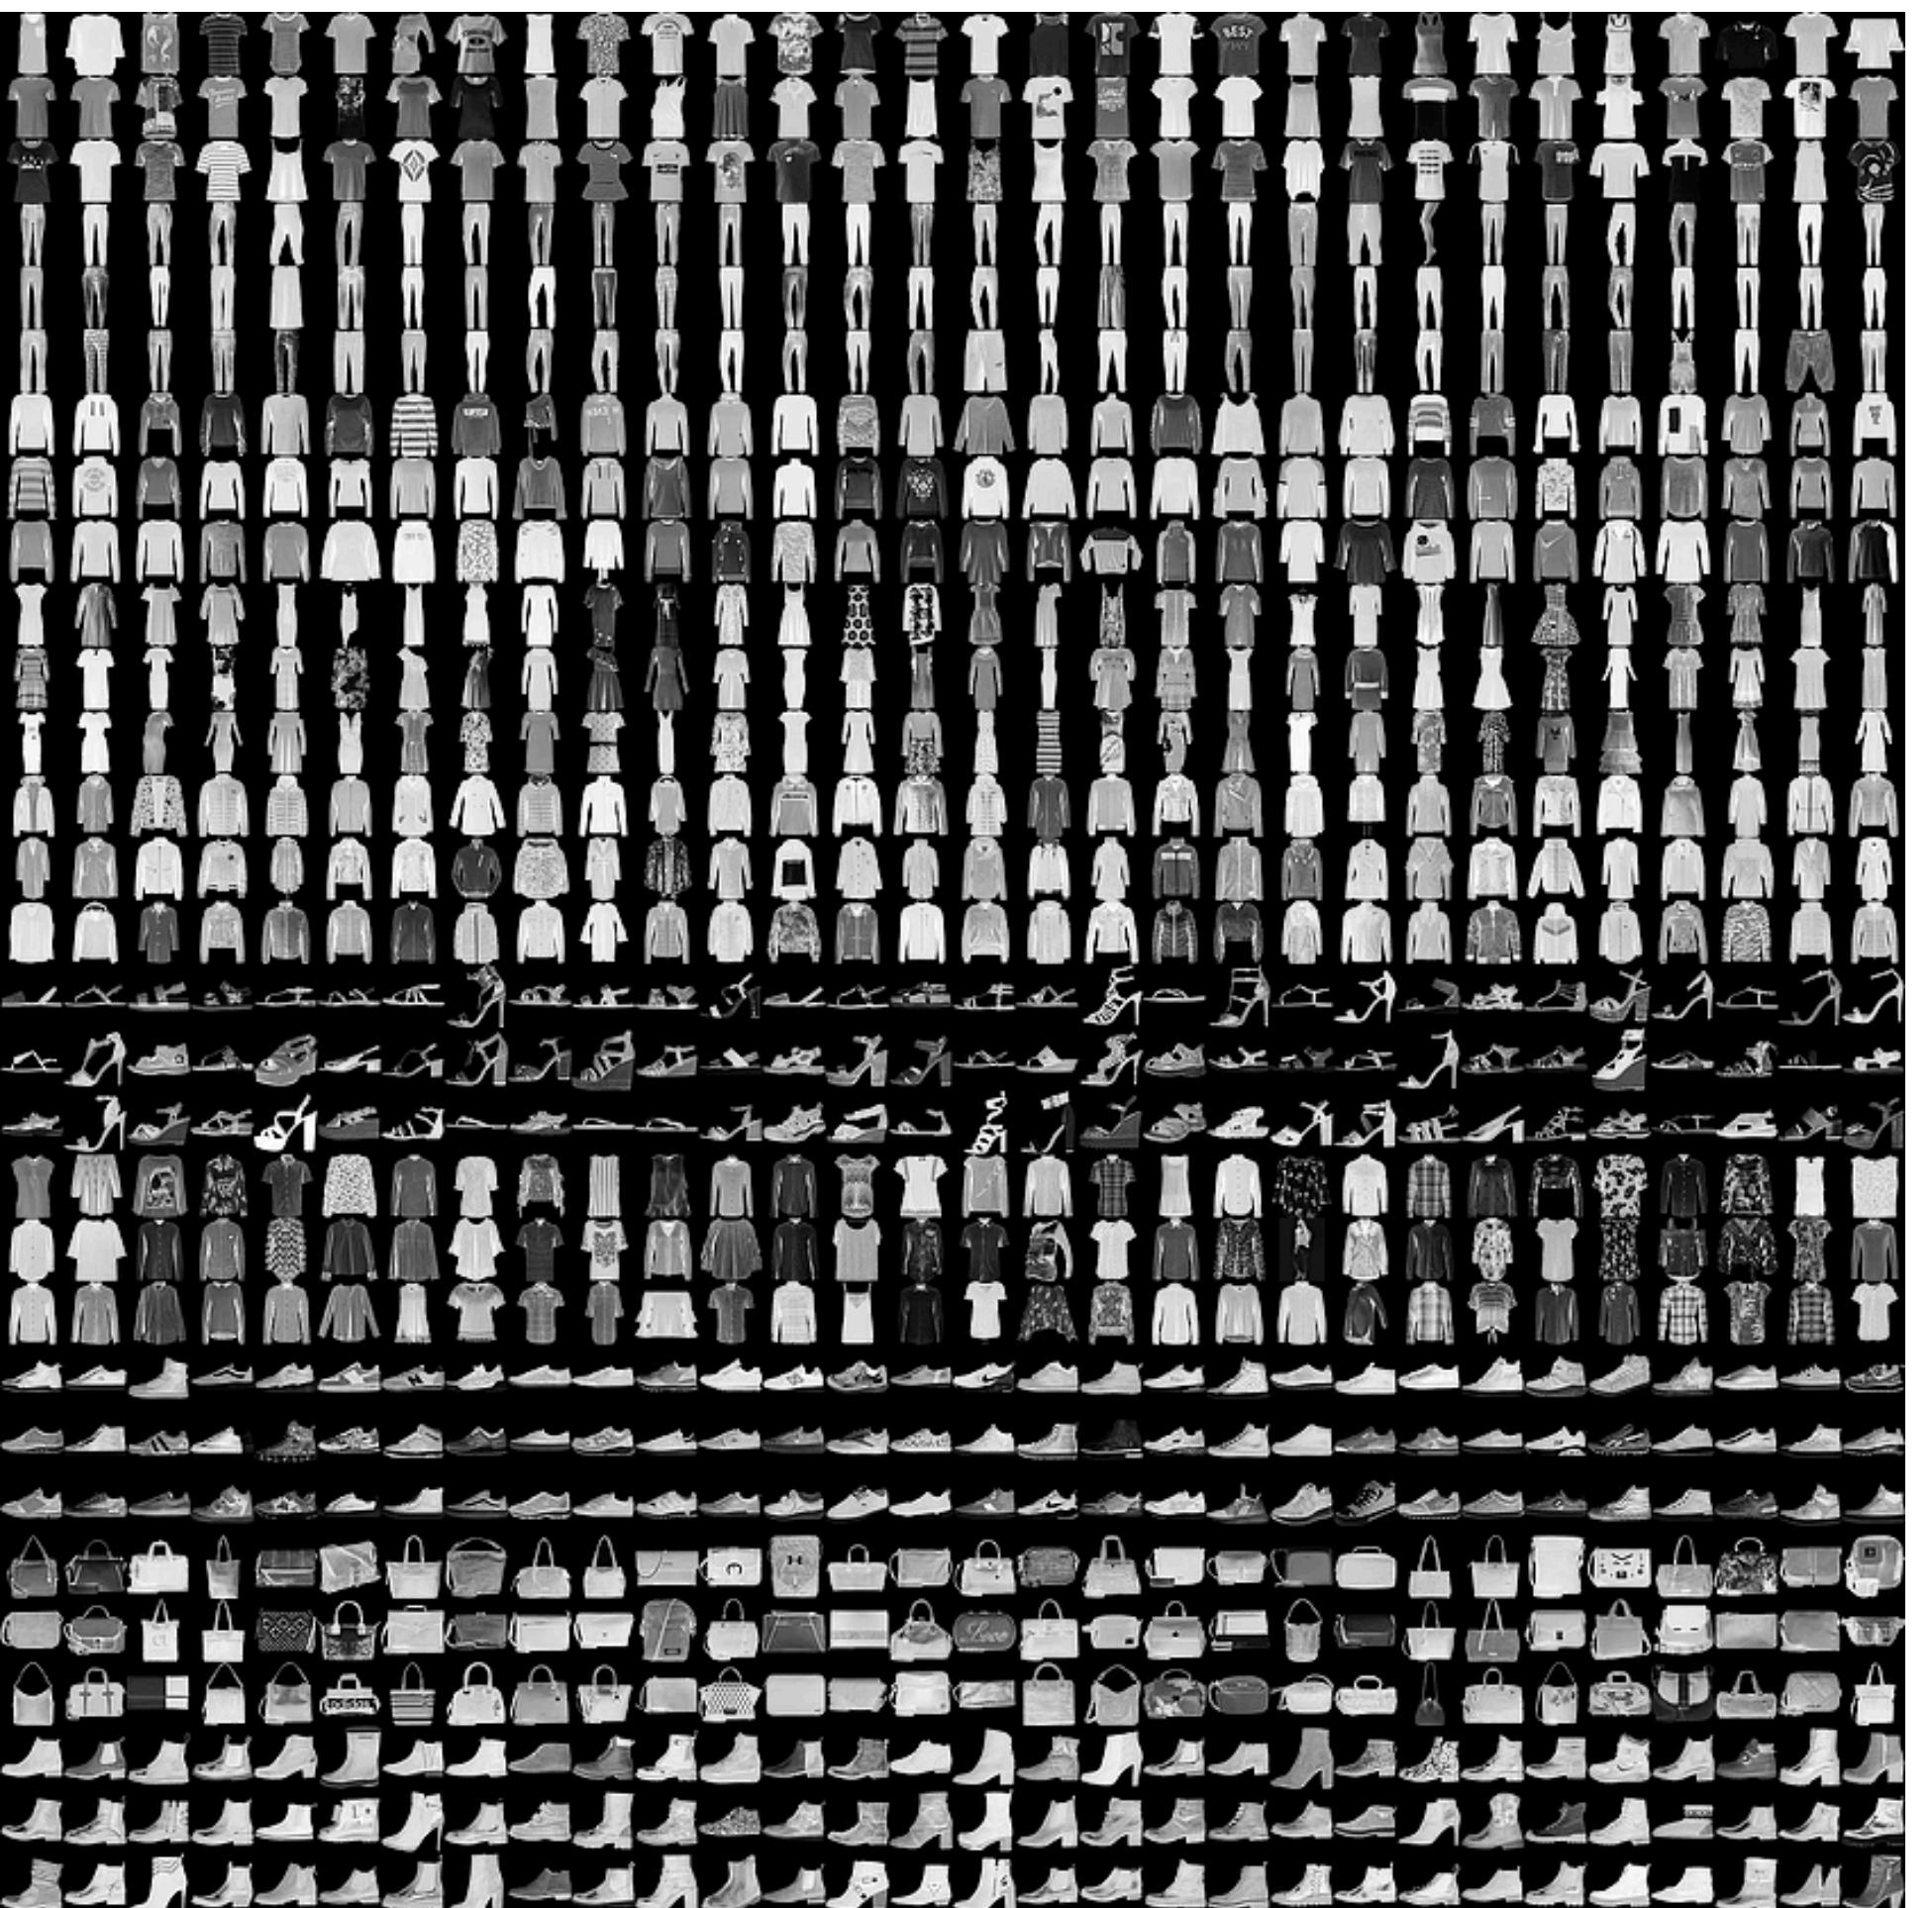
[https://twitter.com/momiji\\_fullmoon/status/1110316960611368960](https://twitter.com/momiji_fullmoon/status/1110316960611368960)

学習率  $\eta$  は歩幅のイメージ ( $\eta$ 小=すり足、 $\eta$ 大=巨人の一歩)

# Fashion MNIST Dataset

- 7万画像
- 10カテゴリ
- 28×28 pixels
- 実験用データセット

ラベル	記述
0	T-shirt/top
1	Trouser
2	Pullover
3	Dress
4	Coat
5	Sandal
6	Shirt
7	Sneaker
8	Bag
9	Ankle boot

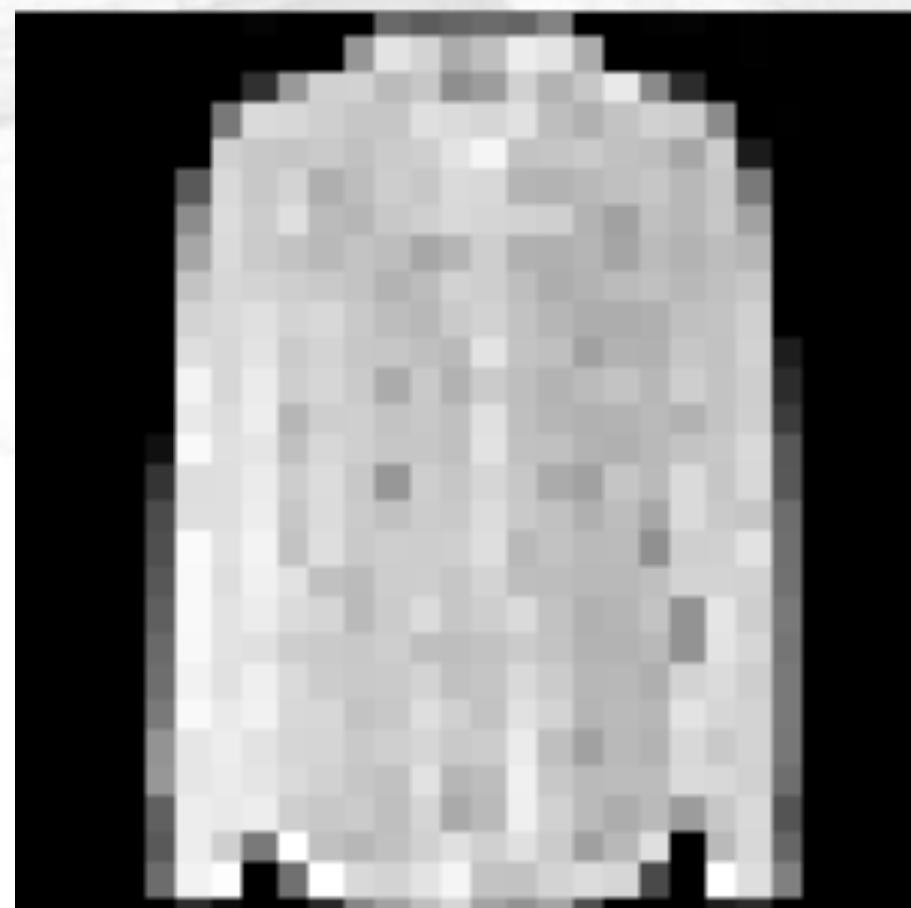


<https://github.com/zalandoresearch/fashion-mnist>

```
1 import tensorflow as tf
2 import matplotlib.pyplot as plt
3
4 fashion_mnist = tf.keras.datasets.fashion_mnist
5 (train_images, train_labels), (test_images, test_labels) = fashion_mnist.load_data()
6 train_images = train_images / 255.
7 test_images = test_images / 255.
```

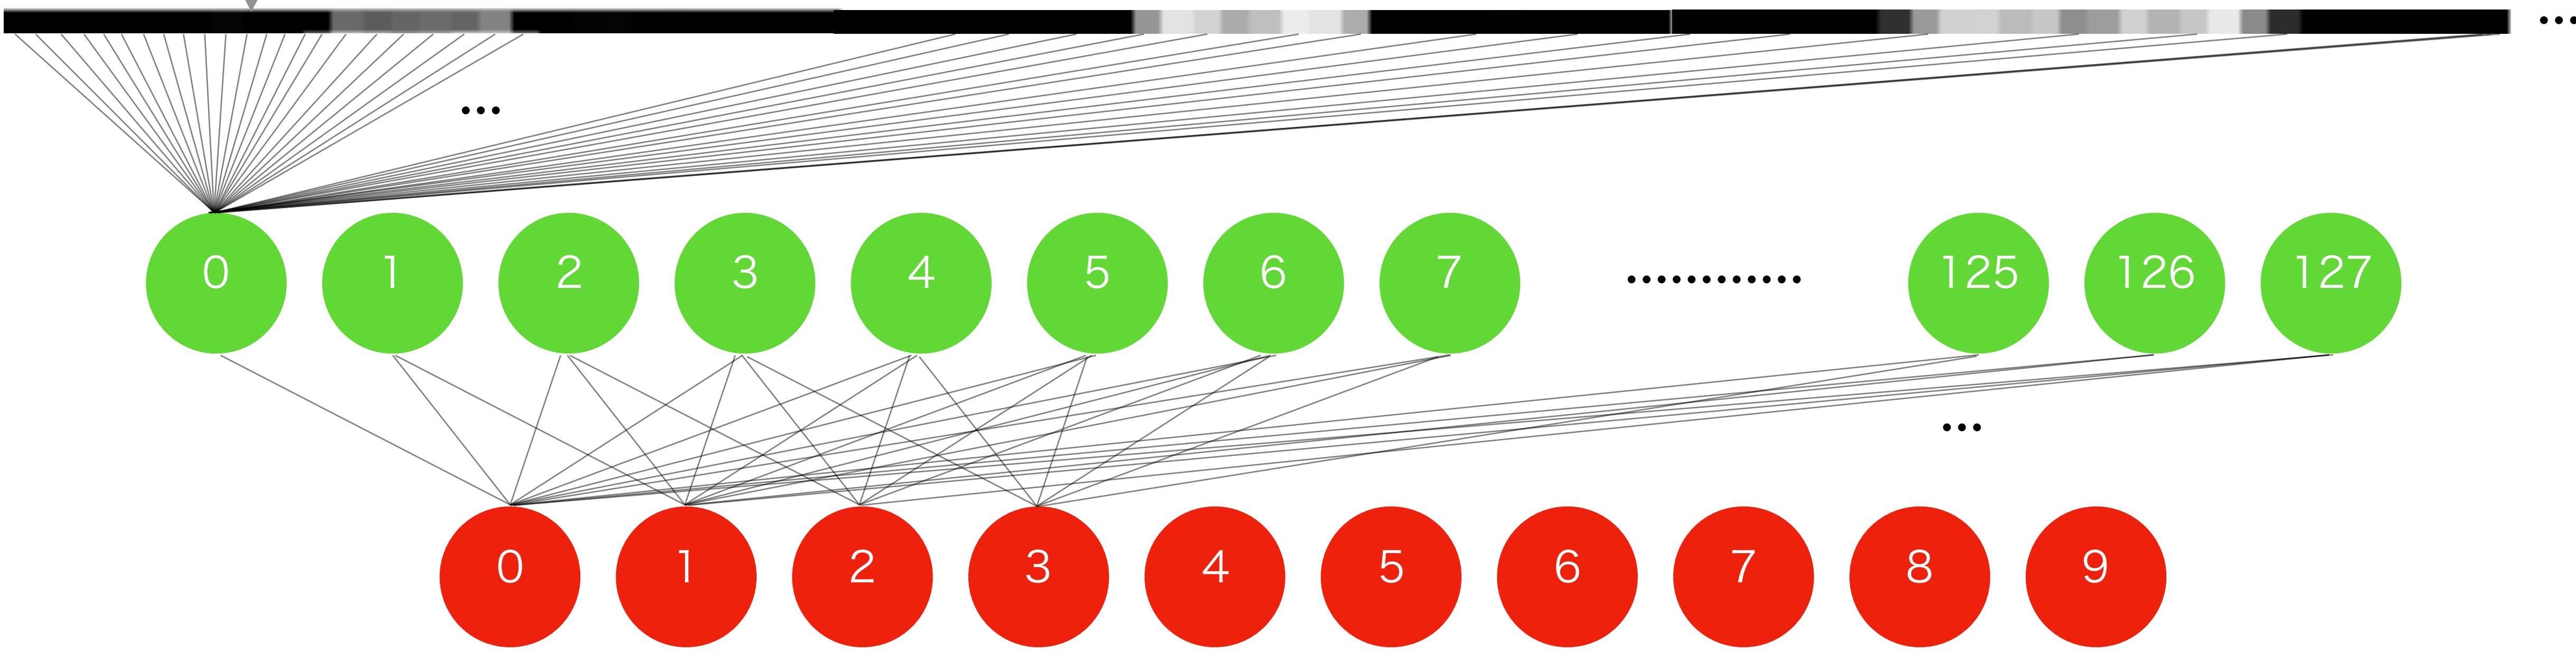
```
1 import tensorflow as tf
2 import matplotlib.pyplot as plt
3
4 fashion_mnist = tf.keras.datasets.fashion_mnist
5 (train_images, train_labels), (test_images, test_labels) = fashion_mnist.load_data()
6 train_images = train_images / 255.
7 test_images = test_images / 255.
8
9 model = tf.keras.Sequential([
10     tf.keras.layers.Flatten(input_shape=(28, 28)),
11     tf.keras.layers.Dense(128, activation=tf.nn.relu),
12     tf.keras.layers.Dense(10, activation=tf.nn.softmax)
13 ])
```

```
1 import tensorflow as tf
2 import matplotlib.pyplot as plt
3
4 fashion_mnist = tf.keras.datasets.fashion_mnist
5 (train_images, train_labels), (test_images, test_labels) = fashion_mnist.load_data()
6 train_images = train_images / 255.
7 test_images = test_images / 255.
8
9 model = tf.keras.Sequential([
10     tf.keras.layers.Flatten(input_shape=(28, 28)),
11     tf.keras.layers.Dense(128, activation=tf.nn.relu),
12     tf.keras.layers.Dense(10, activation=tf.nn.softmax)
13 ])
14
15 model.compile(optimizer=tf.train.AdamOptimizer(),
16                 loss='sparse_categorical_crossentropy', metrics=['accuracy'])
16 model.fit(train_images, train_labels, epochs=5, batch_size=128, verbose=1,
17             validation_data=(test_images, test_labels))
```



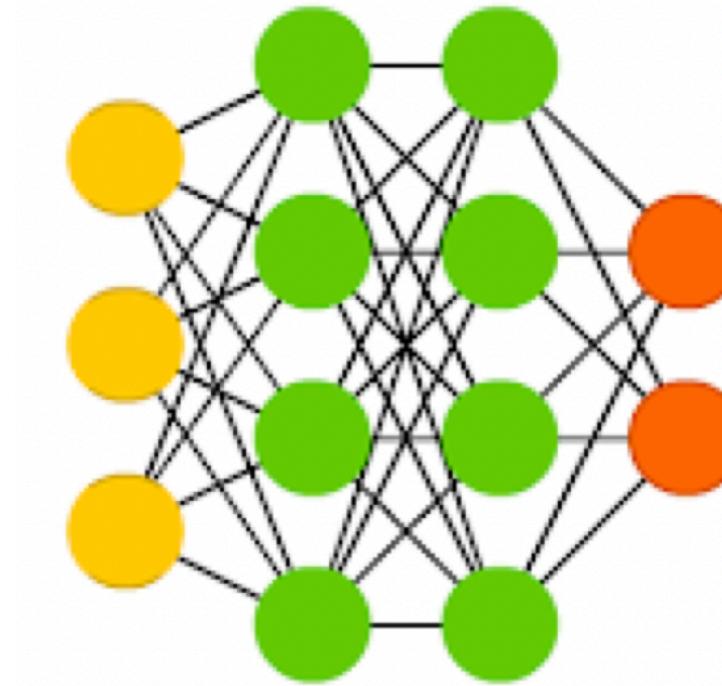
```
9 model = tf.keras.Sequential([
10     tf.keras.layers.Flatten(input_shape=(28, 28)),
11     tf.keras.layers.Dense(128, activation=tf.nn.relu),
12     tf.keras.layers.Dense(10, activation=tf.nn.softmax)
13 ])
```

Flatten: (28, 28) => (784)



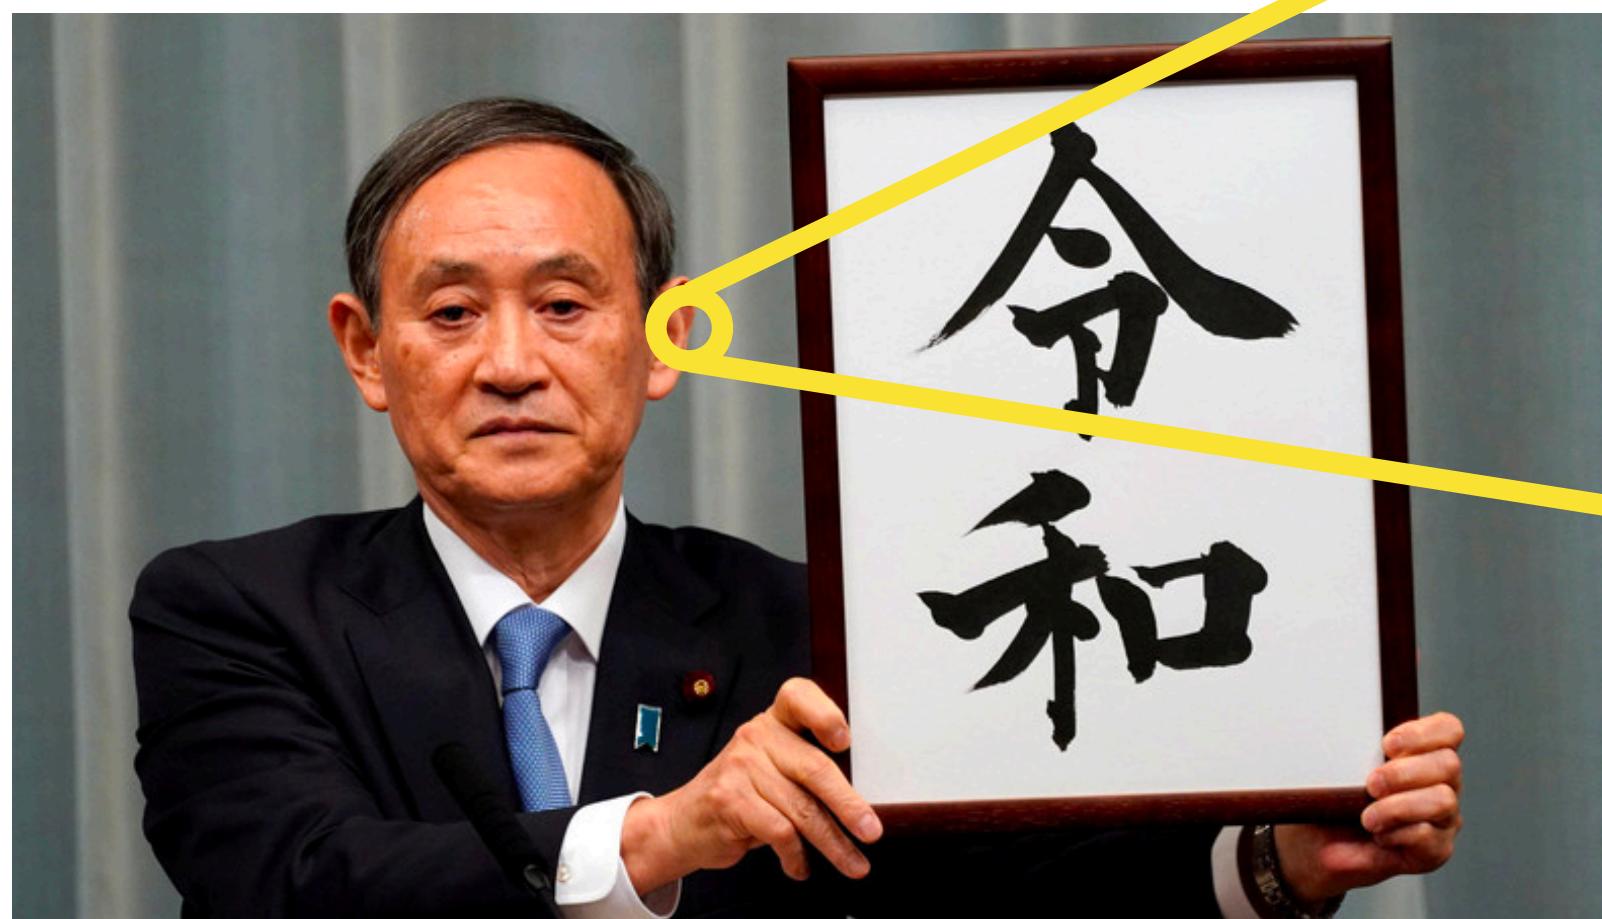
# Dense Layer の欠点

- 入力のベクトルの**全要素**の相関をみている
  - > 住宅価格予測みたいな話ならまだいい
  - > もう作ってる特徴量と特徴量の組み合わせ
    - 例) 東京墨田区 & 床面積 30m<sup>2</sup> & 1K & 風呂トイレ別 & 新築 => 家賃月10万円
- 「**画像の特徴量**」を抽出してからDense Layerに渡せば効率的



**畳込みニューラルネットワーク (Convolutional Neural Network; CNN)**

# 畳み込み (Convolution)



144	60	19
188	82	32
156	55	27

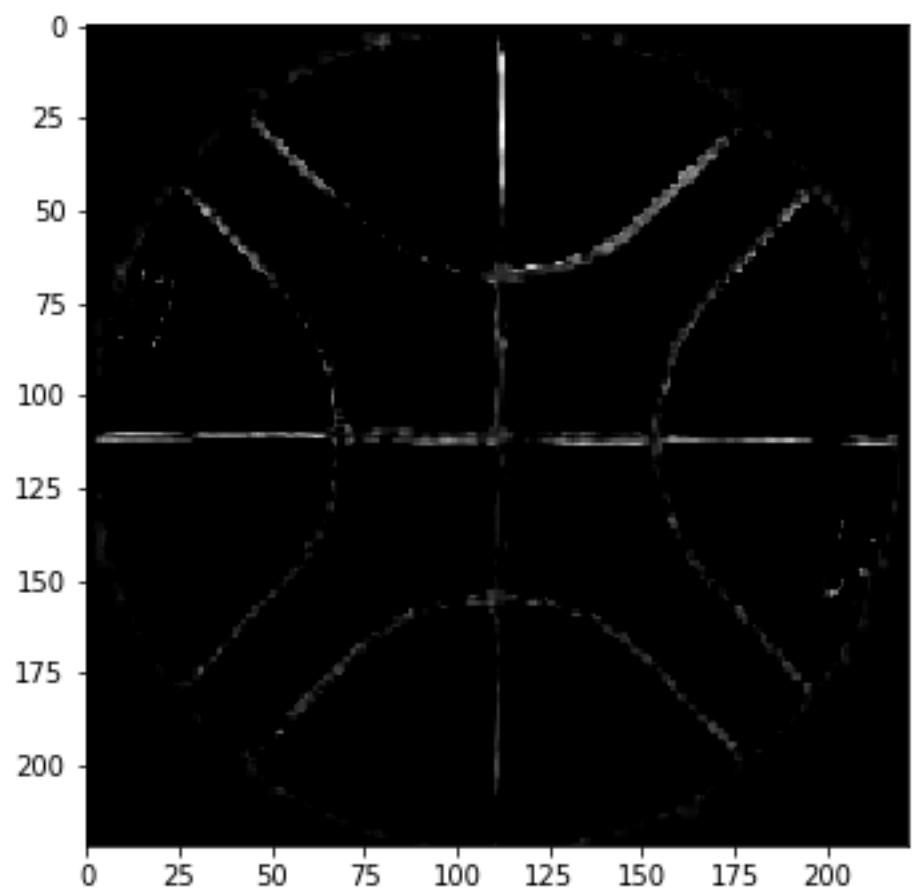
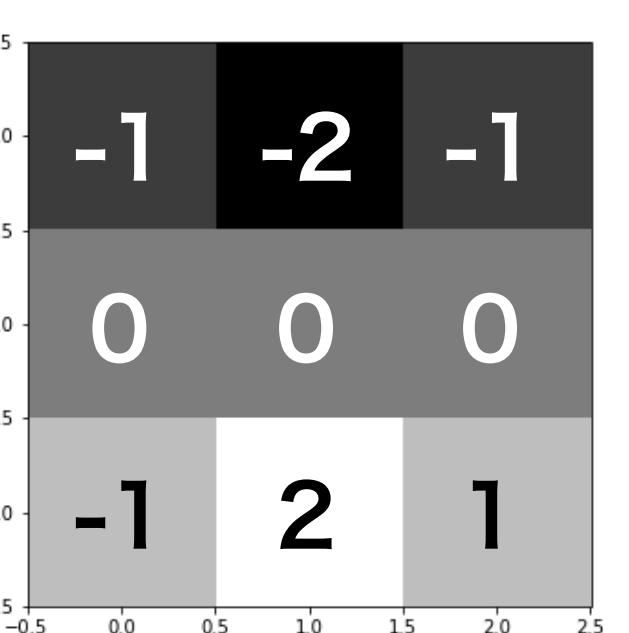
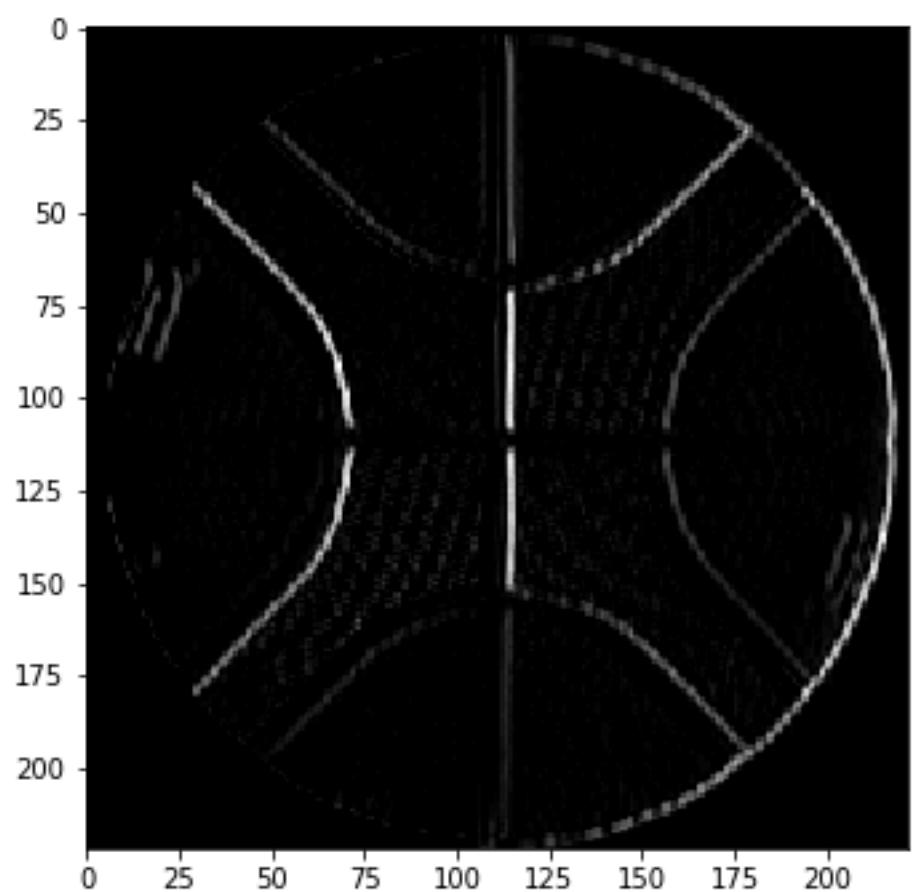
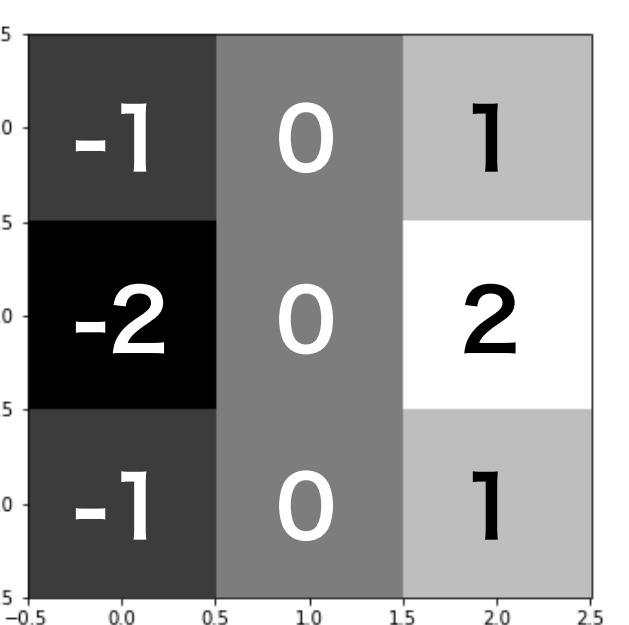
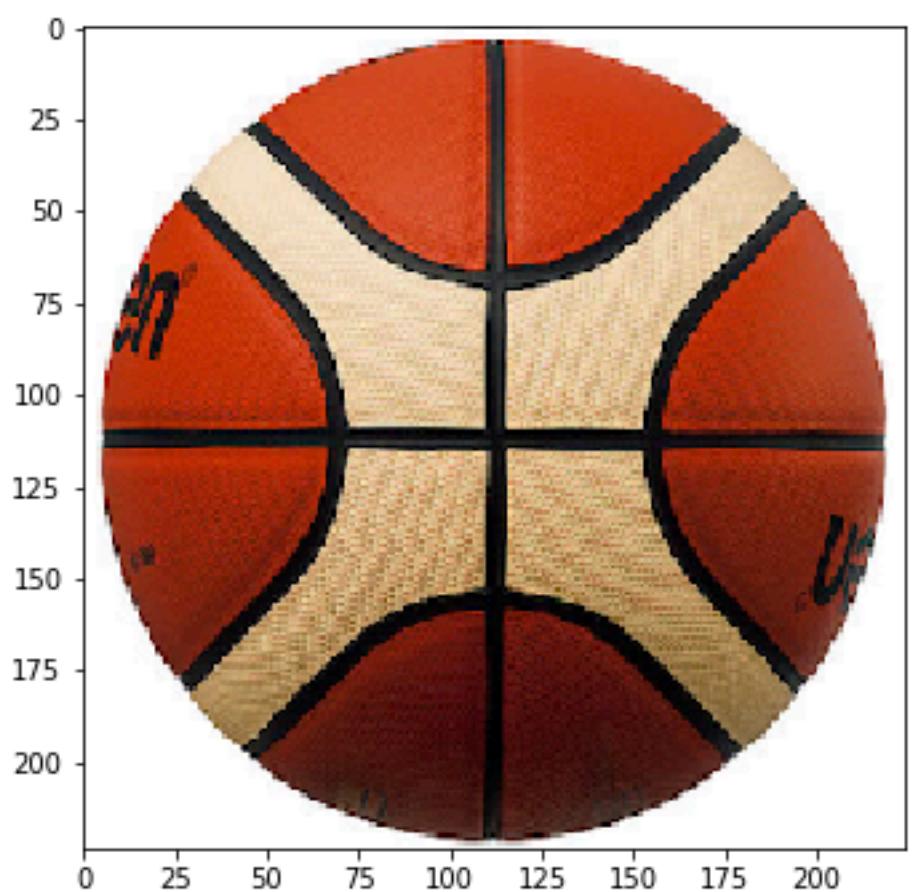
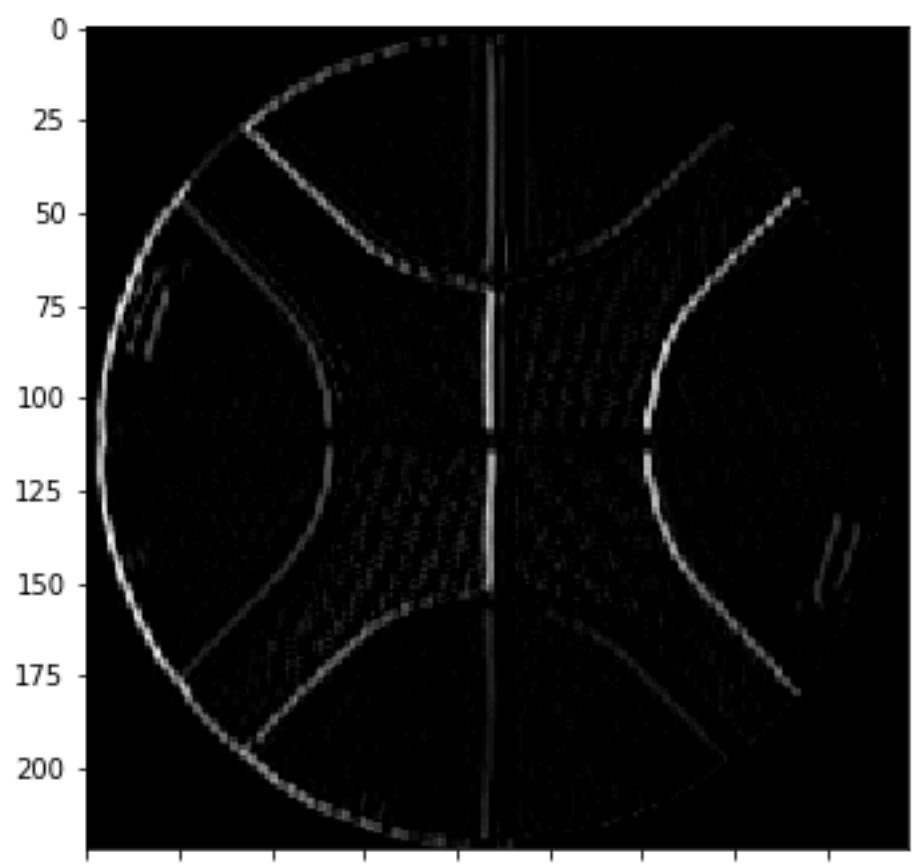
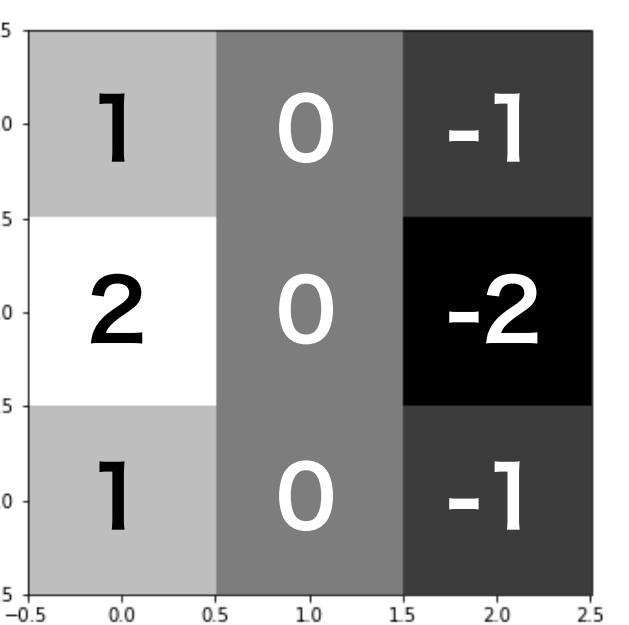
-1	0	-2
0.5	4.5	-1.5
1.5	2	-3

CURRENT\_PIXEL\_VALUE = 82

NEW\_PIXEL\_VALUE =

$$\begin{aligned} & (-1 * 144) + (0 * 60) + (-2 * 19) \\ & + (0.5 * 188) + (4.5 * 82) + (-1.5 * 32) \\ & + (1.5 * 156) + (2 * 55) + (-3 * 27) \end{aligned}$$

$$u_{ijm} = \sum_{k=0}^{K-1} \sum_{p=0}^{W-1} \sum_{q=0}^{H-1} z_{i+p, j+q, k}^{(l-1)} h_{pqkm} + b_{ijm}$$



# stride=(1, 1), padding='valid'



# stride=(1, 1), padding='valid'



stride=(1, 1), padding='valid'



stride=(1, 1), padding='valid'



stride=(1, 1), padding='valid'



# stride=(1, 1), padding='valid'



# stride=(1, 1), padding='valid'



stride=(2, 2), padding='valid'



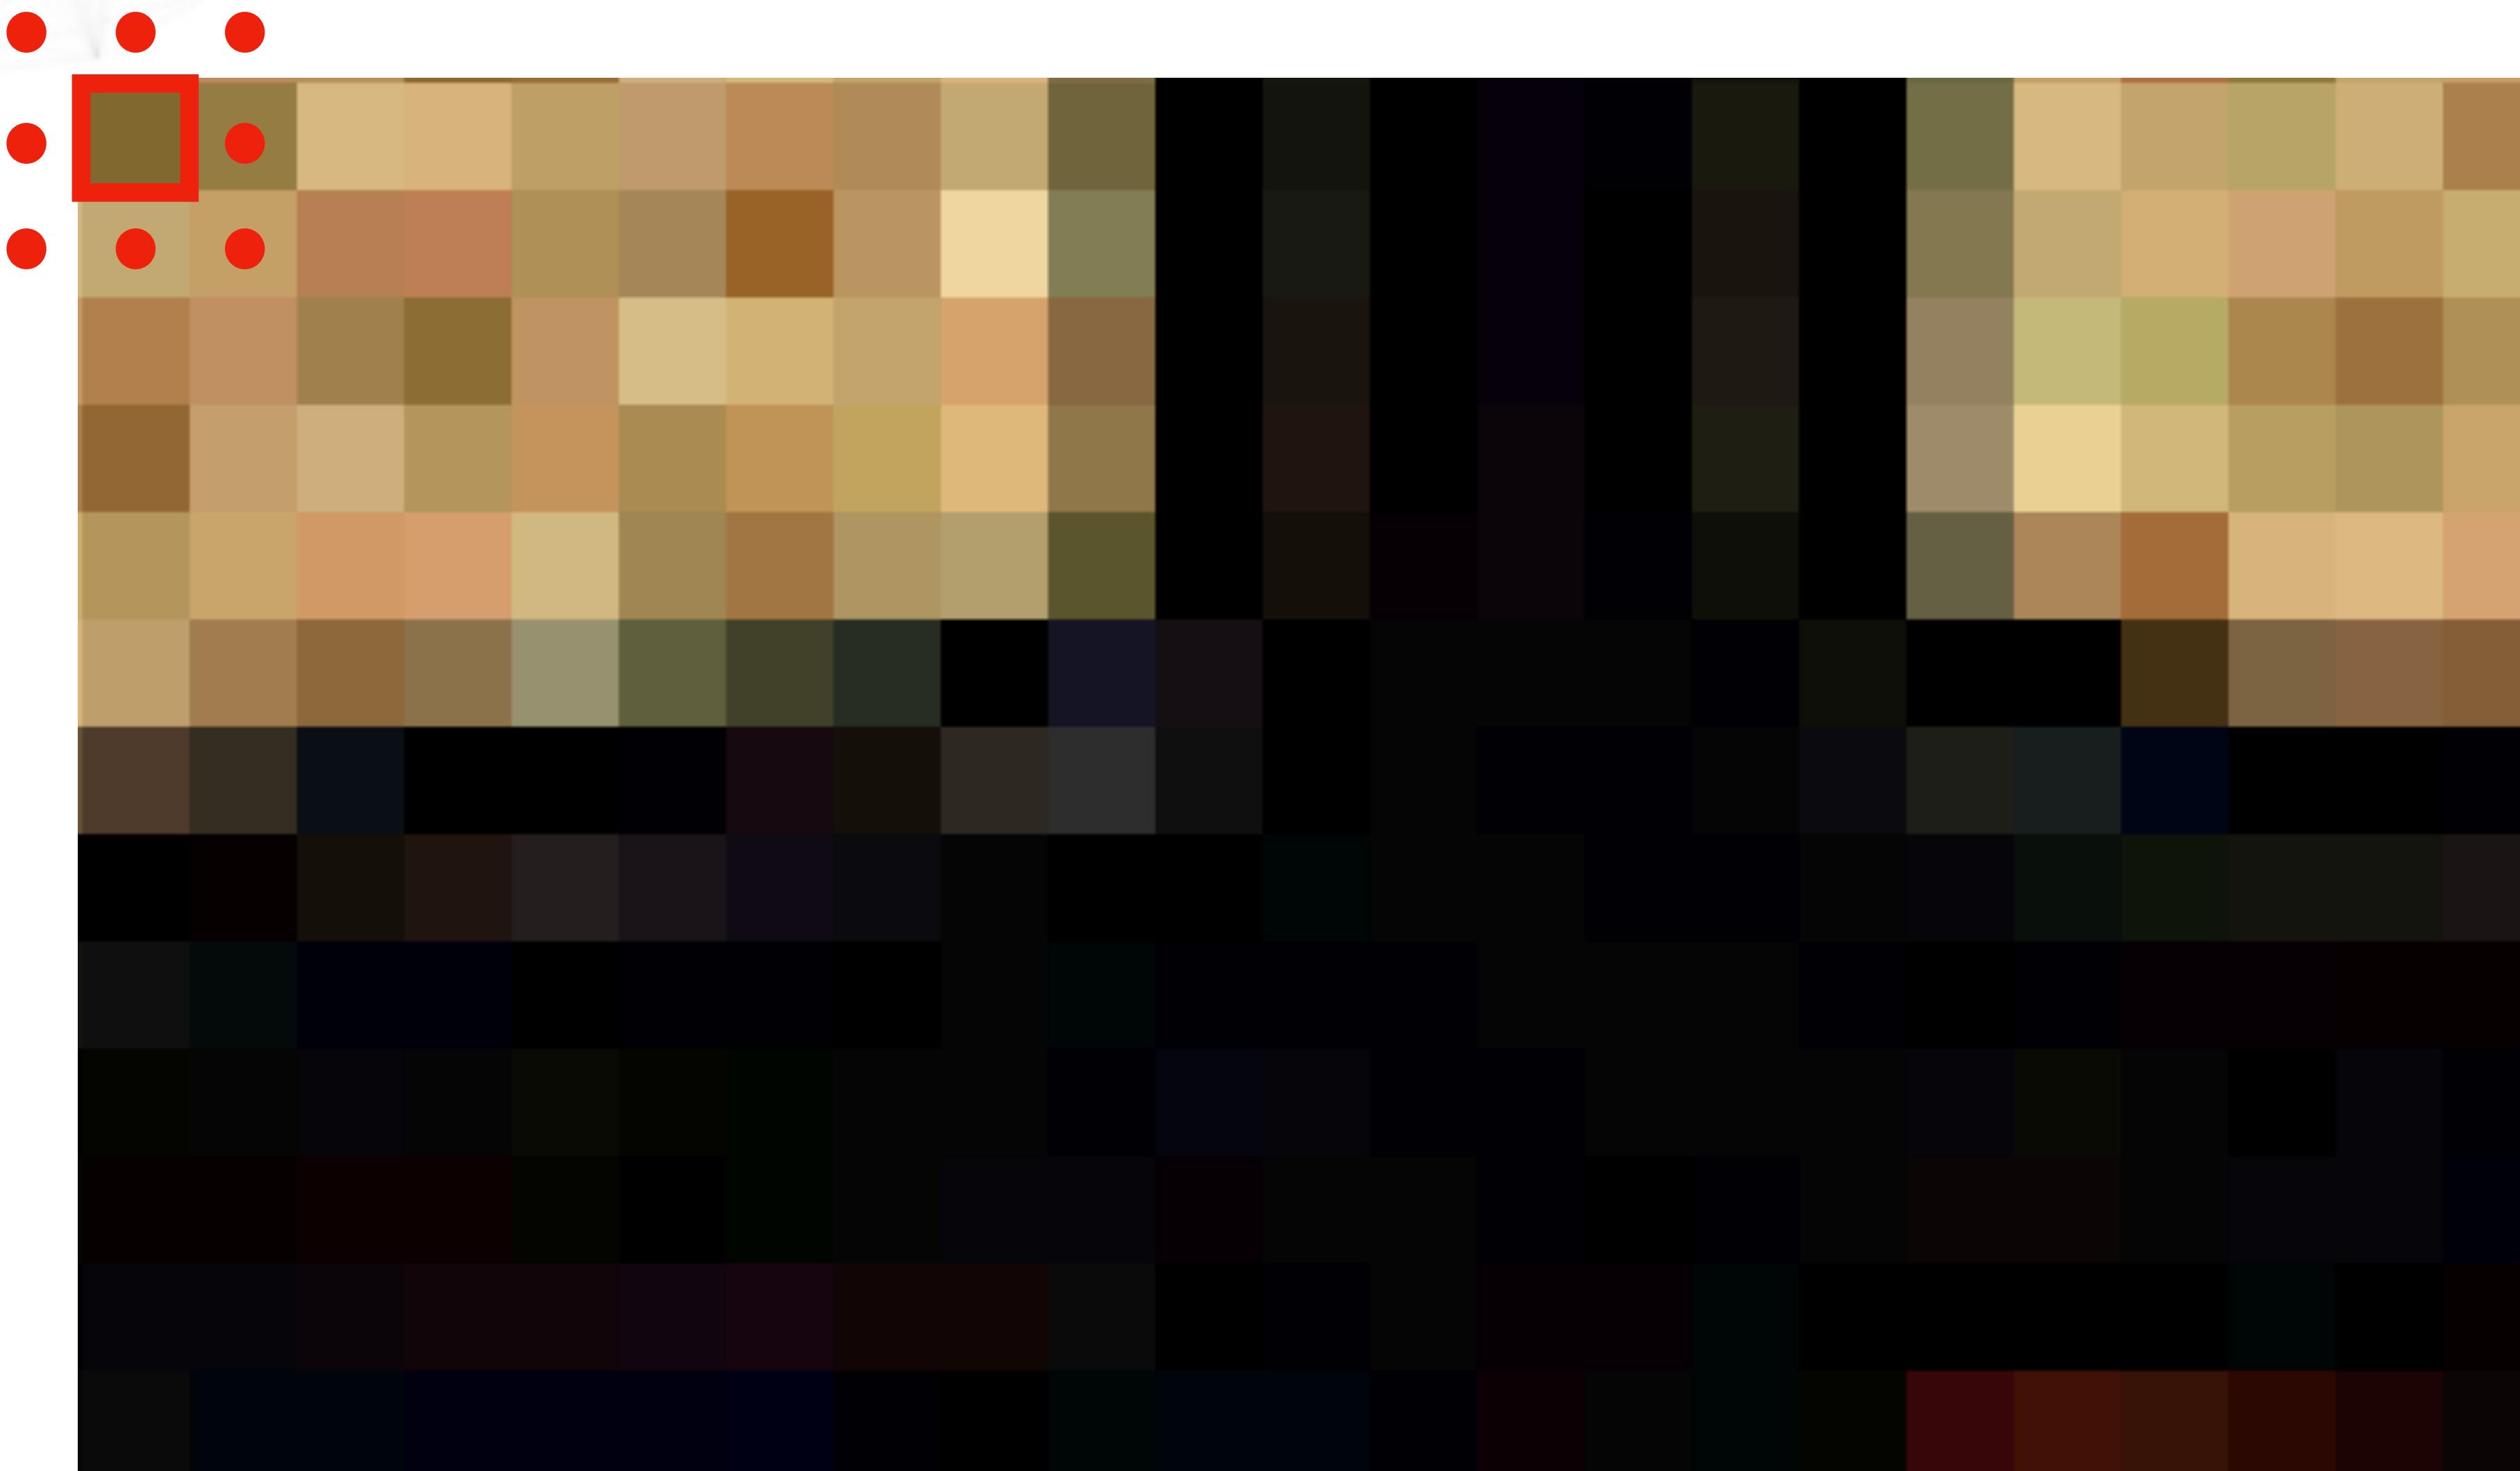
stride=(2, 2), padding='valid'



stride=(2, 2), padding='valid'



stride=(1, 1), padding='same'



# stride=(1, 1), padding='same'



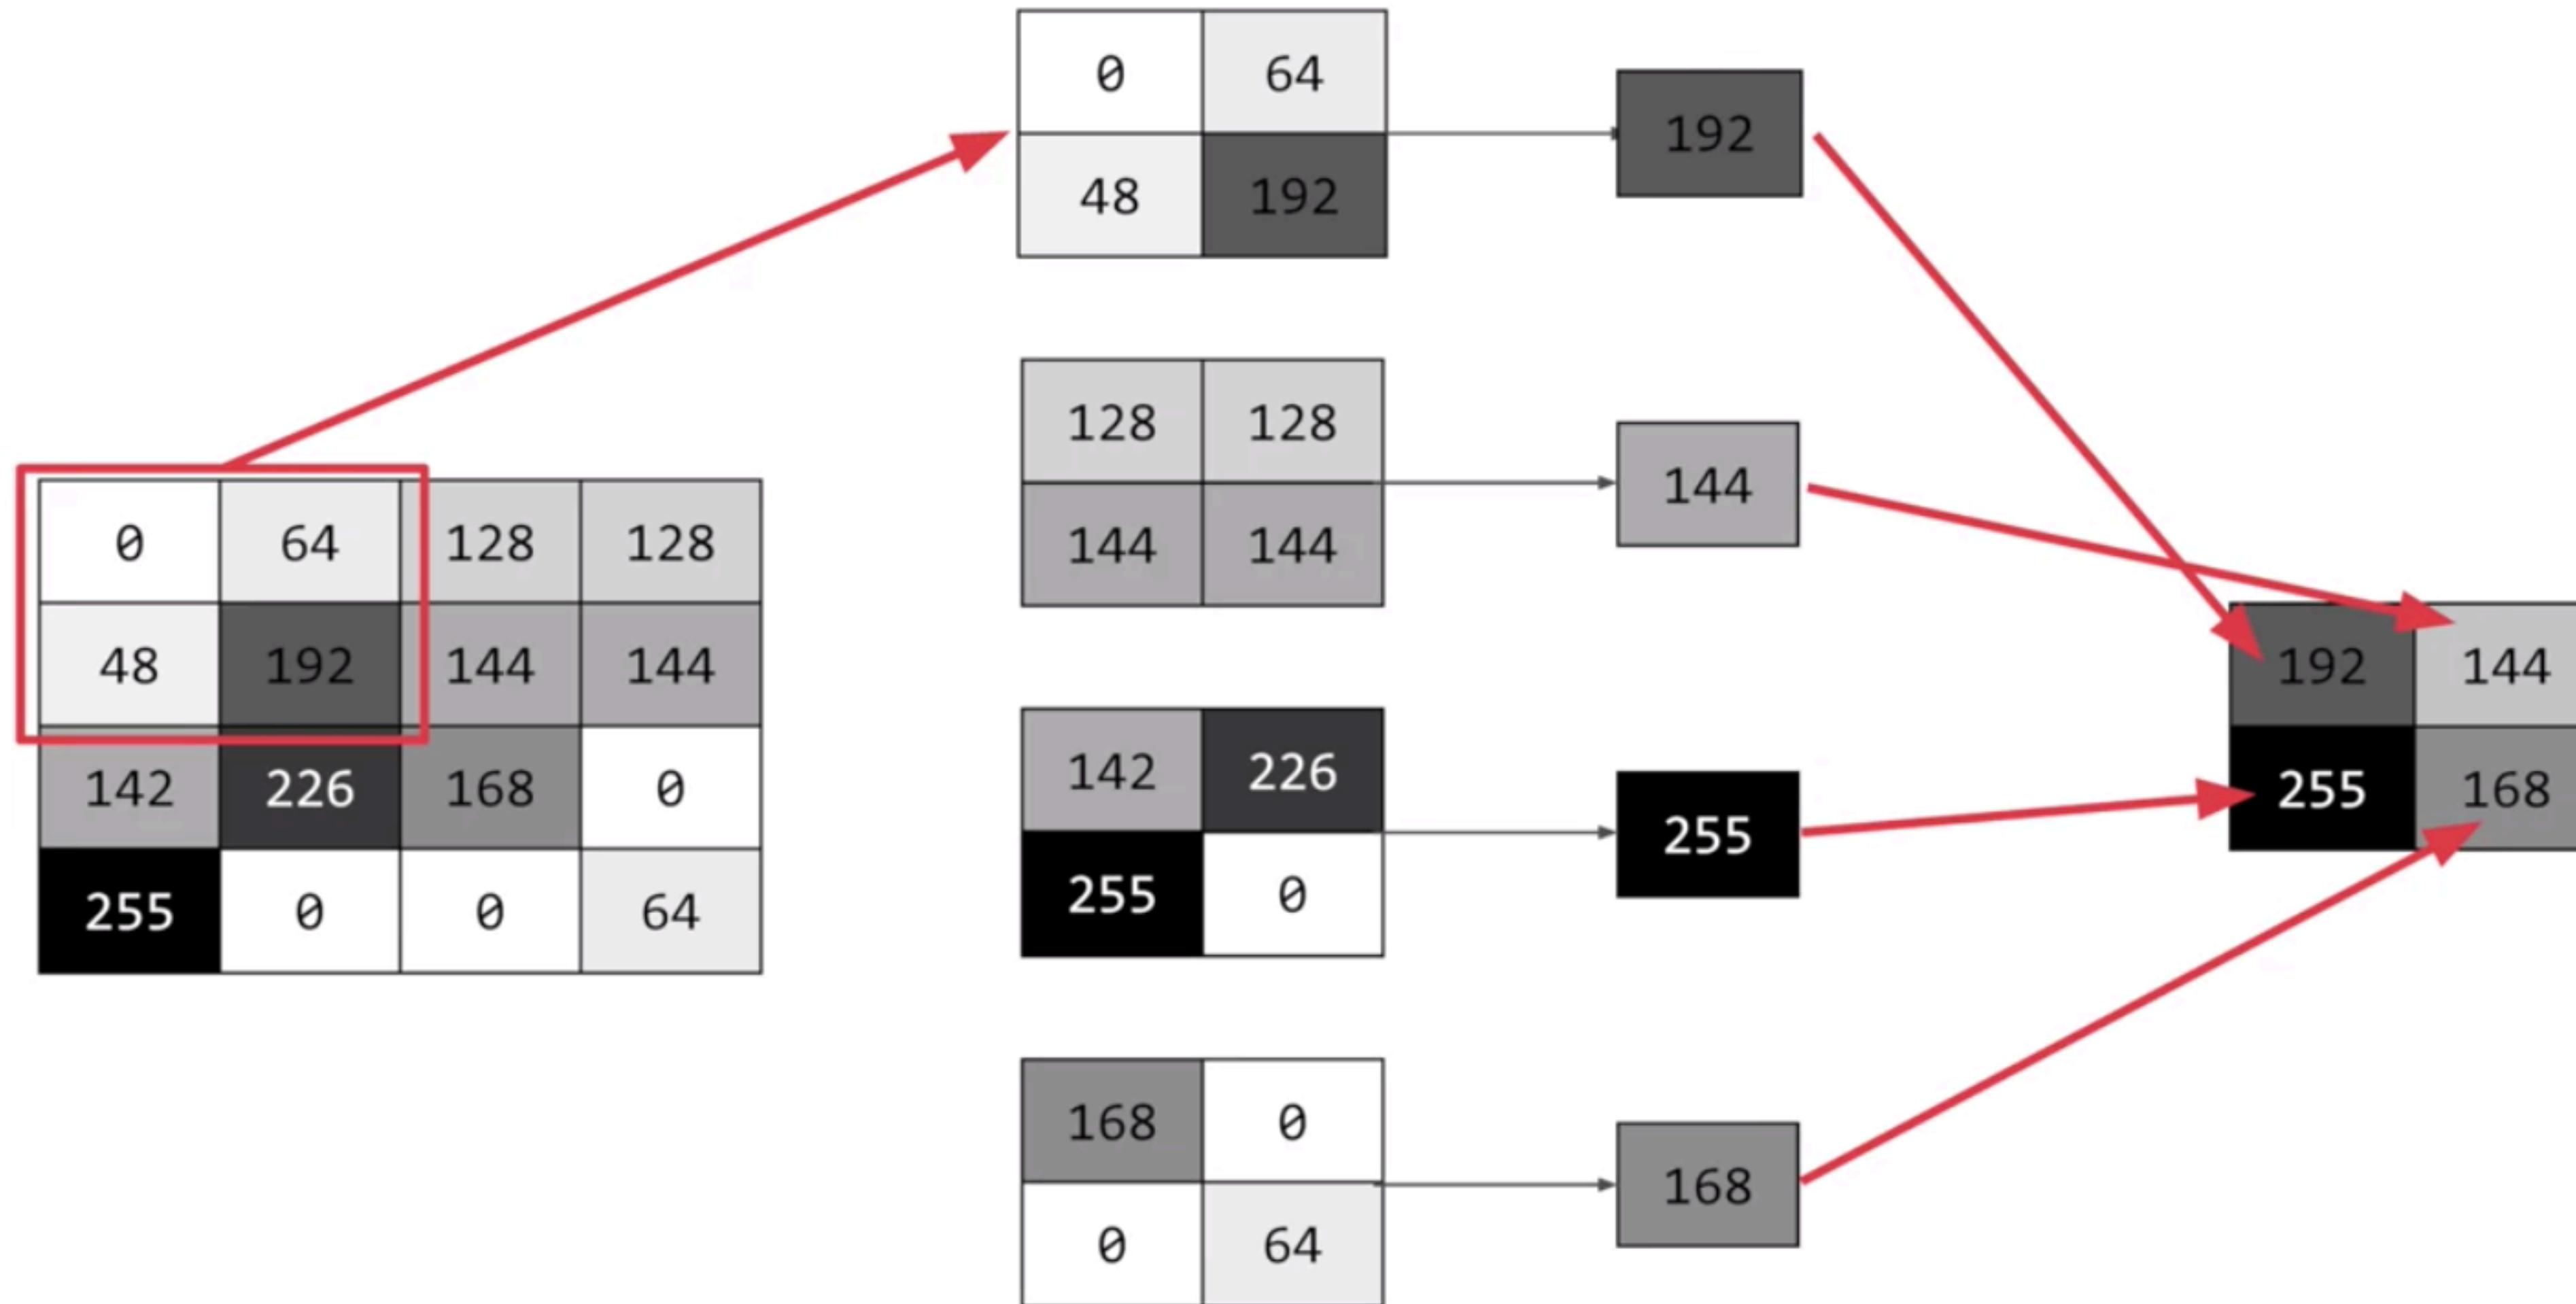
stride=(1, 1), padding='same'



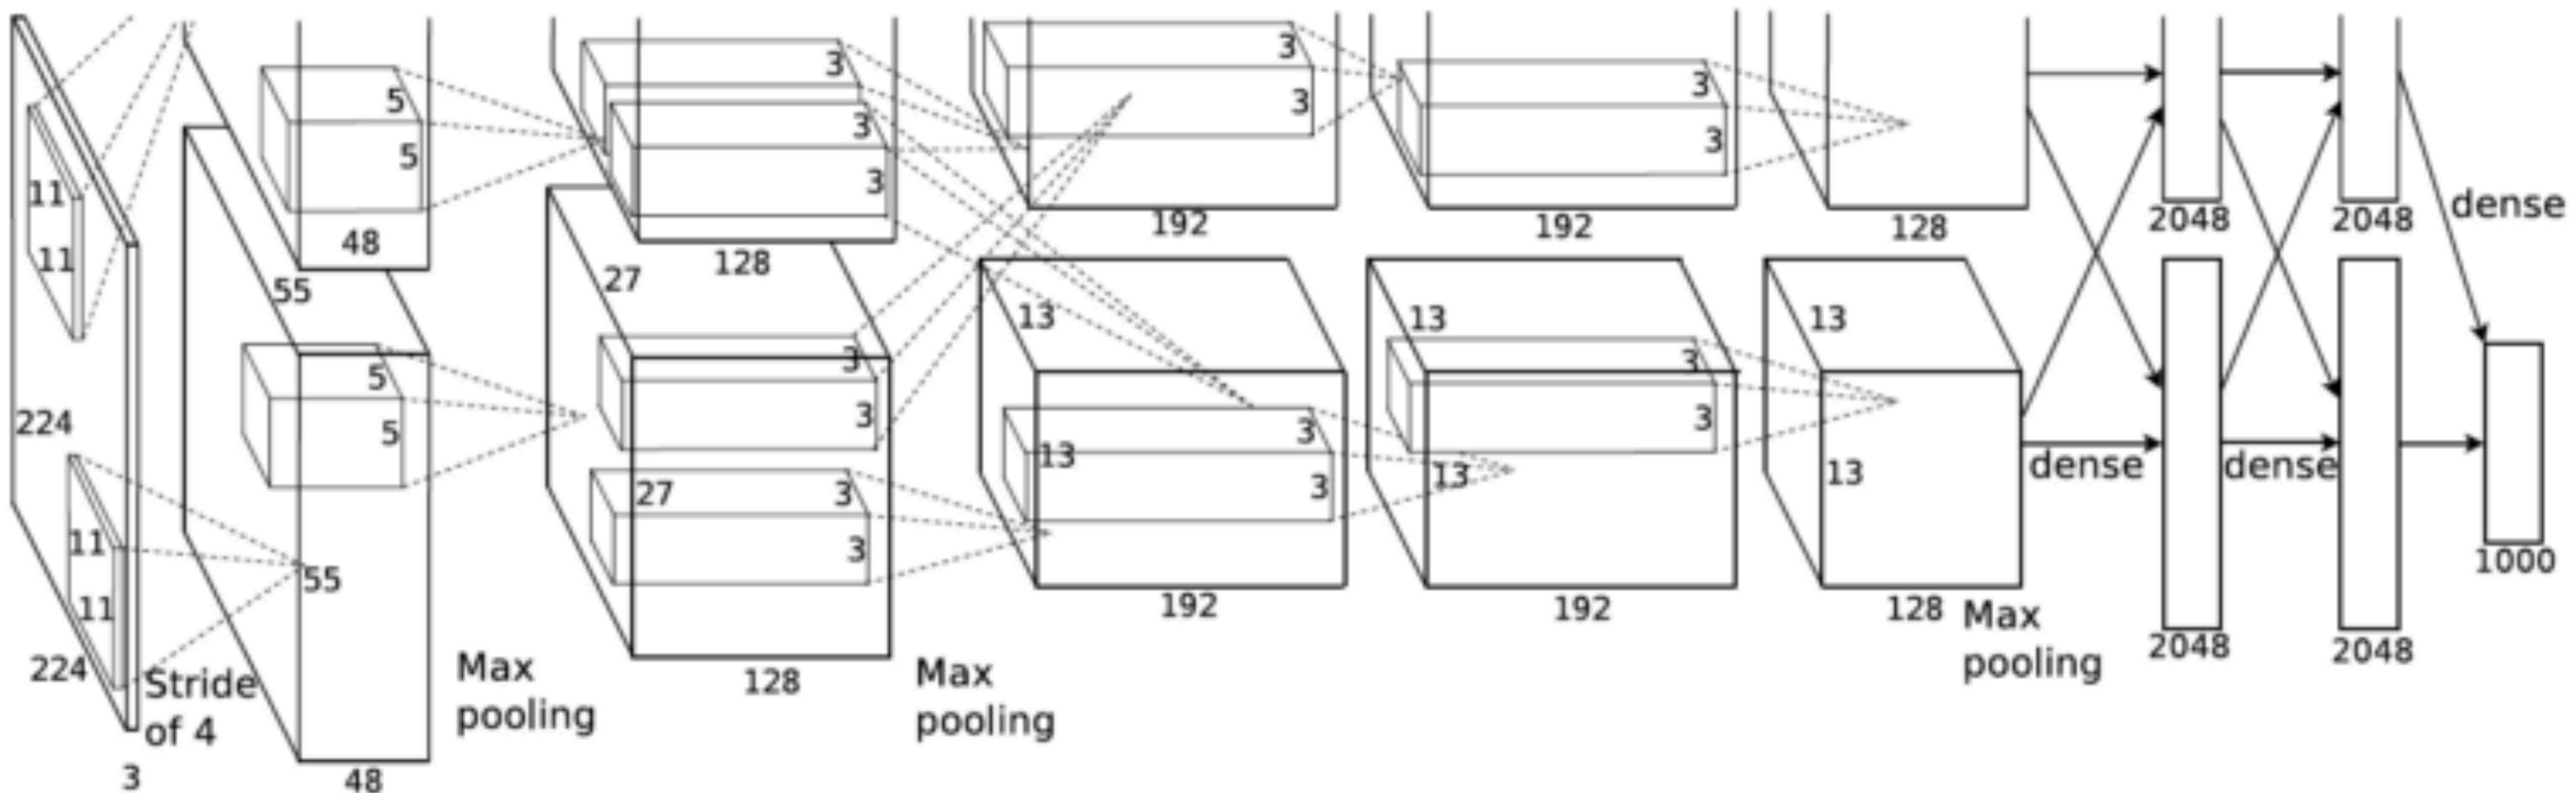
stride=(1, 1), padding='same'

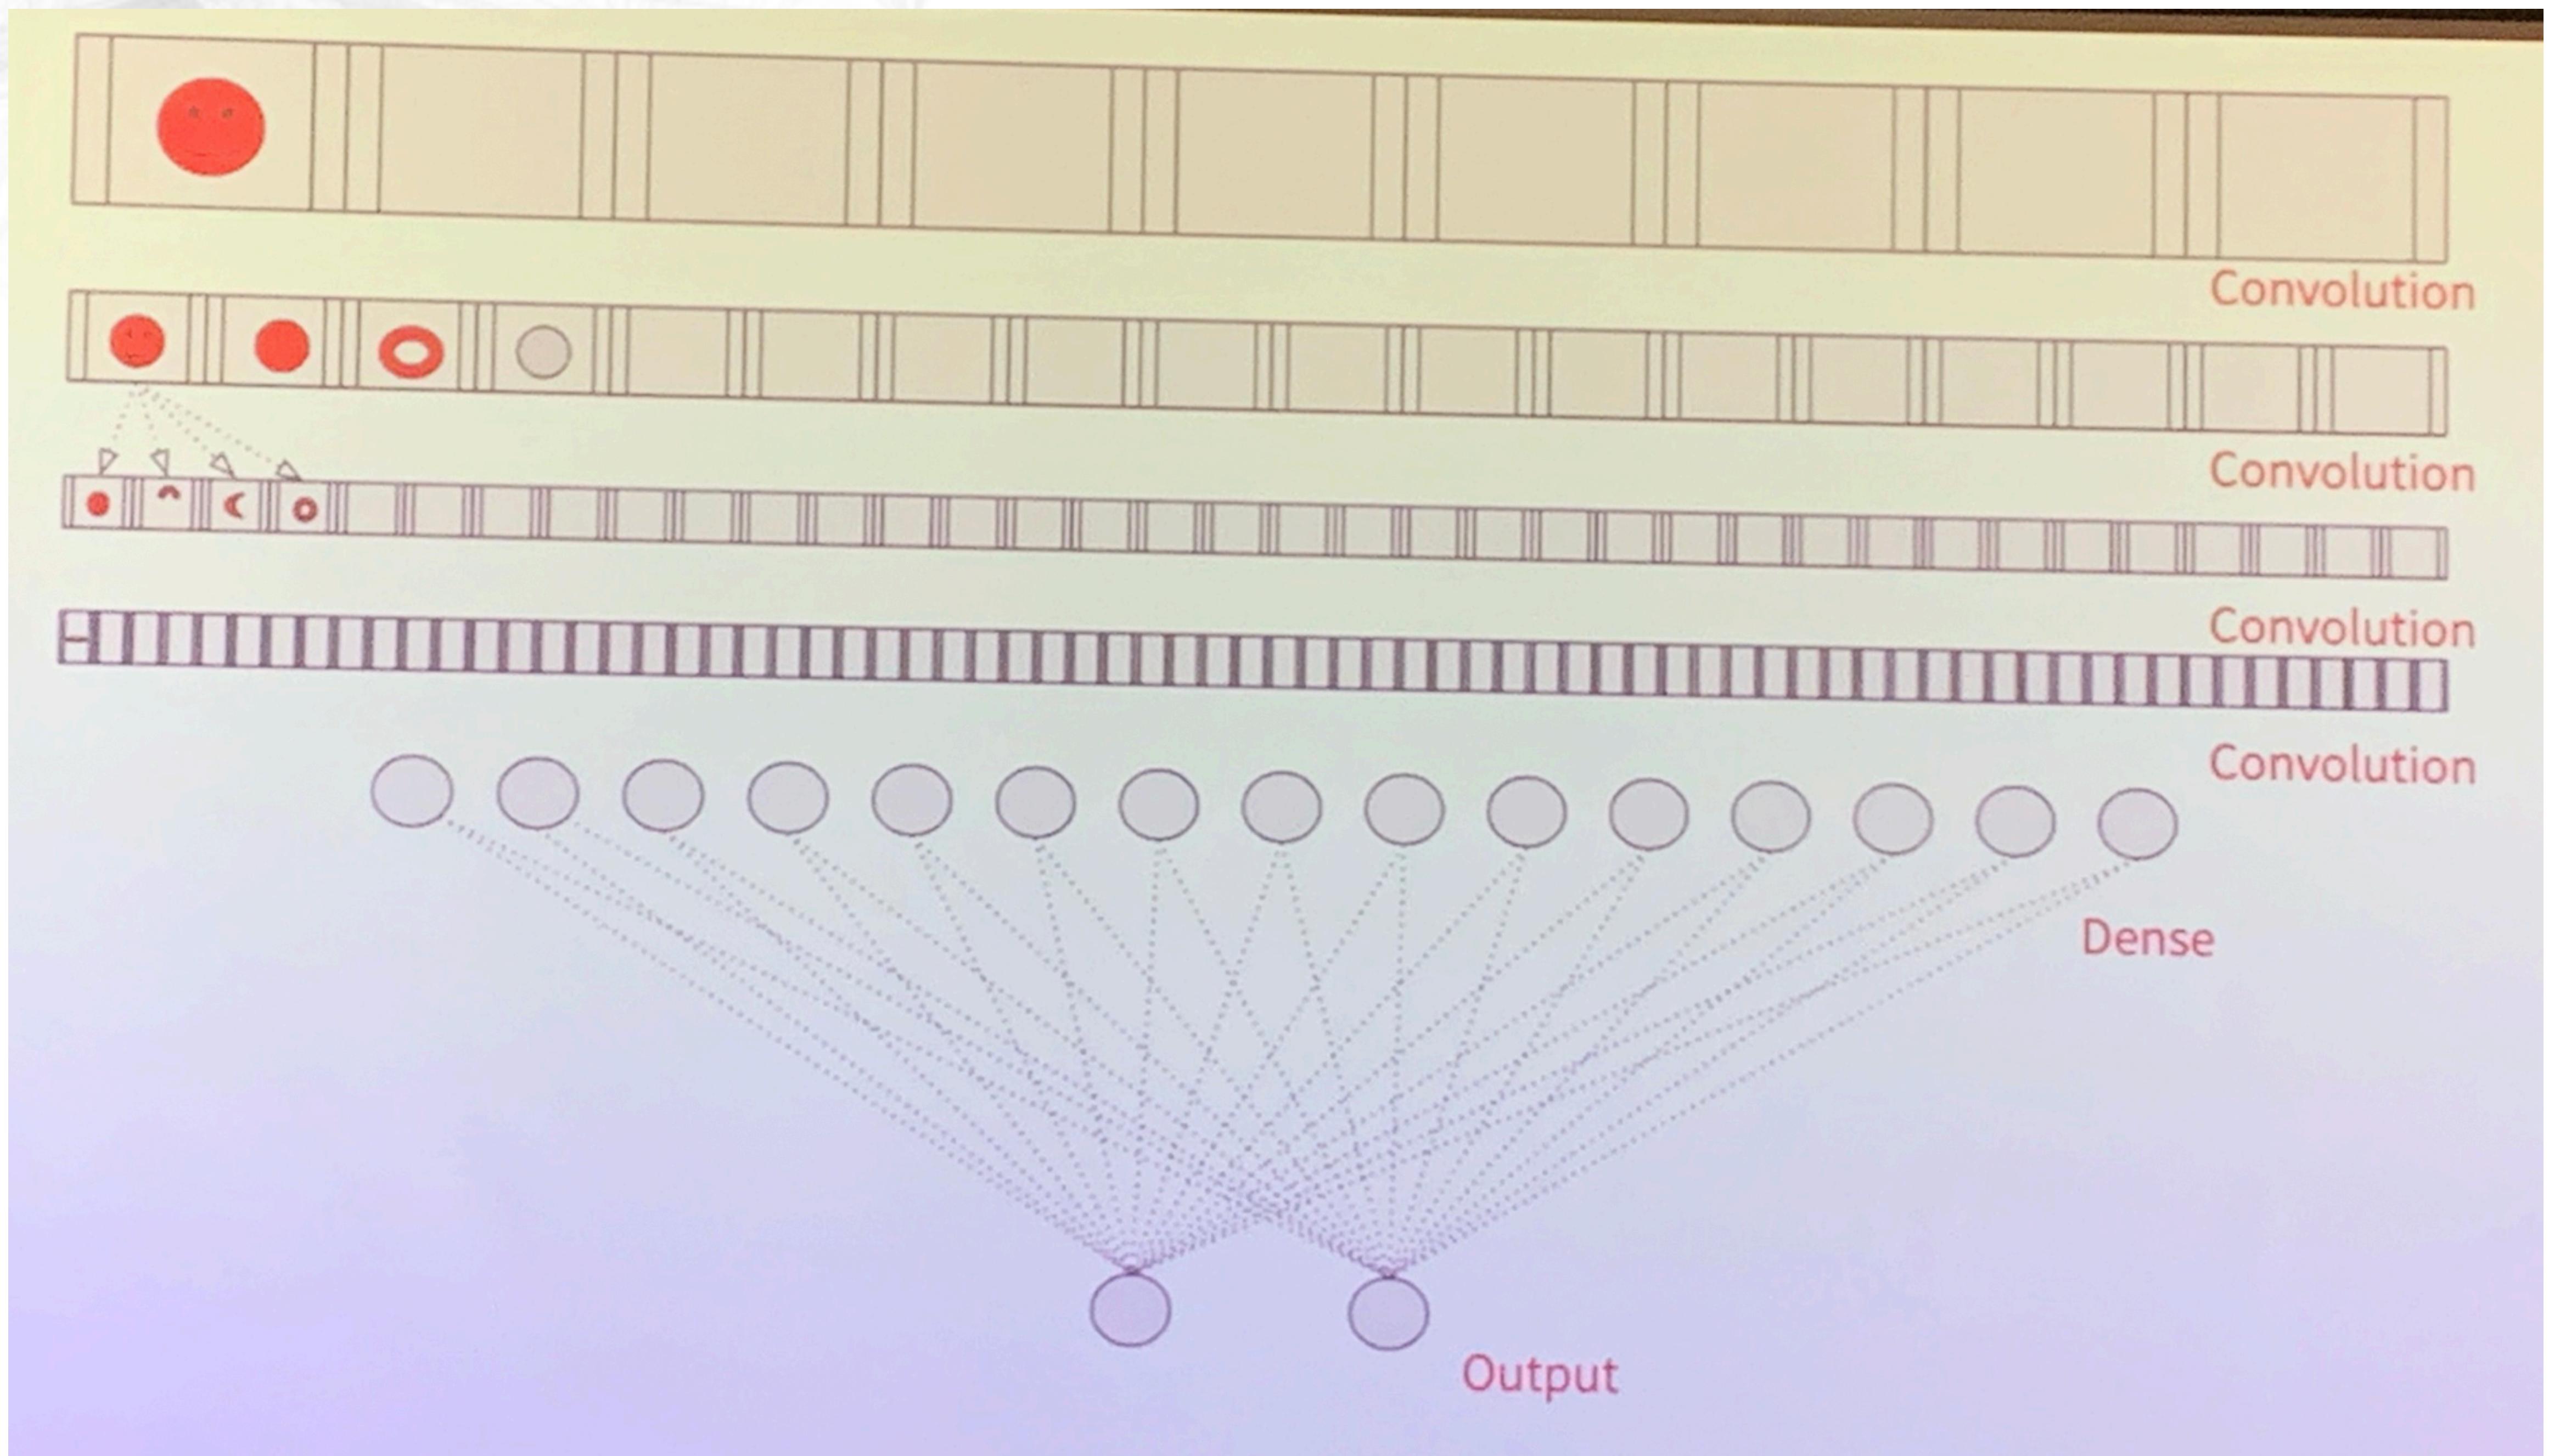


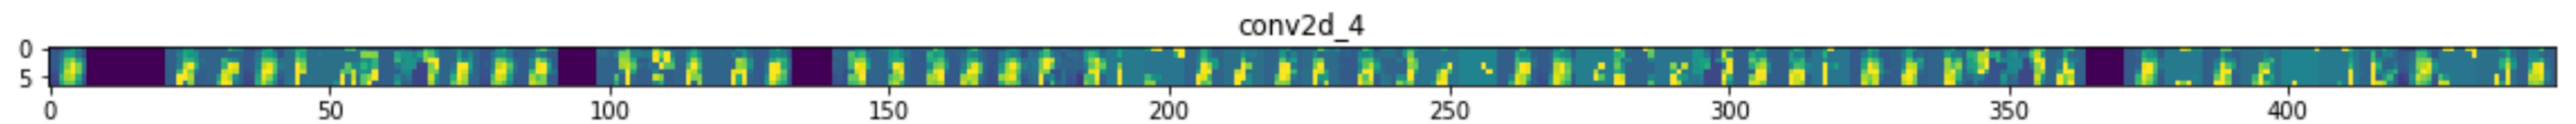
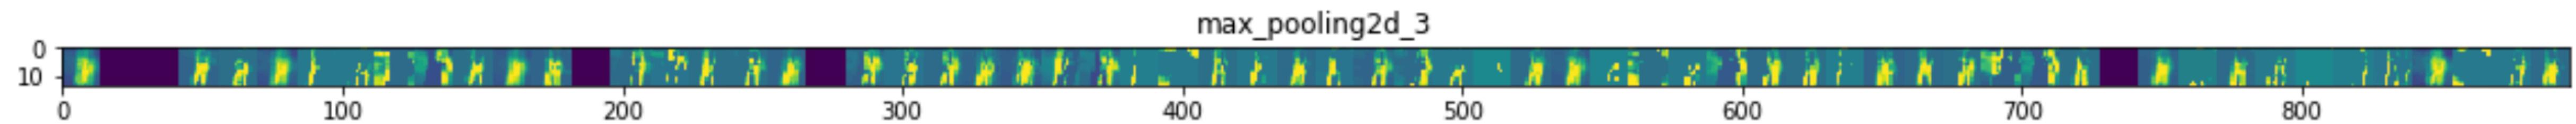
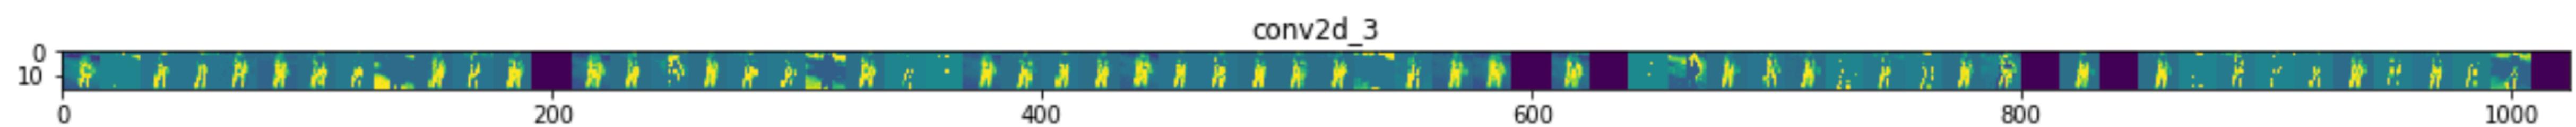
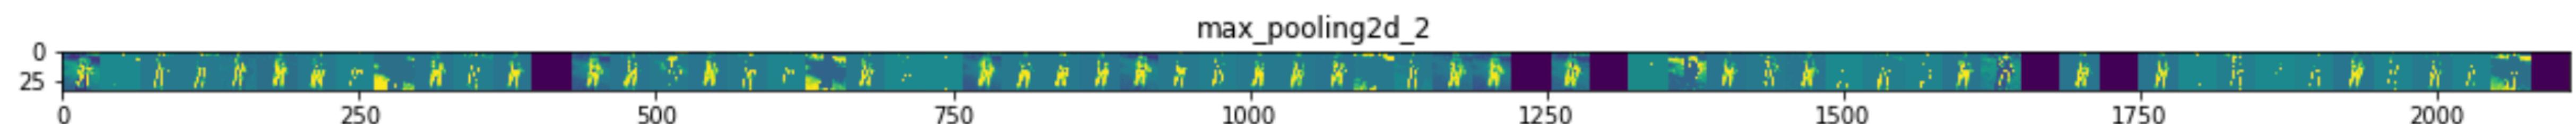
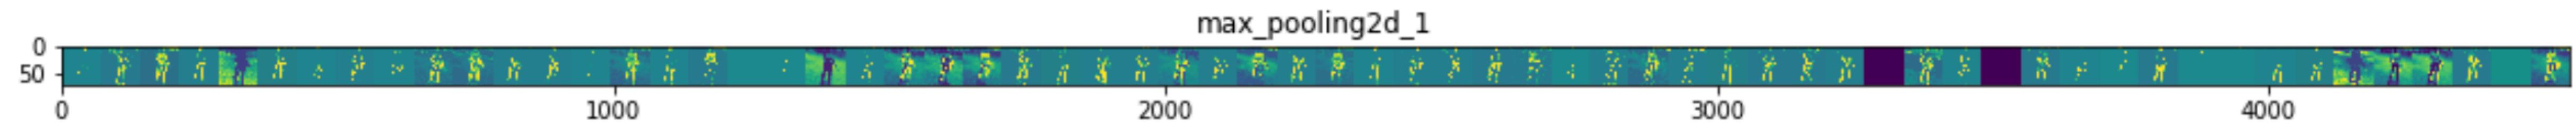
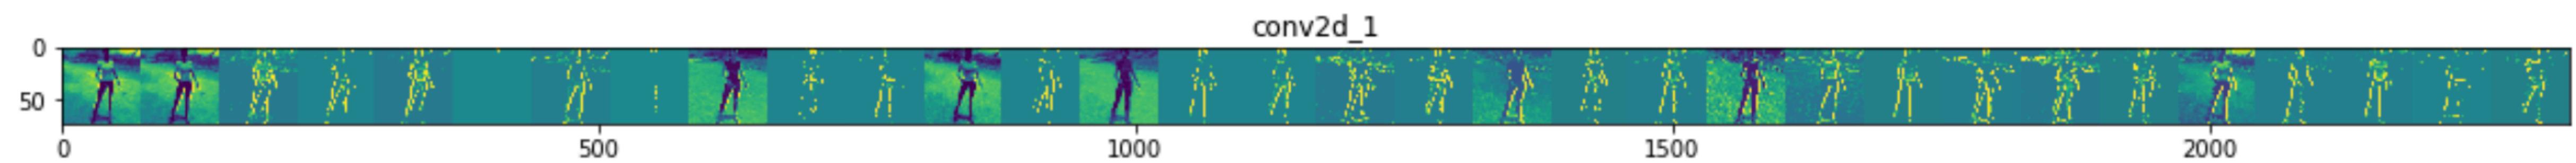
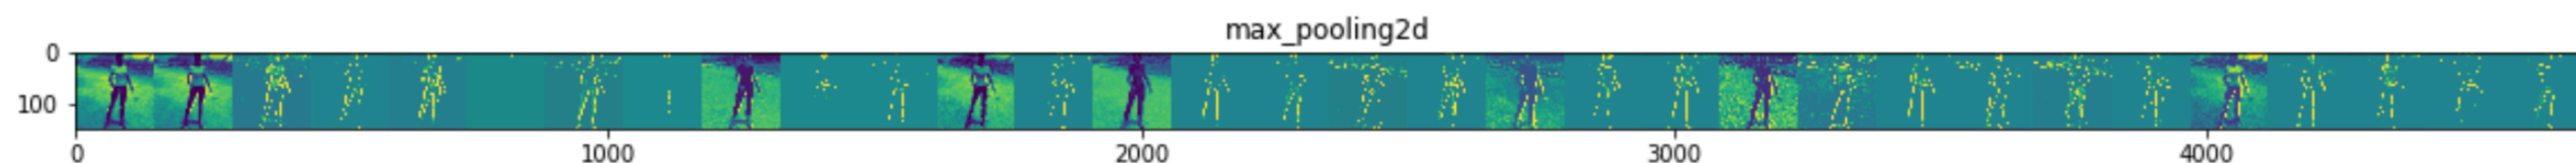
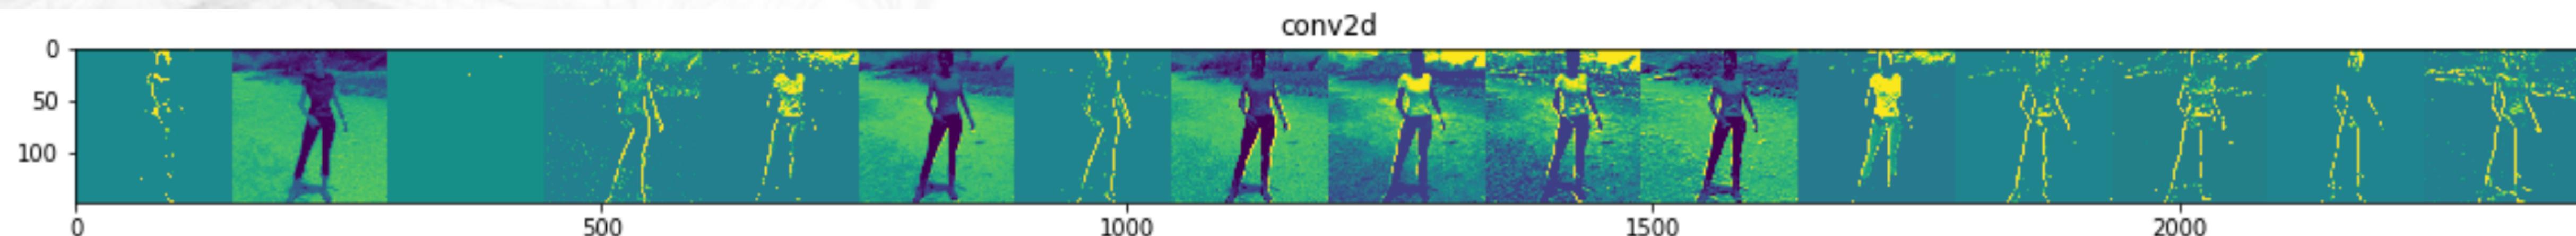
# Max Pooling

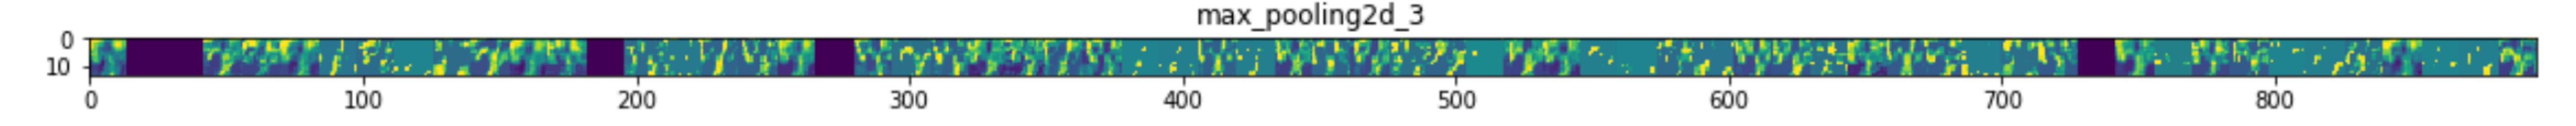
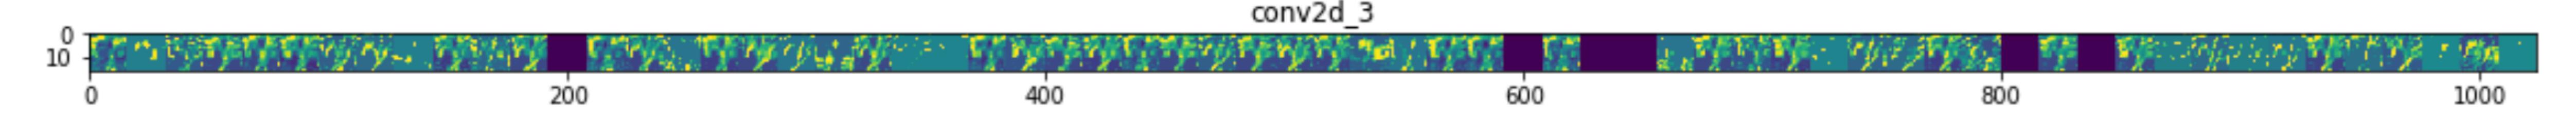
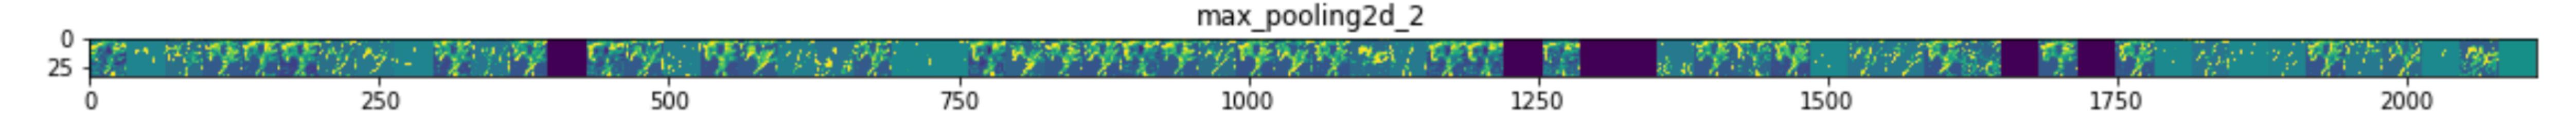
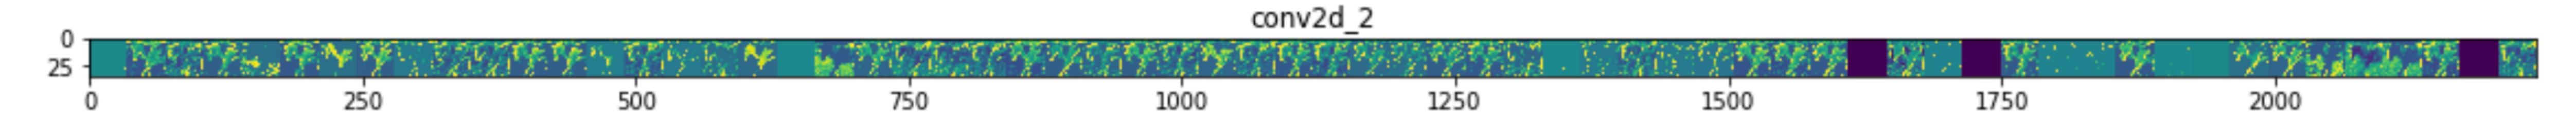
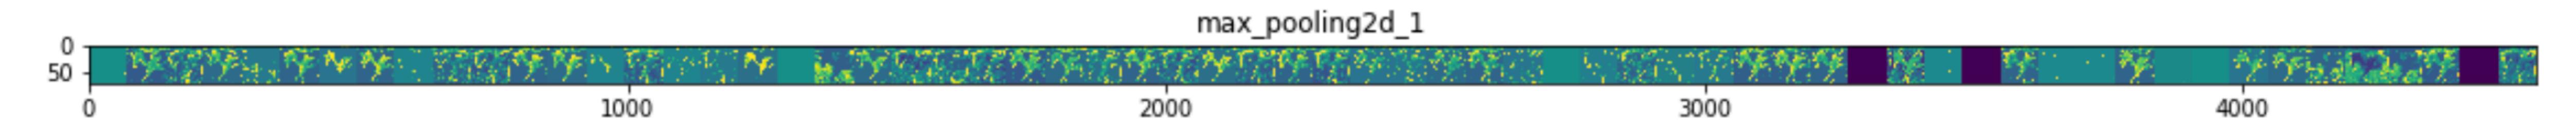
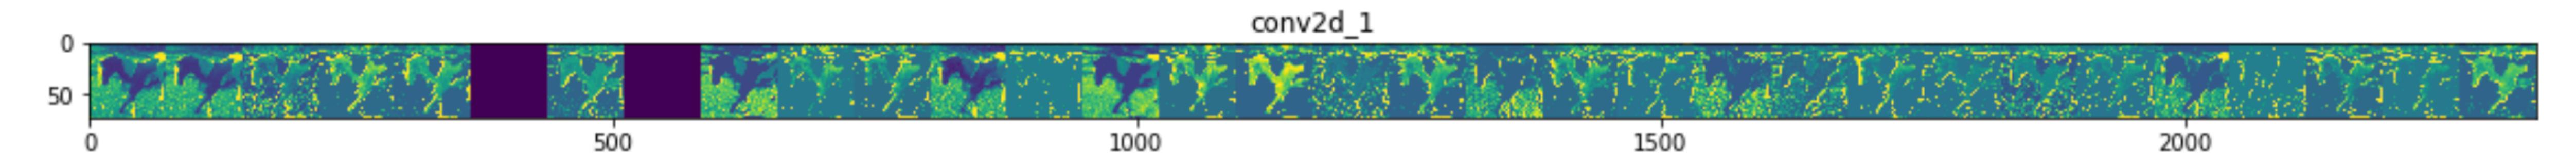
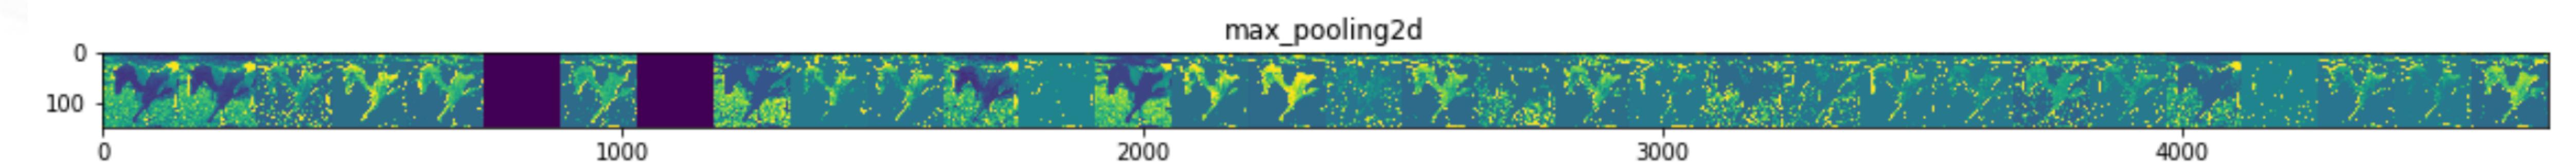
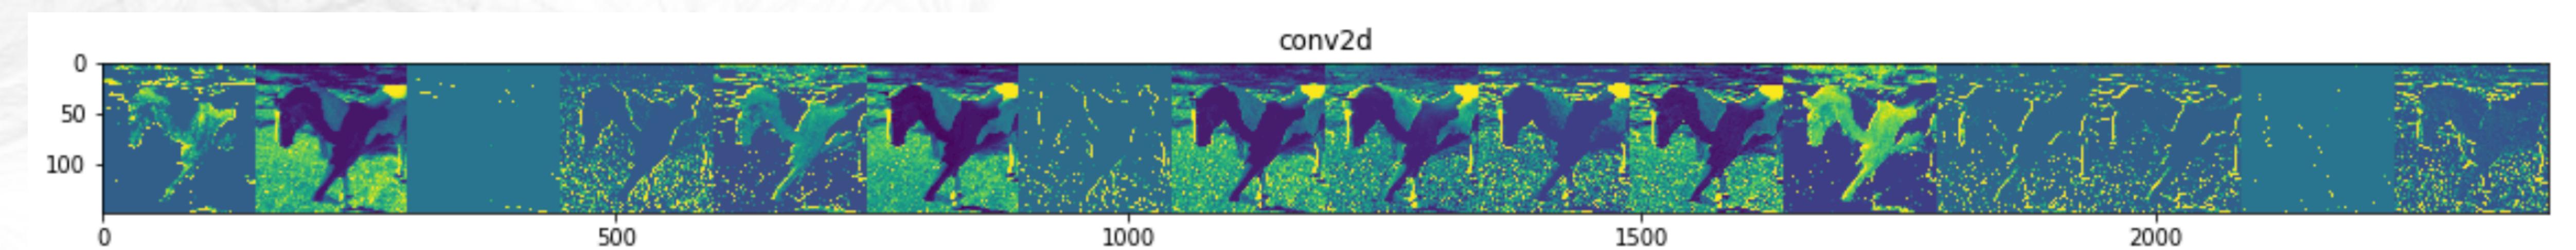


# AlexNet

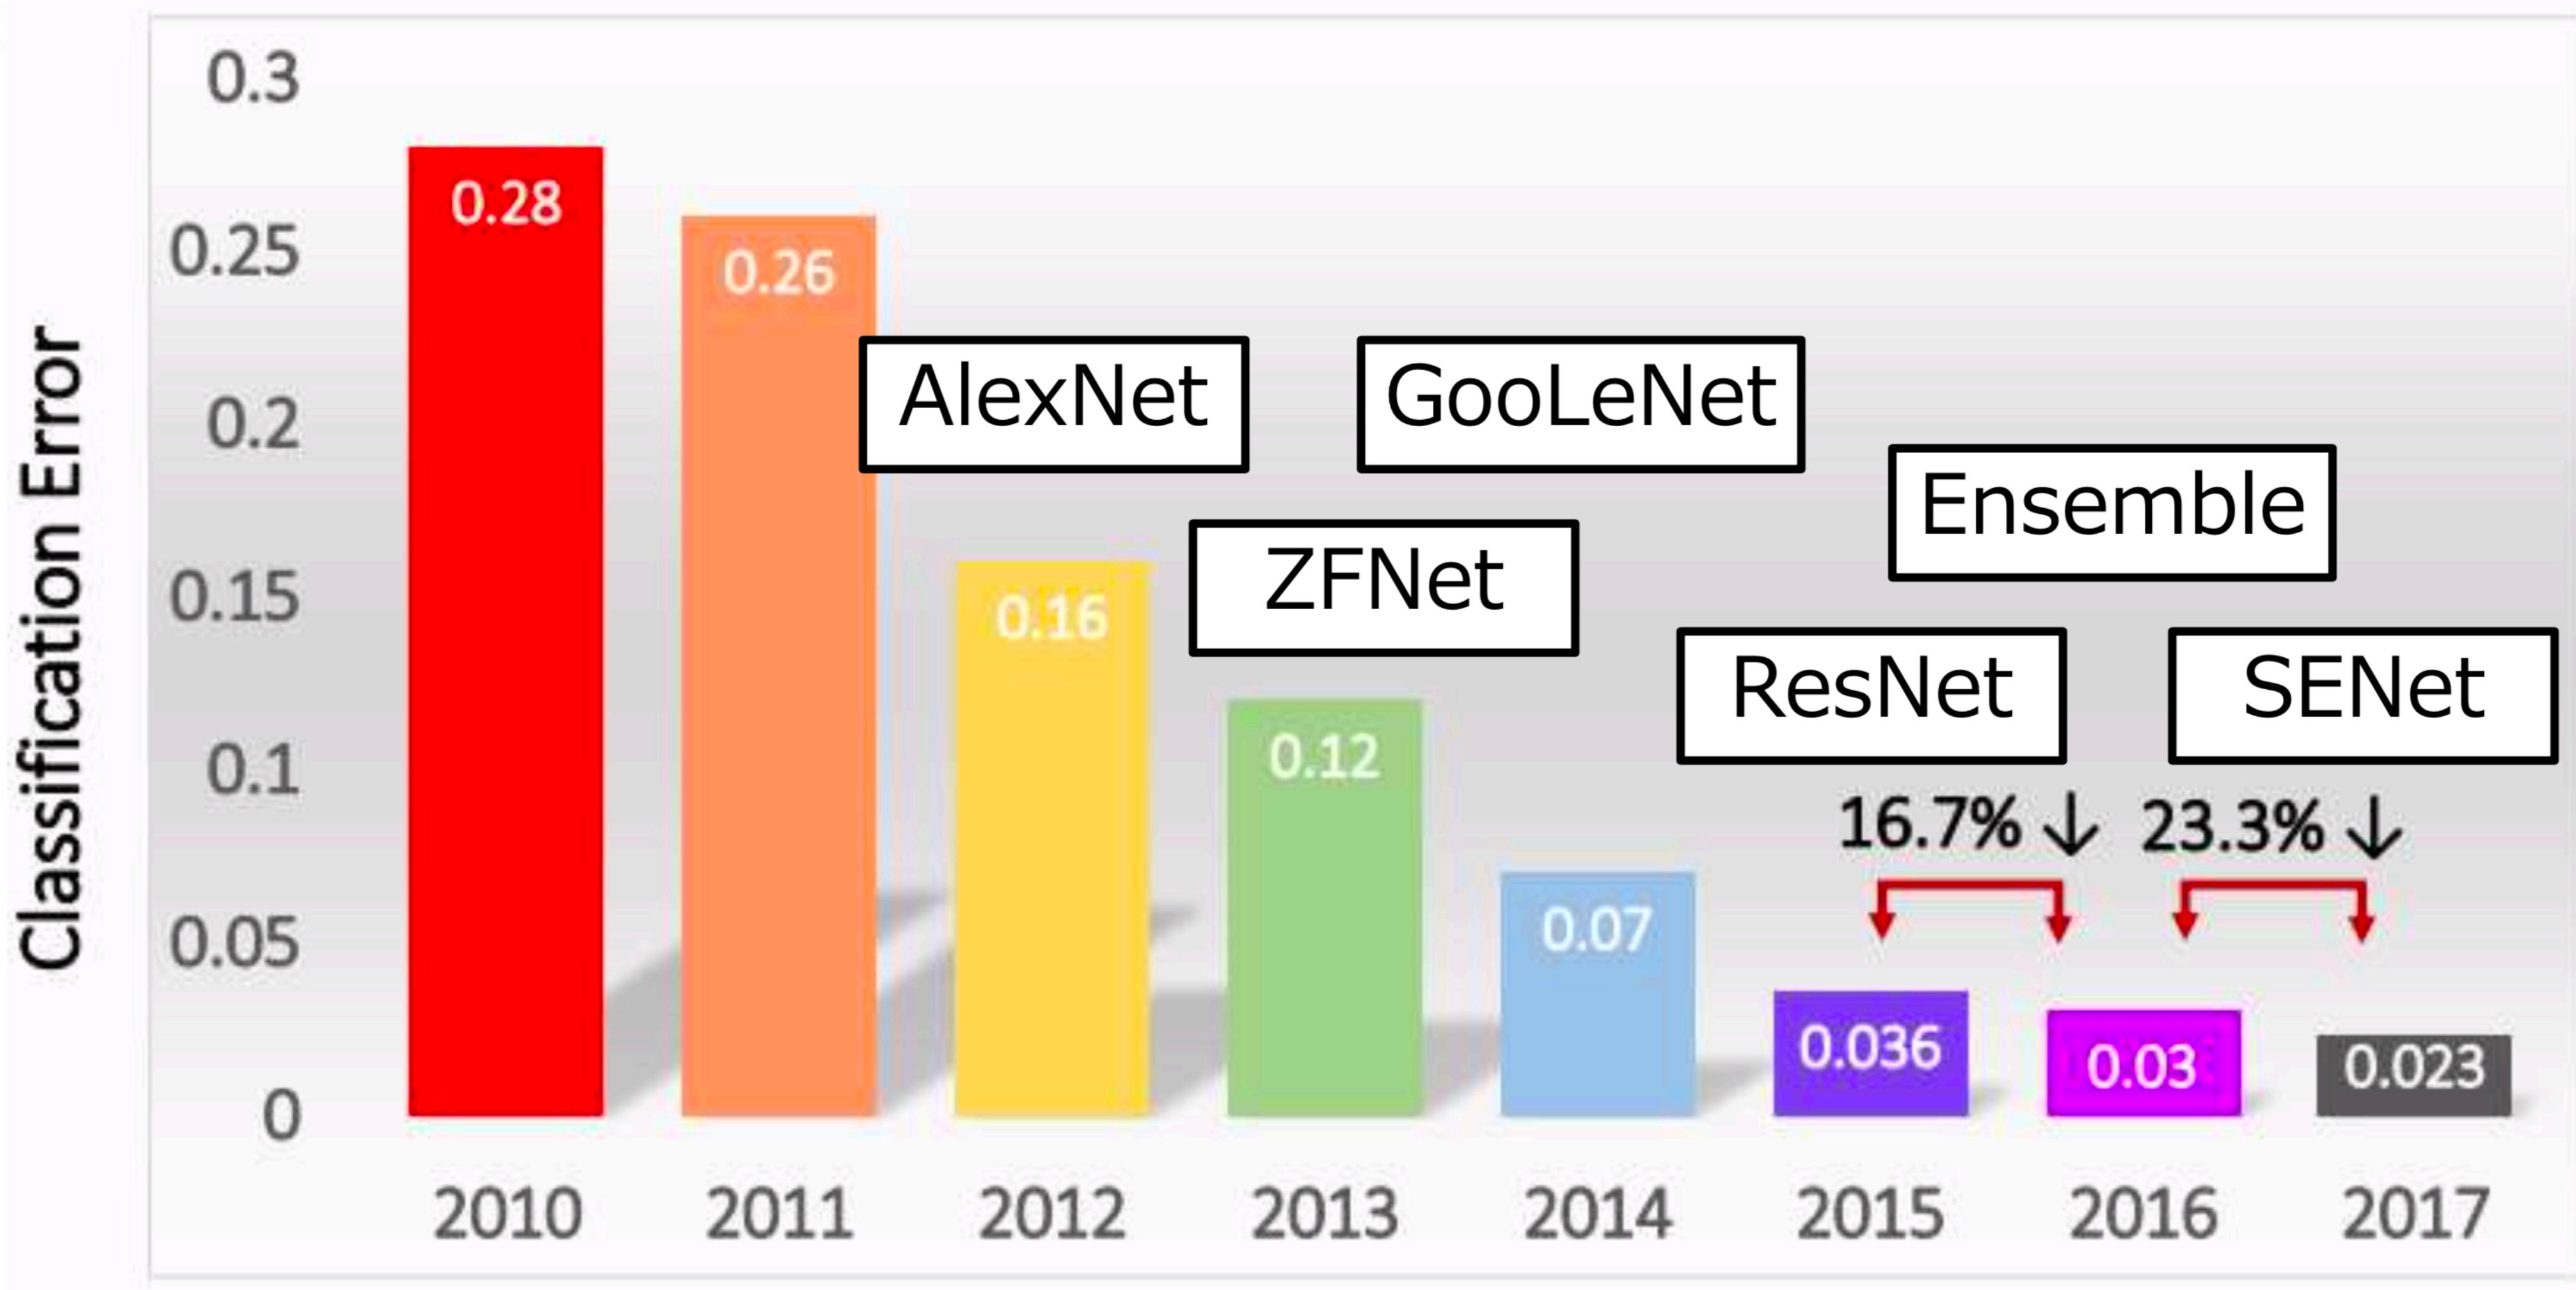








# ImageNet コンペの優勝モデル





# ResNet



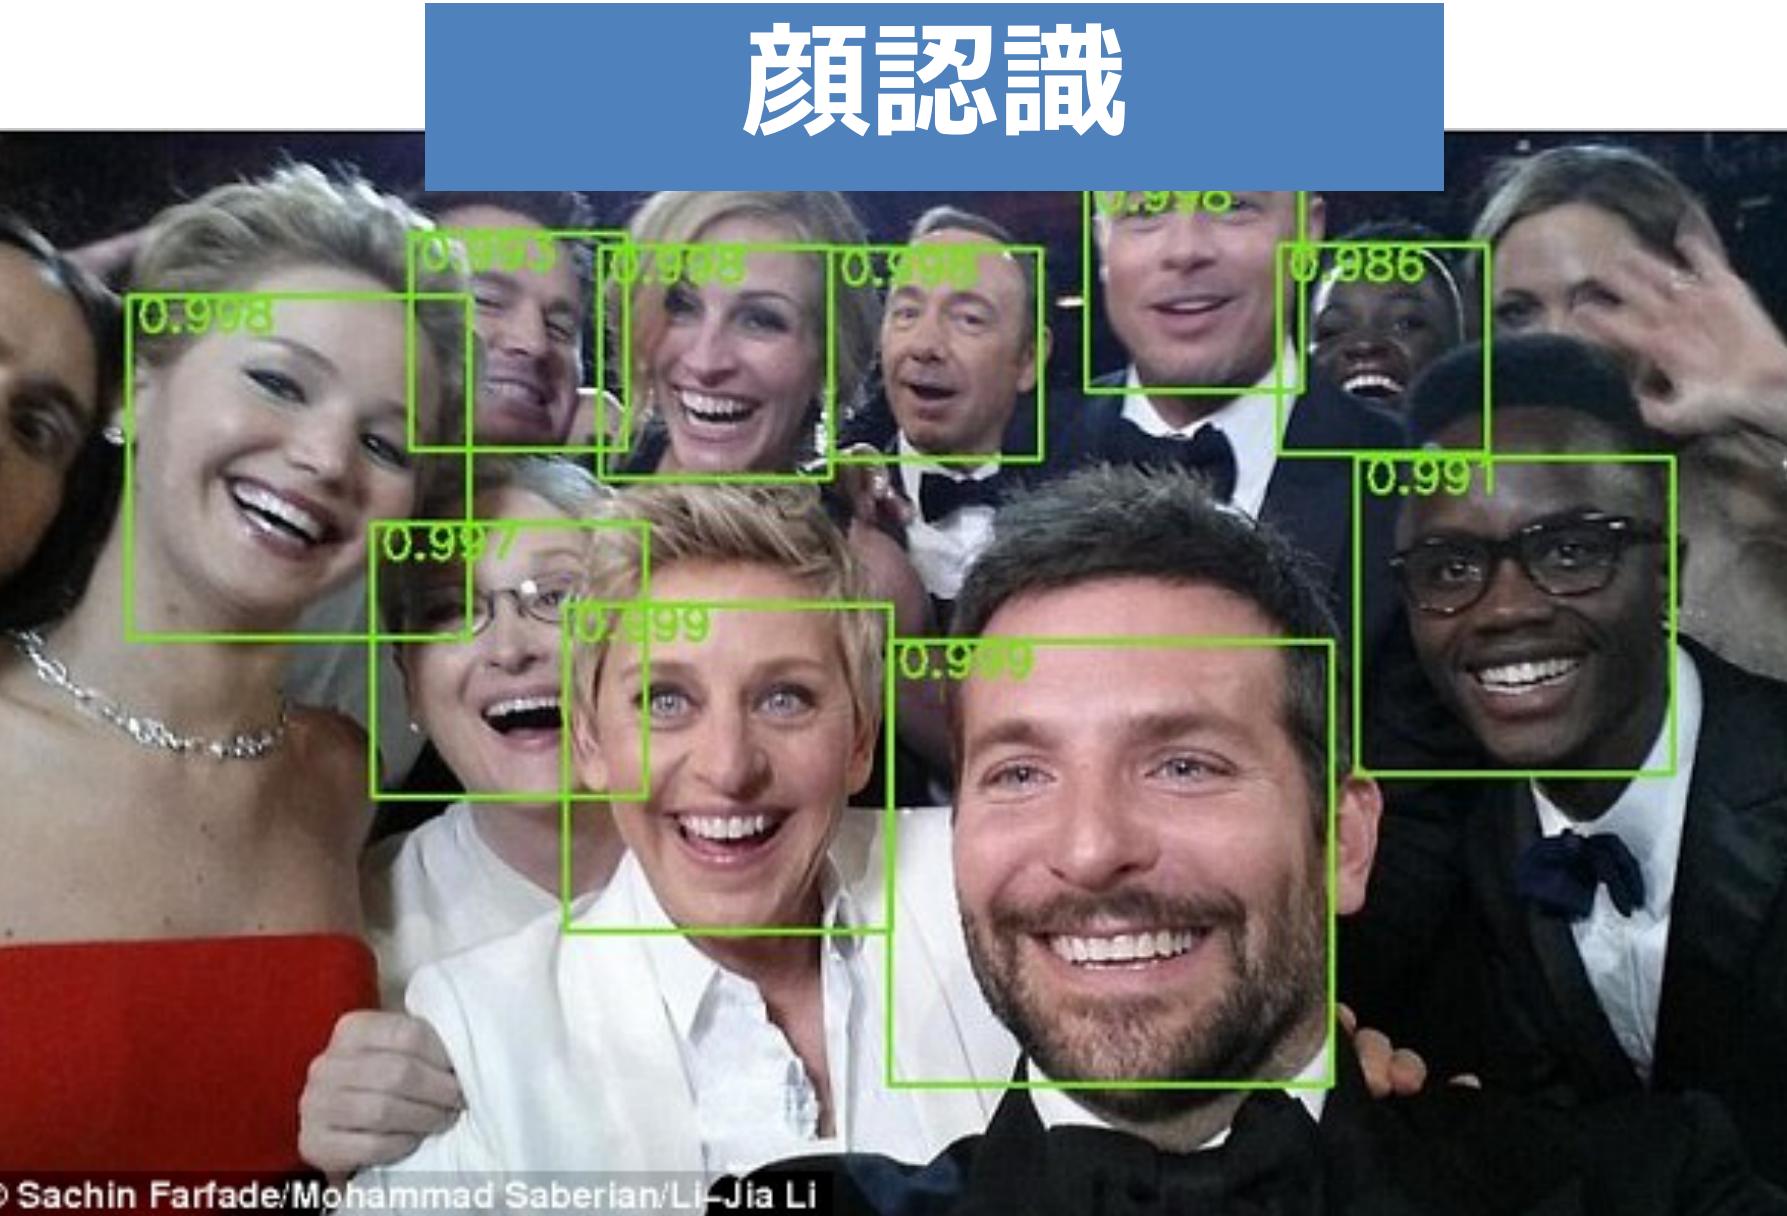
# YOLO

# skymind | 画像認識のアプリケーション

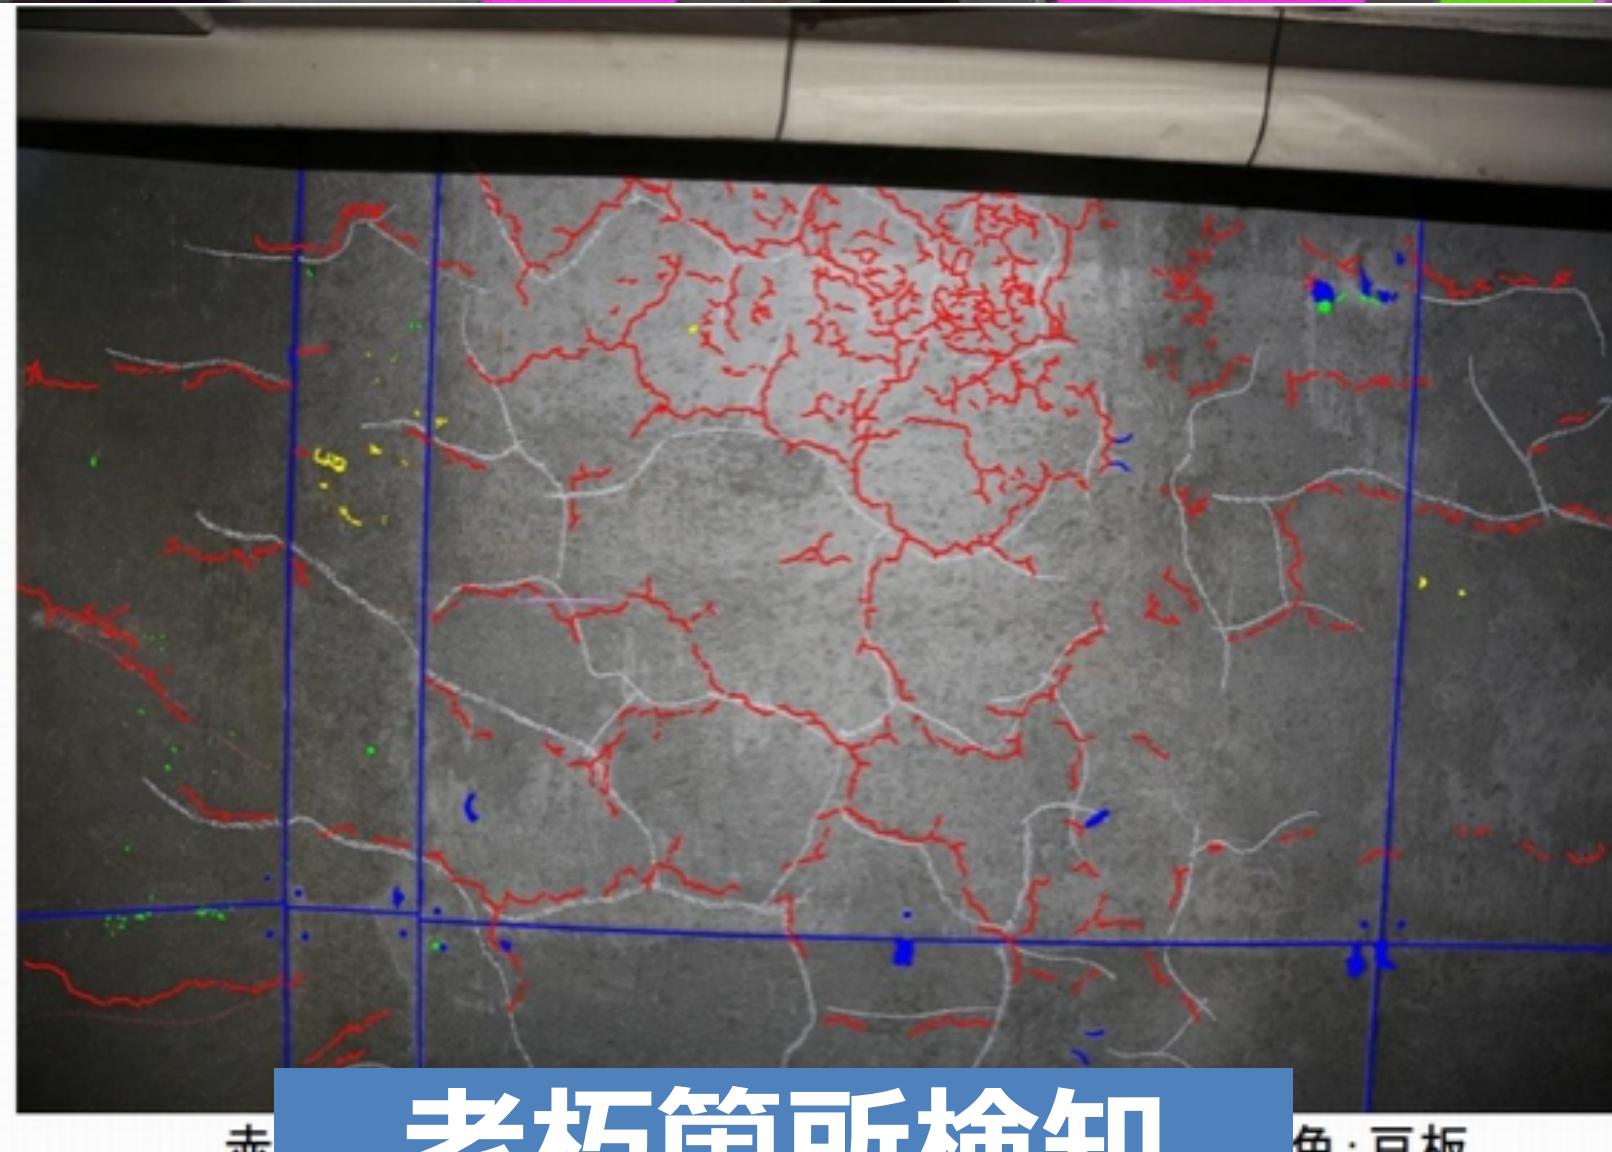
## 一般物体認識



## 顔認識



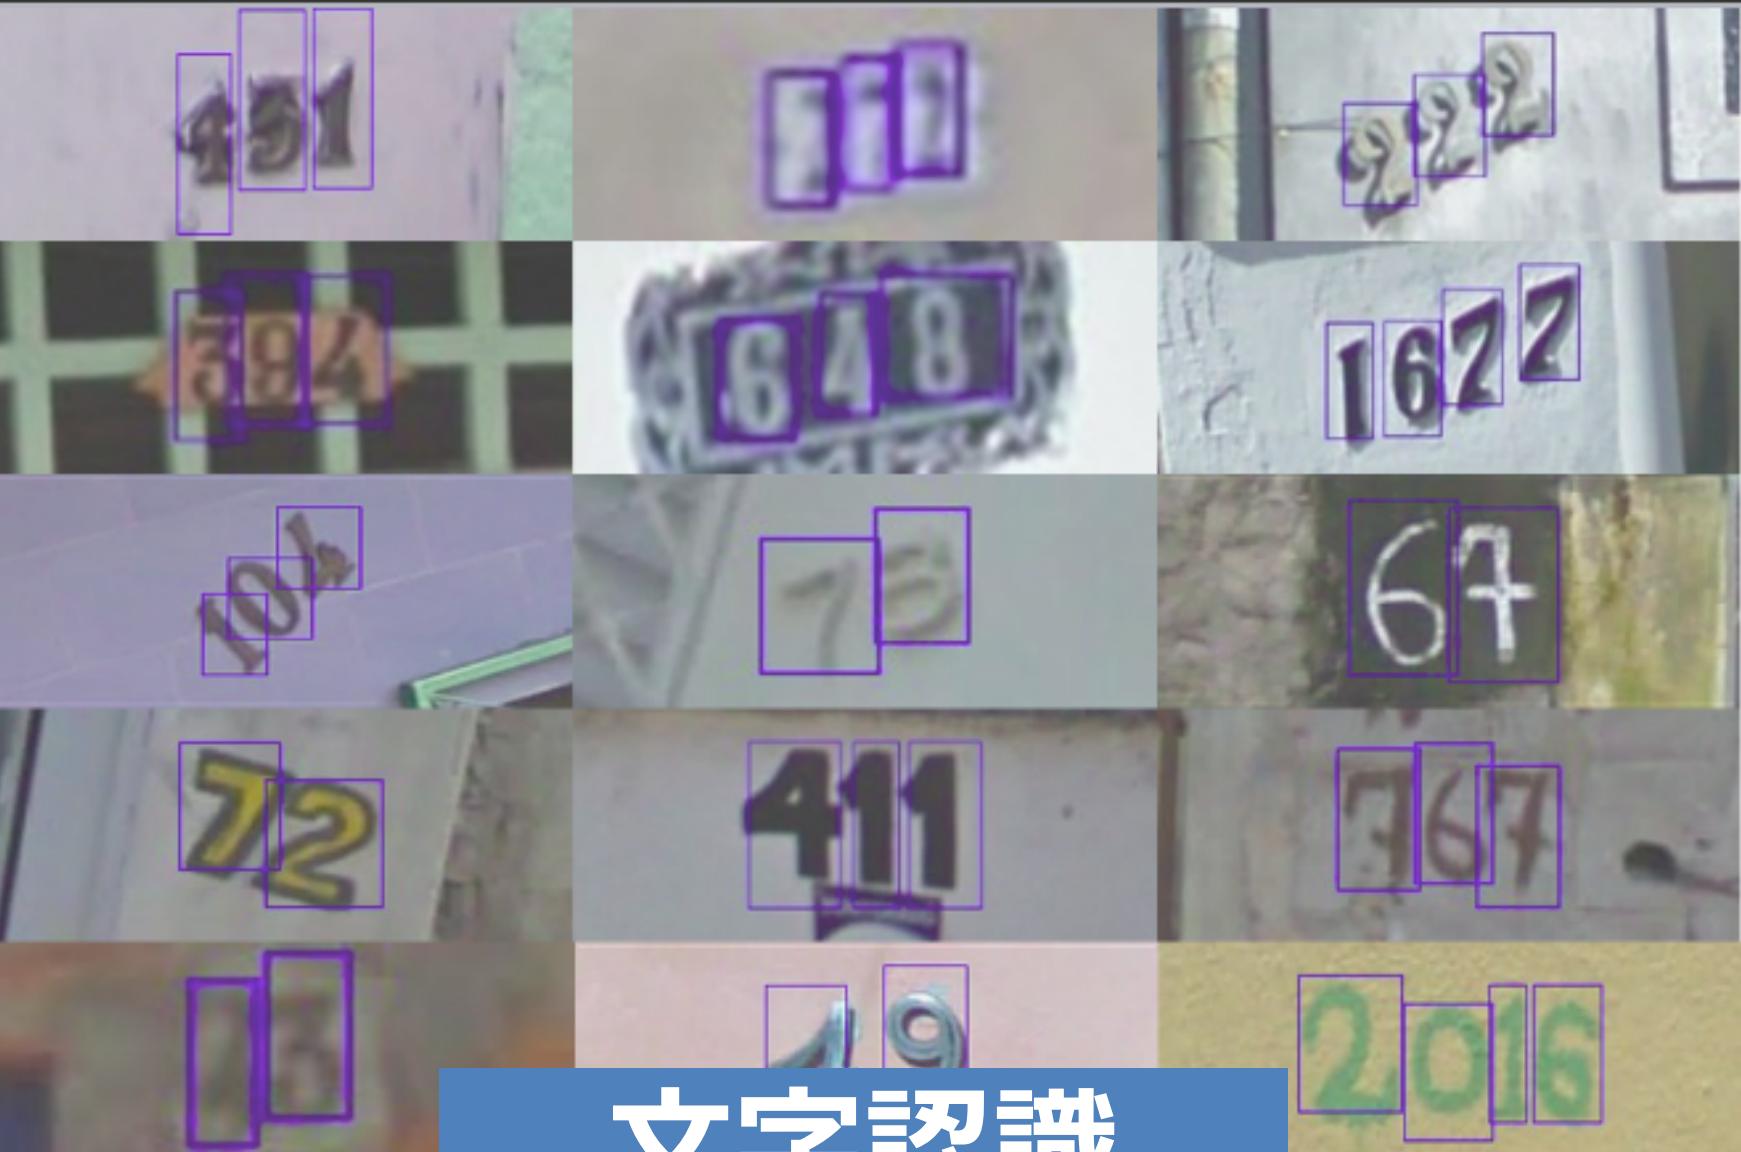
## 年齢推定



## 老朽箇所検知



## セグメンテーション



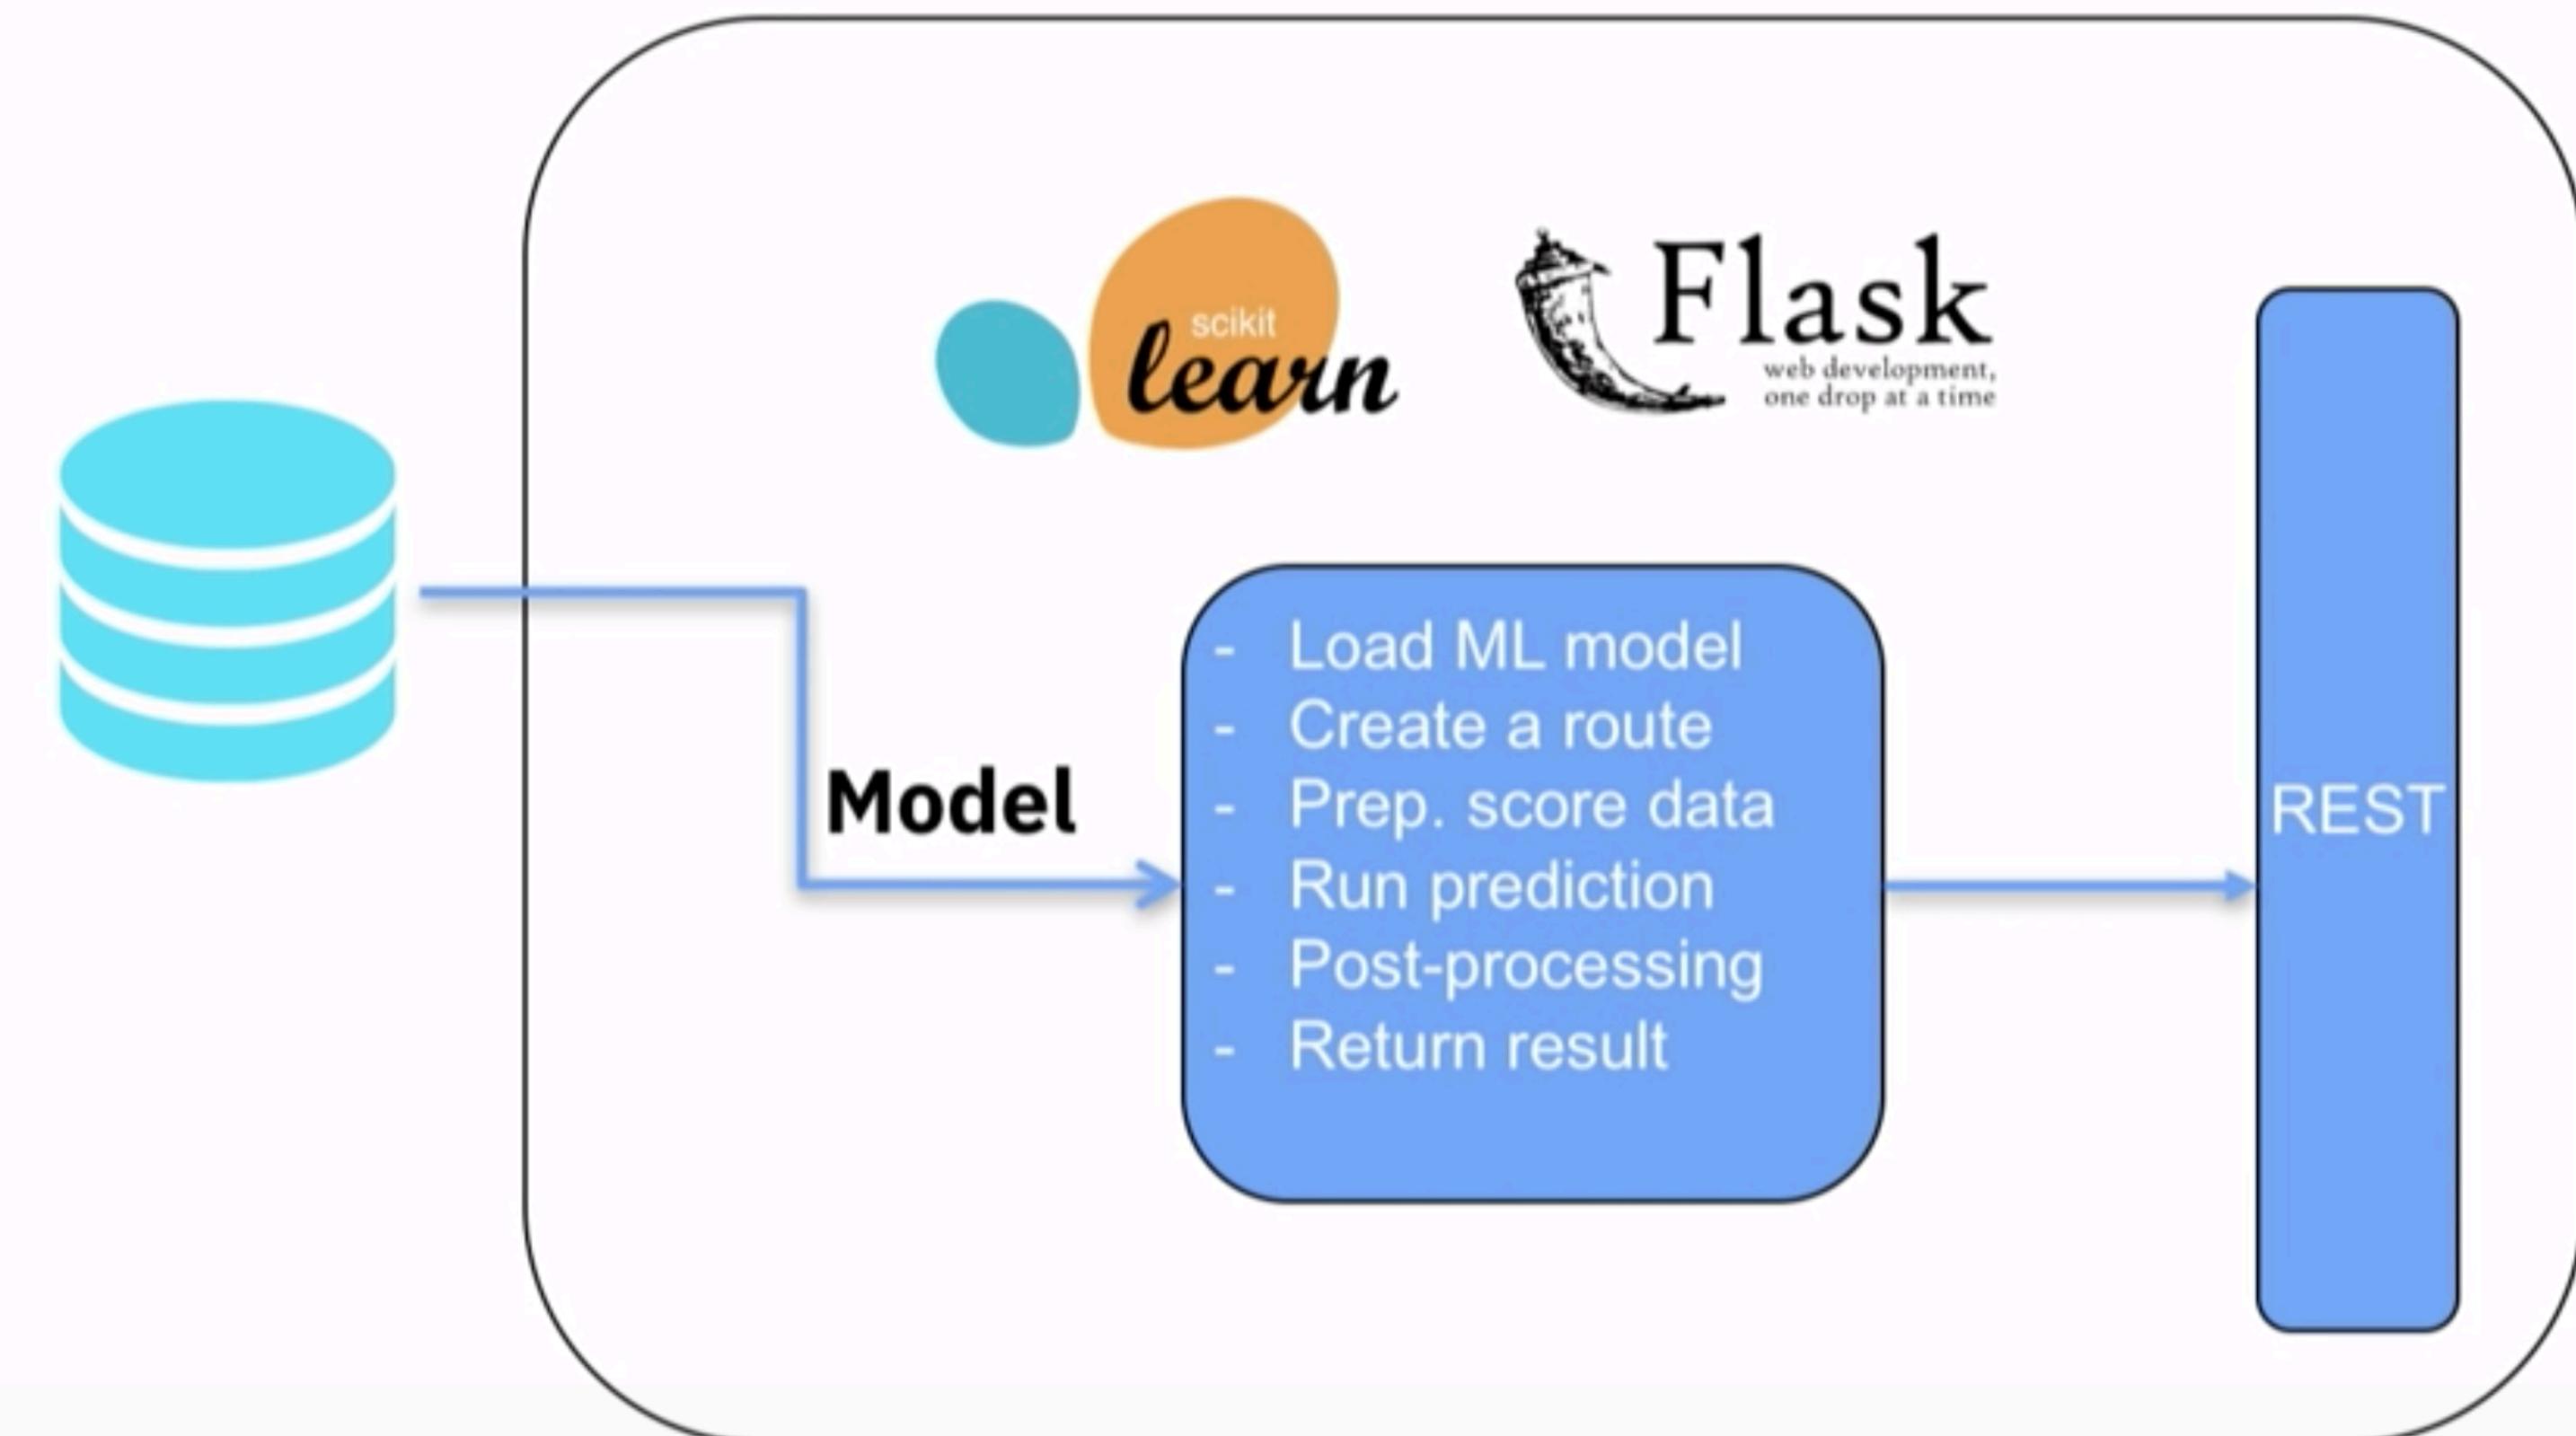
## 文字認識

# Deployment with Cloud Foundry

Python App - REST API  
Endpoint: /v1/xyz/prediction : POST

## Questions

- How to configure the app ?
- How many instances ?
- How much RAM ?
- Zero downtime deployment ?
- Connection to DB?
- Logging ?
- Crashes and recovery?
- Routing ?
- Workload scaling ?
- Integration with other tools?
- ...



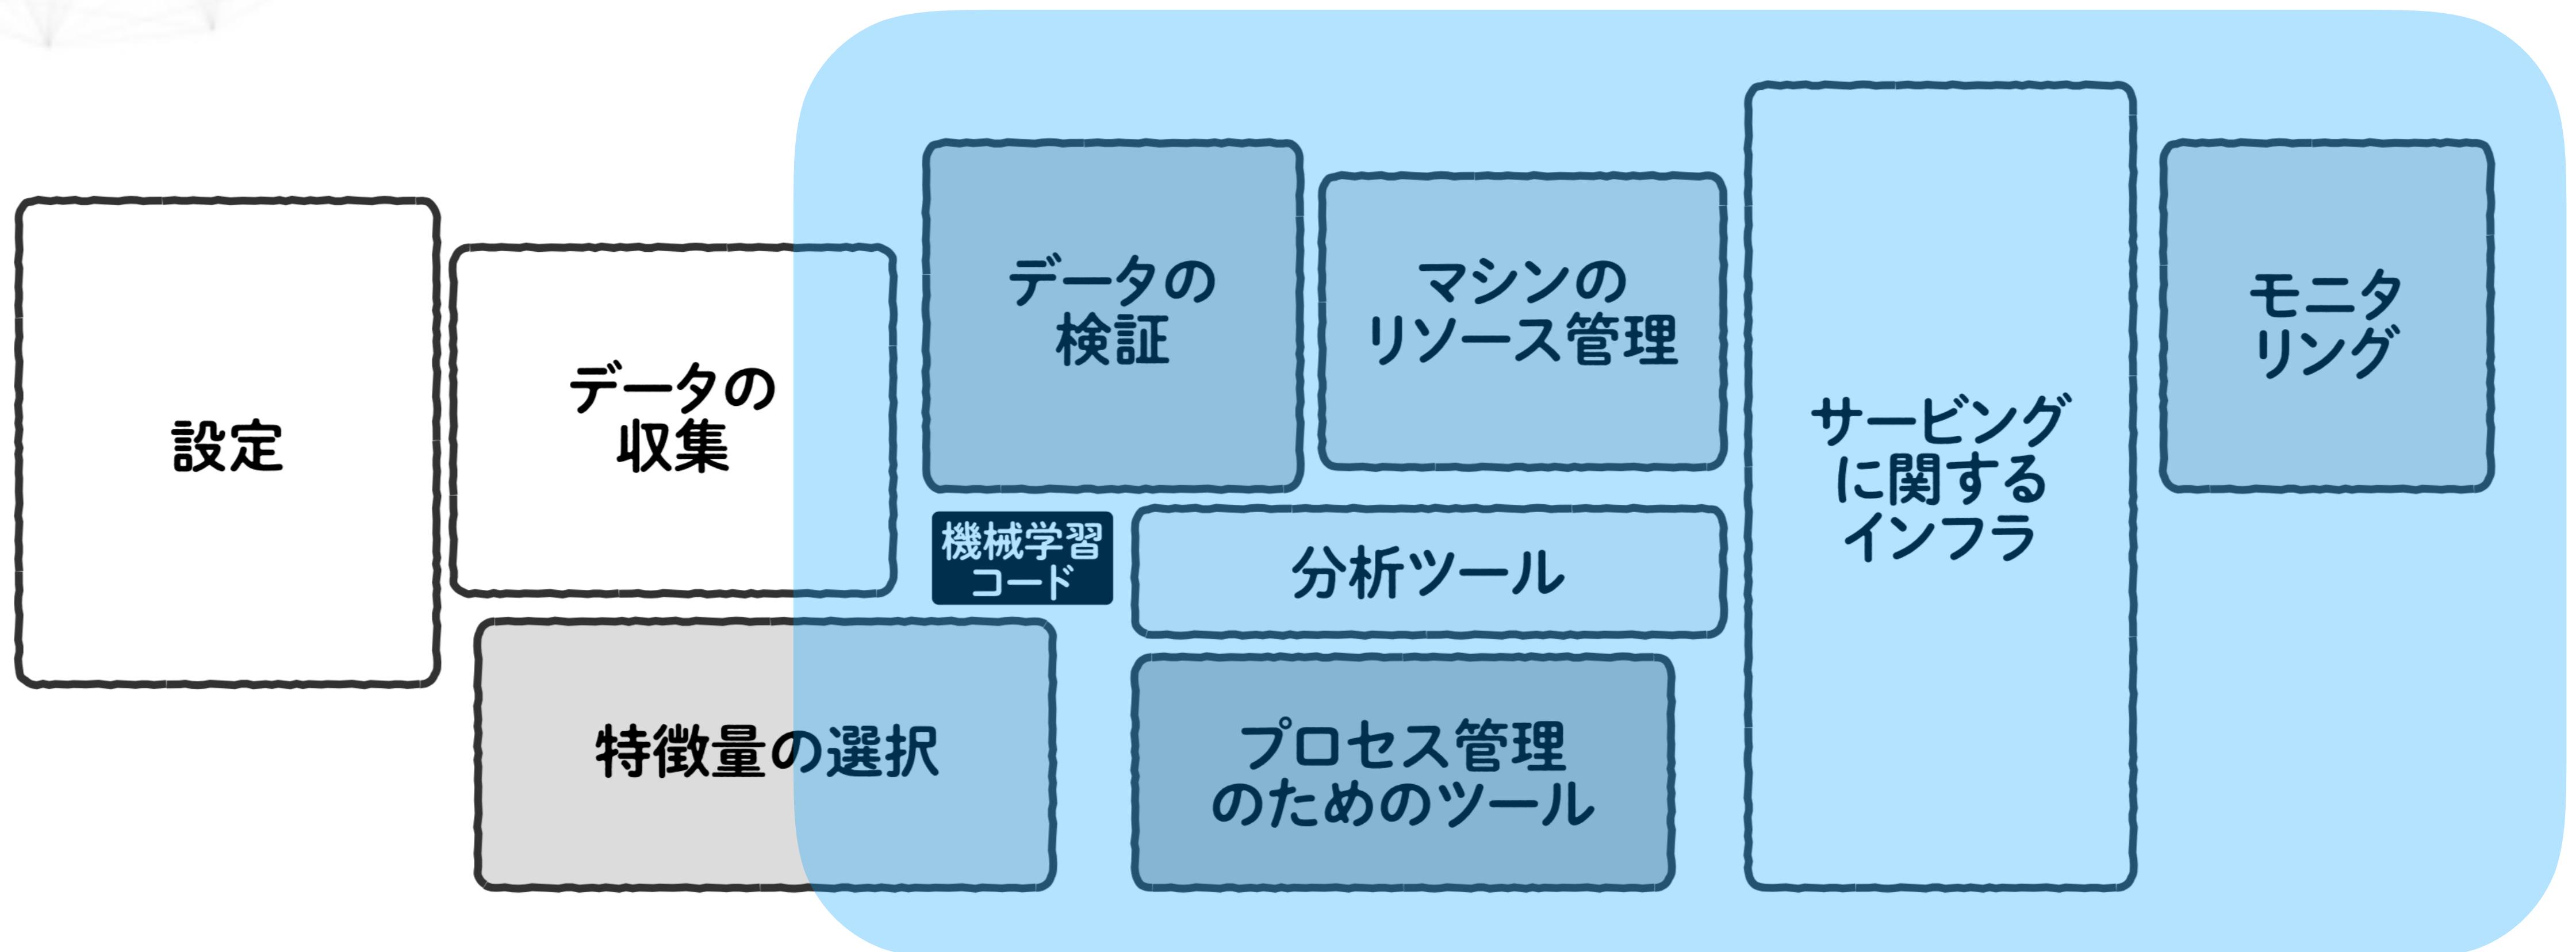
# モデルの利用



▶ 図 3.1 機械学習のモデリングのコードの割合は少ない ([2] 中の図を翻訳して引用)

n月間ラムダノート Vol.1, No.1 より

# モデルの利用

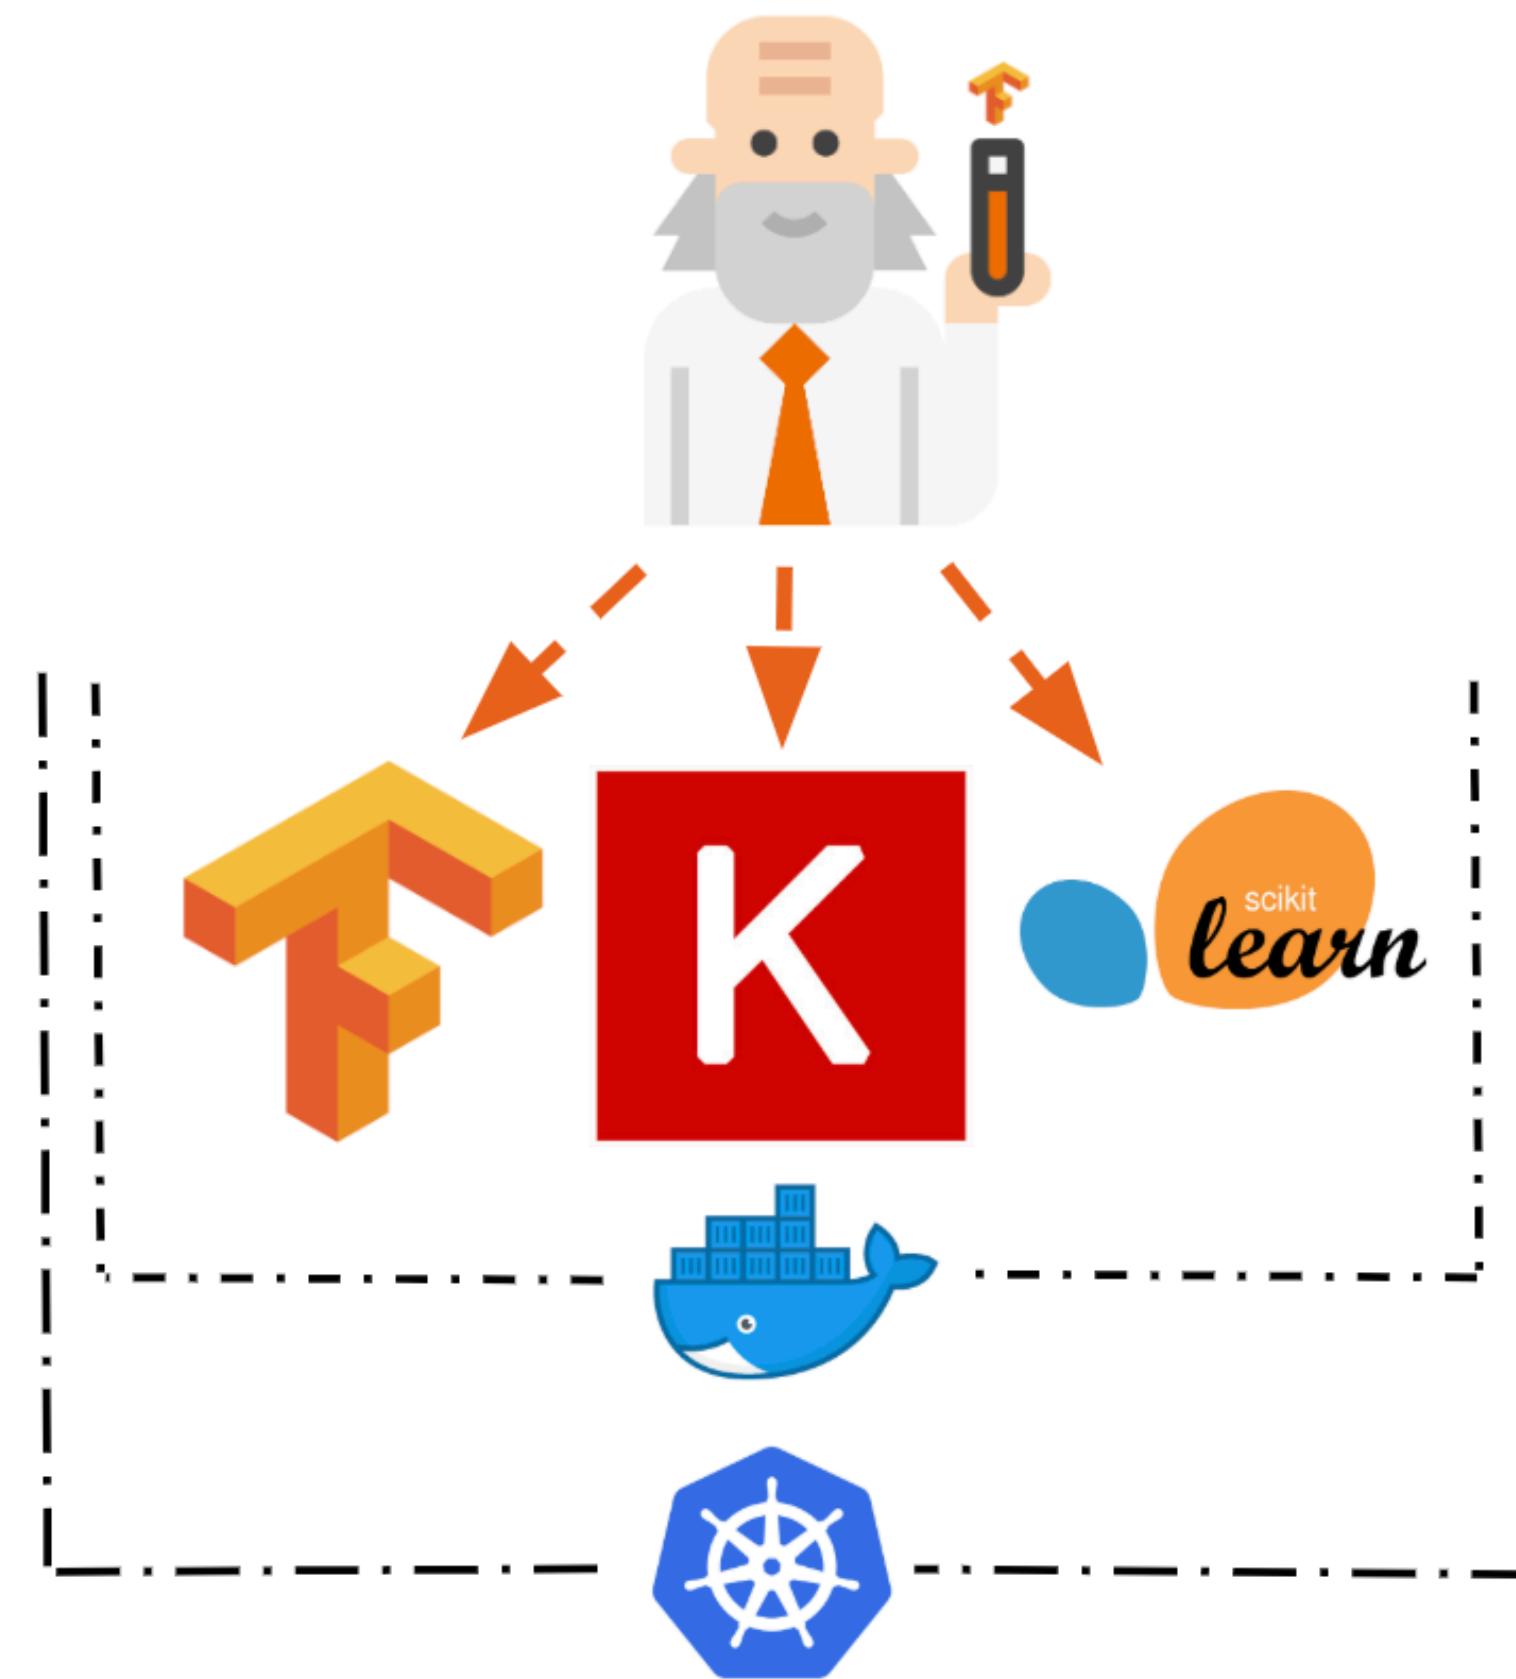


▶ 図 3.1 機械学習のモデリングのコードの割合は少ない ([2] 中の図を翻訳して引用)

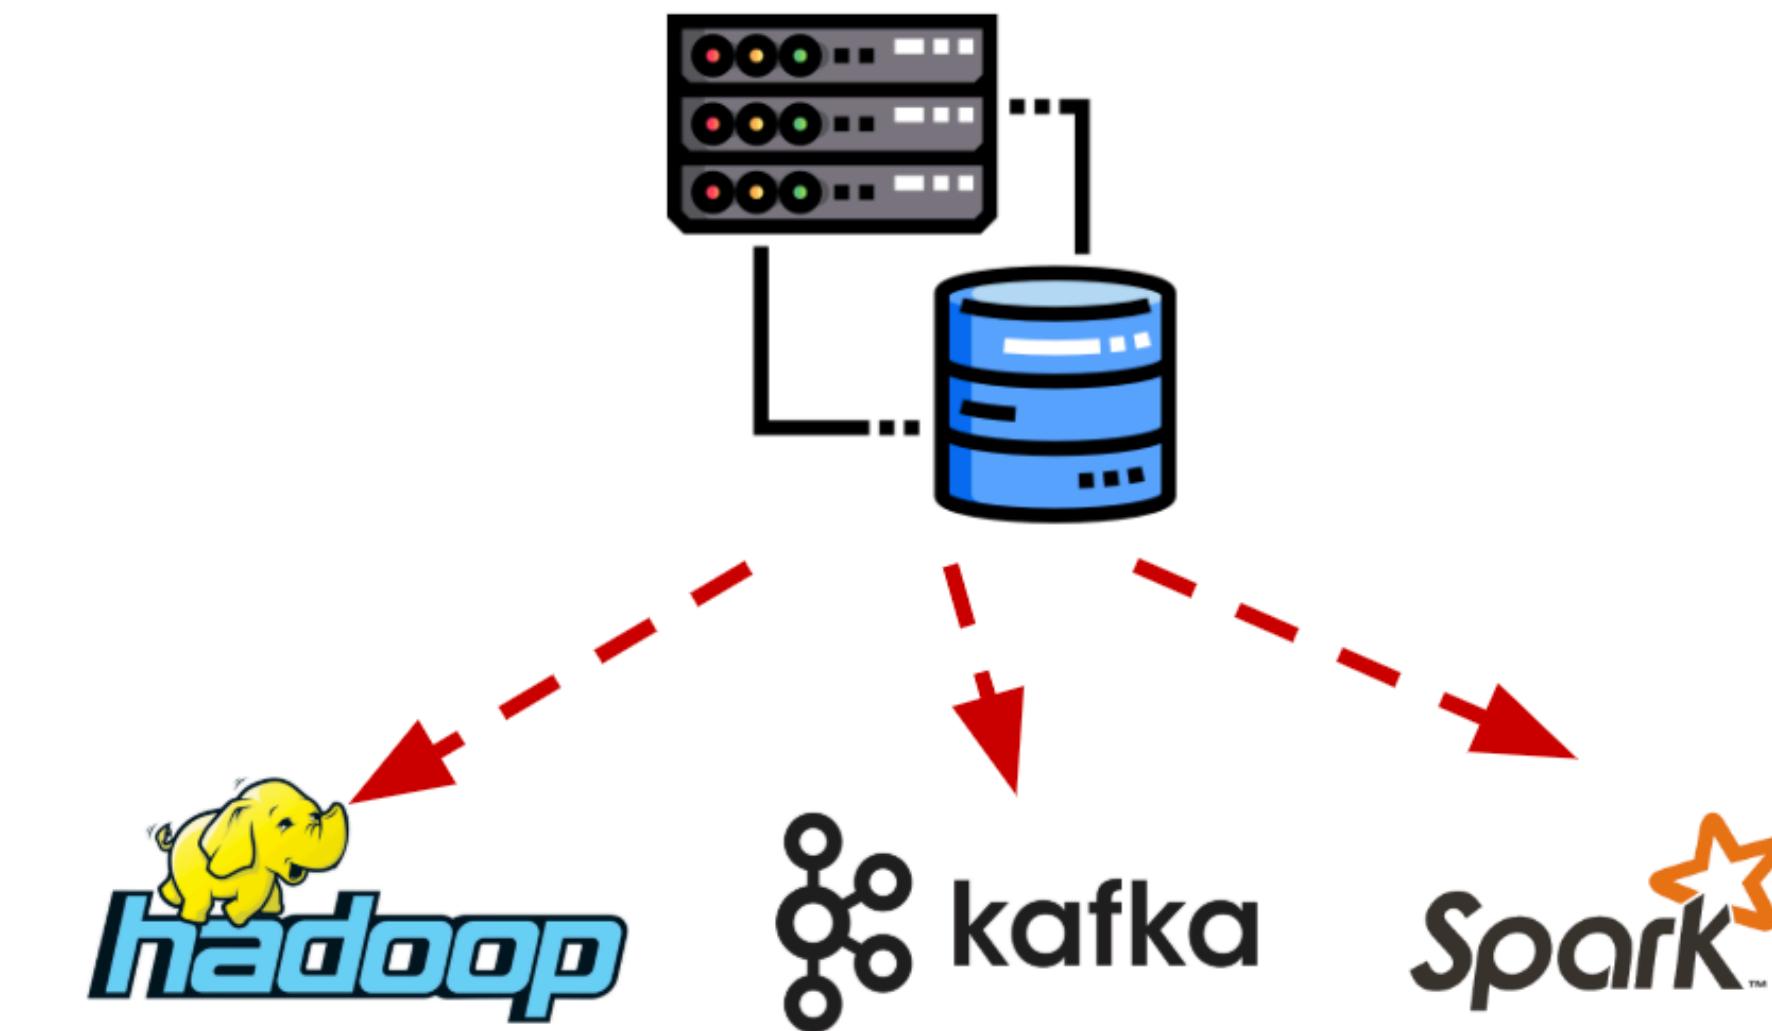
n月間ラムダノート Vol.1, No.1 より

# Most AI Frameworks Not Built for Enterprise

## Development

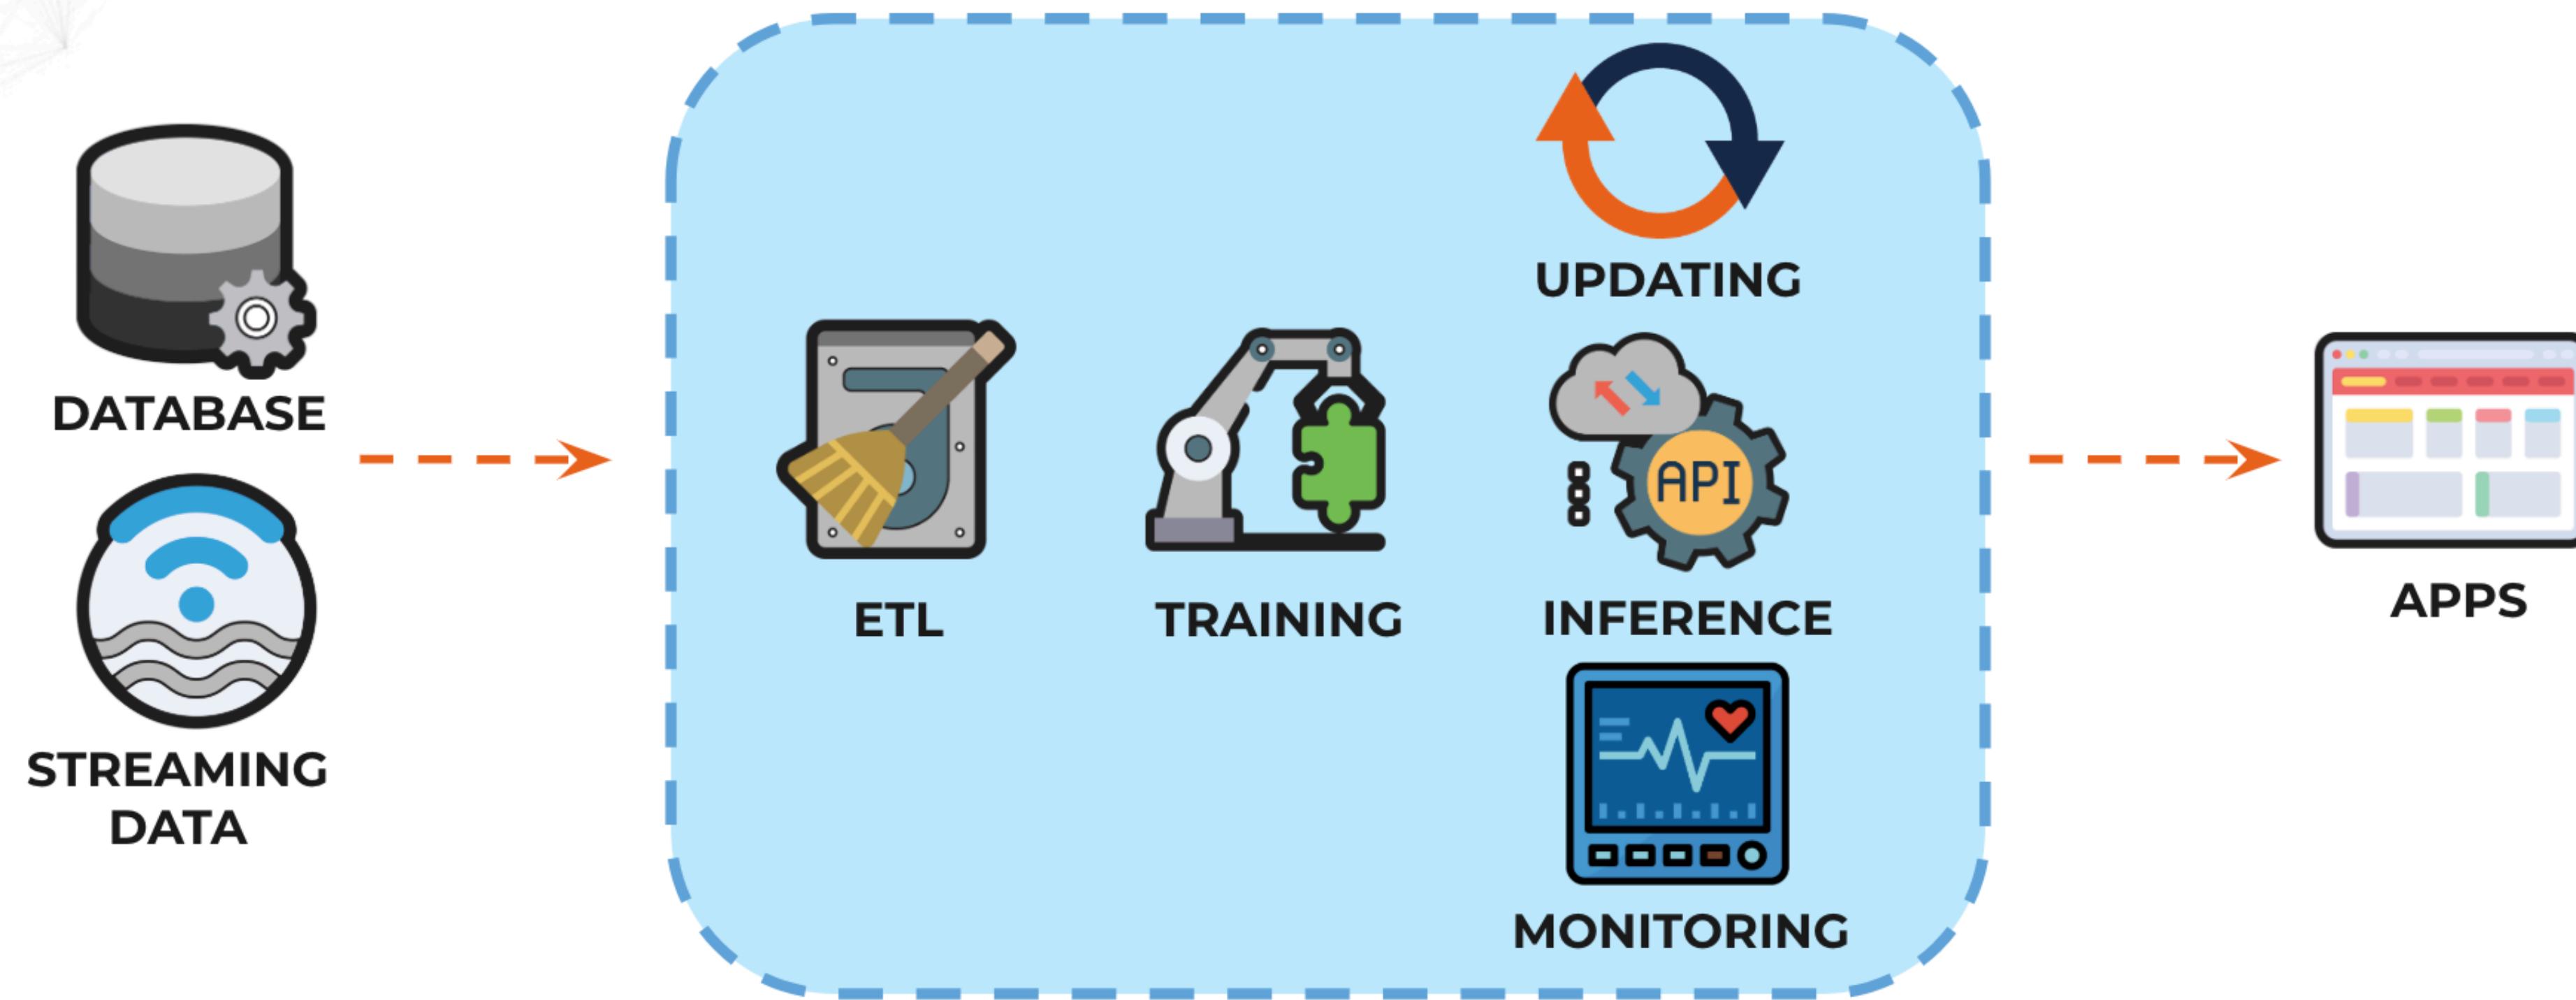


## Production



- Petabytes of data
- Distributed clusters for data processing
- SLAs mandatory
  - Uptime guarantees!

# SKIL provides a solution



Standardized AI infrastructure lets teams unify practices.

Configurable clusters: Adapts to opinionated ops teams and messy environments

A managed software distribution for popular ML/AI frameworks

Ingests AI models from any ML vendor/tool, allowing data scientists to keep favorite tools

## Workspaces, Experiments

- **Workspaces** are a shared lab to conduct a set of “experiments” centered around a particular project.
- **Experiments** are different configurations of neural net models and data pipelines applied to a given problem

The screenshot displays the skymind SKIL User Interface. On the left, a sidebar menu lists: WORKSPACES (highlighted with a red box), DEPLOYMENTS, JOBS, CLUSTER NODES, PROCESSES, and PLUGINS. The main area shows two tabs: 'Workspaces' and 'Experiments'. The 'Workspaces' tab is active, showing a table with columns 'NAME' and 'CREATED'. One row is visible: Demo (ID: b84e7216-ff20-4ad0-bf54-c1d07d298d5e) created on Tue, Mar 12, 2019 1:43 PM, 2 minutes ago. The 'Experiments' tab shows a table with columns 'EXPERIMENT' and 'LAST UPDATED'. One row is visible: demo (Last updated: Tue, Mar 12, 2019 1:46 PM, 2 minutes ago). A 'NOTEBOOK LINK' button is next to it, and an 'Open Notebook' button is at the bottom right. The top navigation bar shows 'Workspaces / Demo'.

NAME	CREATED
Demo b84e7216-ff20-4ad0-bf54-c1d07d298d5e	Tue, Mar 12, 2019 1:43 PM 2 minutes ago

EXPERIMENT	LAST UPDATED
demo	Tue, Mar 12, 2019 1:46 PM 2 minutes ago

## Deployments

- When mark a model as “deployed”, it serve as **REST API endpoint** to use by client applications
- Endpoint example: *http://[host]:9008/endpoints/deployment/model/kerasmodel/default*

The screenshot shows the skymind SKIL User Interface. On the left, there is a sidebar with the following navigation options:

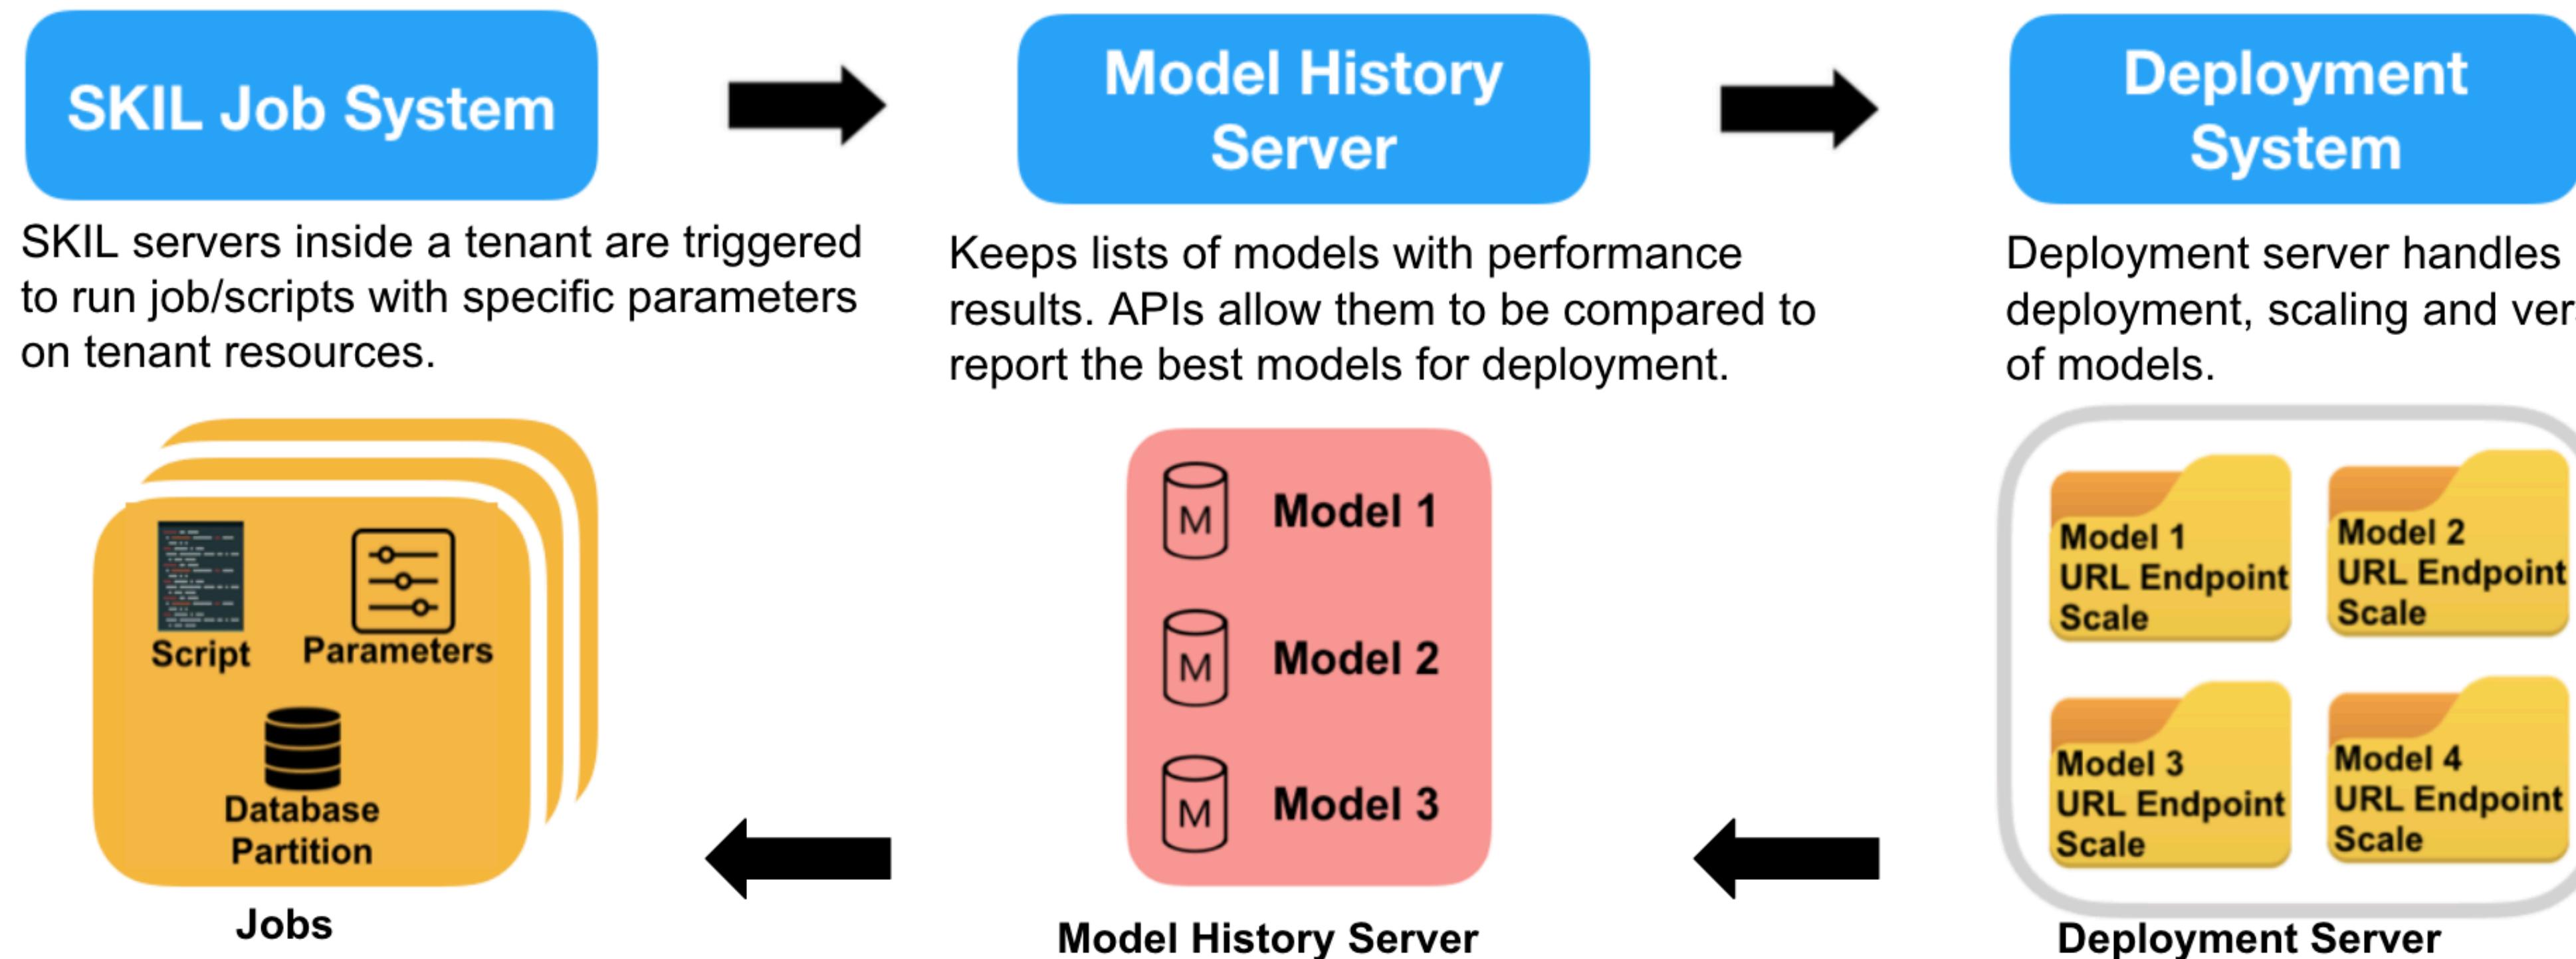
- WORKSPACES
- DEPLOYMENTS** (highlighted with a red border)
- JOBSS
- CLUSTER NODES
- PROCESSES
- PLUGINS

The main content area is titled "Deployments" and displays the message "Currently deployed models, transforms, and KNN". It contains a table with the following data:

NAME	STATUS	DETAILS	ACTIONS
mnist_demo	Not Deployed	1 Model , 0 Transforms, 0 KNN	<button>Edit</button> <button>Delete</button>

On the top right of the main content area, there are three icons: a gear (settings), a question mark (help), and a user icon labeled "Admin".

# skymind | Key Components of Model Management



Best model on real data can  
Be used with transfer learning  
To fine tune model with latest  
data.

Real-time feedback requests  
are stored back in DB to monitor  
model performance on real data stream



# YOLOのデプロイ



# まとめ