

畳み込みニューラル ネットワークの基礎と応用

実践者向けディープラーニング勉強会 第二回

<https://dl4-practitioners.connpass.com/event/124498/>

2019-04-17 (Wed) @ SB C&S

Agenda

19:00 ~ 19:05 opening

19:05 ~ 19:20 機械学習の導入

19:20 ~ 19:30 **walkthrough** Kerasを用いたコーディングサンプル

19:30 ~ 19:45 置み込みニューラルネットワークの基礎

19:45 ~ 20:00 **walkthrough** CNNのフィルタの理解

20:00 ~ 20:10 休憩

20:10 ~ 20:25 **breakout** CNNアーキテクチャの紹介

20:25 ~ 20:40 **Skymind Intelligent Layer (SKIL) の解説**

20:40 ~ 20:55 **walkthrough** SKILを用いたモデルのデプロイ

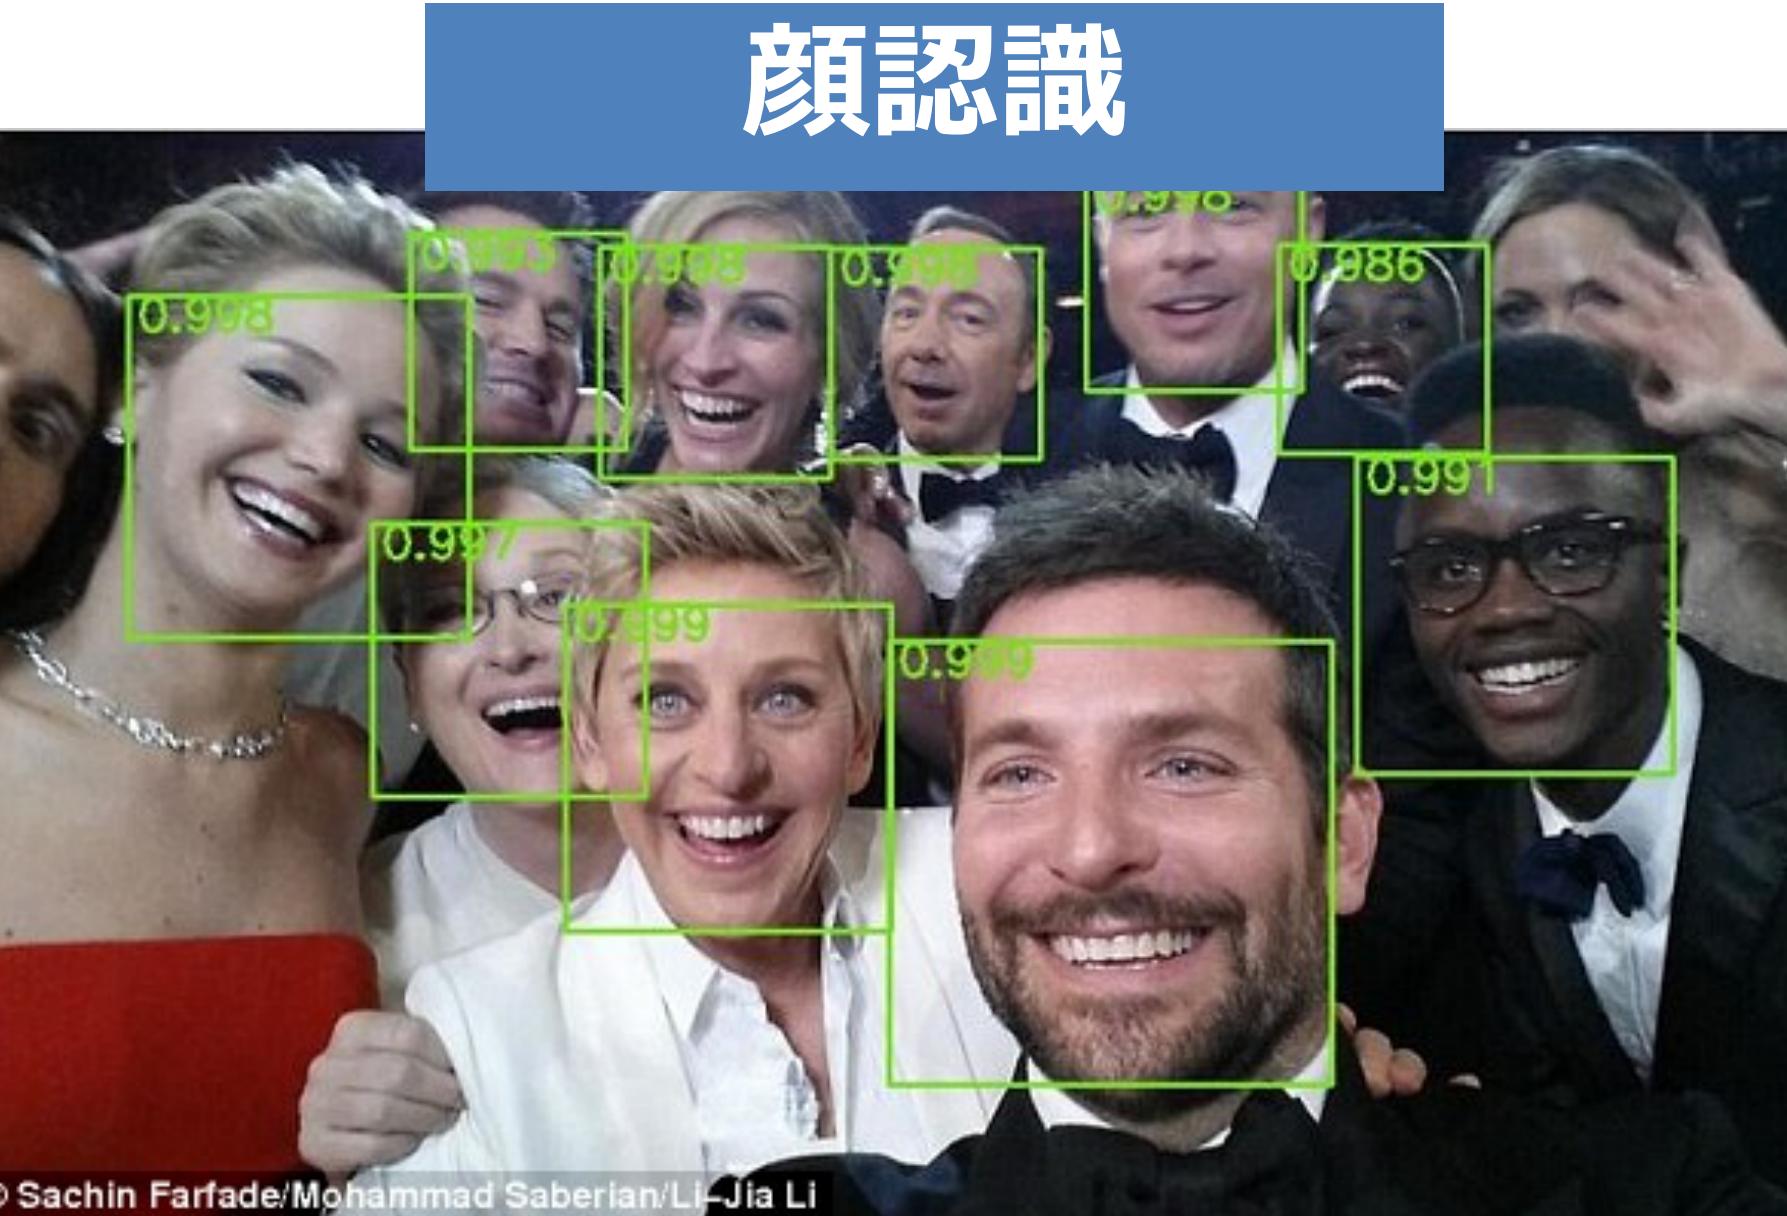
20:55 ~ 21:00 closing

skymind | 画像認識のアプリケーション

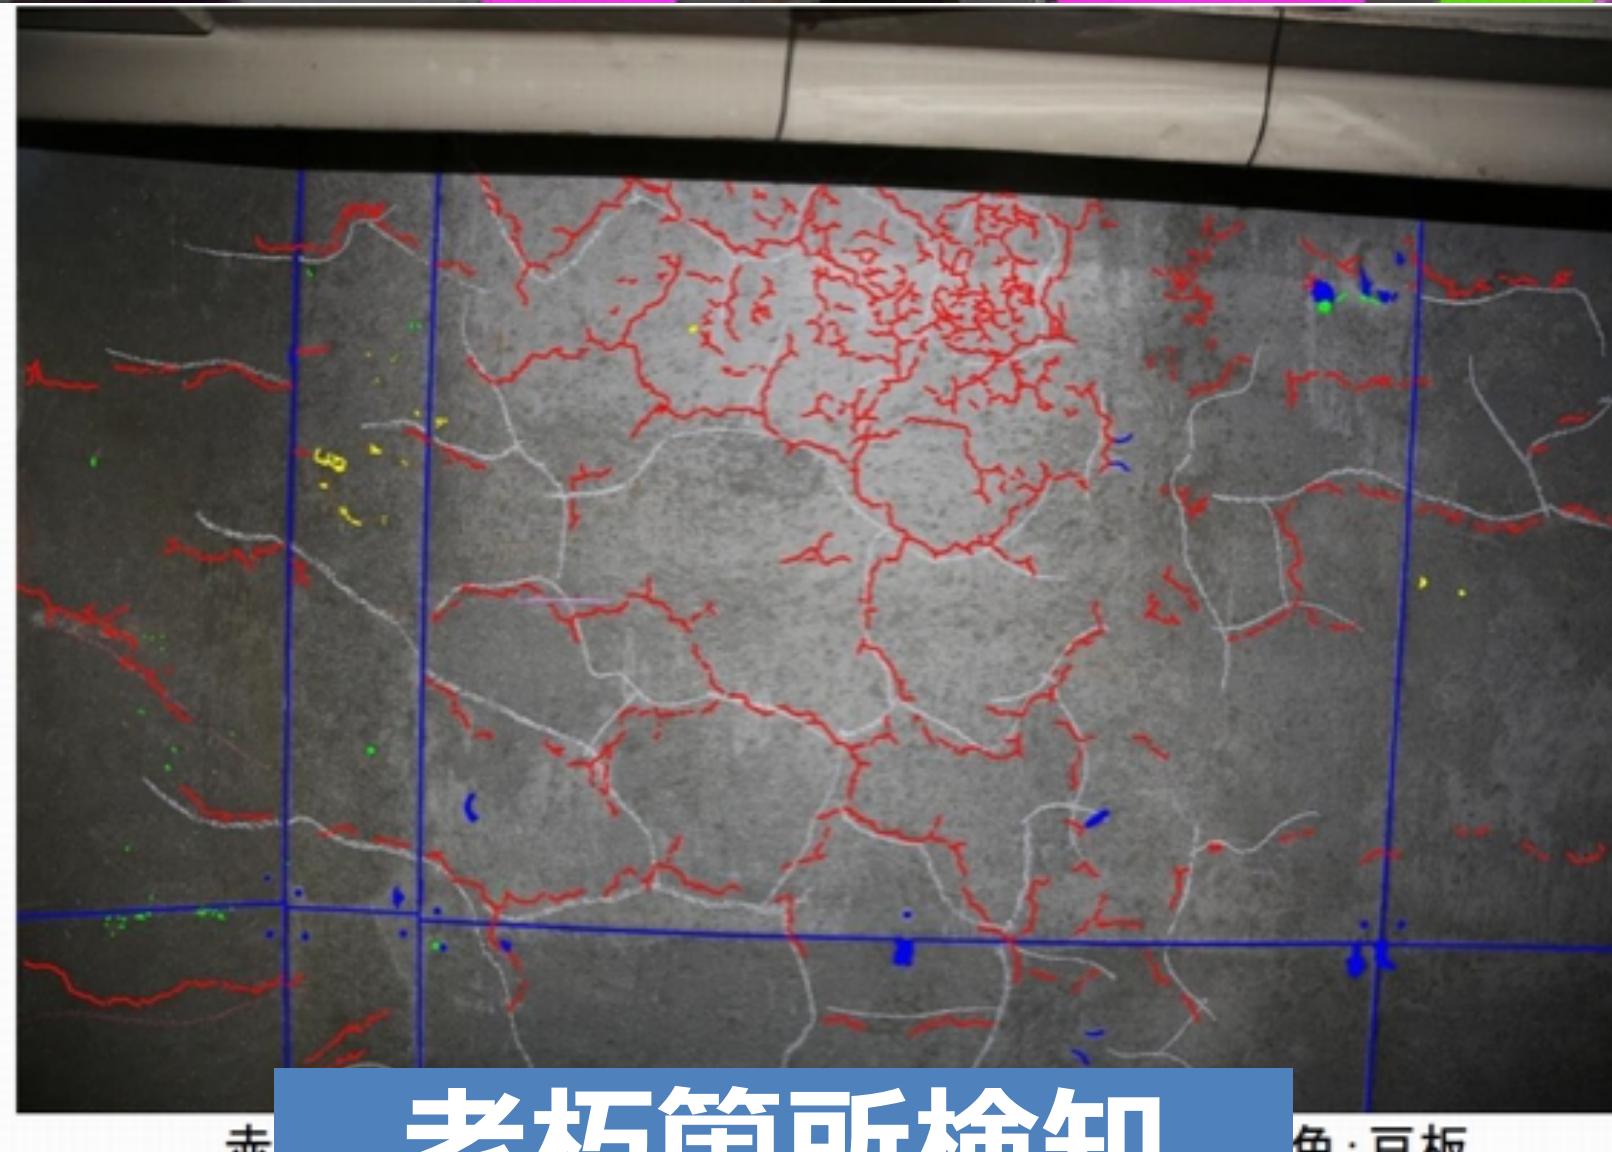
一般物体認識



顔認識



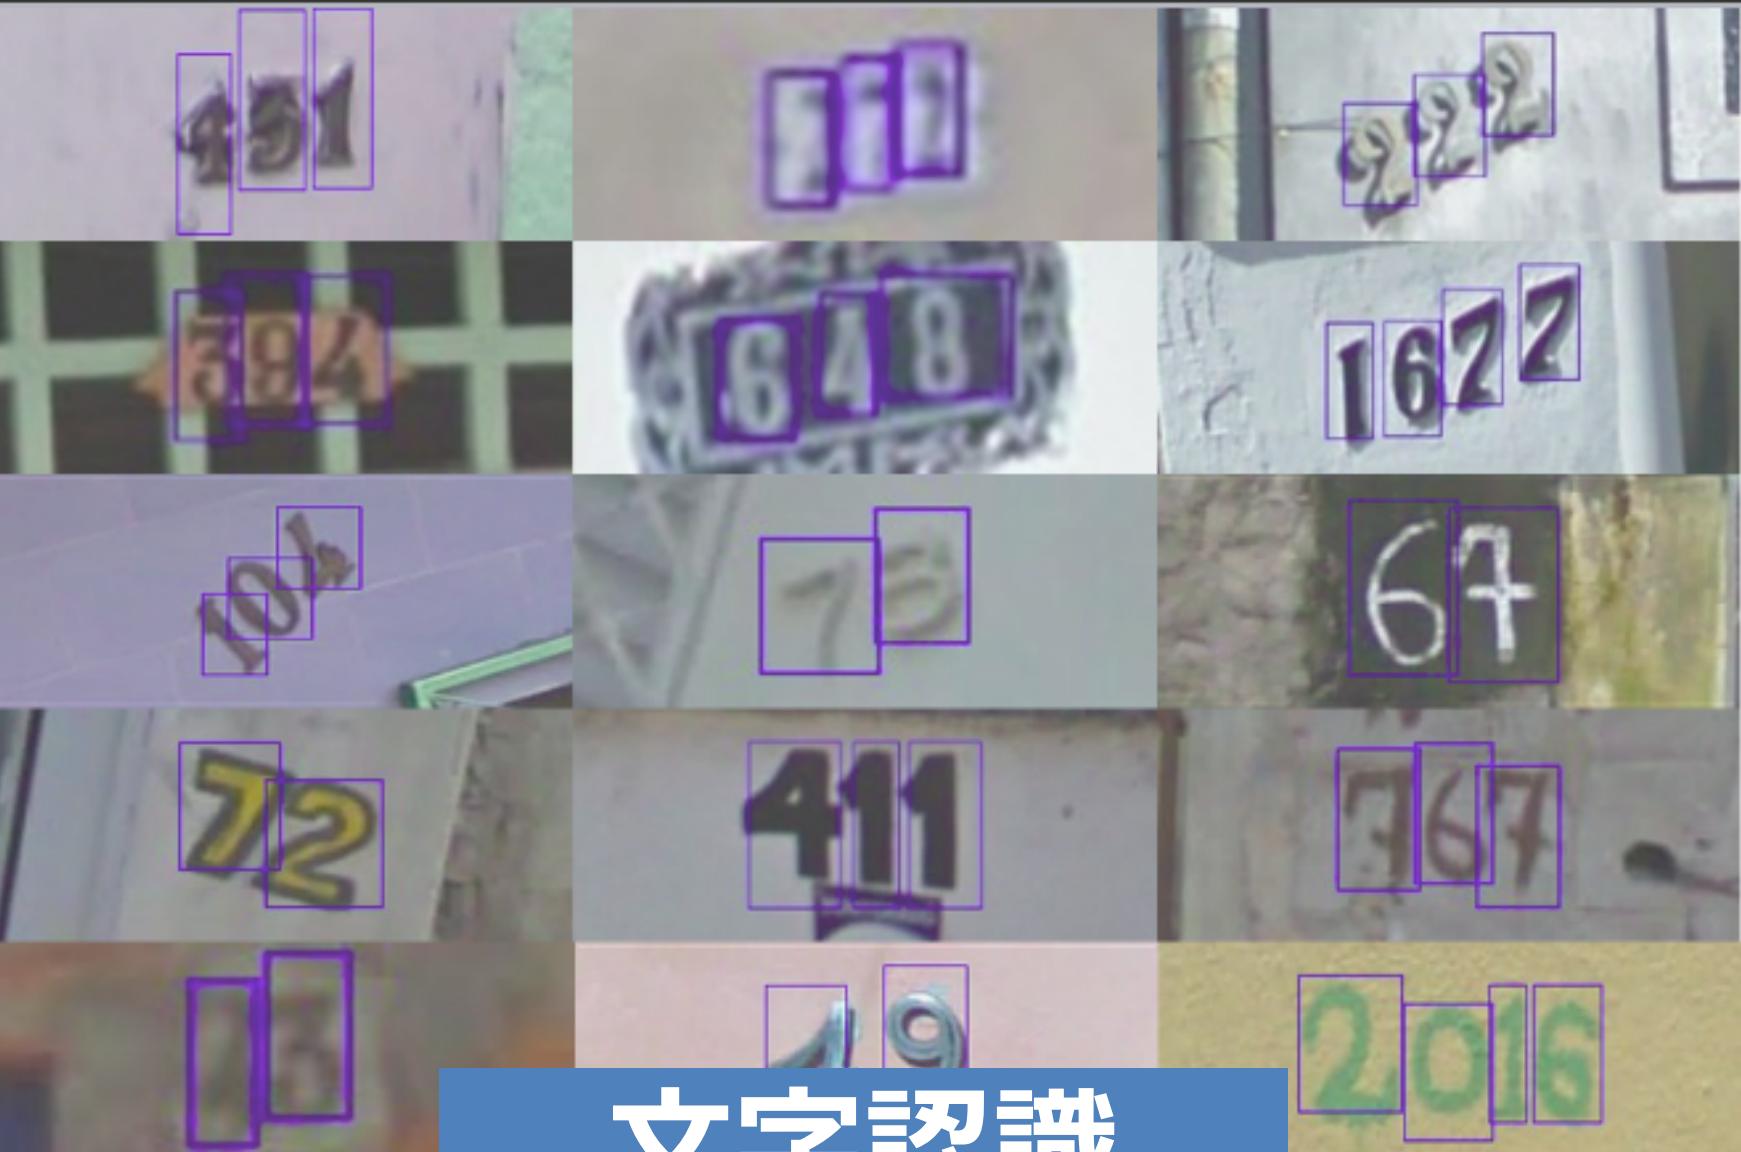
年齢推定



老朽箇所検知



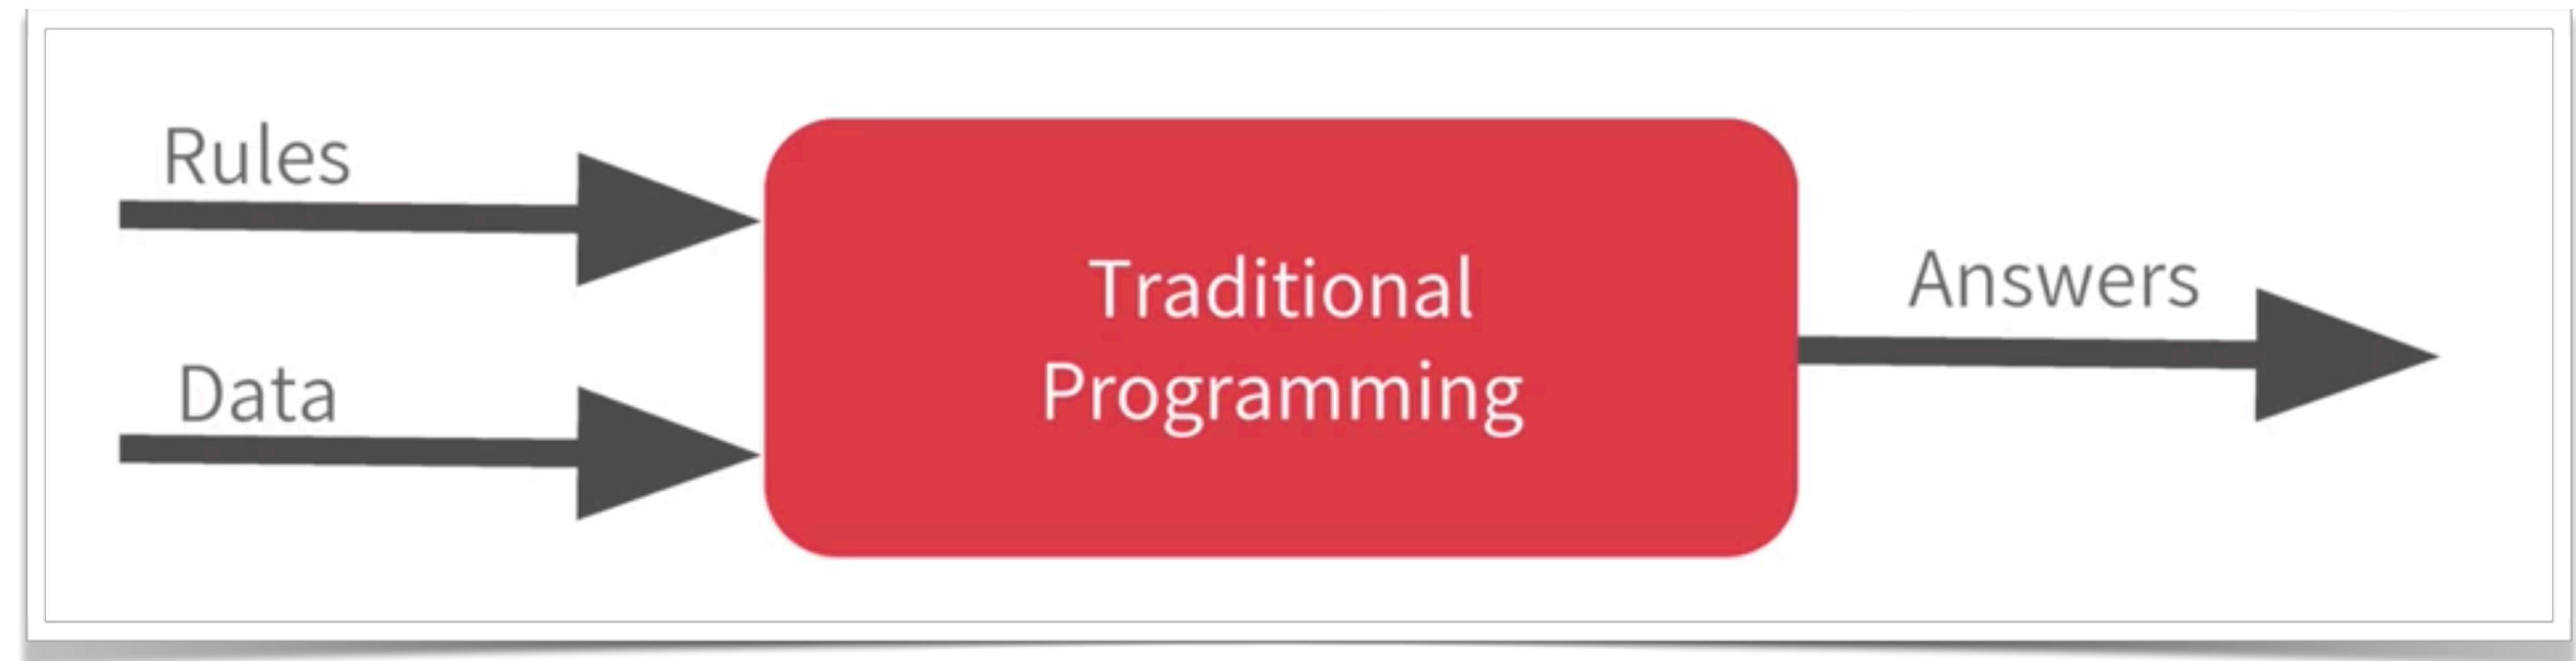
セグメンテーション



文字認識



ニューラルネットワークの基礎



<https://www.coursera.org/learn/introduction-tensorflow>

Activity Recognition



```
if(speed<4){  
    status=WALKING;  
}
```

<https://www.coursera.org/learn/introduction-tensorflow>

Activity Recognition



```
if(speed<4){  
    status=WALKING;  
}
```

```
if(speed<4){  
    status=WALKING;  
} else {  
    status=RUNNING;  
}
```

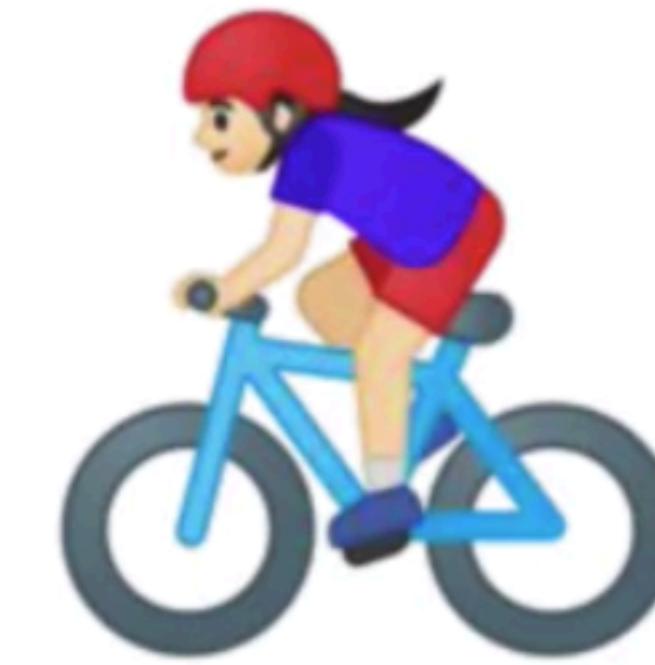
Activity Recognition



```
if(speed<4){  
    status=WALKING;  
}
```



```
if(speed<4){  
    status=WALKING;  
} else {  
    status=RUNNING;  
}
```



```
if(speed<4){  
    status=WALKING;  
} else if(speed<12){  
    status=RUNNING;  
} else {  
    status=BIKING;  
}
```

Activity Recognition



```
if(speed<4){  
    status=WALKING;  
}
```



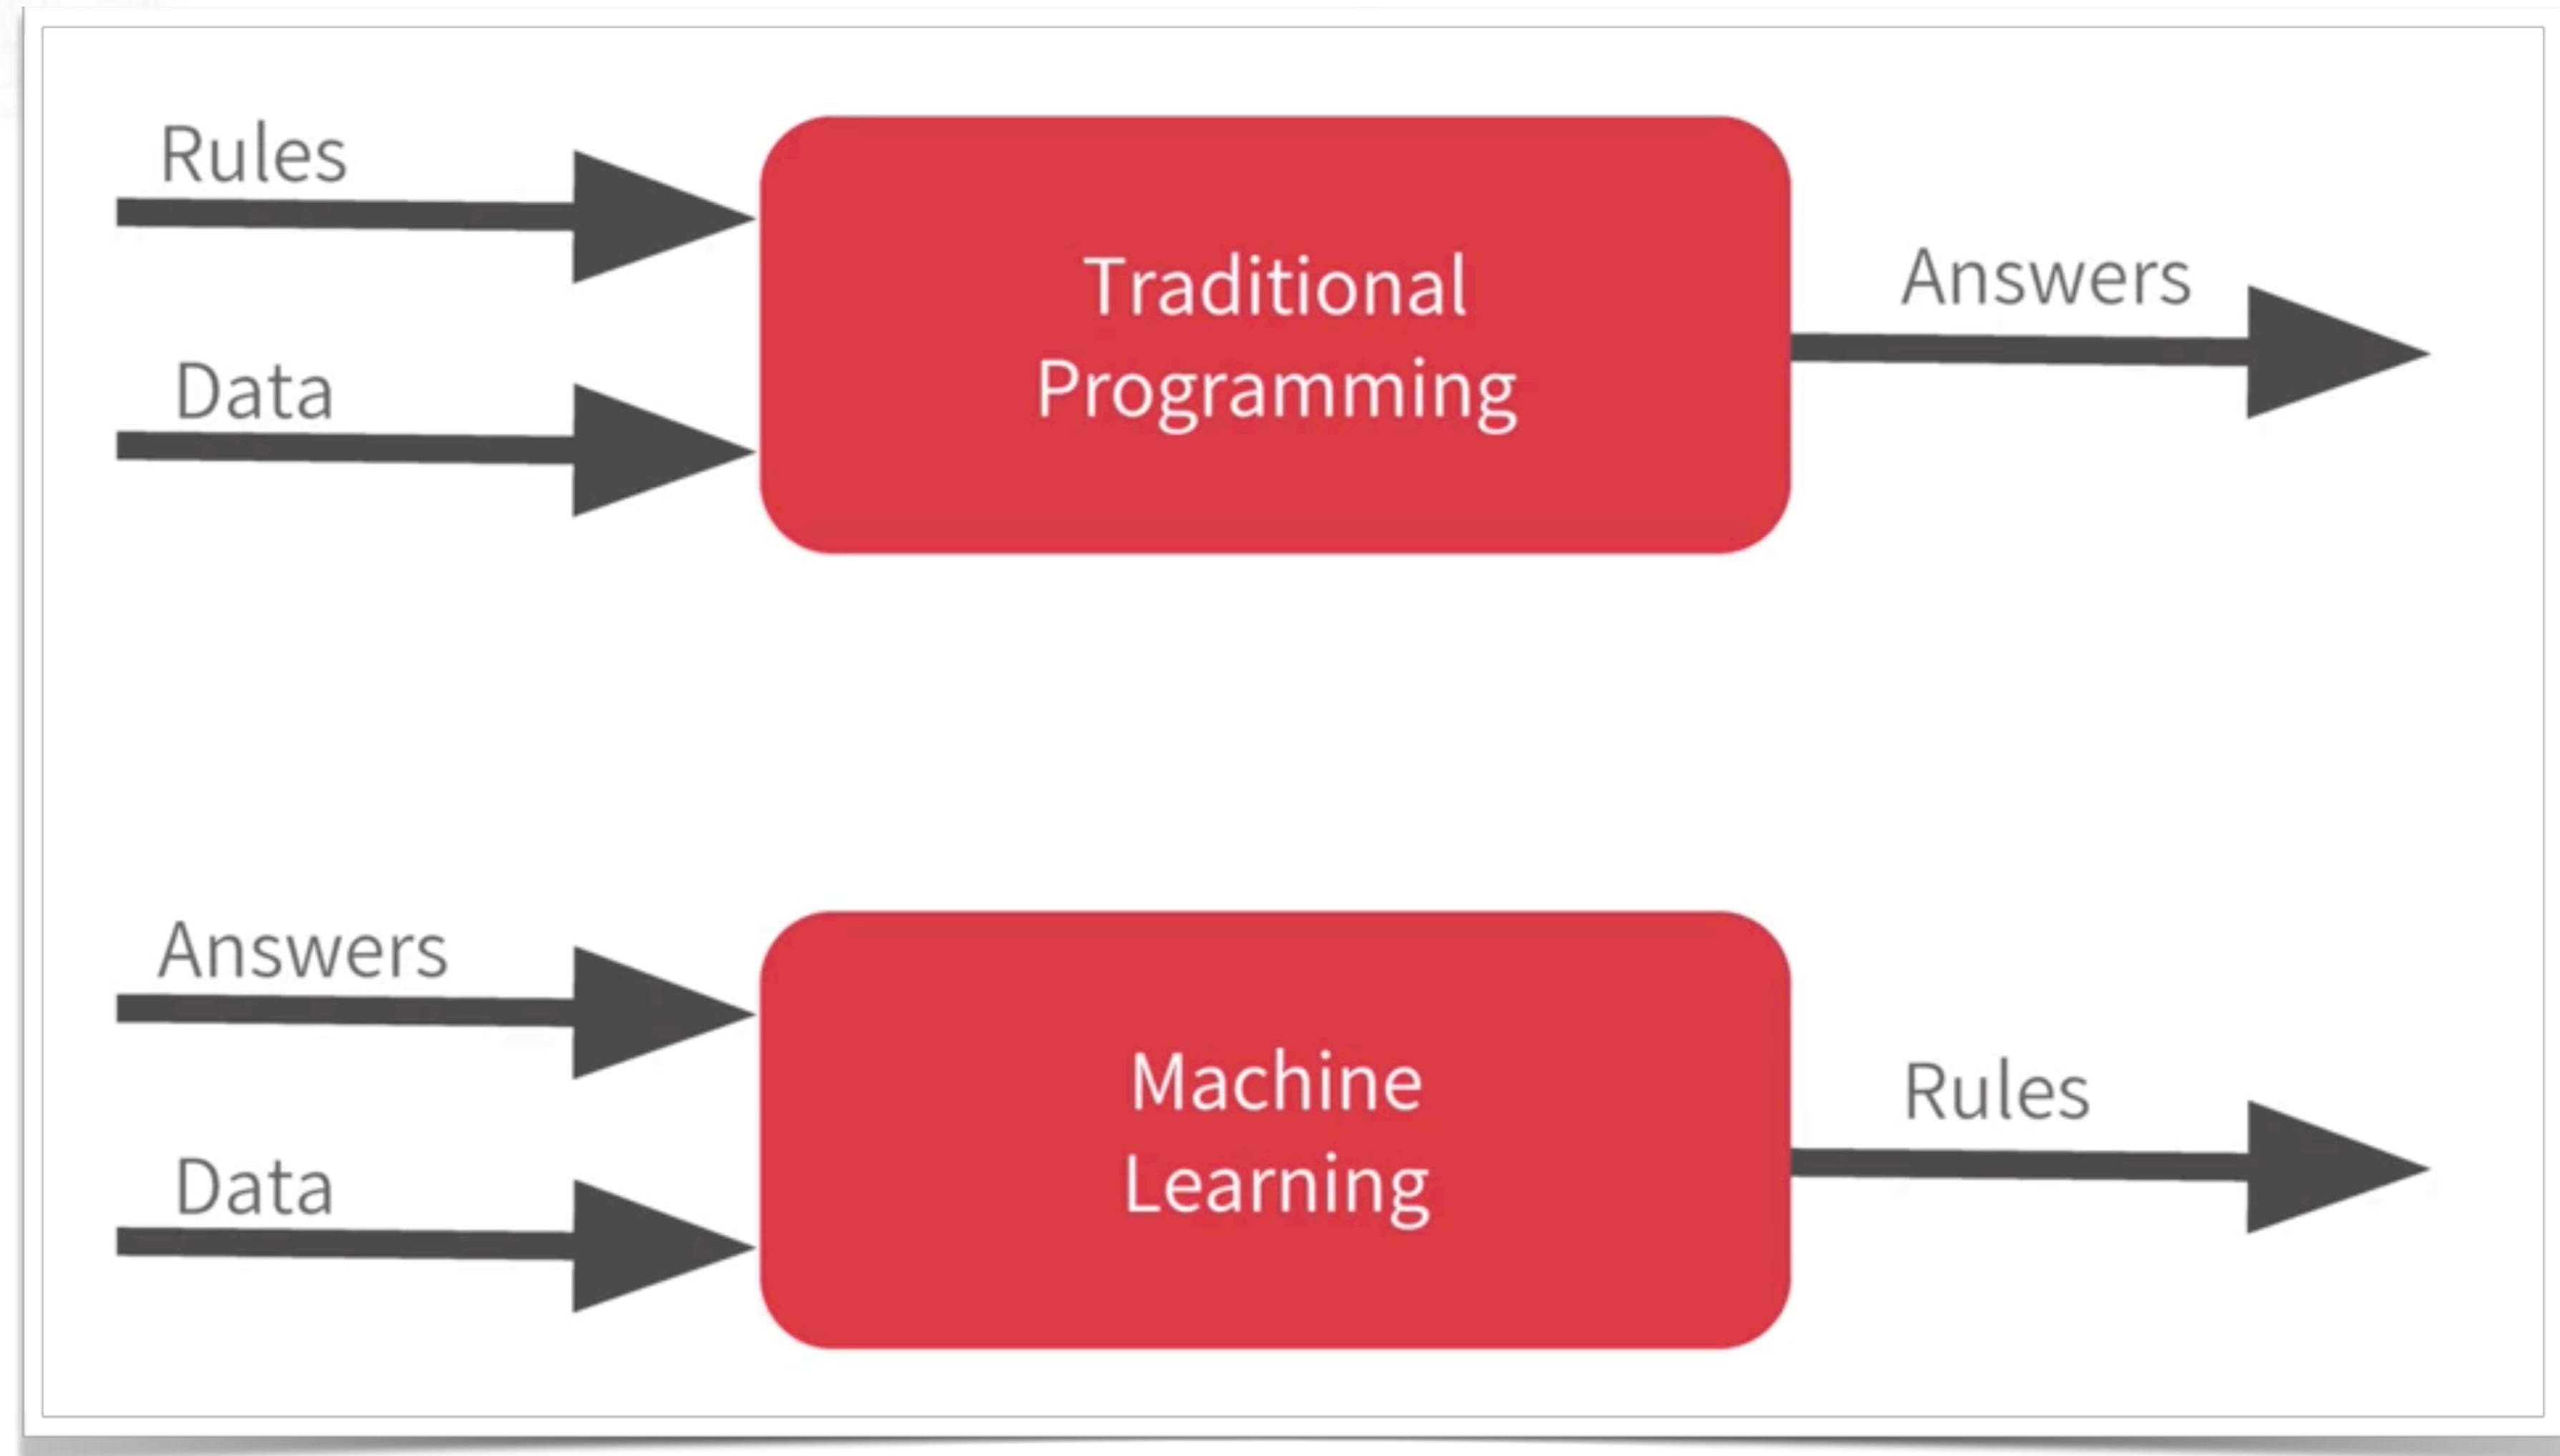
```
if(speed<4){  
    status=WALKING;  
} else {  
    status=RUNNING;  
}
```



```
if(speed<4){  
    status=WALKING;  
} else if(speed<12){  
    status=RUNNING;  
} else {  
    status=BIKING;  
}
```

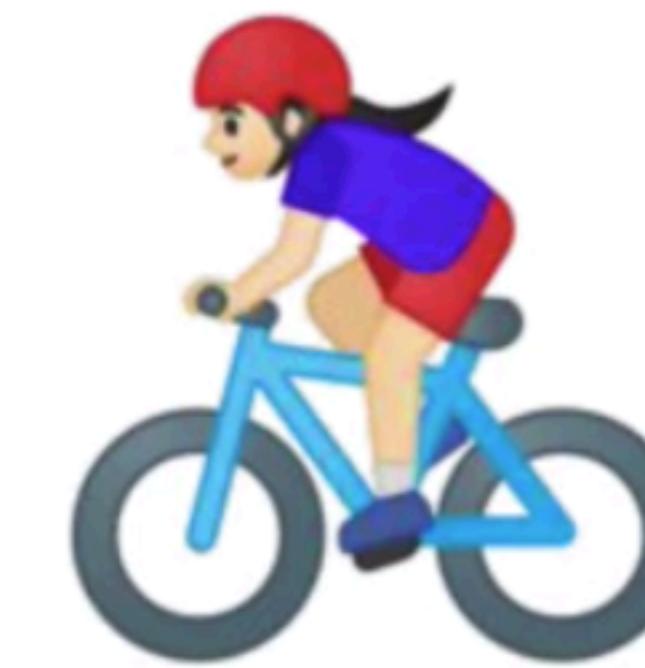


// Oh crap



<https://www.coursera.org/learn/introduction-tensorflow>

Activity Recognition



0101001010100101010
1001010101001011101
0100101010010101001
0101001010100101010

Label = WALKING

1010100101001010101
0101010010010010001
0010011111010101111
1010100100111101011

Label = RUNNING

1001010011111010101
1101010111010101110
1010101111010101011
1111110001111010101

Label = BIKING

111111111010011101
0011111010111110101
0101110101010101110
1010101010100111110

Label = GOLFING
(Sort of)

機械学習の"Hello World"

$$\begin{aligned}x &= -2, -1, 0, 1, 2, 3, 4 \\y &= -3, -1, 1, 3, 5, 7, 9\end{aligned}$$

$$y = f(x)$$

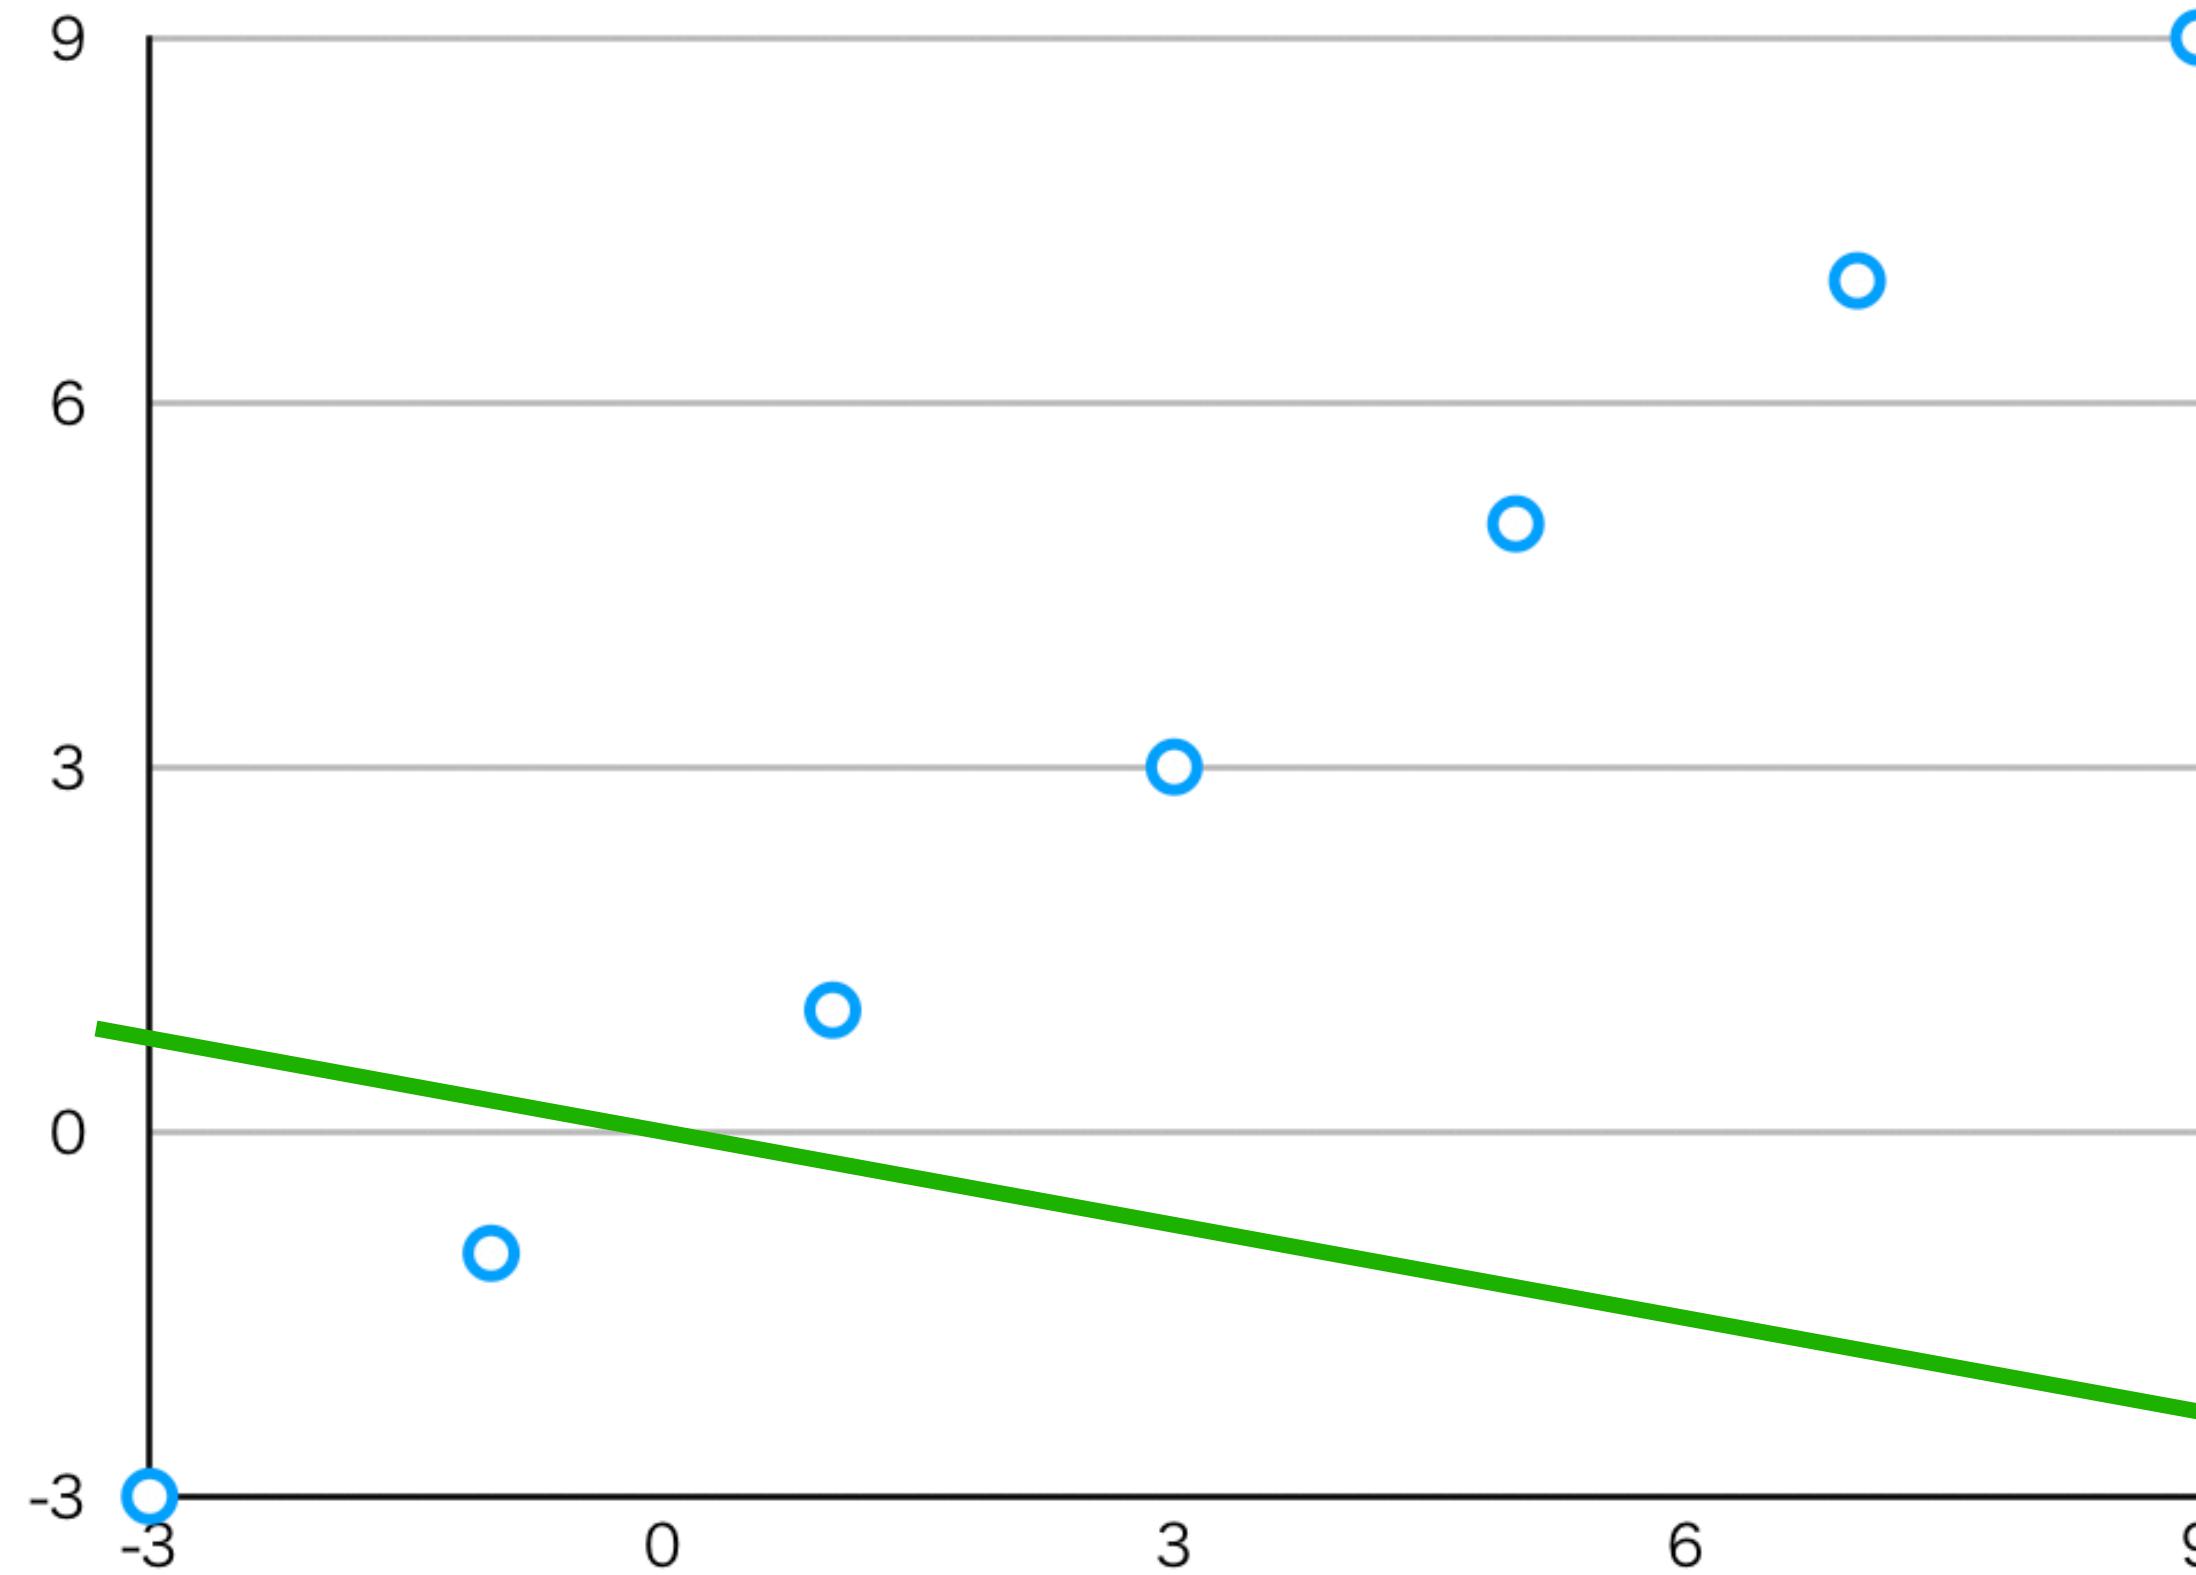
機械学習の"Hello World"

$$\begin{aligned}x &= -2, -1, 0, 1, 2, 3, 4 \\y &= -3, -1, 1, 3, 5, 7, 9\end{aligned}$$

$$y = f(x) = 2x + 1$$

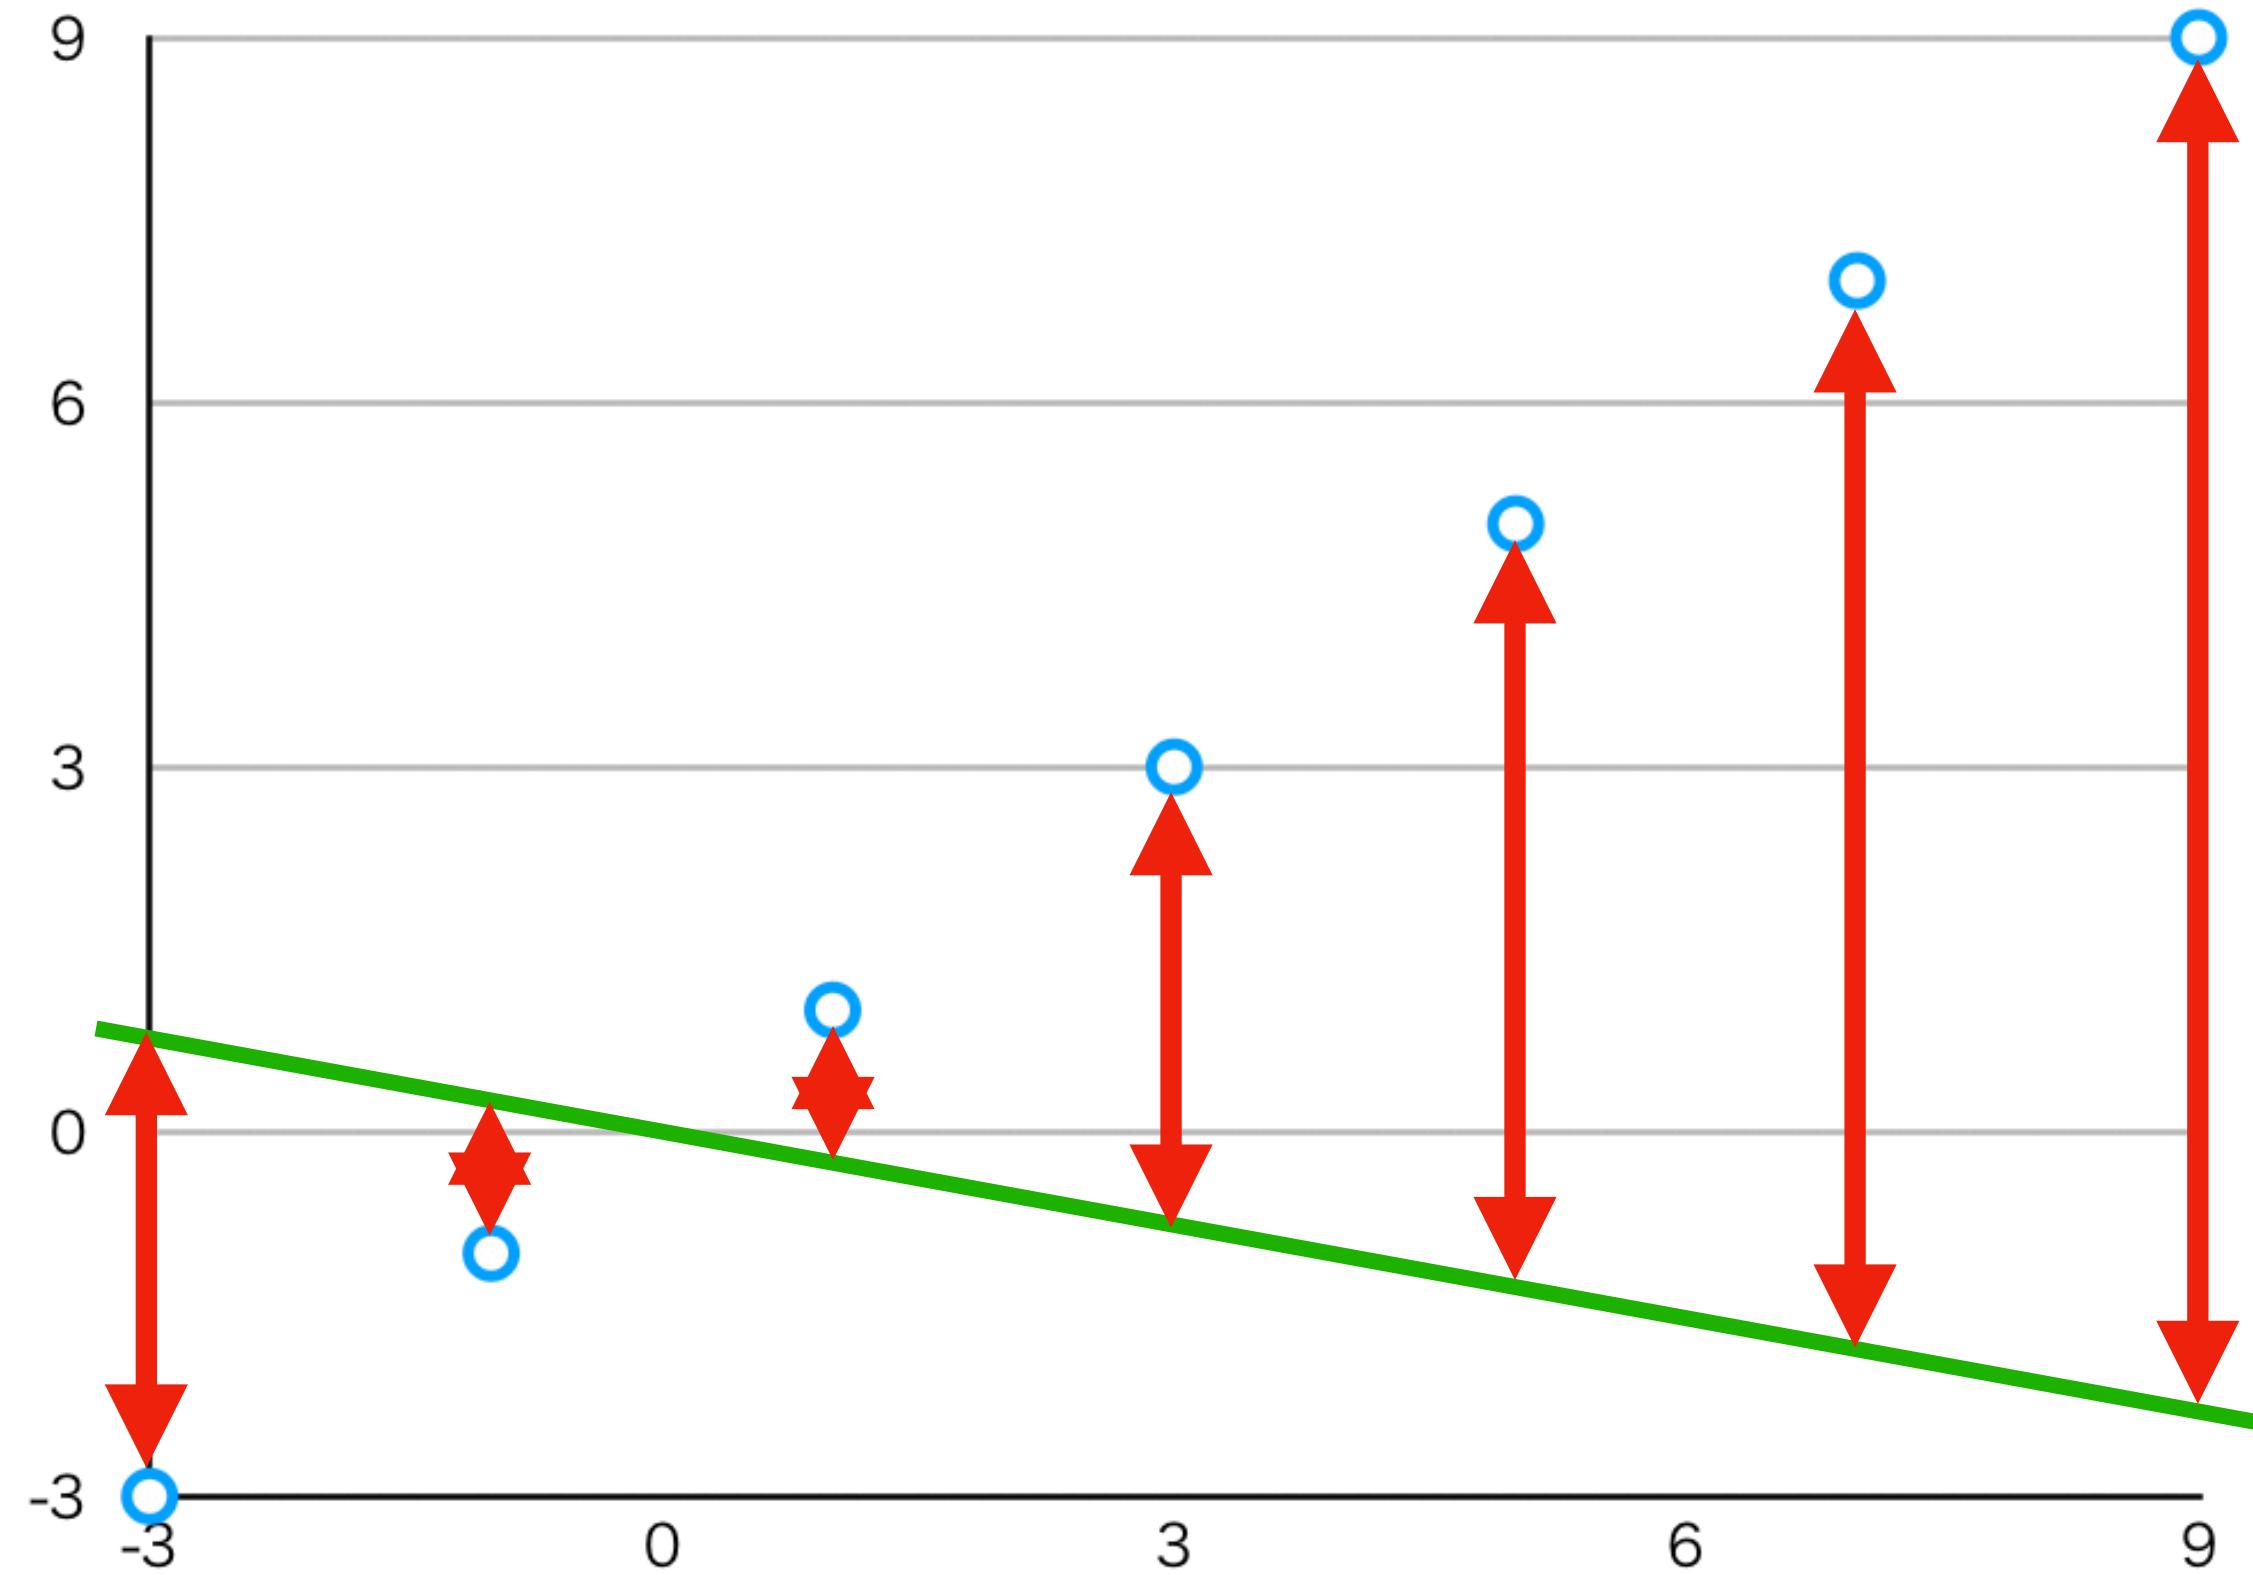
機械学習のアプローチ

- 適当にモデルを初期化 ($y_-=ax+b$)

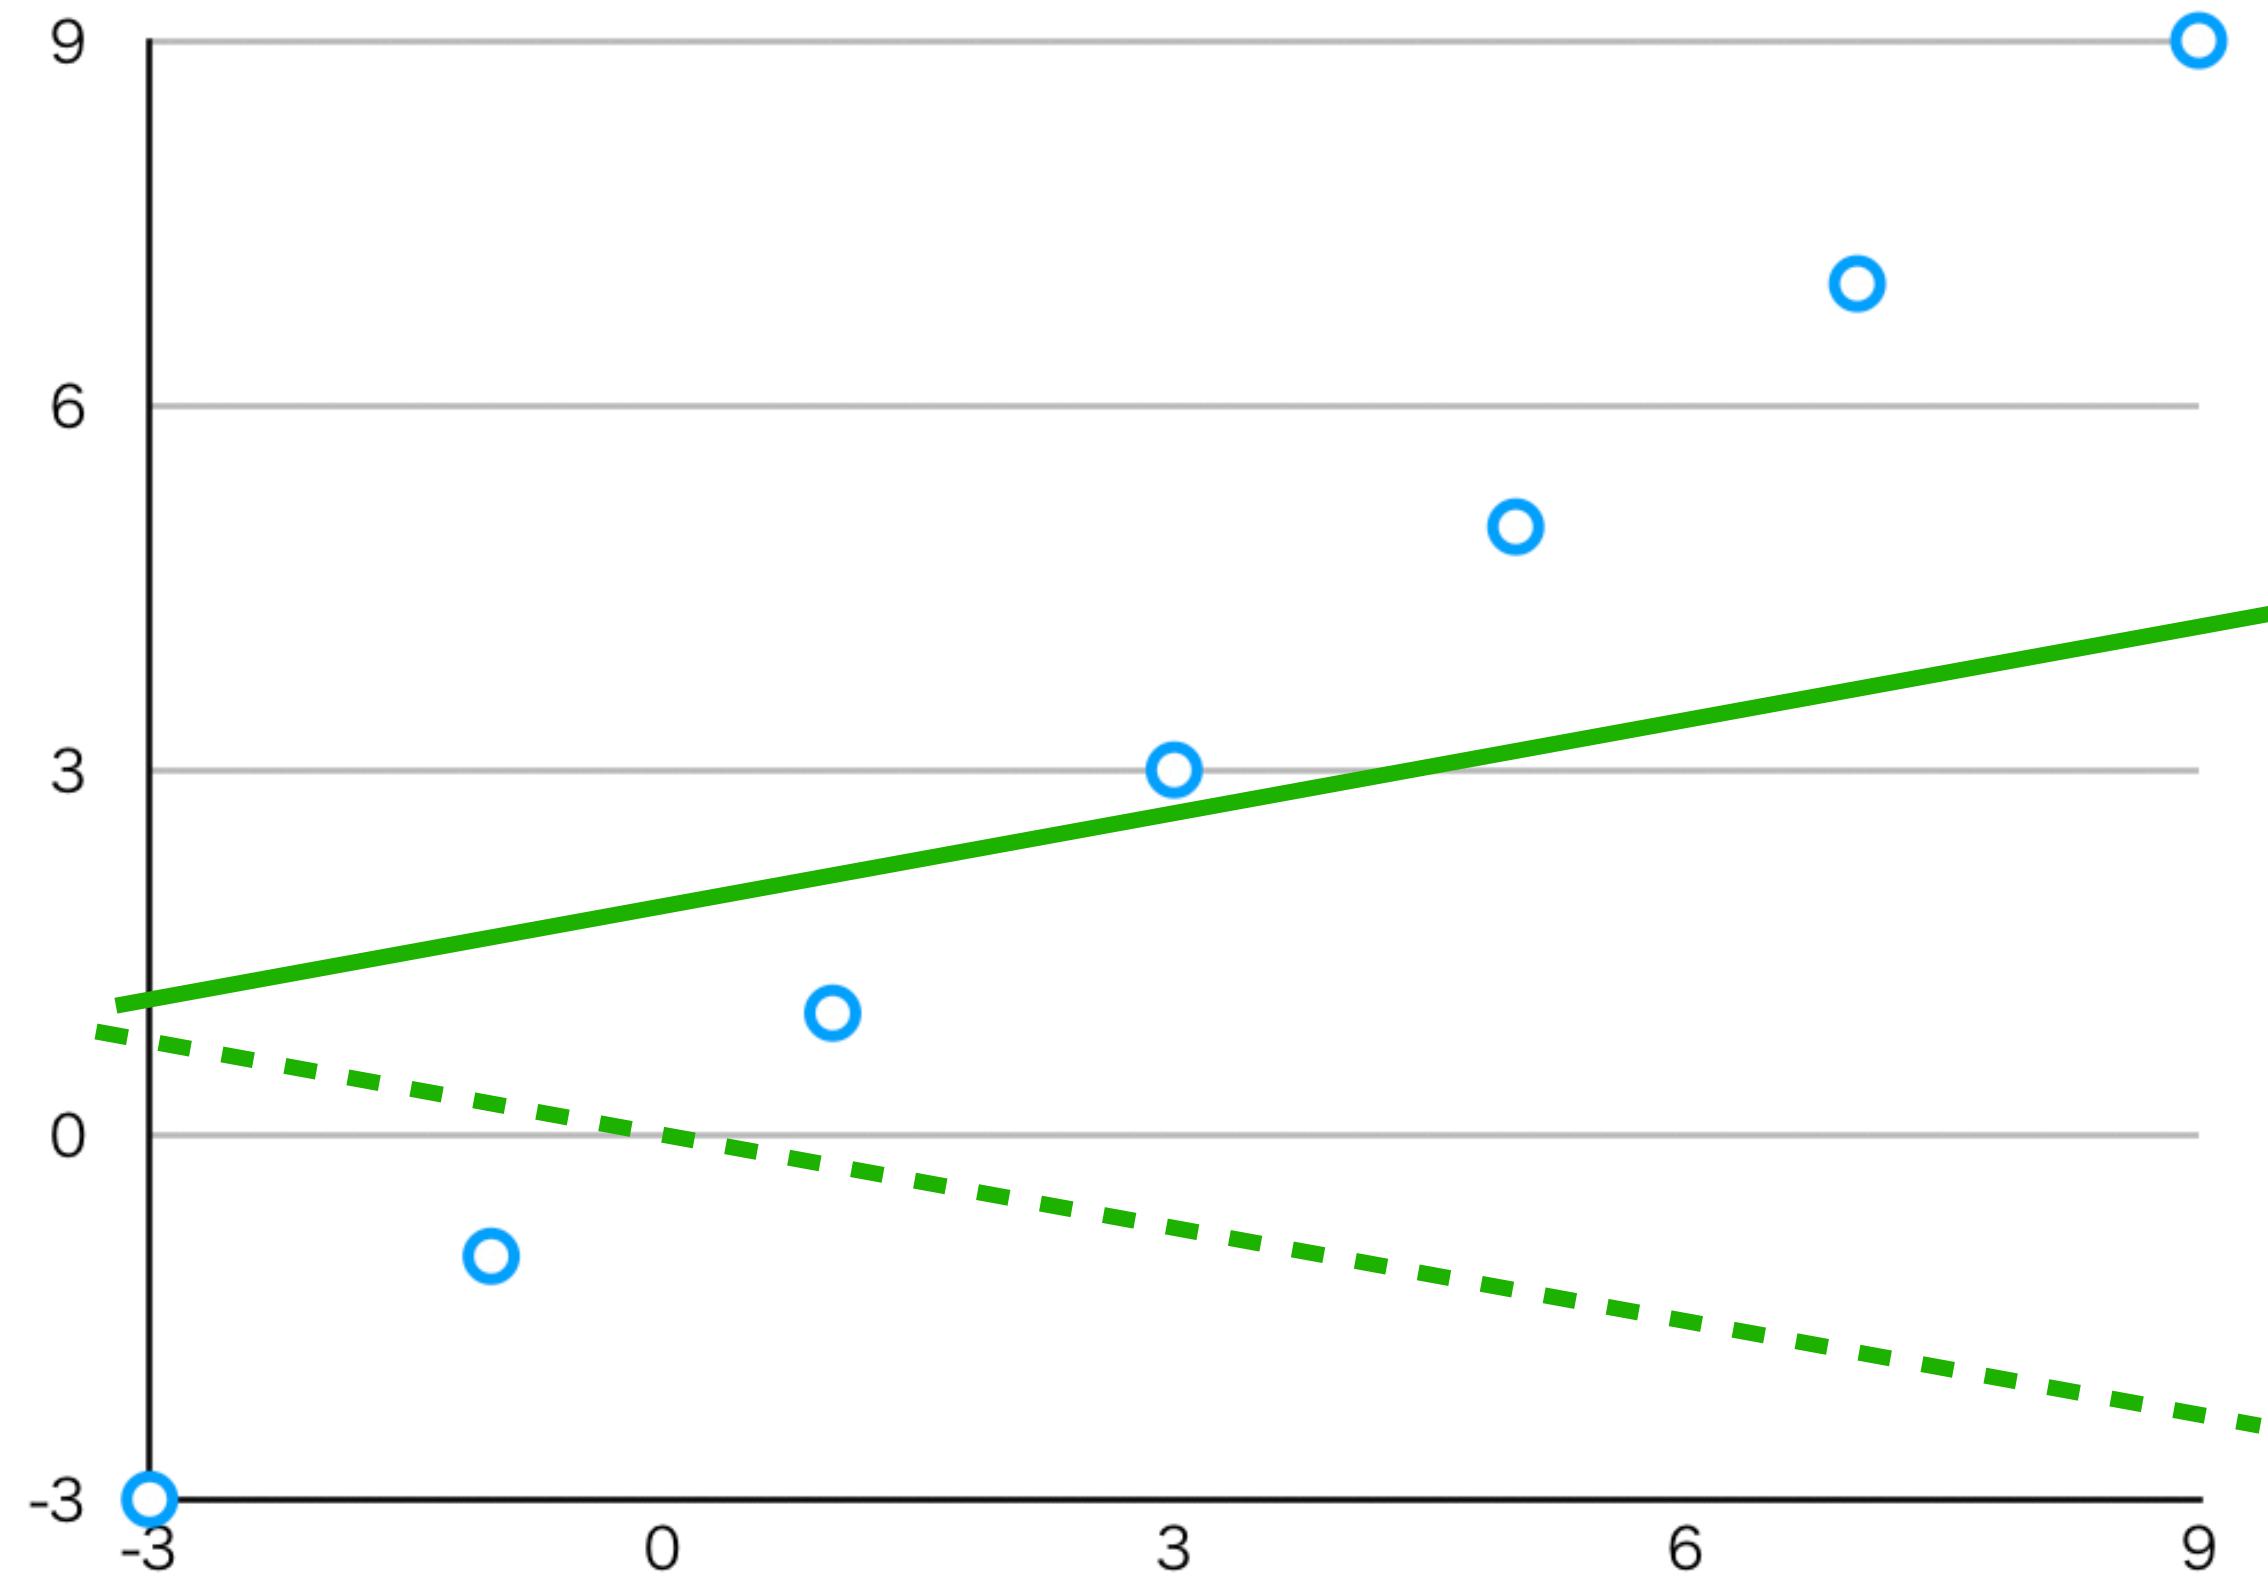


機械学習のアプローチ

- 適当にモデルを初期化 ($y_{_} = ax + b$)
- 誤差（損失）を計算 ($L = \frac{1}{N} \sum (y - y_{_})^2$)

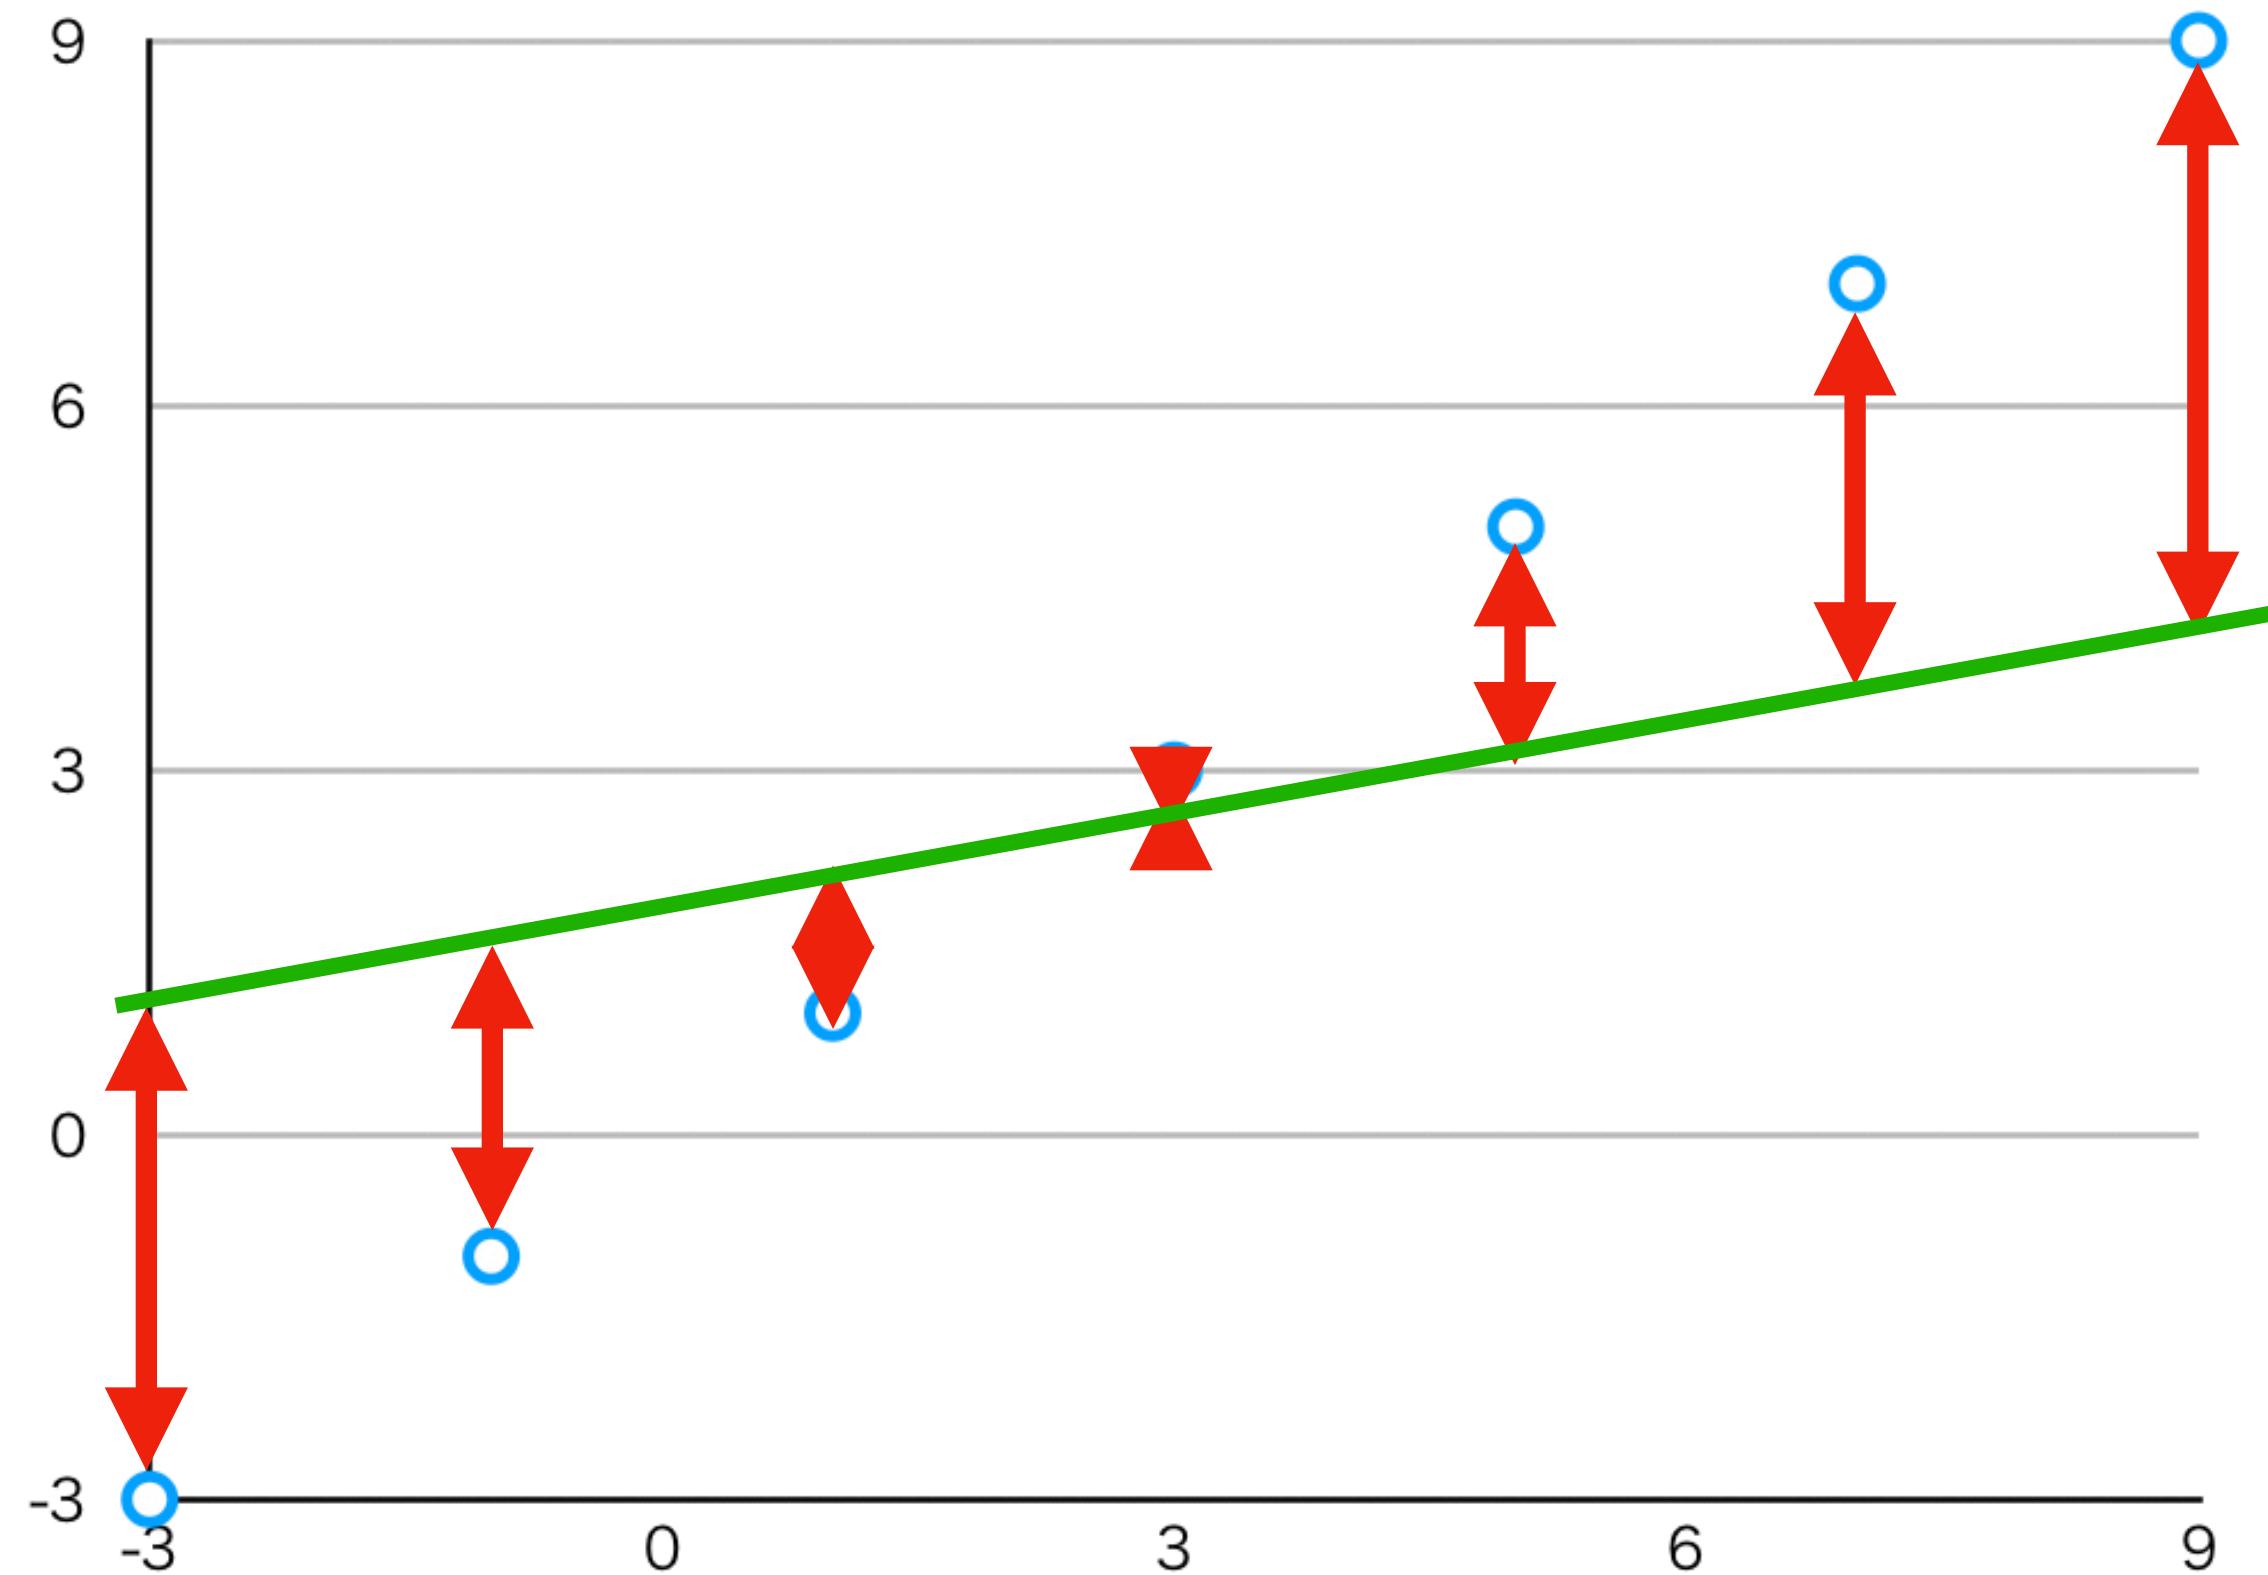


機械学習のアプローチ



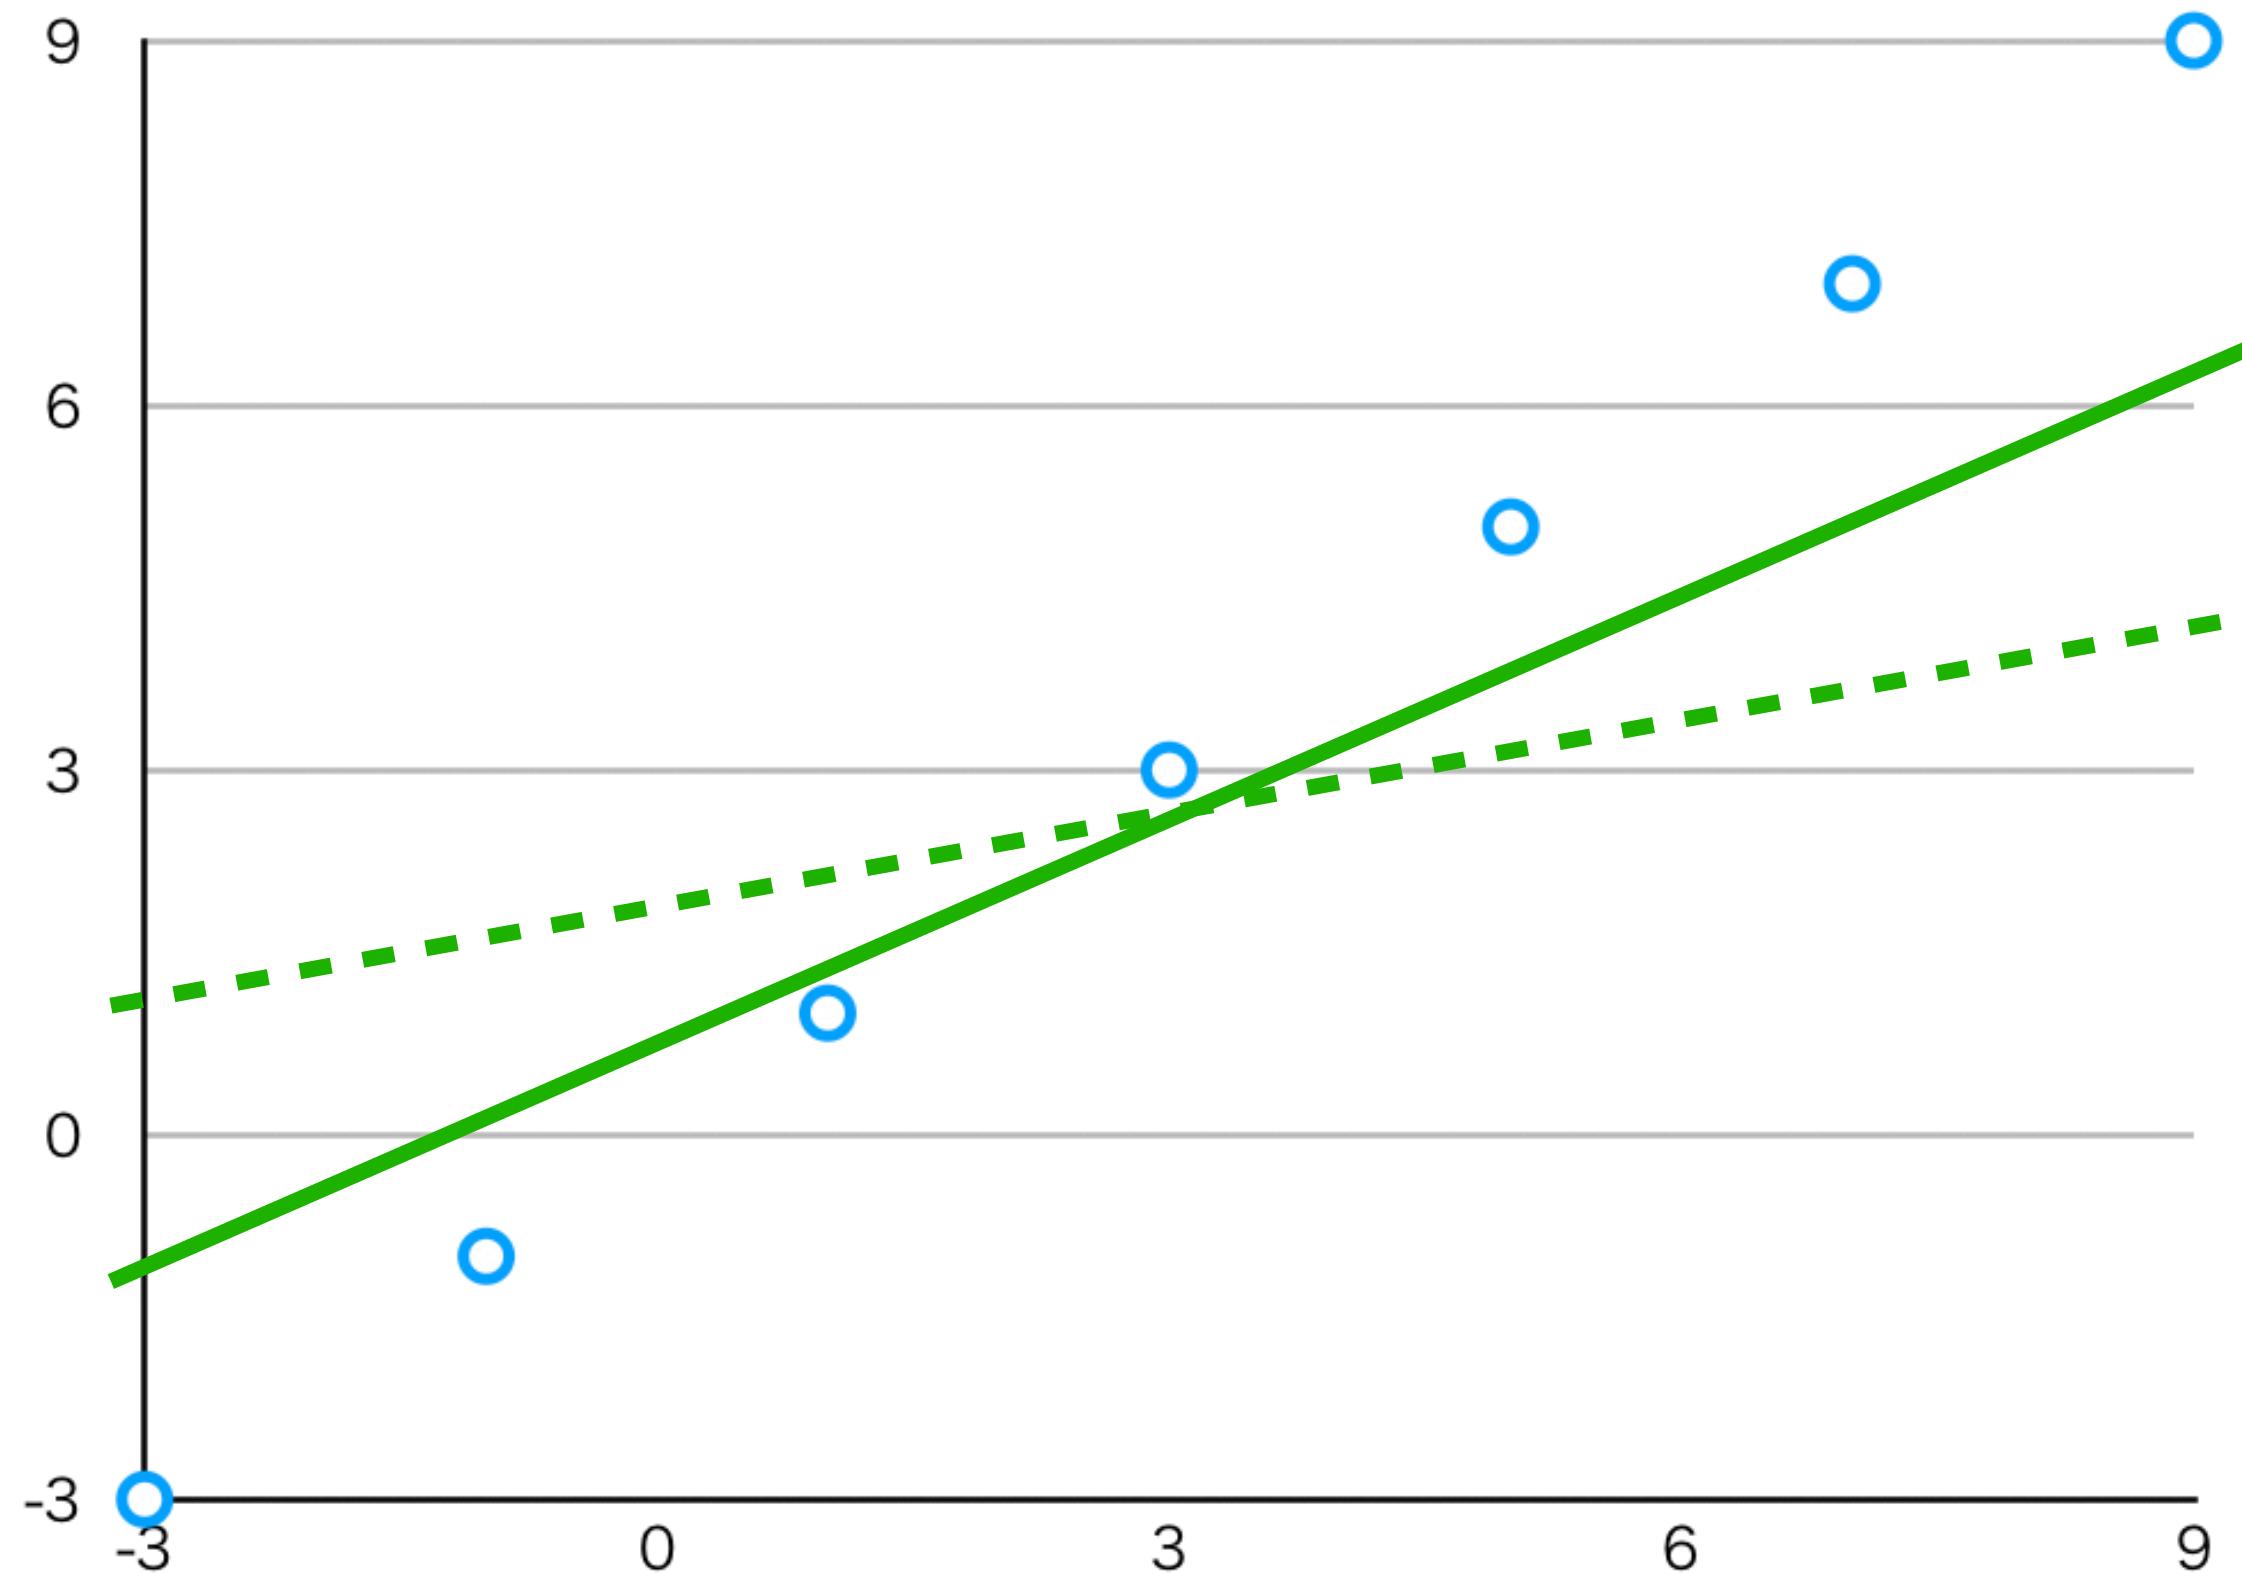
- 適当にモデルを初期化 ($y_- = ax + b$)
- 誤差（損失）を計算 ($L = \frac{1}{N} \sum (y - y_-)^2$)
- 誤差が小さくなるようにパラメータ (a, b) を少し更新 ($a \leftarrow a - \eta \frac{\partial L}{\partial a}$)

機械学習のアプローチ



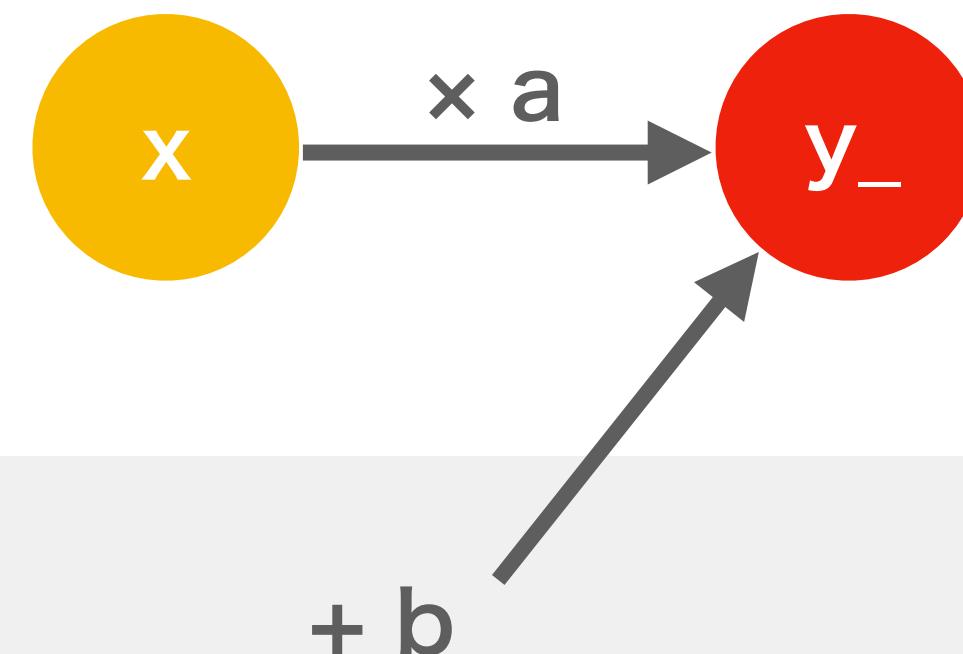
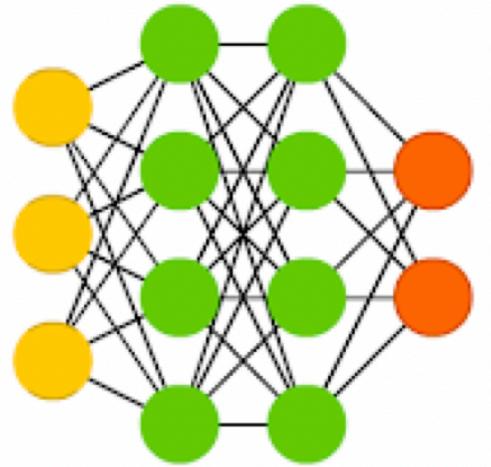
- 適当にモデルを初期化 ($y_{_} = ax + b$)
- 誤差（損失）を計算 ($L = 1/N \sum (y - y_{_})^2$)
- 誤差が小さくなるようにパラメータ (a, b) を少し更新 ($a \leftarrow a - \eta \partial L / \partial a$)
- 誤差（損失）を計算 ($L = 1/N \sum (y - y_{_})^2$)

機械学習のアプローチ



- 適当にモデルを初期化 ($y_{_} = ax + b$)
- 誤差（損失）を計算 ($L = 1/N \sum (y - y_{_})^2$)
- 誤差が小さくなるようにパラメータ (a, b) を少し更新 ($a \leftarrow a - \eta \frac{\partial L}{\partial a}$)
- 誤差（損失）を計算 ($L = 1/N \sum (y - y_{_})^2$)
- 誤差が小さくなるようにパラメータ (a, b) を少し更新 ($a \leftarrow a - \eta \frac{\partial L}{\partial a}$)
- ...

```
1 import tensorflow as tf
2
3
4 model = tf.keras.Sequential([tf.keras.layers.Dense(units=1, input_shape=[1])])
5 model.compile(optimizer='sgd', loss='mean_squared_error')
6
7 xs = np.array([-1.0, 0.0, 1.0, 2.0, 3.0, 4.0], dtype=float)
8 ys = np.array([-3.0, -1.0, 1.0, 3.0, 5.0, 7.0], dtype=float)
9
10 model.fit(xs, ys, epochs=500)
11
12 >>>print(model.predict([10.0]))
13 [[18.977331]]
```



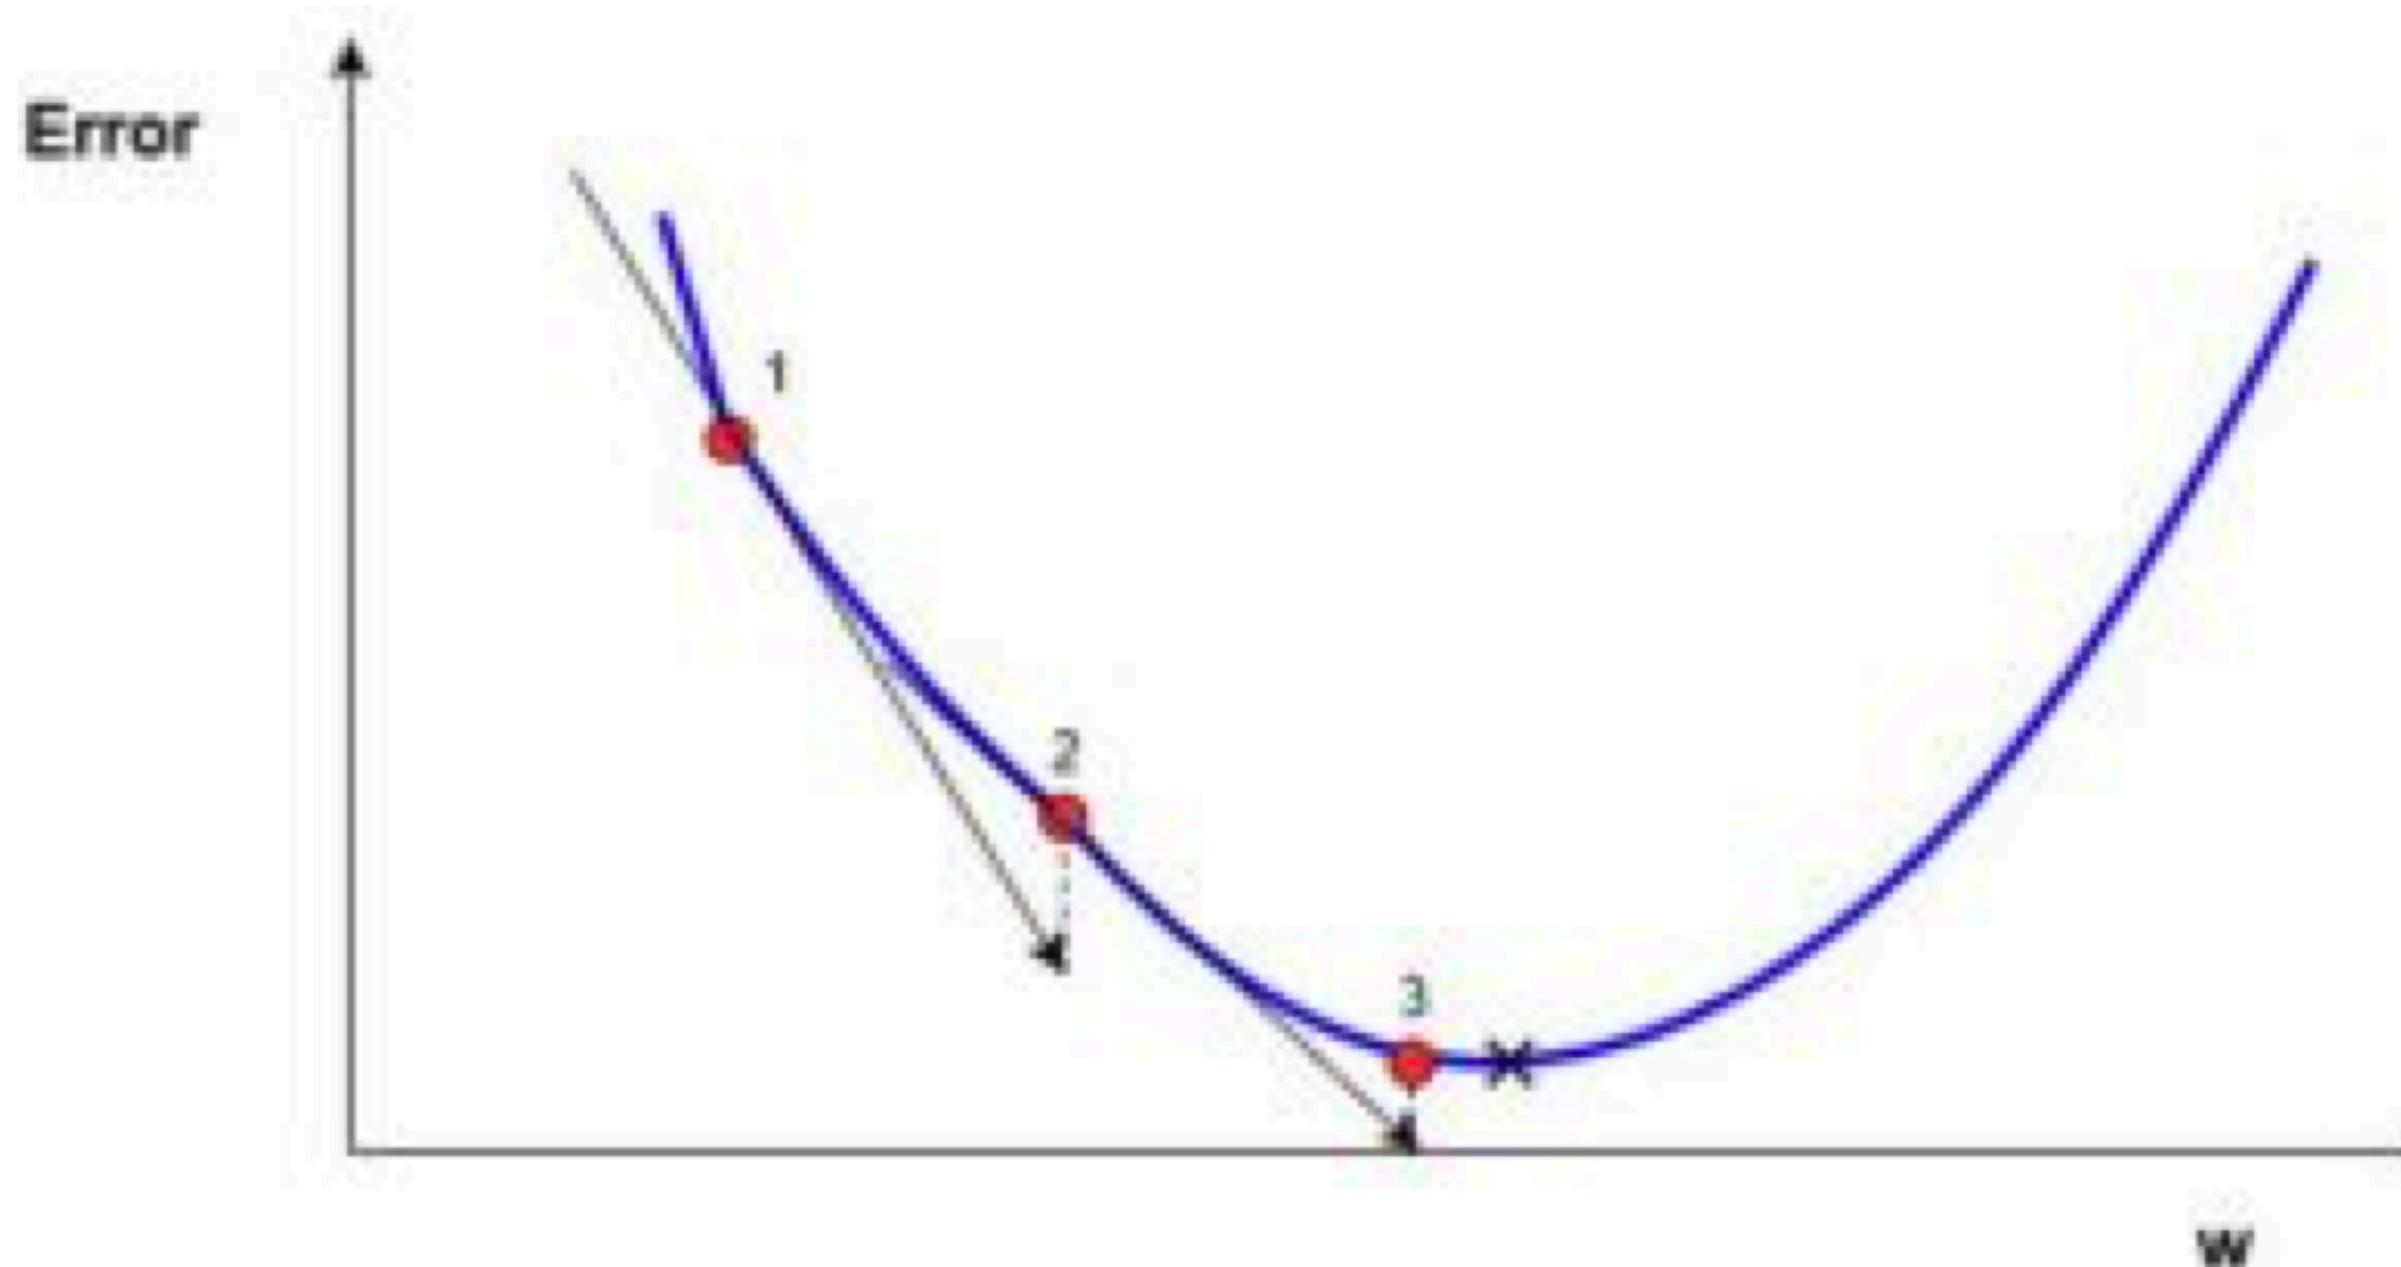
```
1 import tensorflow as tf
2
3
4 model = tf.keras.Sequential([tf.keras.layers.Dense(units=1, input_shape=[1])])
5 model.compile(optimizer='sgd', loss='mean_squared_error')
6
7 xs = np.array([-1.0, 0.0, 1.0, 2.0, 3.0, 4.0], dtype=float)
8 ys = np.array([-3.0, -1.0, 1.0, 3.0, 5.0, 7.0], dtype=float)
9
10 model.fit(xs, ys, epochs=500)
11
12 >>>print(model.predict([10.0]))
13 [[18.977331]]
```

```
1 import tensorflow as tf
2
3
4 model = tf.keras.Sequential([tf.keras.layers.Dense(units=1, input_shape=[1])])
5 model.compile(optimizer='sgd', loss='mean_squared_error')
6
7 xs = np.array([-1.0, 0.0, 1.0, 2.0, 3.0, 4.0], dtype=float)
8 ys = np.array([-3.0, -1.0, 1.0, 3.0, 5.0, 7.0], dtype=float)
9
10 model.fit(xs, ys, epochs=500)
11
12 >>>print(model.predict([10.0]))
13 [[18.977331]]
```

確率的勾配降下法 SGD, stochastic gradient descent.

デファクトスタンダード

(改良版 : momentum, Nesterovの加速法, RMSProp, Adam, ...)



$$w \leftarrow w - \eta \frac{\partial L}{\partial w}$$

勾配降下法



<https://www.yamakei-online.com/journal/detail.php?id=3185>

$$w \leftarrow w - \eta \frac{\partial L}{\partial w}$$

「目隠しで足元の勾配情報のみを使って山の頂上を目指すようなもの」

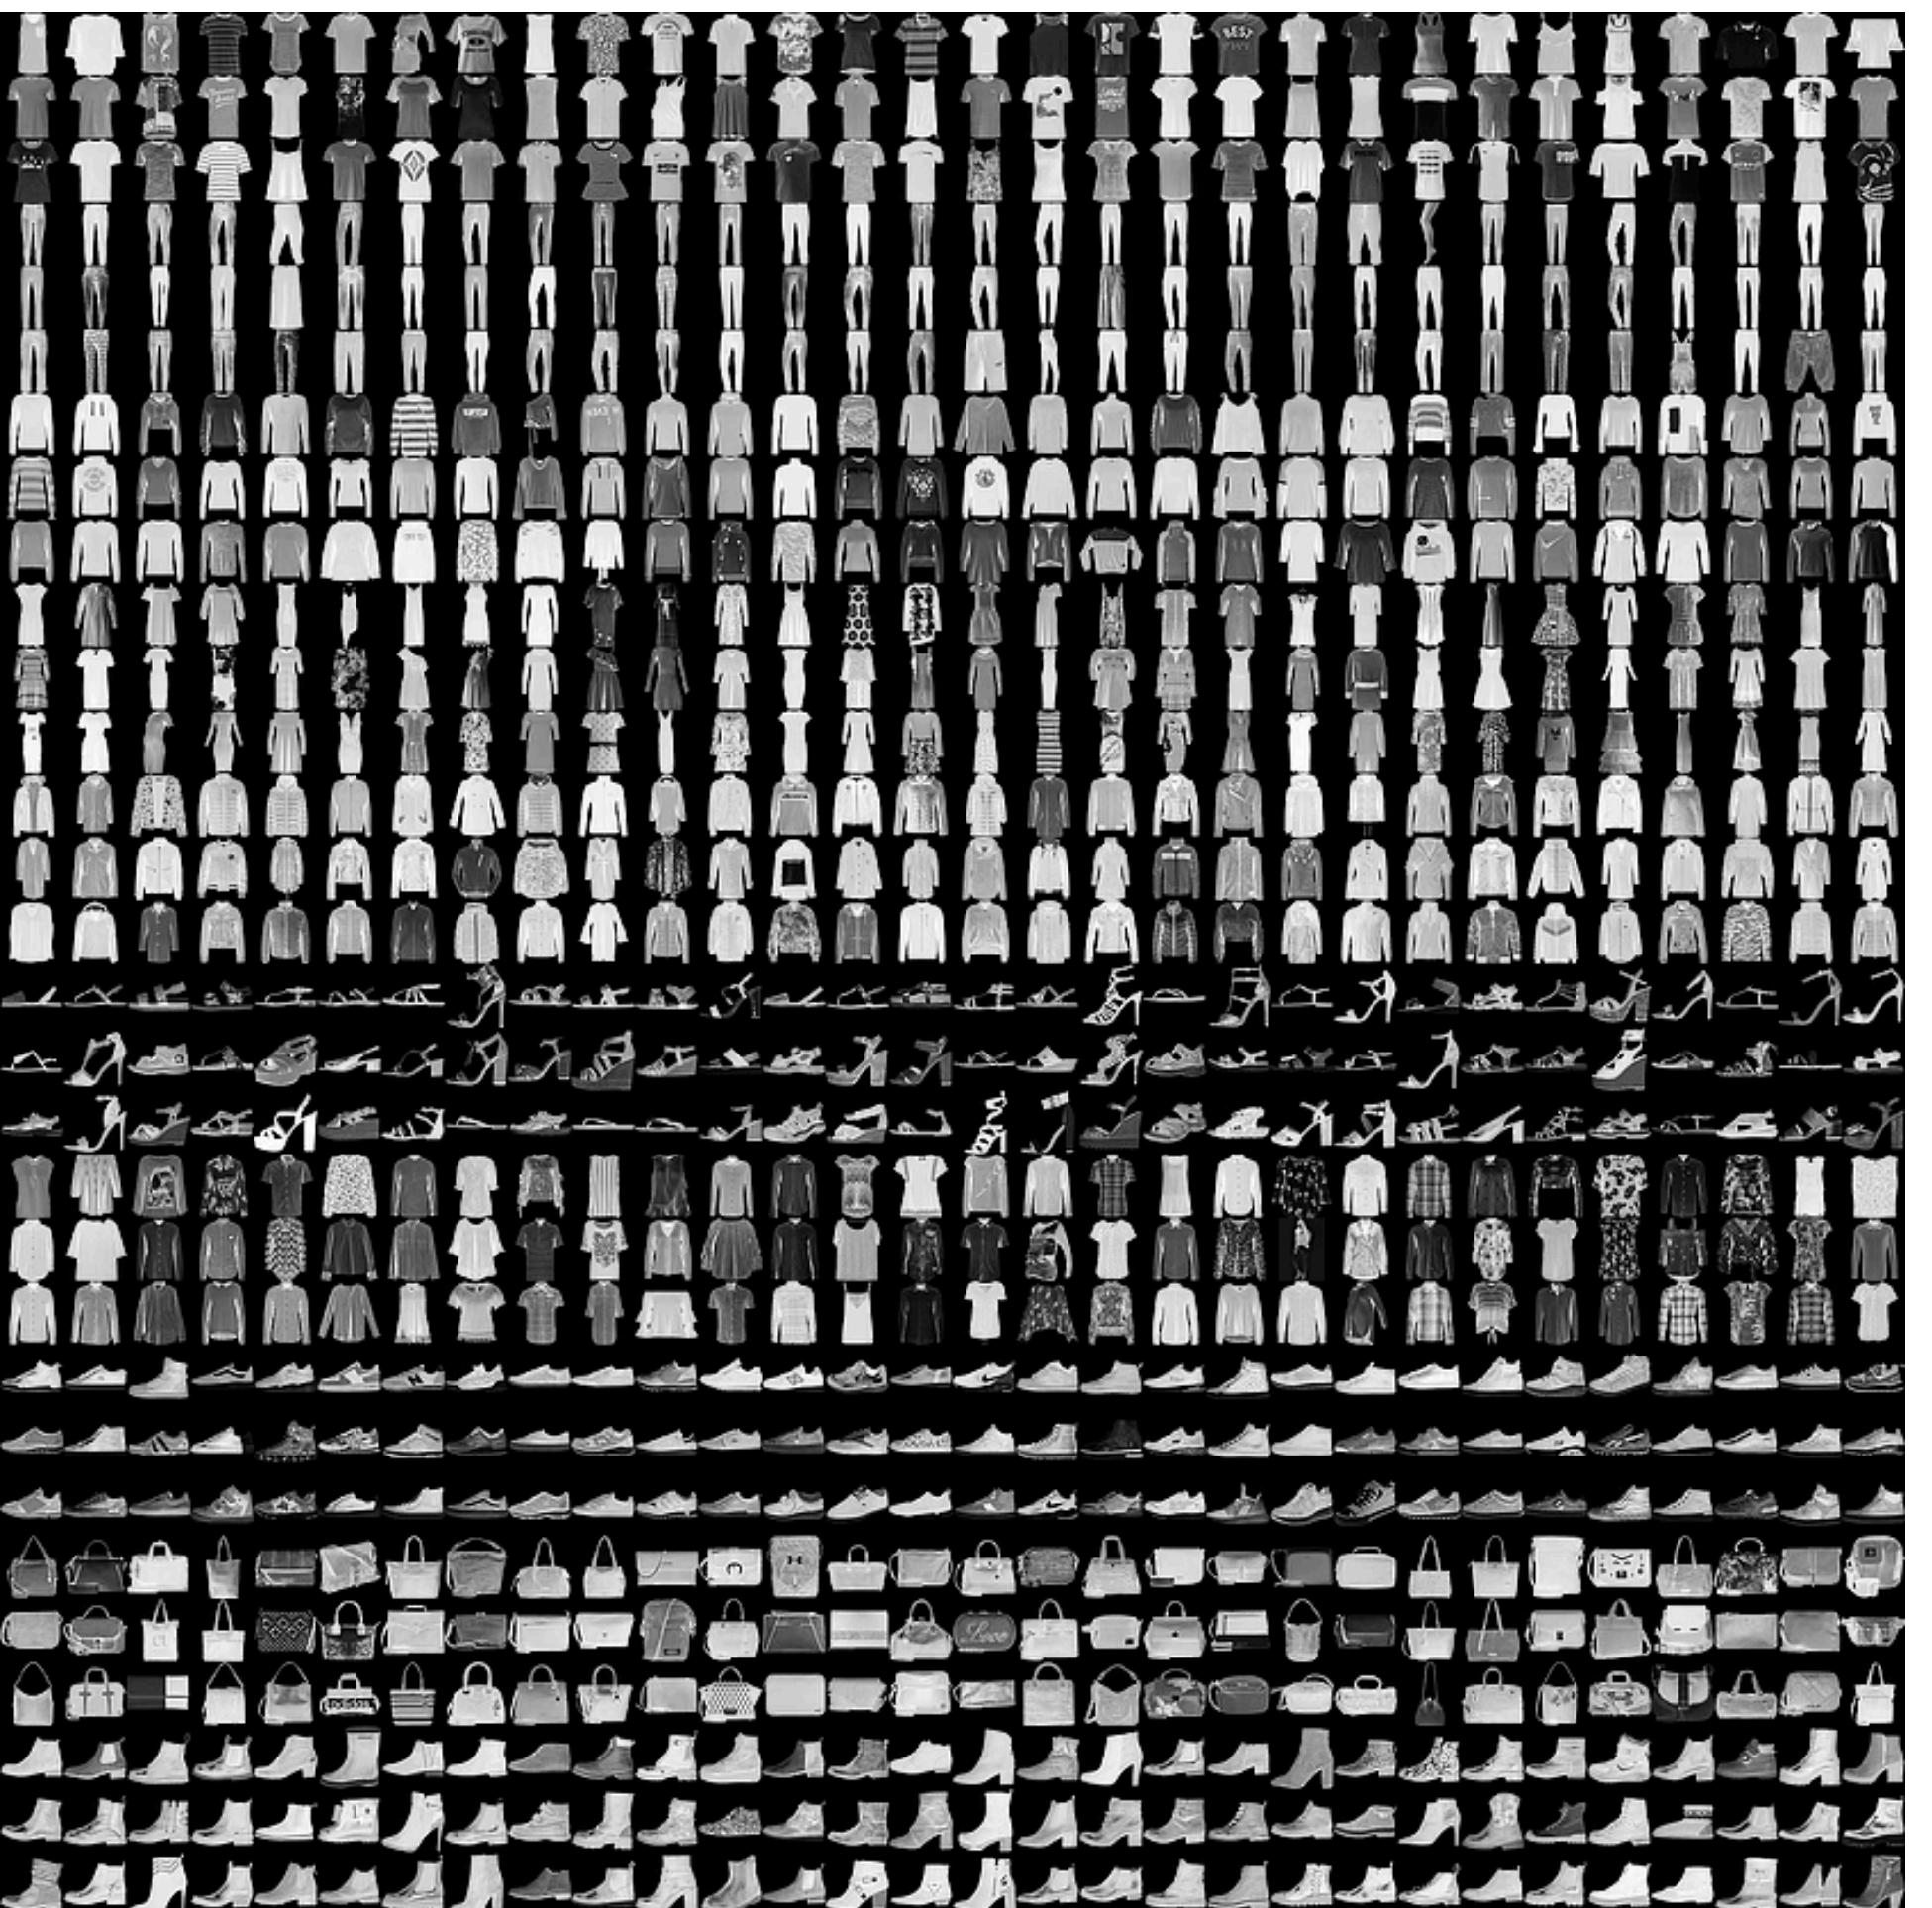
https://twitter.com/momiji_fullmoon/status/1110316960611368960

学習率 η は歩幅のイメージ (η 小=すり足、 η 大=巨人の一歩)

Fashion MNIST Dataset

- 7万画像
- 10カテゴリ
- 28×28 pixels
- 実験用データセット

ラベル	記述
0	T-shirt/top
1	Trouser
2	Pullover
3	Dress
4	Coat
5	Sandal
6	Shirt
7	Sneaker
8	Bag
9	Ankle boot

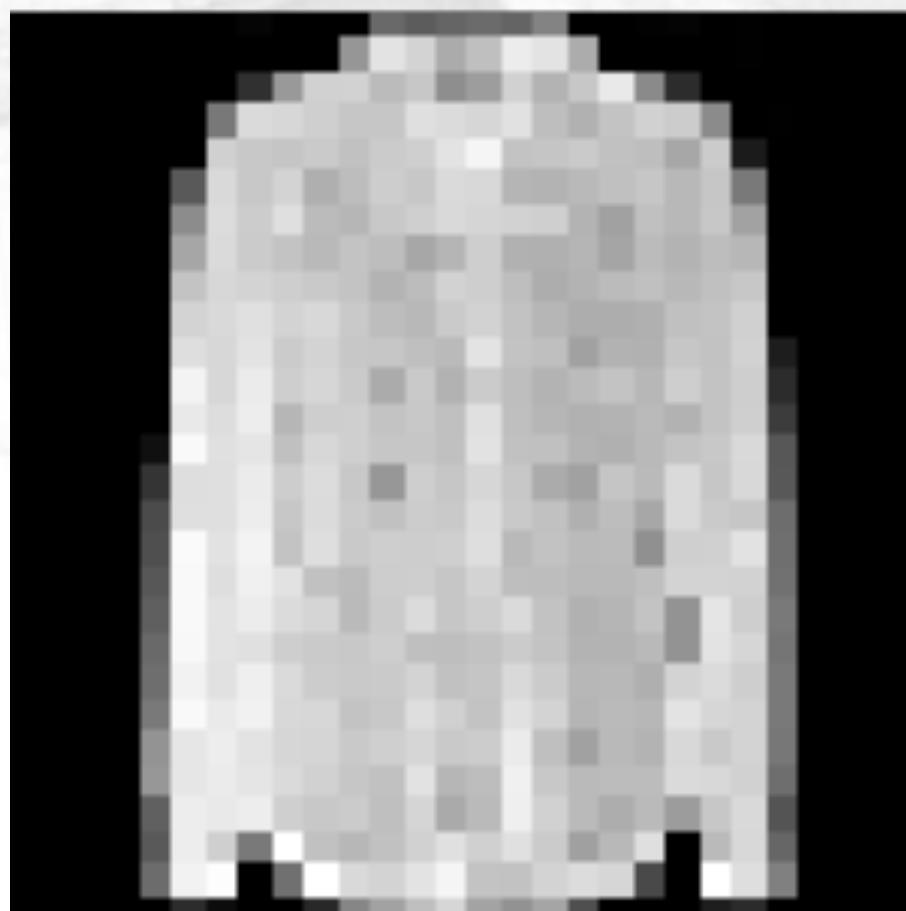


<https://github.com/zalandoresearch/fashion-mnist>

```
1 import tensorflow as tf
2 import matplotlib.pyplot as plt
3
4 fashion_mnist = tf.keras.datasets.fashion_mnist
5 (train_images, train_labels), (test_images, test_labels) = fashion_mnist.load_data()
6 train_images = train_images / 255.
7 test_images = test_images / 255.
```

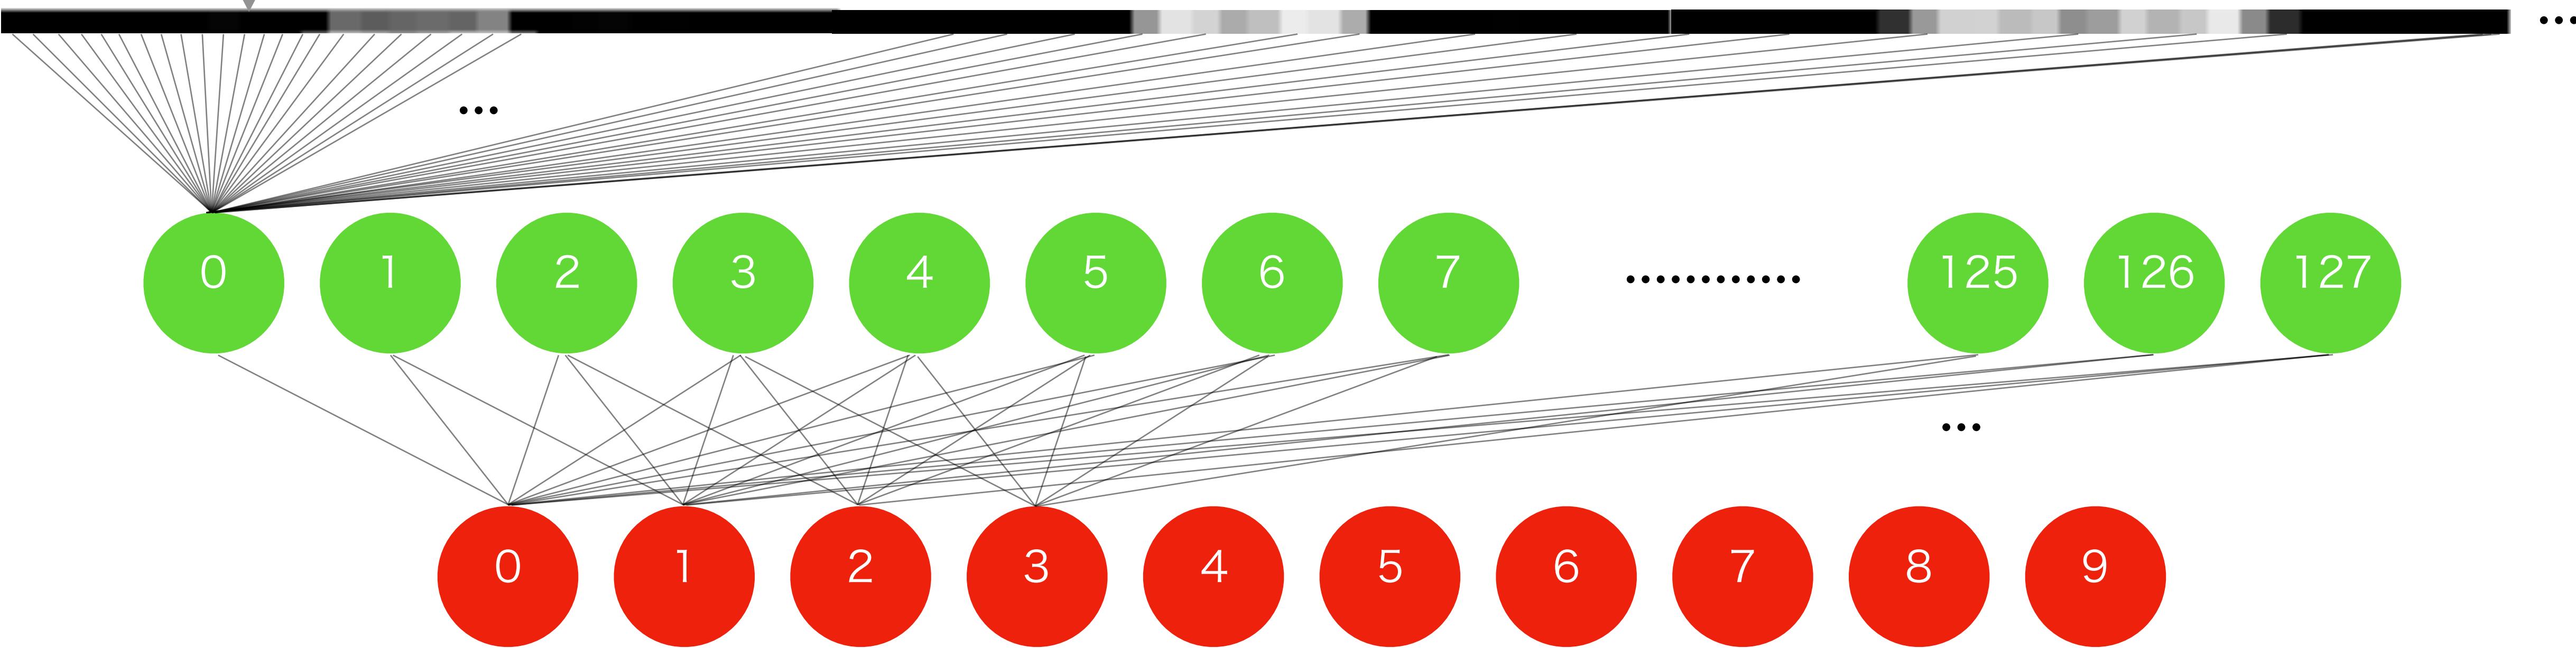
```
1 import tensorflow as tf
2 import matplotlib.pyplot as plt
3
4 fashion_mnist = tf.keras.datasets.fashion_mnist
5 (train_images, train_labels), (test_images, test_labels) = fashion_mnist.load_data()
6 train_images = train_images / 255.
7 test_images = test_images / 255.
8
9 model = tf.keras.Sequential([
10     tf.keras.layers.Flatten(input_shape=(28, 28)),
11     tf.keras.layers.Dense(128, activation=tf.nn.relu),
12     tf.keras.layers.Dense(10, activation=tf.nn.softmax)
13 ])
```

```
1 import tensorflow as tf
2 import matplotlib.pyplot as plt
3
4 fashion_mnist = tf.keras.datasets.fashion_mnist
5 (train_images, train_labels), (test_images, test_labels) = fashion_mnist.load_data()
6 train_images = train_images / 255.
7 test_images = test_images / 255.
8
9 model = tf.keras.Sequential([
10     tf.keras.layers.Flatten(input_shape=(28, 28)),
11     tf.keras.layers.Dense(128, activation=tf.nn.relu),
12     tf.keras.layers.Dense(10, activation=tf.nn.softmax)
13 ])
14
15 model.compile(optimizer=tf.train.AdamOptimizer(),
16                 loss='sparse_categorical_crossentropy', metrics=['accuracy'])
16 model.fit(train_images, train_labels, epochs=5, batch_size=128, verbose=1,
17             validation_data=(test_images, test_labels))
```



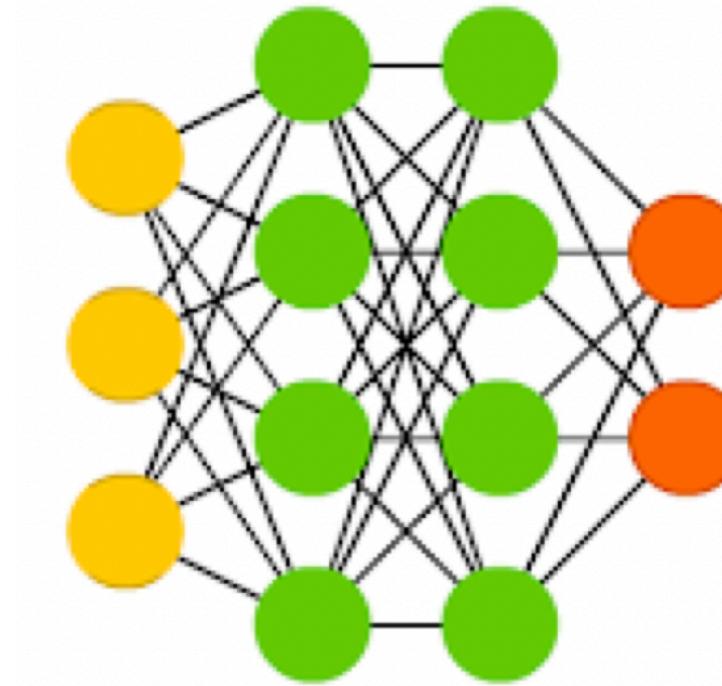
```
9 model = tf.keras.Sequential([
10     tf.keras.layers.Flatten(input_shape=(28, 28)),
11     tf.keras.layers.Dense(128, activation=tf.nn.relu),
12     tf.keras.layers.Dense(10, activation=tf.nn.softmax)
13 ])
```

Flatten: (28, 28) => (784)



Dense Layer の欠点

- 入力のベクトルの**全要素**の相関をみている
 - > 住宅価格予測みたいな話ならまだいい
 - > もう作ってる特徴量と特徴量の組み合わせ
 - 例) 東京墨田区 & 床面積 30m² & 1K & 風呂トイレ別 & 新築 => 家賃月10万円
- 「**画像の特徴量**」を抽出してからDense Layerに渡せば効率的

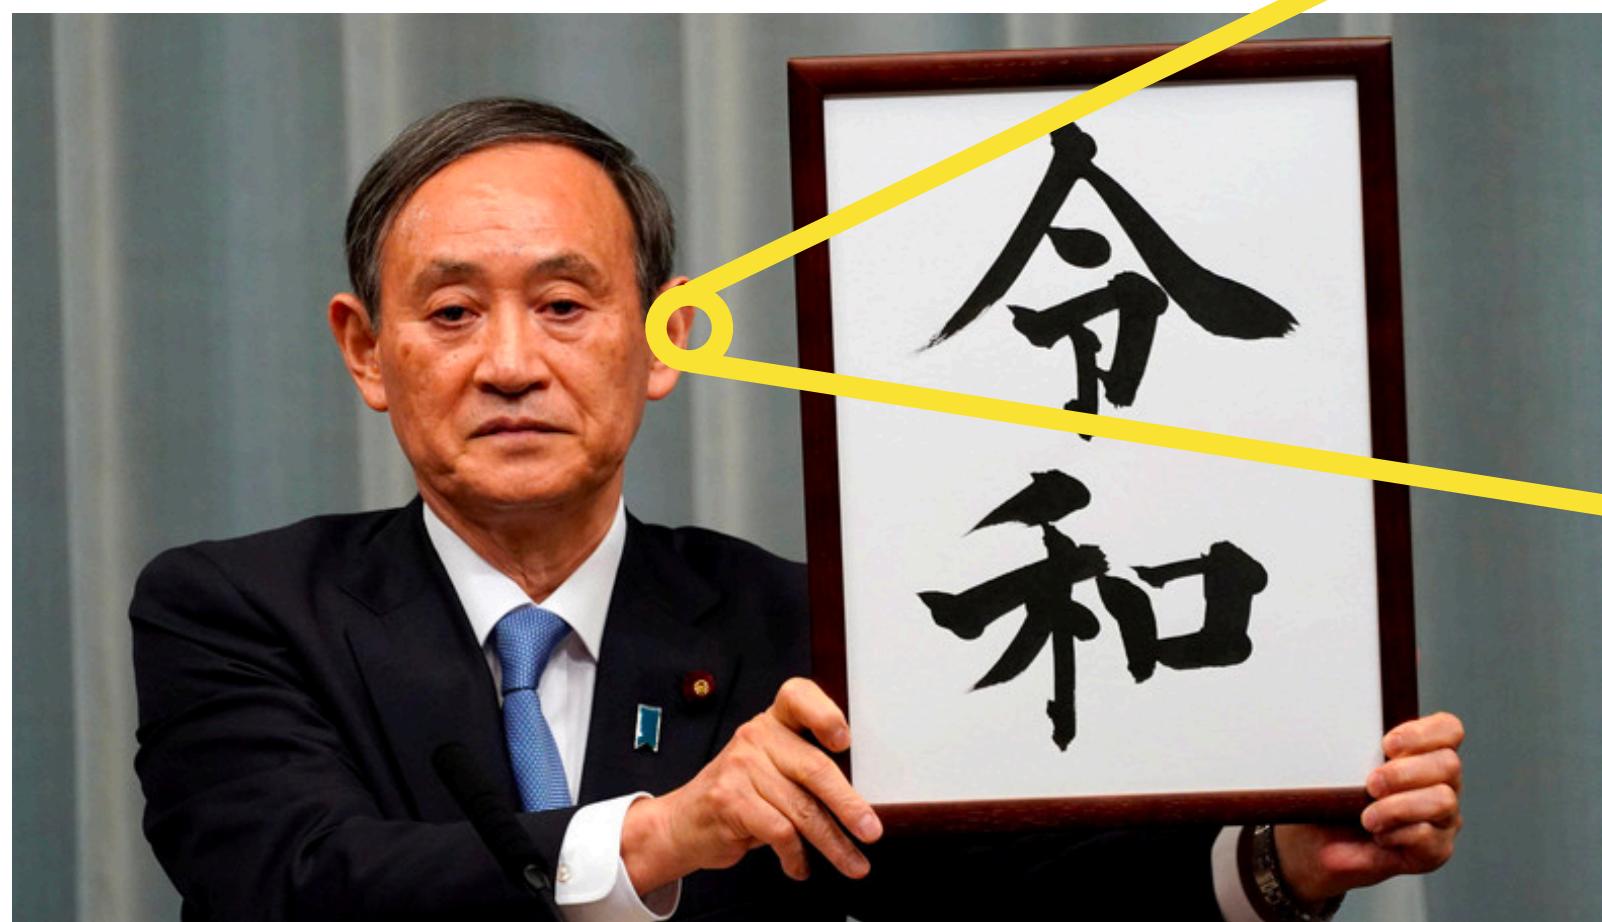


畳込みニューラルネットワーク (Convolutional Neural Network; CNN)



畳み込みニューラルネットワークの基礎

畳み込み (Convolution)



<https://news.yahoo.co.jp/byline/yuasamakoto/20190403-00120722/>

144	60	19
188	82	32
156	55	27

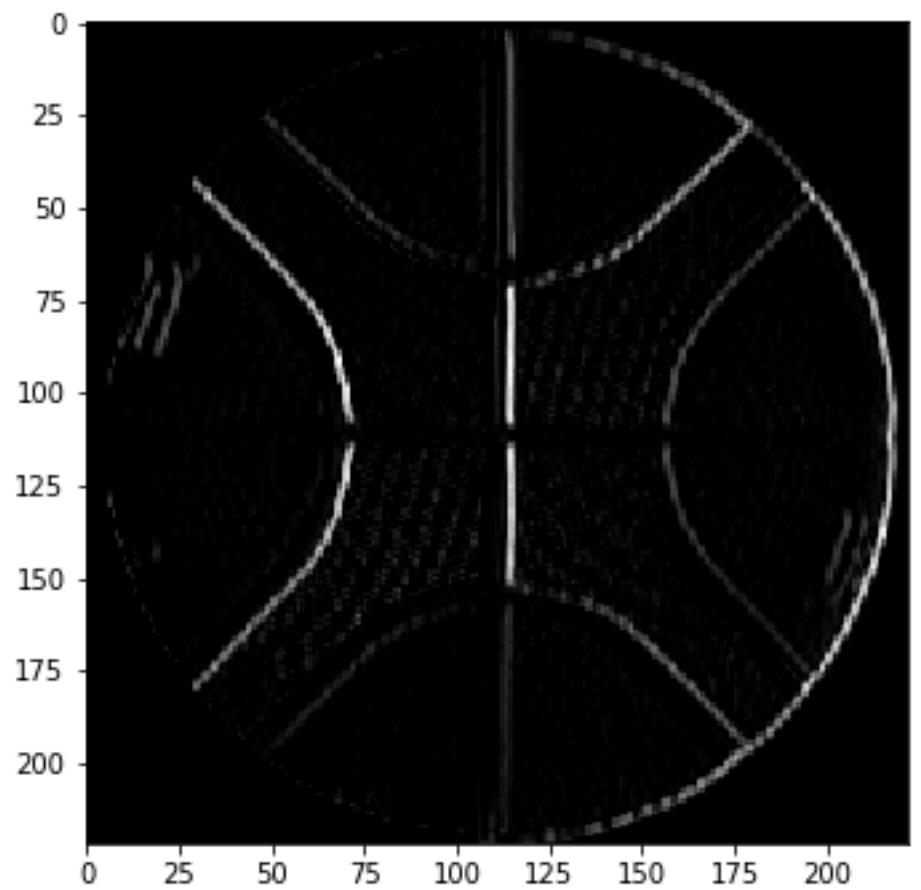
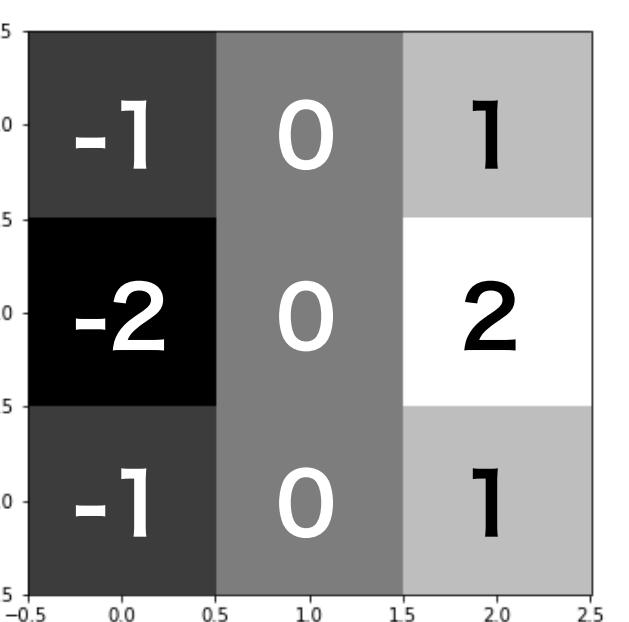
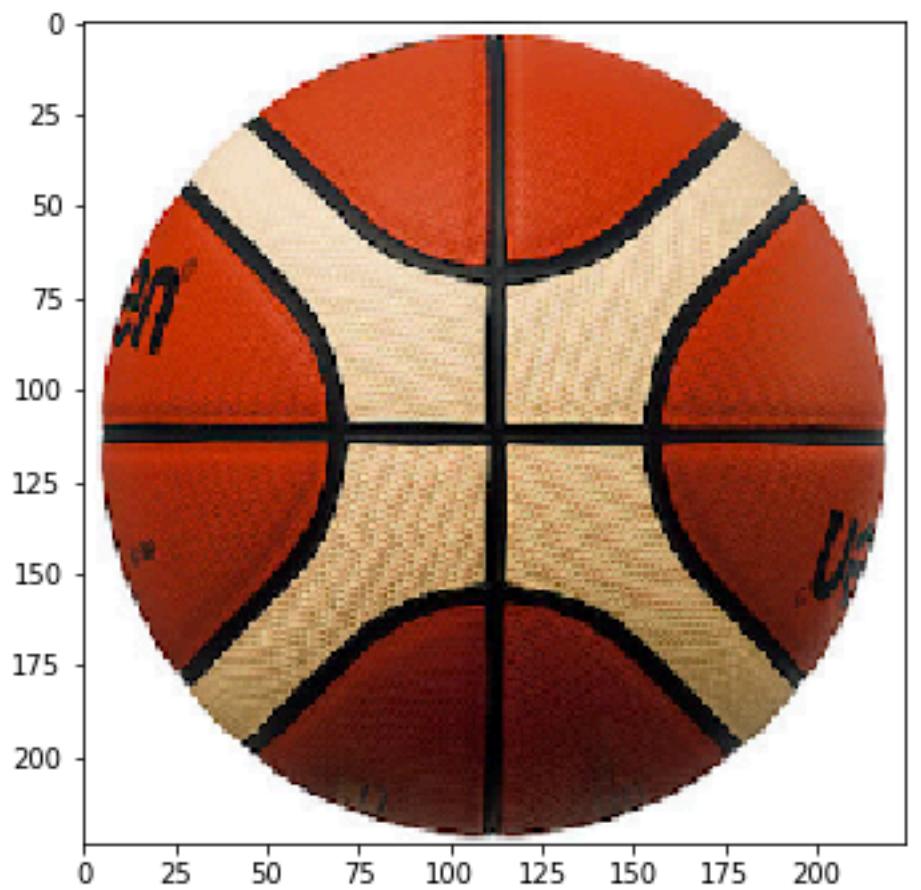
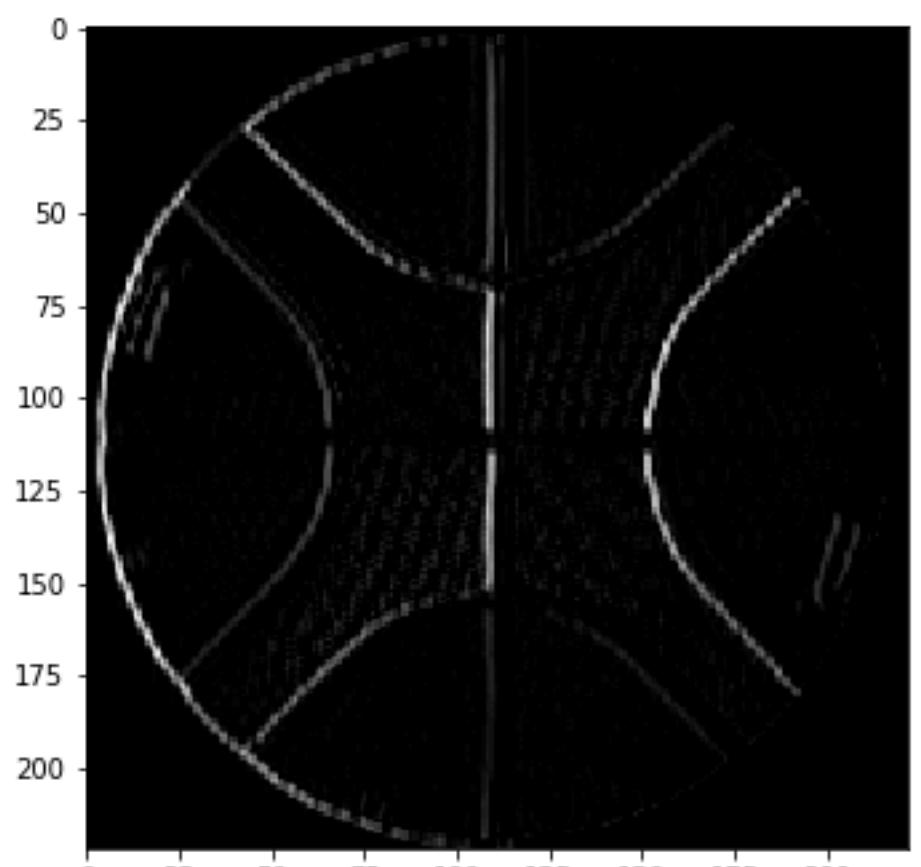
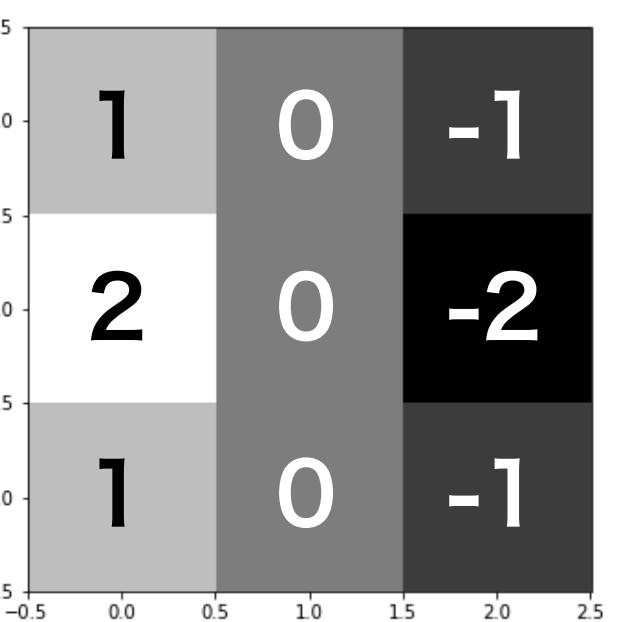
-1	0	-2
0.5	4.5	-1.5
1.5	2	-3

CURRENT_PIXEL_VALUE = 82

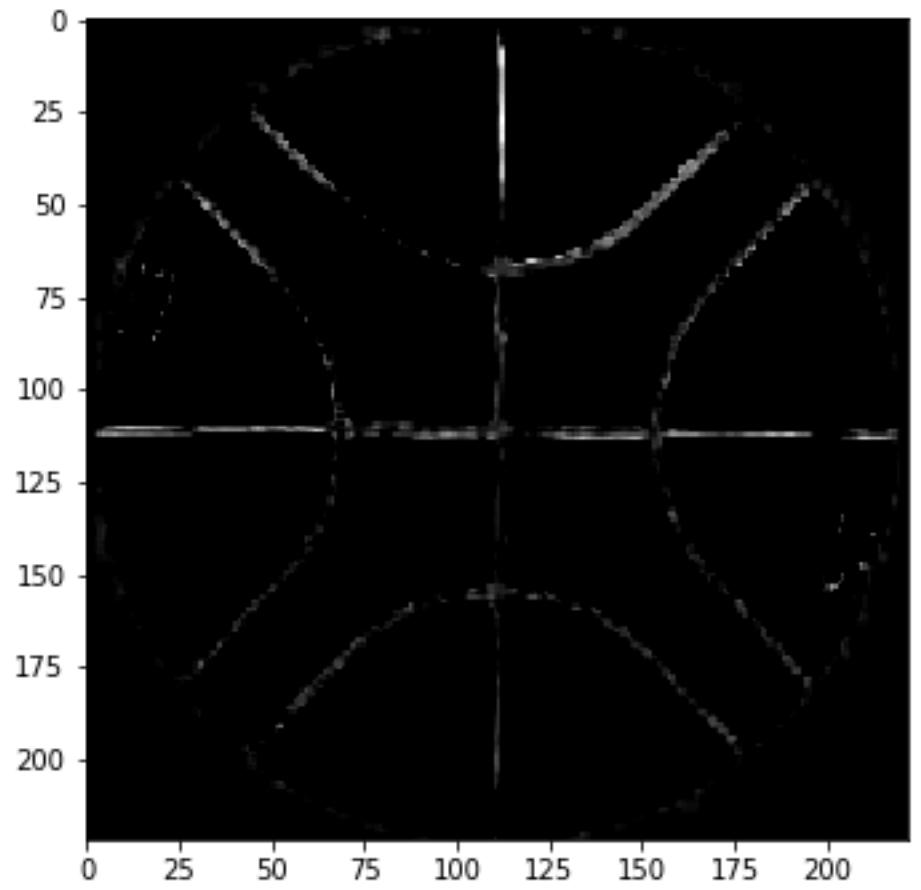
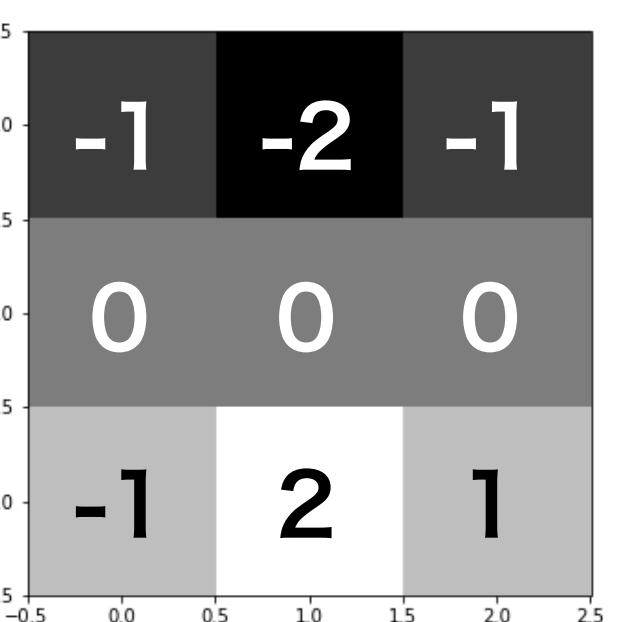
NEW_PIXEL_VALUE =

$$\begin{aligned} & (-1 * 144) + (0 * 60) + (-2 * 19) \\ & + (0.5 * 188) + (4.5 * 82) + (-1.5 * 32) \\ & + (1.5 * 156) + (2 * 55) + (-3 * 27) \end{aligned}$$

$$u_{ijm} = \sum_{k=0}^{K-1} \sum_{p=0}^{W-1} \sum_{q=0}^{H-1} z_{i+p, j+q, k}^{(l-1)} h_{pqkm} + b_{ijm}$$



https://www.bbkong.net/fs/alleyoop/molten_BGL7



stride=(1, 1), padding='valid'



stride=(1, 1), padding='valid'



stride=(1, 1), padding='valid'



stride=(1, 1), padding='valid'



stride=(1, 1), padding='valid'



stride=(1, 1), padding='valid'



stride=(1, 1), padding='valid'



stride=(2, 2), padding='valid'



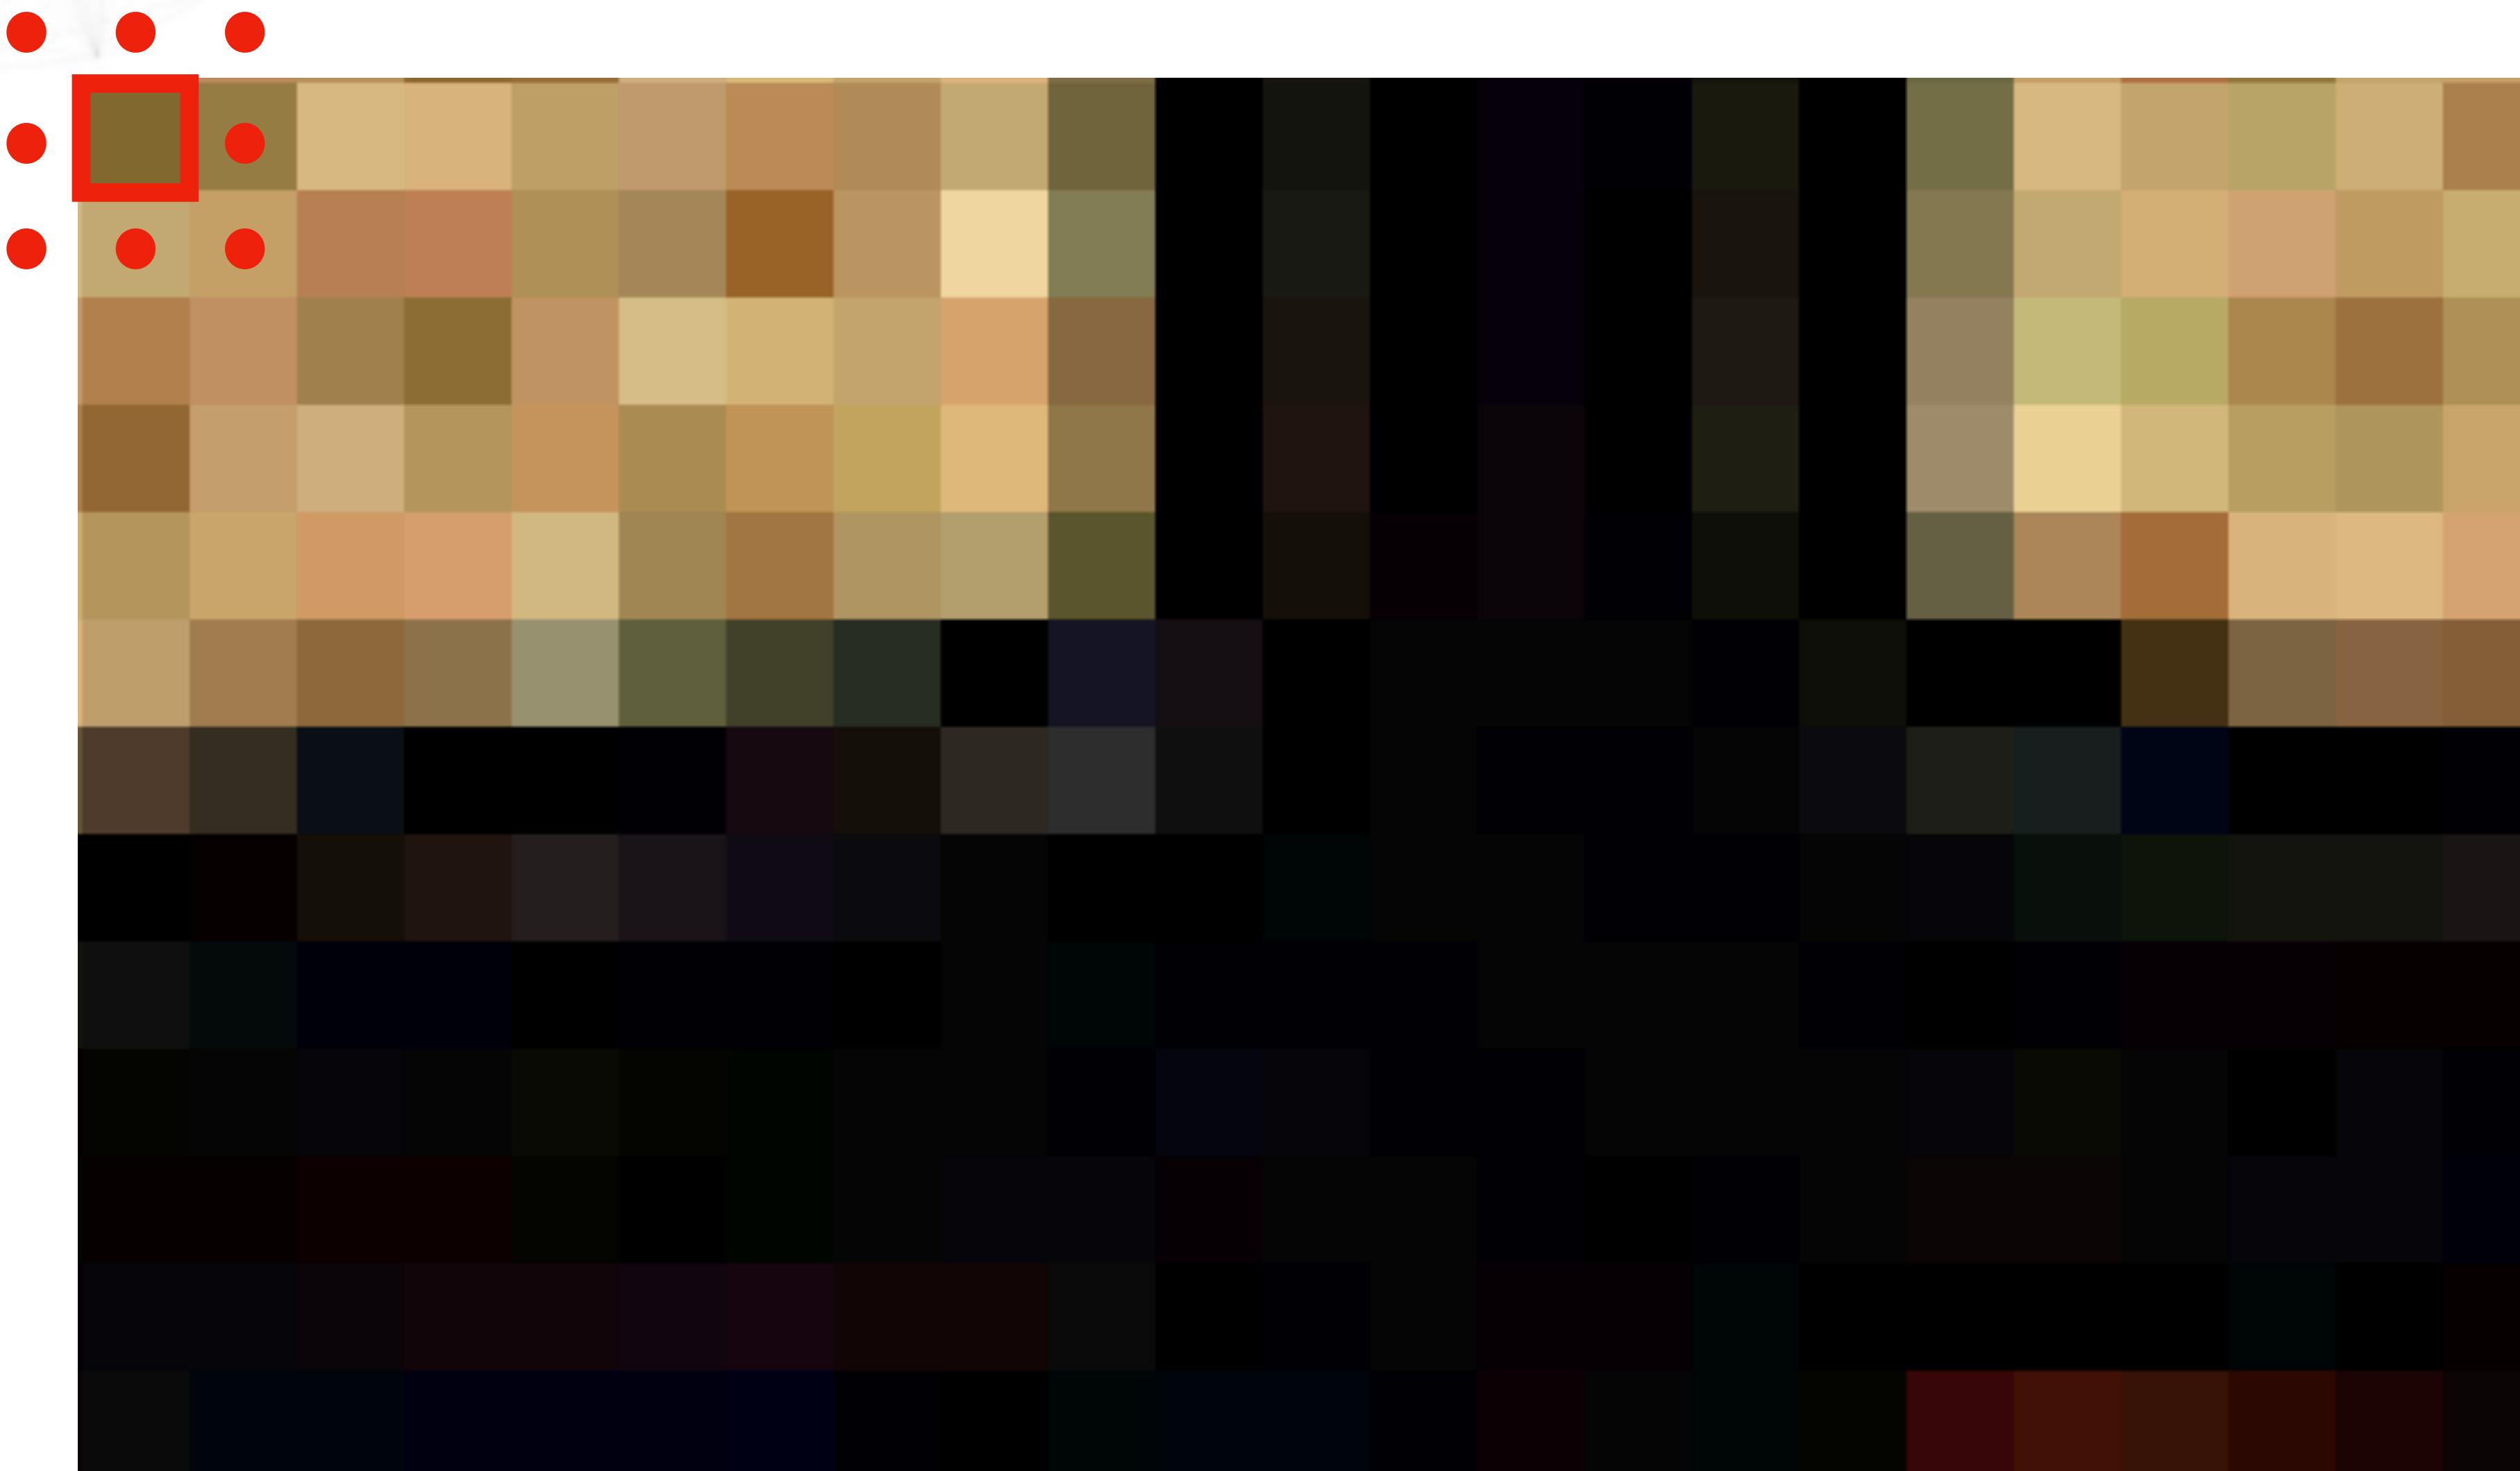
stride=(2, 2), padding='valid'



stride=(2, 2), padding='valid'



stride=(1, 1), padding='same'



stride=(1, 1), padding='same'



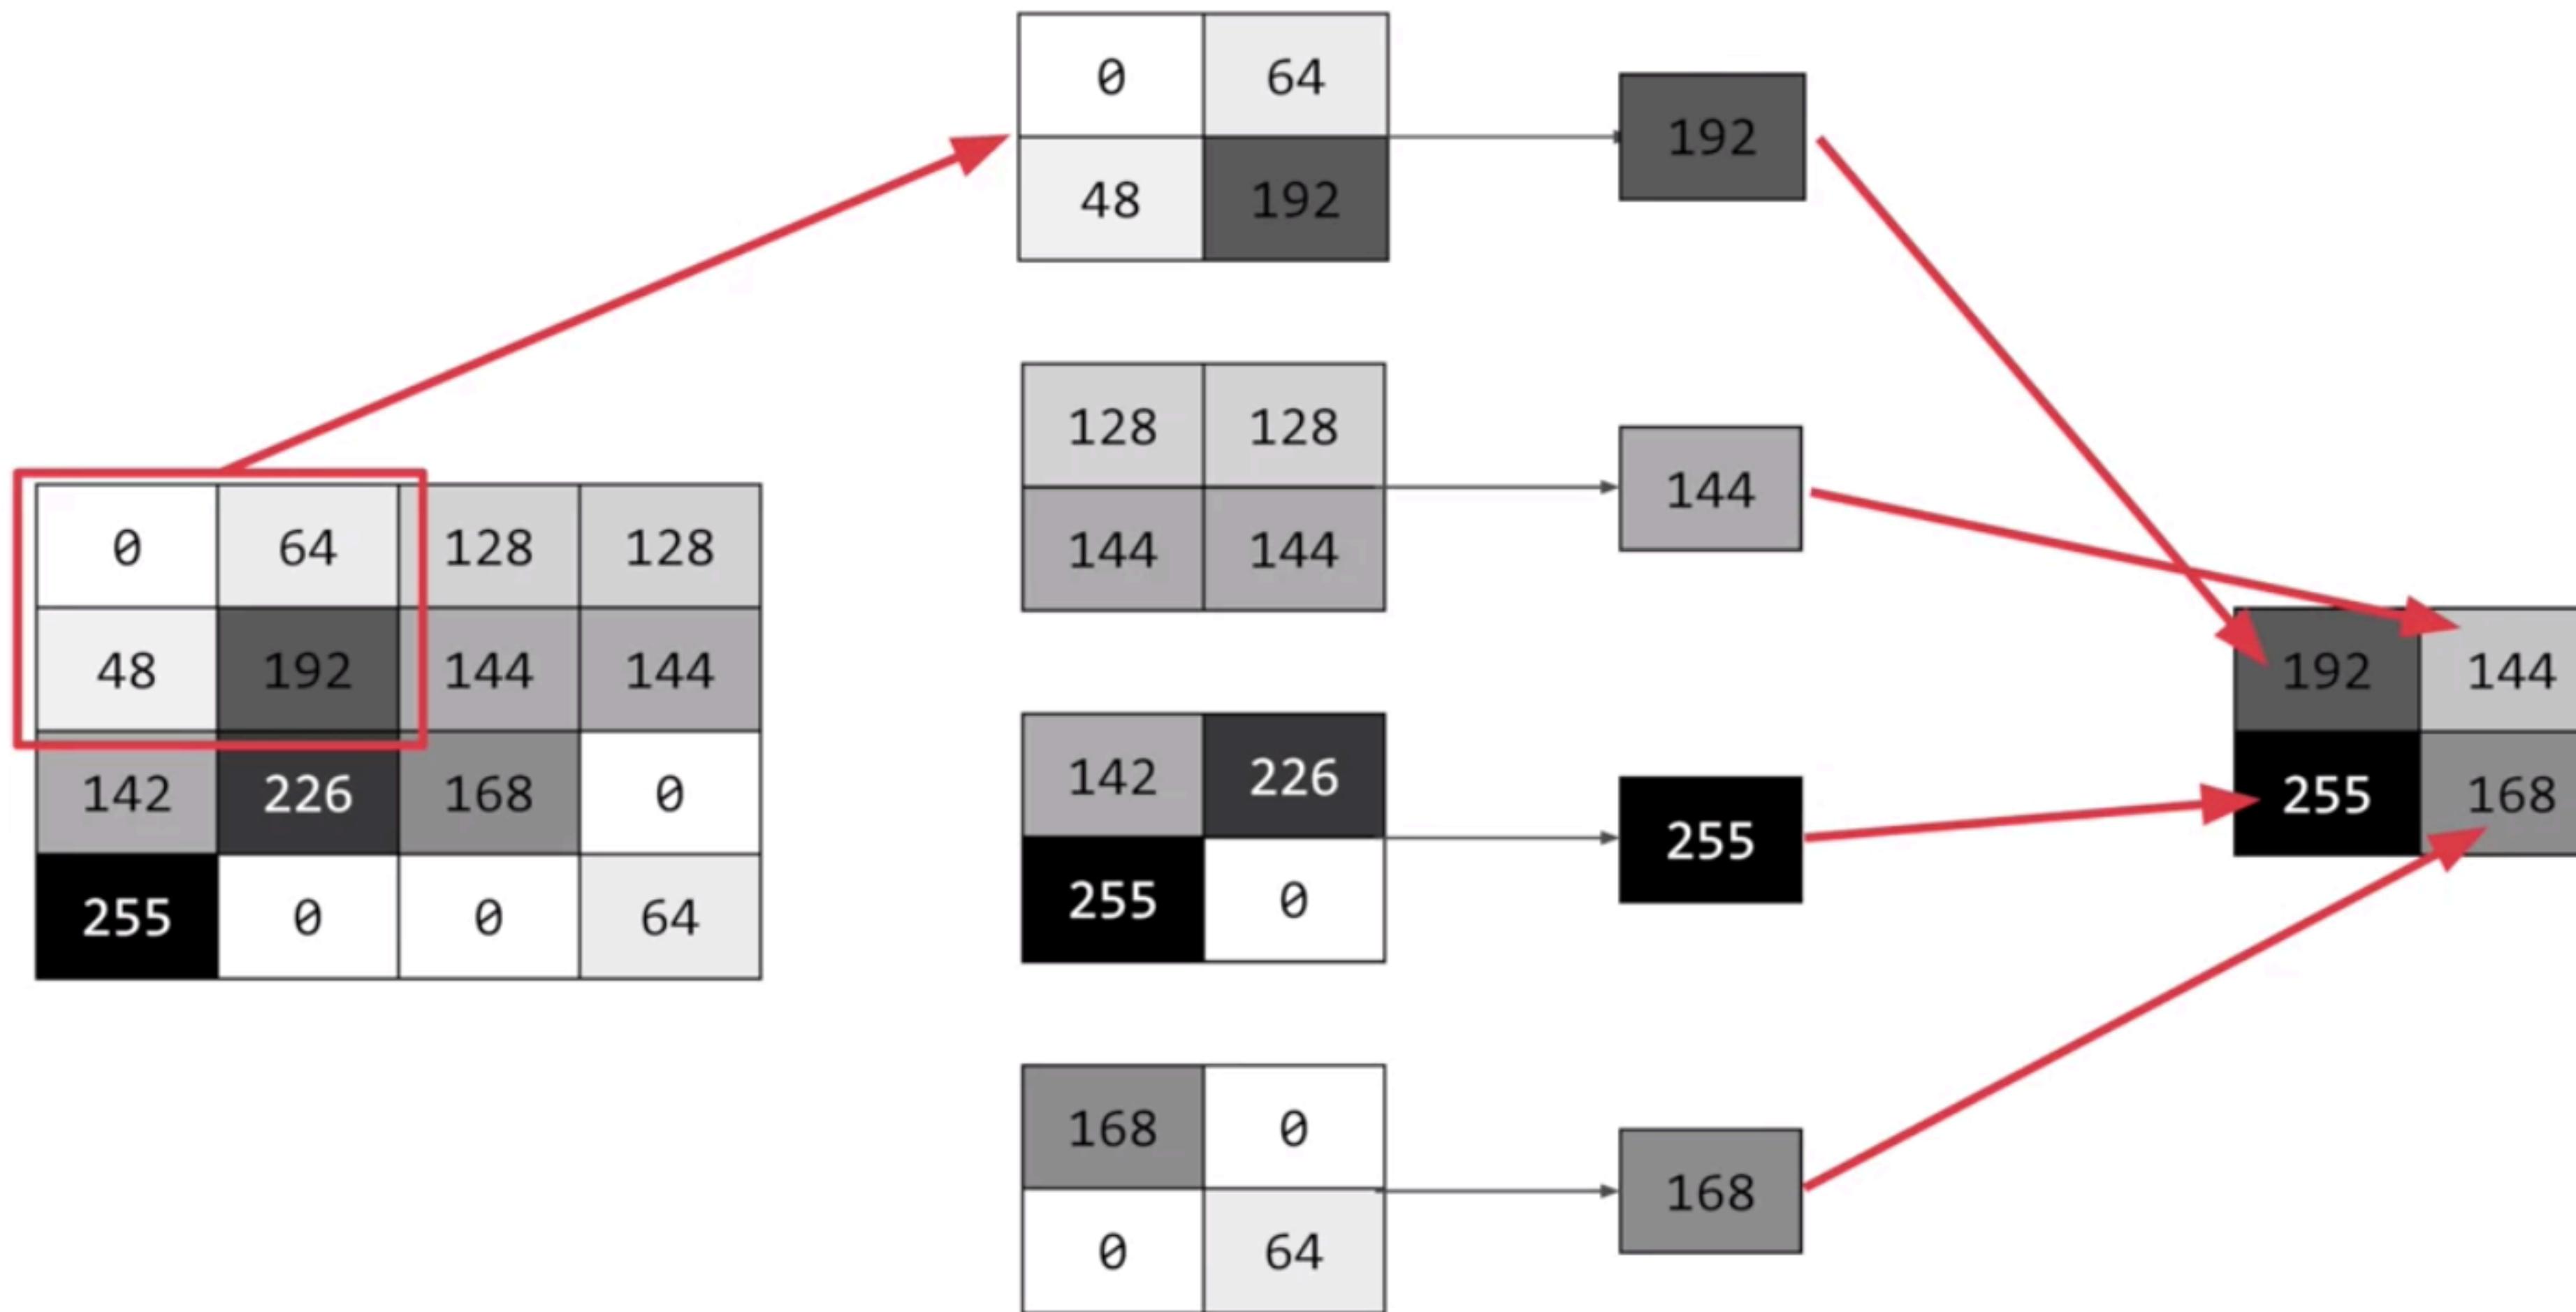
stride=(1, 1), padding='same'



stride=(1, 1), padding='same'

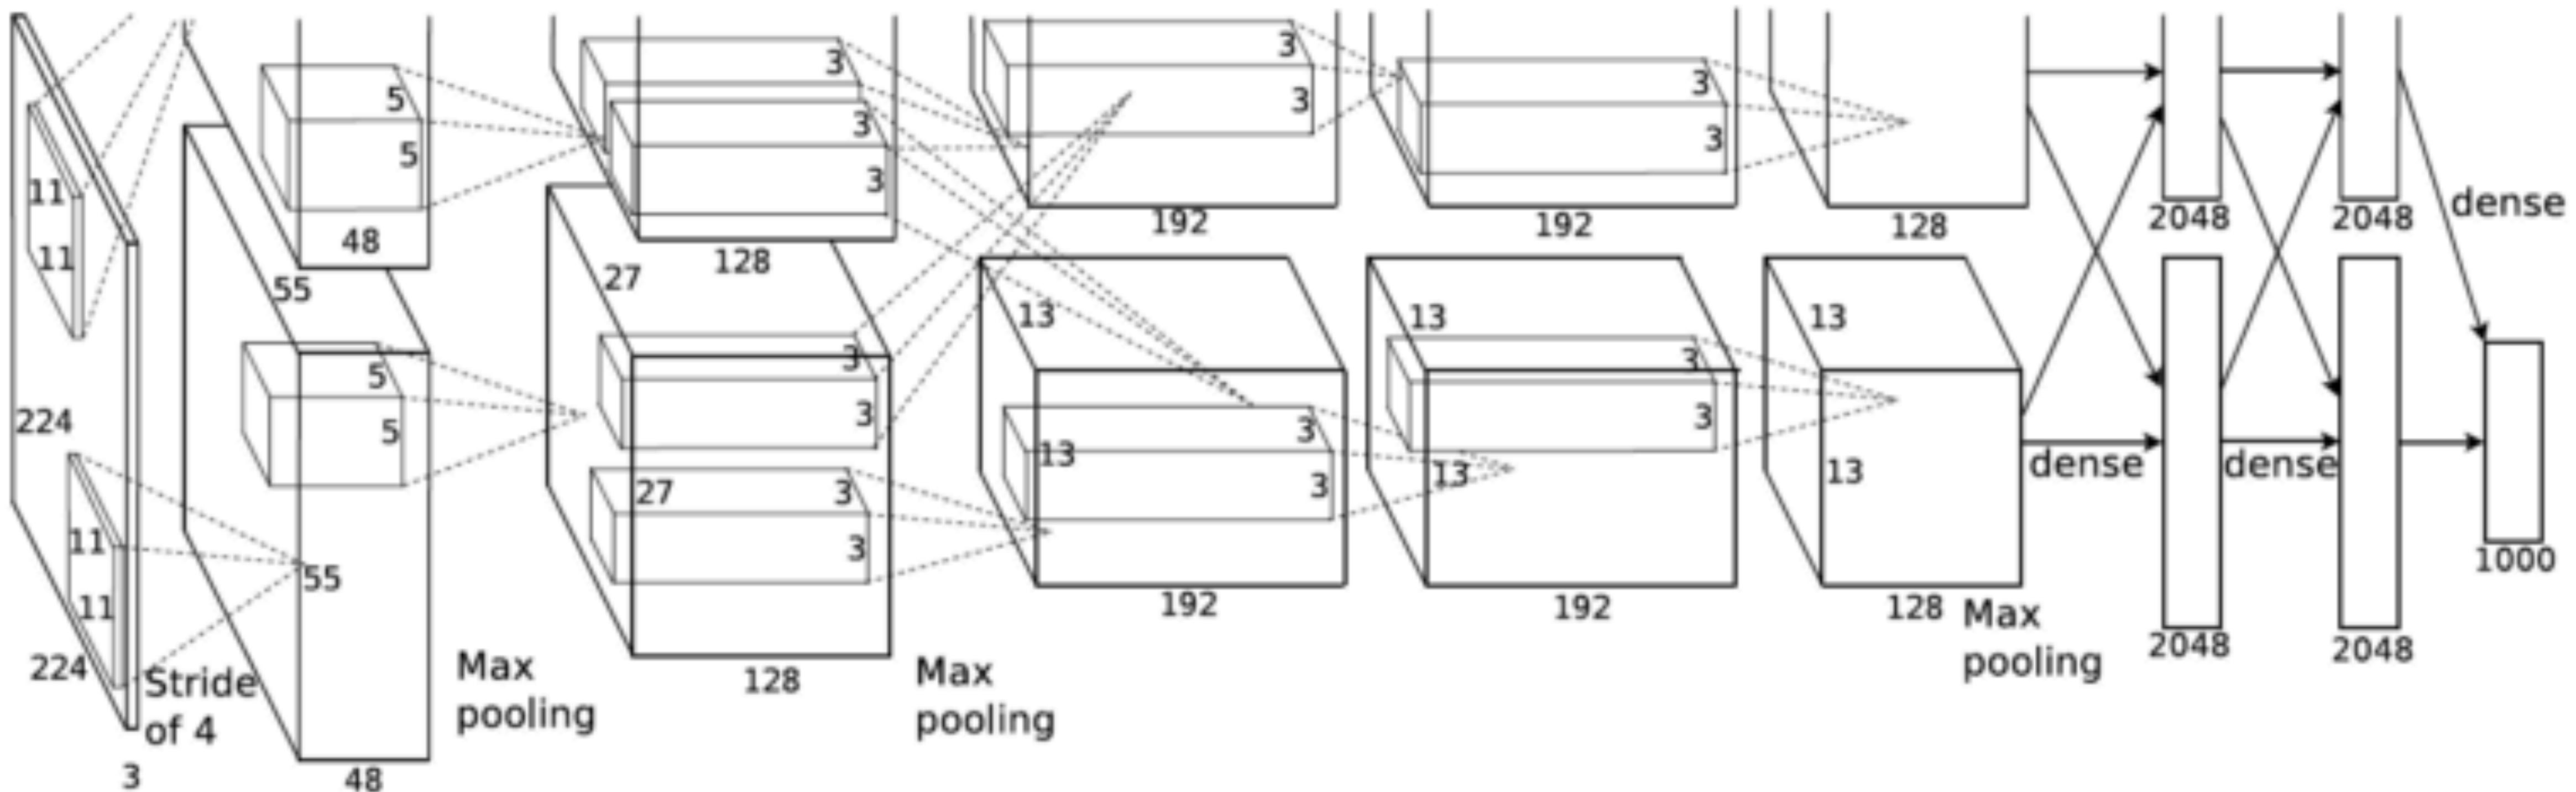


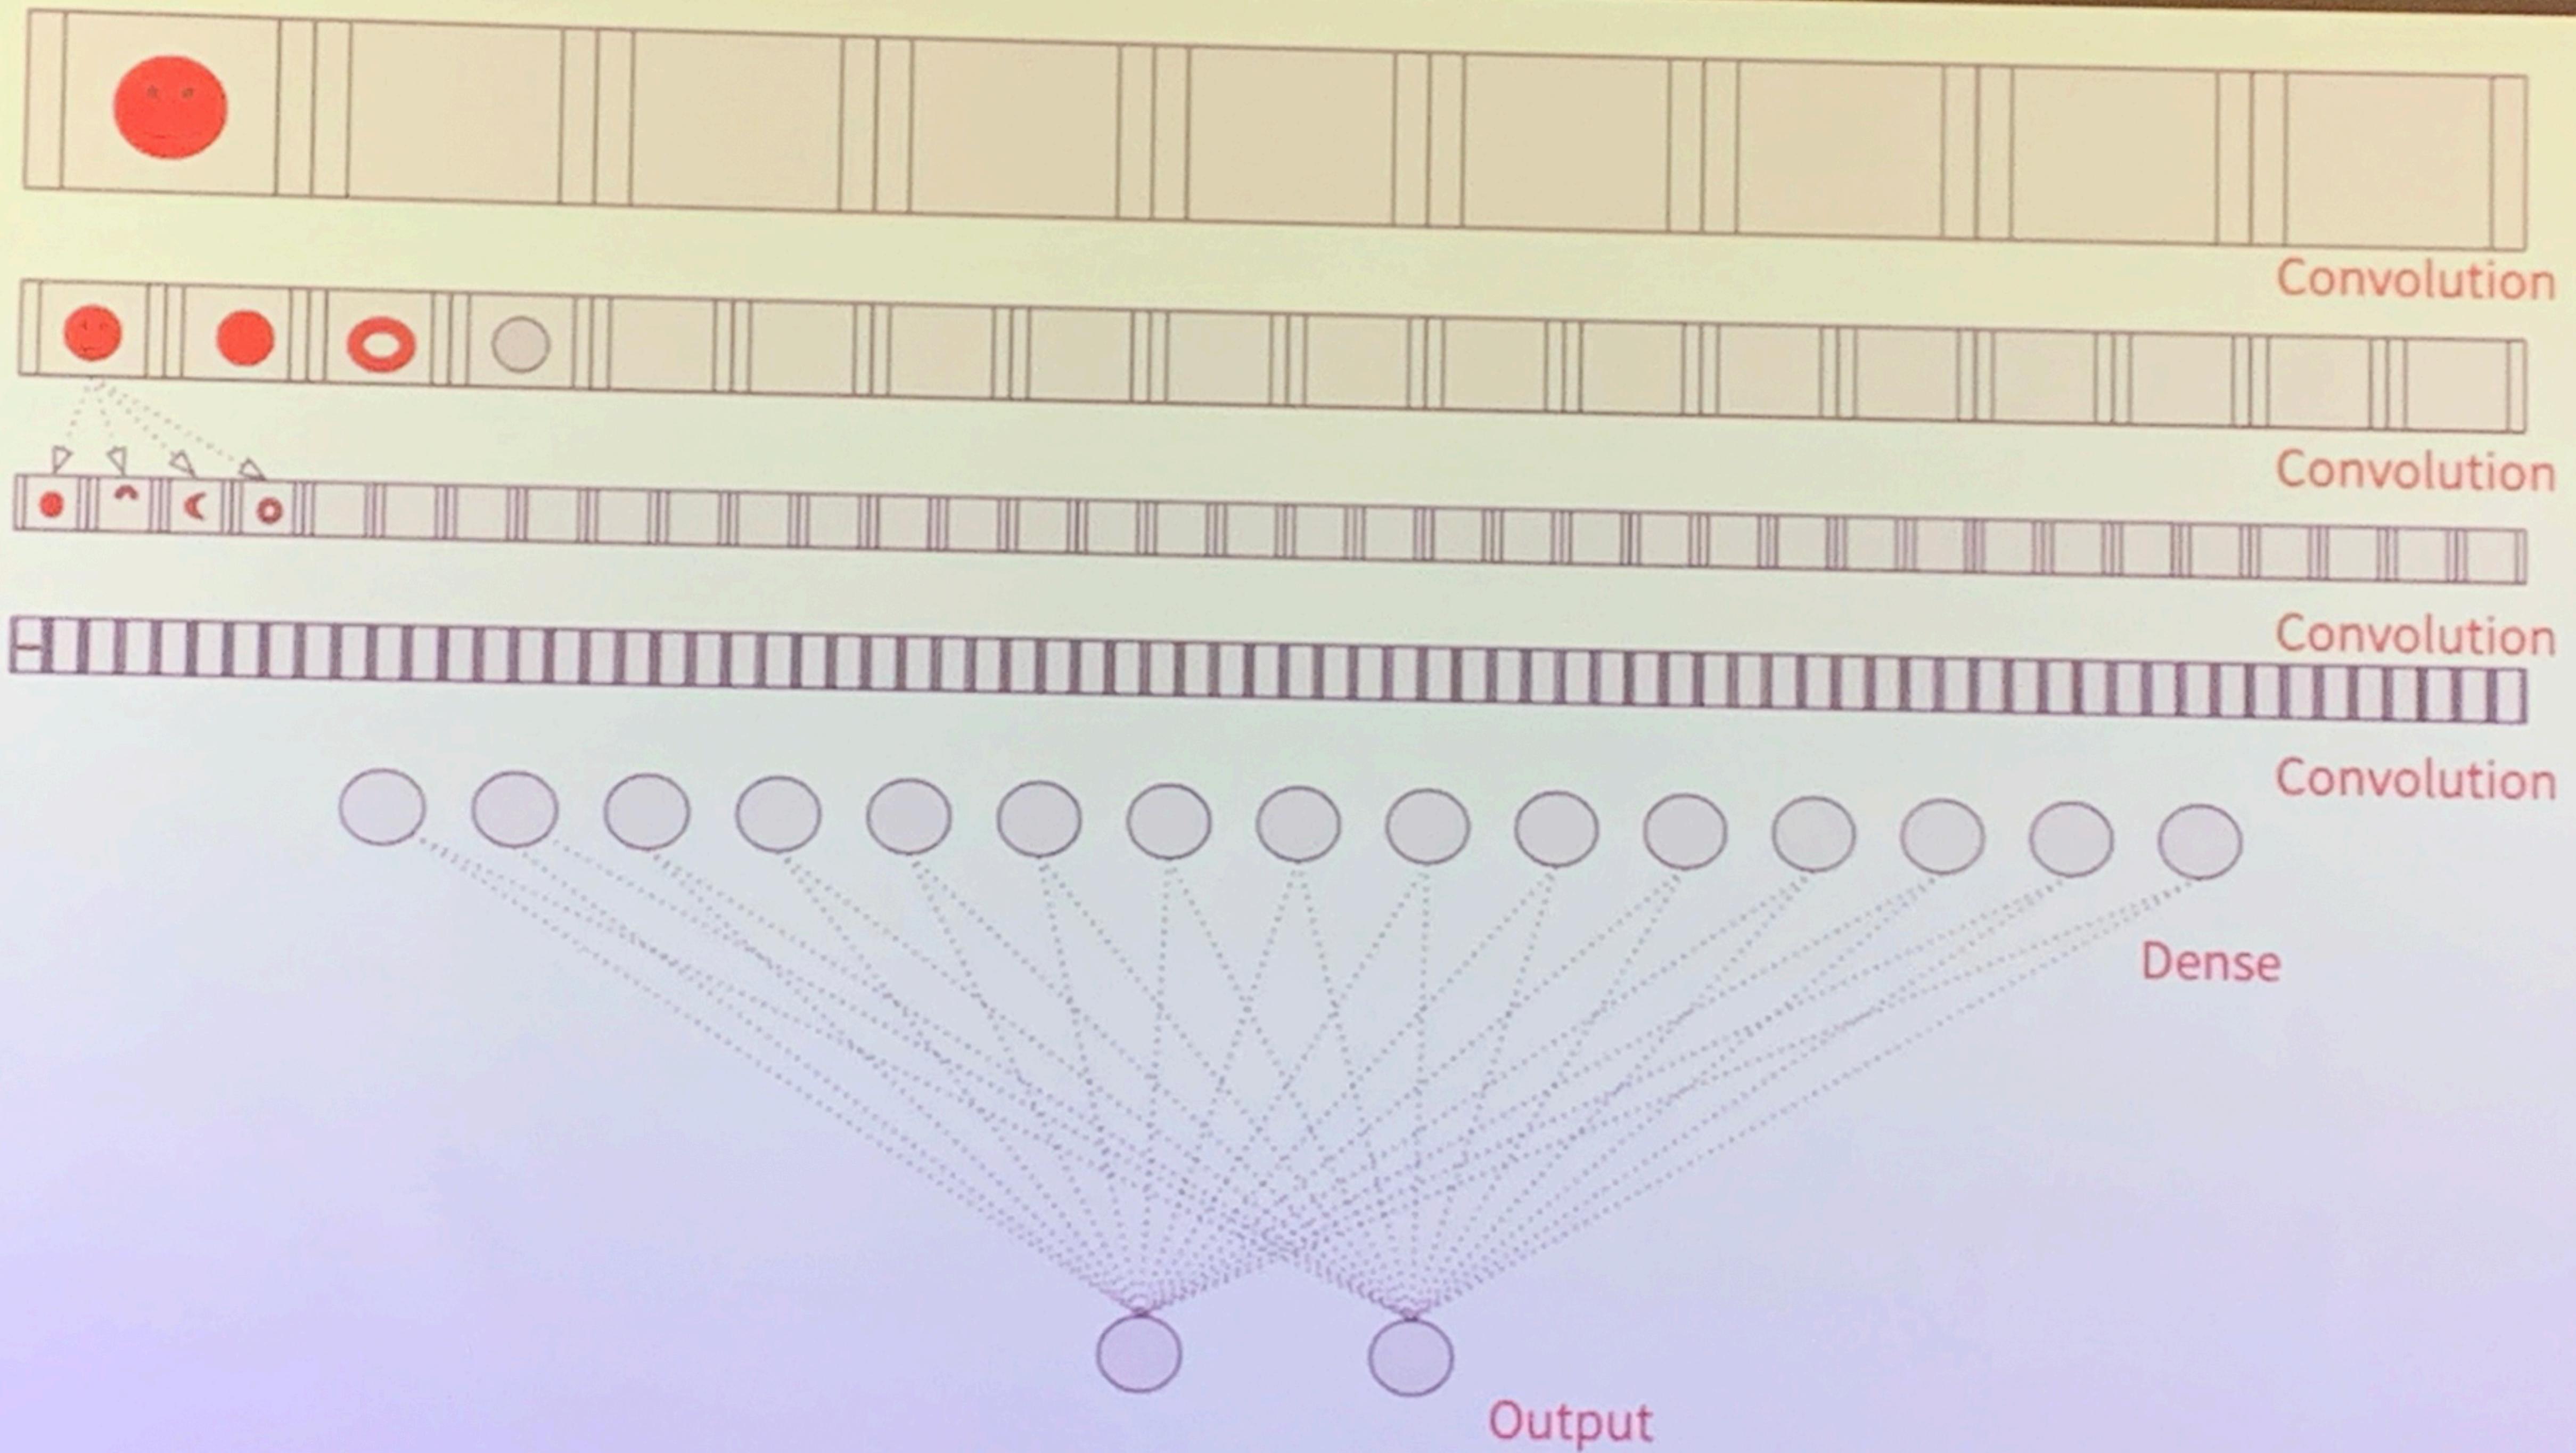
Max Pooling

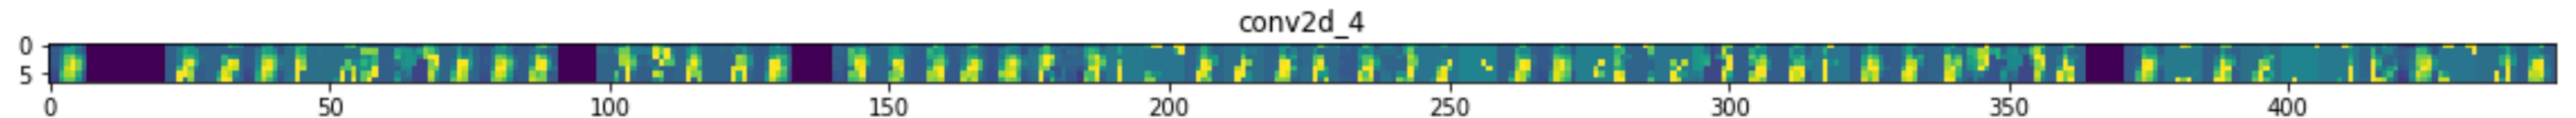
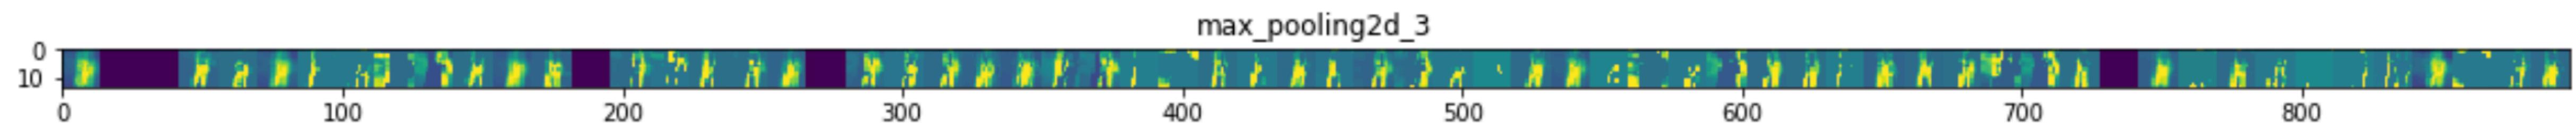
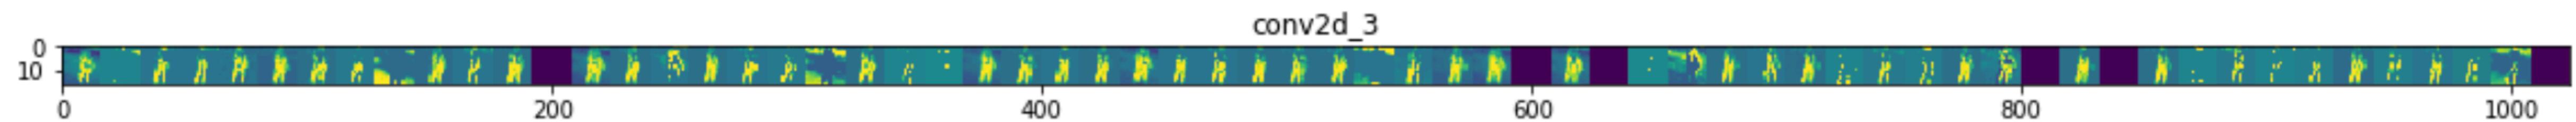
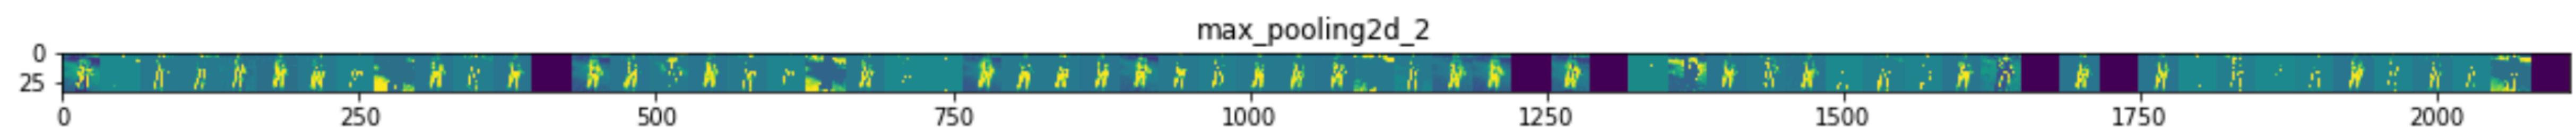
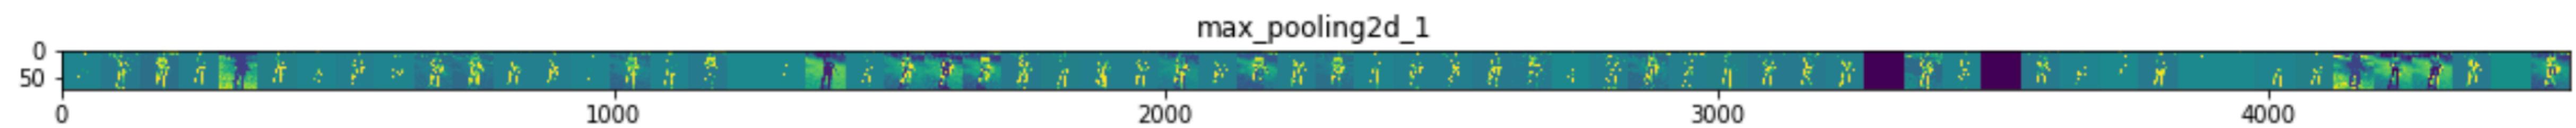
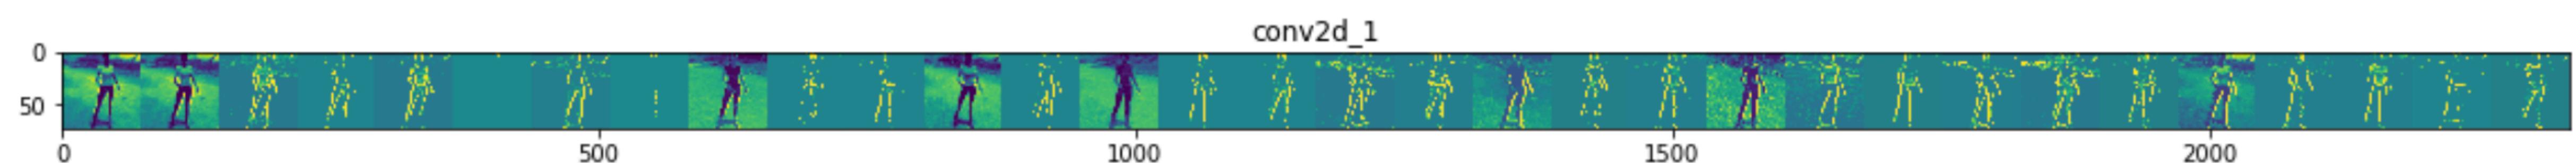
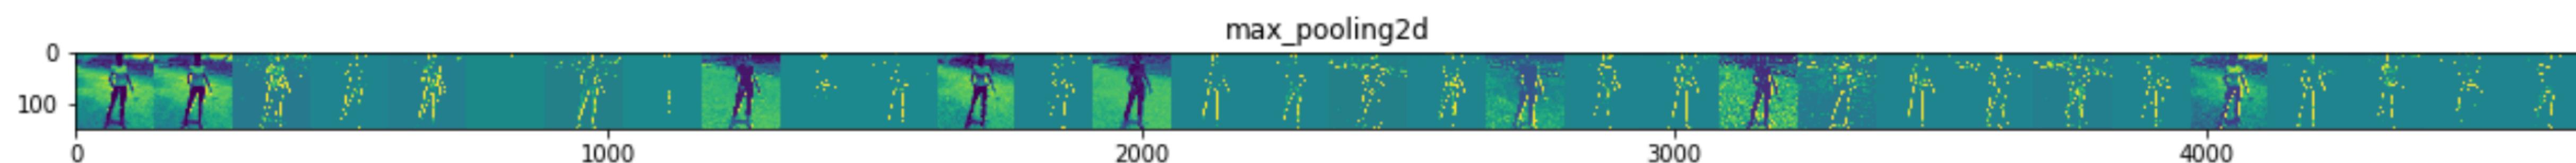
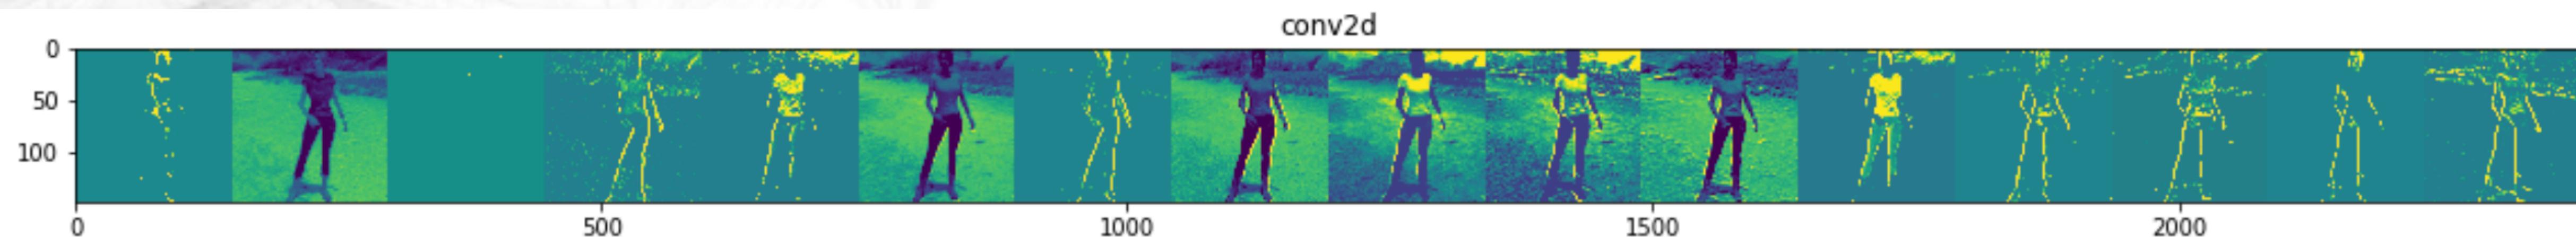


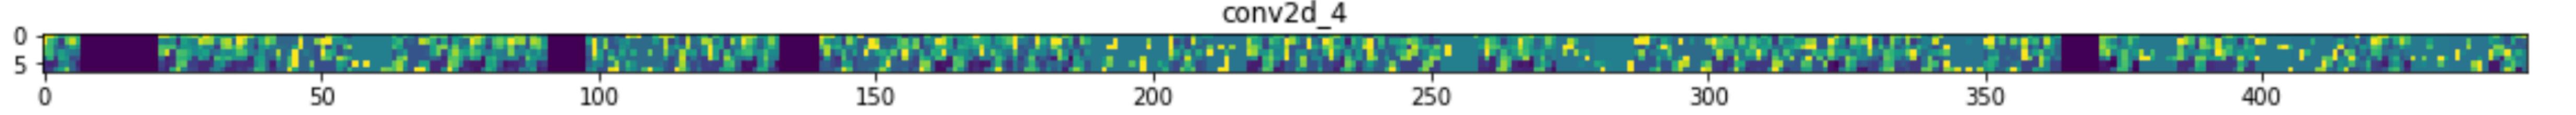
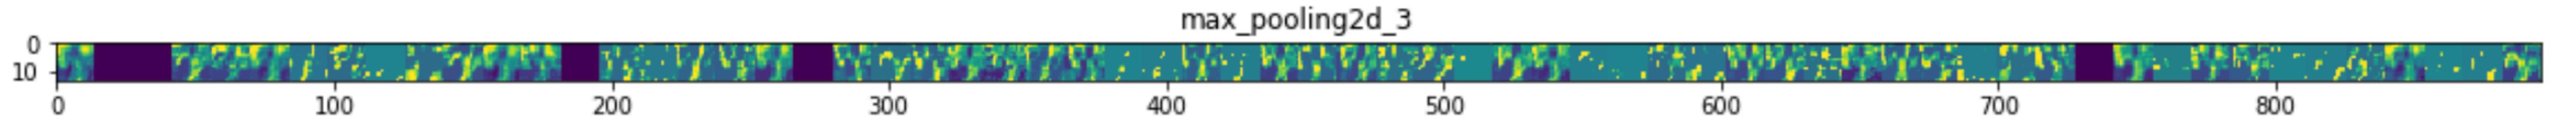
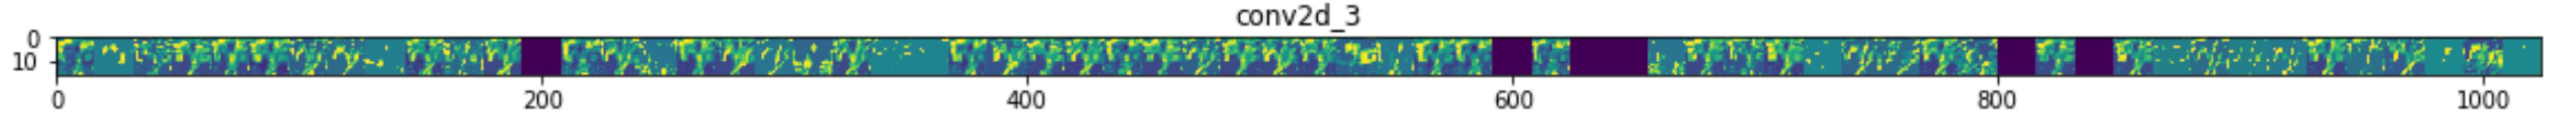
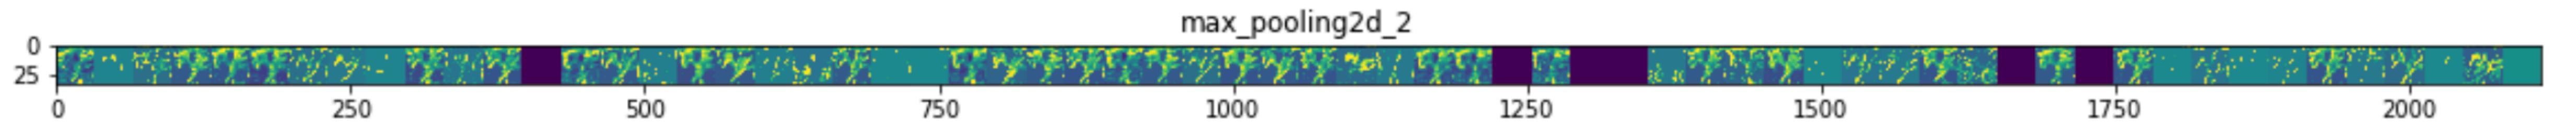
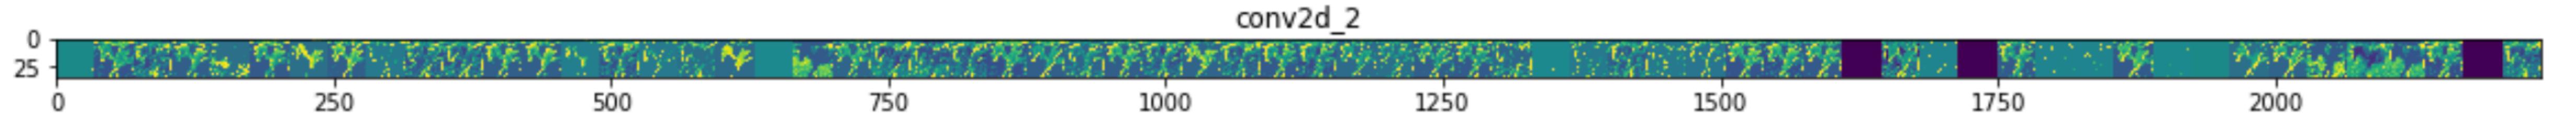
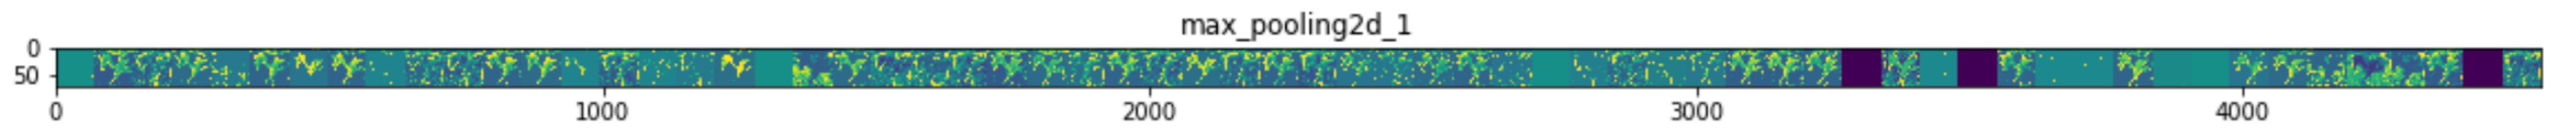
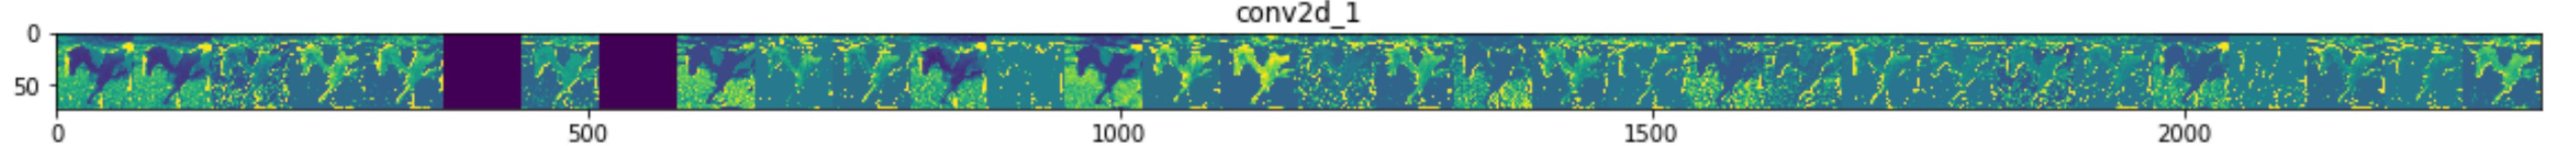
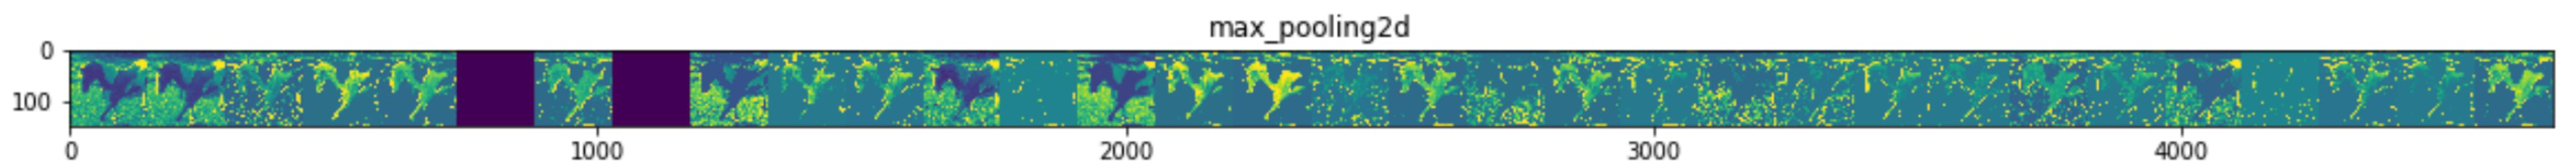
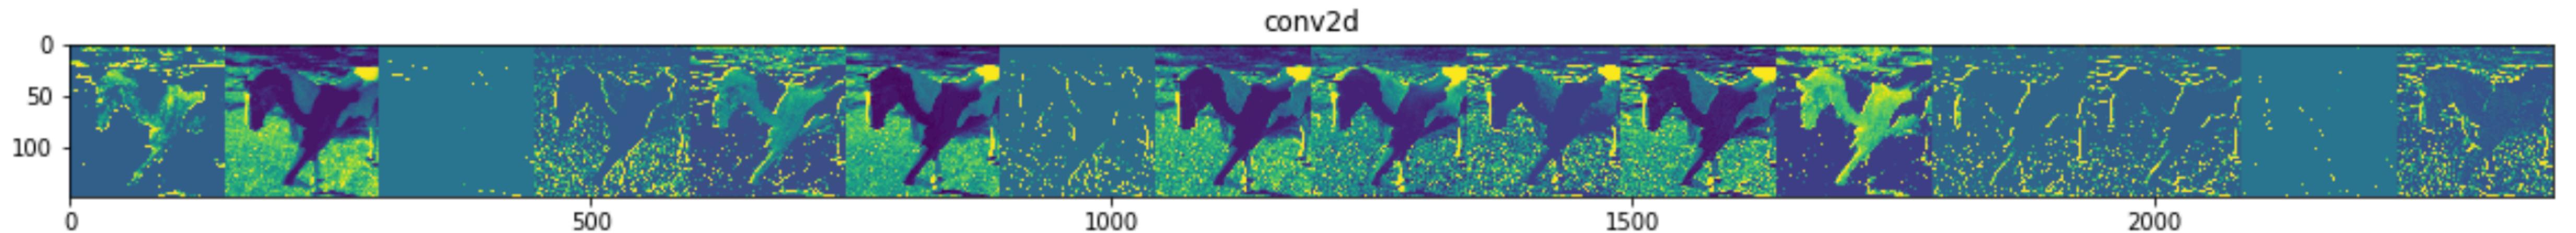
<https://www.coursera.org/learn/introduction-tensorflow>

AlexNet





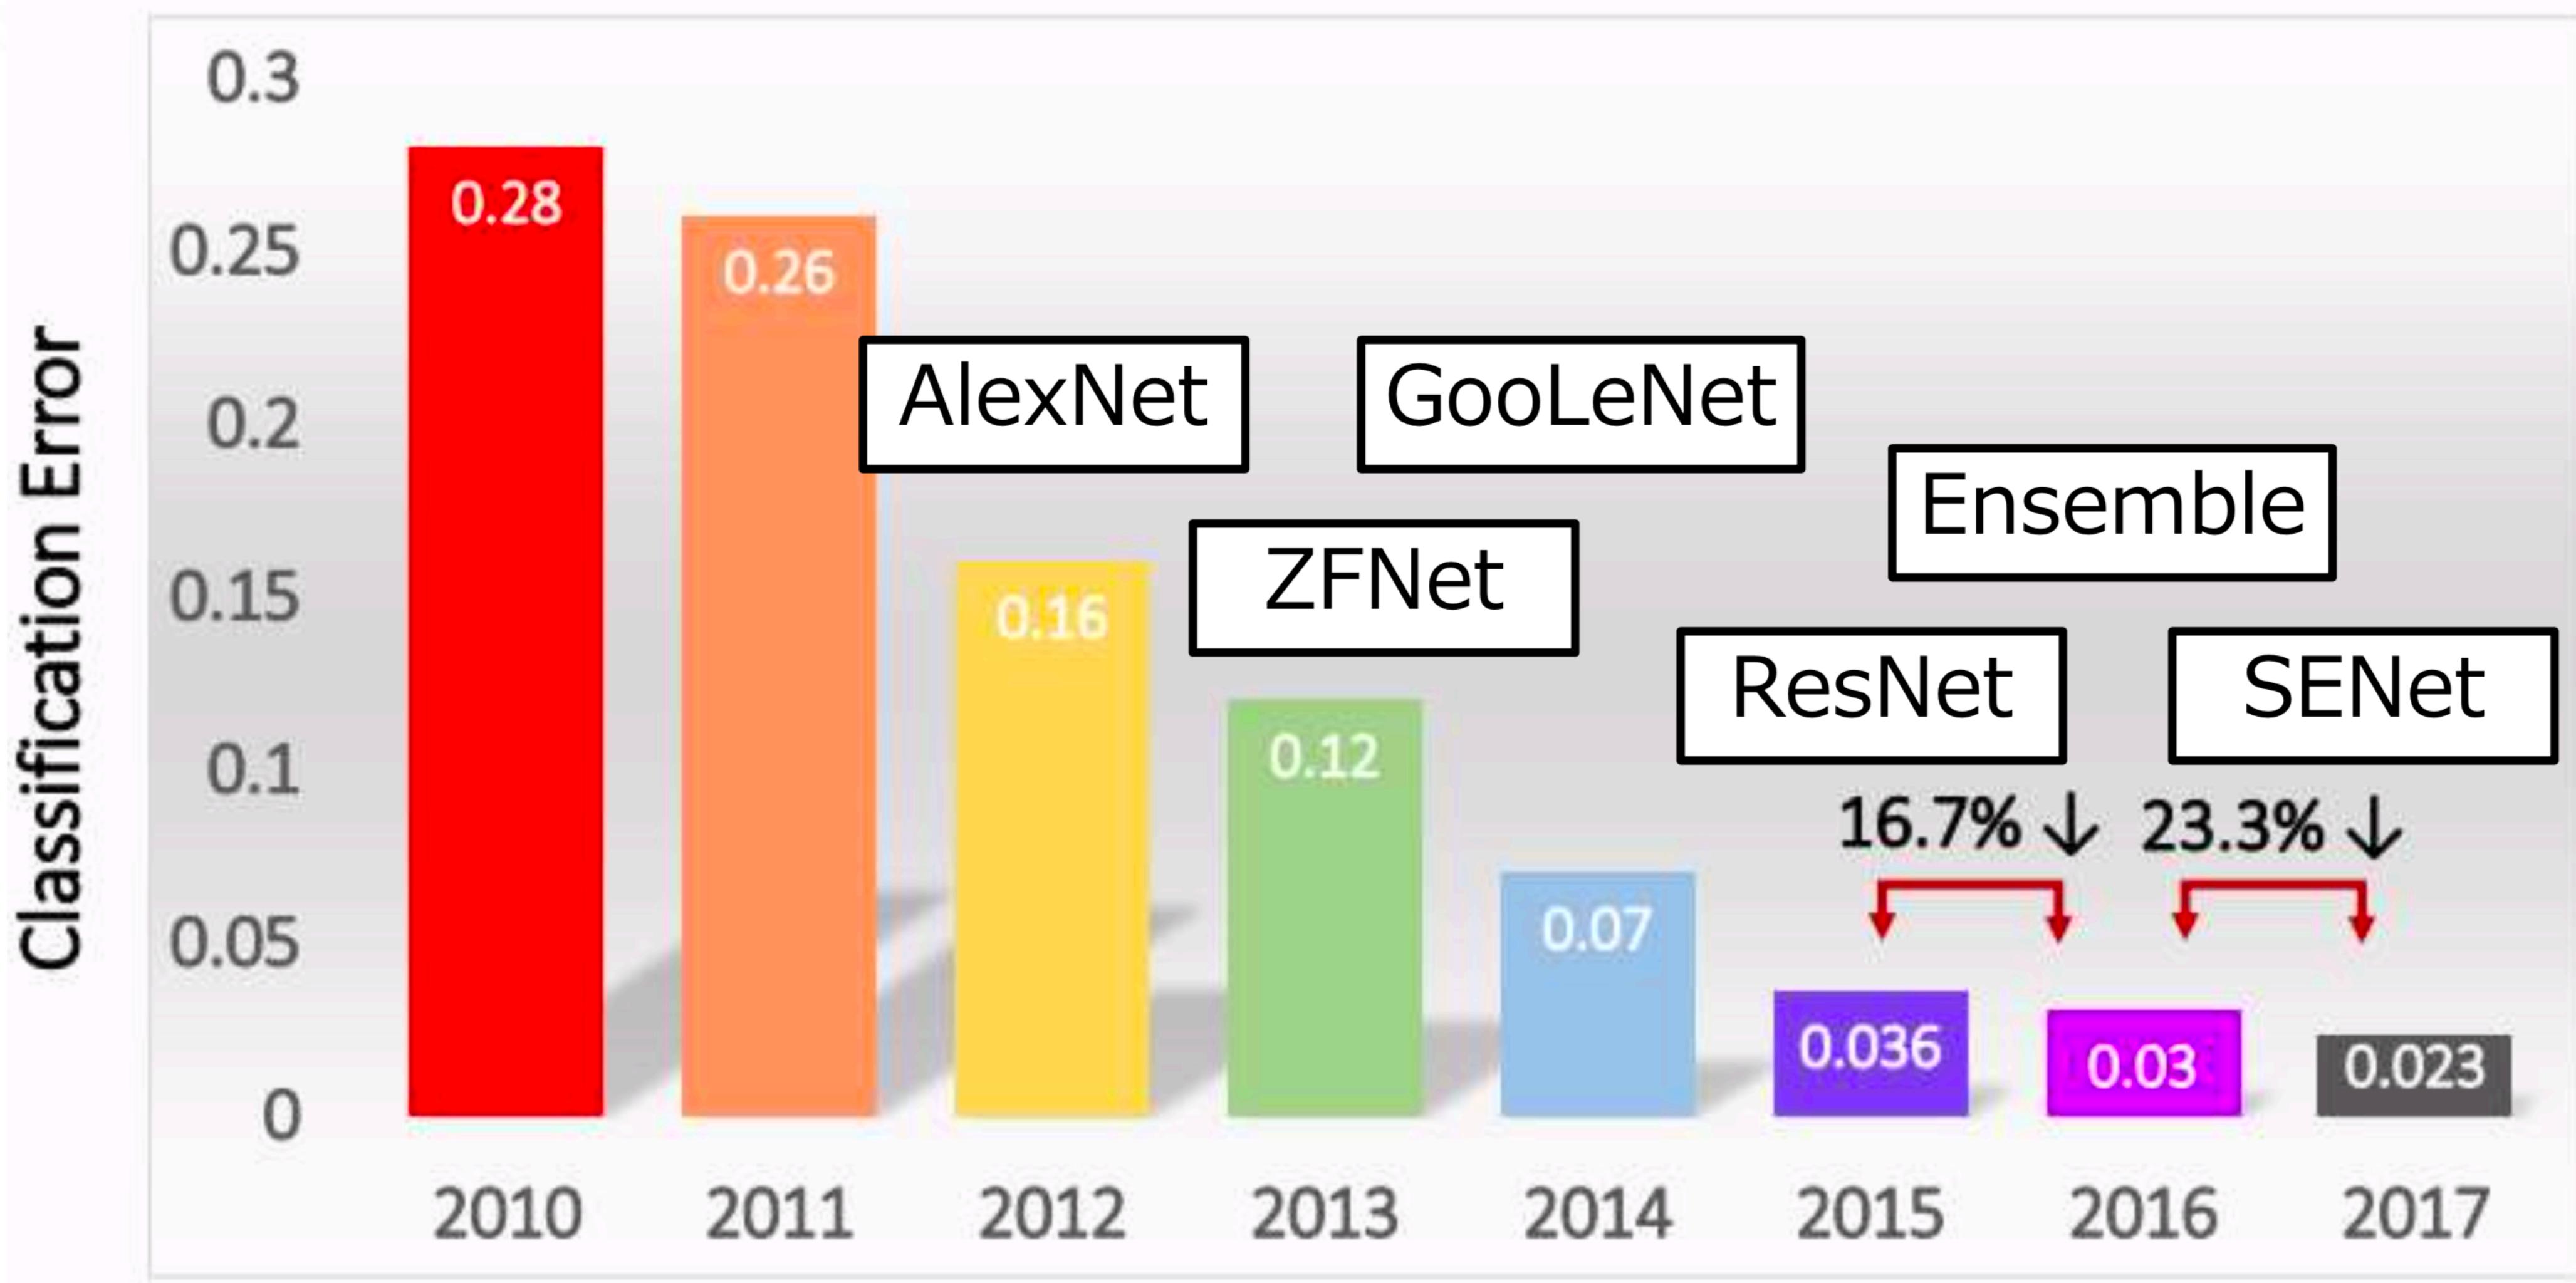






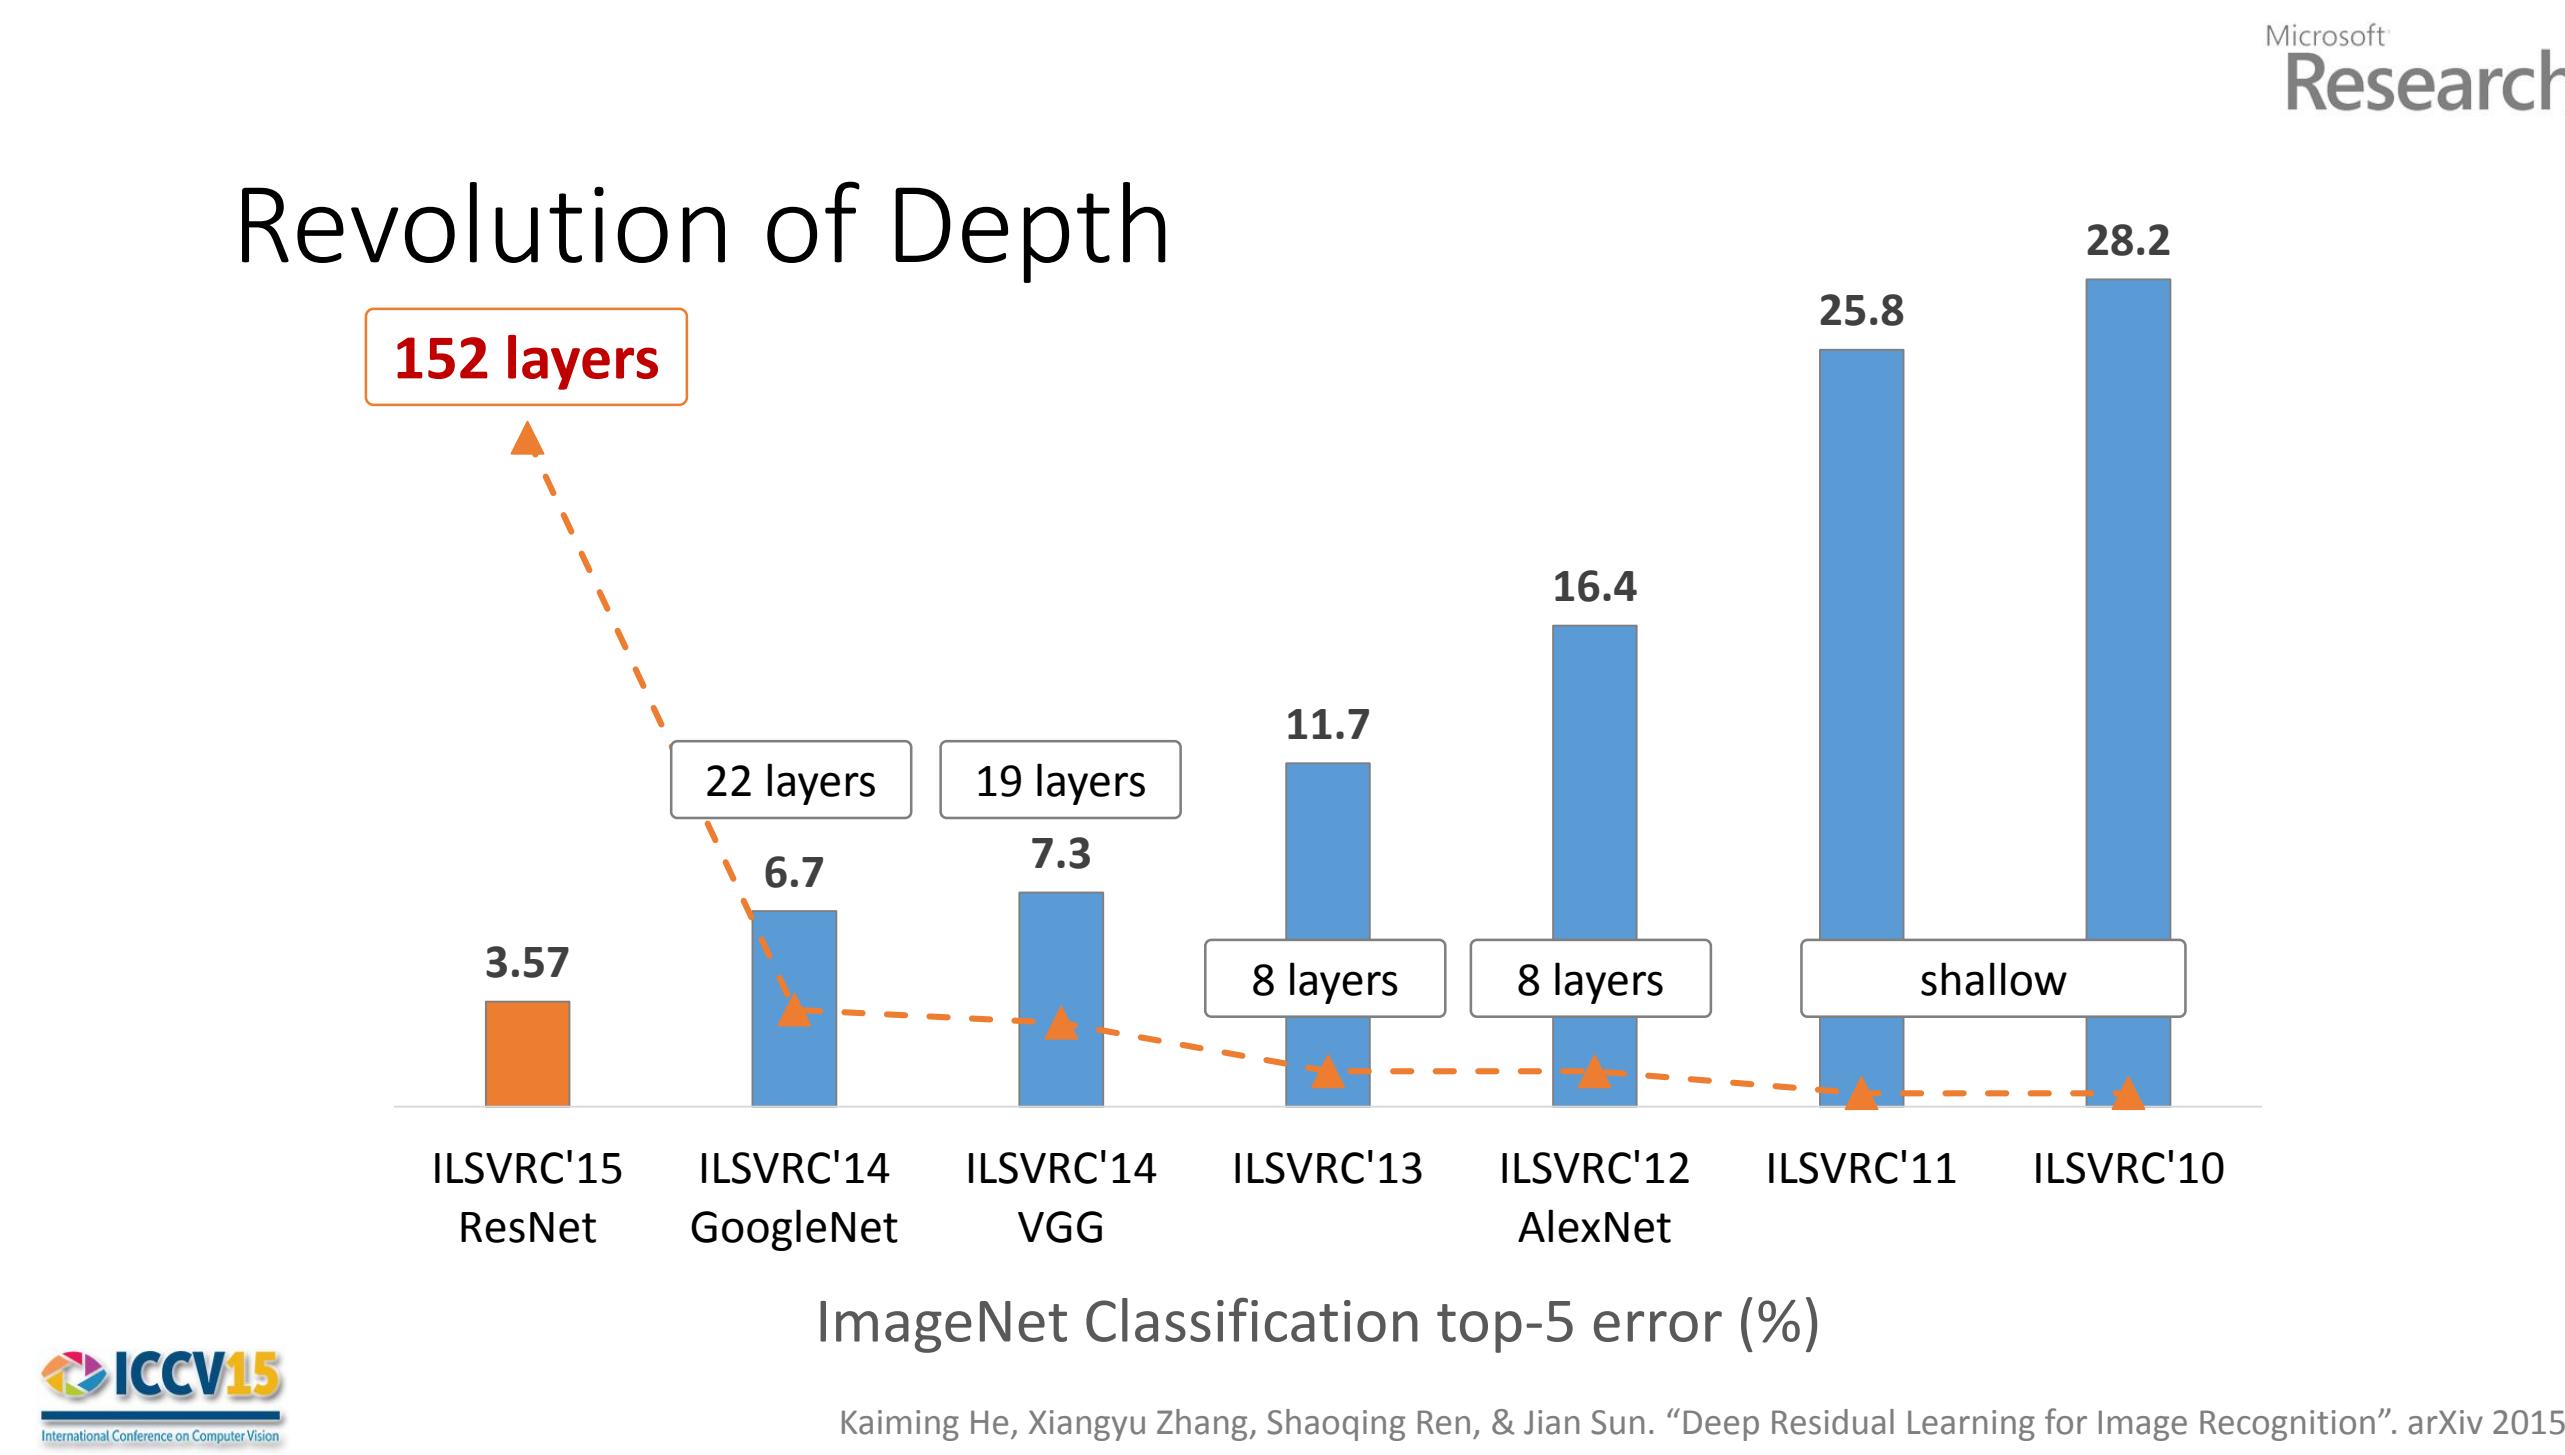
応用的なCNNアーキテクチャの紹介

ImageNet コンペの優勝モデル



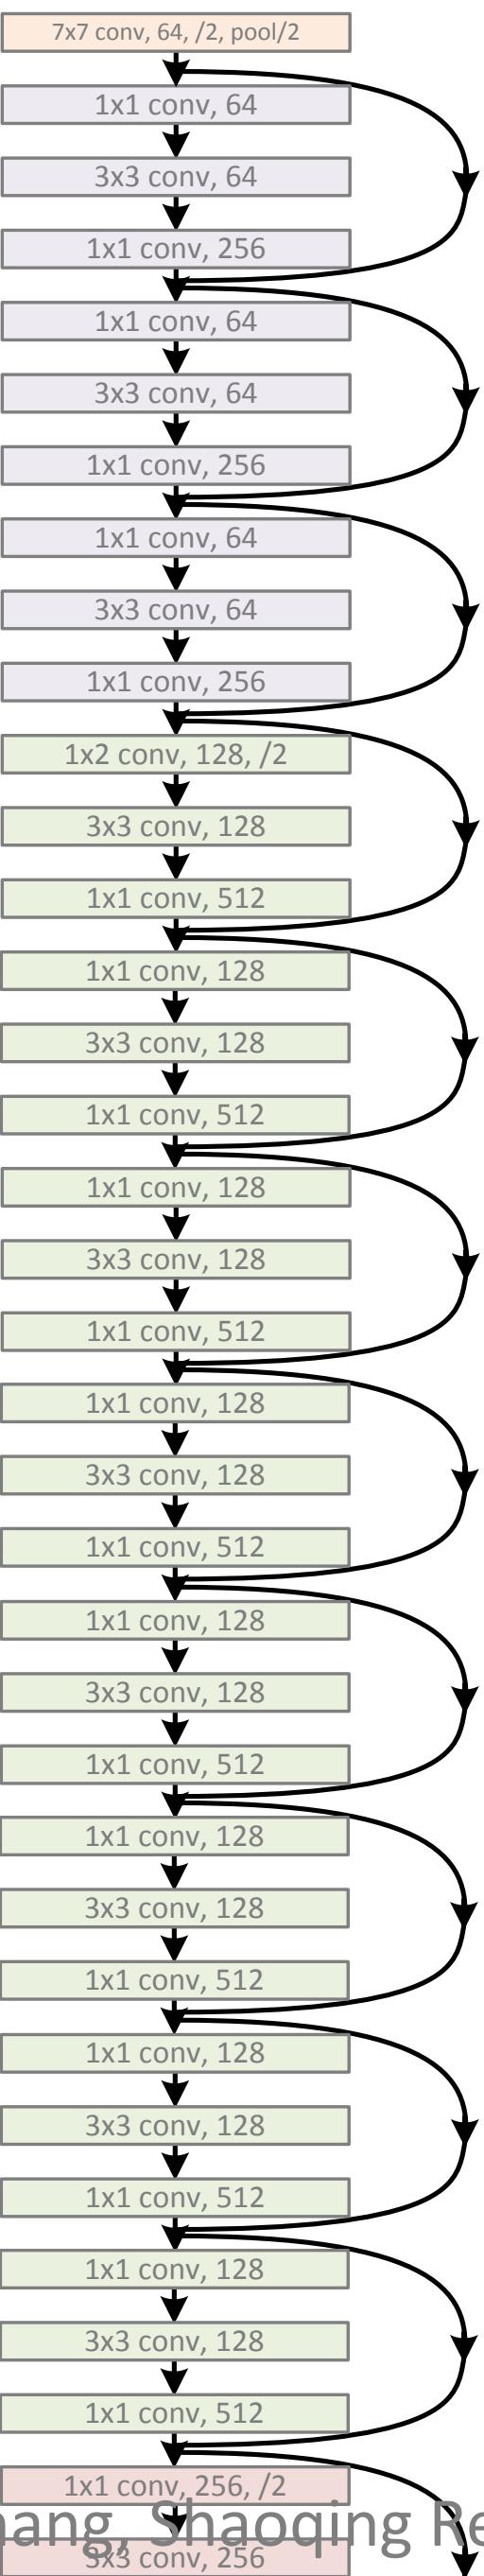
ResNet

- 2015年のImageNetコンペ (ILSVRC) 優勝モデル
- Residualモジュール（ショートカット機構）の導入



Revolution of Depth

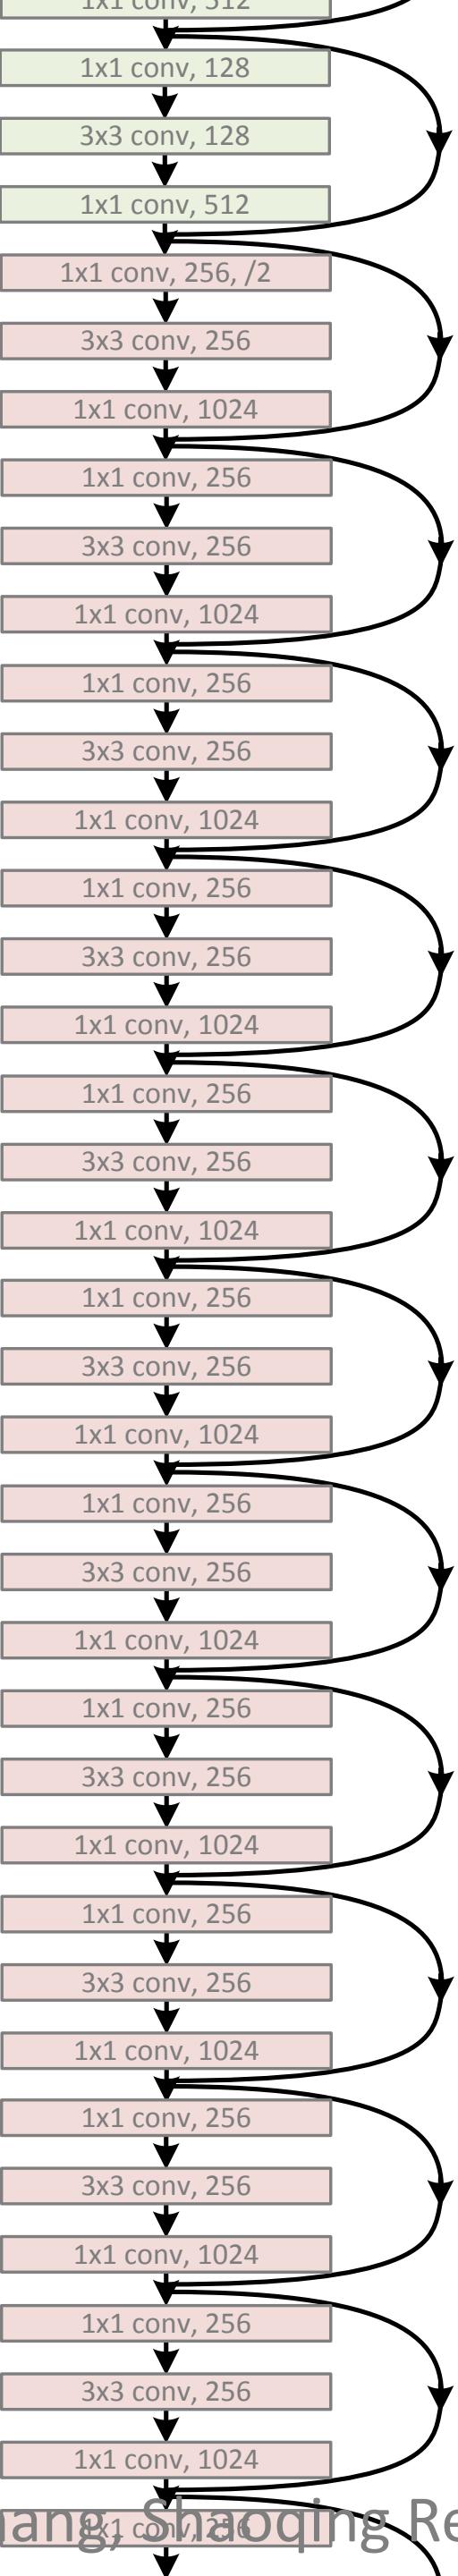
ResNet, 152 layers



(there was an animation here)

Revolution of Depth

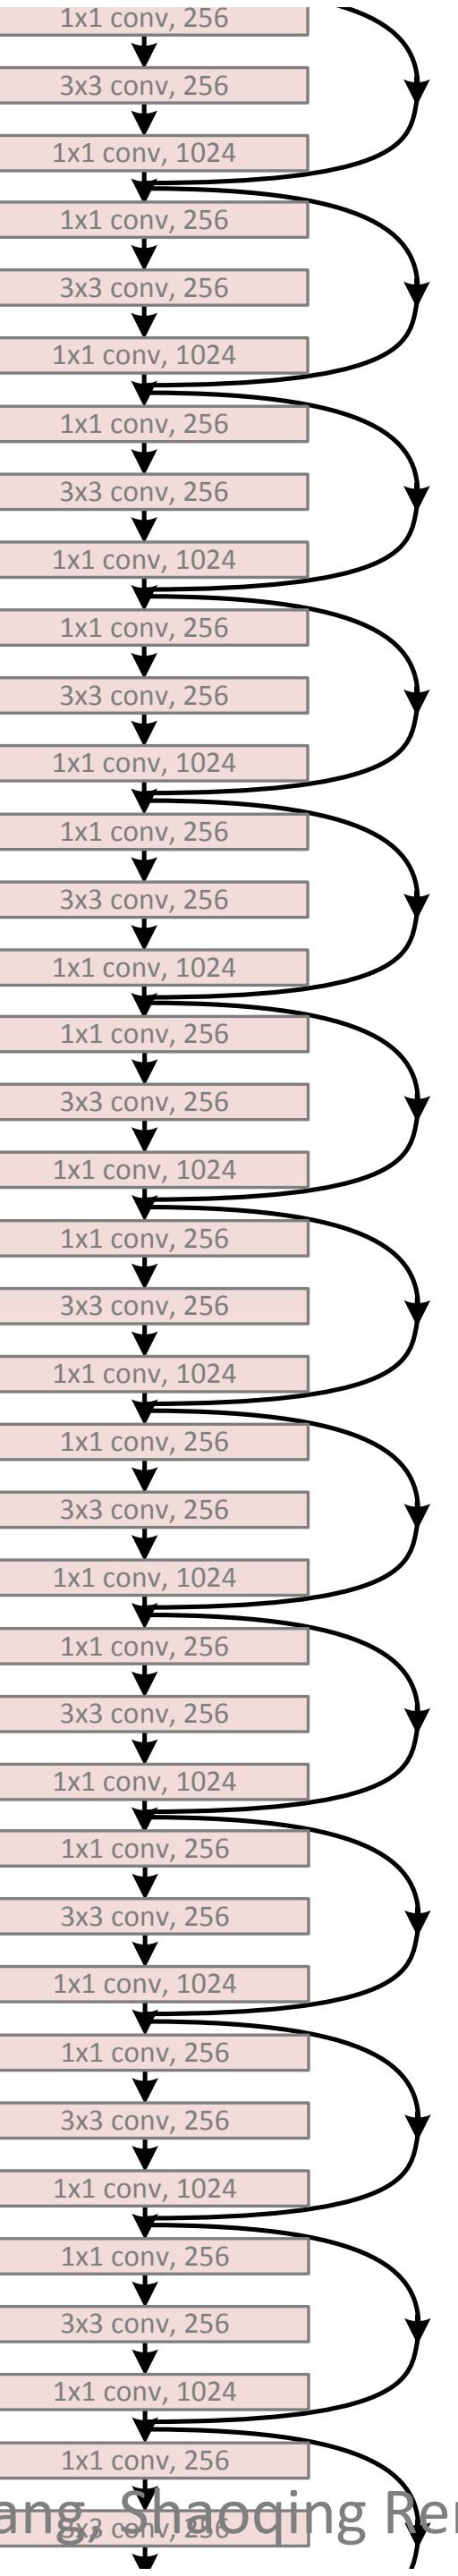
ResNet, 152 layers



(there was an animation here)

Revolution of Depth

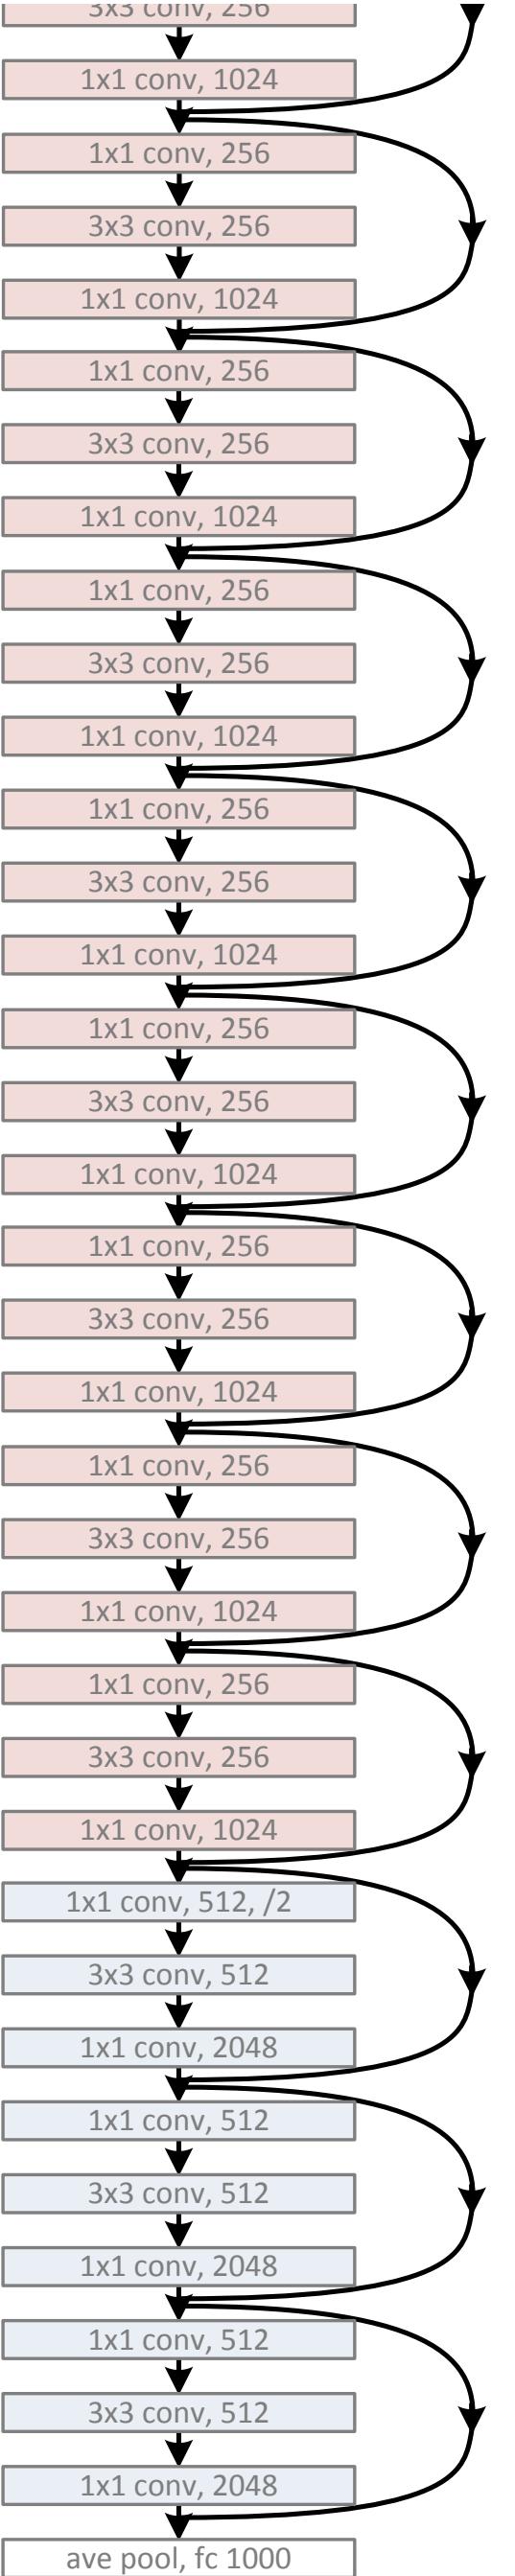
ResNet, 152 layers



(there was an animation here)

Revolution of Depth

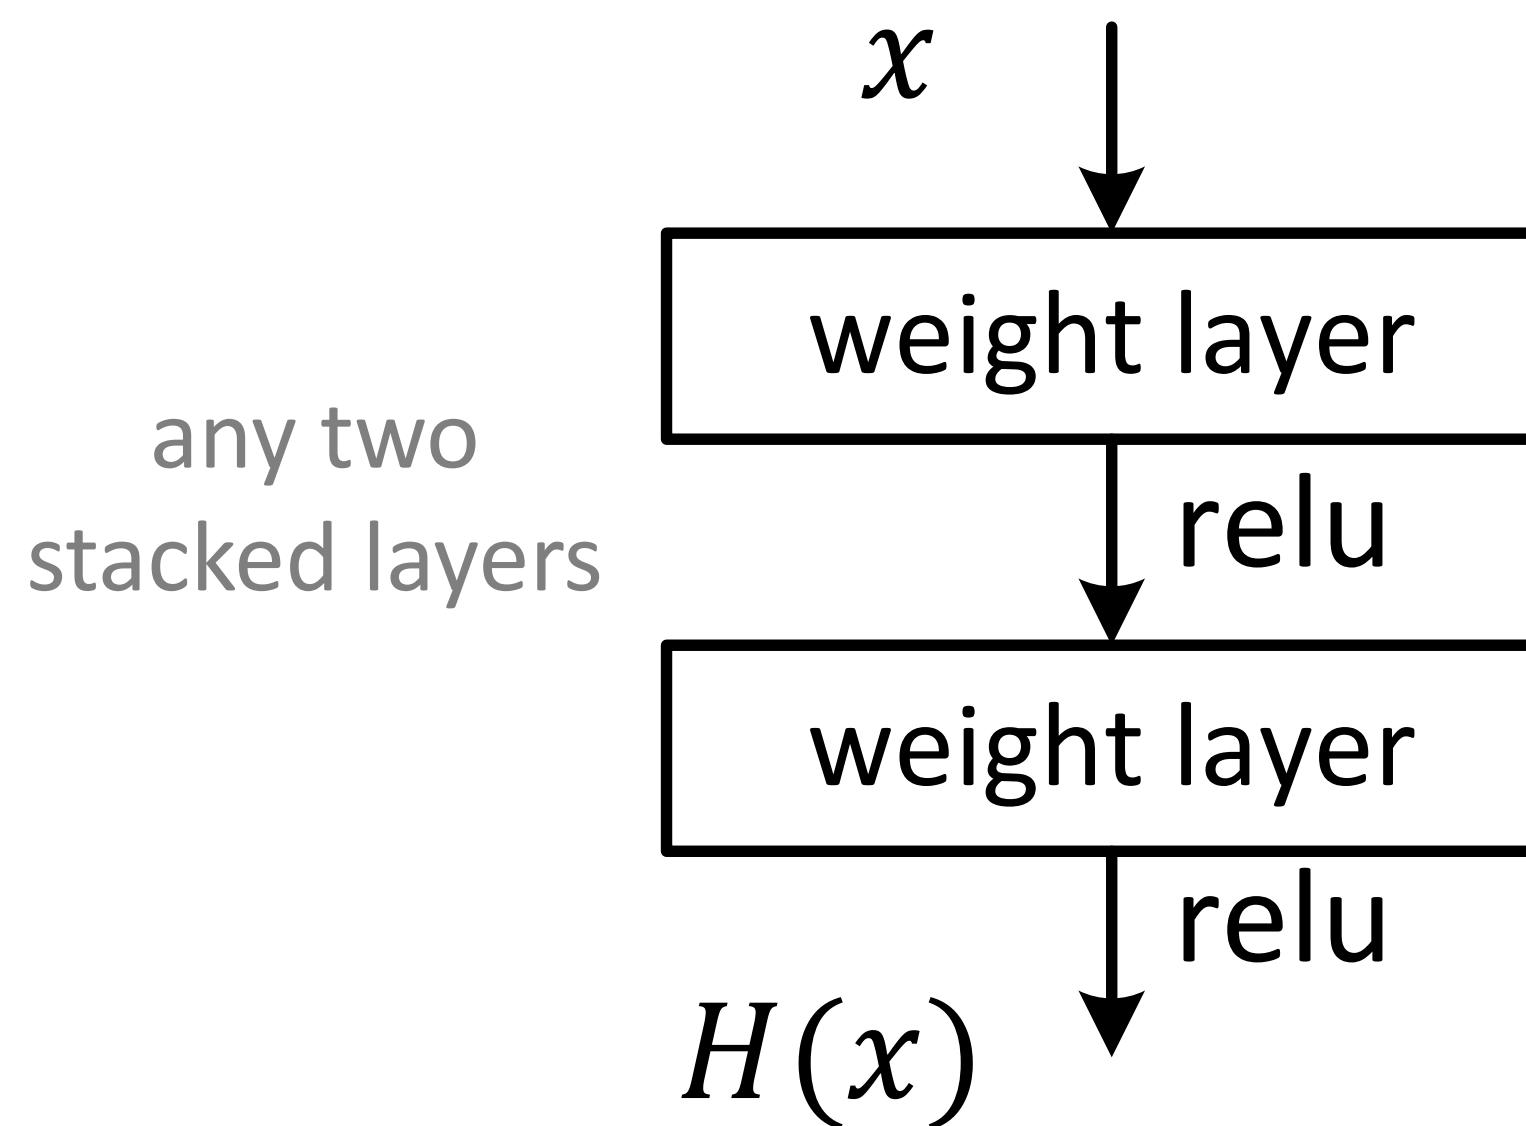
ResNet, 152 layers



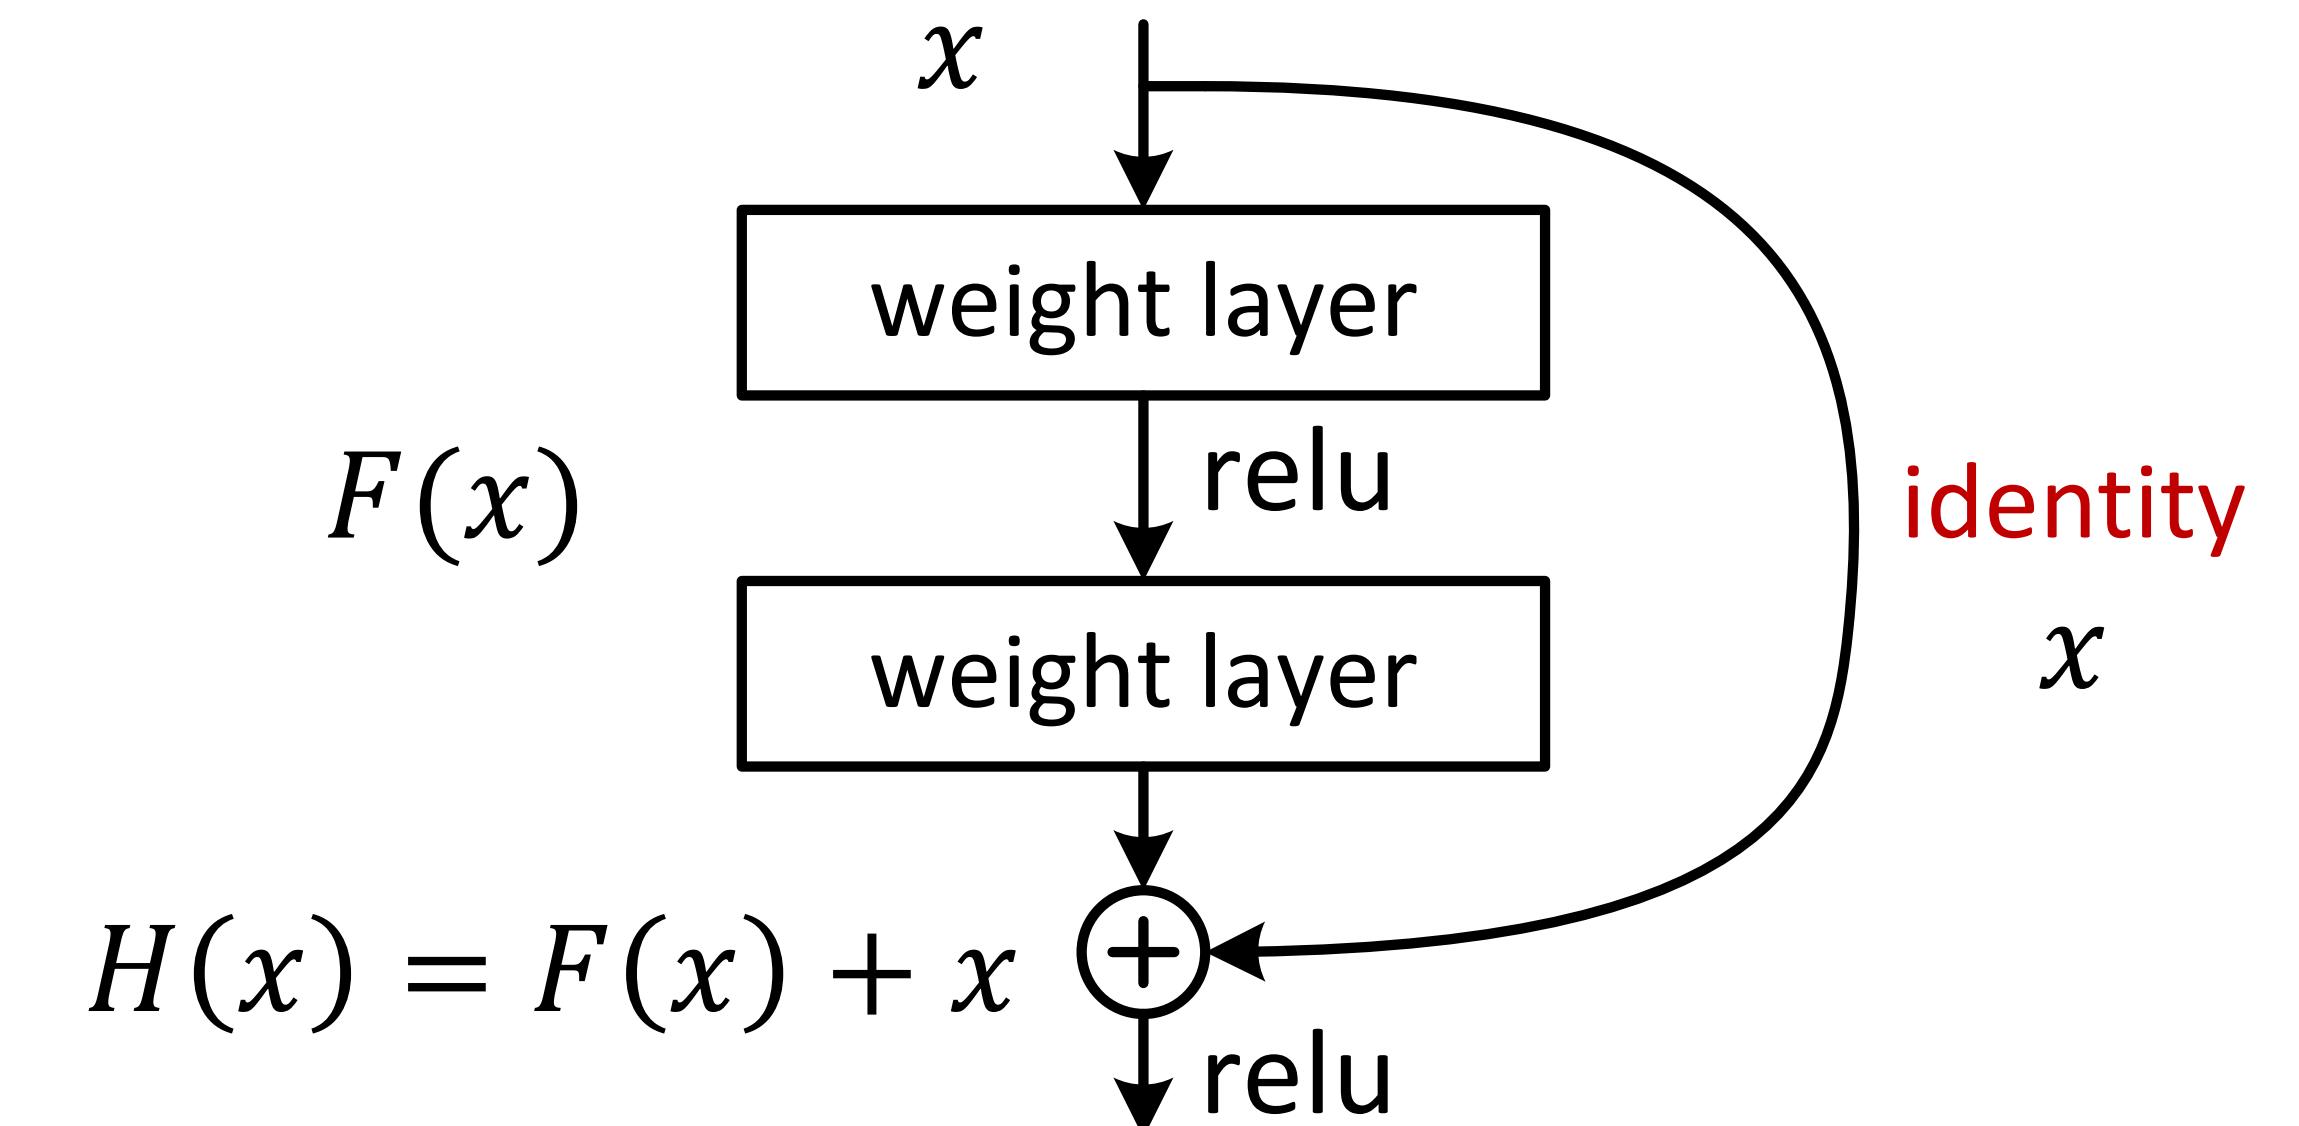
(there was an animation here)

ResNet

- Plain net

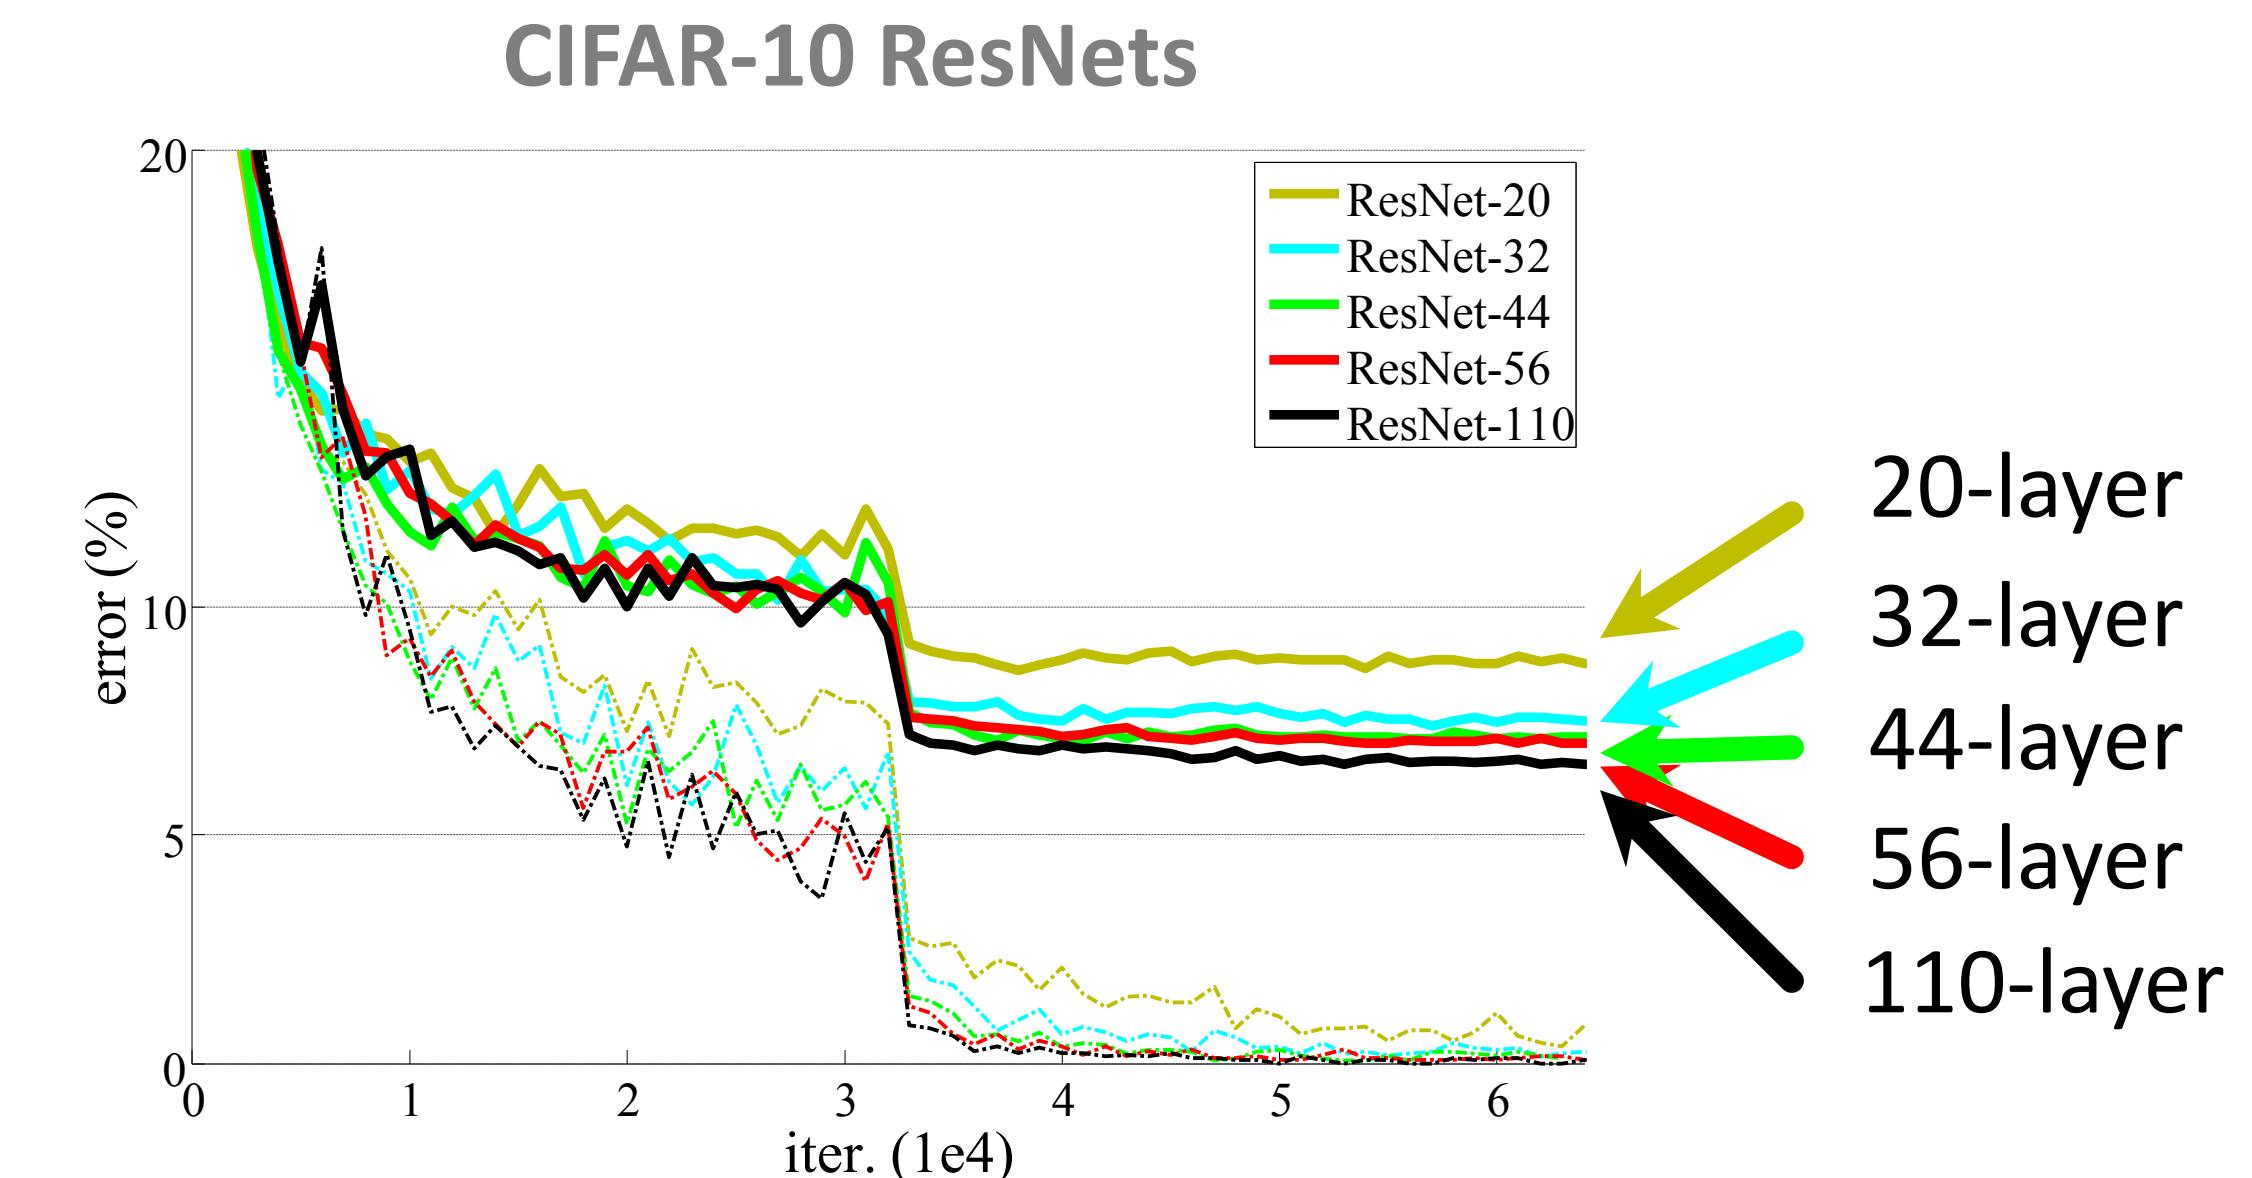
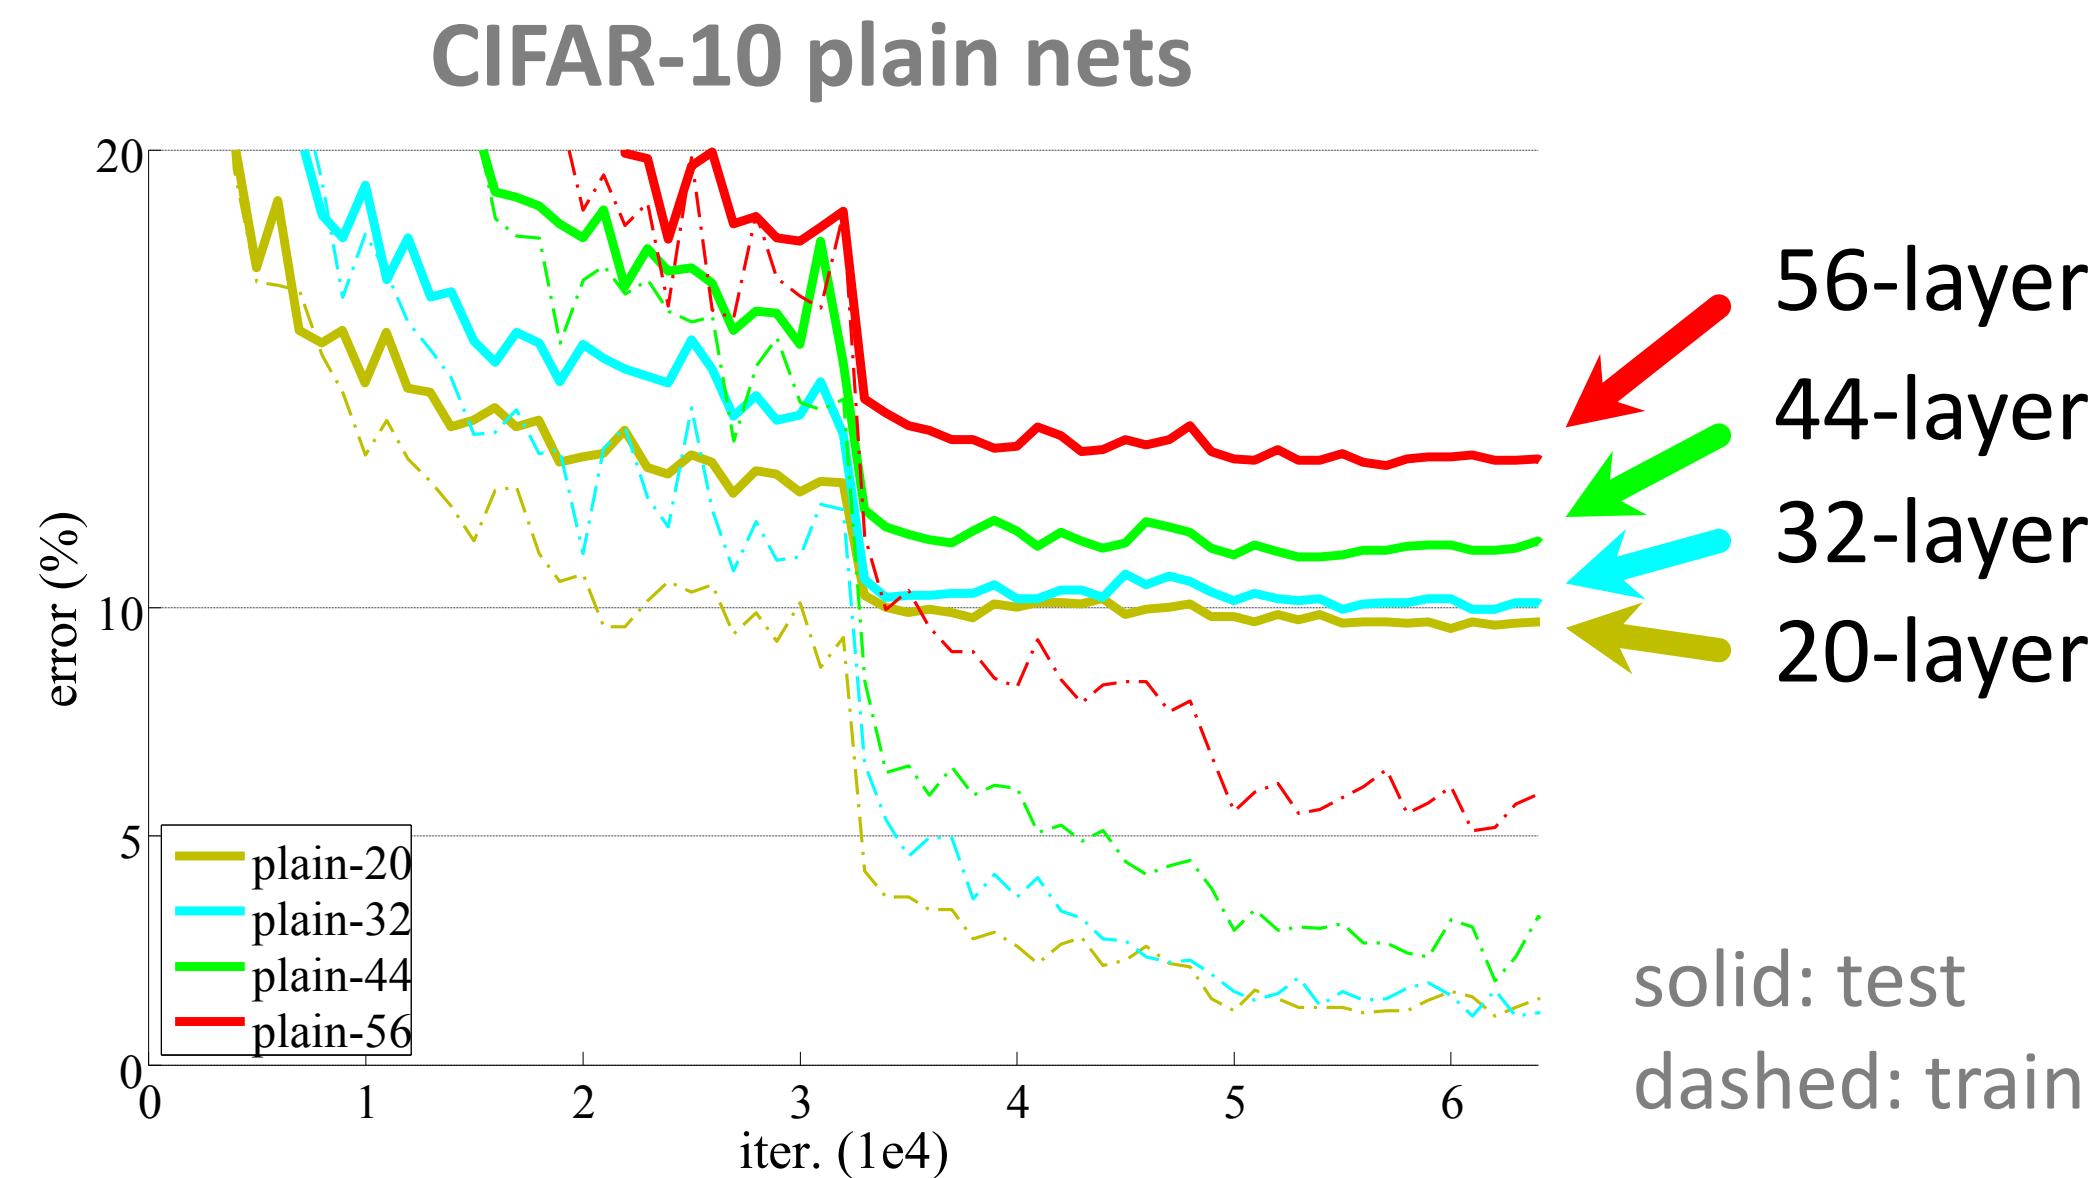


- Residual net



http://image-net.org/challenges/talks/ilsvrc2015_deep_residual_learning_kaiminghe.pdf

CIFAR-10 experiments

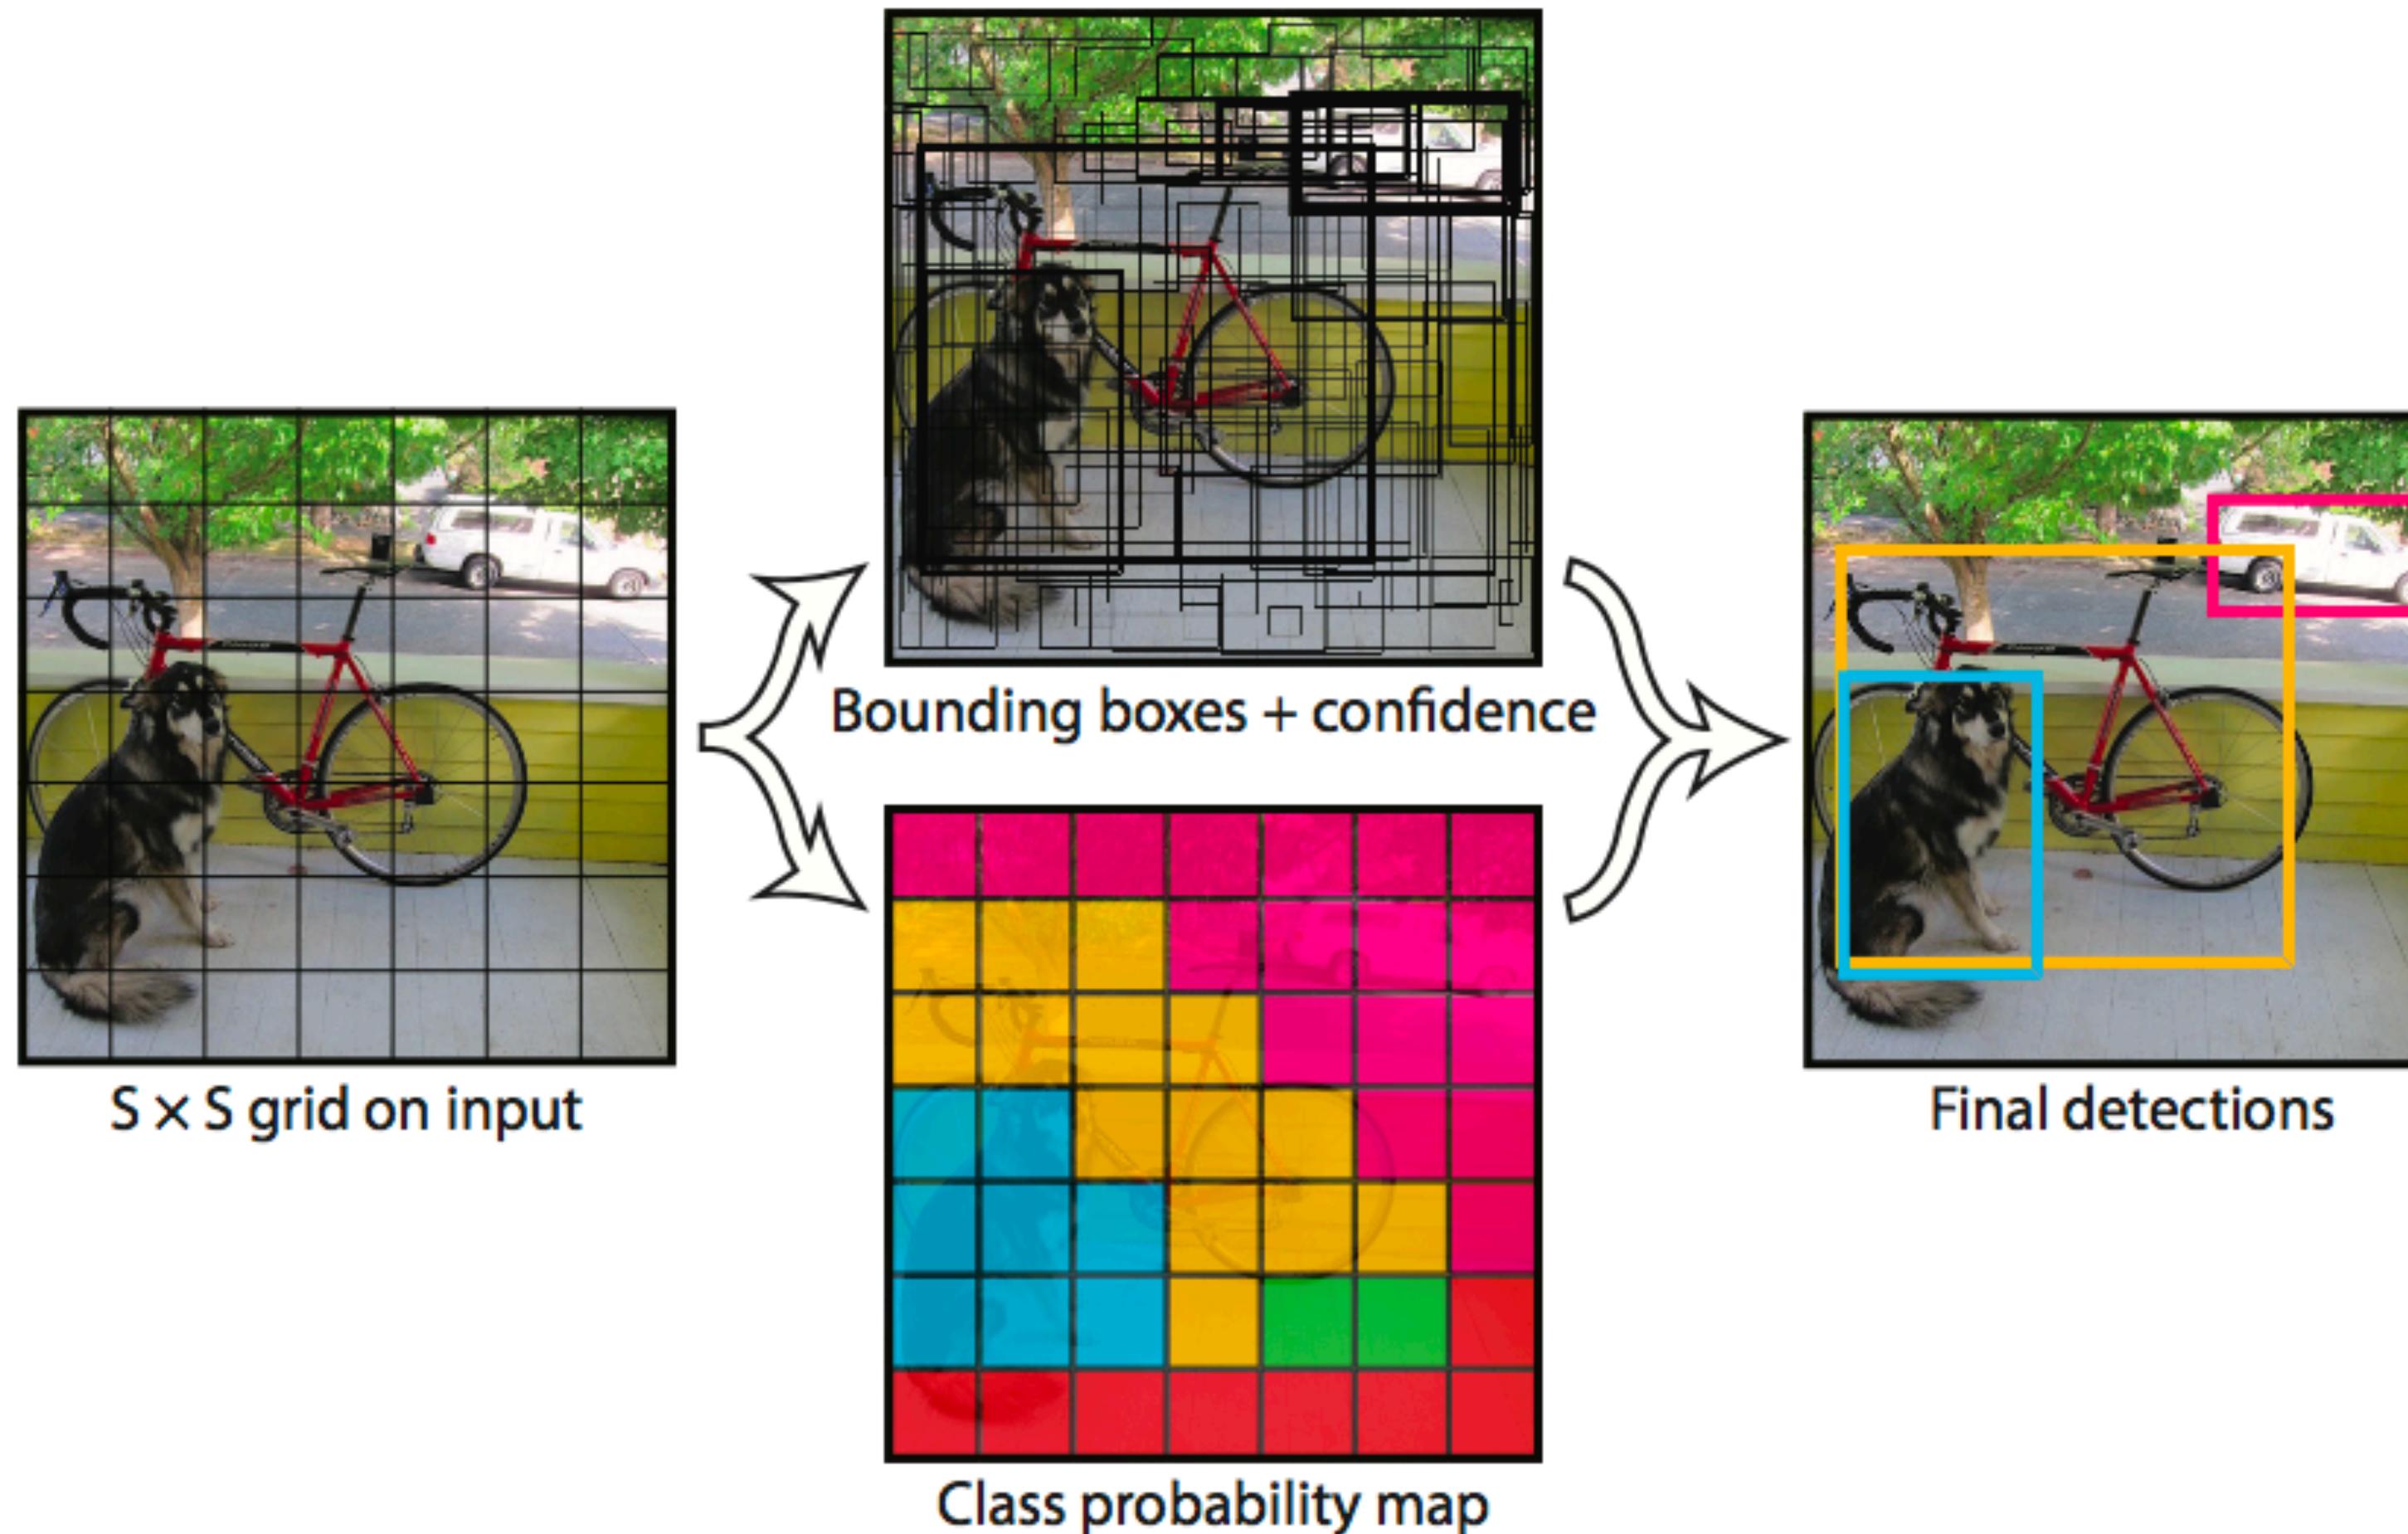


- Deep ResNets can be trained without difficulties
- Deeper ResNets have **lower training error**, and also lower test error

Further Reading

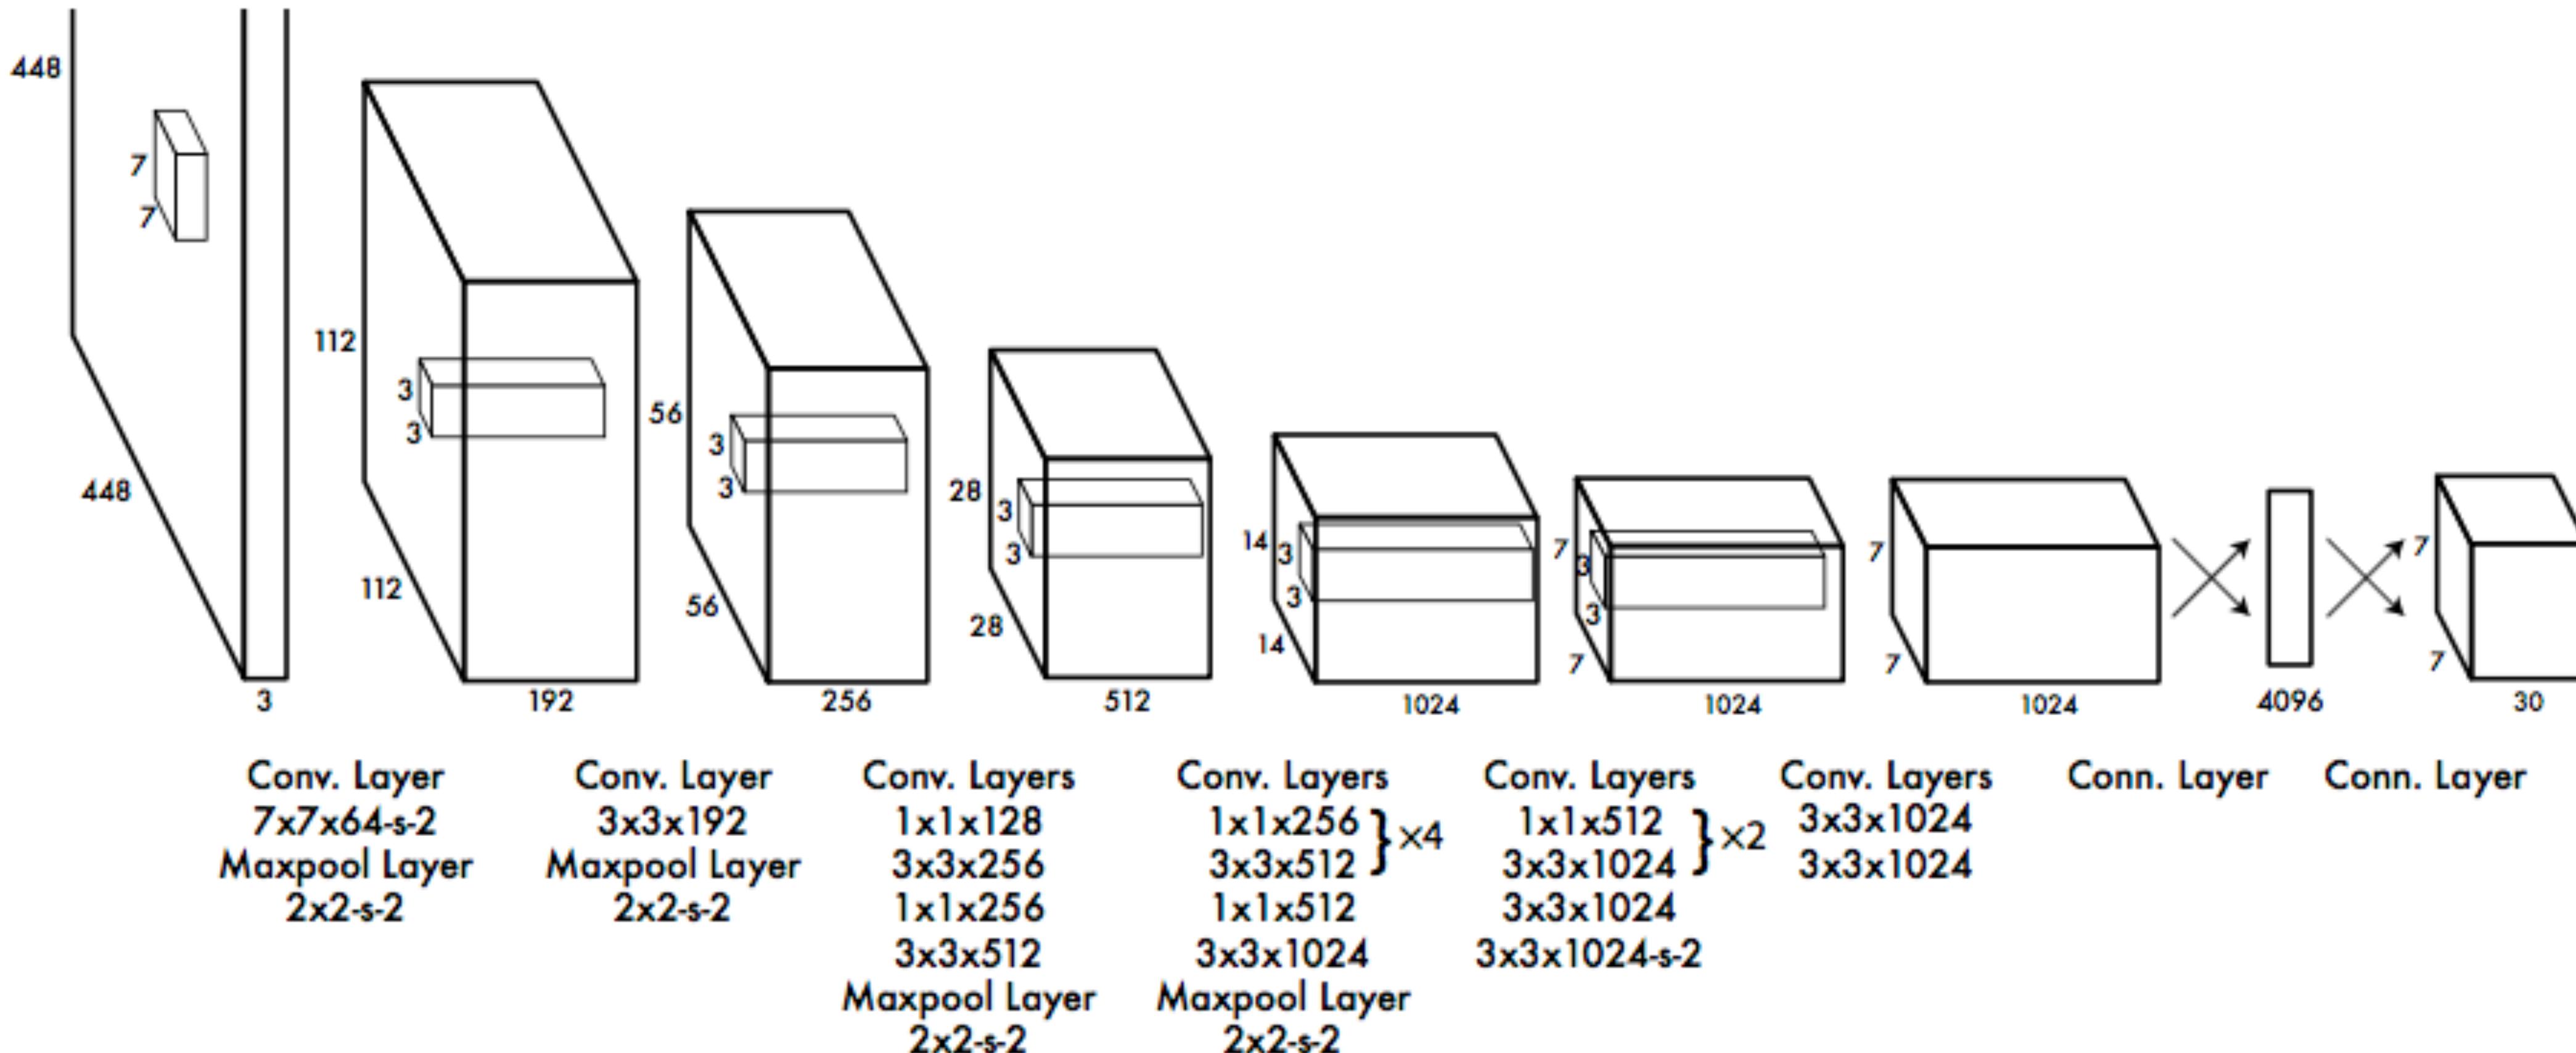
- 一年半前 (2017-12) の資料ですが素晴らしいスライドです
- 置み込みニューラルネットワークの研究動向 by DeNA 内田さん
- <https://www.slideshare.net/ren4yu/ss-84282514>

YOLO (You Only Look Once)

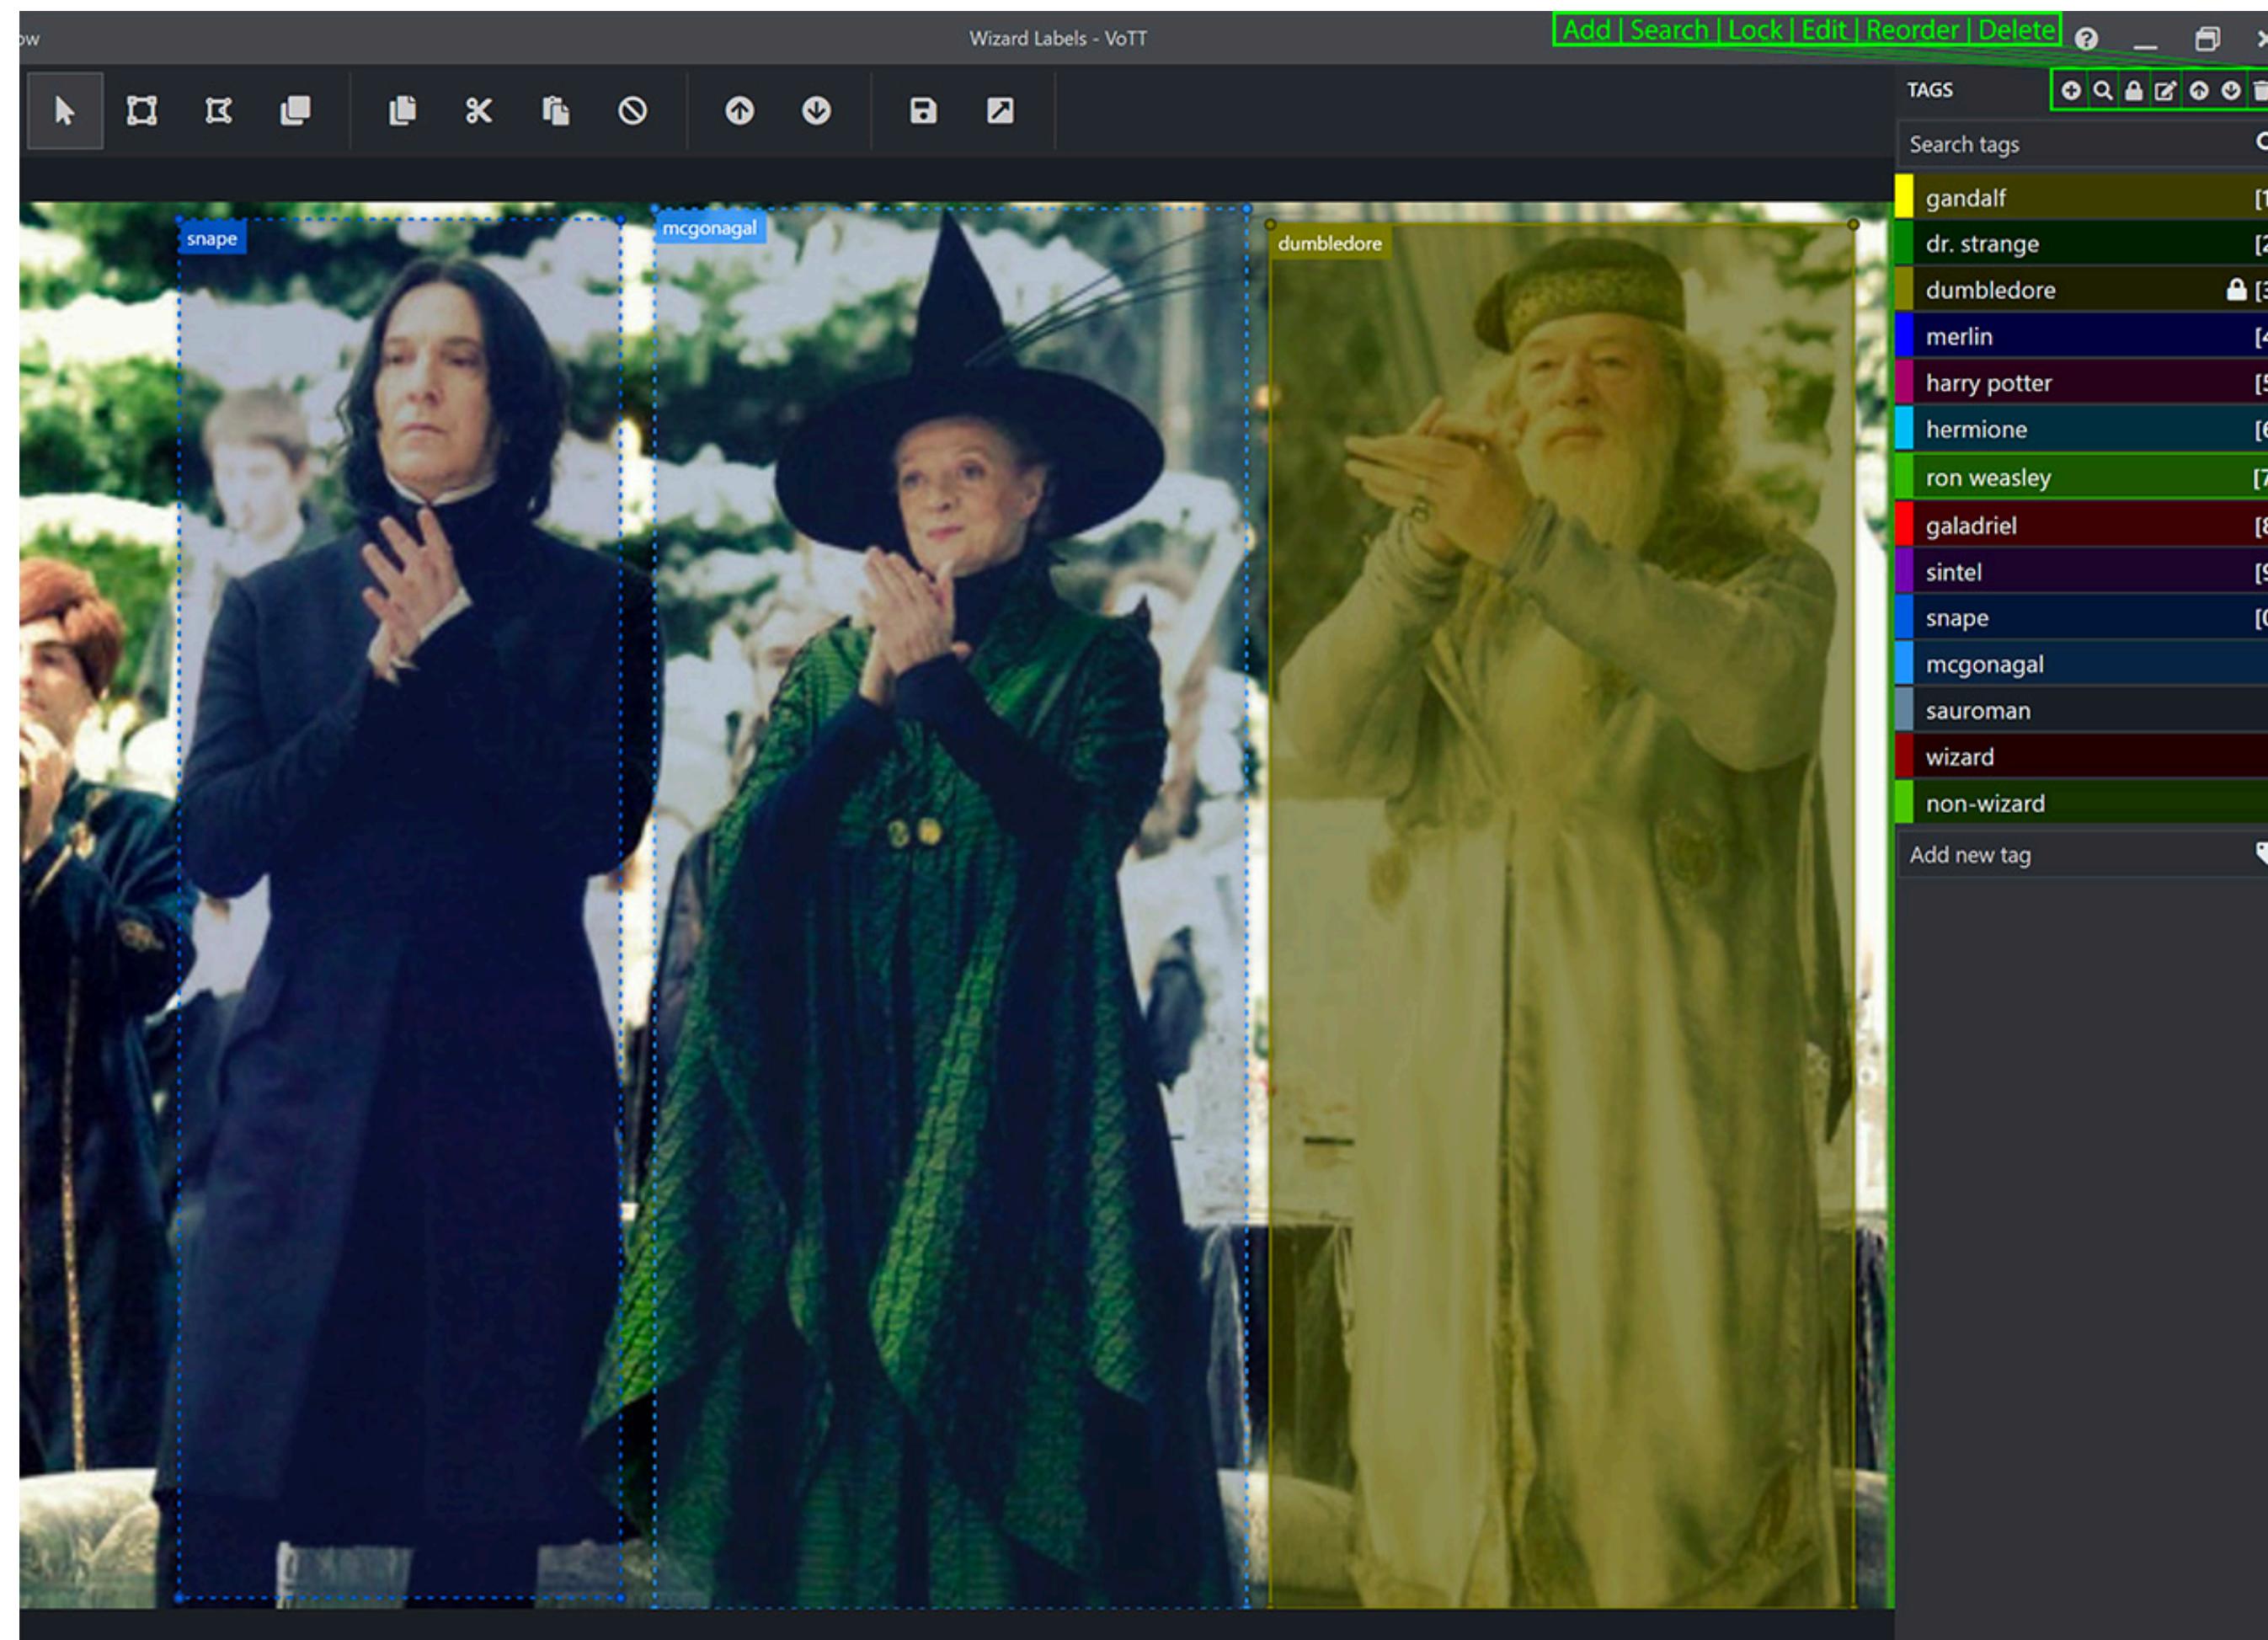


<https://www.youtube.com/watch?v=MPU2HistivI>

YOLOのアーキテクチャ



Labeling / Annotation



<https://github.com/Microsoft/VoTT>



AIモデルの運用・デプロイ

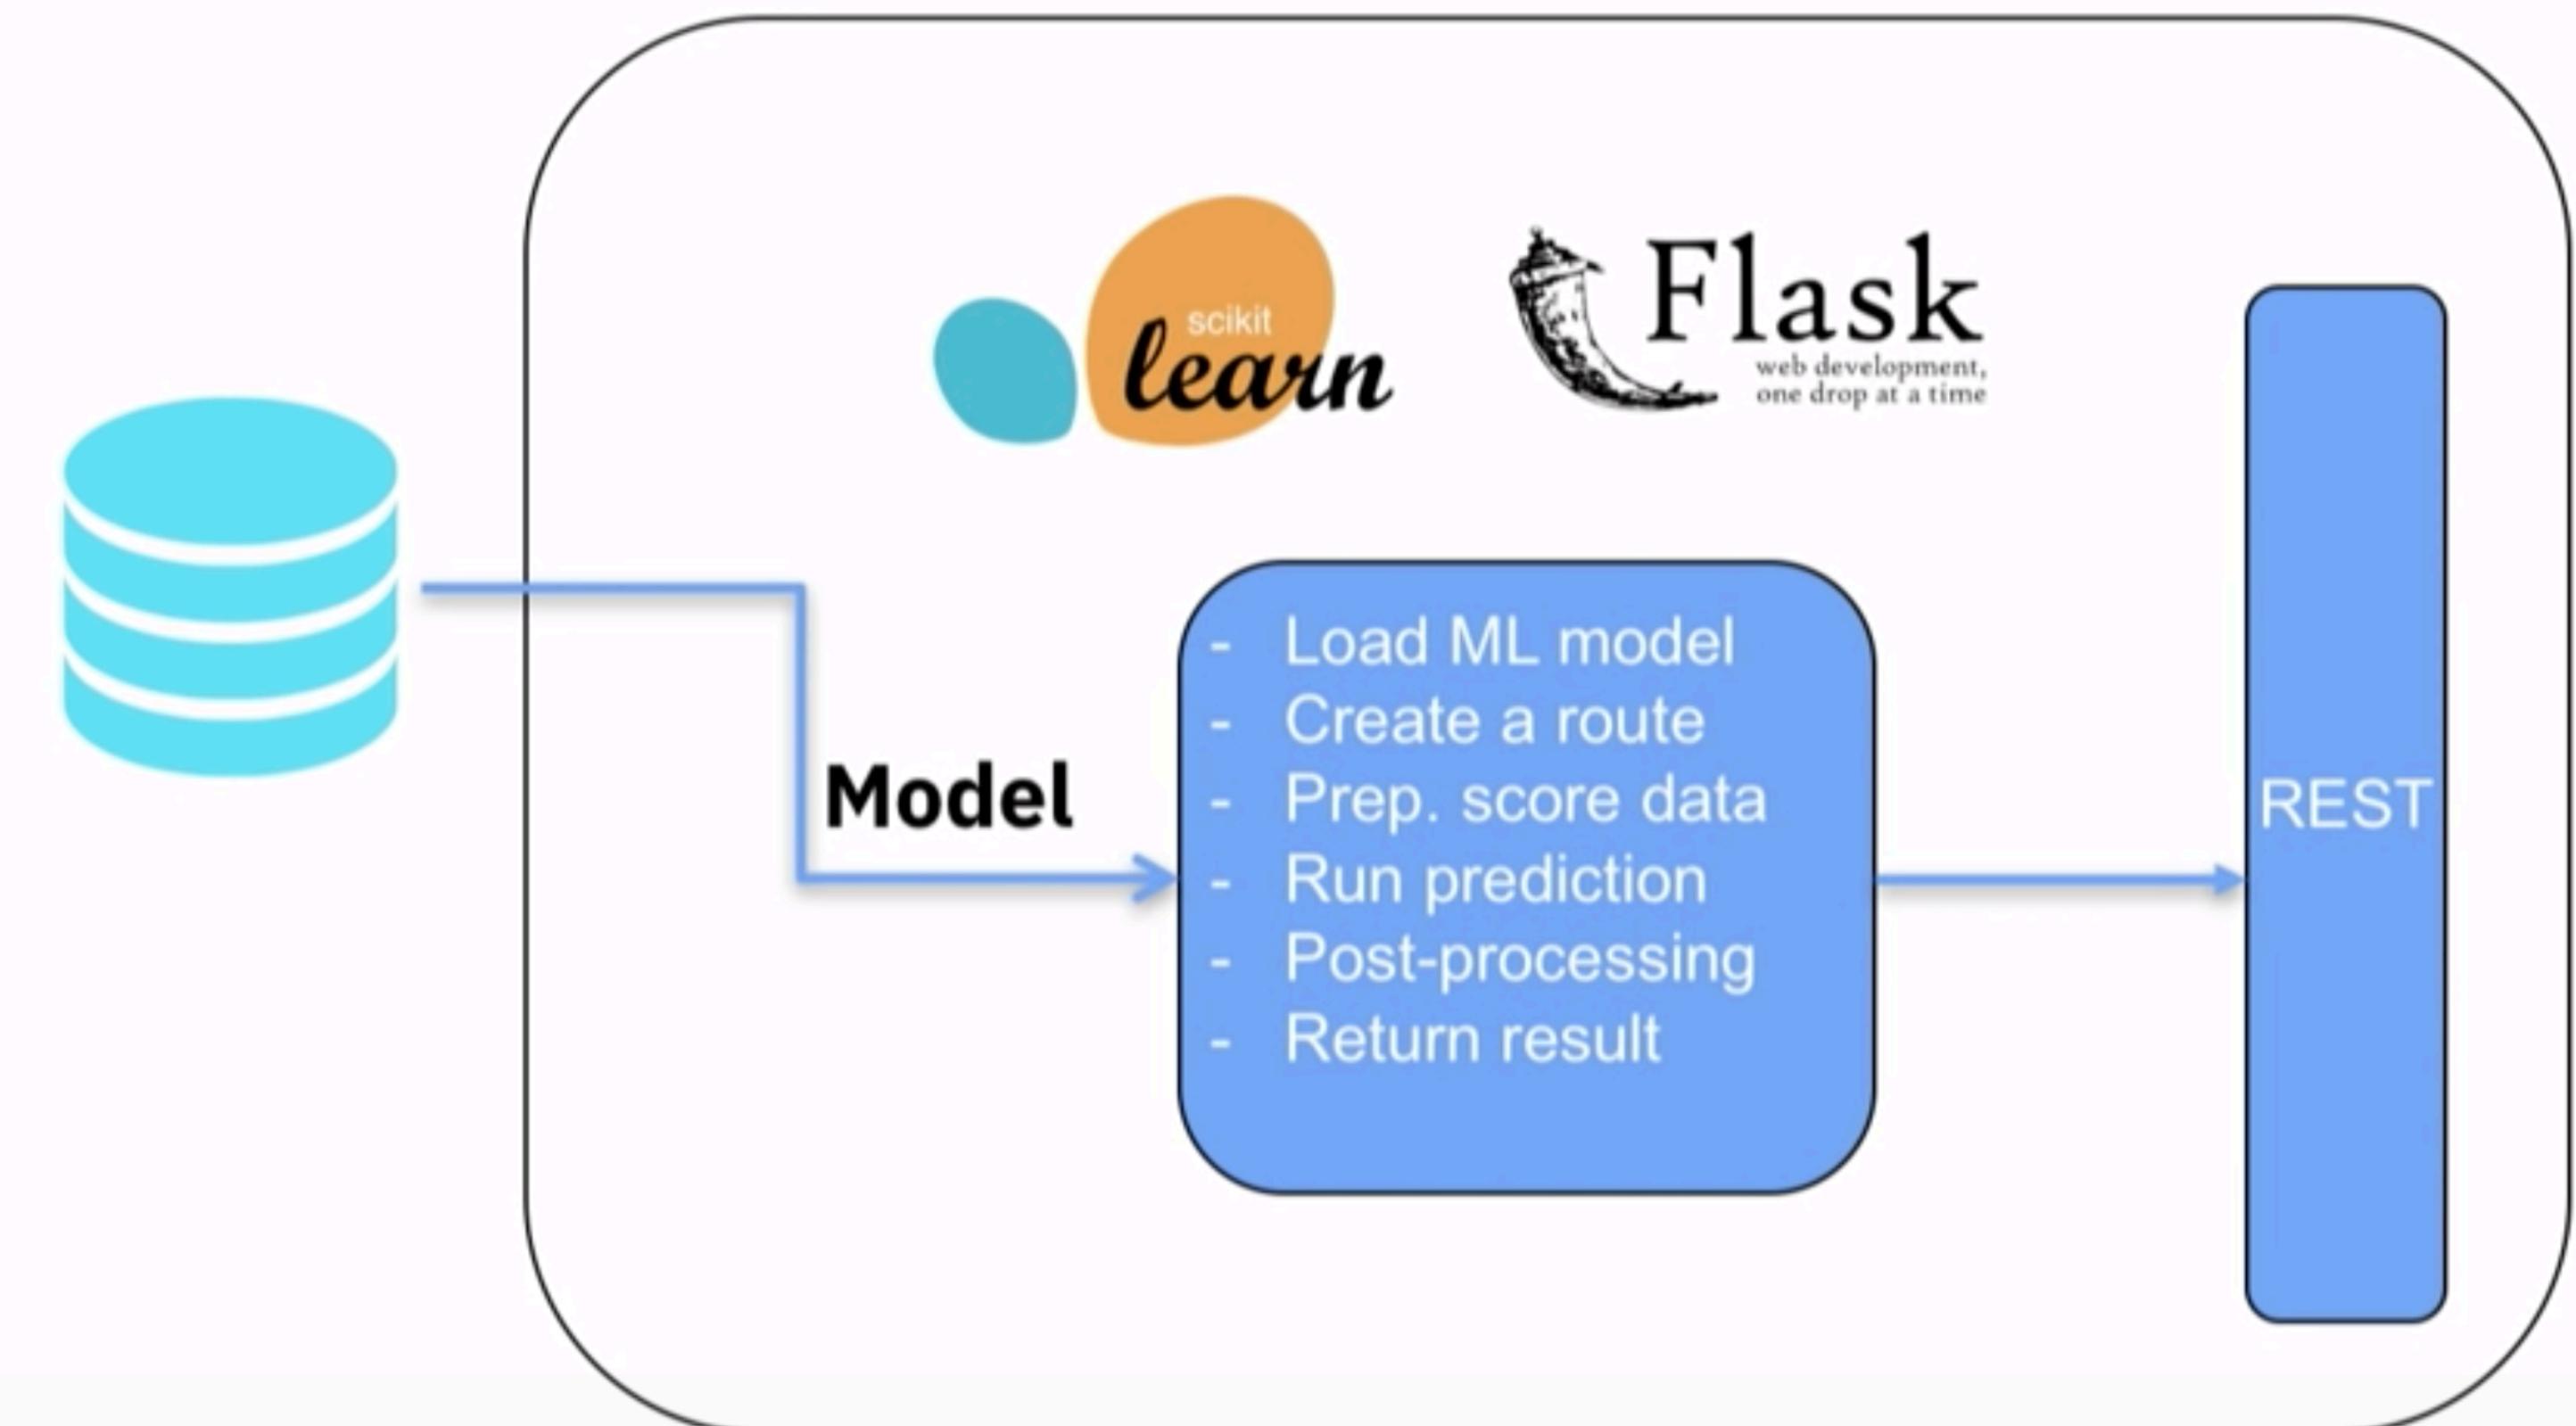


Deployment with Cloud Foundry

Python App - REST API
Endpoint: /v1/xyz/prediction : POST

Questions

- How to configure the app ?
- How many instances ?
- How much RAM ?
- Zero downtime deployment ?
- Connection to DB?
- Logging ?
- Crashes and recovery?
- Routing ?
- Workload scaling ?
- Integration with other tools?
- ...



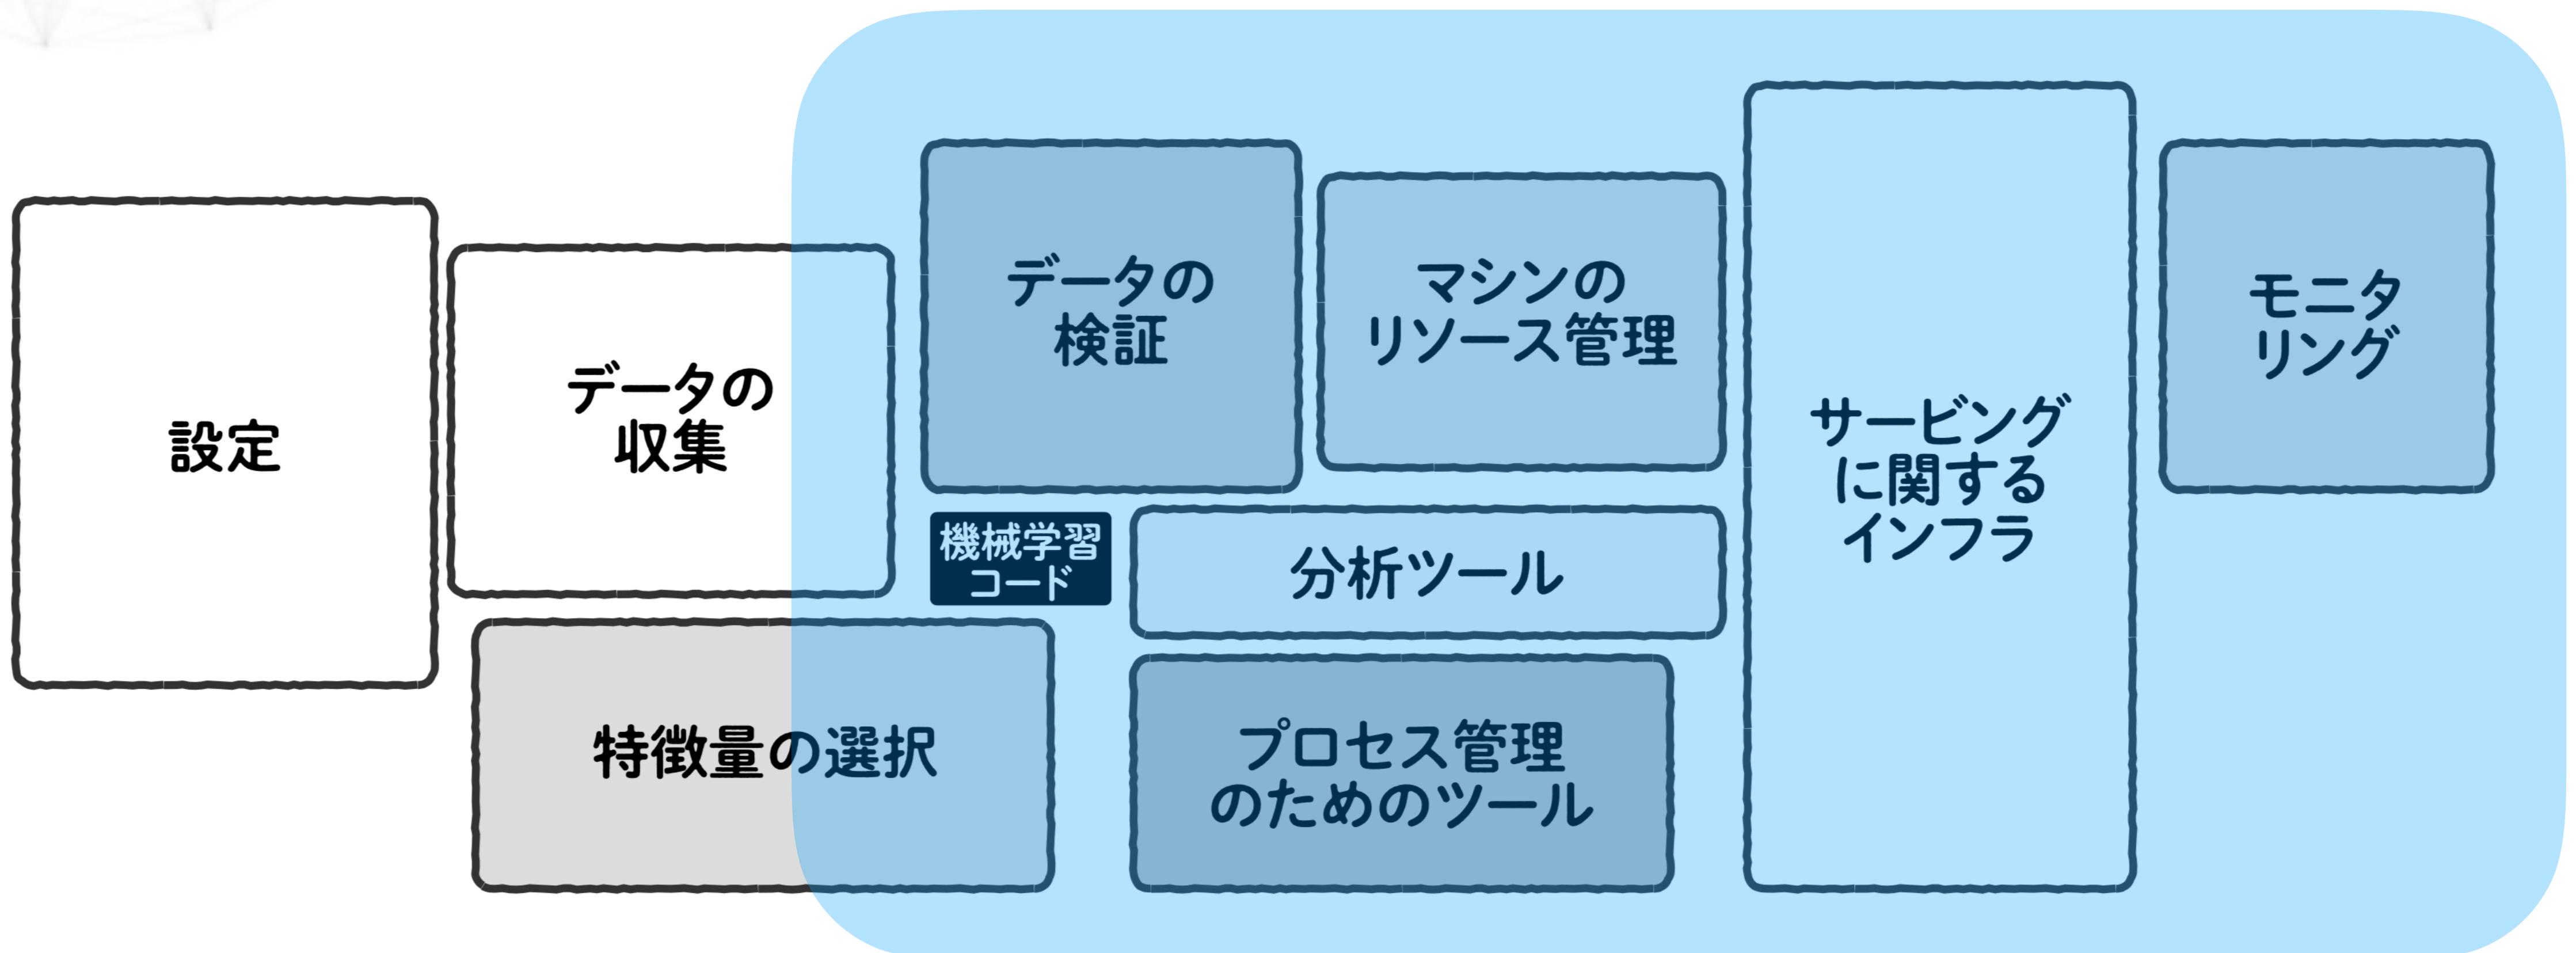
モデルの利用



▶ 図 3.1 機械学習のモデリングのコードの割合は少ない ([2] 中の図を翻訳して引用)

n月間ラムダノートVol.1, No.1より

モデルの利用

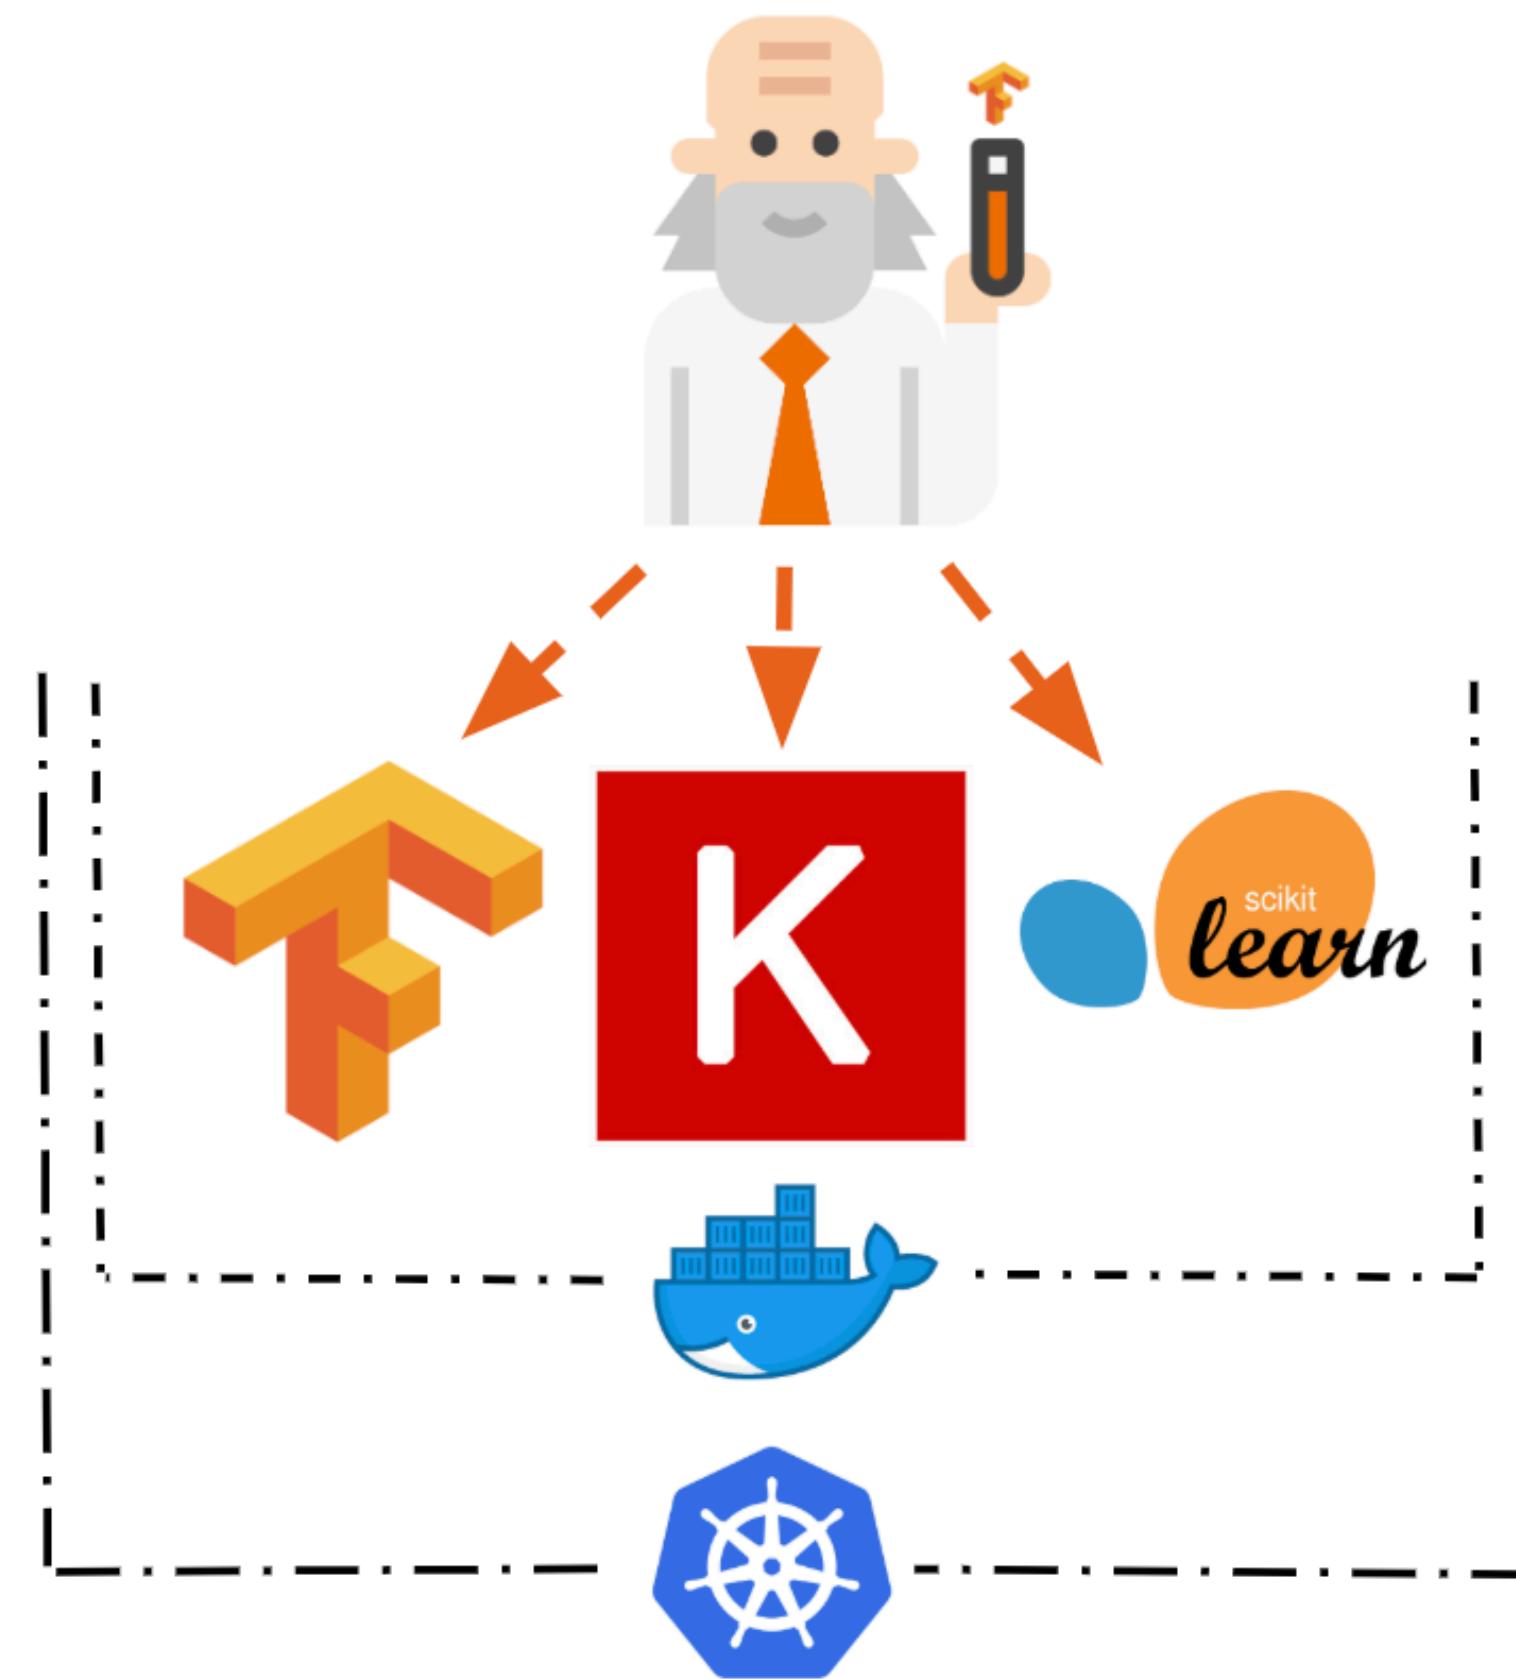


▶ 図 3.1 機械学習のモデリングのコードの割合は少ない ([2] 中の図を翻訳して引用)

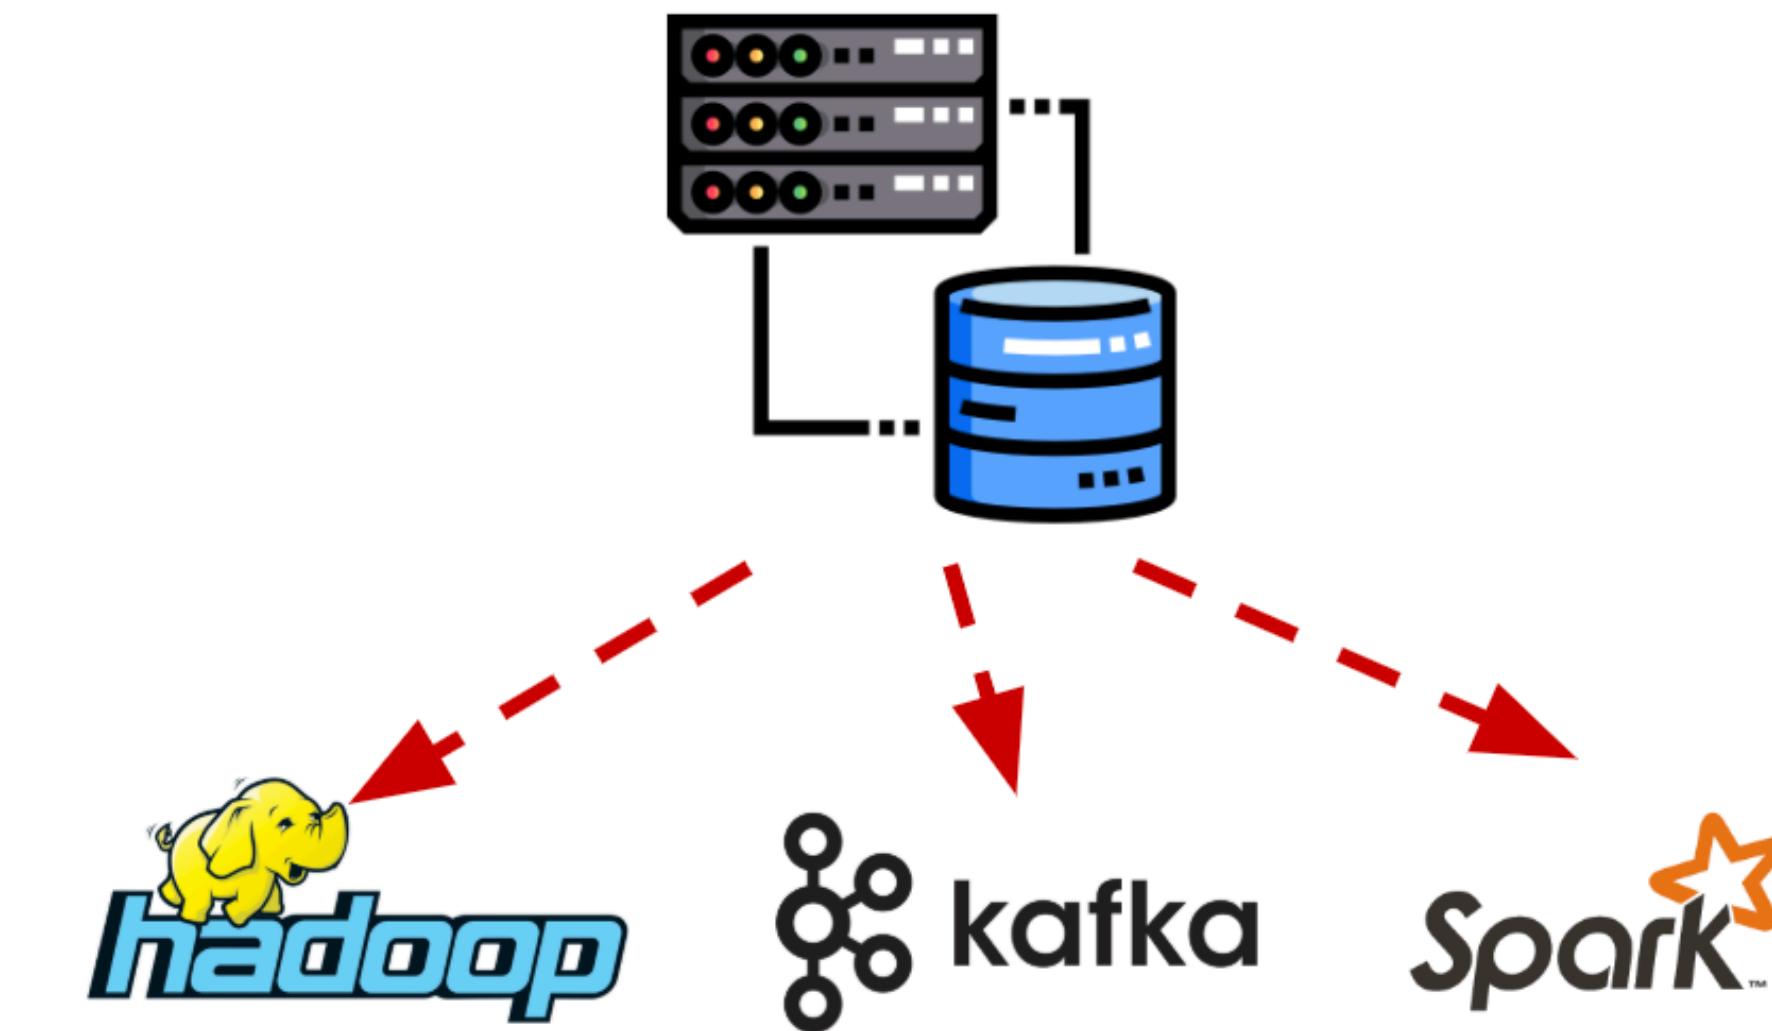
n月間ラムダノート Vol.1, No.1 より

Most AI Frameworks Not Built for Enterprise

Development

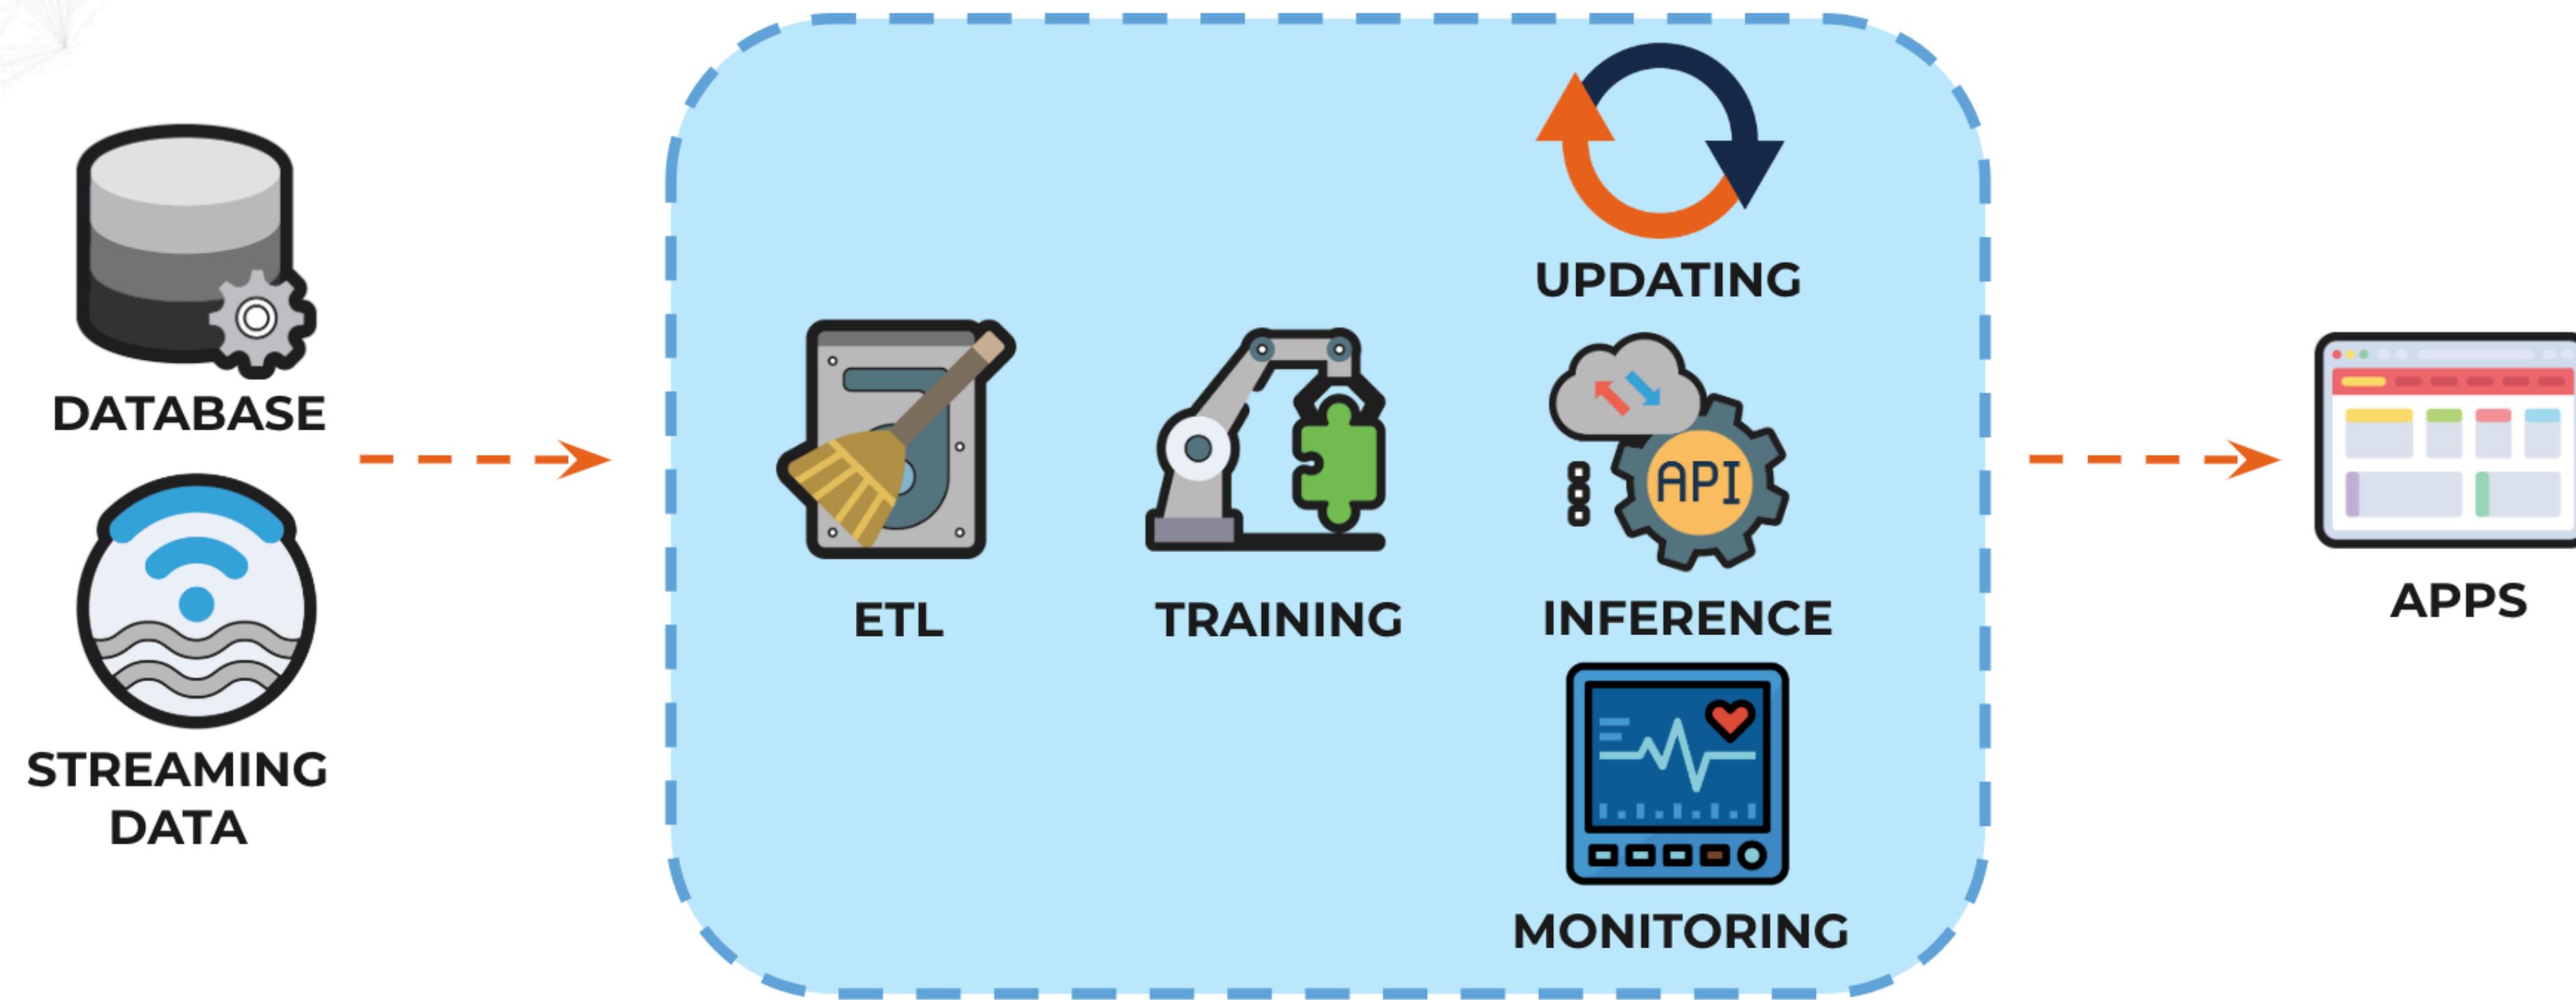


Production



- Petabytes of data
- Distributed clusters for data processing
- SLAs mandatory
 - Uptime guarantees!

SKIL provides a solution



Standardized AI infrastructure lets teams unify practices.

Configurable clusters: Adapts to opinionated ops teams and messy environments

A managed software distribution for popular ML/AI frameworks

Ingests AI models from any ML vendor/tool, allowing data scientists to keep favorite tools

Workspaces, Experiments

- **Workspaces** are a shared lab to conduct a set of “experiments” centered around a particular project.
- **Experiments** are different configurations of neural net models and data pipelines applied to a given problem

The screenshot displays the skymind SKIL User Interface. On the left, a sidebar menu lists: WORKSPACES (highlighted with a red box), DEPLOYMENTS, JOBS, CLUSTER NODES, PROCESSES, and PLUGINS. The main area shows two tabs: 'Workspaces' and 'Experiments'. The 'Workspaces' tab is active, showing a table with columns 'NAME' and 'CREATED'. One row is visible: Demo (ID: b84e7216-ff20-4ad0-bf54-c1d07d298d5e) created on Tue, Mar 12, 2019 1:43 PM, 2 minutes ago. The 'Experiments' tab shows a table with columns 'EXPERIMENT' and 'LAST UPDATED'. One row is visible: demo (Last updated: Tue, Mar 12, 2019 1:46 PM, 2 minutes ago). A 'NOTEBOOK LINK' button is next to it, and an 'Open Notebook' button is at the bottom right. The top navigation bar shows 'Workspaces / Demo'.

NAME	CREATED
Demo b84e7216-ff20-4ad0-bf54-c1d07d298d5e	Tue, Mar 12, 2019 1:43 PM 2 minutes ago

EXPERIMENT	LAST UPDATED
demo	Tue, Mar 12, 2019 1:46 PM 2 minutes ago

Deployments

- When mark a model as “deployed”, it serve as **REST API endpoint** to use by client applications
- Endpoint example: *http://[host]:9008/endpoints/deployment/model/kerasmodel/default*

The screenshot shows the skymind SKIL User Interface. On the left, there is a sidebar with the following navigation options:

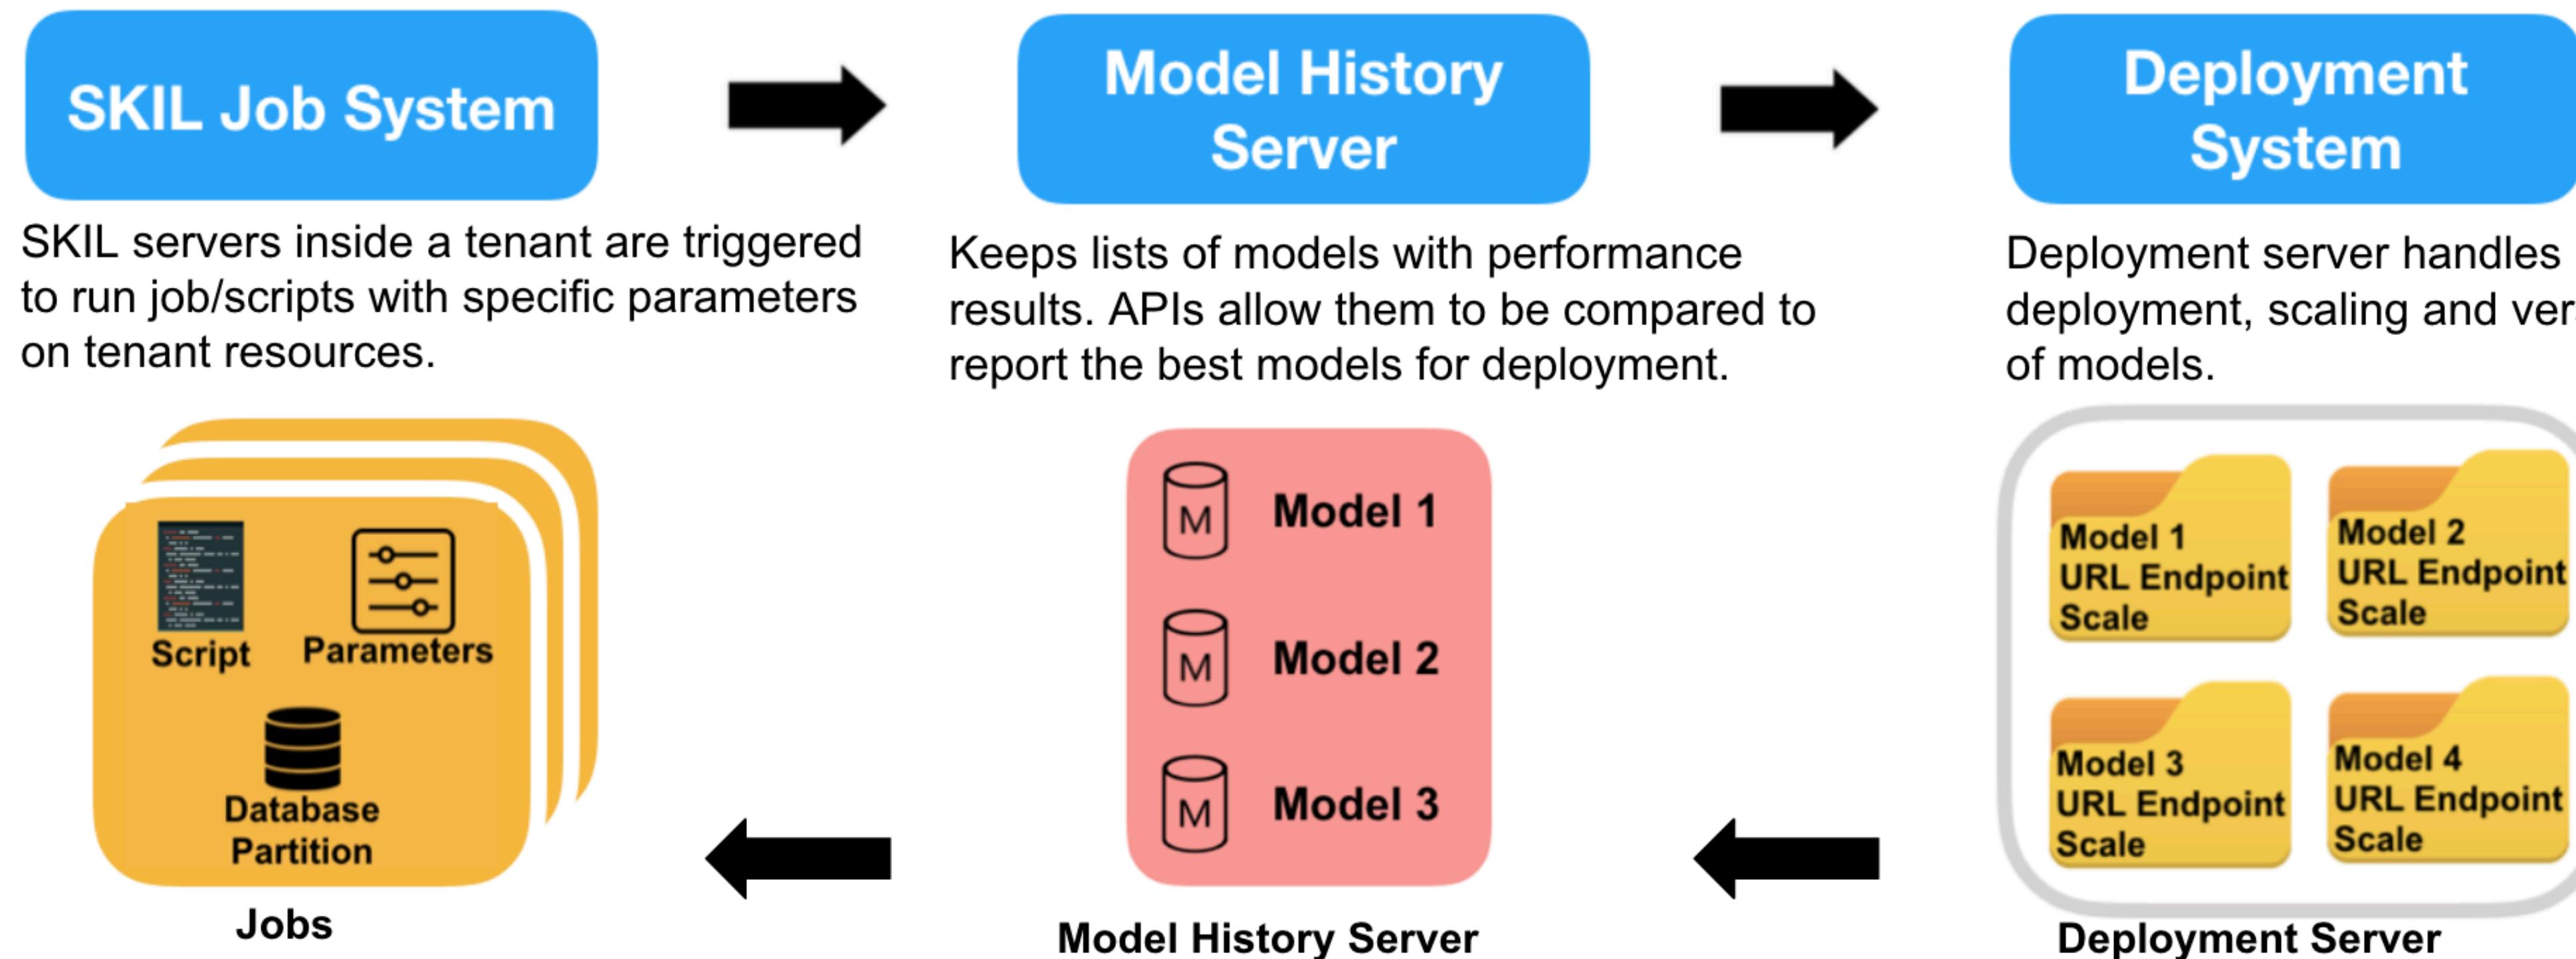
- WORKSPACES
- DEPLOYMENTS** (highlighted with a red border)
- JOBSS
- CLUSTER NODES
- PROCESSES
- PLUGINS

The main content area is titled "Deployments" and displays the message "Currently deployed models, transforms, and KNN". It contains a table with the following data:

NAME	STATUS	DETAILS	ACTIONS
mnist_demo	Not Deployed	1 Model , 0 Transforms, 0 KNN	Edit Delete

On the top right of the main content area, there are three icons: a gear (settings), a question mark (help), and a user icon labeled "Admin".

skymind | Key Components of Model Management



Best model on real data can
Be used with transfer learning
To fine tune model with latest
data.

Real-time feedback requests
are stored back in DB to monitor
model performance on real data stream

CNNモデルのディープ□イ

Best Model

NAME, ID

CREATED

LABELS

Keras Fashion MNIST Model
67389ec8-60ad-11e9-8e8a-
5d8b0e8c00e1

Wed, Apr 17, 2019 10:09
AM
7 hours ago

Models

Recently created models

NAME, ID

CREATED

LABELS

Keras Fashion MNIST Model
67389ec8-60ad-11e9-8e8a-
5d8b0e8c00e1

Wed, Apr 17, 2019 10:09
AM
7 hours ago

Details - Keras Fashion MNIST Model

 Mark Best Delete Deploy**ID** 67389ec8-60ad-11e9-8e8a-5d8b0e8c00e1**Created** 7 hours ago**Labels****ETL JSON****Input
Formats****URI** file:///var/skil/storage/models/675e8d9a-60ad-11e9-
8e8a-5d8b0e8c00e1.h5**Evaluation Results**

CREATED	NAME	ACCURACY	PRECISION	F1	AUC	RECALL	RM
---------	------	----------	-----------	----	-----	--------	----

7 hours ago	Keras mlp 10 epochs	0.9115	0	0.9122	0	0.9115	0
-------------	---------------------------	--------	---	--------	---	--------	---

Revision History

CREATED

NAME, ID

skymind SKIL

WORKSPACES

DEPLOYMENTS

J OBS

CLUSTER NODES

PROCESSES

PLUGINS

fashionmnist Stopped 1 Stopped http://52.15.10.10:8080

Start

Models

NAME **STATUS** **SCALE** **METRICS** **ENDPOINT**

fashionmnist Stopped 1 Stopped http://52.15.10.10:8080

Start

Transforms

NAME **STATUS** **SCALE** **METRICS** **ENDPOINT**

No results

fashionmnist

STATUS	SCALE	ID
Starting	1	5

Usage

Endpoints Java Python RPA

http://52.15.103.124:9008/endpoints/fashionmnist/model/fashionmnist/default

```
curl -F image=@image.jpg -H "Authorization: Bearer <auth_token>" "http://52.15.103.1
```

Copy

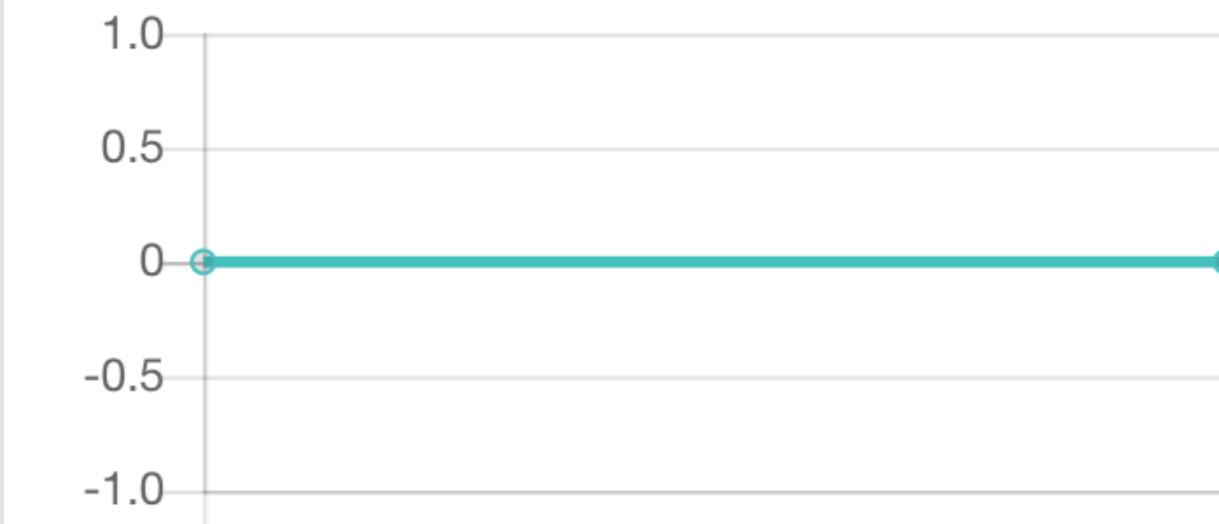
Version History

VERSION NUMBER CREATED UPDATED

0	Wed, Apr 17, 2019 10:14 AM	a few seconds ago
	7 hours ago	

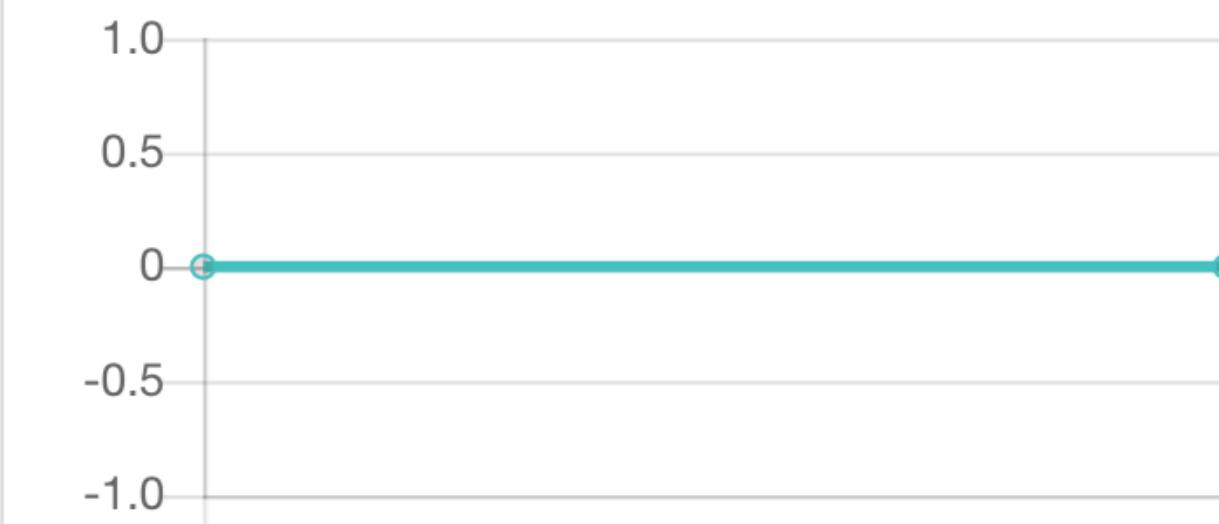
```
curl --request POST \
--url http://<SKILhostIP>:9008/login \
--header 'accept: application/json' \
--header 'content-type: application/json' \
--data '{"userId":<userId>,"password":<password>}'
```

Transactions



Not started

Latency



Links

- SKIL Download Page
 - <https://skymind.ai/download>
- SKIL Documentation
 - <https://docs.skymind.ai/docs>
- SKIL REST API reference
 - <https://docs.skymind.ai/reference?newApiExplorer=true>
- SKIL Clients
 - <https://github.com/SkymindIO/skil-clients>



まとめ