

畳み込みニューラル ネットワークの基礎と応用

Tableau データサイエンス勉強会第4回 一画像認識とAIの巻一

<https://techplay.jp/event/750555>

2019-10-21 (Mon) @ TECH PLAY SHIBUYA

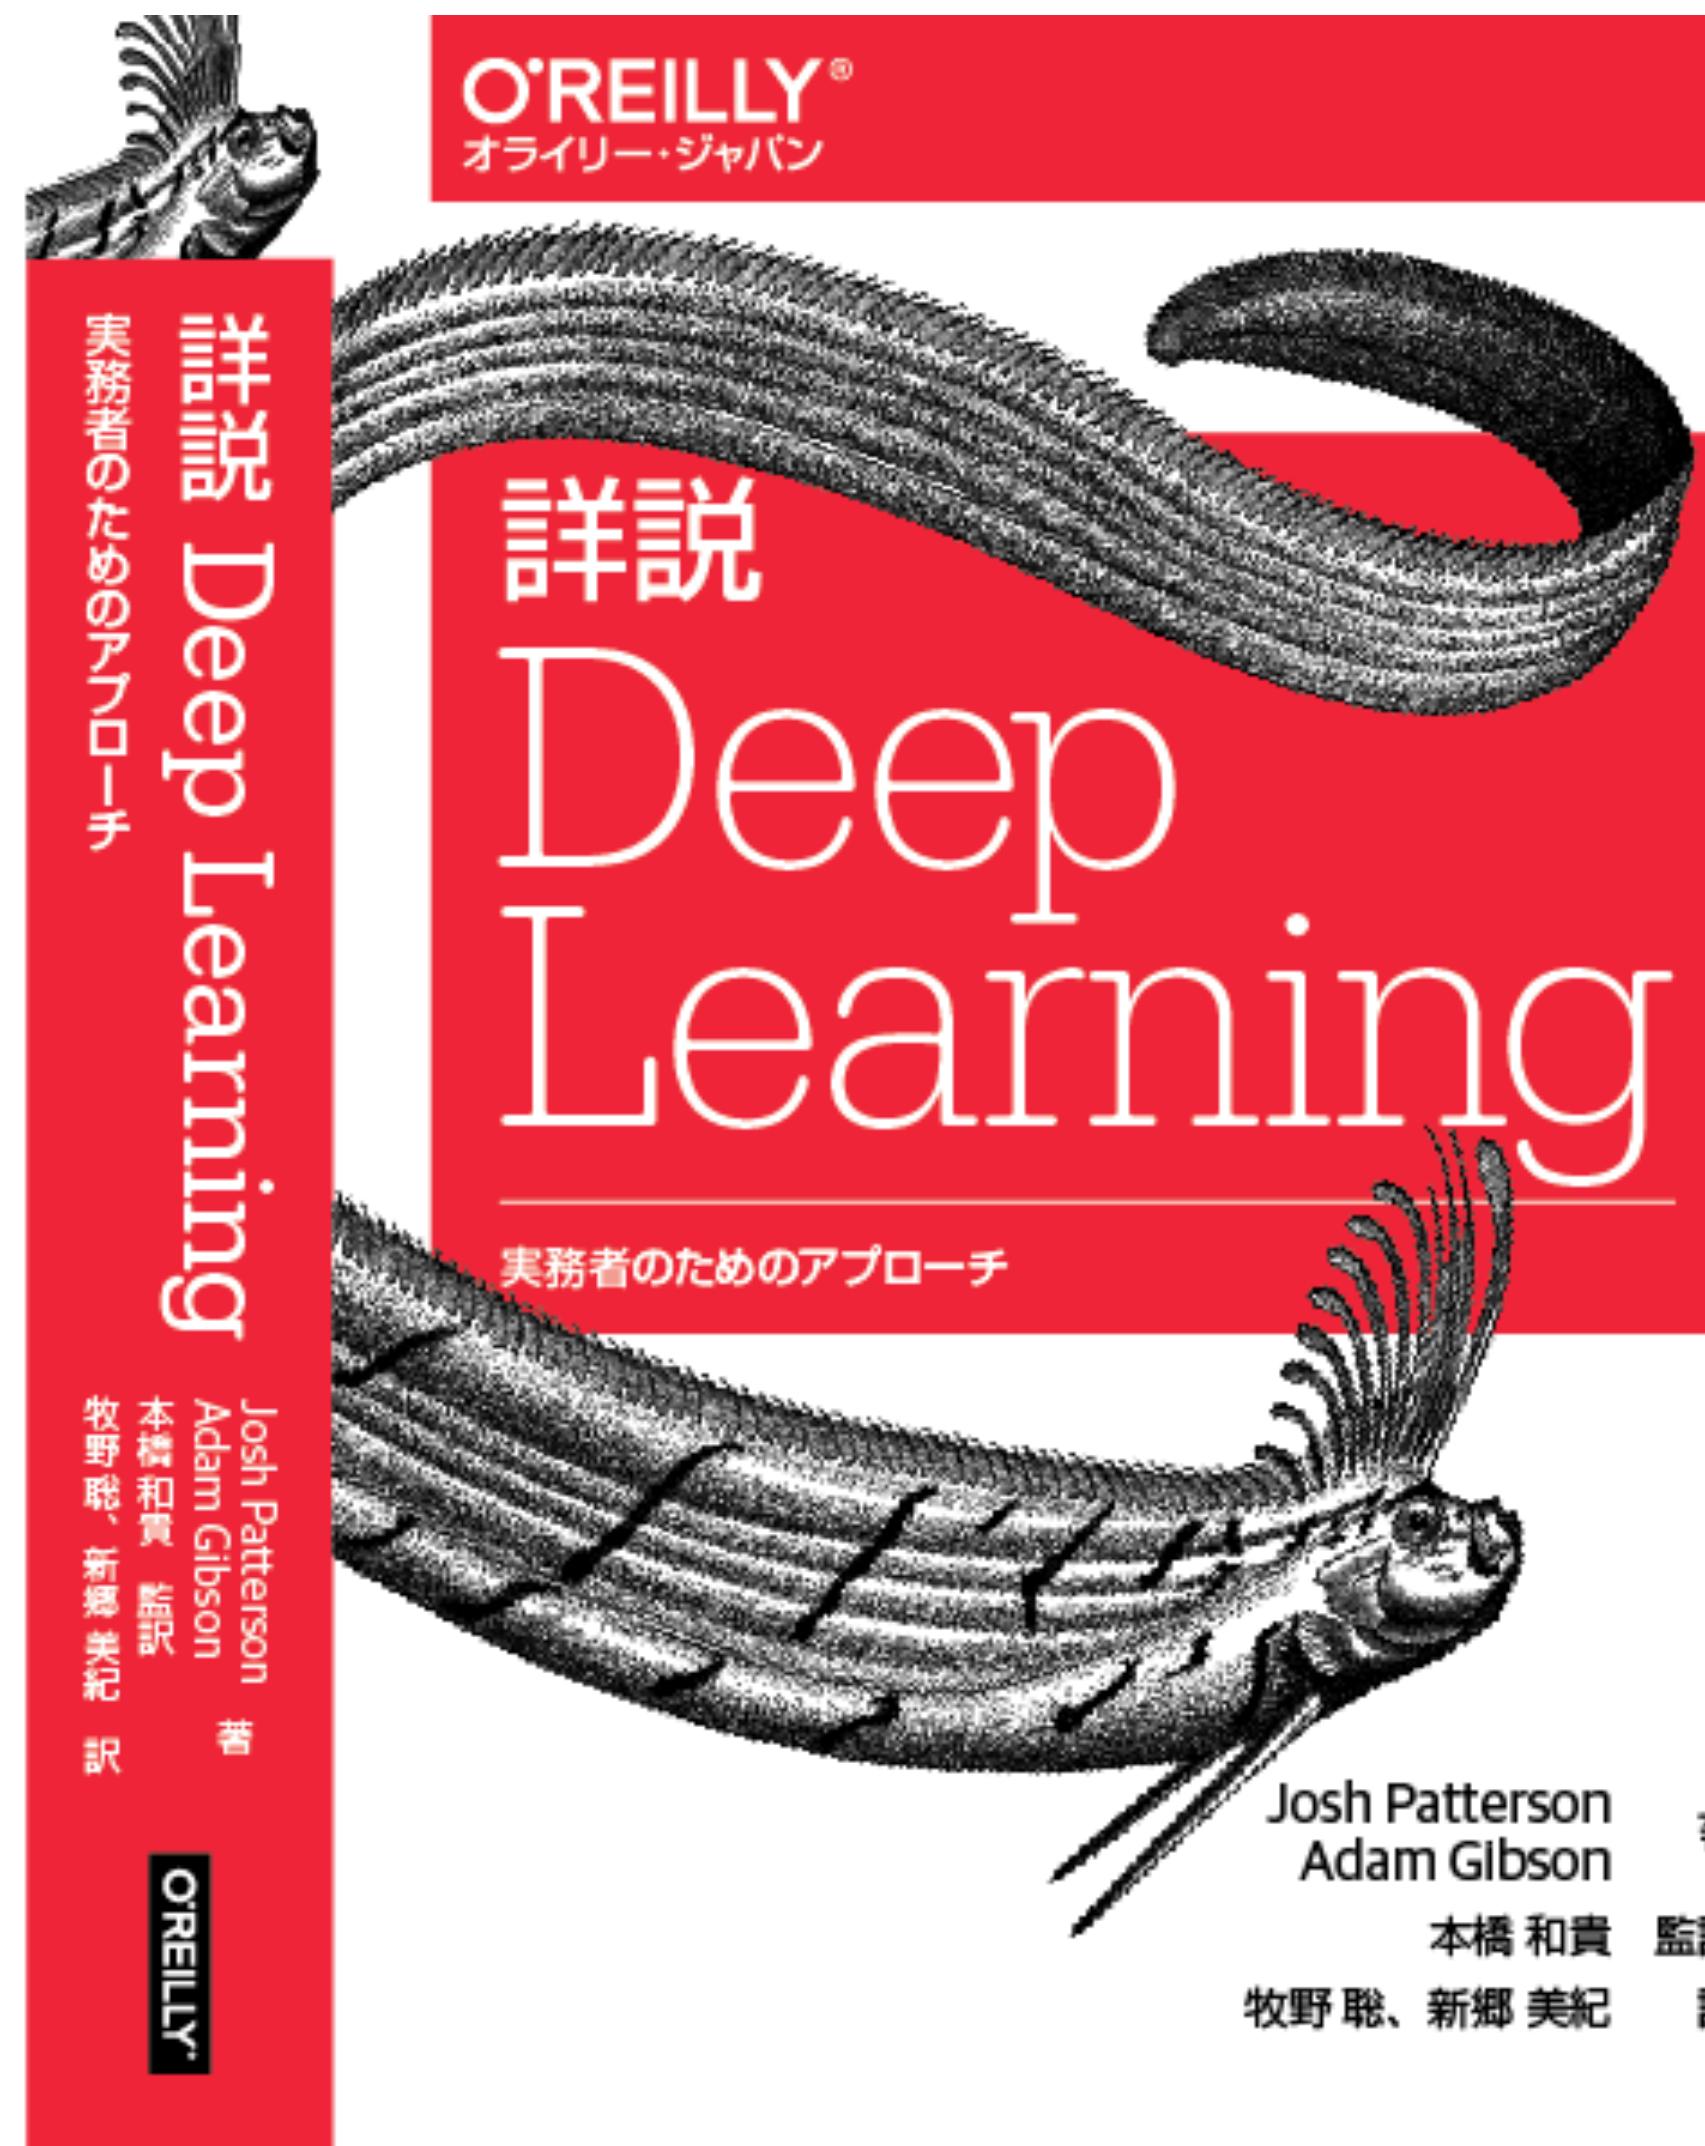
自己紹介

▶ 本橋 和貴 @kmotohas

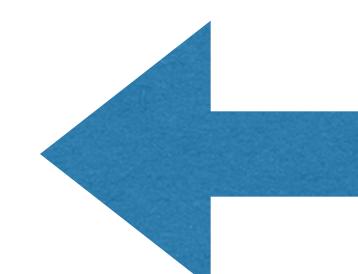
- スカイマインド株式会社
 - Deep Learning Engineer (前職ではDL+ROS)
- 素粒子物理学実験 (LHC-ATLAS実験) 出身
 - 博士 (理学)
- 好きな本 : 詳説 Deep Learning – 実務者のためのアプローチ



2019年8月9日発売



- 原著 “Deep Learning – A Practitioner’s Approach” は 2017年8月発売
- JVM言語用ディープラーニング開発フレームワーク Deeplearning4j (DL4J) を用いた解説書
 - 著者は DL4J の開発者 Adam Gibson、Skymind Inc を創業
 - ソフトウェア/アプリケーション/システム・エンジニアなどがメインターゲット
 - ディープラーニングの基礎からHadoop/Sparkといったビッグデータ分析基盤との連携まで解説



Agenda

機械学習の導入

Kerasを用いたコーディングサンプル

畳み込みニューラルネットワークの基礎

CNNのフィルタの理解

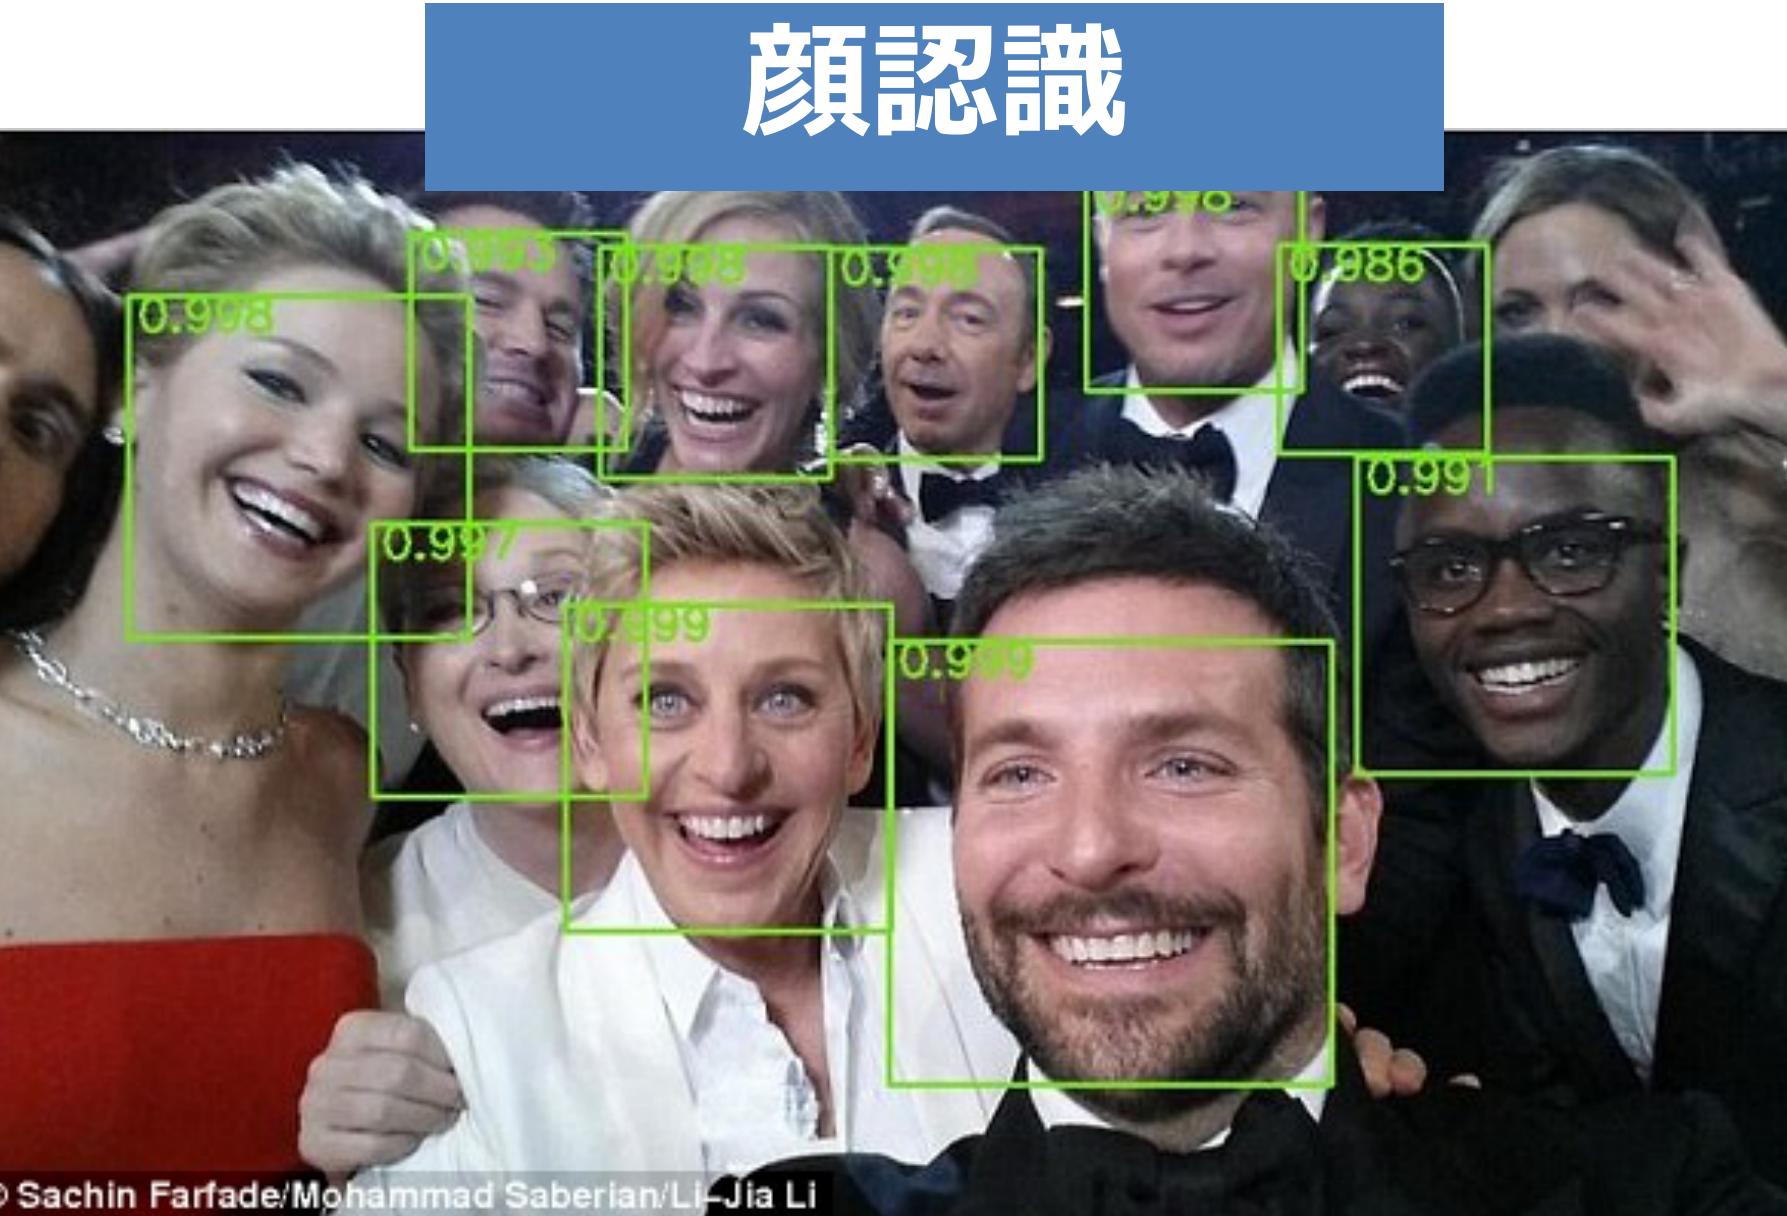
CNNアーキテクチャの紹介

skymind | 画像認識のアプリケーション

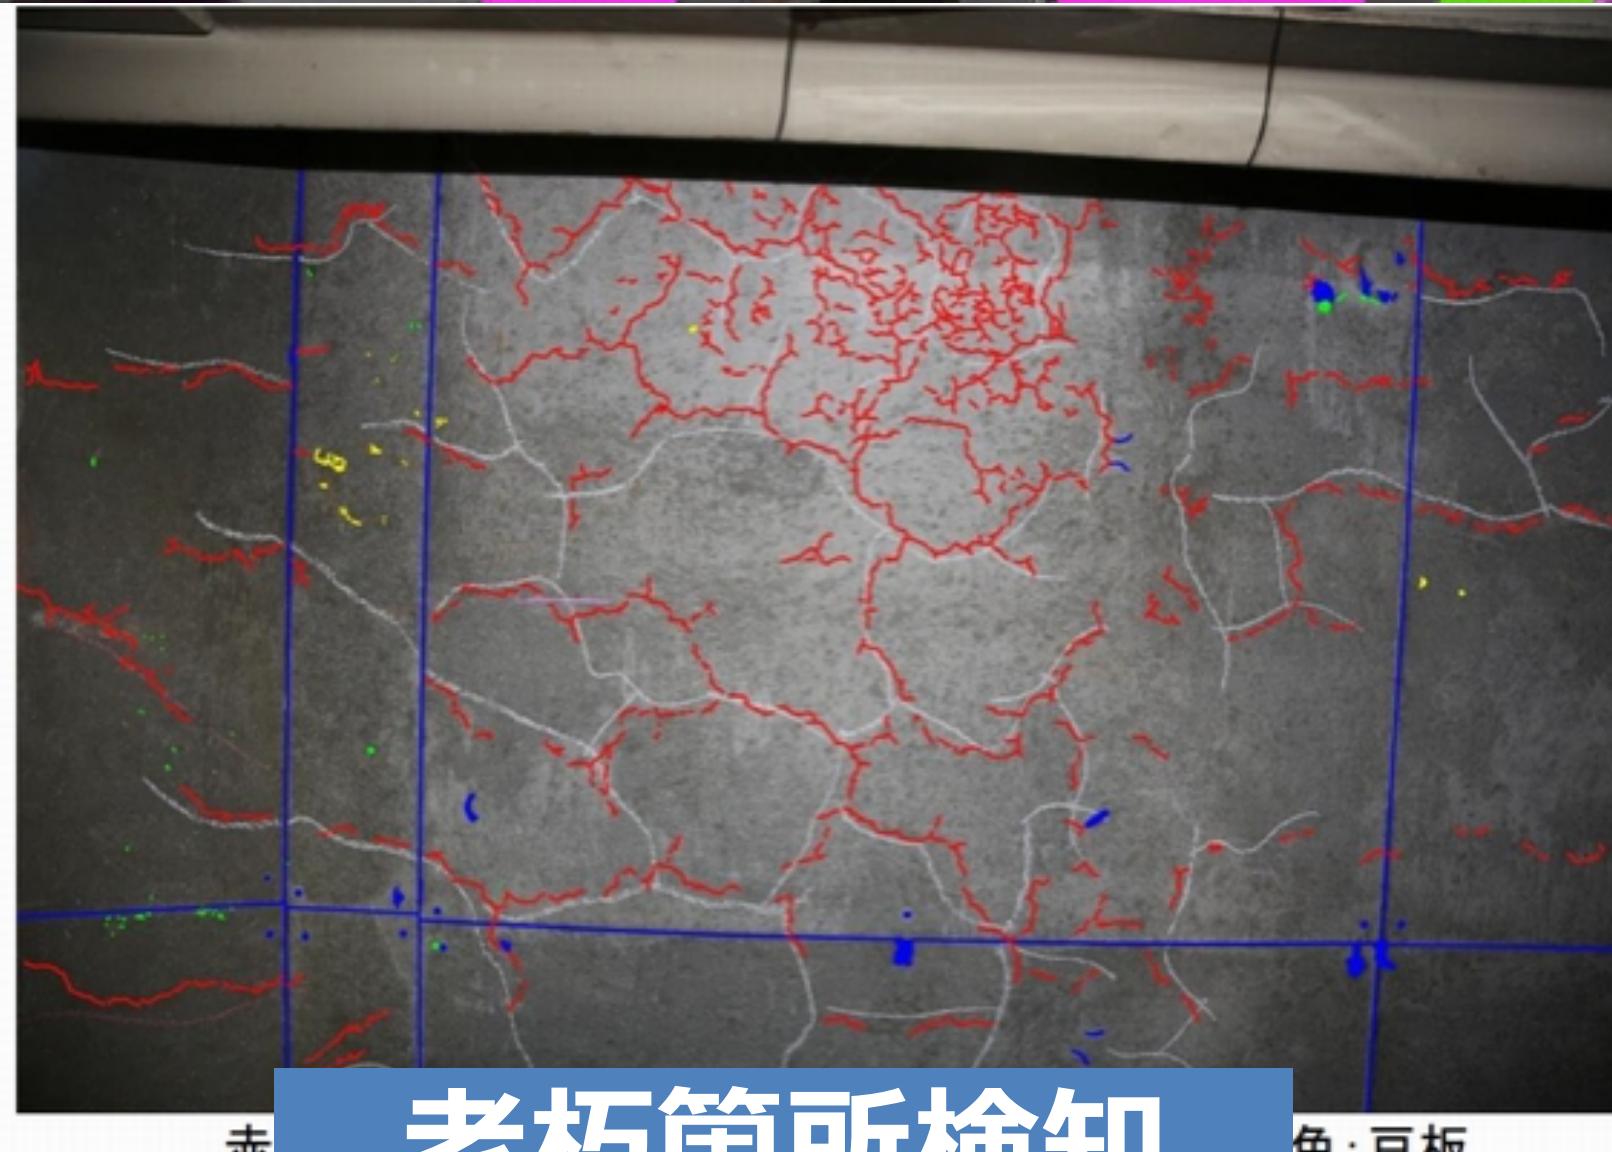
一般物体認識



顔認識



年齢推定

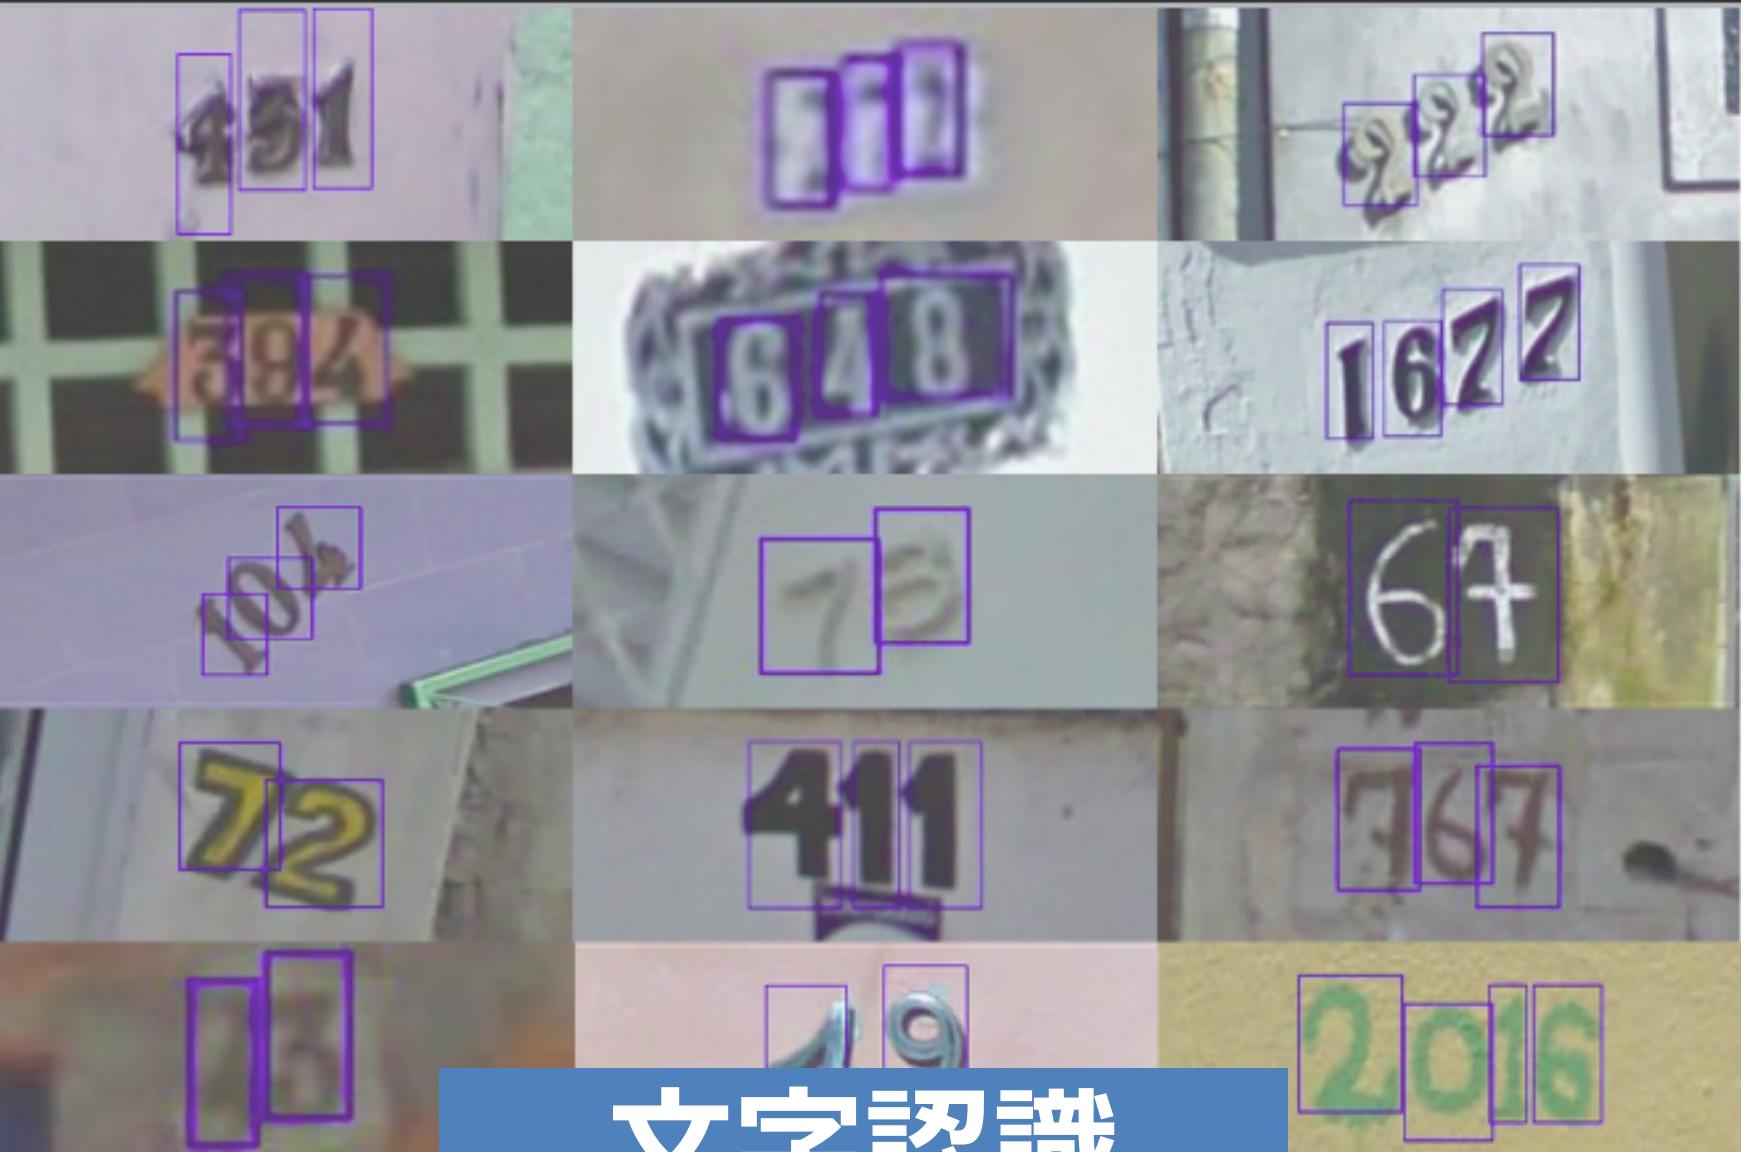


老朽箇所検知

色:豆板



セグメンテーション



文字認識



ニューラルネットワークの基礎



<https://www.coursera.org/learn/introduction-tensorflow>

Activity Recognition



```
if(speed<4){  
    status=WALKING;  
}
```

<https://www.coursera.org/learn/introduction-tensorflow>

Activity Recognition



```
if(speed<4){  
    status=WALKING;  
}
```

```
if(speed<4){  
    status=WALKING;  
} else {  
    status=RUNNING;  
}
```

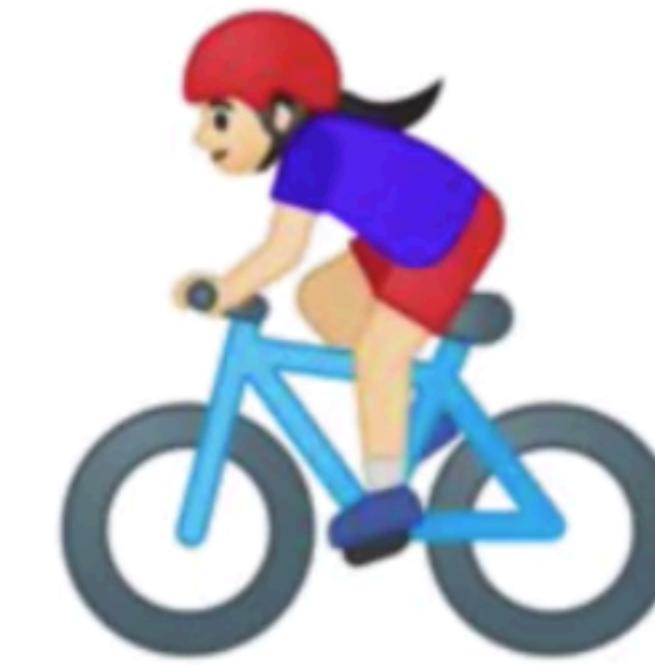
Activity Recognition



```
if(speed<4){  
    status=WALKING;  
}
```



```
if(speed<4){  
    status=WALKING;  
} else {  
    status=RUNNING;  
}
```



```
if(speed<4){  
    status=WALKING;  
} else if(speed<12){  
    status=RUNNING;  
} else {  
    status=BIKING;  
}
```

Activity Recognition



```
if(speed<4){  
    status=WALKING;  
}
```



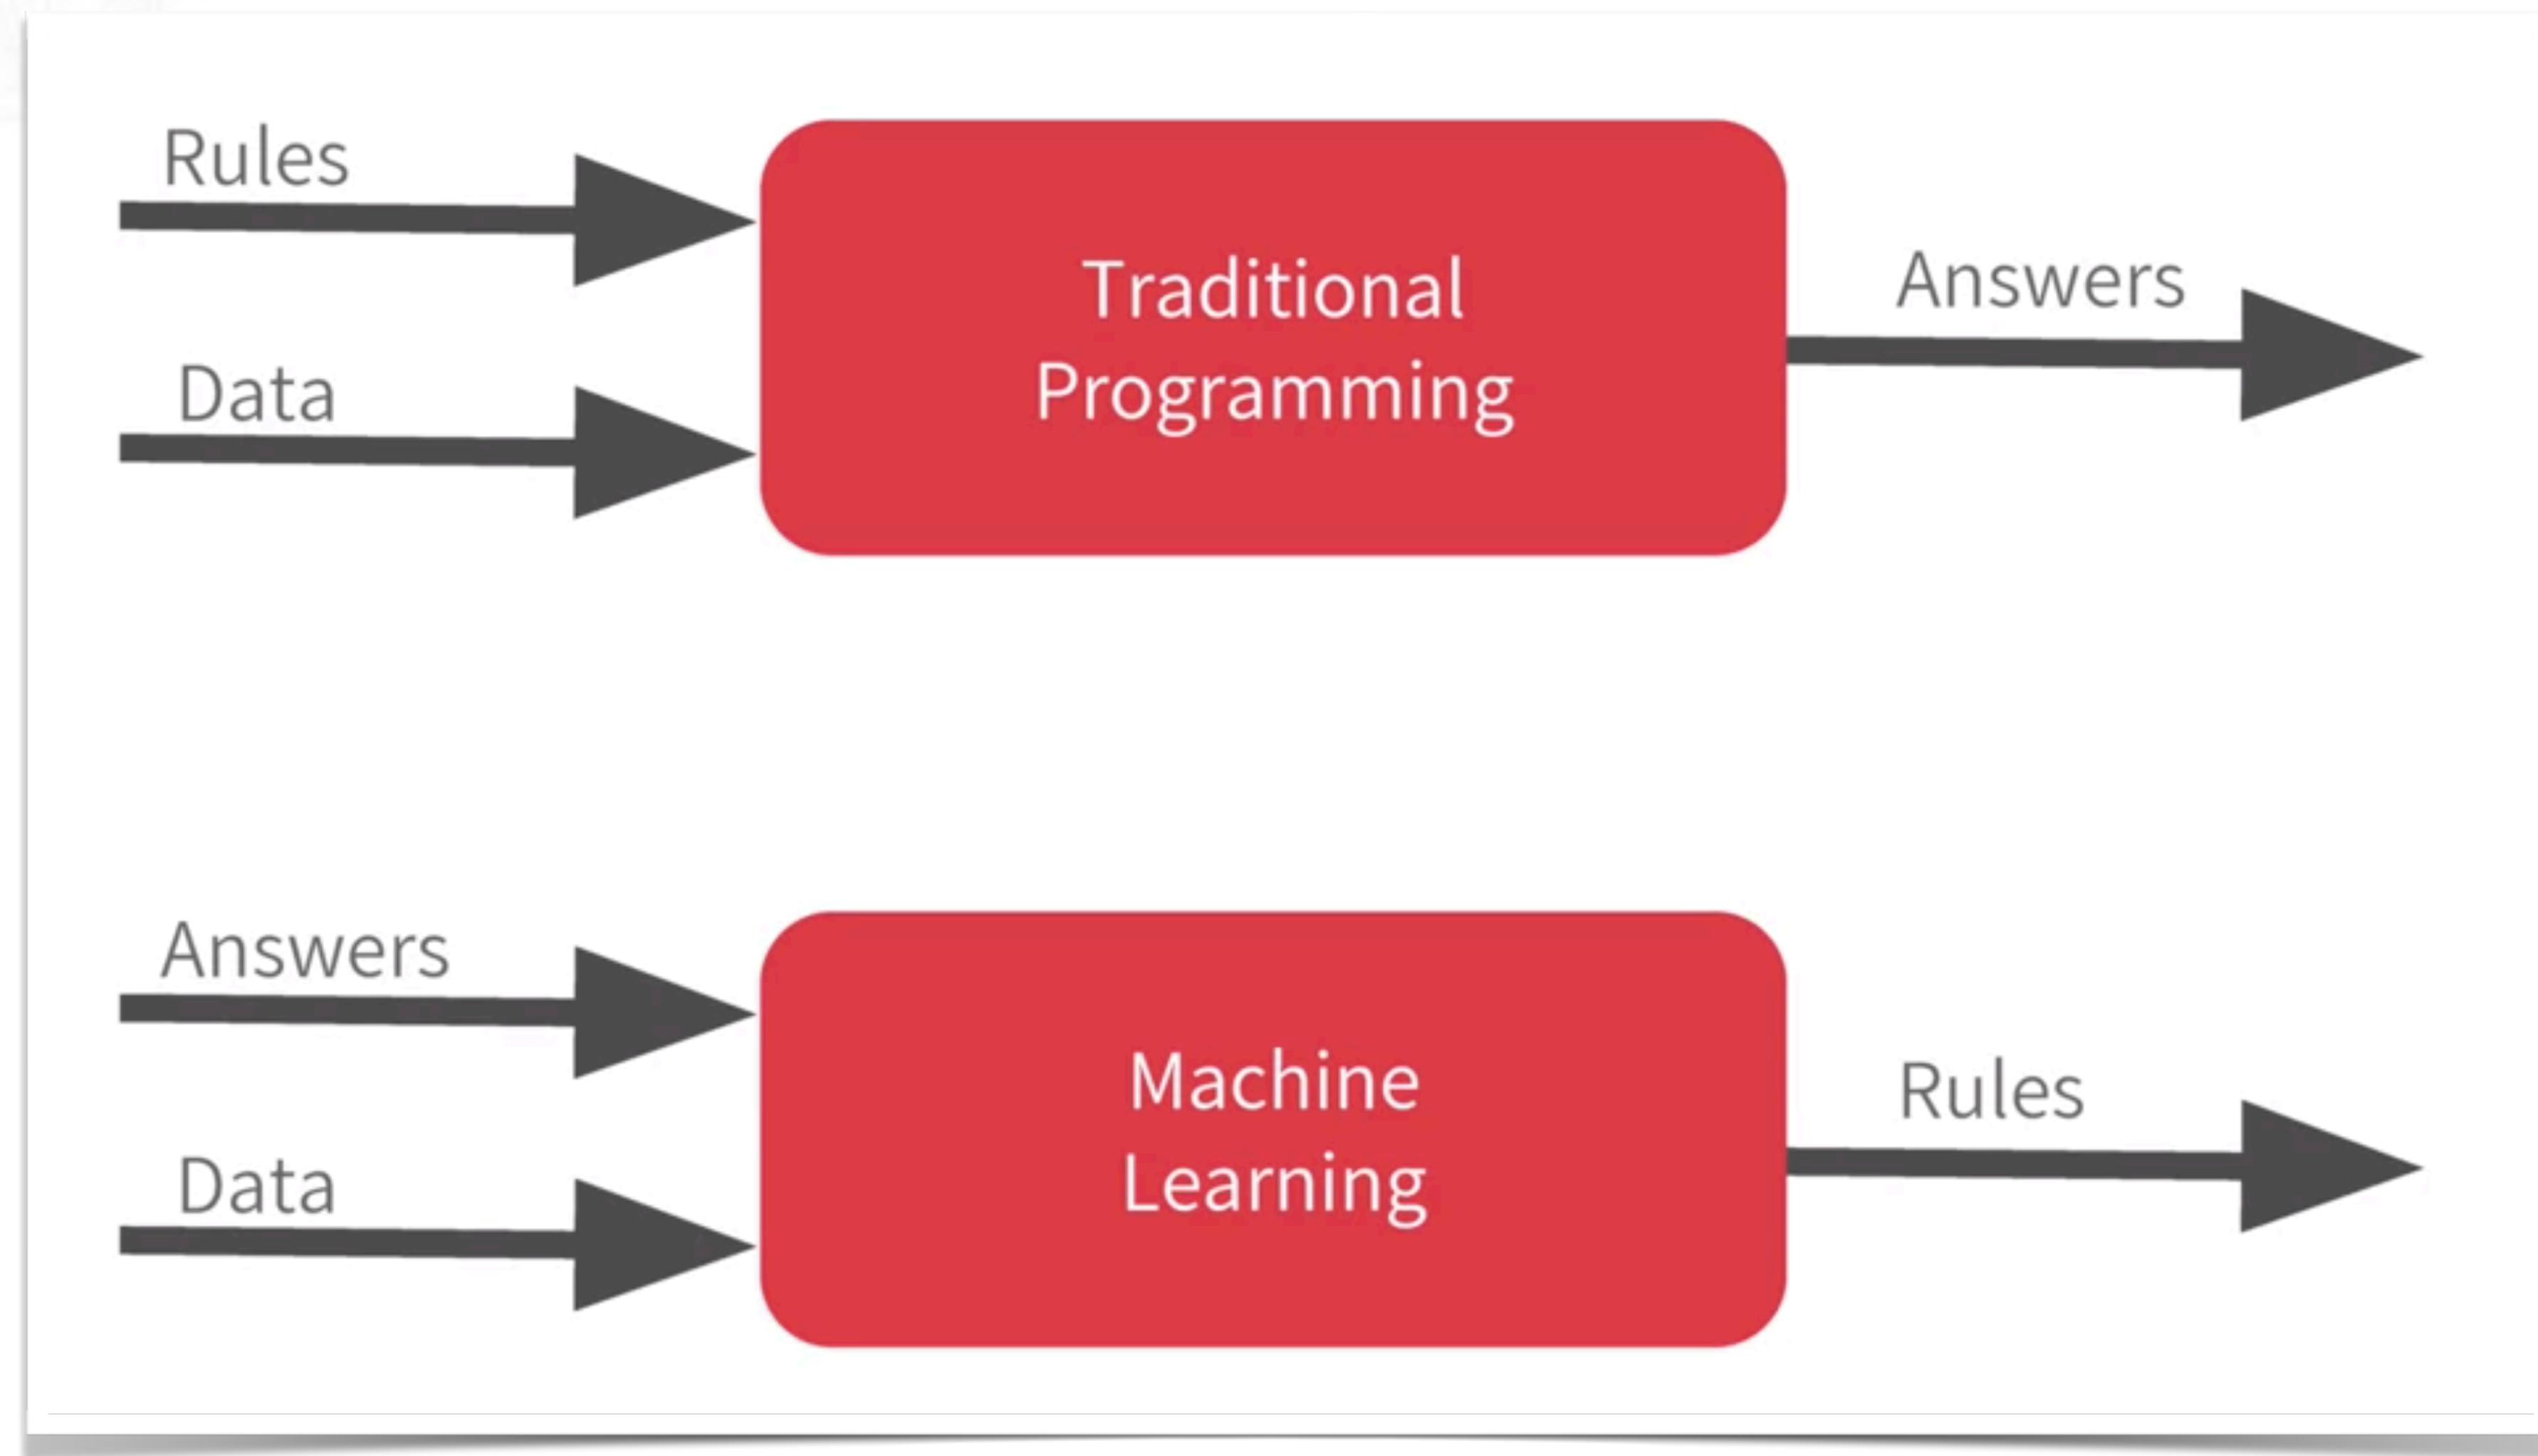
```
if(speed<4){  
    status=WALKING;  
} else {  
    status=RUNNING;  
}
```



```
if(speed<4){  
    status=WALKING;  
} else if(speed<12){  
    status=RUNNING;  
} else {  
    status=BIKING;  
}
```



// Oh crap



<https://www.coursera.org/learn/introduction-tensorflow>

Activity Recognition



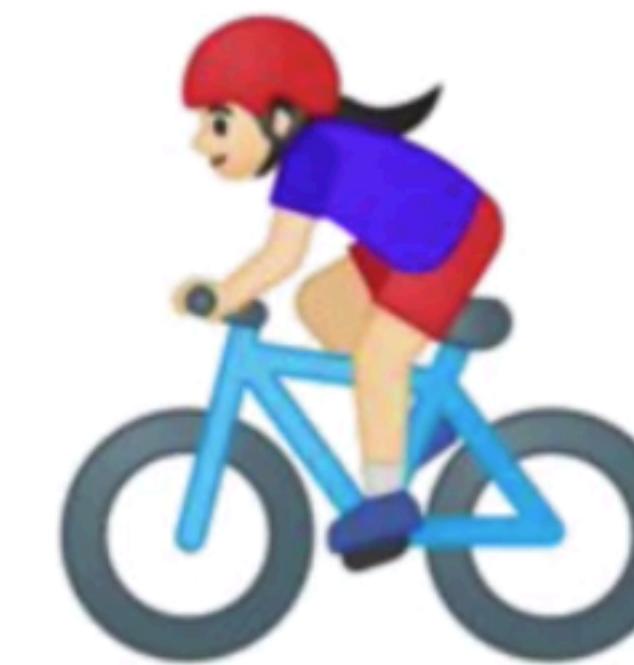
0101001010100101010
1001010101001011101
0100101010010101001
0101001010100101010

Label = WALKING



1010100101001010101
0101010010010010001
00100111101010111
1010100100111101011

Label = RUNNING



1001010011111010101
1101010111010101110
1010101111010101011
1111110001111010101

Label = BIKING



111111111010011101
0011111010111110101
0101110101010101110
1010101010100111110

Label = GOLFING
(Sort of)

機械学習の"Hello World"

$x = -2, -1, 0, 1, 2, 3, 4$

$y = -3, -1, 1, 3, 5, 7, 9$

$$y = f(x)$$

機械学習の"Hello World"

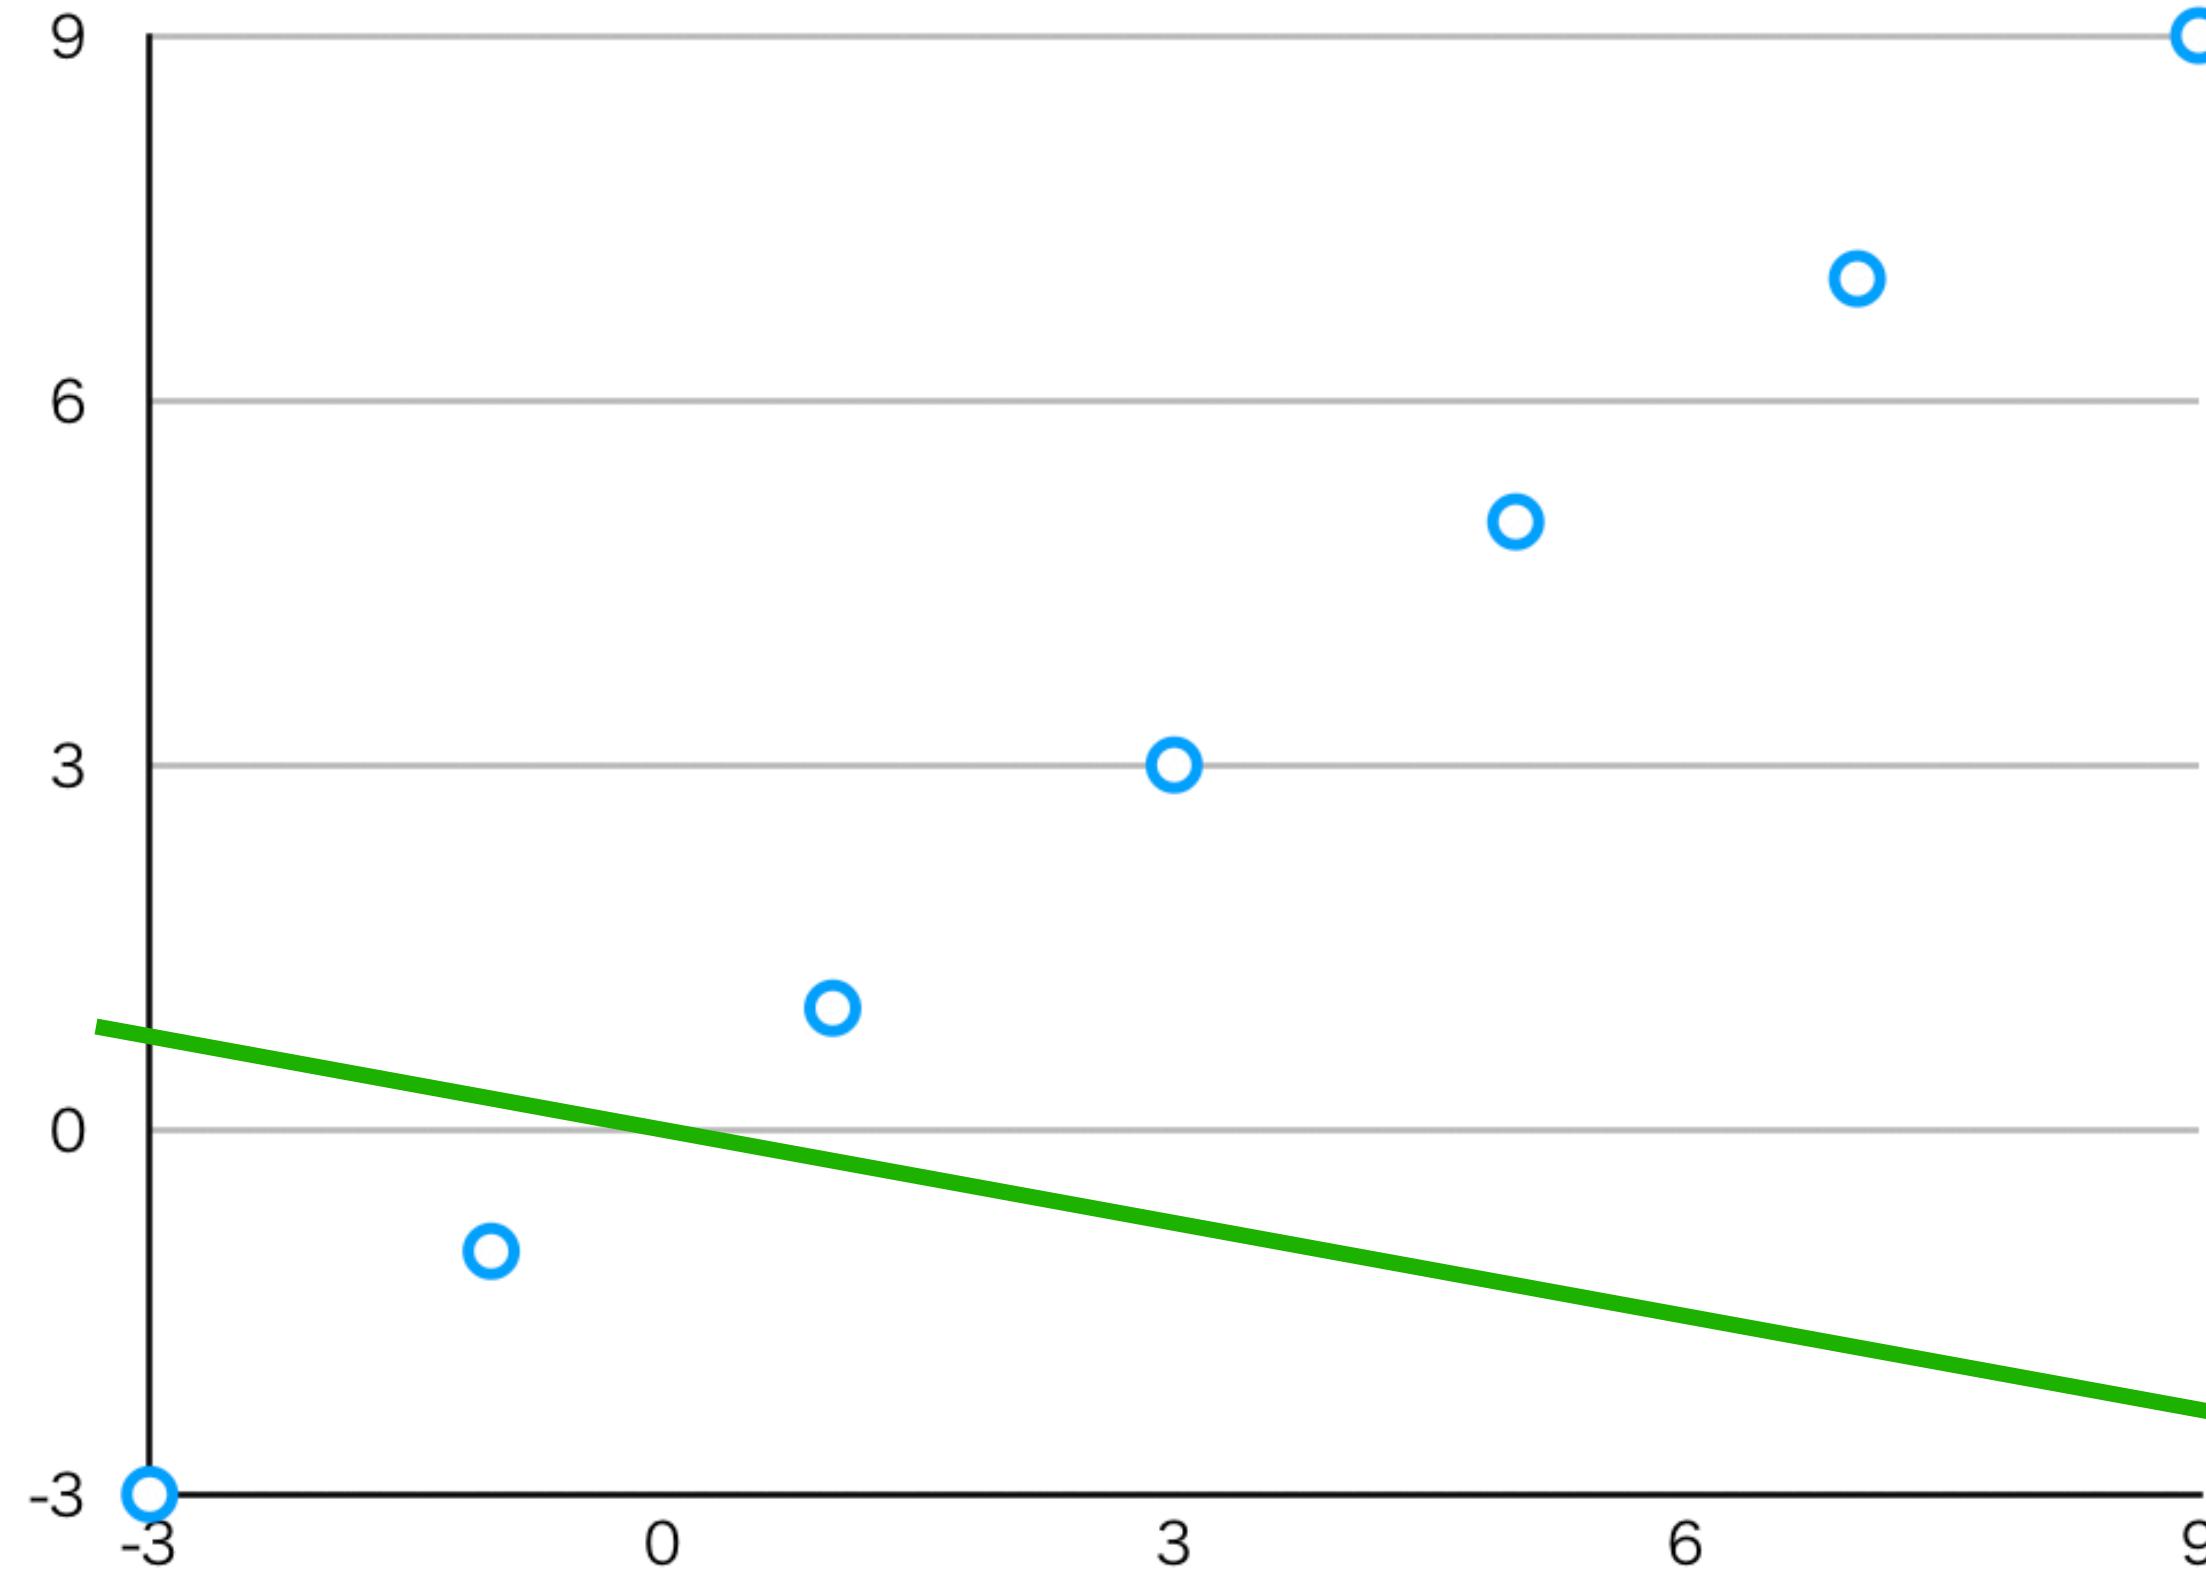
$x = -2, -1, 0, 1, 2, 3, 4$

$y = -3, -1, 1, 3, 5, 7, 9$

$y = f(x) = 2x + 1$

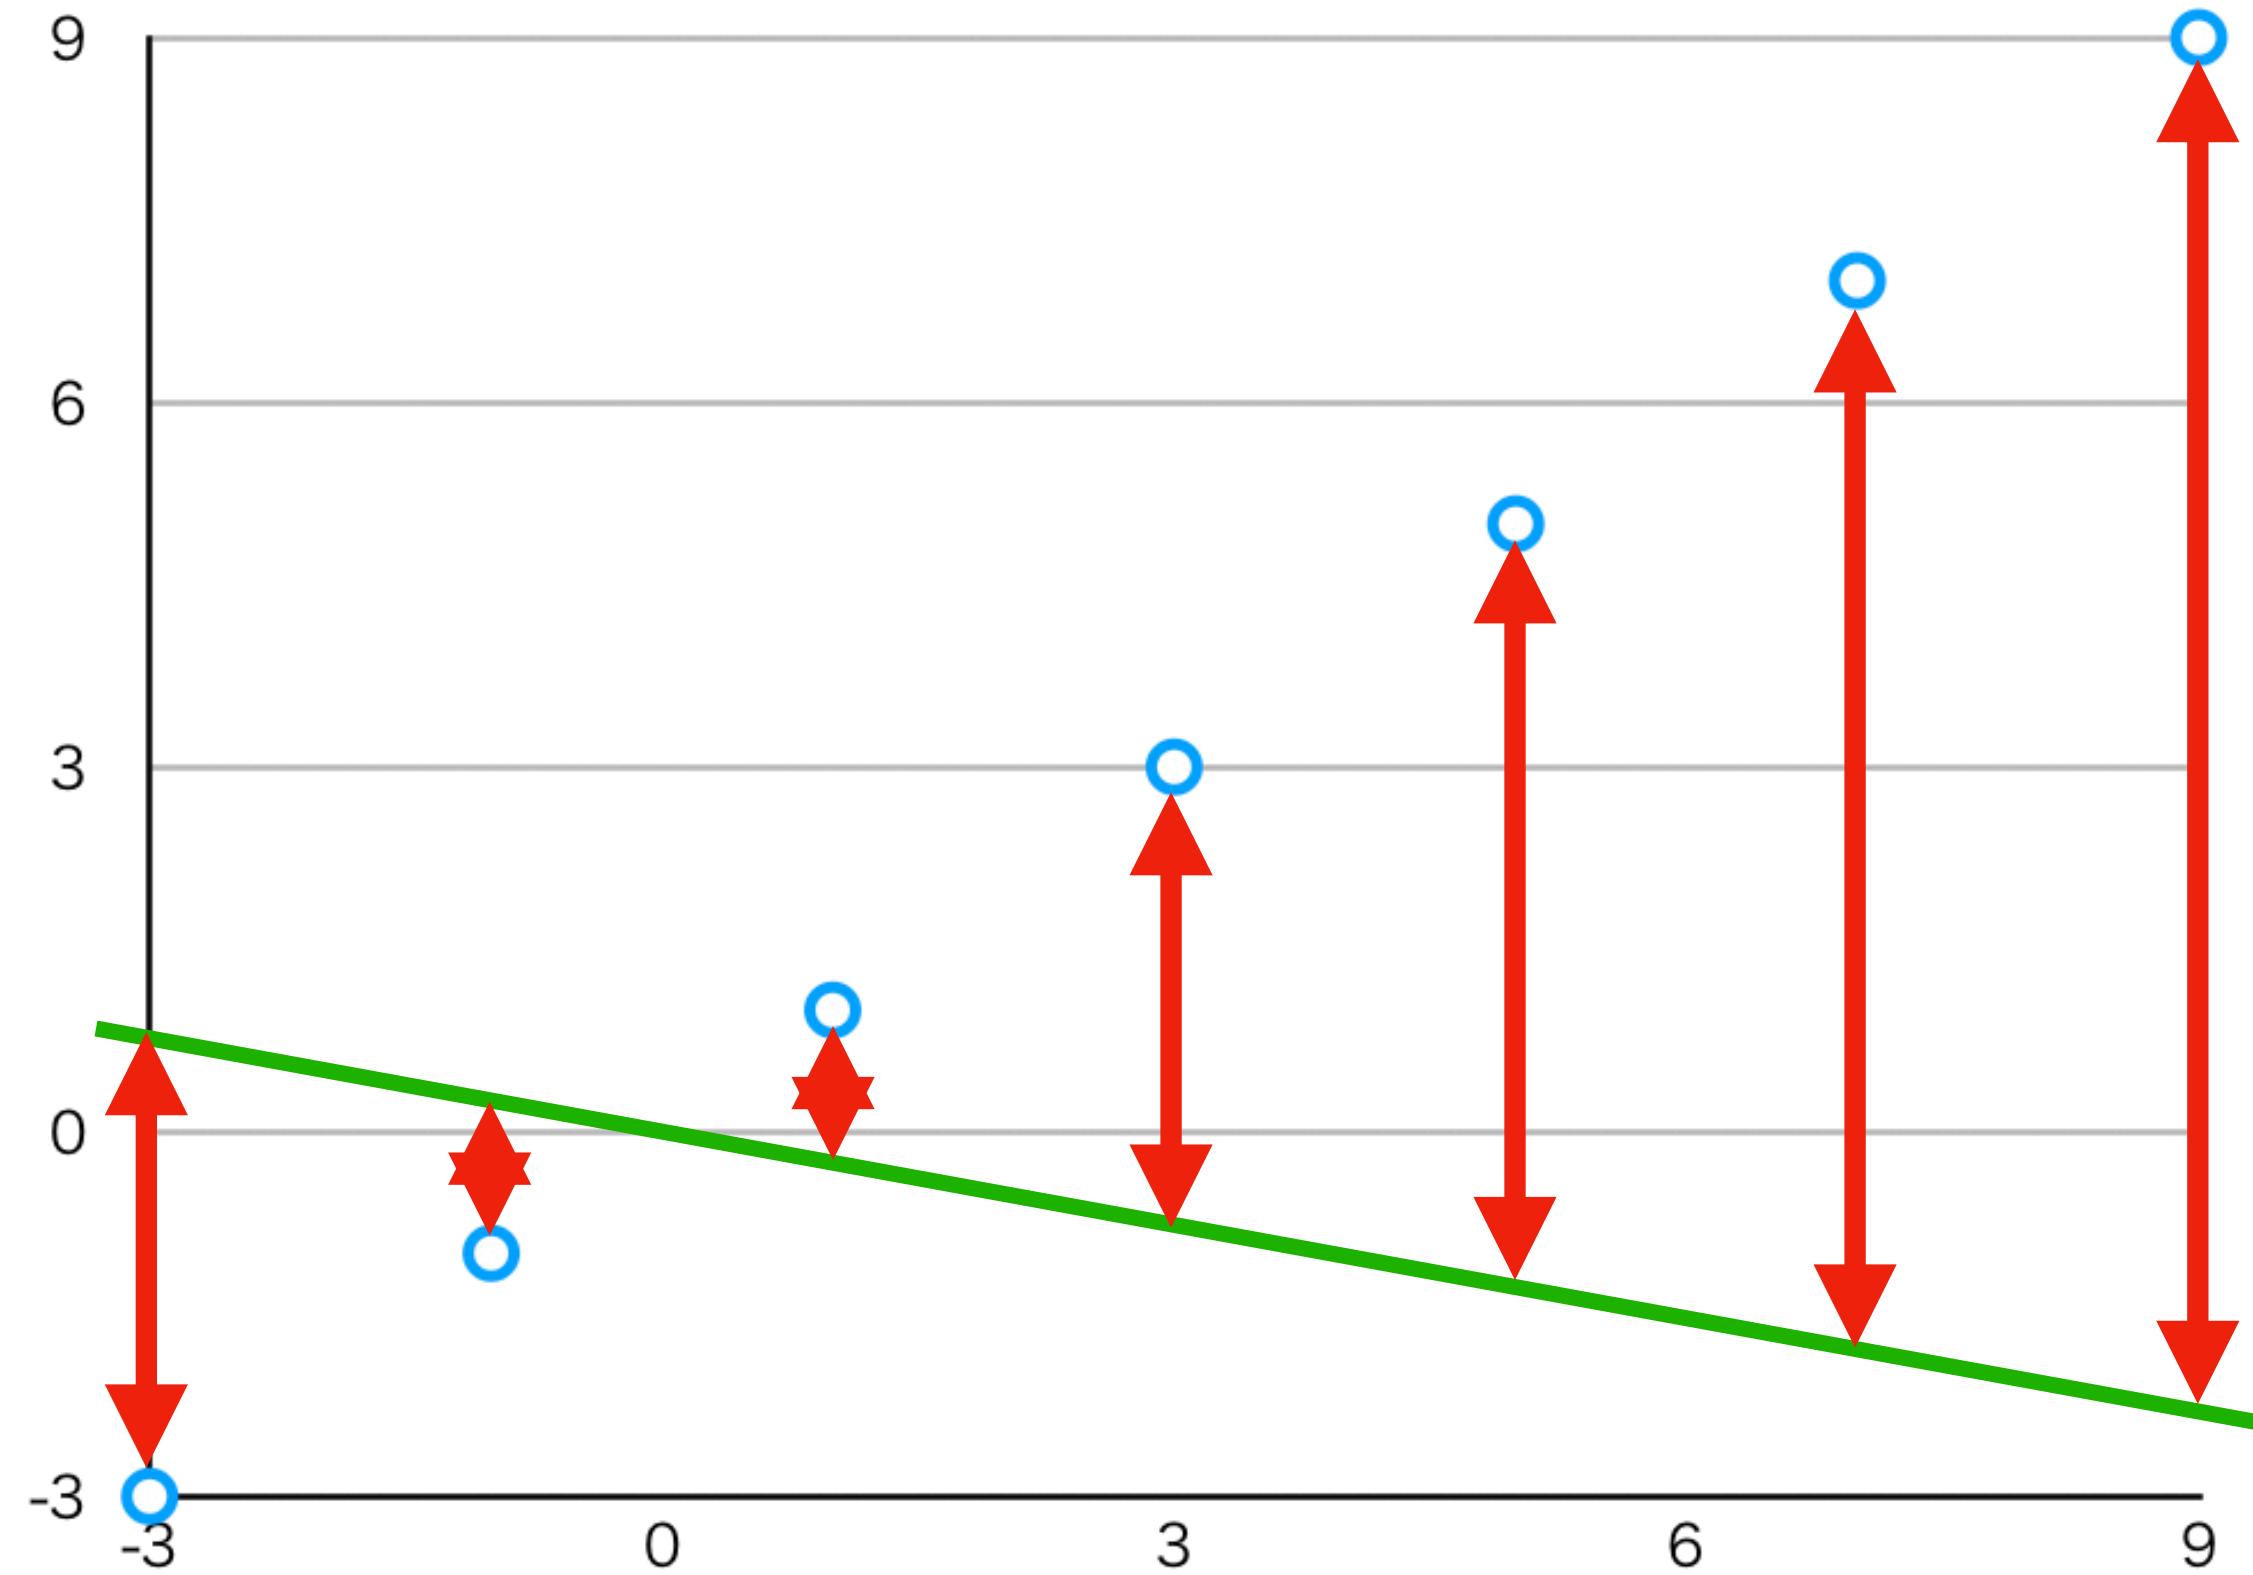
機械学習のアプローチ

- 適当にモデルを初期化 ($y = ax + b$)

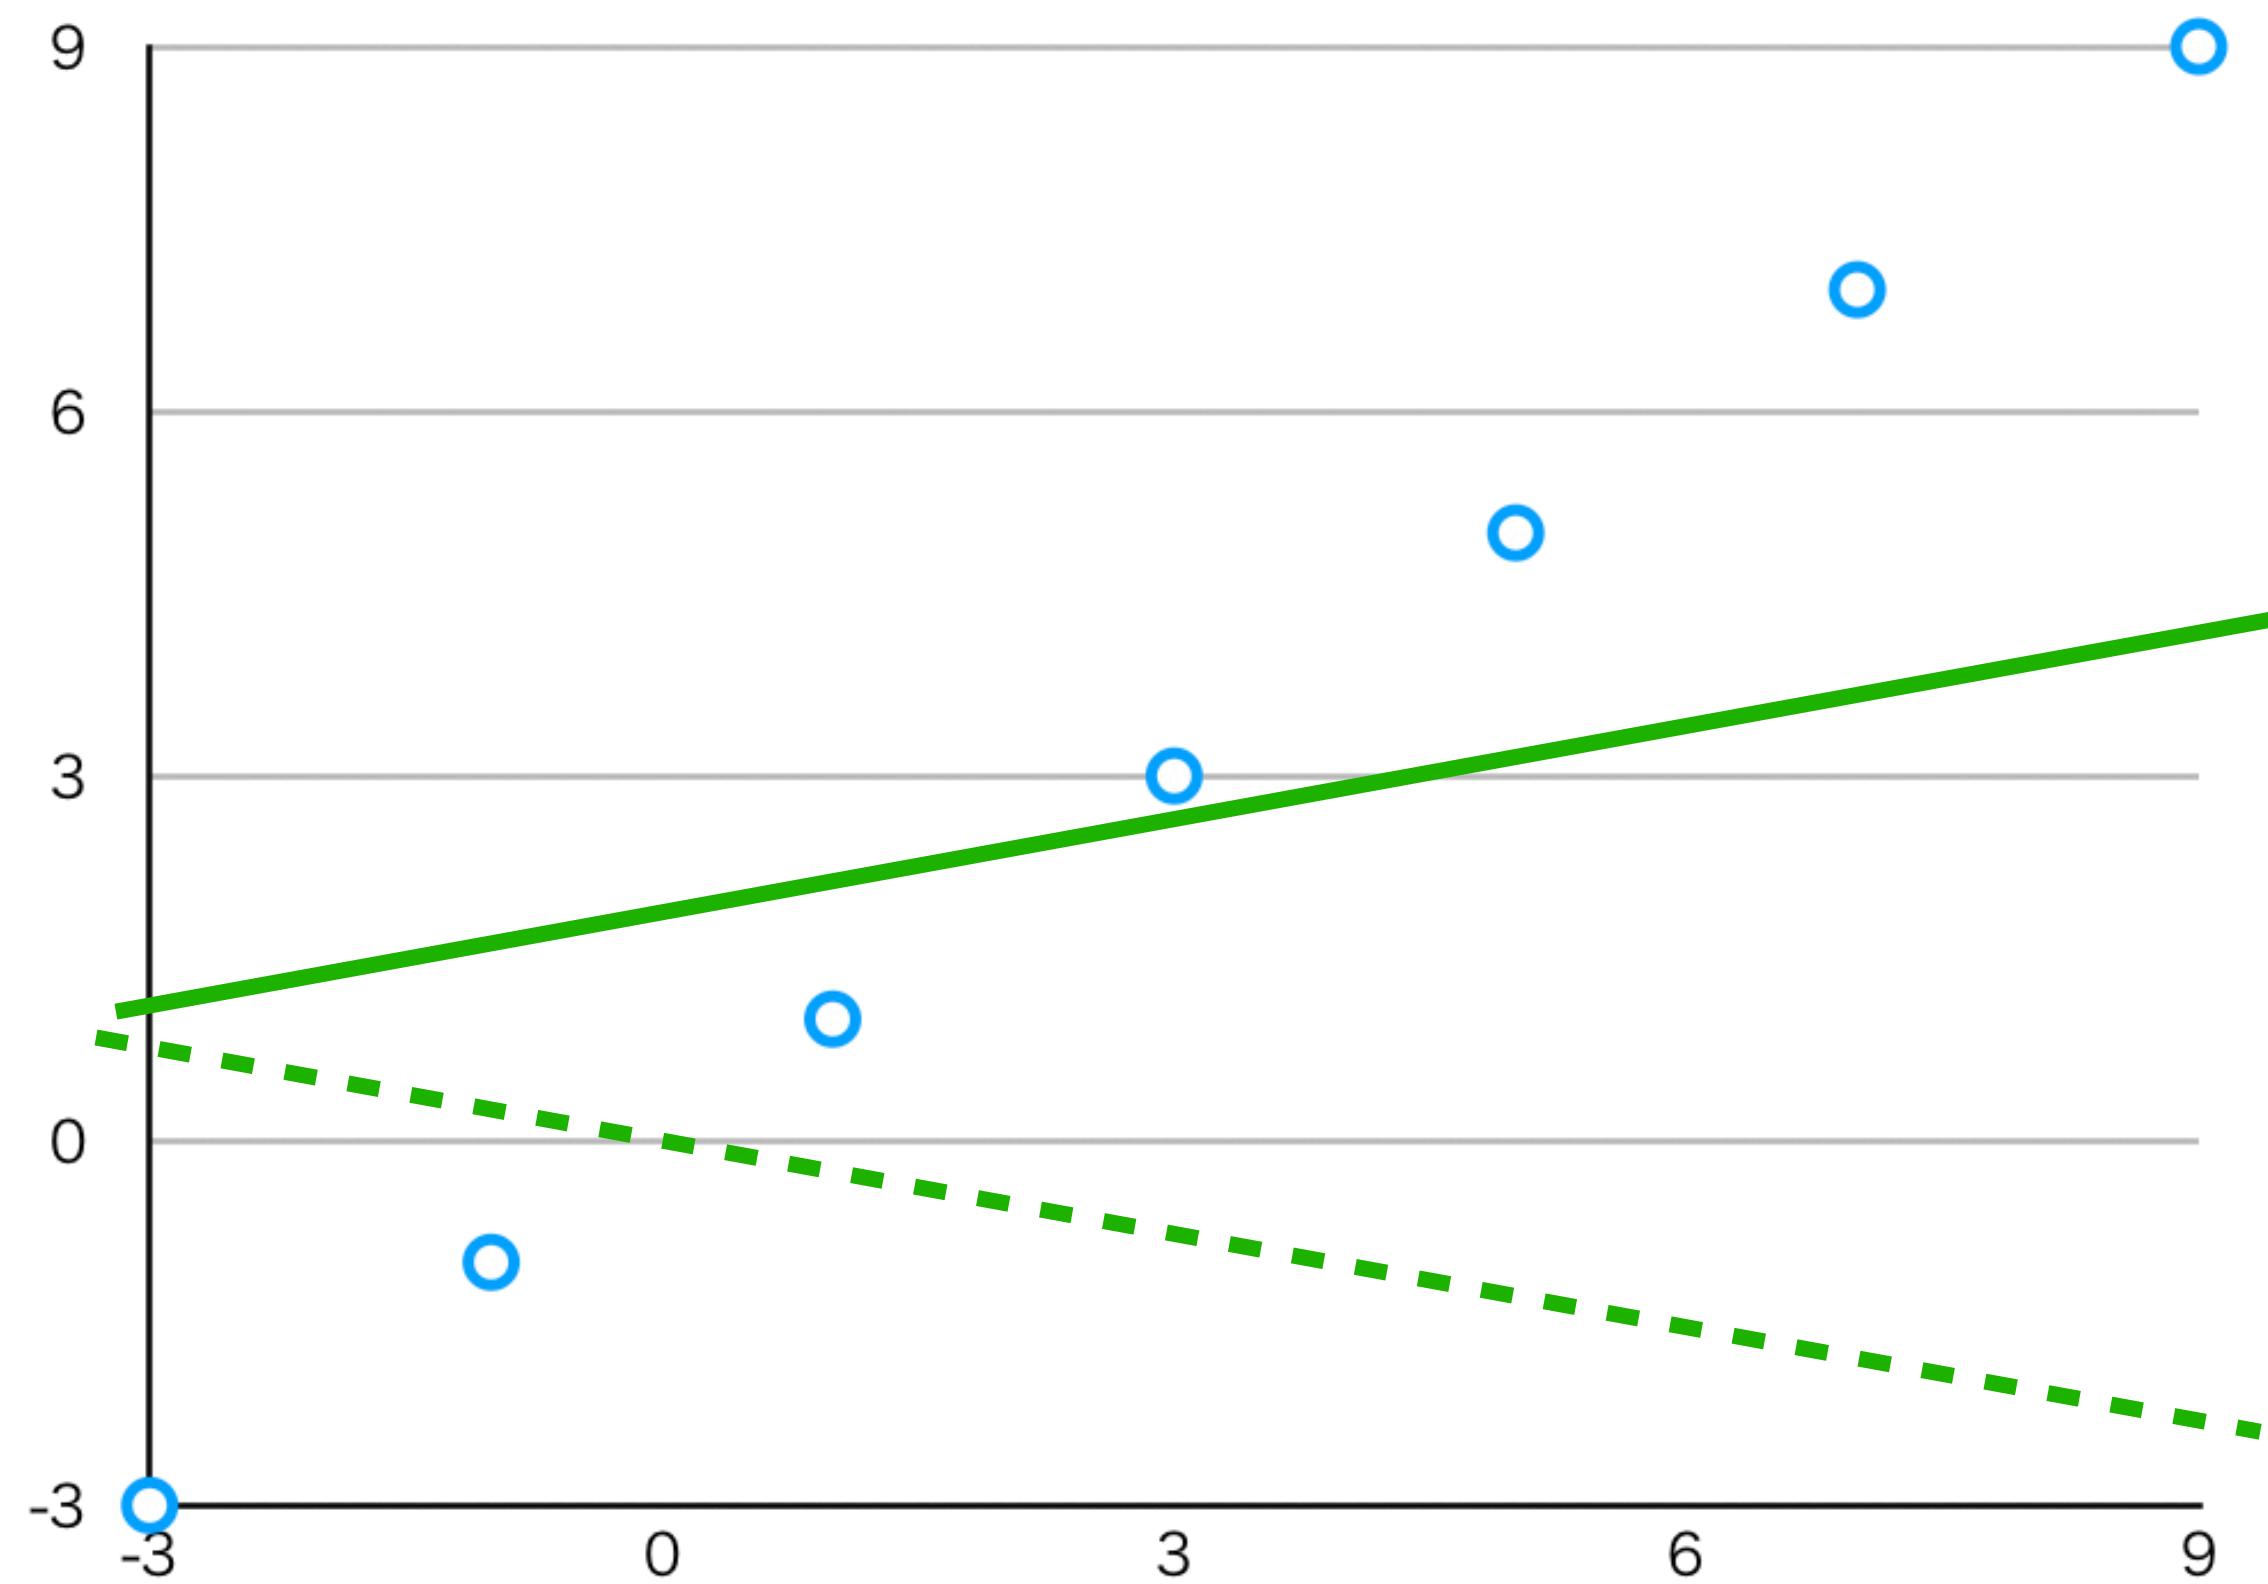


機械学習のアプローチ

- 適当にモデルを初期化 ($y_{_} = ax + b$)
- 誤差（損失）を計算 ($L = \frac{1}{N} \sum (y - y_{_})^2$)

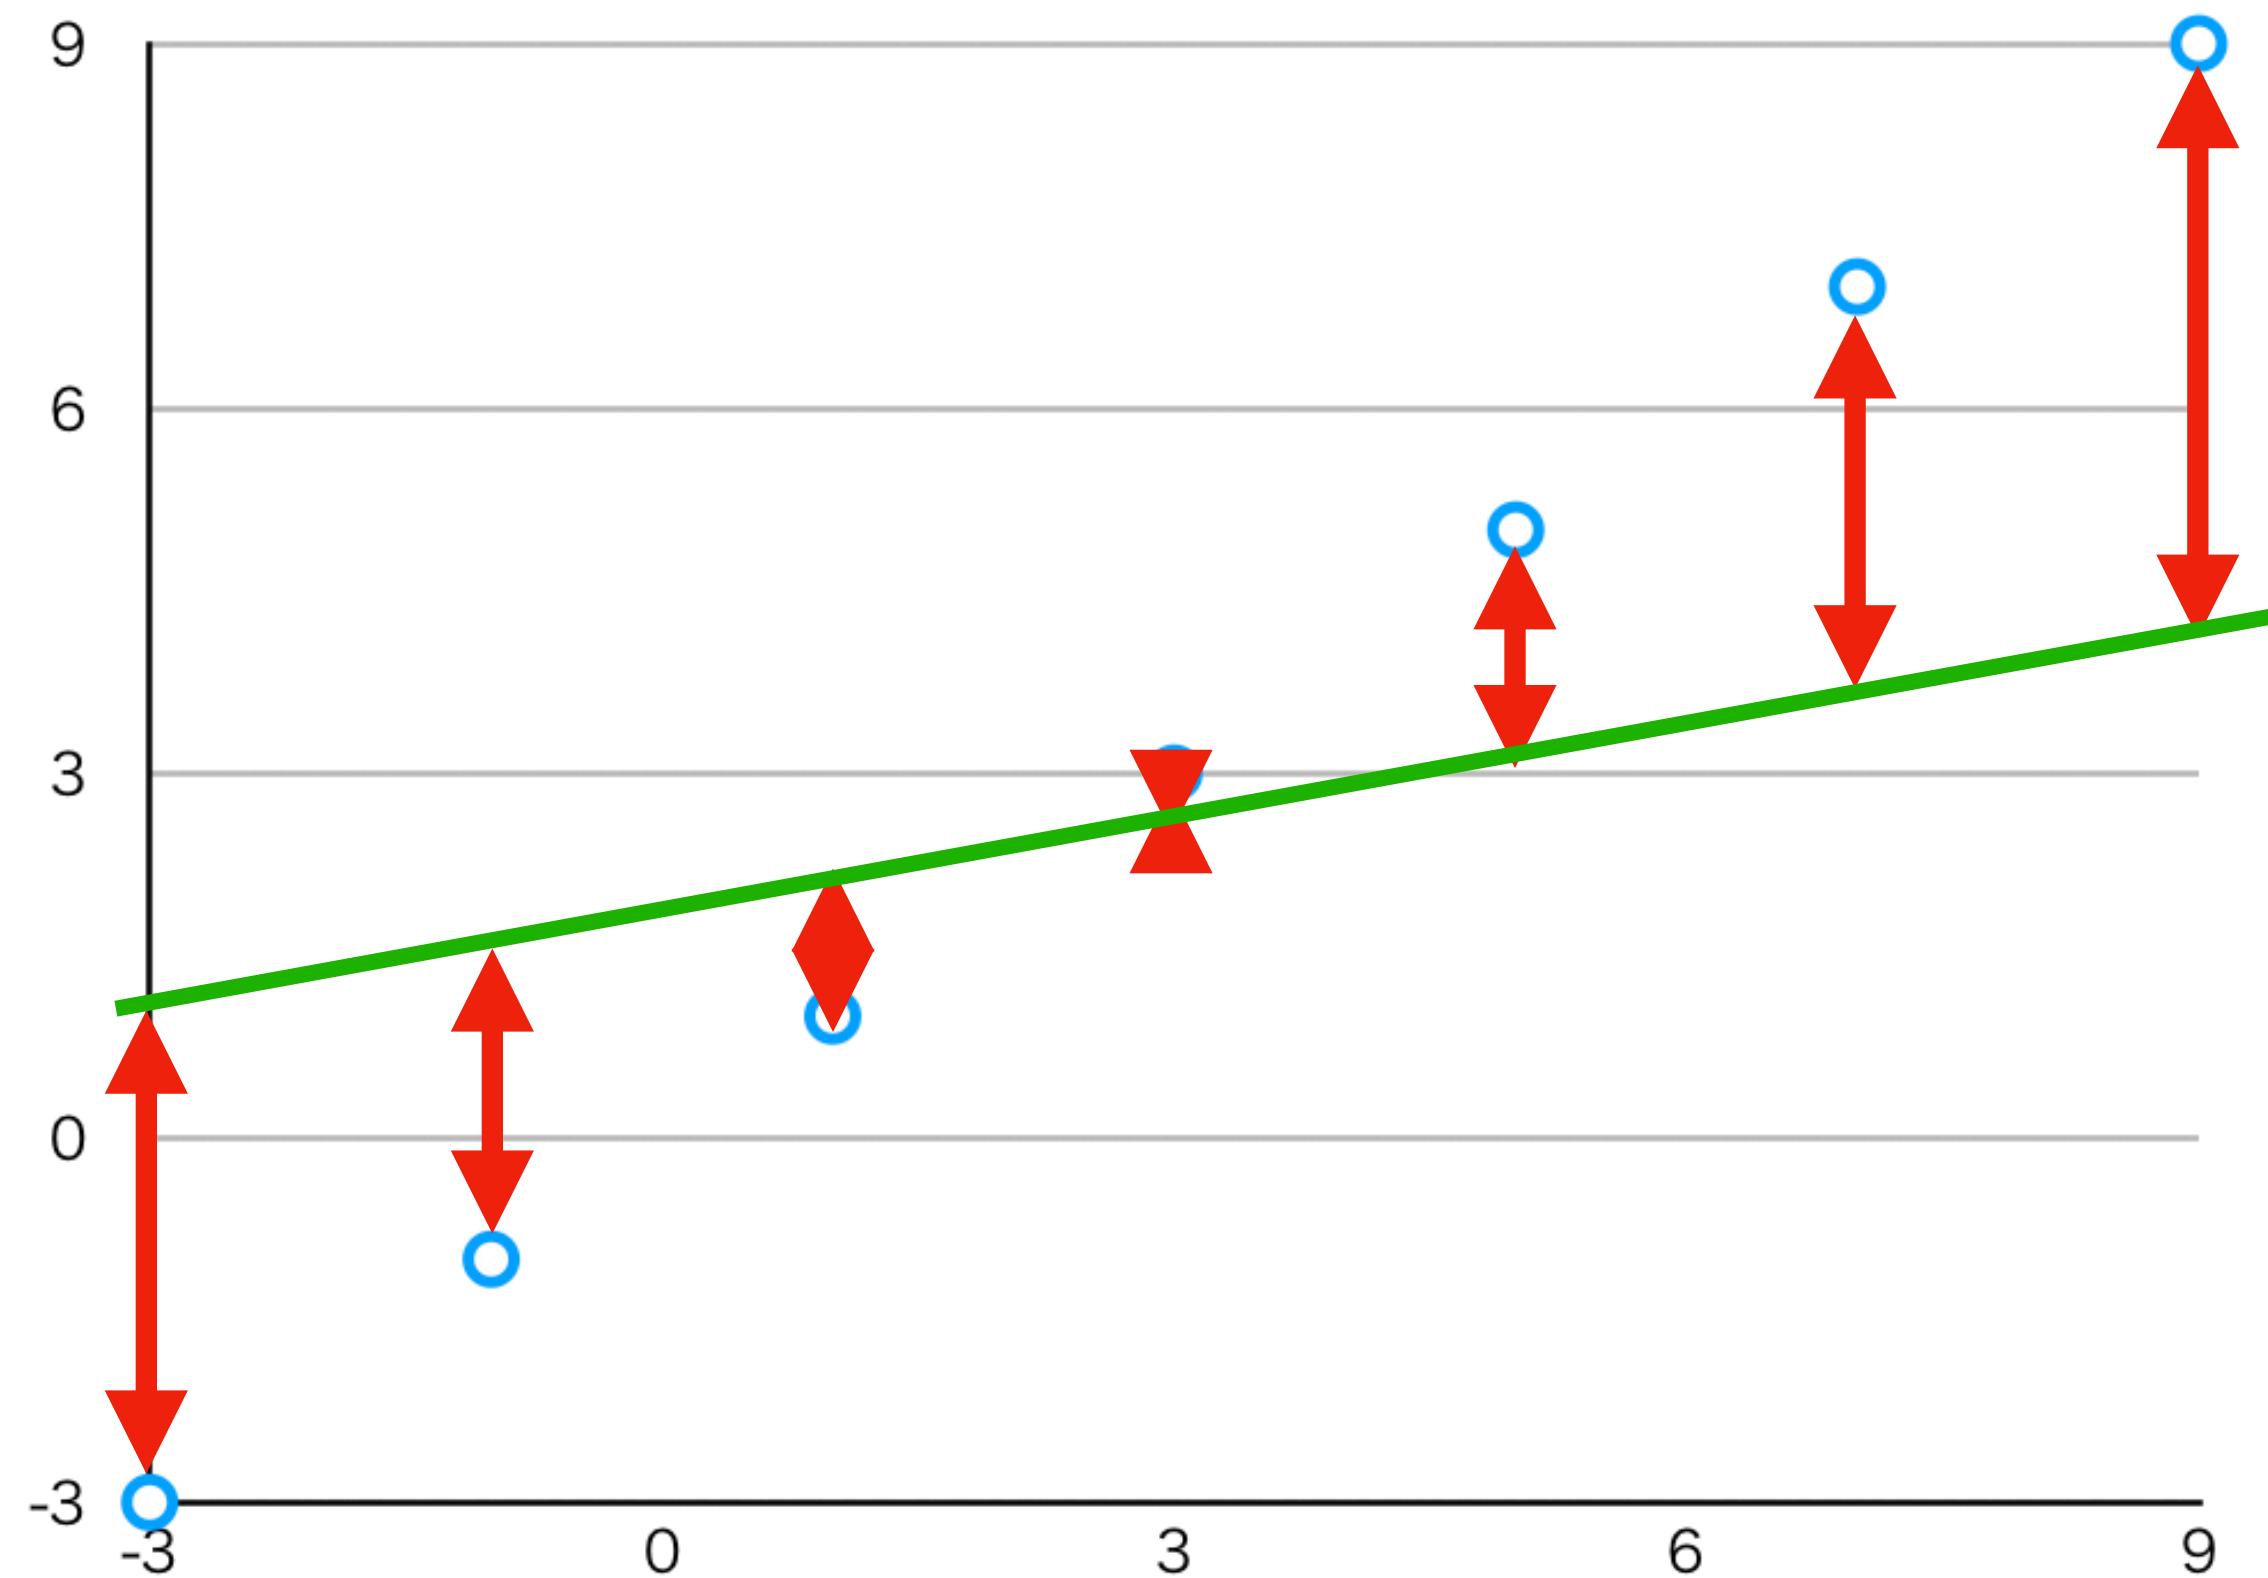


機械学習のアプローチ



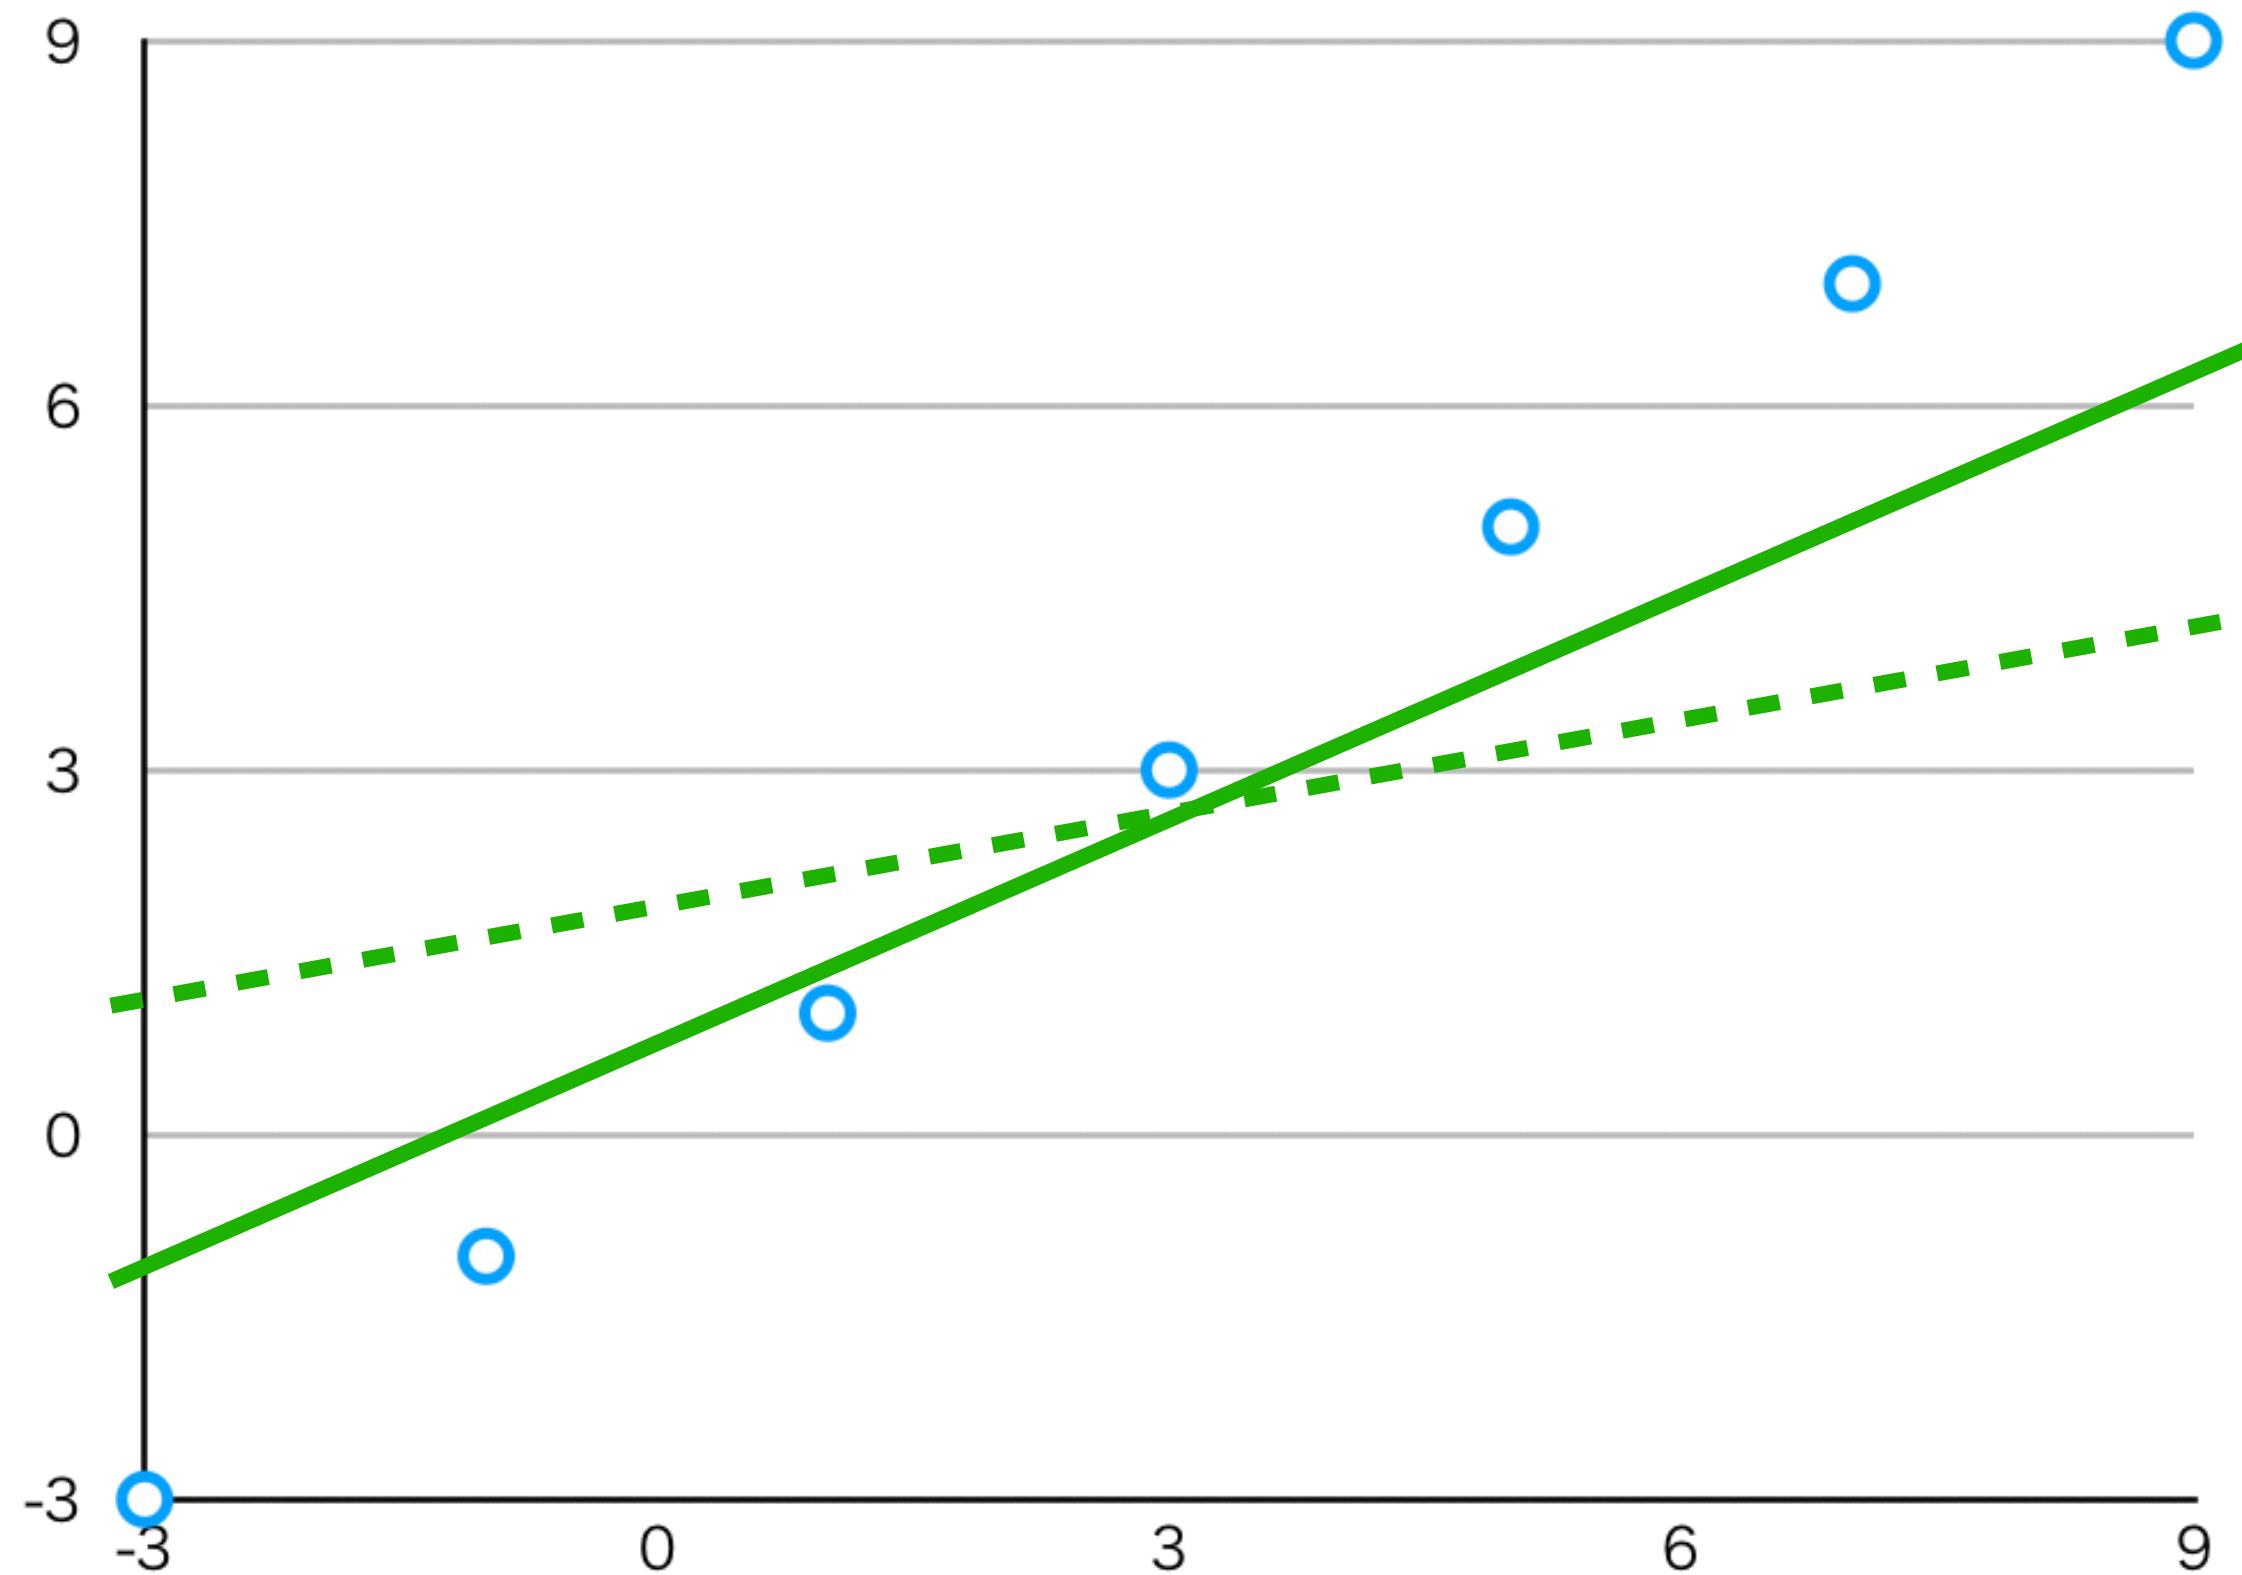
- 適当にモデルを初期化 ($y_ = ax + b$)
- 誤差（損失）を計算 ($L = 1/N \sum (y - y_)^2$)
- 誤差が小さくなるようにパラメータ (a, b) を少し更新 ($a \leftarrow a - \eta \partial L / \partial a$)

機械学習のアプローチ



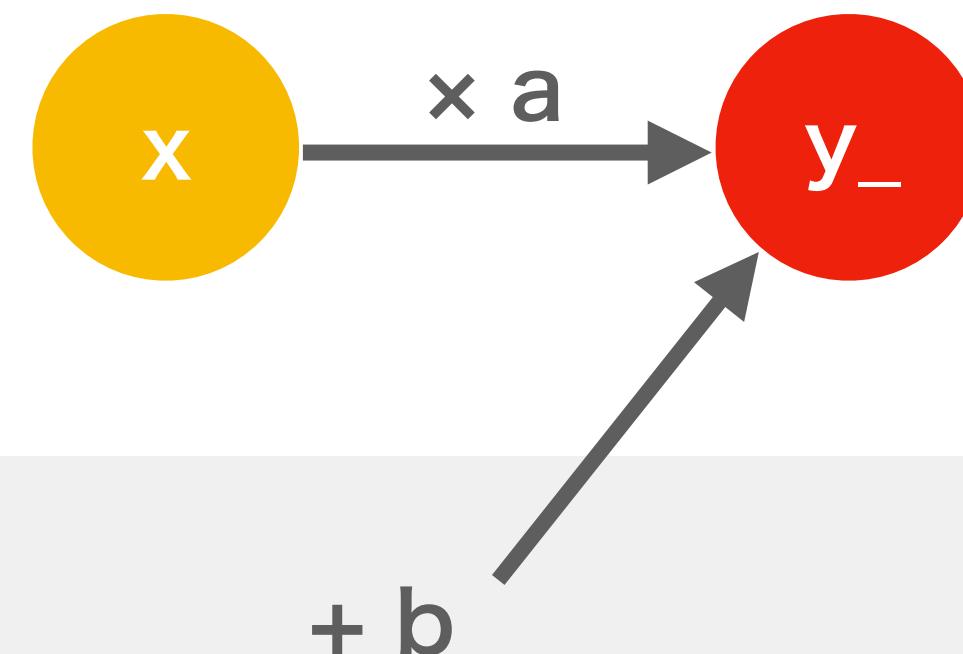
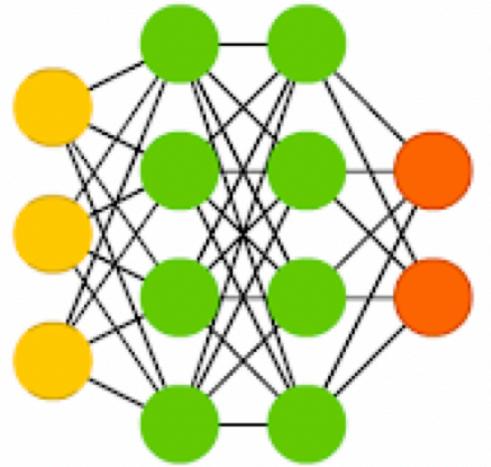
- 適当にモデルを初期化 ($y_{_} = ax + b$)
- 誤差（損失）を計算 ($L = 1/N \sum (y - y_{_})^2$)
- 誤差が小さくなるようにパラメータ (a, b) を少し更新 ($a \leftarrow a - \eta \partial L / \partial a$)
- 誤差（損失）を計算 ($L = 1/N \sum (y - y_{_})^2$)

機械学習のアプローチ



- 適当にモデルを初期化 ($y_{_} = ax + b$)
- 誤差（損失）を計算 ($L = 1/N \sum (y - y_{_})^2$)
- 誤差が小さくなるようにパラメータ (a, b) を少し更新 ($a \leftarrow a - \eta \frac{\partial L}{\partial a}$)
- 誤差（損失）を計算 ($L = 1/N \sum (y - y_{_})^2$)
- 誤差が小さくなるようにパラメータ (a, b) を少し更新 ($a \leftarrow a - \eta \frac{\partial L}{\partial a}$)
- ...

```
1 import tensorflow as tf
2 import numpy as np
3
4 model = tf.keras.Sequential([tf.keras.layers.Dense(units=1, input_shape=[1])])
5 model.compile(optimizer='sgd', loss='mean_squared_error')
6
7 xs = np.array([-1.0, 0.0, 1.0, 2.0, 3.0, 4.0], dtype=float)
8 ys = np.array([-3.0, -1.0, 1.0, 3.0, 5.0, 7.0], dtype=float)
9
10 model.fit(xs, ys, epochs=500)
11
12 >>>print(model.predict([10.0]))
13 [[18.977331]]
```



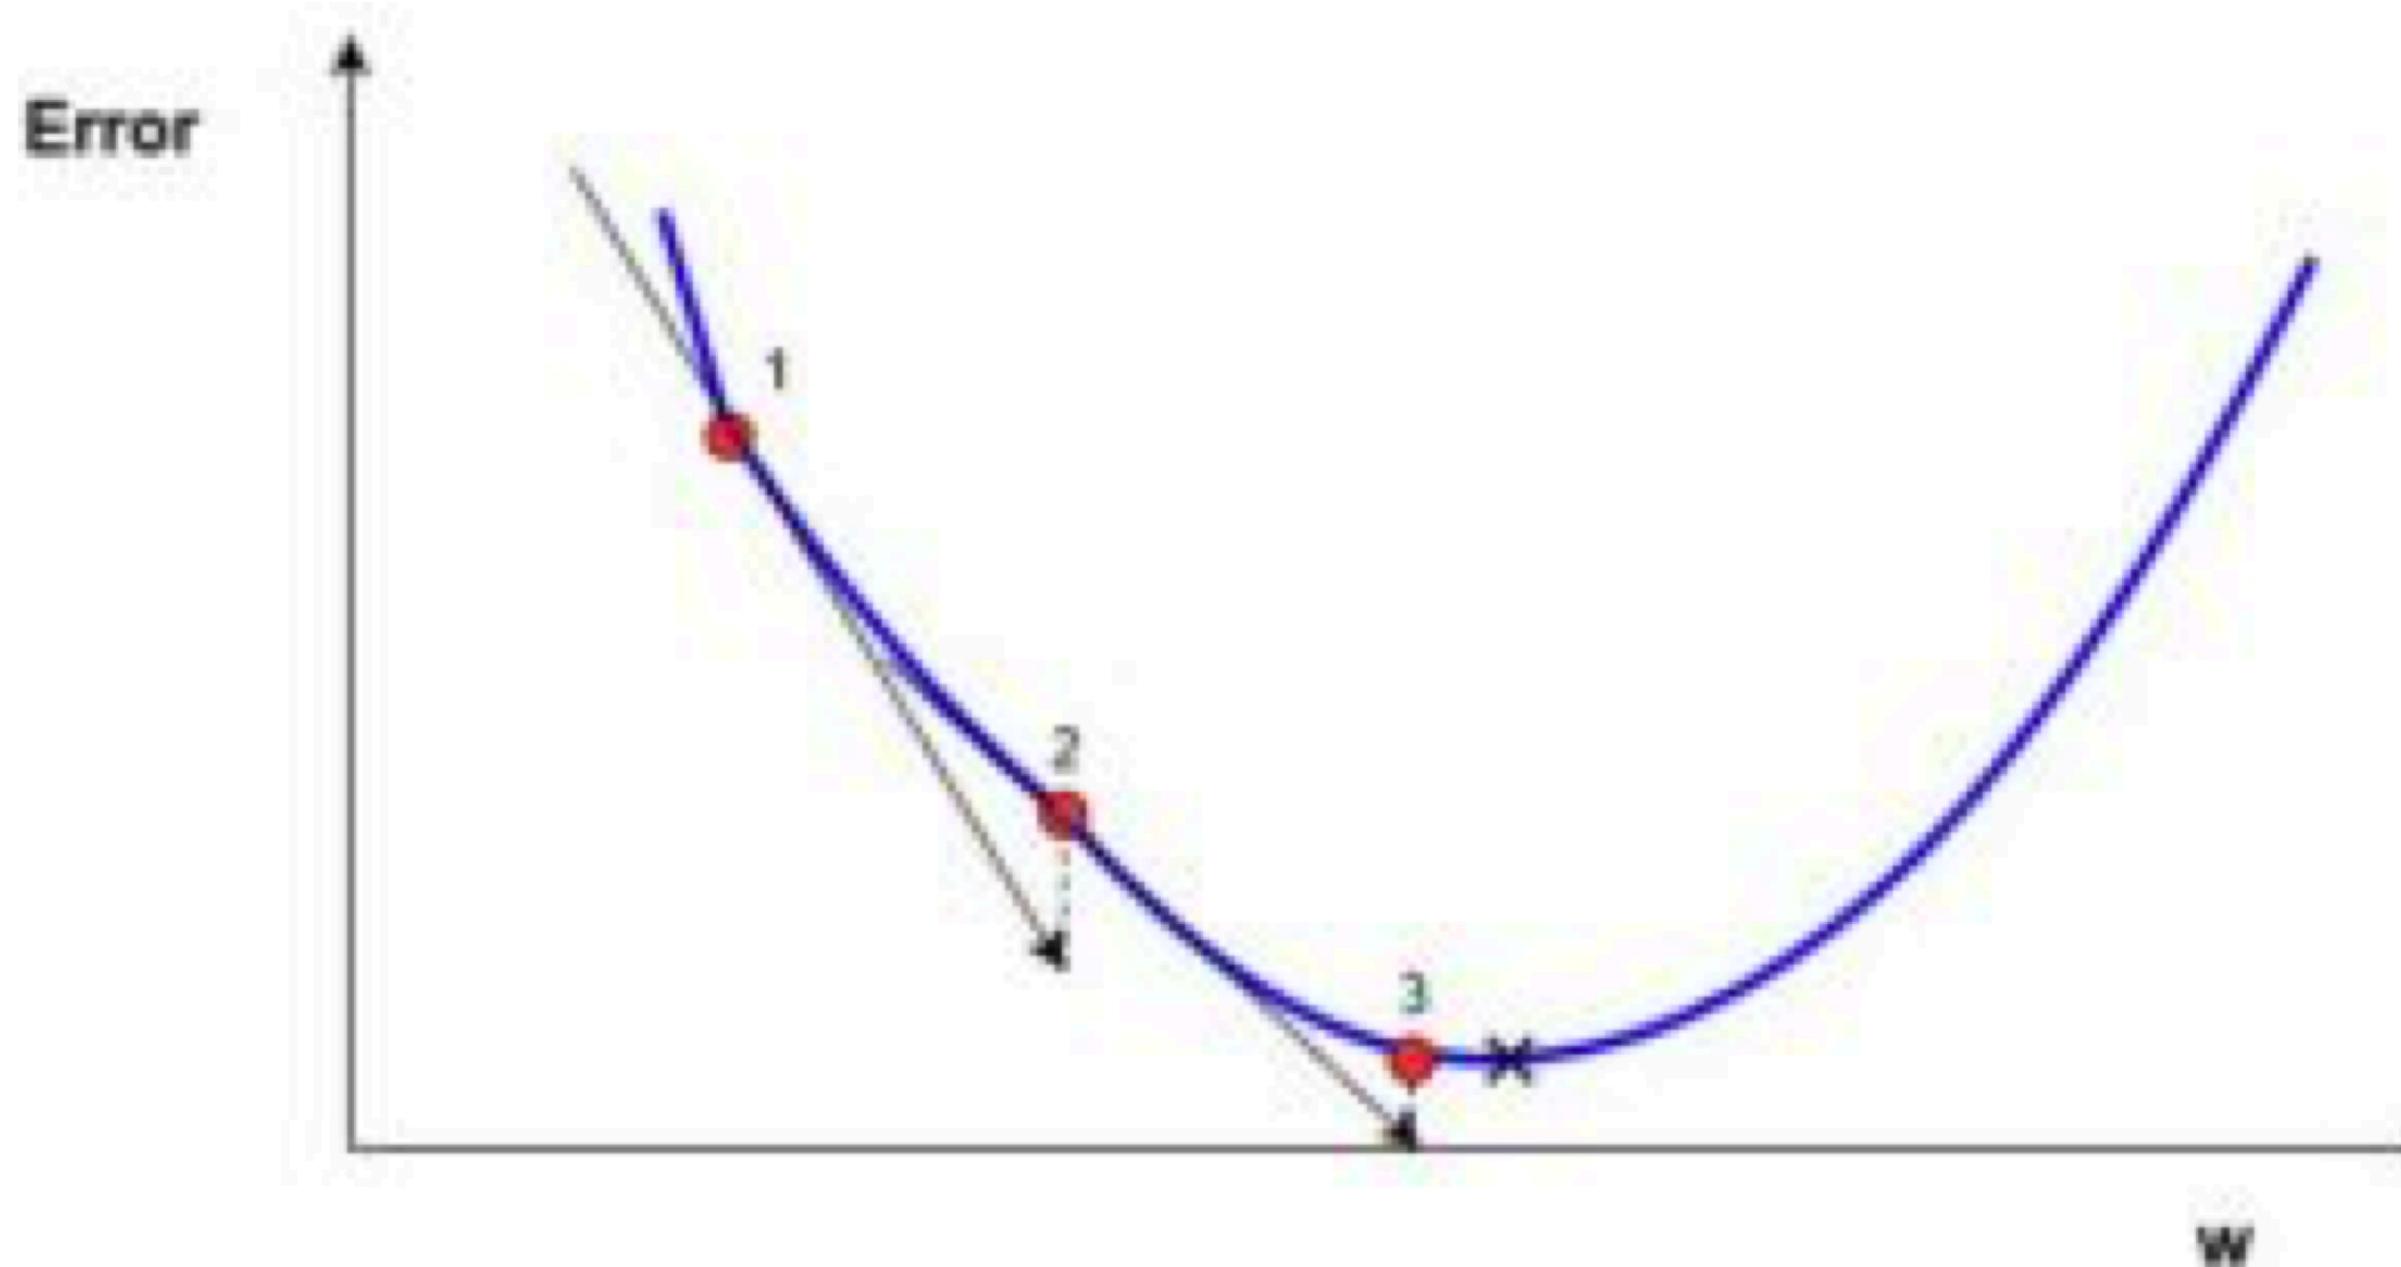
```
1 import tensorflow as tf
2 import numpy as np
3
4 model = tf.keras.Sequential([tf.keras.layers.Dense(units=1, input_shape=[1])])
5 model.compile(optimizer='sgd', loss='mean_squared_error')
6
7 xs = np.array([-1.0, 0.0, 1.0, 2.0, 3.0, 4.0], dtype=float)
8 ys = np.array([-3.0, -1.0, 1.0, 3.0, 5.0, 7.0], dtype=float)
9
10 model.fit(xs, ys, epochs=500)
11
12 >>>print(model.predict([10.0]))
13 [[18.977331]]
```

```
1 import tensorflow as tf
2 import numpy as np
3
4 model = tf.keras.Sequential([tf.keras.layers.Dense(units=1, input_shape=[1])])
5 model.compile(optimizer='sgd', loss='mean_squared_error')
6
7 xs = np.array([-1.0, 0.0, 1.0, 2.0, 3.0, 4.0], dtype=float)
8 ys = np.array([-3.0, -1.0, 1.0, 3.0, 5.0, 7.0], dtype=float)
9
10 model.fit(xs, ys, epochs=500)
11
12 >>>print(model.predict([10.0]))
13 [[18.977331]]
```

確率的勾配降下法 SGD, stochastic gradient descent.

デファクトスタンダード

(改良版 : momentum, Nesterovの加速法, RMSProp, Adam, ...)



$$w \leftarrow w - \eta \frac{\partial L}{\partial w}$$

勾配降下法



<https://www.yamakei-online.com/journal/detail.php?id=3185>

$$w \leftarrow w - \eta \frac{\partial L}{\partial w}$$

「目隠しで足元の勾配情報のみを使って山の頂上を目指すようなもの」

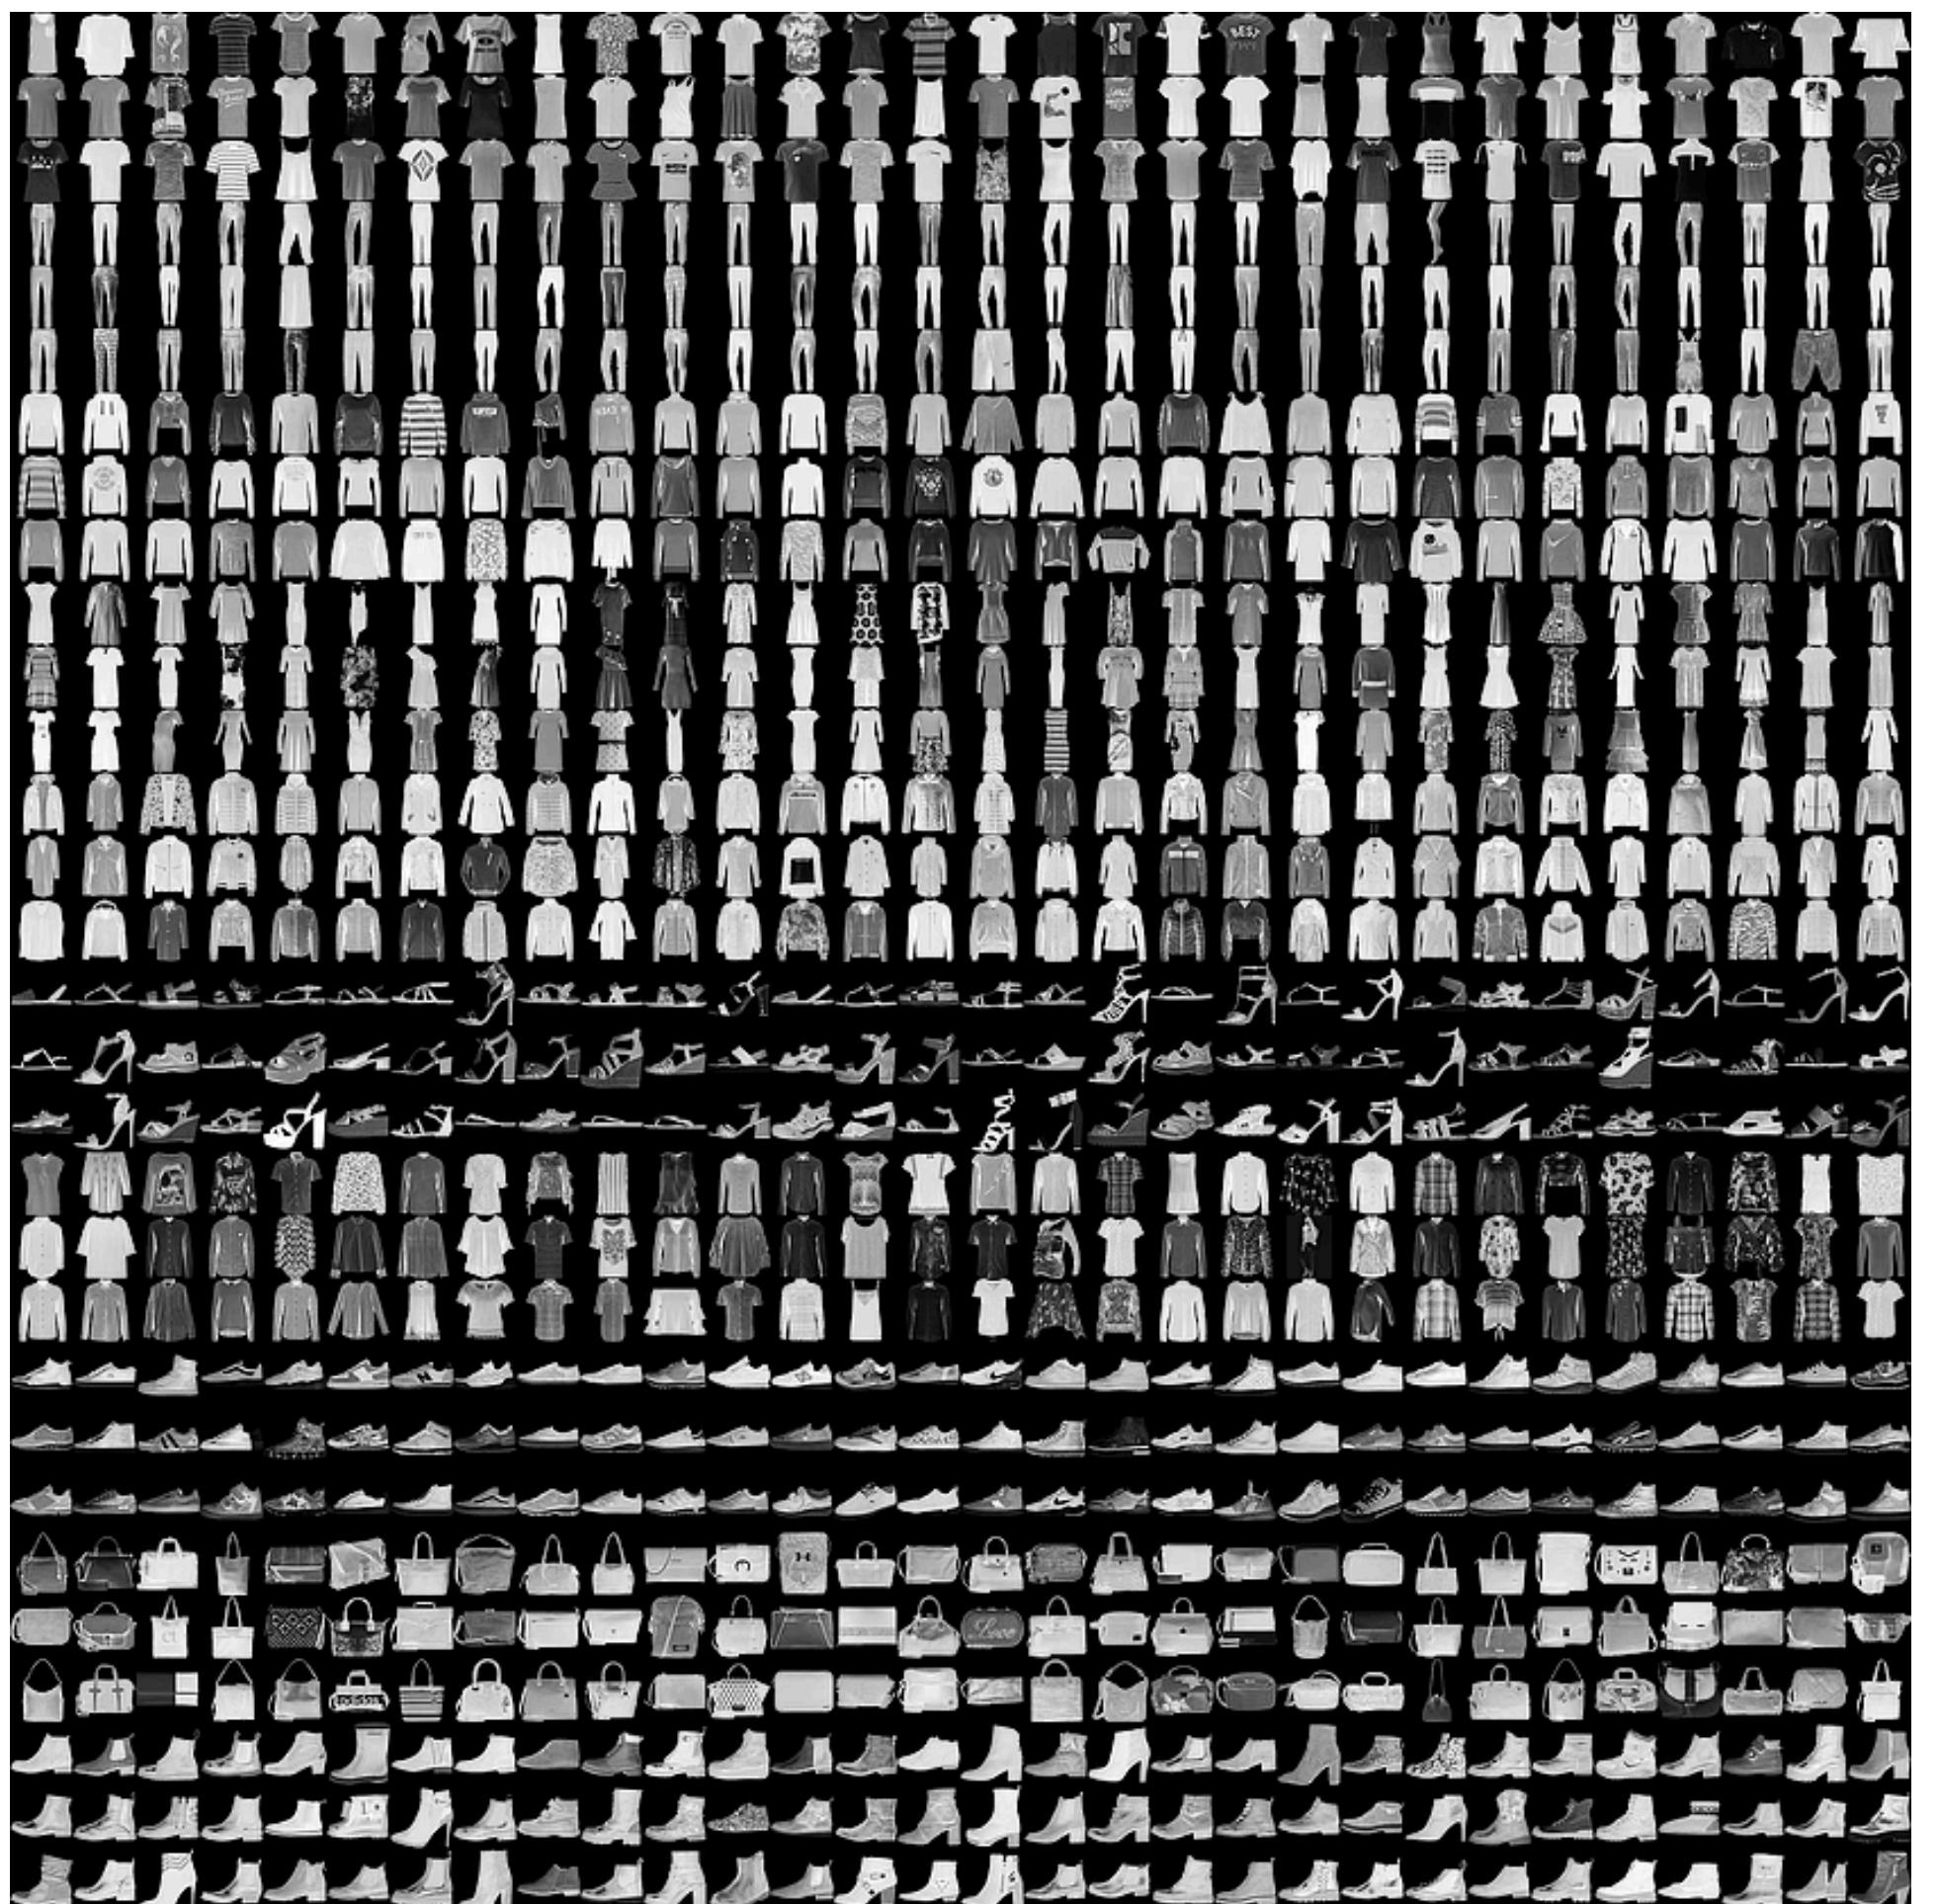
https://twitter.com/momiji_fullmoon/status/1110316960611368960

学習率 η は歩幅のイメージ (η 小=すり足、 η 大=巨人の一歩)

Fashion MNIST Dataset

- 7万画像
- 10カテゴリ
- 28×28 pixels
- 実験用データセット

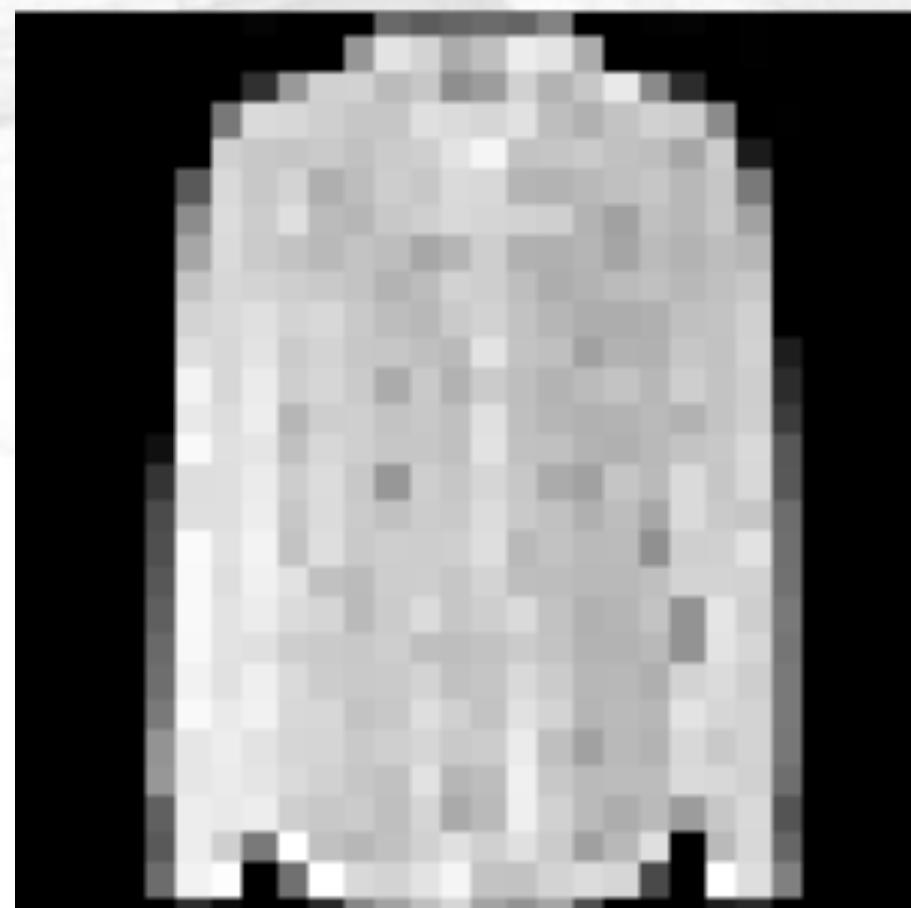
ラベル	記述
0	T-shirt/top
1	Trouser
2	Pullover
3	Dress
4	Coat
5	Sandal
6	Shirt
7	Sneaker
8	Bag
9	Ankle boot



<https://github.com/zalandoresearch/fashion-mnist>

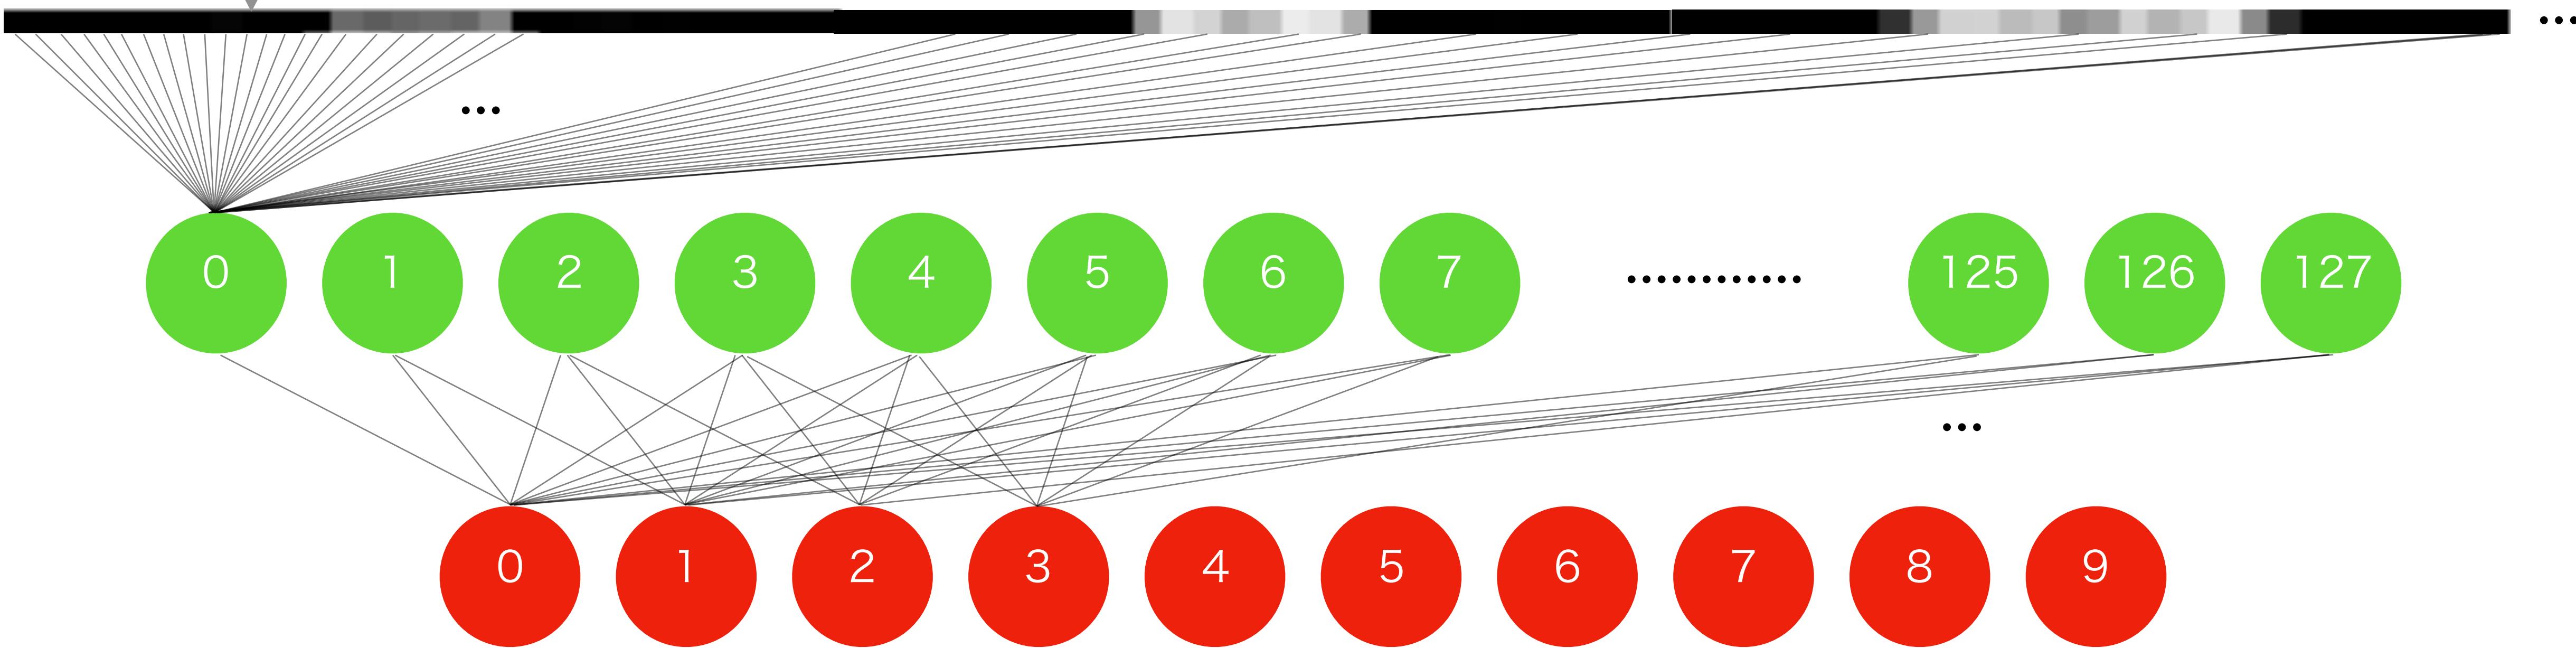
```
1 import tensorflow as tf
2 import matplotlib.pyplot as plt
3
4 fashion_mnist = tf.keras.datasets.fashion_mnist
5 (train_images, train_labels), (test_images, test_labels) = fashion_mnist.load_data()
6 train_images = train_images / 255.
7 test_images = test_images / 255.
```

```
1 import tensorflow as tf
2 import matplotlib.pyplot as plt
3
4 fashion_mnist = tf.keras.datasets.fashion_mnist
5 (train_images, train_labels), (test_images, test_labels) = fashion_mnist.load_data()
6 train_images = train_images / 255.
7 test_images = test_images / 255.
8
9 model = tf.keras.Sequential([
10     tf.keras.layers.Flatten(input_shape=(28, 28)),
11     tf.keras.layers.Dense(128, activation=tf.nn.relu),
12     tf.keras.layers.Dense(10, activation=tf.nn.softmax)
13 ])
```



```
9 model = tf.keras.Sequential([
10     tf.keras.layers.Flatten(input_shape=(28, 28)),
11     tf.keras.layers.Dense(128, activation=tf.nn.relu),
12     tf.keras.layers.Dense(10, activation=tf.nn.softmax)
13 ])
```

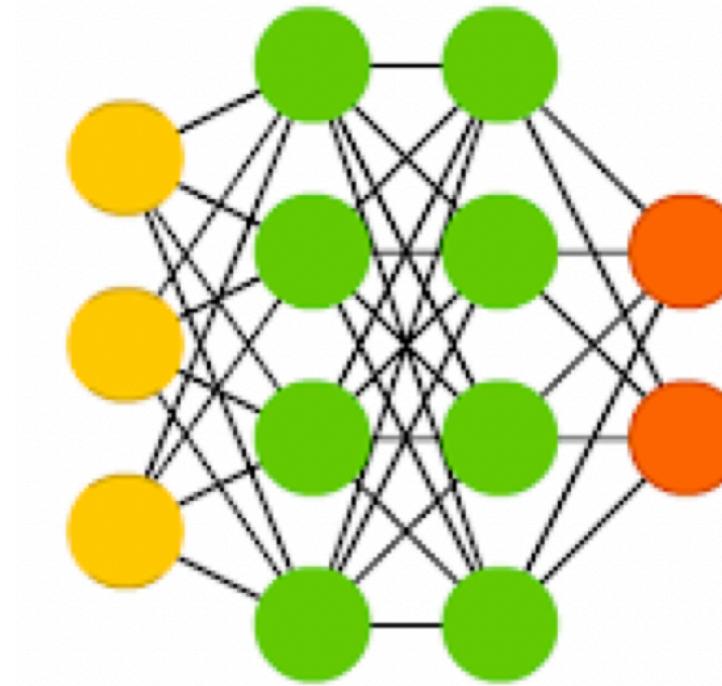
Flatten: (28, 28) => (784)



```
1 import tensorflow as tf
2 import matplotlib.pyplot as plt
3
4 fashion_mnist = tf.keras.datasets.fashion_mnist
5 (train_images, train_labels), (test_images, test_labels) = fashion_mnist.load_data()
6 train_images = train_images / 255.
7 test_images = test_images / 255.
8
9 model = tf.keras.Sequential([
10     tf.keras.layers.Flatten(input_shape=(28, 28)),
11     tf.keras.layers.Dense(128, activation=tf.nn.relu),
12     tf.keras.layers.Dense(10, activation=tf.nn.softmax)
13 ])
14
15 model.compile(optimizer=tf.train.AdamOptimizer(),
16                 loss='sparse_categorical_crossentropy', metrics=['accuracy'])
16 model.fit(train_images, train_labels, epochs=5, batch_size=128, verbose=1,
17             validation_data=(test_images, test_labels))
```

Dense Layer の欠点

- 入力のベクトルの**全要素**の相関をみている
 - > 住宅価格予測みたいな話ならまだいい
 - > もう作ってる特徴量と特徴量の組み合わせ
 - 例) 東京墨田区 & 床面積 30m² & 1K & 風呂トイレ別 & 新築 => 家賃月10万円
- 「**画像の特徴量**」を抽出してからDense Layerに渡せば効率的



畳込みニューラルネットワーク (Convolutional Neural Network; CNN)



畳み込みニューラルネットワークの基礎

畳み込み (Convolution)



144	60	19
188	82	32
156	55	27

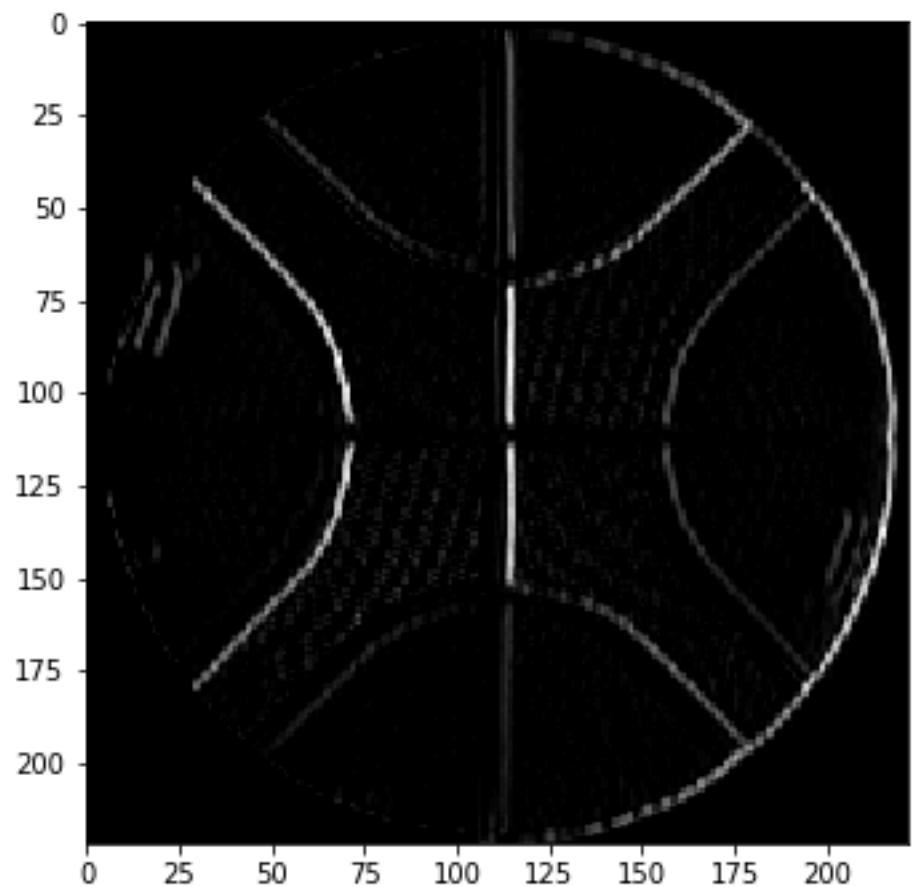
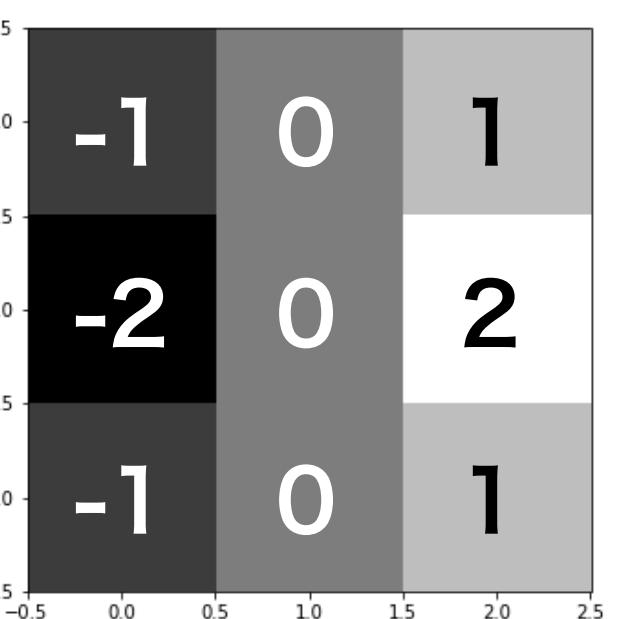
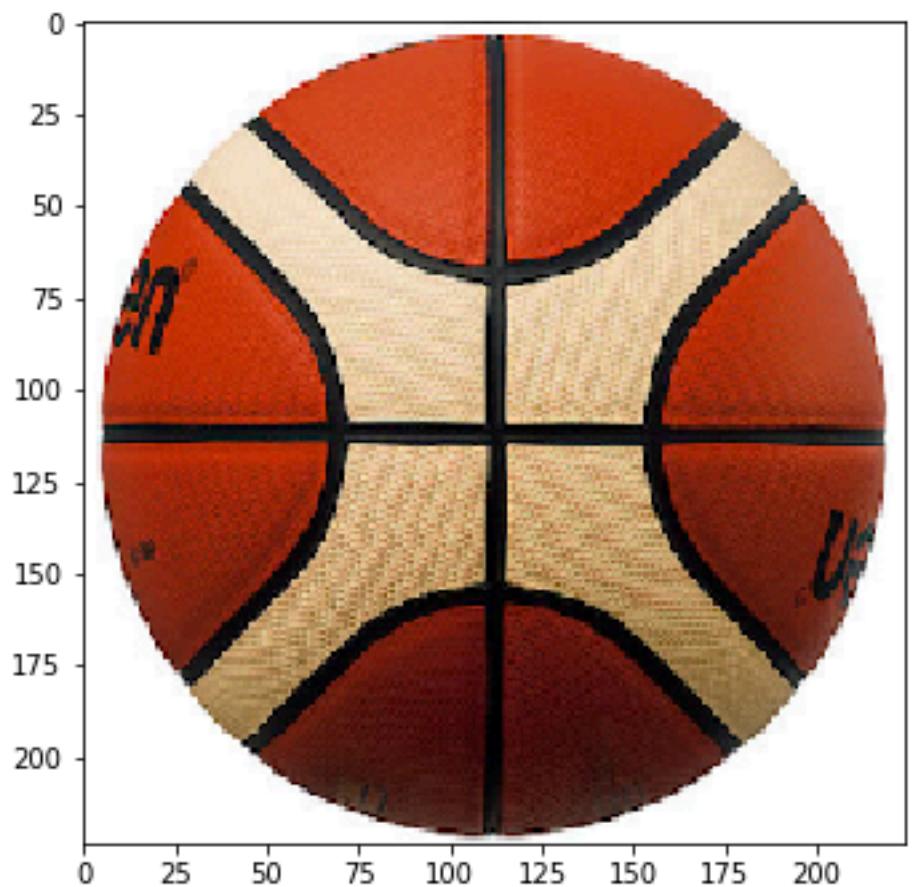
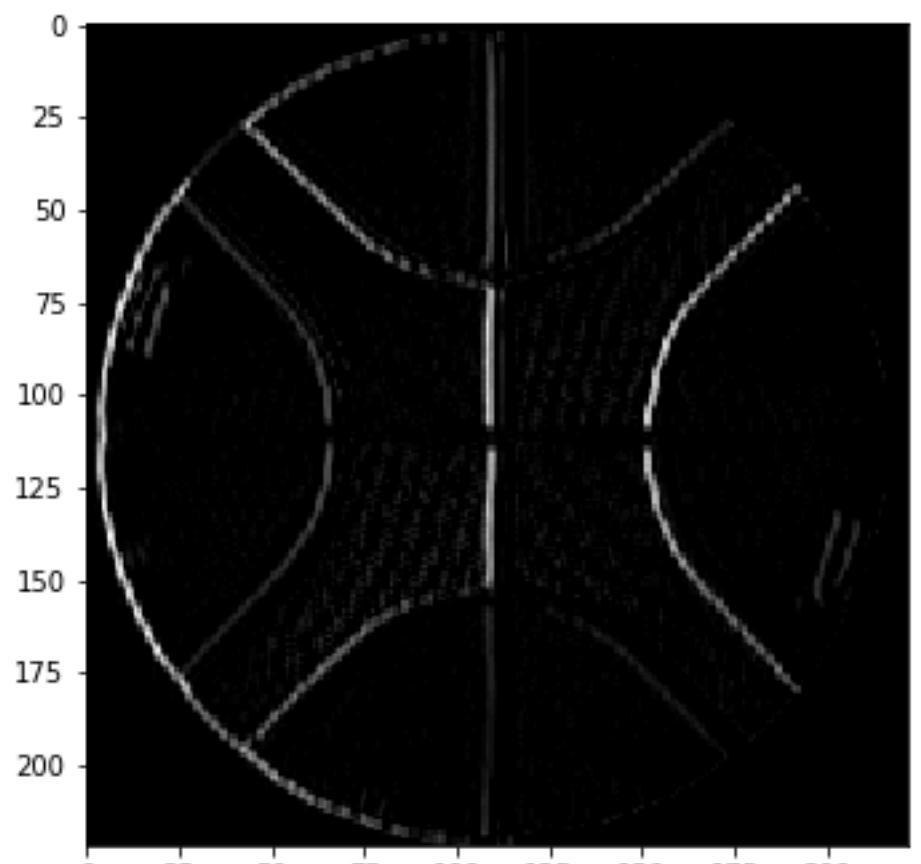
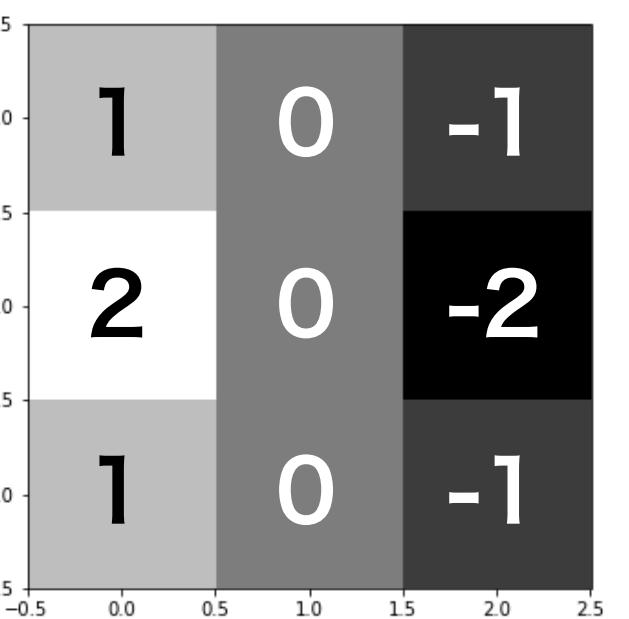
-1	0	-2
0.5	4.5	-1.5
1.5	2	-3

CURRENT_PIXEL_VALUE = 82

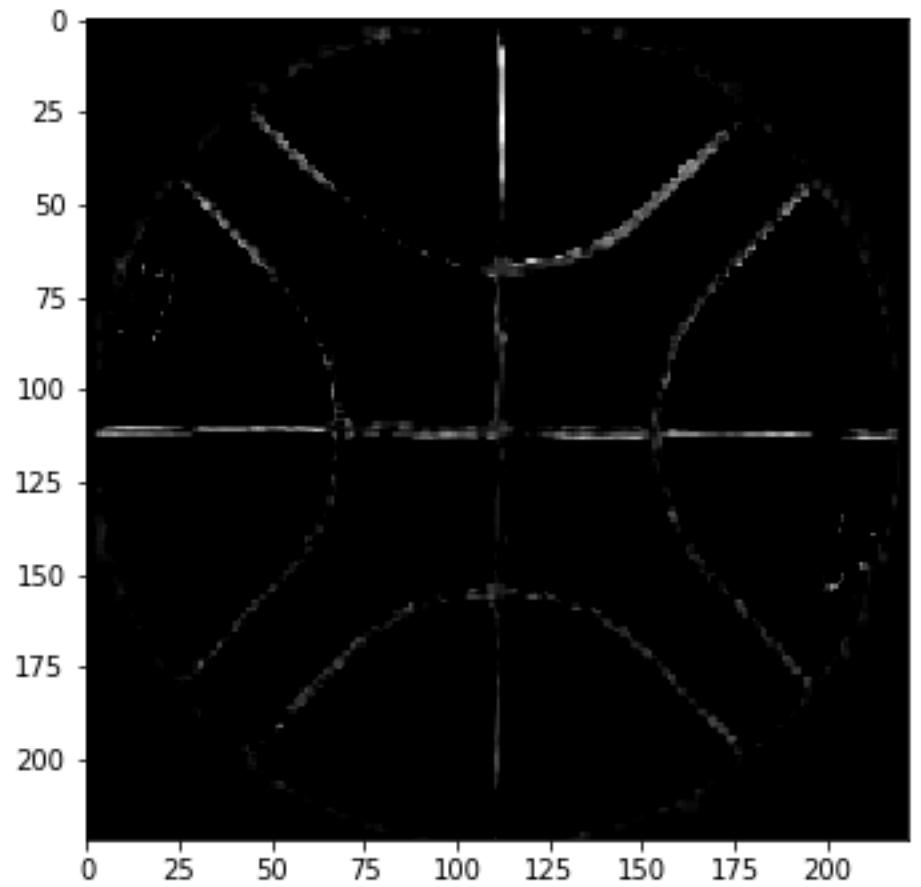
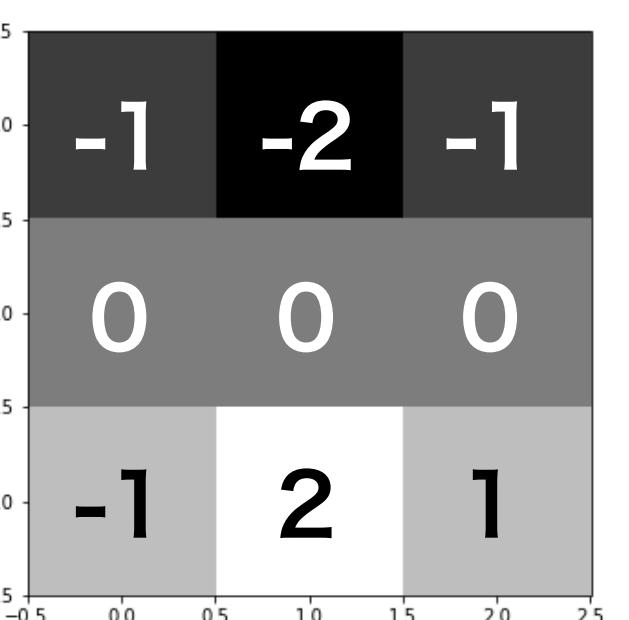
NEW_PIXEL_VALUE =

$$\begin{aligned} u_{ijm} = & \sum_{k=0}^{K-1} \sum_{p=0}^{W-1} \sum_{q=0}^{H-1} z_{i+p, j+q, k}^{(l-1)} h_{pqkm} + b_{ijm} \\ & (-1 * 144) + (0 * 60) + (-2 * 19) \\ & + (0.5 * 188) + (4.5 * 82) + (-1.5 * 32) \\ & + (1.5 * 156) + (2 * 55) + (-3 * 27) \end{aligned}$$

$$u_{ijm} = \sum_{k=0}^{K-1} \sum_{p=0}^{W-1} \sum_{q=0}^{H-1} z_{i+p, j+q, k}^{(l-1)} h_{pqkm} + b_{ijm}$$



https://www.bbkong.net/fs/alleyoop/molten_BGL7



stride=(1, 1), padding='valid'



stride=(1, 1), padding='valid'



stride=(1, 1), padding='valid'



stride=(1, 1), padding='valid'



stride=(1, 1), padding='valid'



`stride=(1, 1), padding='valid'`



stride=(1, 1), padding='valid'



stride=(2, 2), padding='valid'



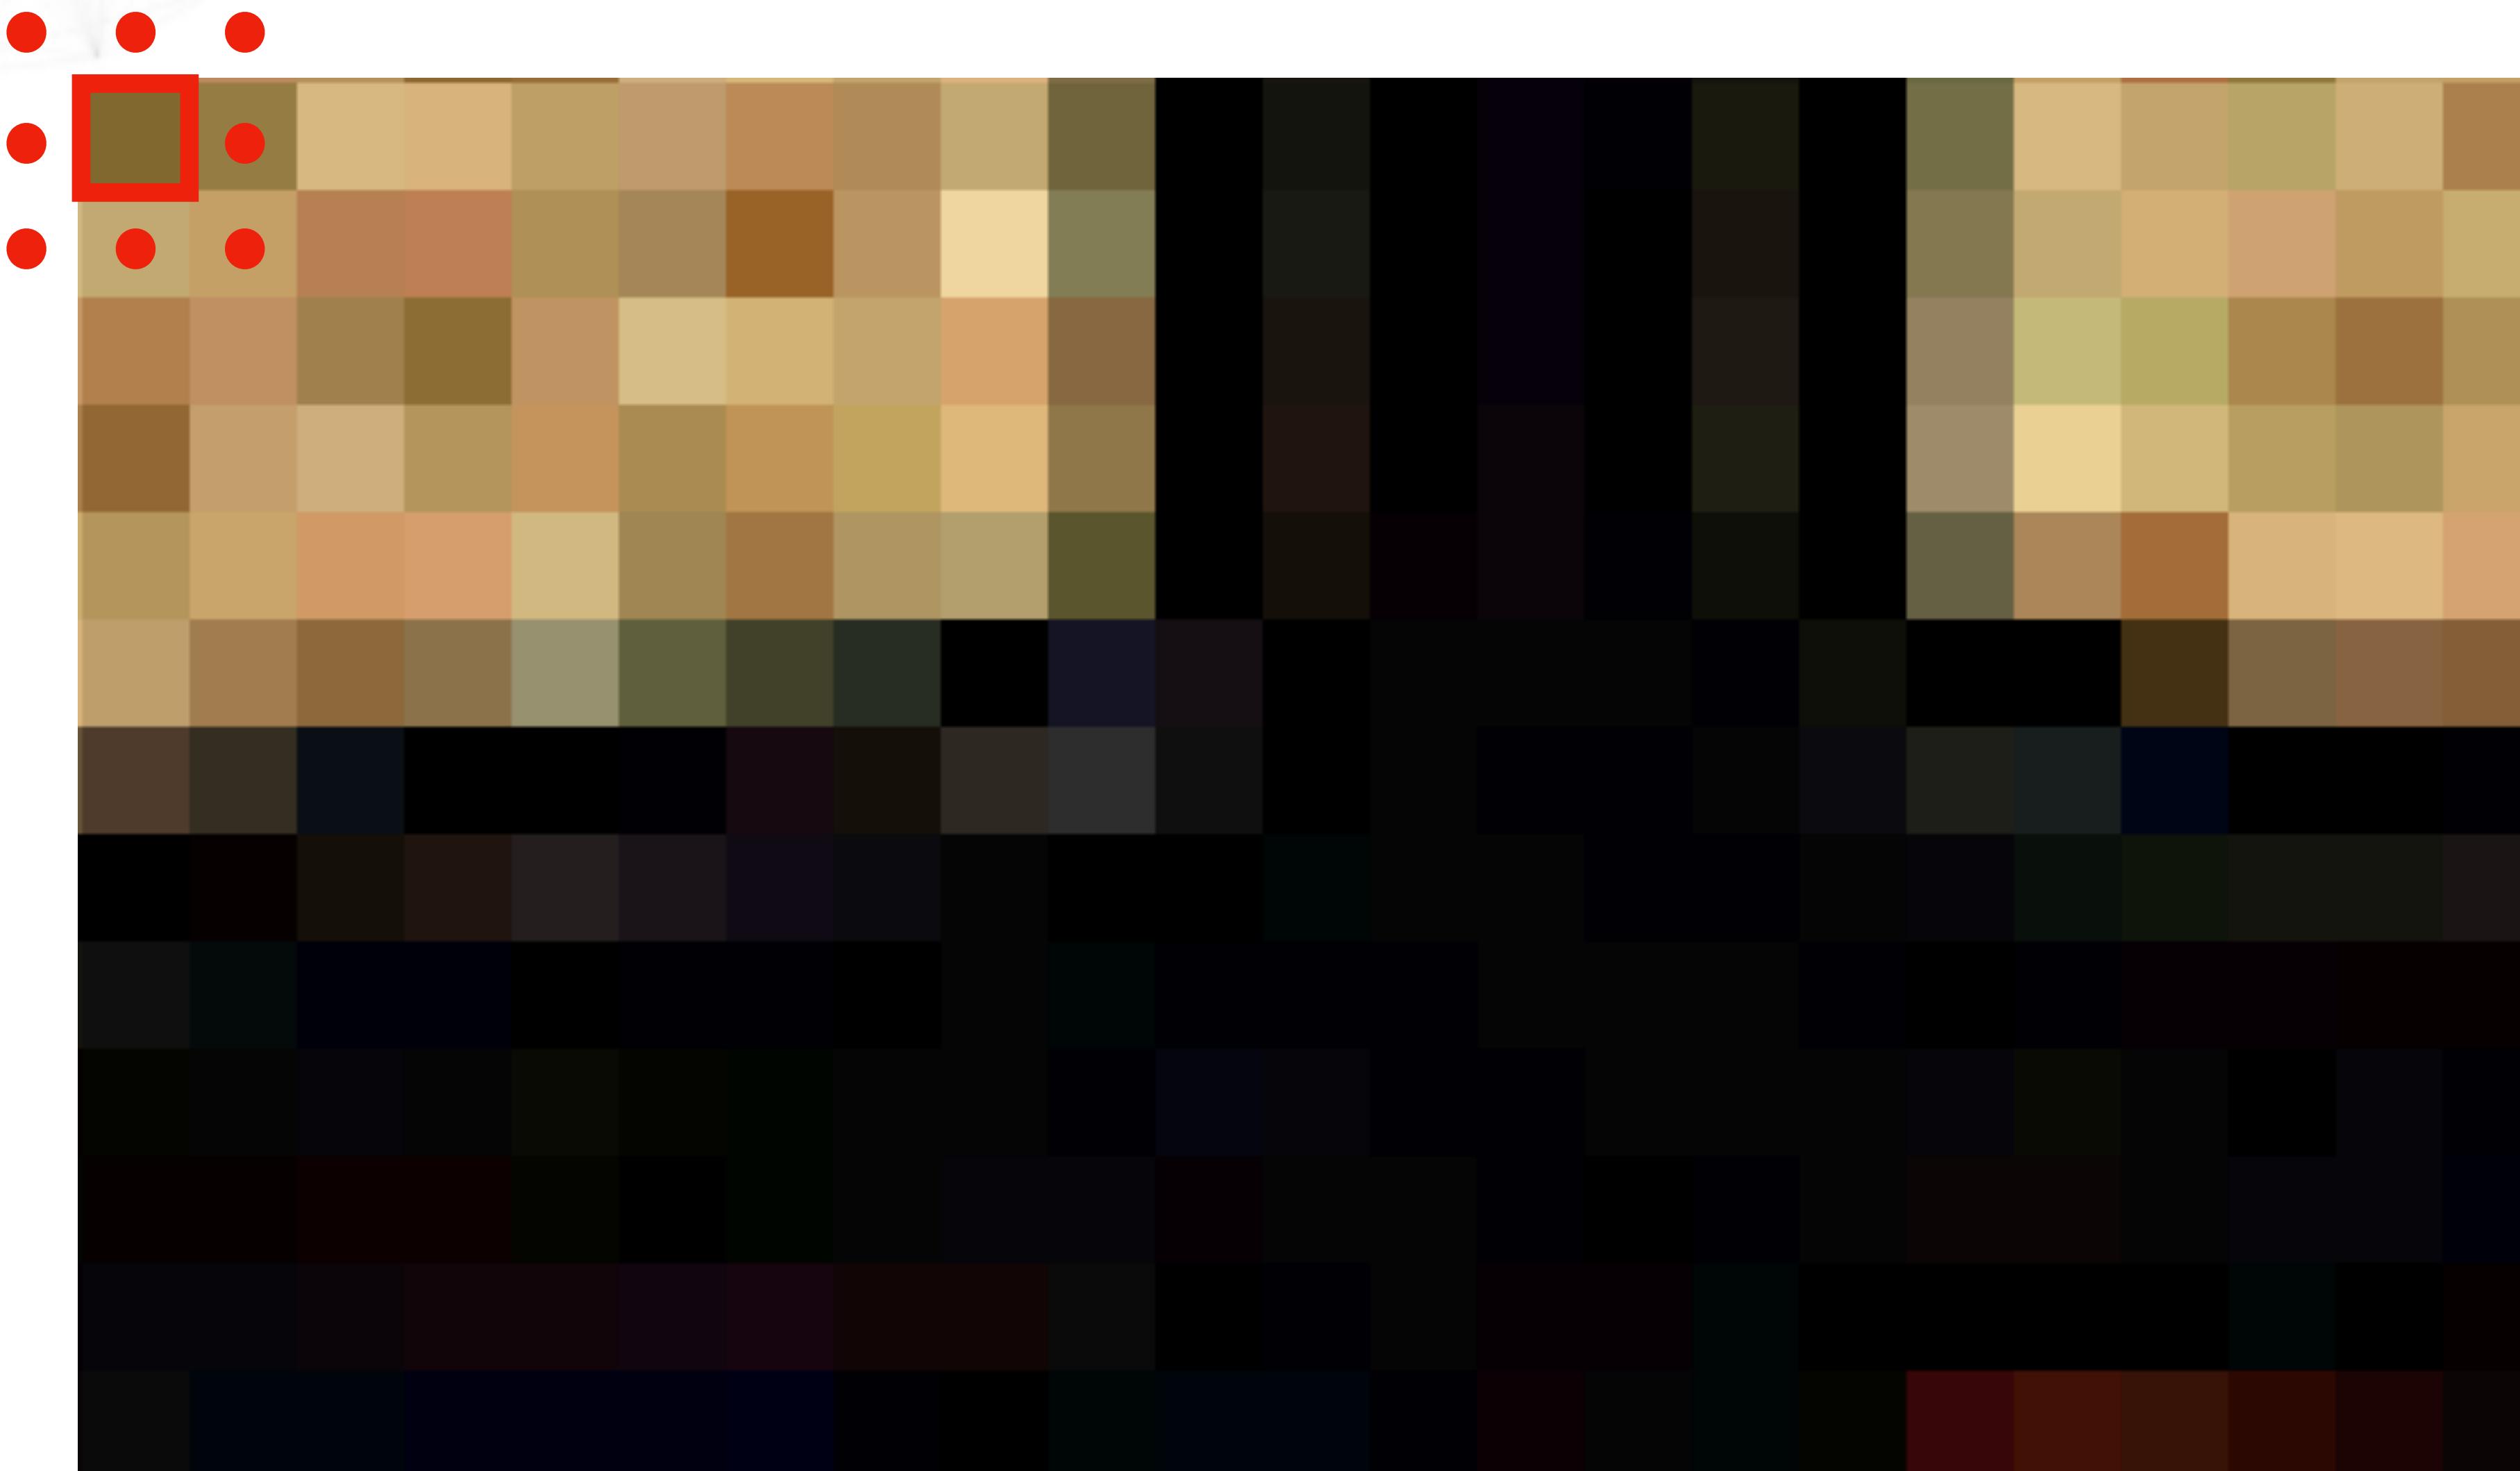
stride=(2, 2), padding='valid'



stride=(2, 2), padding='valid'



stride=(1, 1), padding='same'



stride=(1, 1), padding='same'



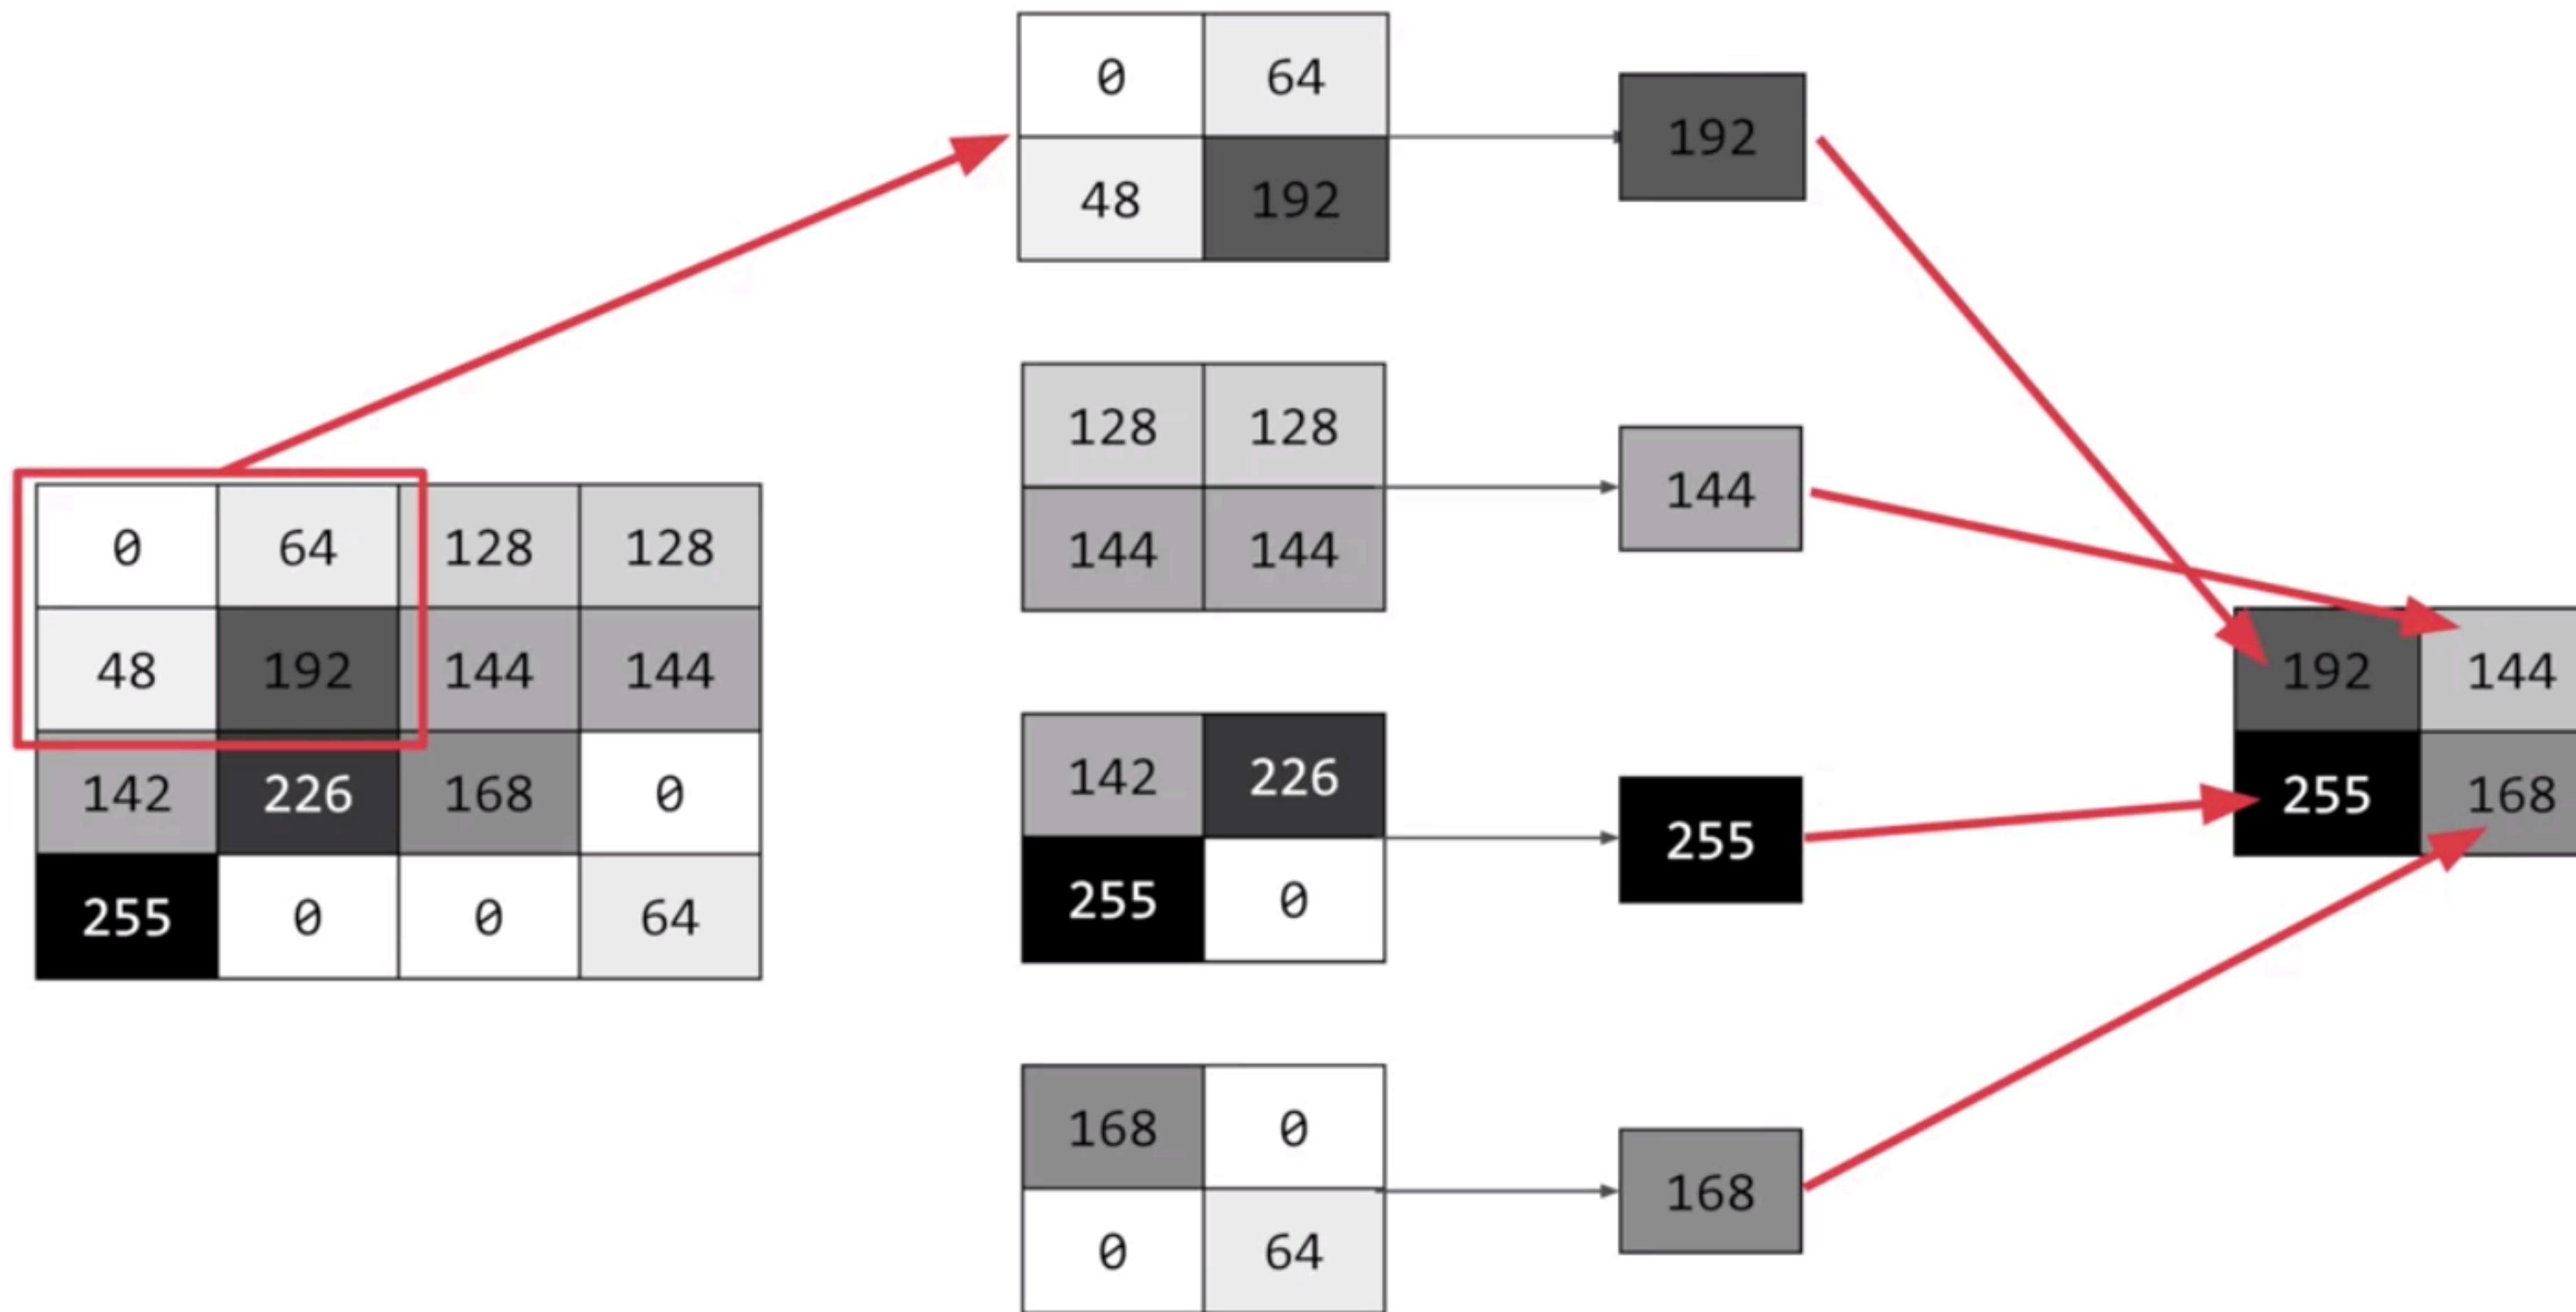
stride=(1, 1), padding='same'



stride=(1, 1), padding='same'

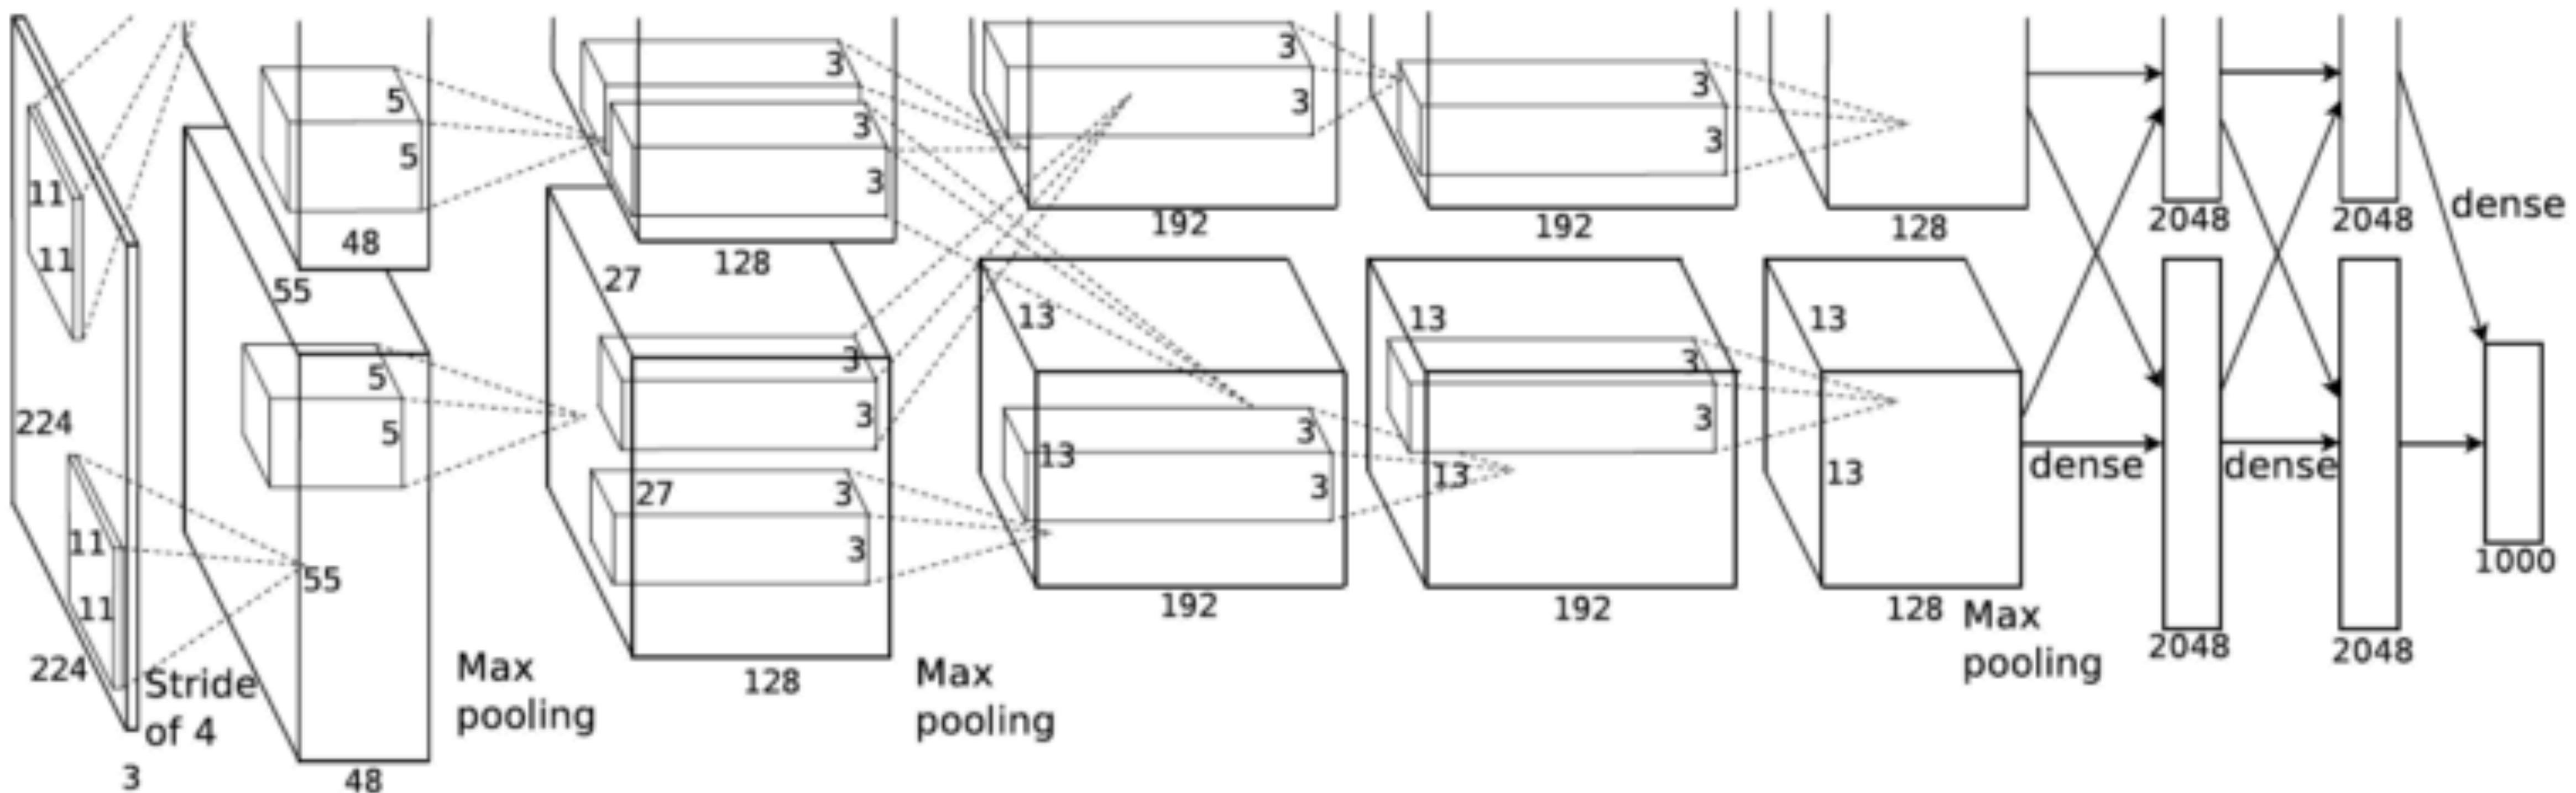


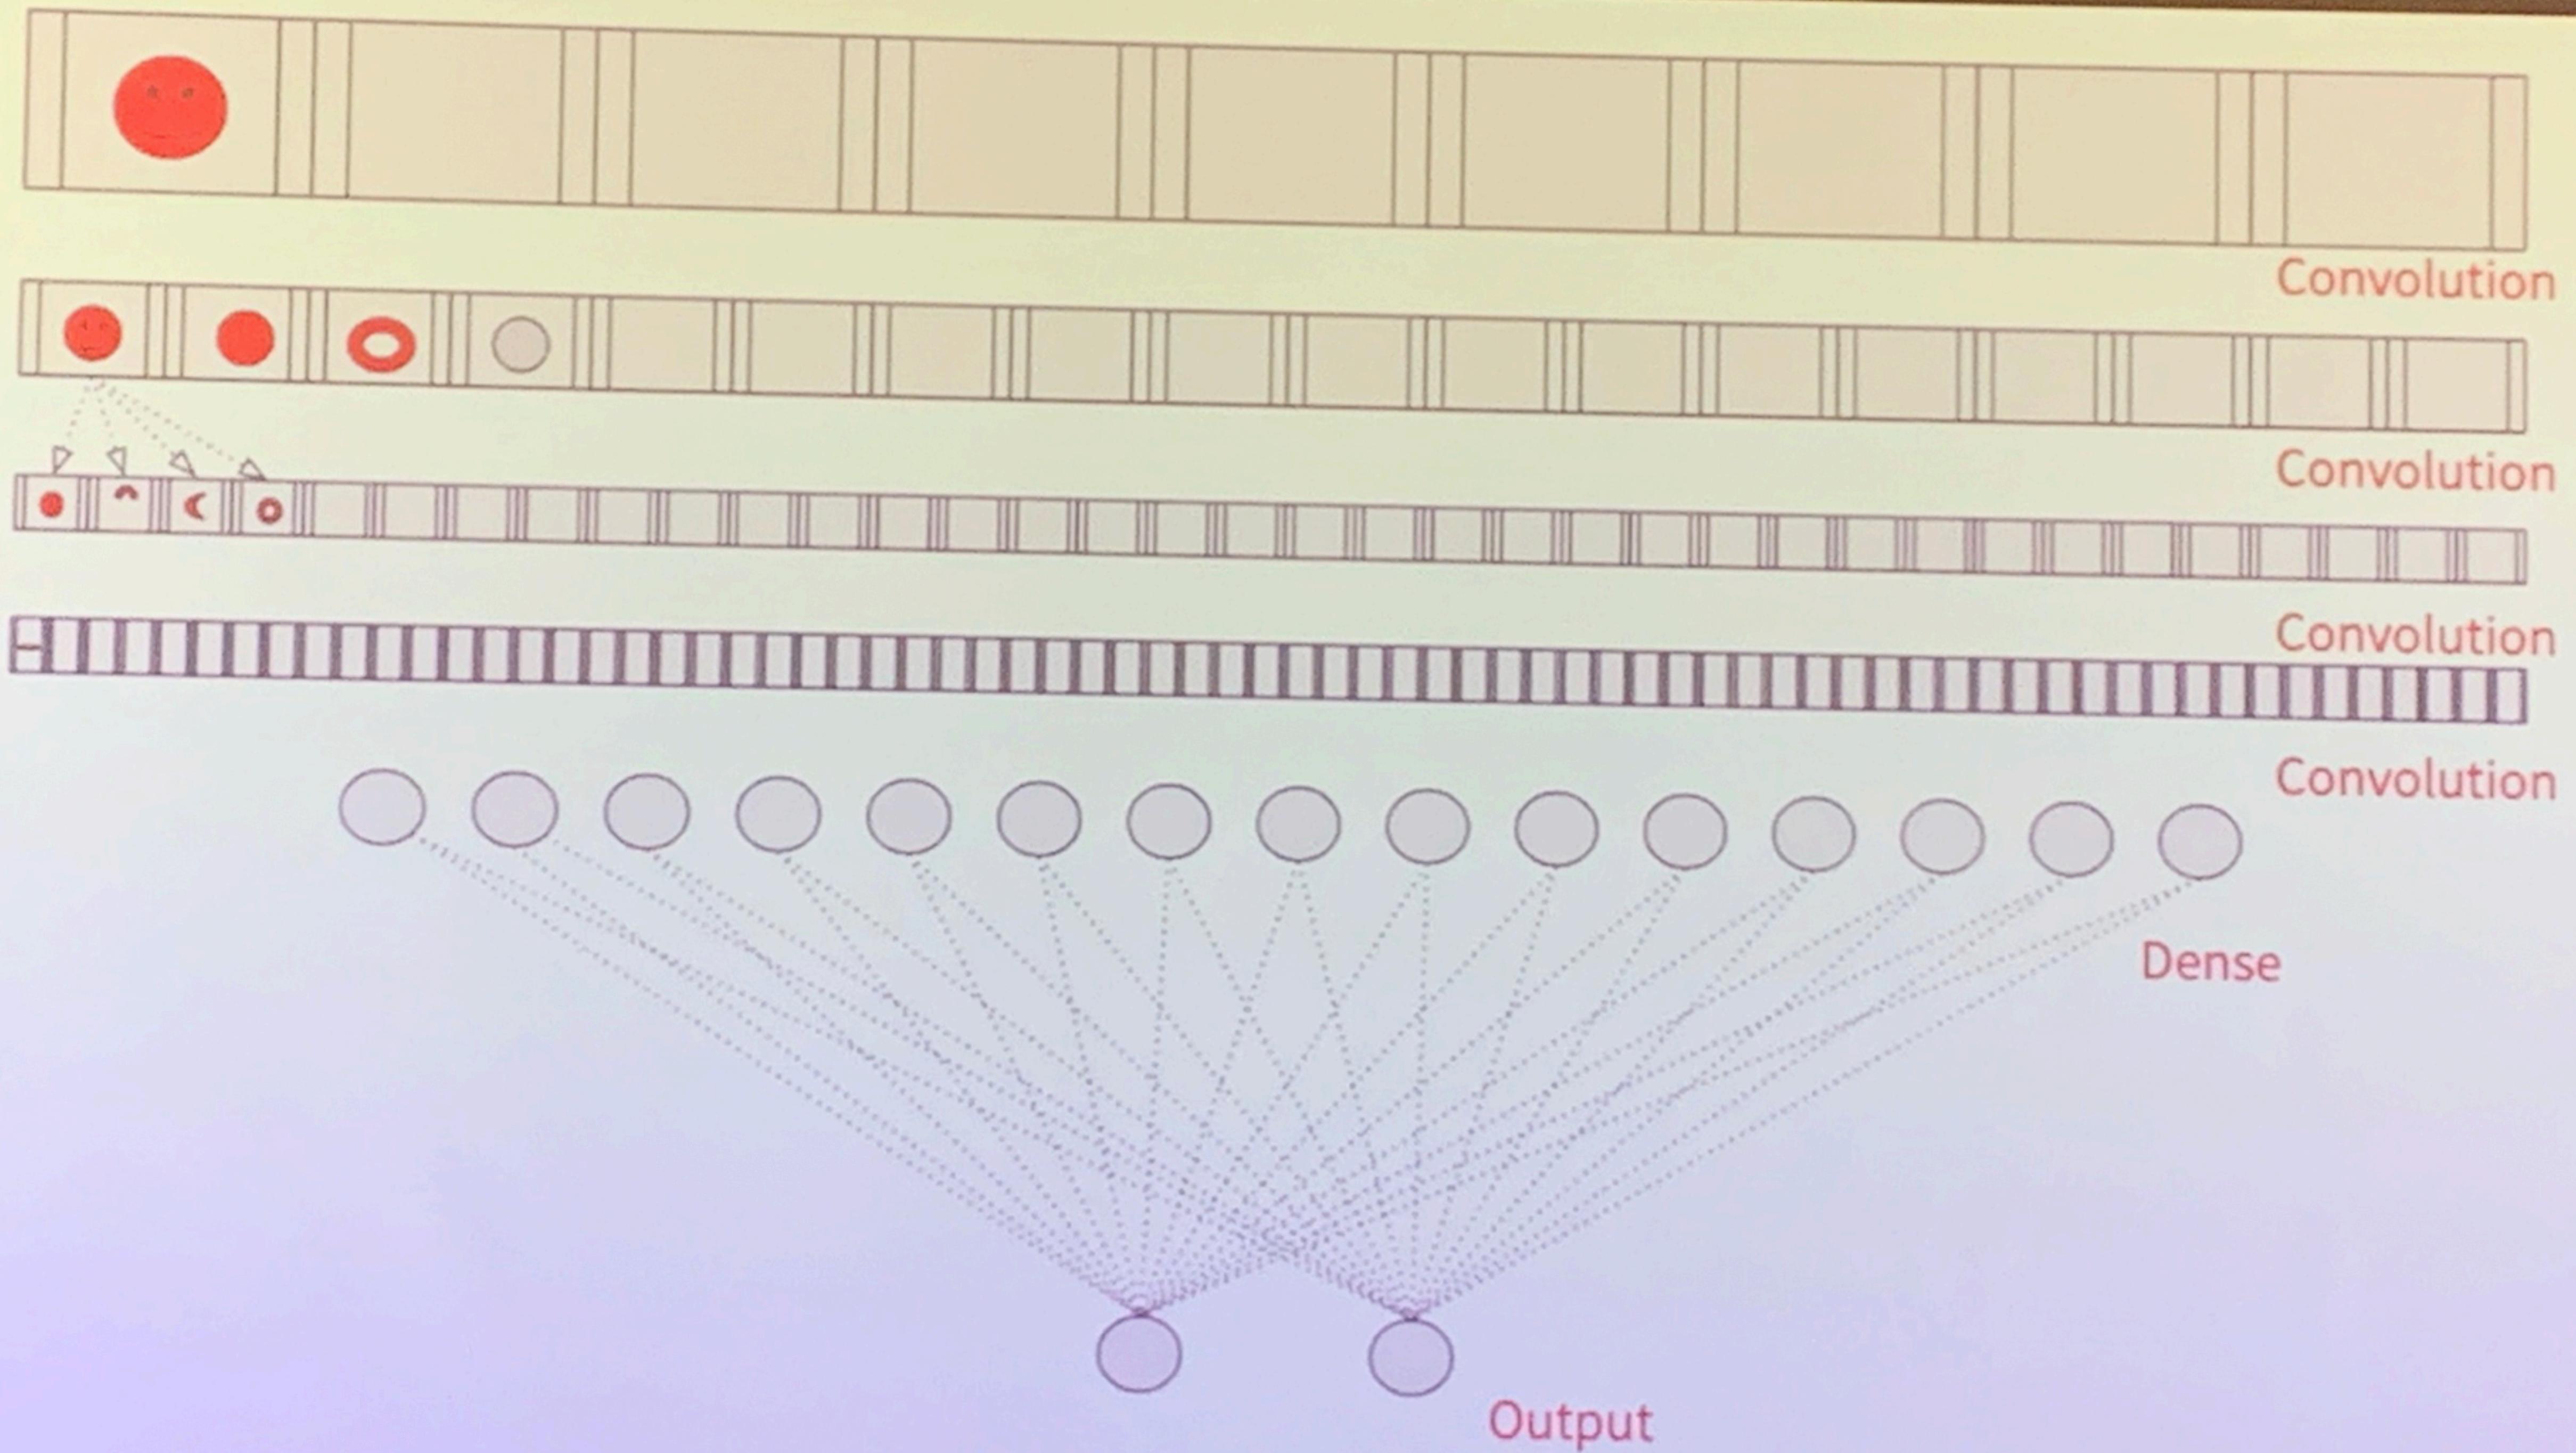
Max Pooling

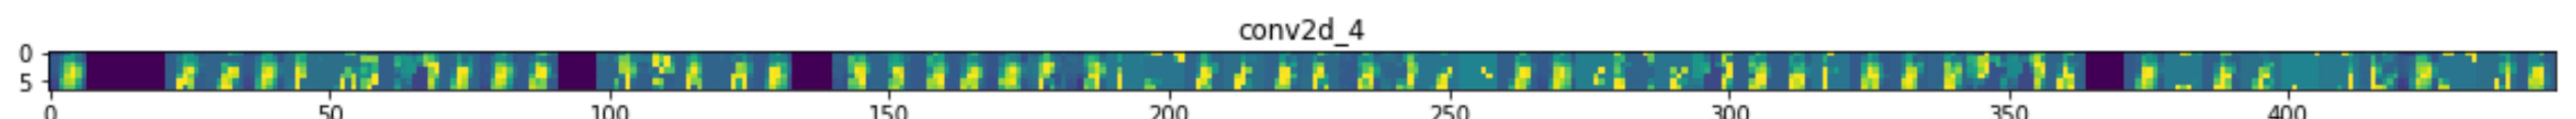
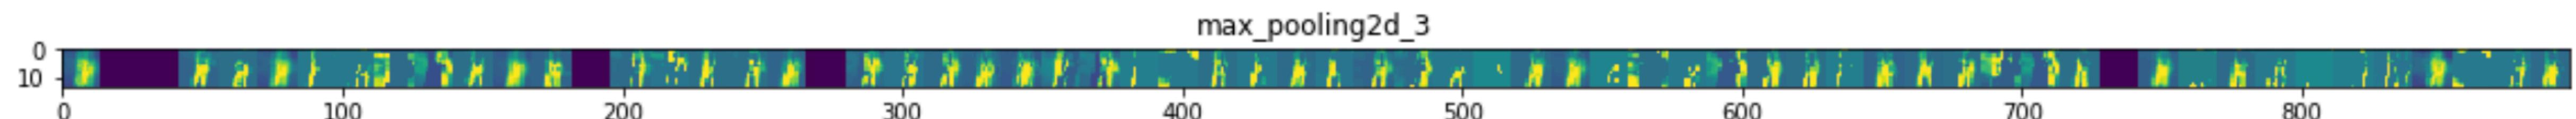
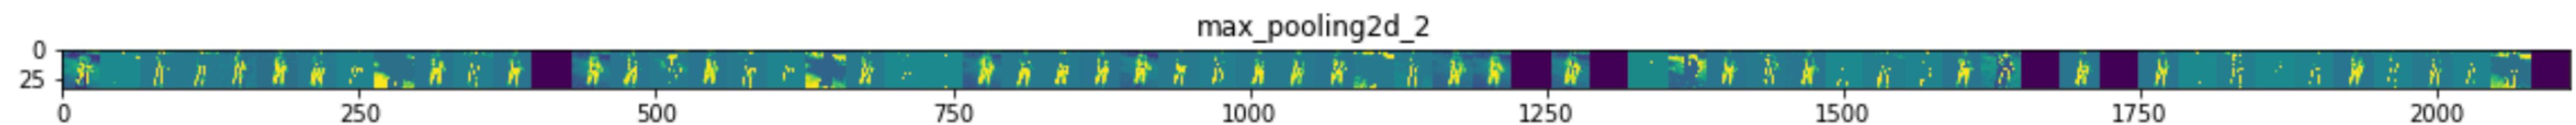
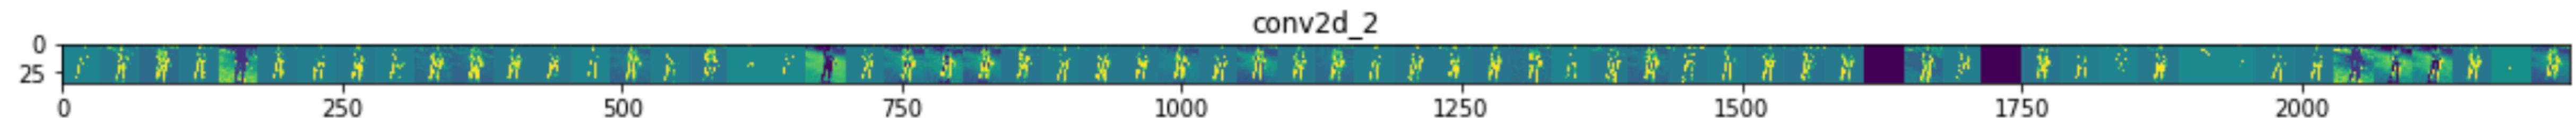
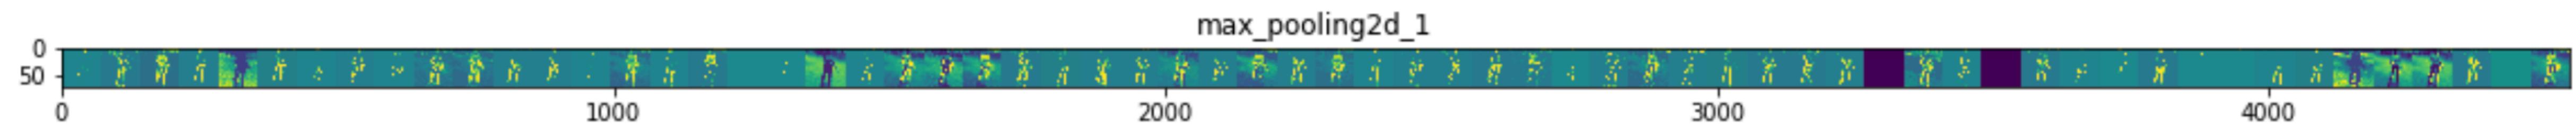
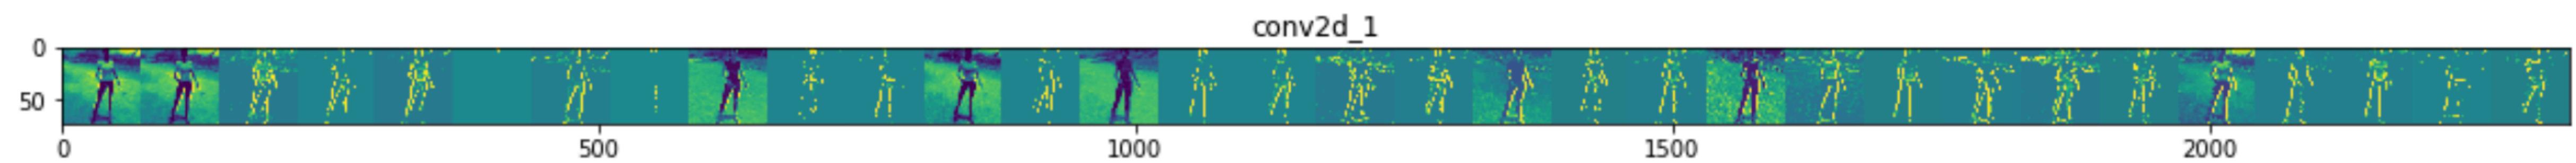
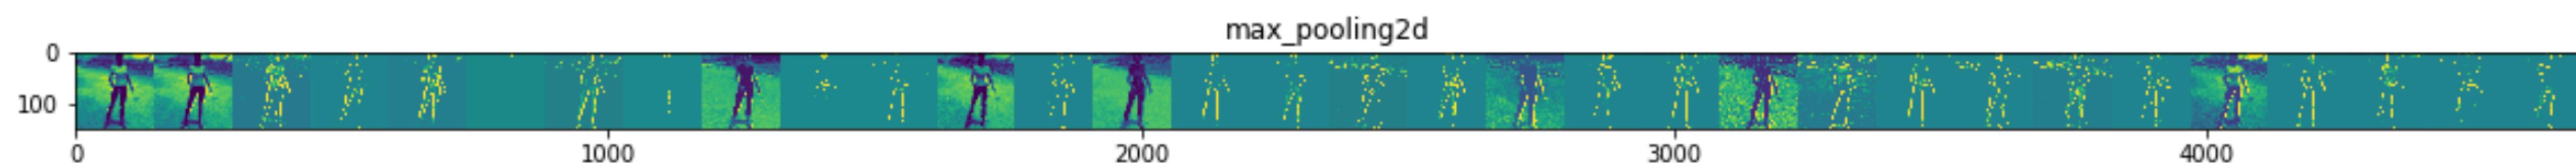
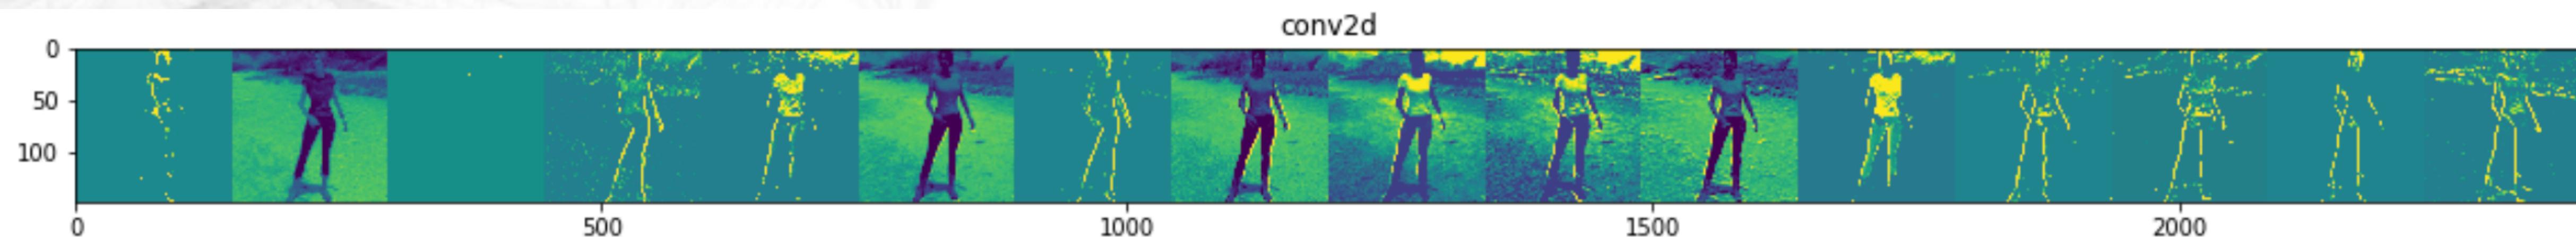


<https://www.coursera.org/learn/introduction-tensorflow>

AlexNet



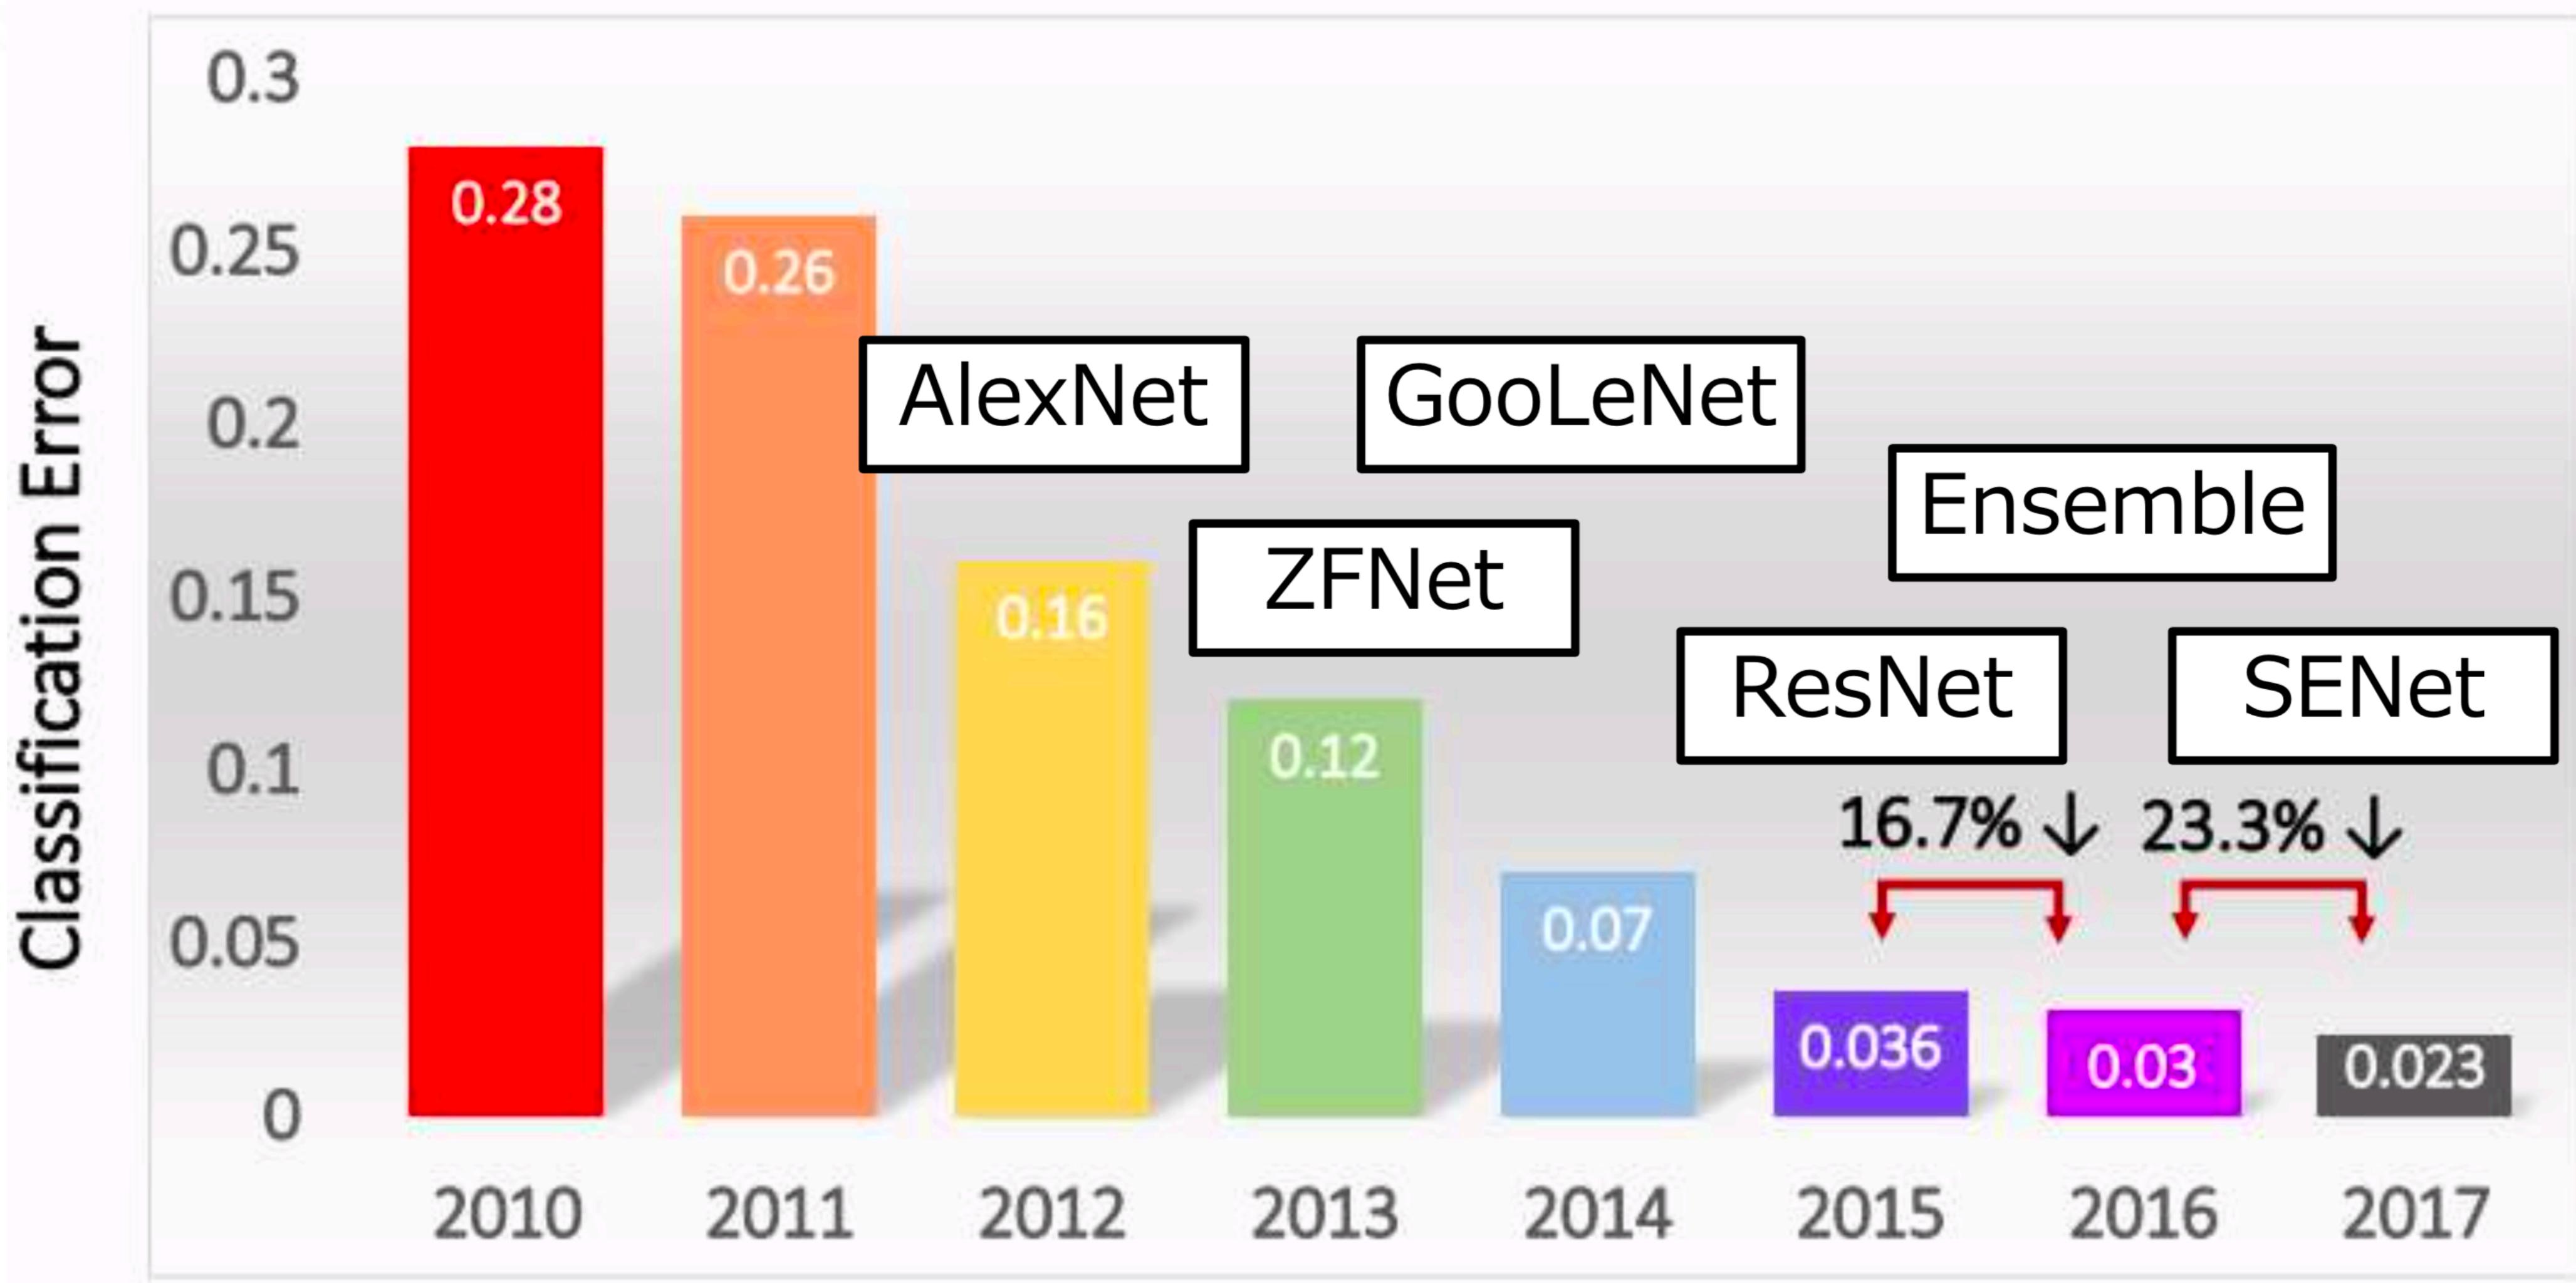






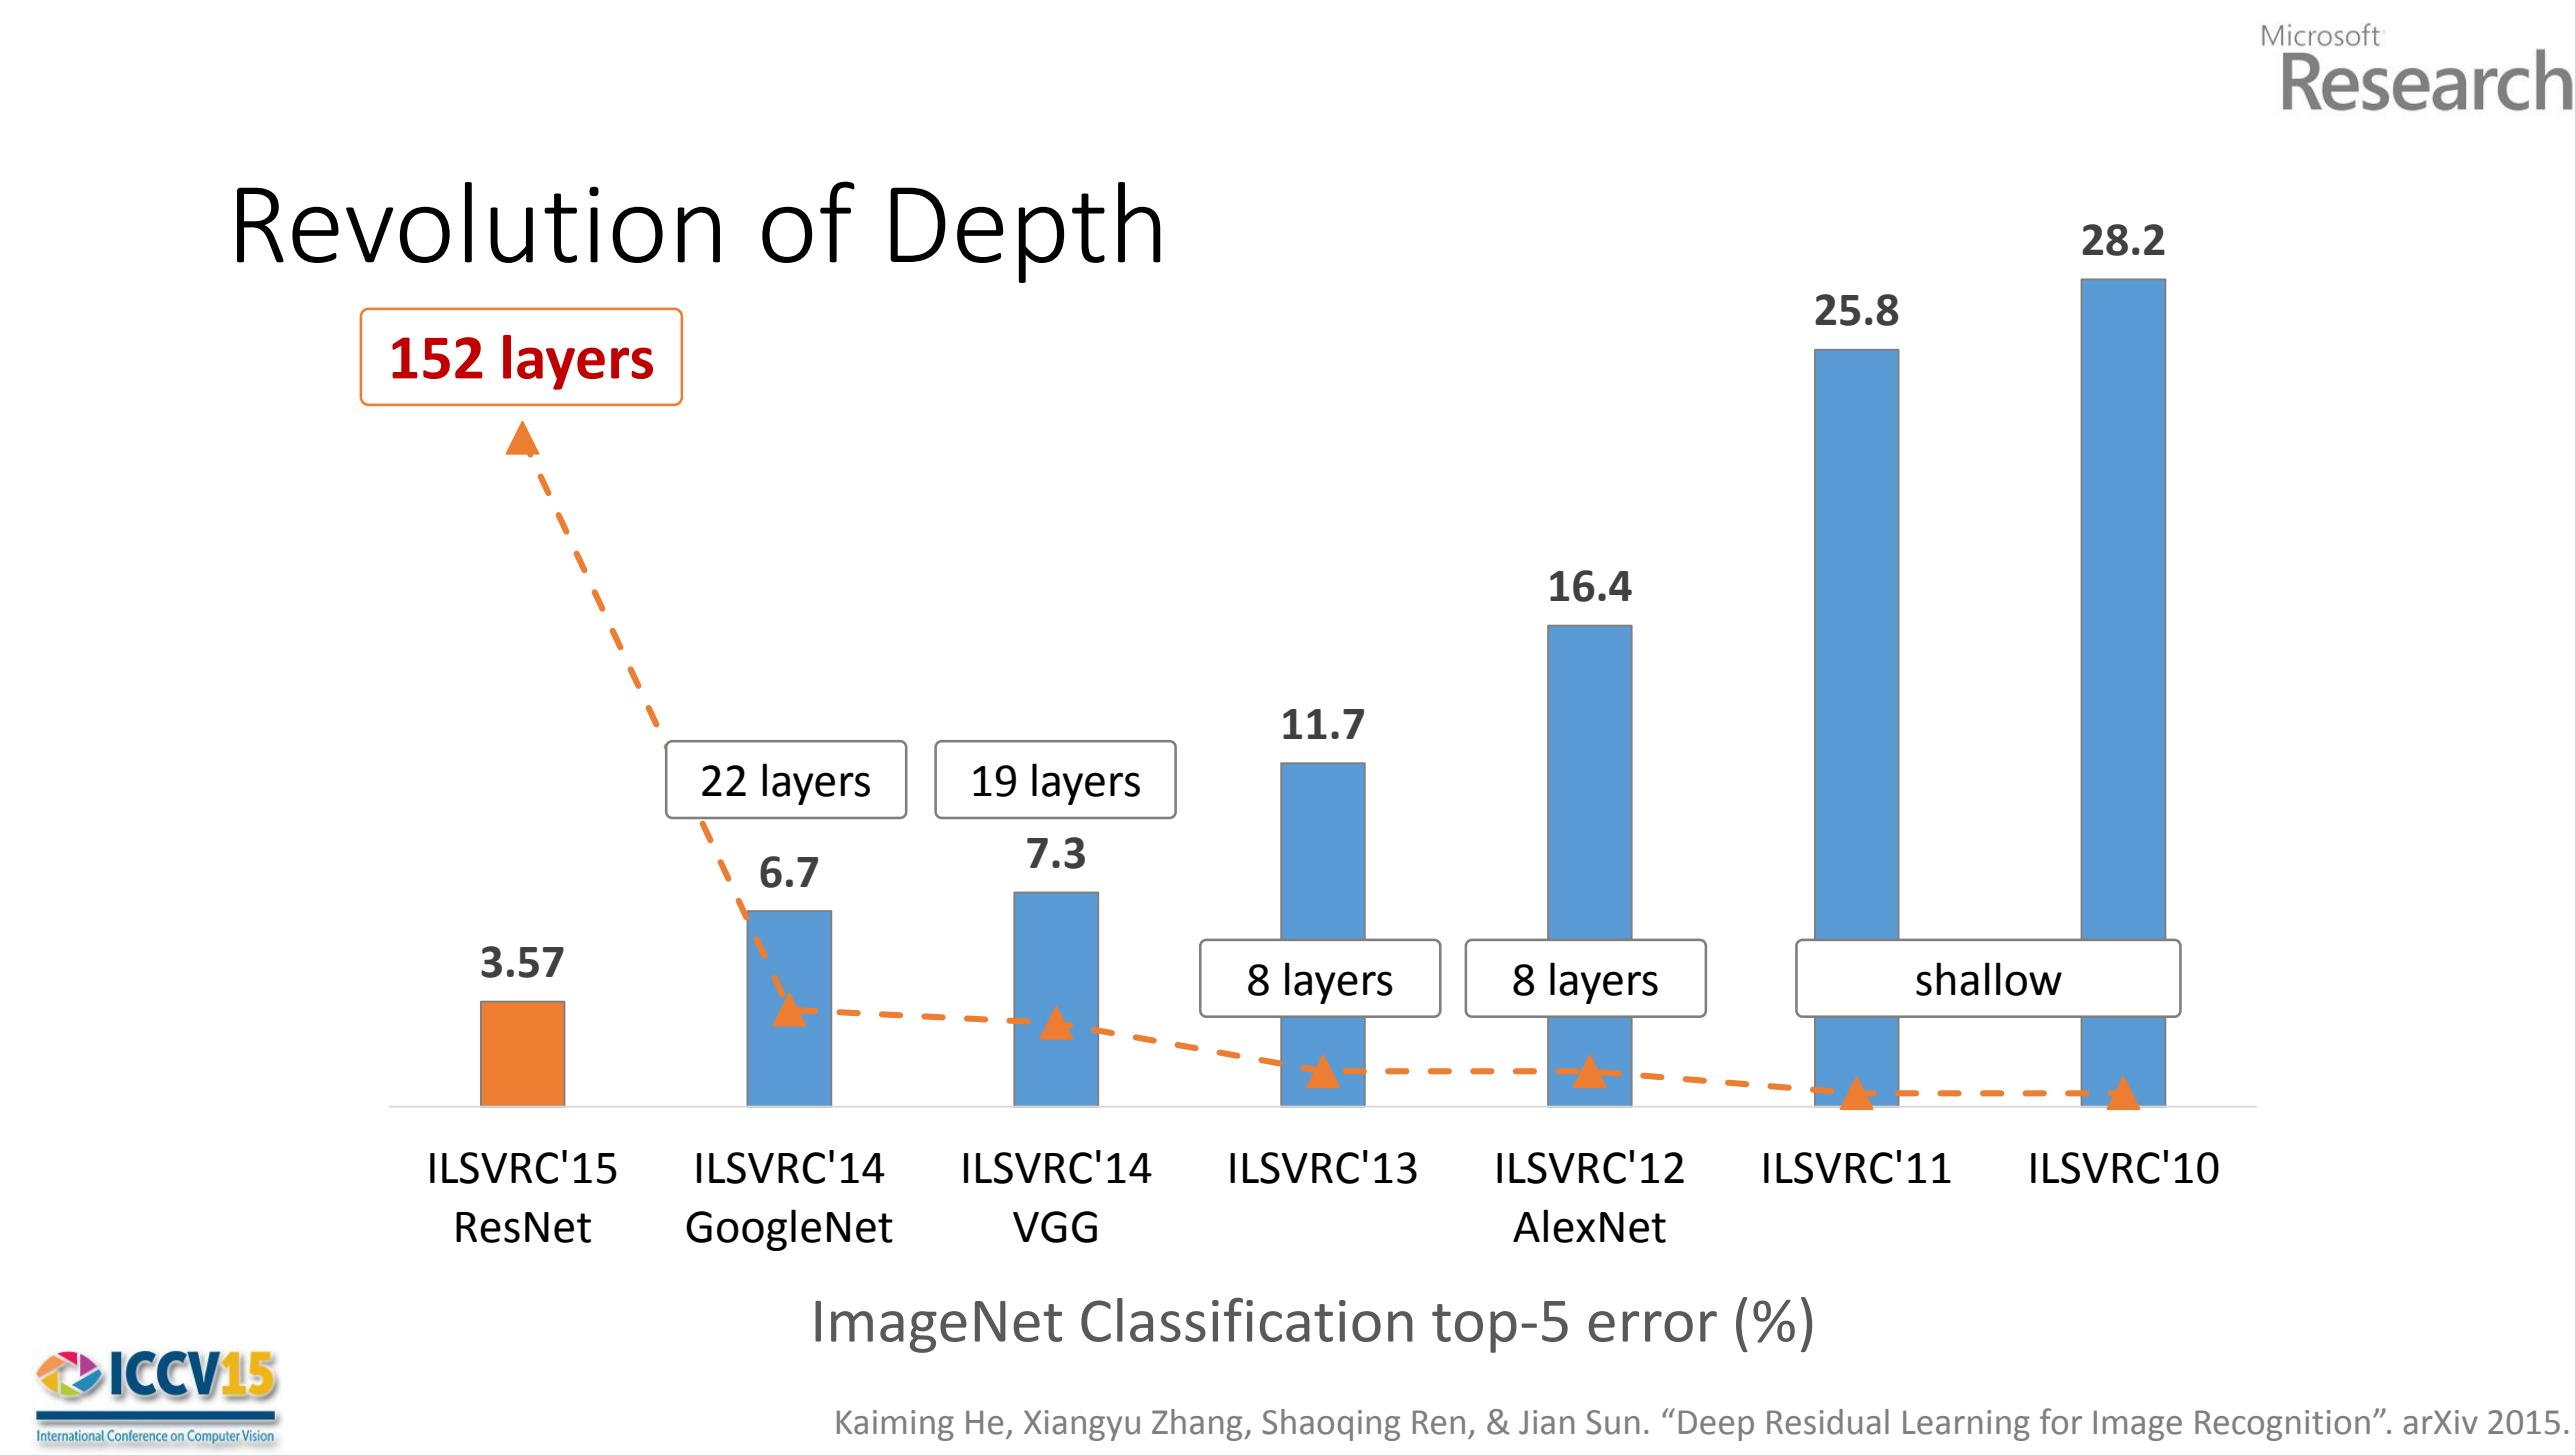
応用的なCNNアーキテクチャの紹介

ImageNet コンペの優勝モデル



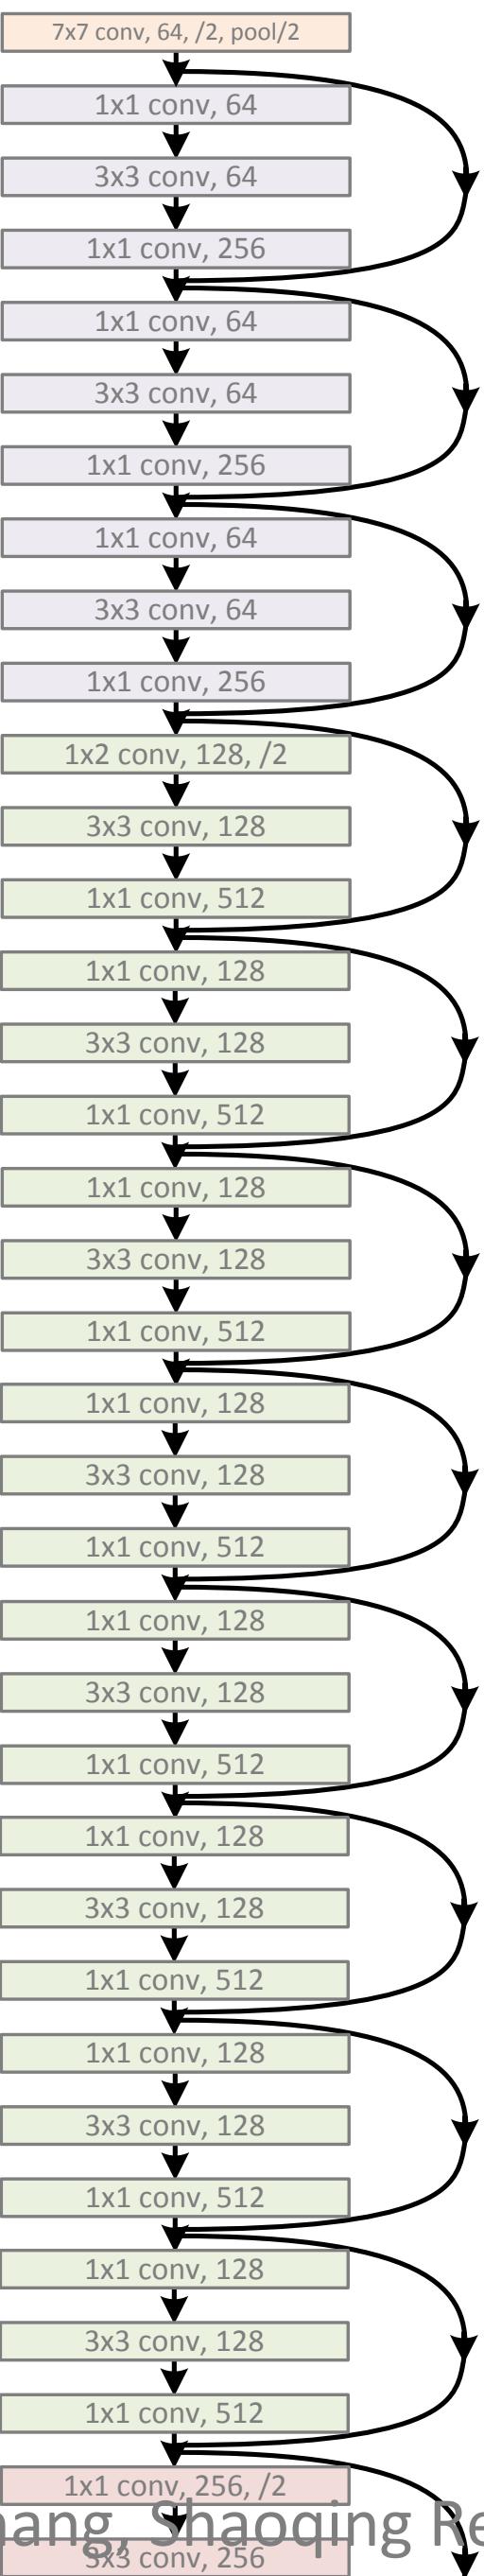
ResNet

- 2015年のImageNetコンペ (ILSVRC) 優勝モデル
- Residualモジュール（ショートカット機構）の導入



Revolution of Depth

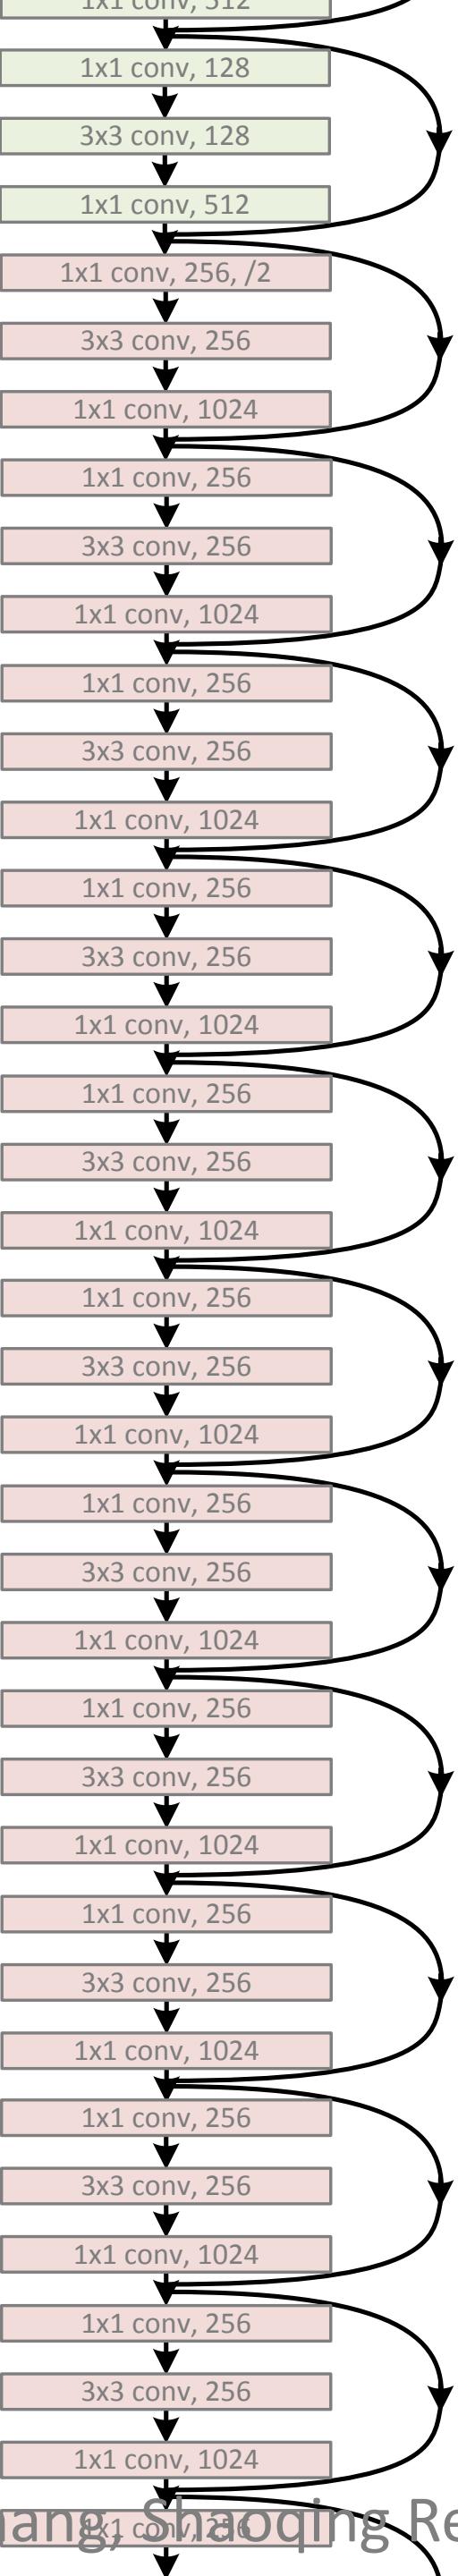
ResNet, 152 layers



(there was an animation here)

Revolution of Depth

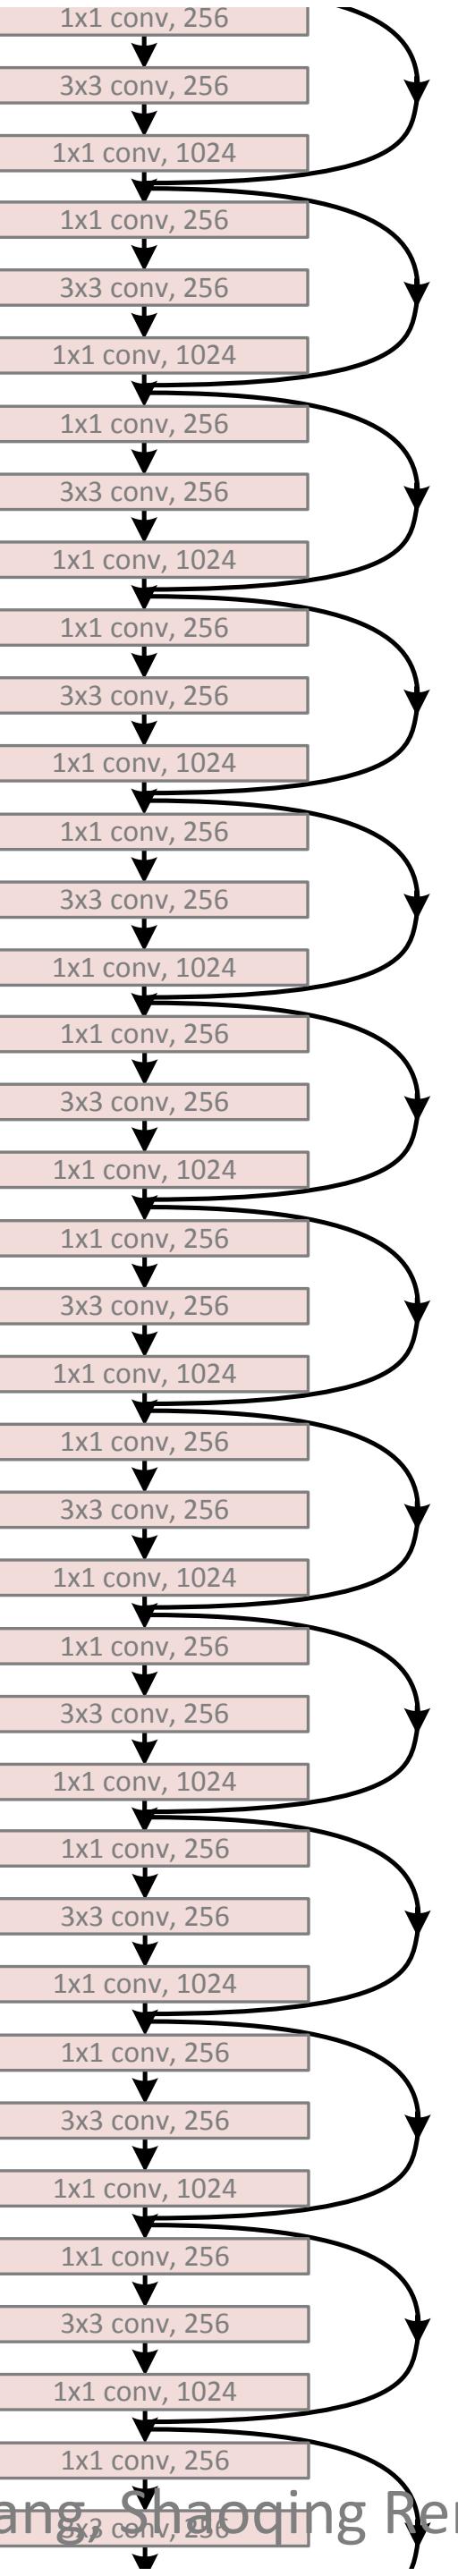
ResNet, 152 layers



(there was an animation here)

Revolution of Depth

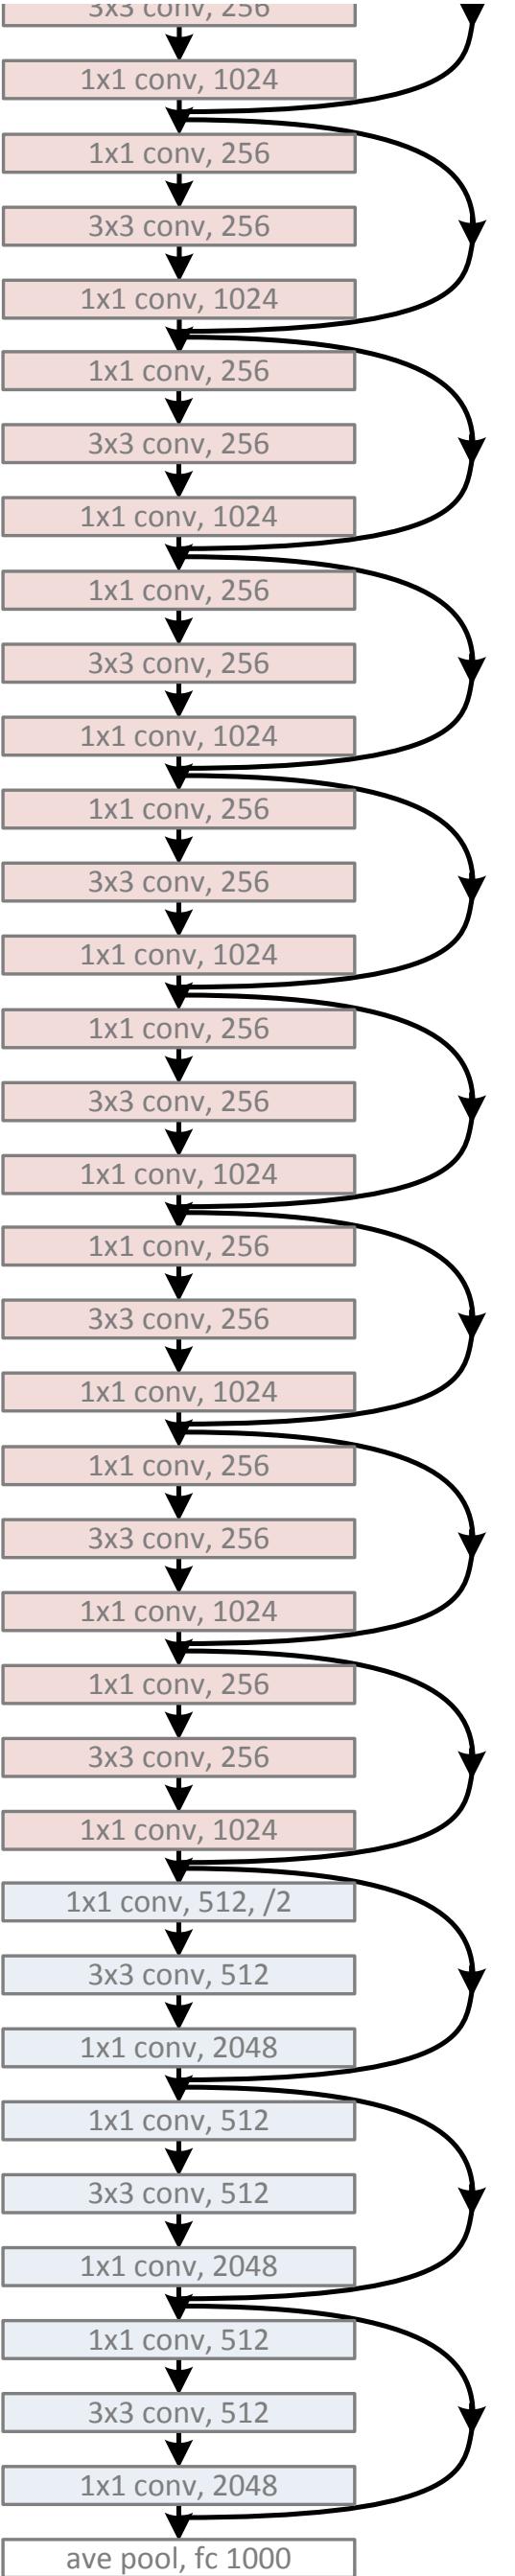
ResNet, 152 layers



(there was an animation here)

Revolution of Depth

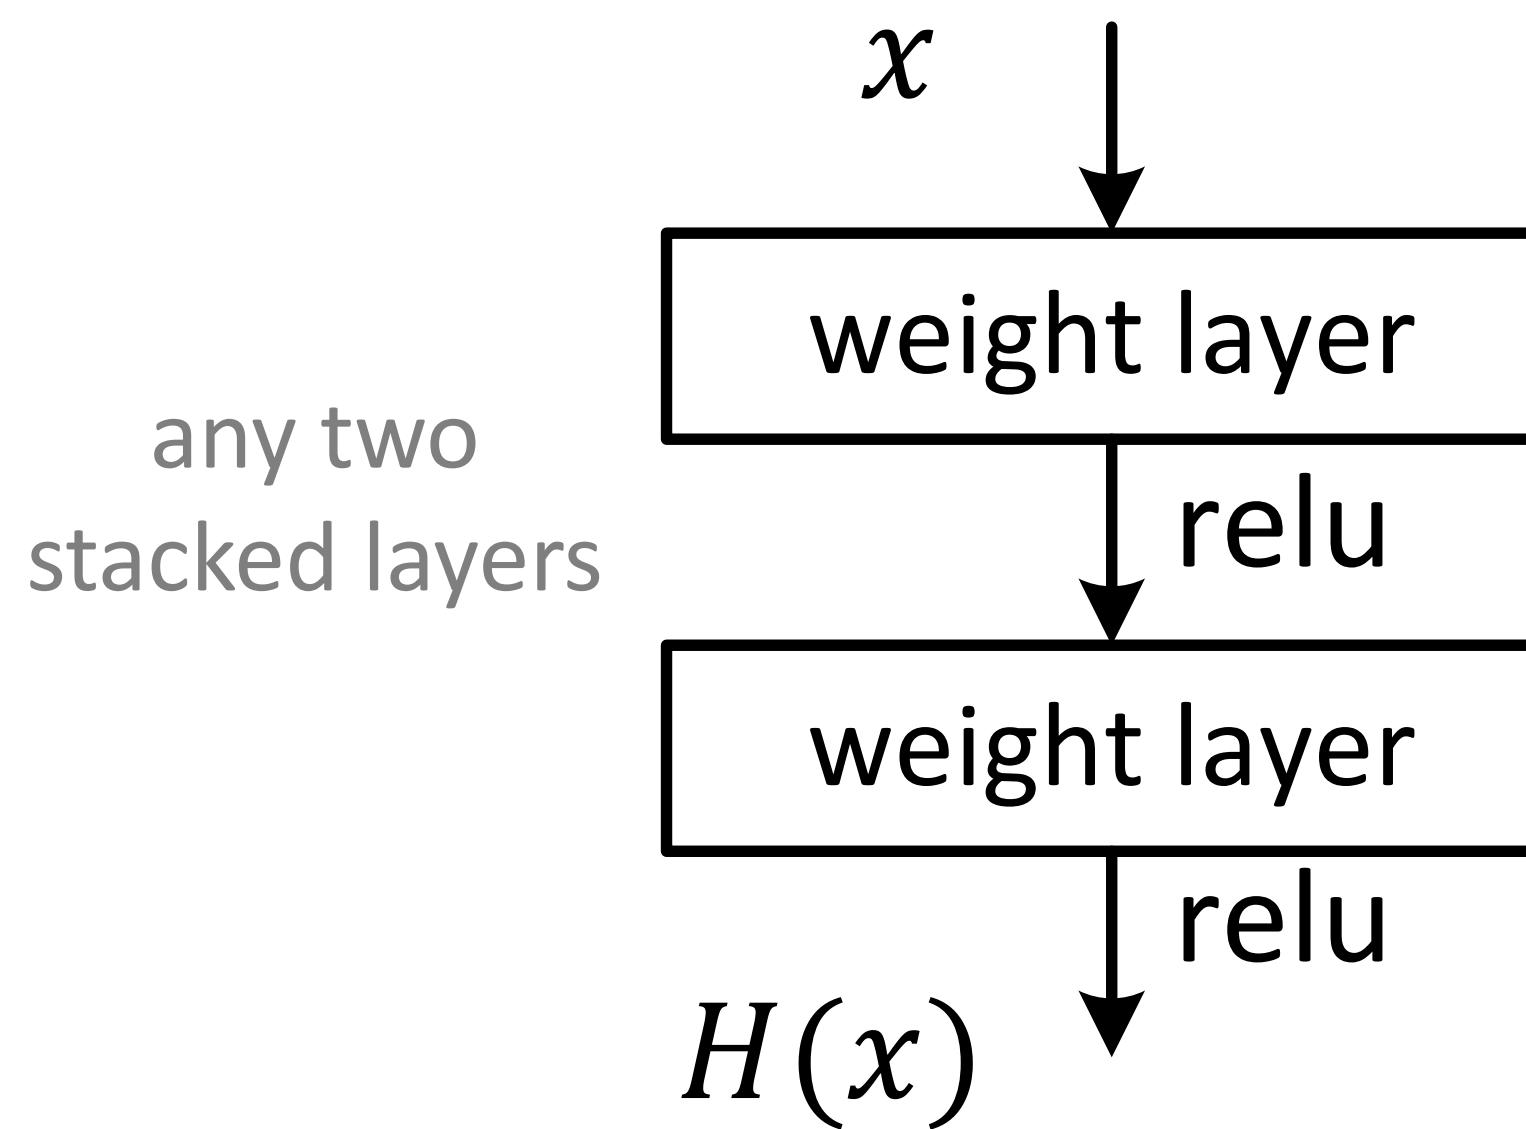
ResNet, 152 layers



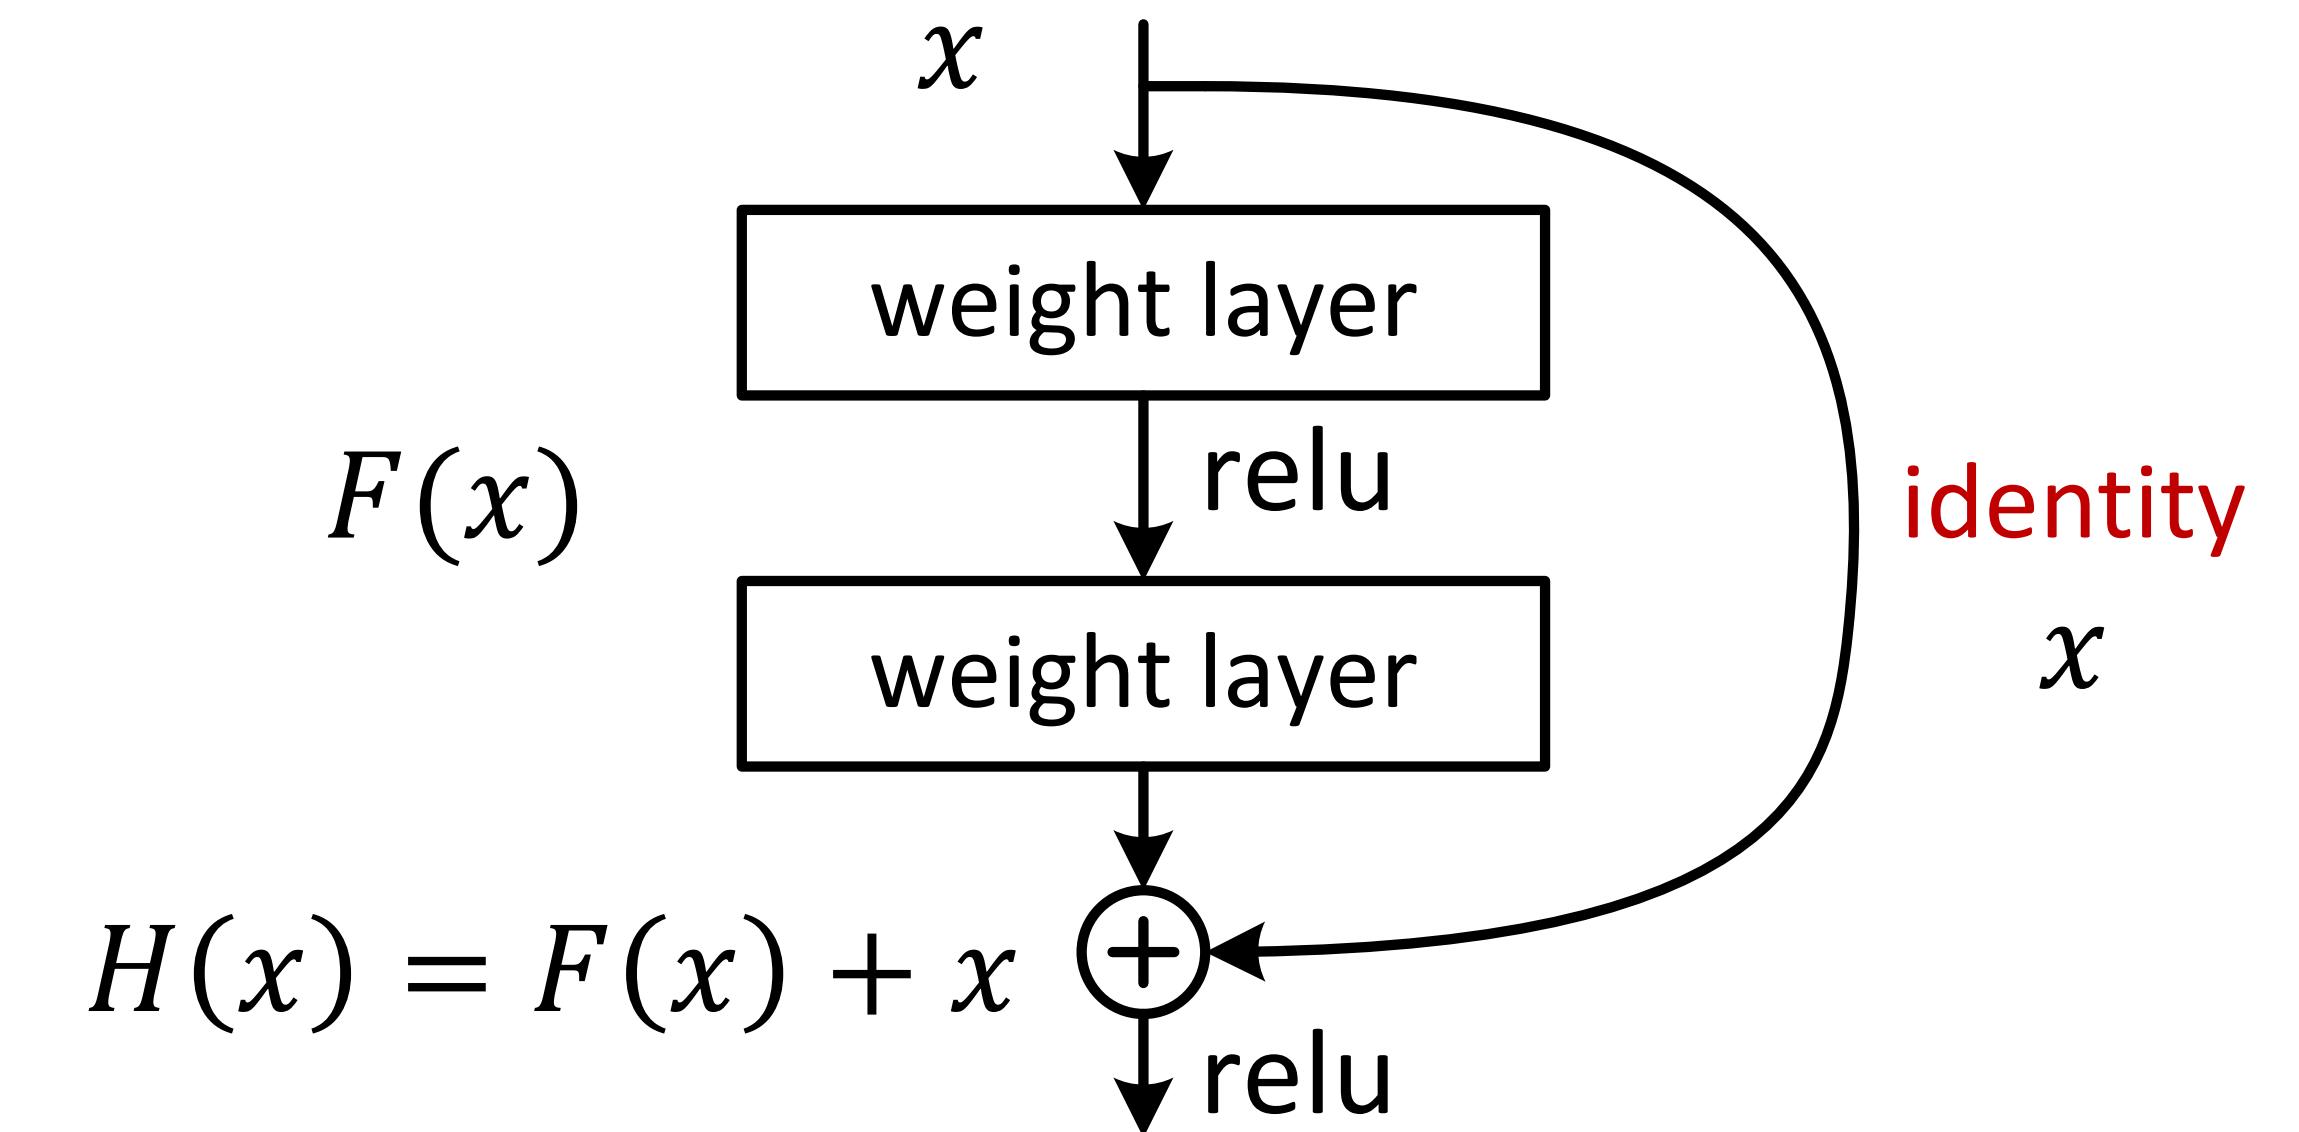
(there was an animation here)

ResNet

- Plain net

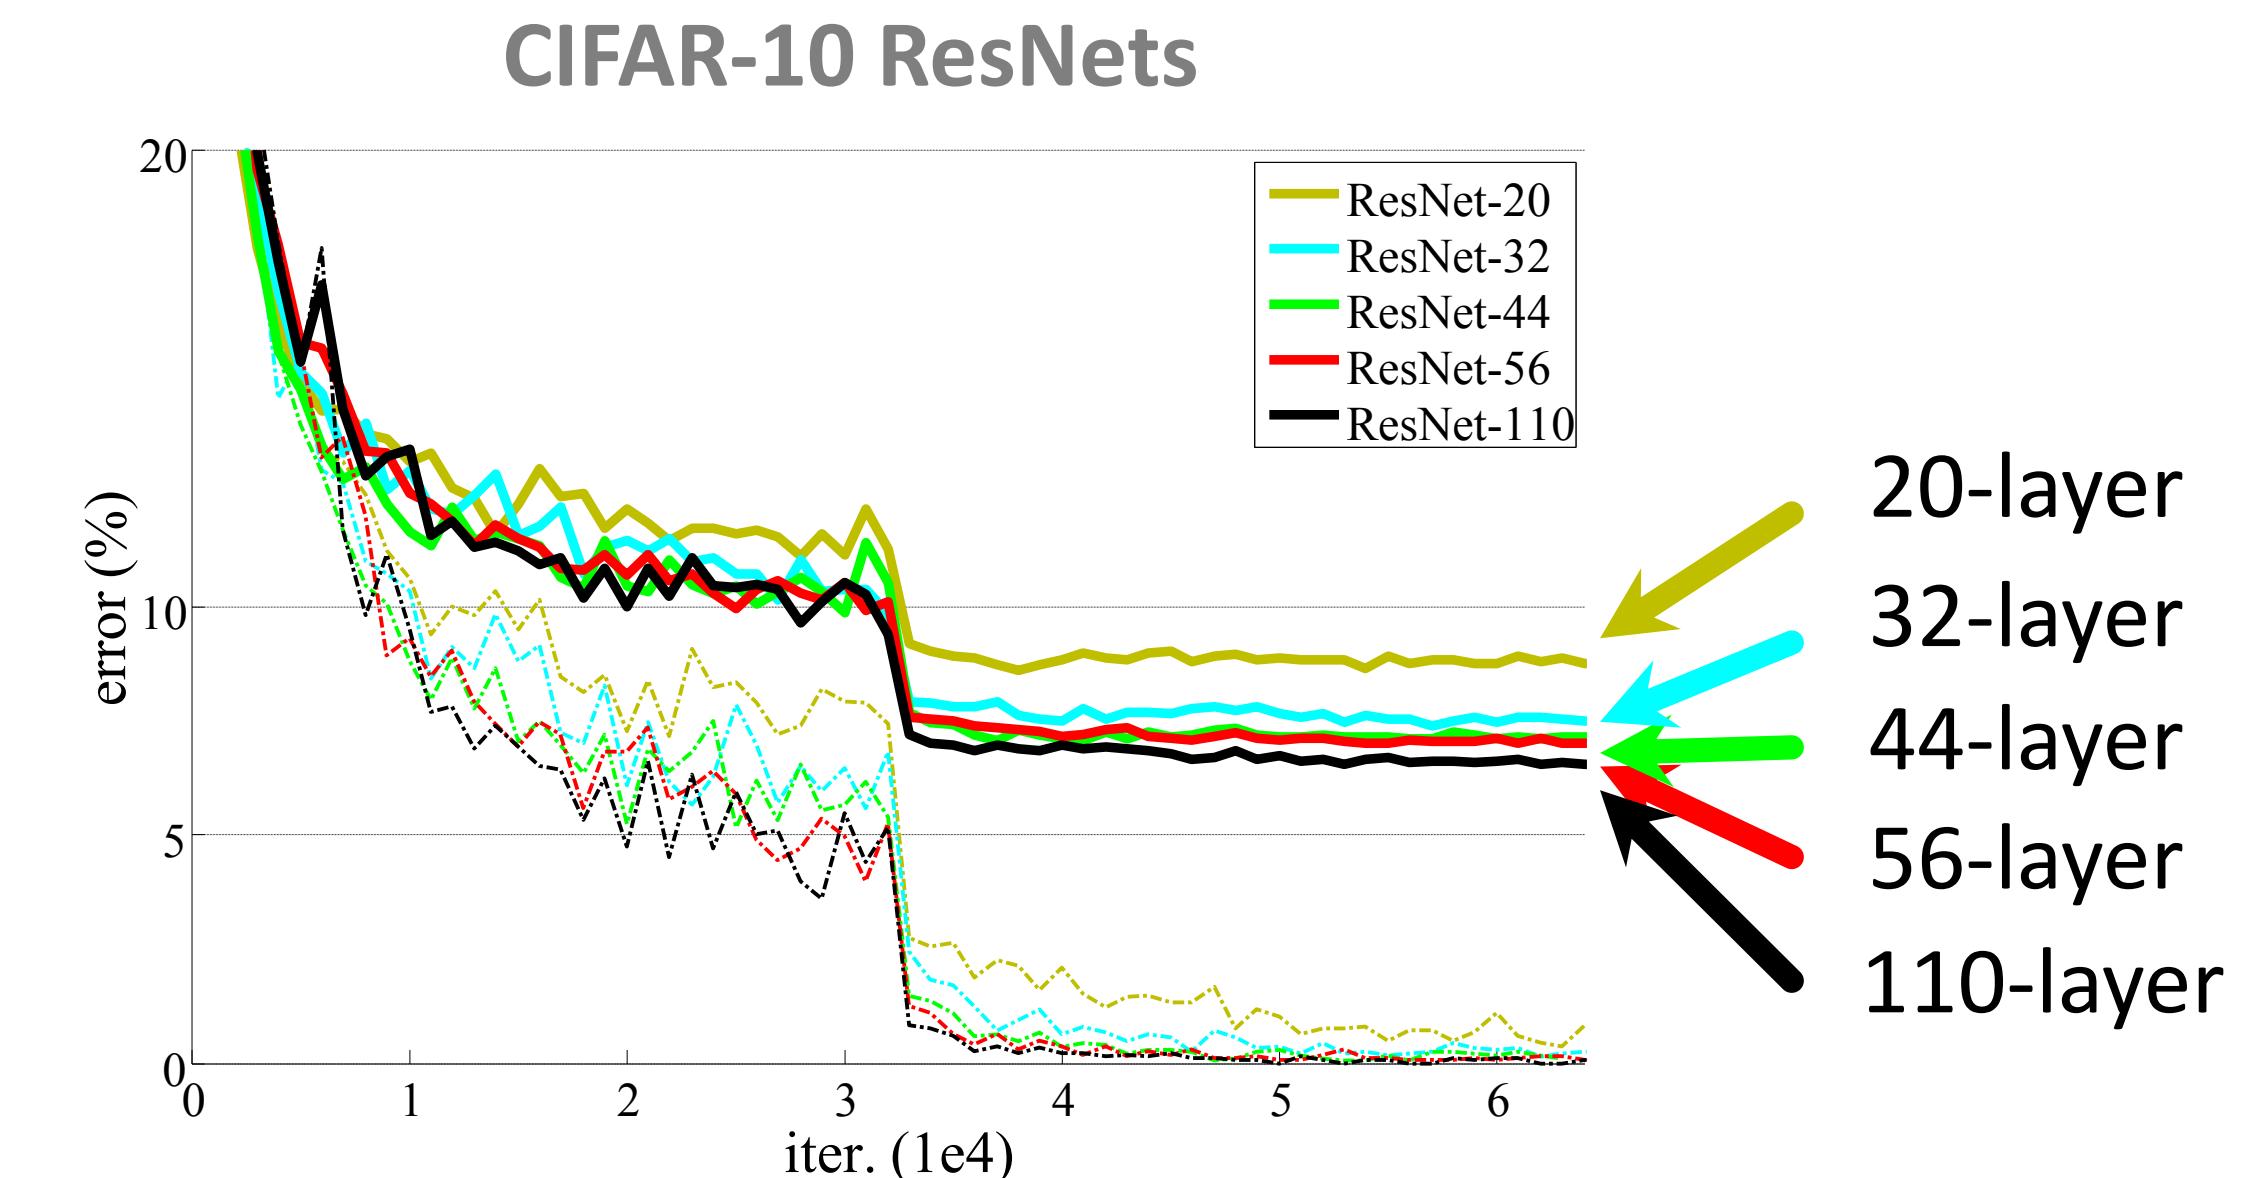
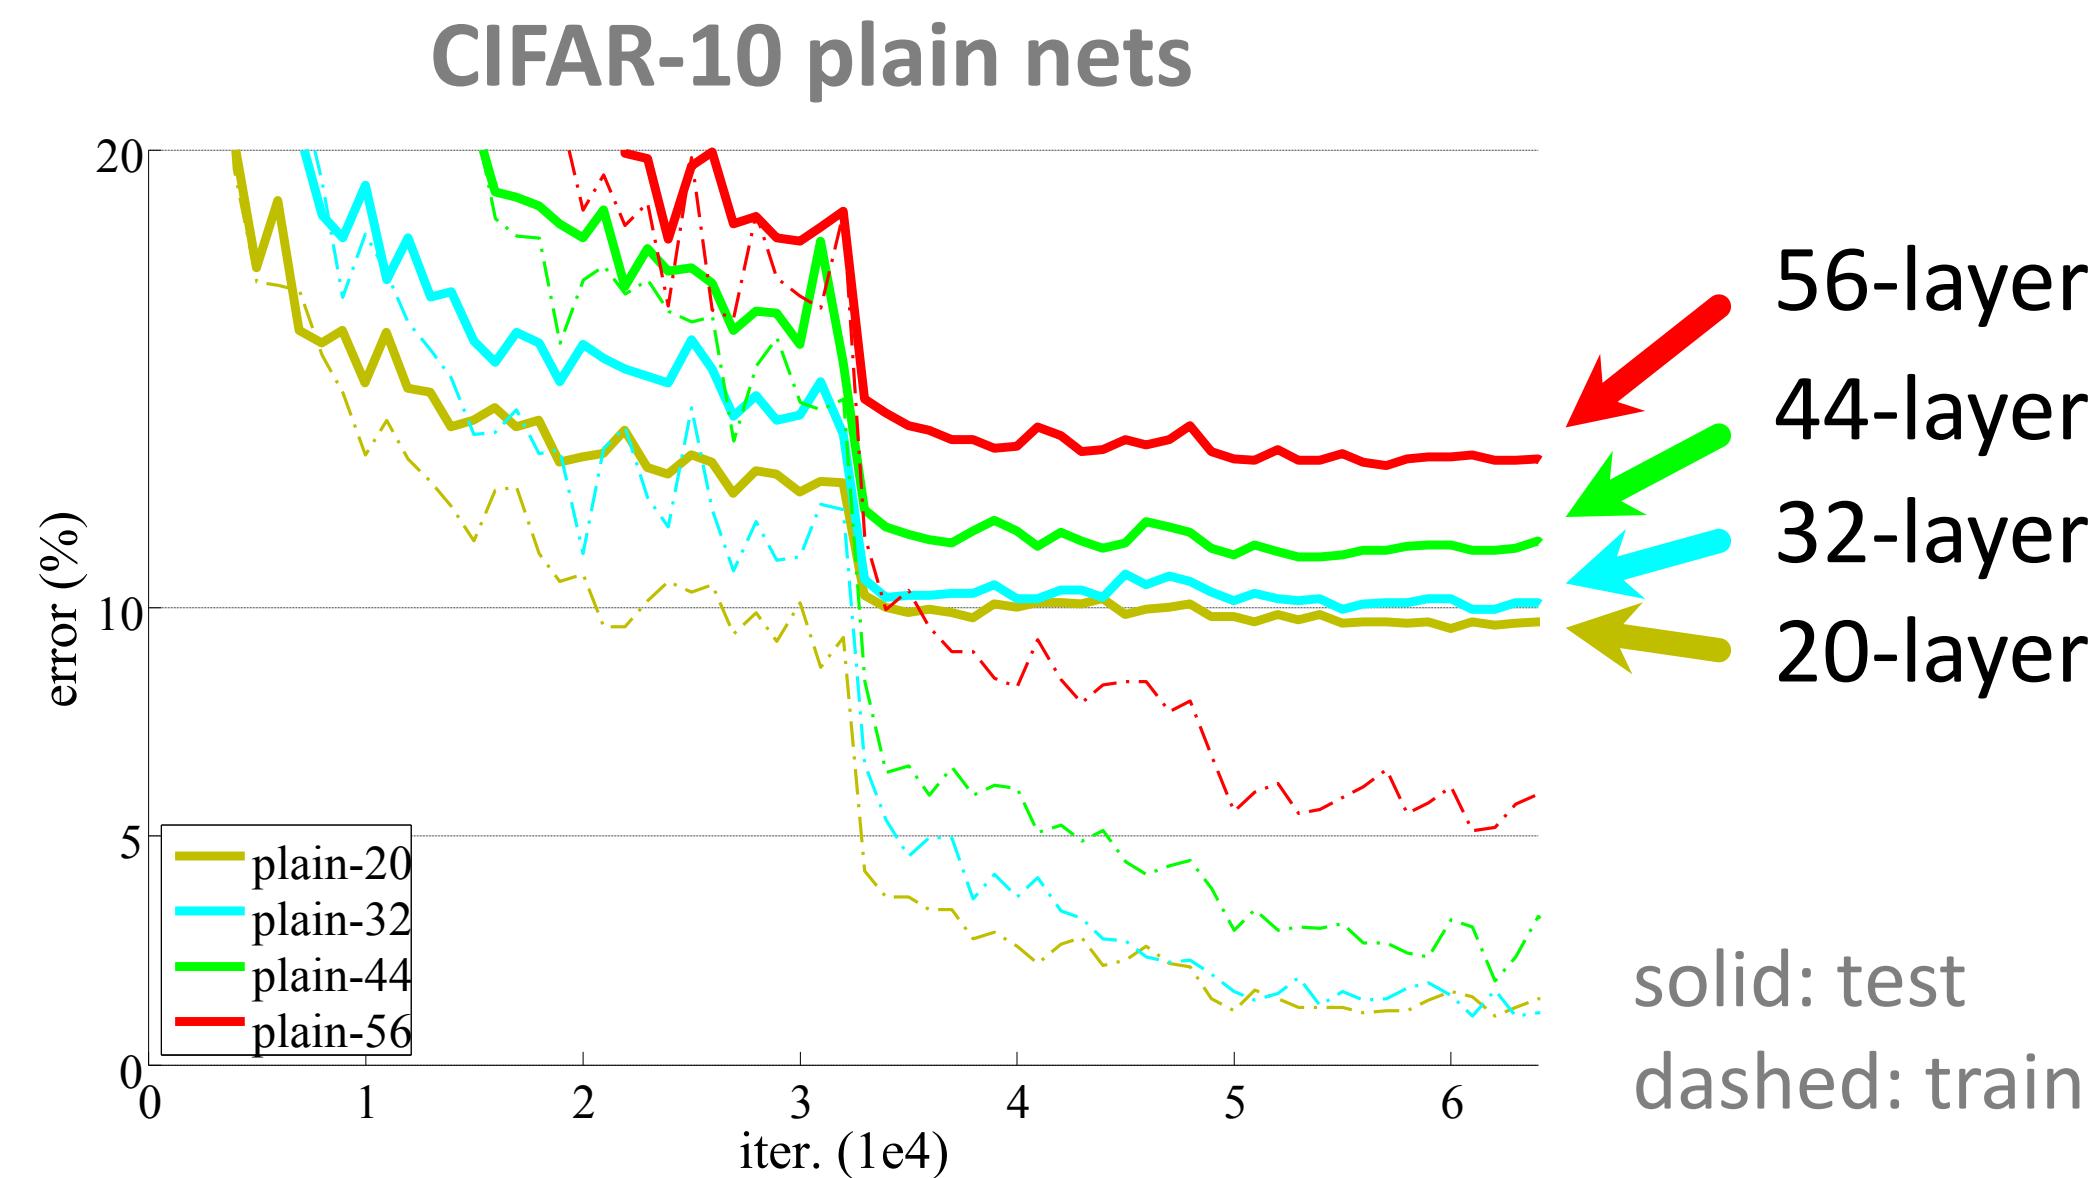


- Residual net



http://image-net.org/challenges/talks/ilsvrc2015_deep_residual_learning_kaiminghe.pdf

CIFAR-10 experiments

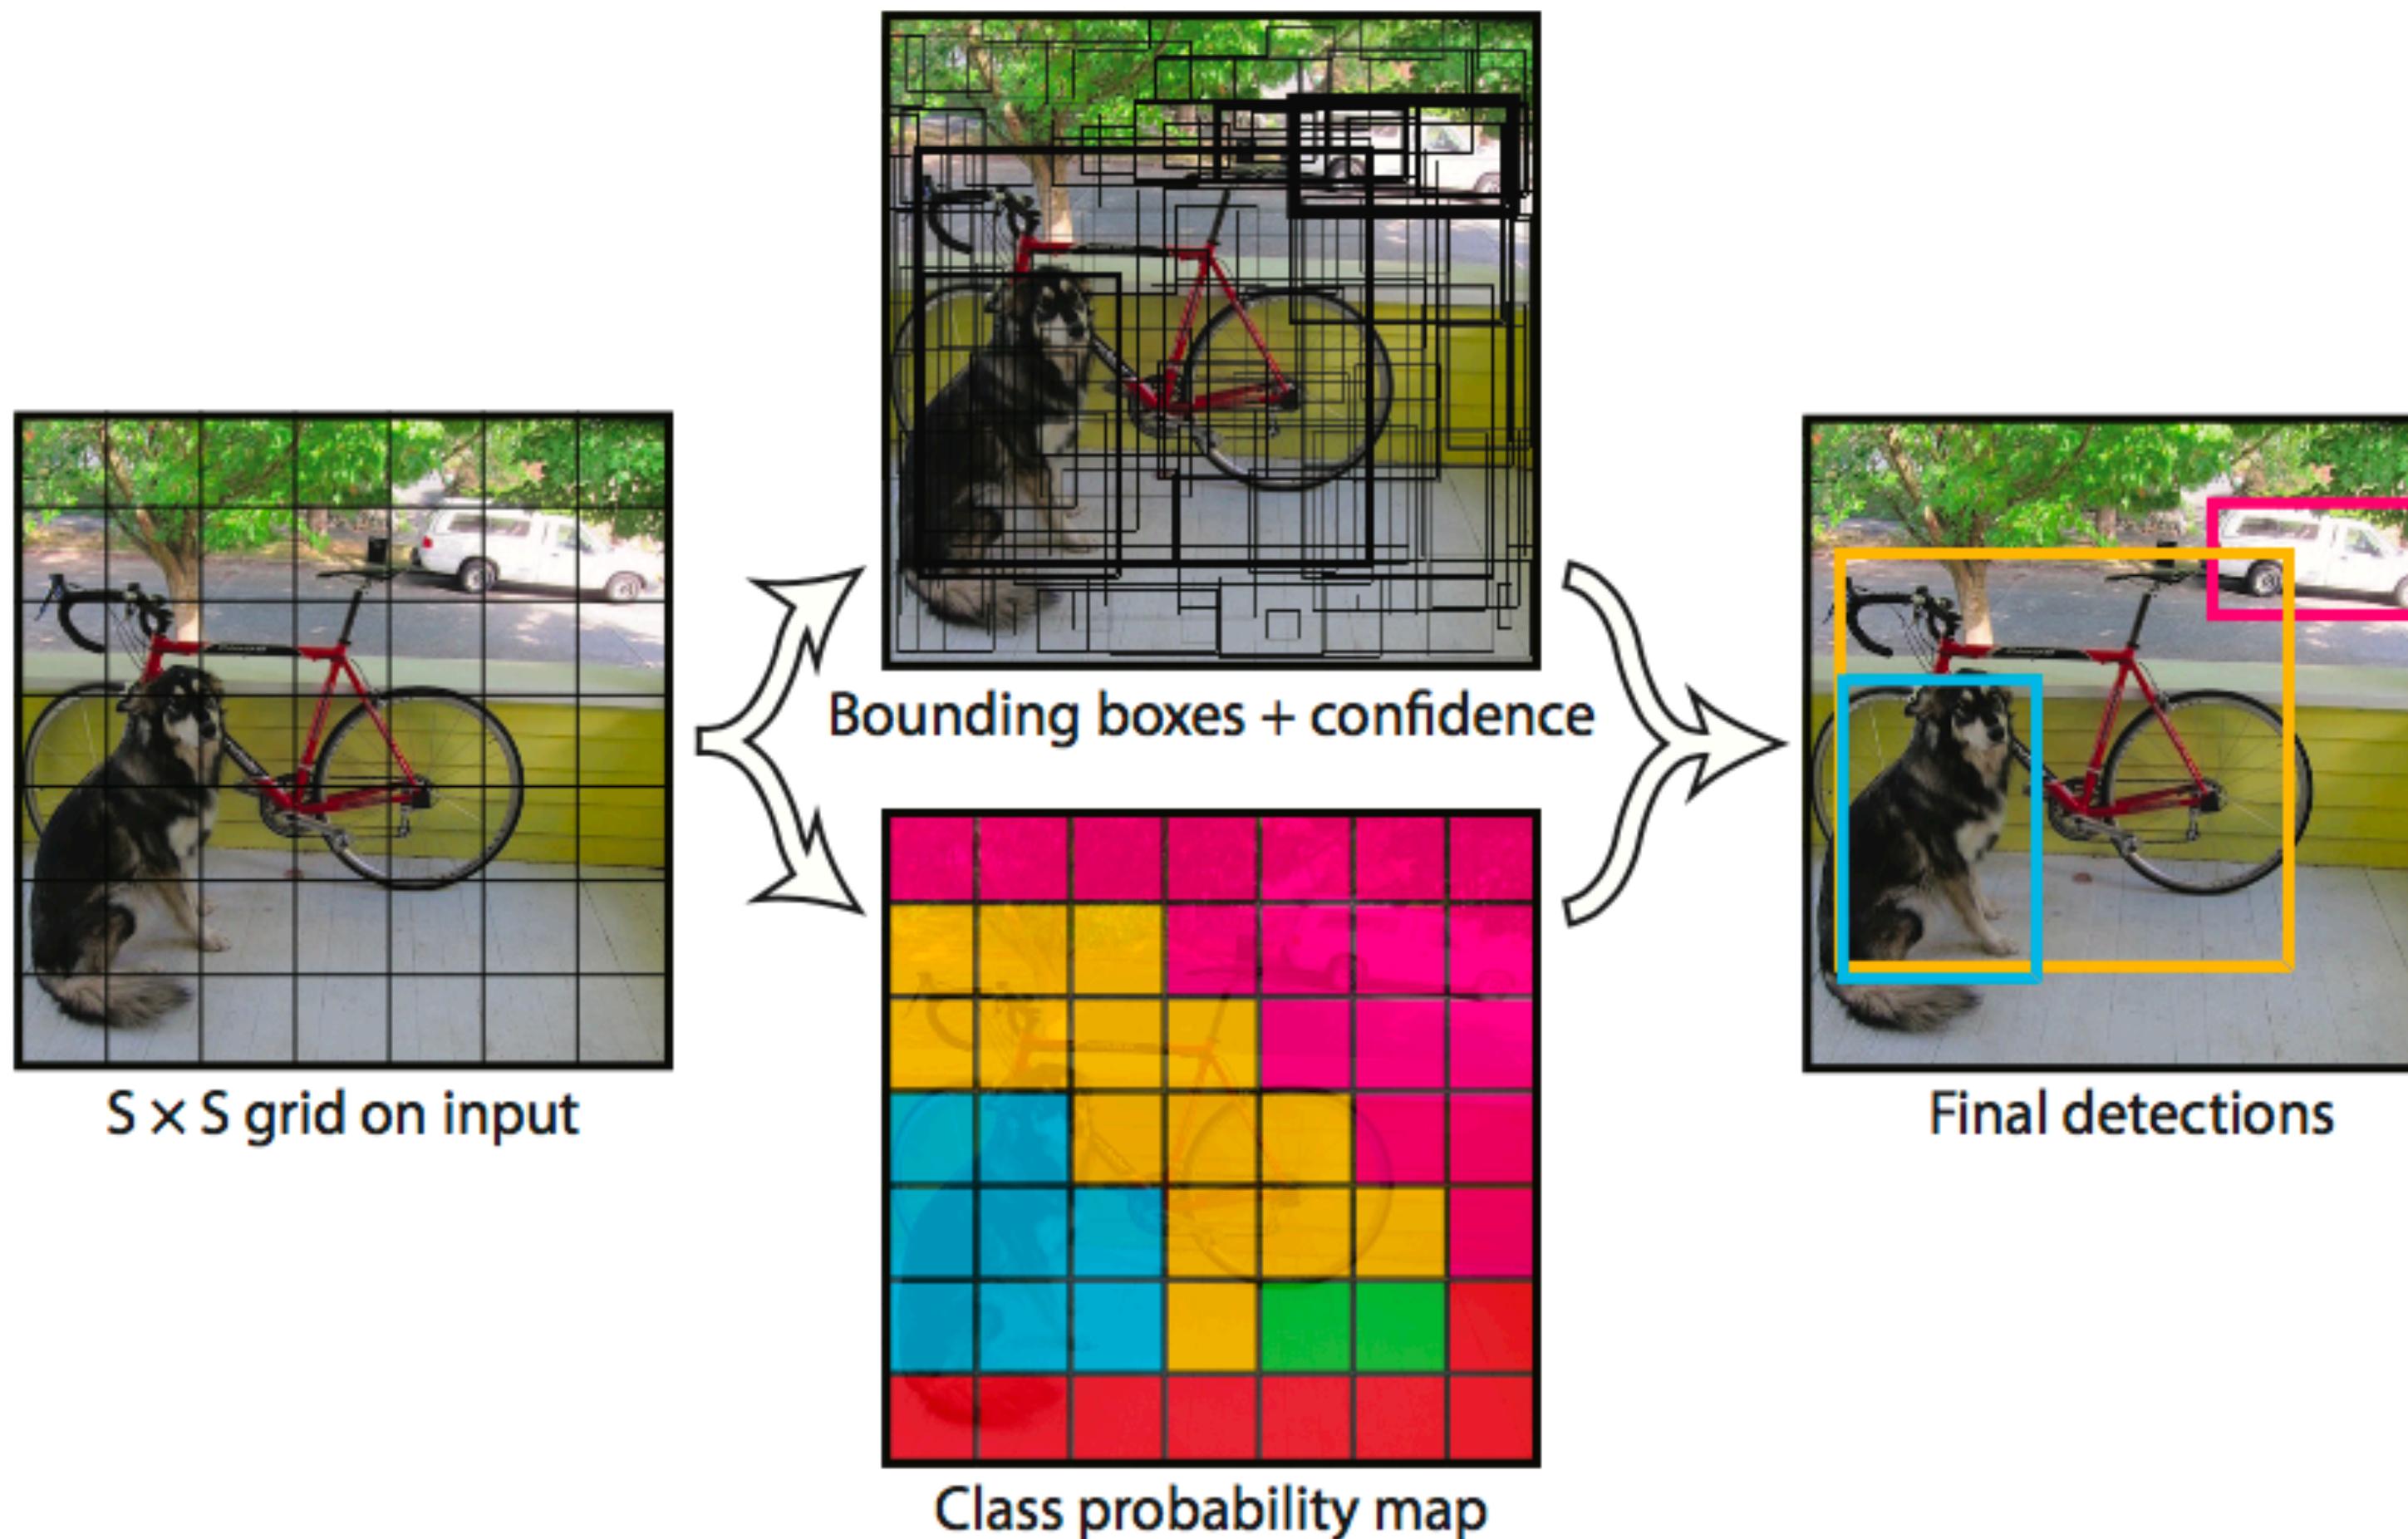


- Deep ResNets can be trained without difficulties
- Deeper ResNets have **lower training error**, and also lower test error

Further Reading

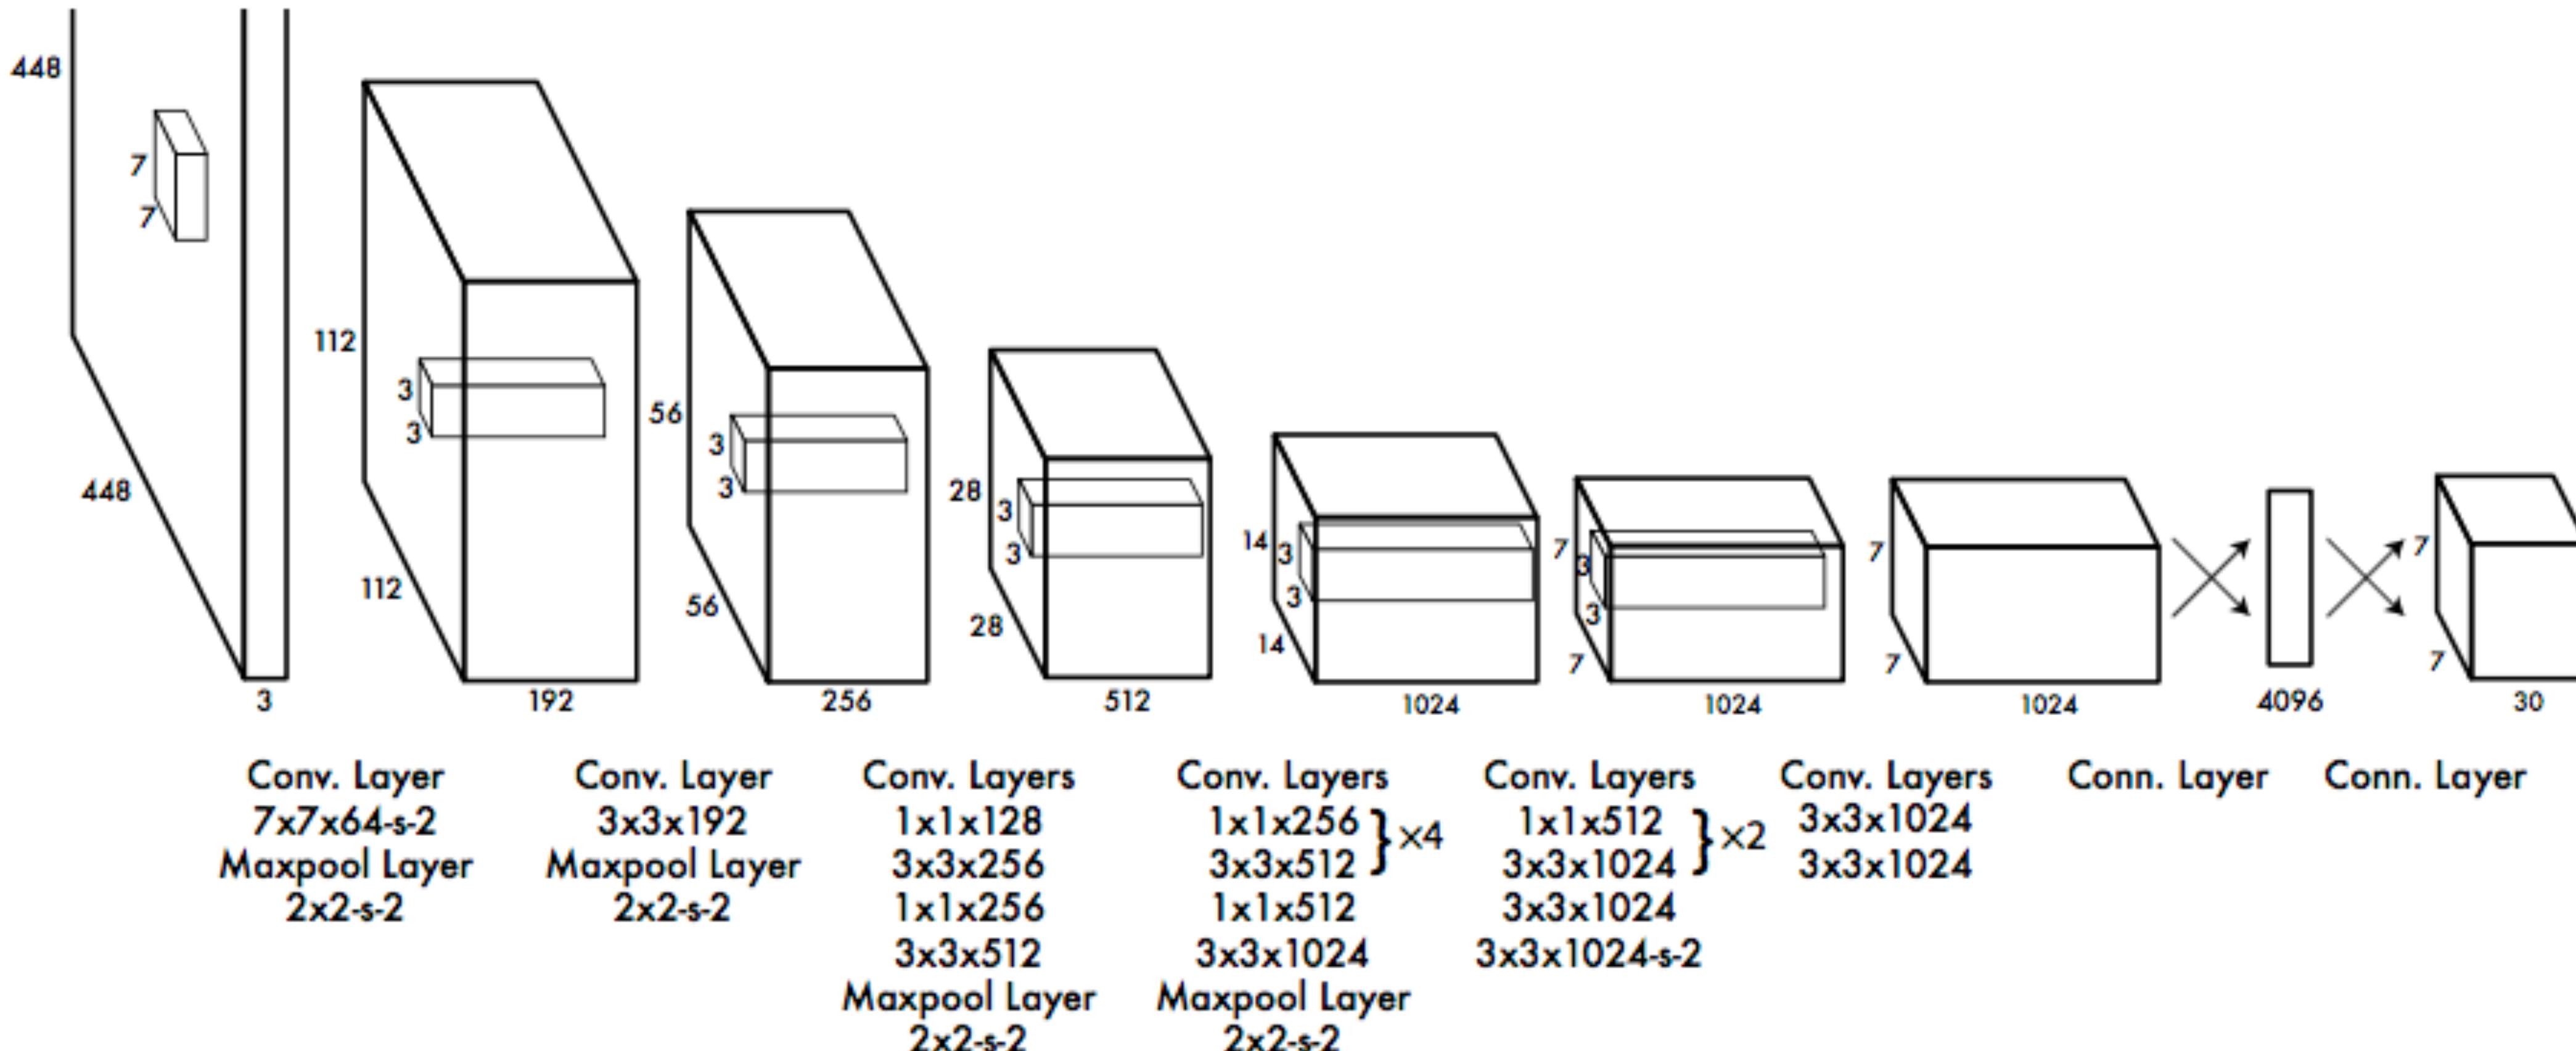
- 一年半前 (2017-12) の資料ですが素晴らしいスライドです
- 置み込みニューラルネットワークの研究動向 by DeNA 内田さん
- <https://www.slideshare.net/ren4yu/ss-84282514>

YOLO (You Only Look Once)

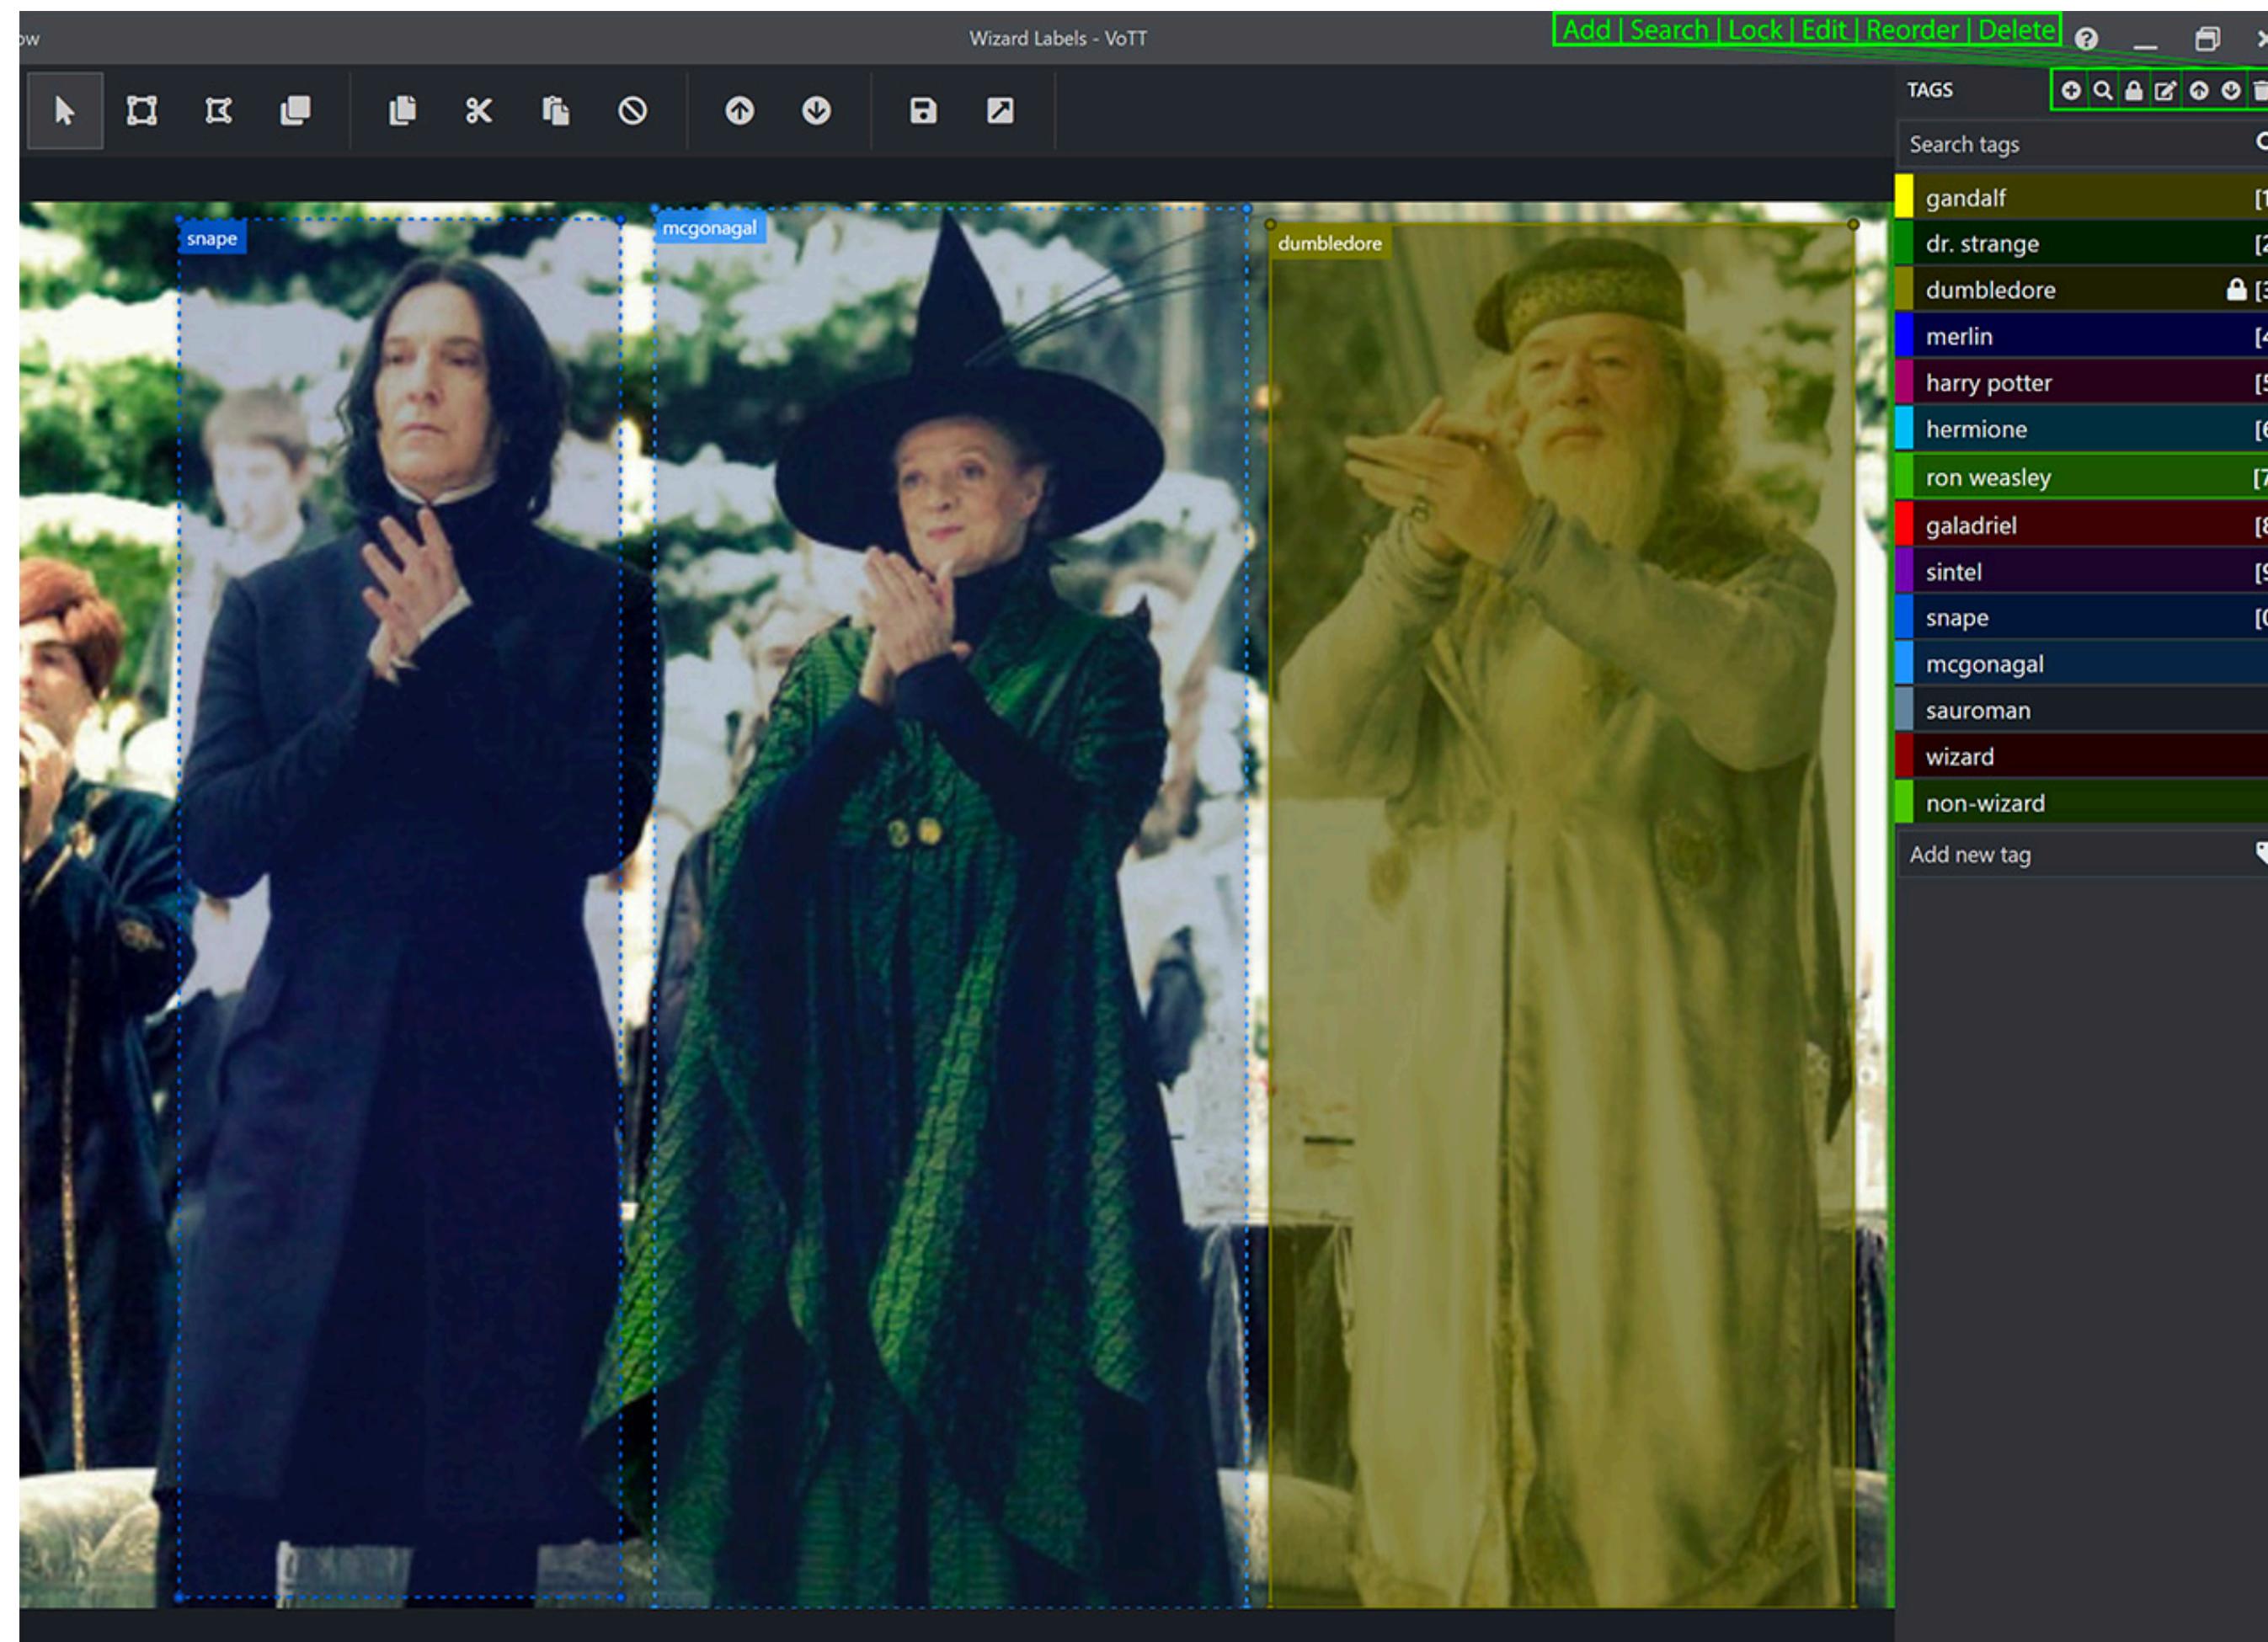


<https://www.youtube.com/watch?v=MPU2HistivI>

YOLOのアーキテクチャ



Labeling / Annotation



<https://github.com/Microsoft/VoTT>

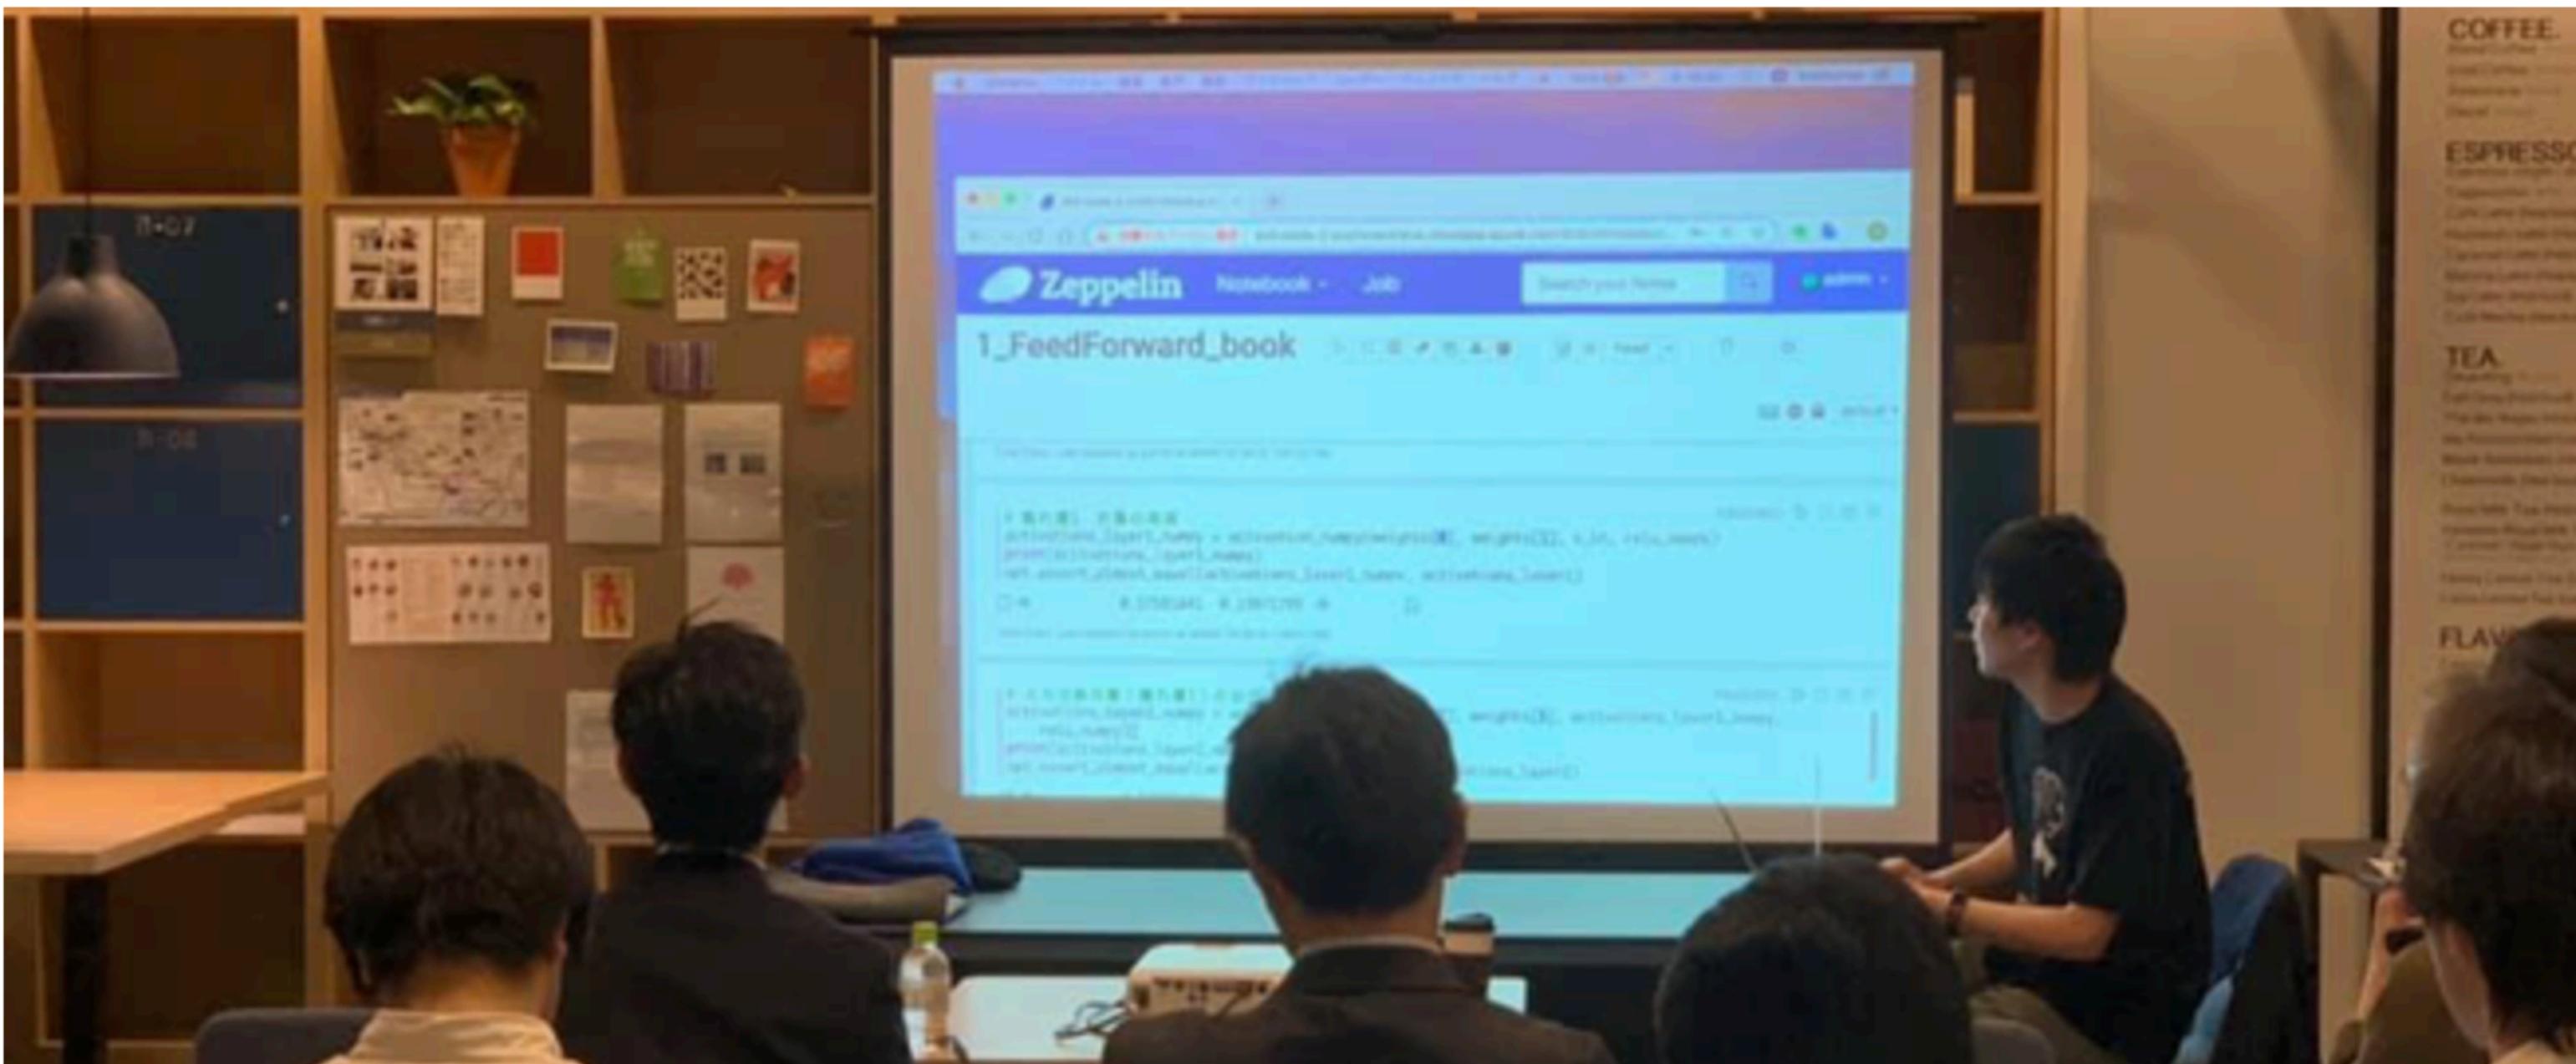
まとめ

- 機械学習ではデータからルールを確率的に学ぶ
- 画像に対して全結合の重みを最適化するのは非効率
- CNNで画像の特徴量抽出方法まで自動で行うアプローチ
- ResNetなど、さらに効率的に行うモデルが提案されている
- 何がうまくいくかは正直データによるので、経験的にうまくいきそうなアーキテクチャの組み合わせだけ用意してあとは勝手に選んで欲しい
→ Auto ML

10月
23

実践者向けディープラーニング勉強会 第七回

主催 : Miki_Shingo



- <https://dl4-practitioners.connpass.com/event/149810/>
- 毎月開催、来月11/18はGoogleの人がAutoMLの話をしてくれる予定
- 無料